Microsoft 70-463 Exam Questions & Answers

ACTUAL EXAMS

Microsoft 70-463 Exam Questions & Answers

Exam Name: Implementing a Data Warehouse with Microsoft SQL Server 2012

Exam A

QUESTION 1
You are implementing a SQL Server Integration Services (SSIS) package that loads data hosted in a SQL Azure database into a data warehouse. The source system contains redundant or inconsistent data. When the package finds invalid data, the row containing the invalid data must be omitted but it must also be written to a text file for further analysis. You need to establish the best technique to log these invalid rows while keeping the amount of development effort to a minimum. What should you do?

A. Add an OnError event handler to the SSIS project.
B. Use an msi file to deploy the package on the server.
C. Open a command prompt and run the gacutil command.
D. Open a command prompt and run the dtutil /copy command.
E. Open a command prompt and run the dtexec /rep /conn command.
F. Open a command prompt and run the dtexec /dumperror /conn command.
G. Run the package by using the dtexecui.exe utility and the SQL Log provider.
H. Create a reusable custom logging component and use it in the SSIS project.
I. Configure the SSIS solution to use the Project Deployment Model.
J. Configure the output of a component in the package data flow to use a data tap.
K. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL Server.

Correct Answer: J
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://technet.microsoft.com/en-us/library/hh230989.aspx
http://www.rafael-salas.com/2012/01/ssis-2012-quick-peek-to-data-taps.html
http://msdn.microsoft.com/en-us/library/ms162820.aspx
http://msdn.microsoft.com/en-us/library/hh231187.aspx
http://technet.microsoft.com/en-us/library/ms140223.aspx
http://msdn.microsoft.com/en-us/library/jj655339.aspx

QUESTION 2
You are troubleshooting an existing SQL Server Integration Services (SSIS) package. On several occasions, the package execution does not finish and no data seems to have been transferred. You need to ensure that package logging occurs. Your solution must minimize deployment and development efforts. What should you do?

A. Add an OnError event handler to the SSIS project.
B. Use an msi file to deploy the package on the server.
C. Open a command prompt and run the gacutil command.
D. Open a command prompt and run the dtutil /copy command.
E. Open a command prompt and run the dtexec /rep /conn command.
F. Open a command prompt and run the dtexec /dumperror /conn command.
G. Run the package by using the dtexecui.exe utility and the SQL Log provider.
H. Create a reusable custom logging component and use it in the SSIS project.
I. Configure the SSIS solution to use the Project Deployment Model.
J. Configure the output of a component in the package data flow to use a data tap.
K. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL

Server.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms141212.aspx
http://www.mssqltips.com/sqlservertip/2450/ssis-package-deployment-model-in-sql-server-2012-part-1-of-2/
http://www.mssqltips.com/sqlservertip/2450/ssis-package-deployment-model-in-sql-server-2012-part-2-of-2/

QUESTION 3
To ease the debugging of packages, you standardize the SQL Server Integration Services (SSIS) package logging methodology. The methodology has the following requirements:
▪ Centralized logging in SQL Server
▪ Simple deployment
▪ Availability of log information through reports or T-SQL
▪ Automatic purge of older log entries
▪ Configurable log details

You need to configure a logging methodology that meets the requirements while minimizing the amount of deployment and development effort. What should you do?

A. Use the gacutil command.
B. Use the dtutil /copy command.
C. Use the Integration Services Deployment Wizard.
D. Create an OnError event handler.
E. Create a reusable custom logging component.
F. Run the package by using the dtexec /rep /conn command.
G. Run the package by using the dtexec /dumperror /conn command.
H. Run the package by using the dtexecui.exe utility and the SQL Log provider.
I. Add a data tap on the output of a component in the package data flow.
J. Deploy the package by using an msi file.
K. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.

Correct Answer: H
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms140246.aspx
http://msdn.microsoft.com/en-us/library/hh231187.aspx

QUESTION 4

A SQL Server Integration Services (SSIS) package on a computer is failing periodically in production. The package was created less than one year ago and was deployed to the SSIS catalog. Sometimes the package is started on a SQL Agent schedule; sometimes the package is started manually by an SSIS developer by using the Object Explorer in SQL Server Management Studio. You need to identify the authenticated user responsible for starting the package each time it failed in the past. Where can you find this information?

A. the SQL Server Log
B. the SSISDB.[catalog].[executions] view
C. the SSISDB.[catalog].[event_messages] view
D. the SQL Agent Job History
E. the SQL Agent Error Log

Correct Answer: B
Section: (none)
Explanation

Explanation/Reference:
Looks Correct

QUESTION 5

You maintain a SQL Server Integration Services (SSIS) package. The package was developed by using SQL Server 2008 Business Intelligence Development Studio (BIDS). The package includes custom scripts that must be upgraded. You need to upgrade the package to SQL Server 2012. Which tool should you use?

A. SSIS Upgrade Wizard in SQL Server 2008 BIDS
B. SSIS Upgrade Wizard in SQL Server Data Tools
C. SQL Server DTExecUI utility (dtexecui.exe)
D. SQL Server dtexec utility (dtexec.exe)

Correct Answer: B
Section: (none)
Explanation

Explanation/Reference:
Look Correct

QUESTION 6
You are designing a data warehouse for a software distribution business that stores sales by software title. It stores sales targets by software category. Software titles are classified into subcategories and categories. Each software title is included in only a single software subcategory, and each subcategory is included in only a single category. The data warehouse will be a data source for an Analysis Services cube. The data warehouse contains two fact tables:
• factSales, used to record daily sales by software title
• factTarget, used to record the monthly sales targets by software category
Reports must be developed against the warehouse that reports sales by software title, category and subcategory, and sales targets. You need to design the software title dimension. The solution should use as few tables as possible while supporting all the requirements. What should you do?

A. Create three software tables, dimSoftware, dimSoftwareCategory. and dimSoftwareSubcategory and a fourth bridge table that joins software titles to their appropriate category and subcategory table records with foreign key constraints. Direct the cube developer to use key granularity attributes.
B. Create three software tables, dimSoftware, dimSoftwareCategory, and dimSoftwareSubcategory. Connect factSales to all three tables and connect factTarget to dimSoftwareCategory with foreign key constraints. Direct the cube developer to use key granularity attributes.
C. Create one table, dimSoftware, which contains Software Detail, Category, and Subcategory columns. Connect factSales to dimSoftware with a foreign key constraint. Direct the cube developer to use a non-key granularity attribute for factTarget.
D. Create two tables, dimSoftware and dimSoftwareCategory. Connect factSales to dimSoftware and factTarget to dimSoftwareCategory with foreign key constraints. Direct the cube developer to use

key granularity attributes.

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
Looks Correct

QUESTION 7
You are reviewing the design of a customer dimension table in an existing data warehouse hosted on SQL Azure. The current dimension design does not allow the retention of historical changes to customer attributes such as Postcode. You need to redesign the dimension to enable the full historical reporting of changes to multiple customer attributes including Postcode. What should you do?

A. Add StartDate and EndDate columns to the customer dimension.
B. Add an IsCurrent column to the customer dimension.
C. Enable Snapshot Isolation on the data warehouse.
D. Add CurrentValue and PreviousValue columns to the customer dimension.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
looks correct

QUESTION 8
You are designing an enterprise star schema that will consolidate data from three independent data marts. One of the data marts is hosted on SQL Azure. Most of the dimensions have the same structure and content. However, the geography dimension is slightly different in each data mart. You need to design a consolidated dimensional structure that will be easy to maintain while ensuring that all dimensional data from the three original solutions is represented. What should you do?

A. Create a junk dimension for the geography dimension.
B. Implement change data capture.
C. Create a conformed dimension for the geography dimension.
D. Create three geography dimensions.

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
checked against both

QUESTION 9
You work as a senior database administrator at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You are running a training exercise for Microsoft SQL Server 2012 junior administrators. You are discussing the use ofSlowly Changing Dimension Transformation Outputs. One of the output options causes Derived Column transformations to create columns for the expired row and the current row indicators.

Which option is the output that causes this?

A. Unchanged Output
B. Inferred Member Updates Output
C. Historical Attributes Inserts Output
D. Fixed Attribute Output
E. Changing Attributes Updates Output

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 10
You work as a senior database administrator at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You are hosting a training exercise for Microsoft SQL Server 2012 administrators. You are currently discussing Multicast transformation.

Which of the following is TRUE with regards to Multicast transformation? (Choose two.)

A. It prevents a package from creating logical copies of data.
B. It allows a package to create logical copies of data.
C. It distributes its input to one or more outputs.
D. It distributes its input to a single output.

Correct Answer: BC
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 11
You work as a database administrator at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You are in the process of creating a SQL Server Integration Services (SSIS) package. The package will be configured to make use of the GenerateAndPersistNewIndex Fuzzy Lookup transformation option.

Which of the following is the purpose of this option?

A. It allows for a current index to be reused.
B. It allows for a new index to be created, without saving it.
C. It allows for a new index to be created and saved, without maintaining it.
D. It allows for a new index to be created, saved, and maintained.

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 12

You work as a senior database administrator at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You are running a training exercise for Microsoft SQL Server 2012 junior administrators. You are discussing data quality projects.

Which of the following are benefits of data quality projects? (Choose all that apply.)

A. It enables you to perform data cleansing on your source data by using the knowledge in a DQS knowledge base.
B. It allows you to perform data matching on your source data by using the knowledge in a DQS knowledge base.
C. It allows you to perform data matching on your source data by using the matching policy in a knowledge base.
D. It enables you to perform data cleansing on your source data by using the matching policy in a knowledge base.

Correct Answer: AC
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 13
You work as a database developer at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You have been asked to revise a SQL Server Integration Services (SSIS) package. The package only includes Execute SQL tasks.

To make sure that, in the event of an Execute SQL task failing, changes made by all of the Execute SQL tasks is rolled back. To do this, you want to alter the TransactionOption property of the package.

Which of the following is the value that must be set for the TransactionOption property of the package?

A. NotSupported.
B. Supported.
C. NotRequired.
D. Required.

Correct Answer: D
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 14
You work as a senior database developer at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You are hosting a training exercise for ABC.com's junior database developers. You are currently discussing an EventHandlerClass property that is used to retrieve or configure a value that indicates whether the container participates in transactions.

Which of the following is the EventHandlerClass property that you are discussing?

A. The TransactionOption property.
B. The Executables property.
C. The LoggingOptions property.
D. The Variables property.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 15
You work as a senior database administrator at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You are running a training exercise for Microsoft SQL Server 2012 junior administrators. You are discussing the effects of setting IsolationLevel transaction property of a Data Flow task to Chaos.

Which of the following suitably describes the effect?

A. It prevents pending changes from more highly isolated transactions from being overwritten.
B. It prevents shared locks from being issued and exclusive locks from being honored.
C. It prevents other users from updating the data.
D. It reduces blocking by storing a version of data that one application can read while another is modifying the same data.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 16
You work as a database administrator at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You have created a SQL Server Integration Services (SSIS) project that you are preparing to deploy using the Integration Services Deployment Wizard.

Which of the following is the third step in the Integration Services Deployment Wizard's process?

A. Review
B. Select Destination
C. Deploy/Results
D. Select Source

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 17
You are using a SQL Server Integration Services (SSIS) project that is stored in the SSIS catalog. An Environment has been defined in the SSIS catalog. You need to add the Environment to the project. Which stored procedure should you use?

A. The catalog.create_execution stored procedure.
B. The catalog.create_environment_variable stored procedure.
C. The catalog.create_environment_reference stored procedure.
D. The catalog.add_data_tap stored procedure.

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 18
You work as a database administrator at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You are in the process of creating a SQL Server Integration Services (SSIS) package. The package will be configured upload a file to a table in a SQL Azure database.

You need to make sure that the code includes an event handler that makes use of an executable to raise the event when its verification is completed.

Which of the following is the event handler that you should use?

A. The OnPostExecute event handler.
B. The OnTaskFailed event handler.
C. The OnPreExecute event handler.
D. The OnPostValidate event handler.

Correct Answer: D
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 19
You work as a senior database administrator at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You are running a training exercise for Microsoft SQL Server 2012 junior administrators. You are discussing replication type that is used when subscribers need to receive data, make changes offline, and later synchronize changes with the Publisher and other Subscribers.

Which of the following is the replication type being discussed?

A. Snapshot replication.
B. Transactional replication.
C. Merge replication.
D. Peer-to-peer replication

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 20
You work as a senior database developer at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You are running a training exercise for Microsoft SQL Server 2012 junior developers. You are discussing the SorABCeyPosition property for a data flow.

Which of the following represents the data type for this property?

A. Integer
B. IDTSExternalMetadataColumn100
C. String
D. Boolean

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 21
You work as a database administrator at ABC.com. The ABC.com network consists of a single domain named ABC.com. ABC.com makes use of Microsoft SQL Server 2012 in their environment.

You have received reports regarding a current existing SQL Server Integration Services (SSIS) package not completing its execution, and not transferring data. You have decided to make use of
the Project Deployment Model.

Which of the following is TRUE with regards to using the Project Deployment Model?

A. CLR integration is required on the database engine.
B. CLR integration is not required on the database engine.
C. Events that are created by the package are captured automatically and saved to the catalog during execution.
D. Packages are executed using the dtExec and DTExecUI execution utilities.

Correct Answer: AC
Section: (none)
Explanation

Explanation/Reference:
Explanation:

QUESTION 22

You install a SQL Server 2012 database engine instance on a production server. A month later, you install SQL Server 2012 Integration Services (SSIS). You must develop an SSIS project and deploy it to the server by using the Project Deployment model. Operations Log records that are outside the configured retention period must be

cleaned automatically. You need to create the SSIS catalog on the production server. What should you do? (Each correct answer presents part of the solution. Choose all that apply.)

A. Enable XP Command Shell.
B. Enable CLR Integration.
C. Enable OLE Automation.
D. Start the SQL Server Browser service.
E. Enable Cross Database Ownership Chaining
F. Start the SQL Server Agent service
G. Enable Ad Hoc Remote Queries.

Correct Answer: BF
Section: (none)
Explanation

Explanation/Reference:

QUESTION 23

You are designing a SQL Server Integration Services (SSIS) solution. The solution will contain an SSIS project that includes several SSIS packages. Each SSIS package will define the same connection managers and variables. You have the following requirements:

• Ensure that the deployment model supports changing the content of connection strings by using parameters at execution time.
• Ensure that the deployment model automatically starts from calls to the catalog.start_execution stored procedure in the SSISDB database.
• Maximize performance at execution time.
• Minimize development effort.
You need to design a solution that meets the requirements.
What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

A. Use a project deployment model. Modify connection manager properties to use project parameters. Ensure that the SSISDB database is created.
B. Use a project deployment model. Configure connections in an XML configuration file referenced by an environment variable that corresponds to the SQL Server environment of each SSIS package.
C. Use a package deployment model. Use a SQL Server package configuration with a common filter. Change the contents of the SSIS Configurations table at runtime.
D. Use a package deployment model. Save each SSIS package to a file share that can be accessed from all environments.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:

QUESTION 24

You are designing an extract, transform, load (ETL) process for loading data from a SQL Server database into a large fact table in a data warehouse each day with the prior day's sales data. The ETL process for the fact table must meet the following requirements:
• Load new data in the shortest possible time.
• Remove data that is more than 36 months old.
• Ensure that data loads correctly.

• Minimize record locking.
• Minimize impact on the transaction log.
You need to design an ETL process that meets the requirements. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

A. Partition the destination fact table by date. Insert new data directly into the fact table and delete old data directly from the fact table.
B. Partition the destination fact table by date. Use partition switching and staging tables both to remove old data and to load new data.
C. Partition the destination fact table by customer. Use partition switching both to remove old data and to load new data into each partition.
D. Partition the destination fact table by date. Use partition switching and a staging table to remove old data. Insert new data directly into the fact table.

Correct Answer: B
Section: (none)
Explanation

Explanation/Reference:

QUESTION 25

You are designing a fact table in a SQL Server database. The fact table must meet the following requirements:
• Include a columnstore index.
• Allow users to choose up to 10 dimension tables and up to five facts at one time.
• Maximize performance of queries that aggregate measures by using any of the 10 dimensions.
• Support billions of rows.
• Use the most efficient design strategy.
You need to design the fact table to meet the requirements. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

A. Design a fact table with 5 dimensional key columns and 10 measure columns. Place the columnstore index on the dimensional key columns.
B. Design a fact table with 5 dimensional key columns and 10 measure columns. Place the columnstore index on the measure columns.
C. Design a fact table with 10 dimensional key columns and 5 measure columns. Place the columnstore index on the dimensional key columns and the measure columns.
D. Design a fact table with 10 dimensional key columns and 5 measure columns. Place the columnstore index on only the measure columns.

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:

QUESTION 26

You are designing a SQL Server Integration Services (SSIS) solution that will load multiple Online Transactional Processing (OLTP) data sources into a SQL Server data mart. You have the following requirements:

• Ensure that the process supports the creation of an exception report that details possible duplicate key values, null ratios within columns, and column-length distributions of values.
• Ensure that users can generate the exception report in an XML format.
• Use the minimum development effort.
You need to design the SSIS solution to meet the requirements. What should you do? (More than one answer

choice may achieve the goal. Select the BEST answer.)

A. Use a Data Profiling task. Use a Data Flow task to extract the XML output of the Data Profiling task into a SQL Server table. Query the table to view the exceptions.
B. Use Data Flow tasks to process the clean data.
C. Use a Data Profiling task. Read the exceptions in Data Profile Viewer.
D. Design a stored procedure that examines data for common dirty data patterns. Use an Execute SQL task.

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:

QUESTION 27

You are designing a SQL Server Integration Services (SSIS) solution. The solution will contain an SSIS project that includes several SSIS packages. Each SSIS package will define the same connection managers and variables. You have the following requirements:
• The deployment model must support changing the content of connection strings by using parameters at execution time.
• The deployment model must automatically log events to the SSISOB database.
• Maximize performance at execution time.
You need to design a solution that meets the requirements. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

A. Use a project deployment model. Modify connection manager properties to use project parameters.
B. Use a package deployment model. Save each SSIS package to a file share that can be accessed from all environments.
C. Use a package deployment model. Configure connections in an XML configuration file referenced by an environment variable that corresponds to the SQL Server environment of each SSIS package.
D. Use a project deployment model. Modify connection manager properties to use package parameters.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:

QUESTION 28

You are designing a partitioning strategy for a large fact table in a data warehouse. Tens of millions of new records are loaded into the data warehouse weekly, outside of business hours. Most queries are generated by reports and by cube processing. Data is frequently queried at the day level and occasionally at the month level. You need to partition the table to maximize the performance of queries. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

A. Partition the fact table by month, and compress each partition.
B. Partition the fact table by week.
C. Partition the fact table by year.
D. Partition the fact table by day, and compress each partition.

Correct Answer: D
Section: (none)

Explanation

Explanation/Reference:

QUESTION 29
You are using SQL Server Data Tools to develop a SQL Server Integration Services (SSIS) project. The first
package that you create in this project contains a package connection that accesses a flat file. Additional
packages in the project must also access this file. You need to define and reuse the flat file connection in all
project packages. What should you do?

A. Convert the package Connection Manager in the first package to a project Connection Manager.
B. Copy the package Connection Manager and paste it into the second package.
C. Convert the project to the Package Deployment model.
D. Set the ProtectionLevel property of the package Connection Manager to DontSaveSensitive to reuse the flat
   file connection.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
checked

QUESTION 30
You are installing the Data Quality Client on user desktops. You need to ensure that the prerequisite software
components are installed. Which components must be present to meet this goal? (Each correct answer
presents part of the solution. Choose all that apply.)

A. SQL Server Management Studio
B. Internet Explorer 6.0 SP1 or later
C. Microsoft Silverlight 5
D. .NET Framework 3.5 SP1
E. .NET Framework 4.0
F. Microsoft Silverlight 4
G. SQL Server Data Tools

Correct Answer: BE
Section: (none)
Explanation

Explanation/Reference:
checked

QUESTION 31
You manage a SQL Server Master Data Services (MDS) environment. A new application requires access to the
product data that is available in the MDS repository. You need to design a solution that gives the application
access to the product data with the least amount of development effort. What should you do?

A. Create a Subscription View in MDS.
B. Access the product entity tables in the MDS database directly.
C. Use SQL Server Integration Services (SSIS) to extract the data and put it in a staging database.
D. Use change data capture on the product entity tables.

Correct Answer: A

QUESTION 32
You manage a SQL Server Master Data Services (MDS) environment. A new application requires access to the product data that is available in the MDS repository. You need to design a solution that gives the application access to the product data with the least amount of development effort. What should you do?

A. Use sp_addlinkedserver to add a linked server to access the MDS database tables directly.
B. Create an OLE DB connection string that sets the Provider property to MDS.
C. Use transactional replication for data synchronization.
D. Create a Subscription View in MDS.

Correct Answer: D

QUESTION 33
You are creating a SQL Server Integration Services (SSIS) package that implements a Type 3 Slowly Changing Dimension (SCD). You need to add a task or component to the package that allows you to implement the SCD logic. What should you use?

A. a Script component
B. an SCD component
C. an Aggregate component
D. a Merge component

Correct Answer: D

QUESTION 34
You are developing a SQL Server Integration Services (SSIS) package that imports data into a data warehouse hosted on SQL Azure. The package uses a Foreach container to process text files found in a folder. The package must be deployed to a single server by using the Project Deployment model. Multiple SQL Server Agent jobs call the package. Each job is executed on a different schedule. Each job passes a different folder path to the package. You need to configure the package to accept the folder path from each job.
Which package configuration should you use?

A. Parent Package Variable
B. XML Configuration File
C. Environment Variable
D. .dtsConfig file
E. Registry Entry

Correct Answer: C

Explanation

Explanation/Reference:

QUESTION 35
You are developing a SQL Server Integration Services (SSIS) package. The package is stored as the myPackage.dtsx file in the root directory of the C drive of the SSIS server. You need to run the package from the command prompt.Which command should you use?

A. dtexec /sql myPackage
B. dtexec /d "C:\File System\myPackage.dtsx"
C. dtexec /f "C:\myPackage.dtsx"
D. dtexec /com "myPackage.dtsx"

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:

QUESTION 36
You are developing a project that contains multiple SQL Server Integration Services (SSIS) packages. The packages will be deployed to the SSIS catalog. One of the steps in each package accesses an FTP site to download data files. You create project parameters to store the username and password that are used to access the FTP site. You need to ensure that the username and password values are encrypted when they are deployed.
What should you do?

A. Convert the parameters to package parameters.
B. Set the Sensitive property of the parameters to True.
C. Set the ProtectionLevel property of the package to EncryptSensitiveWithPassword.
D. Convert the project to the Legacy Deployment model.

Correct Answer: B
Section: (none)
Explanation

Explanation/Reference:

QUESTION 37
You are the data steward for a Business Intelligence project. You must identify duplicate rows stored in a SQL Server table and output discoveries to a CSV file. A Data Quality Services (DQS) knowledge base has been created to support this project. You need to produce the CSV file with the least amount of development effort.
What should you do?

A. Create an Integration Services package and use a Fuzzy Lookup transform.
B. Create a Master Data Services (MDS) model.
C. Create a Data Quality Project.
D. Create an Integration Services package and use a DQS Cleansing transform.
E. Create an Integration Services package and use a Fuzzy Grouping transform.

Correct Answer: C
Section: (none)

Explanation

Explanation/Reference:

QUESTION 38
You are installing SQL Server Data Quality Services (DQS). You need to give users belonging to a specific Active Directory group access to the Data Quality Server. Which SQL Server application should you use?

A. Data Quality Client with administrative credentials
B. SQL Server Configuration Manager with local administrative credentials
C. SQL Server Data Tools with local administrative permissions
D. SQL Server Management Studio with administrative credentials

Correct Answer: D
Section: (none)
Explanation

Explanation/Reference:
not sure

QUESTION 39
You are installing SQL Server Data Quality Services (DQS). You need to give specific users access to the Data Quality Server. Which SQL Server application should you use?

A. SQL Server Configuration Manager
B. SQL Server Data Tools
C. SQL Server Management Studio
D. Data Quality Client

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:

QUESTION 40
You are designing an enterprise star schema that will consolidate data from three independent data marts. One of the data marts is hosted on SQL Azure. Most of the dimensions have the same structure and content. However, the geography dimension is slightly different in each data mart. You need to design a consolidated dimensional structure that will be easy to maintain while ensuring that all dimensional data from the three original solutions is represented. What should you do?

A. Create a conformed dimension for the geography dimension.
B. Implement change tracking.
C. Create a degenerate dimension for the geography dimension.
D. Create a Type 2 slowly changing dimension for the geography dimension.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:

QUESTION 41
You are reviewing the design of a student dimension table in an existing data warehouse hosted on SQL Azure.
The current dimension design does not allow the retention of historical changes to student attributes such as
ParentOccupation. You need to redesign the dimension to enable the full historical reporting of changes to
multiple student attributes including ParentOccupation. What should you do?

A. Add CurrentValue and PreviousValue columns to the student dimension.
B. Enable Snapshot Isolation on the data warehouse.
C. Add an IsCurrent column to the student dimension.
D. Add StartDate and EndDate columns to the student dimension.

Correct Answer: D
Section: (none)
Explanation

Explanation/Reference:

QUESTION 42
You are maintaining a Data Quality Services (DQS) environment. The production server failed and a new server
has been set up. The DQS databases are restored to a new server. All the appropriate permissions are
granted. DQS users are experiencing issues connecting to the new Data Quality Server. You need to enable
users to connect to the new server. Which Surface Area Configuration property should you enable?

A. AdHocRemoteQueriesEnabled
B. SoapEndpointsEnabled
C. ClrIntegrationEnabled
D. RemoteDacEnabled
E. OleAutomationEnabled
F. XpCmdShellEnabled

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
Checked

QUESTION 43

You are designing a data warehouse with two fact tables. The first table contains sales per month and the
second table contains orders per day. Referential integrity must be enforced declaratively. You need to design a
solution that can join a single time dimension to both fact tables. What should you do?

A. Create a time mapping table.
B. Change the level of granularity in both fact tables to be the same.
C. Merge the fact tables.
D. Create a view on the sales table.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:

Not Sure

QUESTION 44
You are developing a SQL Server Integration Services (SSIS) package to load data into a data warehouse.

You need to establish the correct order for loading each table to maximize parallel processing. Which order should you use?

A. 1. DimCurrency, DimScenario, DimAccount in parallel
    2. DimOrganization
    3. FactFinance
B. l. DimCurrency, DimOrganization in parallel
    2. DimScenario, DimAccount in parallel
    3. FactFinance
C. 1. DimCurrency, FactFinance in parallel
    2. DimOrganization, DimScenario, DimAccount in parallel
D. 1. FactFinance
    2. DimOrganization, DimScenario, DimAccount in parallel
    3. DimCurrency
E. 1. DimCurrency
    2. DimOrganization
    3. DimScenario, DimAccount in parallel
    4. FactFinance

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms139892.aspx
http://msdn.microsoft.com/en-us/library/ms141261.aspx

QUESTION 45
You are developing a SQL Server Integration Services (SSIS) package. You need to design a package to change a variable value during package execution by using the least amount of development effort. What should you use?

A. Expression task
B. Script task
C. Execute SQL task
D. Execute Process task
E. Term Extraction transformation

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
According to this reference, this answer looks correct.

Reference:
http://msdn.microsoft.com/en-us/library/hh213137.aspx

QUESTION 46
You are editing a SQL Server Integration Services (SSIS) package that contains three Execute SQL tasks and no other tasks. The three Execute SQL tasks modify products in staging tables in preparation for a data warehouse load. The package and all three Execute SQL product tasks have their TransactionOption property set to Supported. You need to ensure that if any of the three Execute SQL product tasks fail, all three tasks will roll back their changes. What should you do?

A. Change the TransactionOption property of the package to Required.
B. Change the TransactionOption property of all three Execute SQL product tasks to Required.
C. Move the three Execute SQL product tasks into a Foreach Loop container.
D. Move the three Execute SQL product tasks into a Sequence container.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms137690.aspx
http://msdn.microsoft.com/en-us/library/ms141144.aspx

QUESTION 47
You administer a Microsoft SQL Server 2012 server that has SQL Server Integration Services (SSIS) installed. You plan to deploy new SSIS packages to the server. The SSIS packages use the Project Deployment Model together with parameters and Integration Services environment variables. You need to configure the SQL Server environment to support these packages. What should you do?

A. Create SSIS configuration files for the packages.
B. Create an Integration Services catalog.
C. Install Data Quality Services.
D. Install Master Data services.

Correct Answer: B
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference:
http://msdn.microsoft.com/en-us/library/hh479588.aspx
http://msdn.microsoft.com/en-us/library/hh213290.aspx
http://msdn.microsoft.com/en-us/library/hh213373.aspx

QUESTION 48
You are developing a SQL Server Integration Services (SSIS) package that imports data into a data
warehouse. You add an Execute SQL task to the control flow. The task must execute a simple INSERT
statement. The task has the following requirements:
▪ The INSERT statement must use the value of a string package variable. The variable name is StringVar.
▪ The Execute SQL task must use an OLE DB Connection Manager.

In the Parameter Mapping tab of the Execute SQL task, StringVar has been added as the only parameter. You
must configure the SQLStatement property of the Execute SQL task. Which SQL statement should you use?

A. `INSERT INTO dbo.Table (variablevalue) VALUES ($Project::StringVar)`

B. `INSERT INTO dbo.Table (variablevalue) VALUES (@StringVar)`

C. `INSERT INTO dbo.Table (variablevalue) VALUES ($Package::StringVar)`

D. `INSERT INTO dbo.Table (variablevalue) VALUES (?)`

Correct Answer: D
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms141003.aspx
http://msdn.microsoft.com/en-us/library/ms140355.aspx
http://msdn.microsoft.com/en-us/library/cc280502.aspx

QUESTION 49
You are developing a SQL Server Integration Services (SSIS) package to load data into a SQL Server table on
Server A. The package includes a data flow and is executed on ServerB. The destination table has its own
identity column. The destination data load has the following requirements:
▪ The identity values from the source table must be used.
▪ Default constraints on the destination table must be ignored.
▪ Batch size must be 100,000 rows.

You need to add a destination and configure it to meet the requirements. Which destination should you use?

A. ADO NET Destination with Bulk Insert
B. SQL Server Destination
C. OLE DB Destination with Fast Load
D. ADO NET Destination without Bulk Insert
E. OLE DB Destination without Fast Load

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference: http://msdn.microsoft.com/en-us/library/ms141237.aspx

Reference: http://msdn.microsoft.com/en-us/library/ms139821.aspx
Reference: http://msdn.microsoft.com/en-us/library/ms141095.aspx

QUESTION 50
You are creating a SQL Server Master Data Services (MDS) model for a company. The source data for the company is stored in a single table that contains the manager-to-subordinate relationships. You need to create a hierarchy representing the organizational structure of the company. Which hierarchy type should you use?

A. Natural
B. Explicit
C. Parent
D. Recursive

Correct Answer: D
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference: http://technet.microsoft.com/en-us/library/ff487006.aspx
Reference: http://msdn.microsoft.com/en-us/library/ee633747.aspx
Reference: http://technet.microsoft.com/en-us/library/ee633759.aspx

QUESTION 51
You develop a SQL Server Integration Services (SSIS) package that imports SQL Azure data into a data warehouse every night. The SQL Azure data contains many misspellings and variations of abbreviations. To import the data, a developer used the Fuzzy Lookup transformation to choose the closest-matching string from a reference table of allowed values. The number of rows in the reference table is very large. If no acceptable match is found, the Fuzzy Lookup transformation passes a null value. The current setting for the Fuzzy Lookup similarity threshold is 0.50. Many values are incorrectly matched. You need to ensure that more accurate matches are made by the Fuzzy Lookup transformation without degrading performance. What should you do?

A. Change the Exhaustive property to True.
B. Change the similarity threshold to 0.55.
C. Change the similarity threshold to 0.40.
D. Increase the maximum number of matches per lookup.

Correct Answer: B
Section: (none)
Explanation

Explanation/Reference:
According to this reference, this answer looks correct.

Reference: http://msdn.microsoft.com/en-us/library/ms137786.aspx

QUESTION 52
You develop a SQL Server Integration Services (SSIS) package that imports SQL Azure data into a data warehouse every night. The SQL Azure data contains many misspellings and variations of abbreviations. To import the data, a developer used the Fuzzy Lookup transformation to choose the closest-matching string from a reference table of allowed values. The number of rows in the reference table is very large. If no acceptable match is found, the Fuzzy Lookup transformation passes a null value. The current setting for the Fuzzy Lookup similarity threshold is 0.50. Many values are incorrectly matched. You need to ensure that more accurate matches are made by the Fuzzy Lookup transformation without degrading performance. What should you do?

A. Decrease the maximum number of matches per lookup.

B. Change the similarity threshold to 0.55.
C. Change the Exhaustive property to True.
D. Increase the maximum number of matches per lookup.

Correct Answer: B
Section: (none)
Explanation

Explanation/Reference:
According to this reference, this answer looks correct.

Reference: http://msdn.microsoft.com/en-us/library/ms137786.aspx

QUESTION 53
You are designing a SQL Server Integration Services (SSIS) package that uploads a file to a table named Orders in a SQL Azure database. The company's auditing policies have the following requirements:
▪ An entry must be written to a dedicated SQL Server log table named OrderLog.
▪ The entry must be written as soon as the file upload task completes.

You need to meet the company's policy requirements. Which event handler should you use?

A. OnProgress
B. OnInformation
C. OnPostExecute
D. OnComplete

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
According to this reference, this answer looks correct.

Reference: http://msdn.microsoft.com/en-us/library/ms140223.aspx

QUESTION 54
You are developing a SQL Server Integration Services (SSIS) package that imports data from a relational database to a data warehouse. You are importing data from a relational table named Projects. The table has change data capture enabled on all columns. You need to process only the most recent values from rows that have been inserted or updated since the previous execution of the package. Which query should you use as the data source?

A. `SELECT * FROM cdc.fn_cdc_get_all_changes_Projects (@from_lsn, @to_lsn, N' all update old')`
B. `SELECT * FROM cdc.fn_cdc_get_all_changes_Projects (@from_lsn, @to_lsn, N' all')`
C. `SELECT * FROM cdc.fn_cdc_get_net_changes_Projects (@from_lsn, @to_lsn)`
D. `SELECT * FROM cdc.Projects_CT WHERE @from_lsn >= _$start_lsn AND @to_lsn < _$start_lsn`

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference: http://msdn.microsoft.com/en-us/library/bb522511.aspx

Reference: http://msdn.microsoft.com/en-us/library/bb510627.aspx
Reference: http://msdn.microsoft.com/en-us/library/cc645937.aspx

QUESTION 55
You are designing a SQL Server Integration Services (SSIS) data flow to load sales transactions from a source system into a data warehouse hosted on SQL Azure. One of the columns in the data source is named ProductCode. Some of the data to be loaded will reference products that need special processing logic in the data flow. You need to enable separate processing streams for a subset of rows based on the source product code. Which data flow transformation should you use?

A. Source Assistant
B. Audit
C. Destination Assistant
D. Script Component

Correct Answer: D
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference: http://msdn.microsoft.com/en-us/library/ms137640.aspx
Reference: http://msdn.microsoft.com/en-us/library/ms141150.aspx
Reference: http://msdn.microsoft.com/en-us/library/ff929138.aspx
Reference: http://msdn.microsoft.com/en-us/library/ff929116.aspx

QUESTION 56
You are designing a SQL Server Integration Services (SSIS) package that uses the Fuzzy Lookup transformation. The reference data to be used in the transformation does not change. You need to reuse the Fuzzy Lookup match index to increase performance and reduce maintenance. What should you do?

A. Select the GenerateAndPersistNewIndex option in the Fuzzy Lookup Transformation Editor.
B. Select the GenerateNewIndex option in the Fuzzy Lookup Transformation Editor.
C. Select the DropExistingMatchIndex option in the Fuzzy Lookup Transformation Editor.
D. Execute the sp_FuzzyLookupTableMaintenanceUninstall stored procedure.
E. Execute the sp_FuzzyLookupTableMaintenanceInvoke stored procedure.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
According to this reference, this answer looks correct.

Reference: http://msdn.microsoft.com/en-us/library/ms137786.aspx

QUESTION 57
You are creating a SQL Server Integration Services (SSIS) package to retrieve product data from two different sources. One source is hosted in a SQL Azure database. Each source contains products for different distributors. Products for each distributor source must be combined for insertion into a single product table destination. You need to select the appropriate data flow transformation to meet this requirement. Which transformation types should you use? (Each correct answer presents a complete solution. Choose all that apply.)

A. Multicast
B. Merge Join

C. Term Extraction

D. Union All

E. Merge

Correct Answer: DE
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference: http://msdn.microsoft.com/en-us/library/ms141703.aspx
Reference: http://msdn.microsoft.com/en-us/library/ms141775.aspx
Reference: http://msdn.microsoft.com/en-us/library/ms141020.aspx
Reference: http://msdn.microsoft.com/en-us/library/ms141809.aspx
Reference: http://msdn.microsoft.com/en-us/library/ms137701.aspx

QUESTION 58
You are developing a SQL Server Integration Services (SSIS) package that imports data into a data warehouse hosted on SQL Azure. The package uses a Foreach container to process text files found in a folder. The package must be deployed to a single server by using the Project Deployment model. Multiple SQL Server Agent jobs call the package. Each job is executed on a different schedule. Each job passes a different folder path to the package. You need to configure the package to accept the folder path from each job. Which package configuration should you use?

A. .dtsConfig file

B. Registry Entry

C. Parent Package Variable

D. XML Configuration File

E. Environment Variable

Correct Answer: E
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference:
http://msdn.microsoft.com/en-us/library/ms345179.aspx
http://msdn.microsoft.com/en-us/library/ms141708.aspx
http://msdn.microsoft.com/en-us/library/hh213214.aspx
http://msdn.microsoft.com/en-us/library/hh213296.aspx
http://msdn.microsoft.com/en-us/library/hh213293.aspx

QUESTION 59
You are creating a SQL Server Master Data Services (MDS) model. This model is used to store a master list of products. An attribute must be added to the Product entity to define the sales manager responsible for each product. You need to create an attribute in the Product entity that prevents users from entering invalid sales manager values. Which type of attribute should you create?

A. Derived

B. Domain-based

C. User-defined

D. Parent

E. Explicit

F. Recursive

Correct Answer: B
Section: (none)
Explanation

Explanation/Reference:
-- BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/bb190163.aspx
http://msdn.microsoft.com/en-us/library/ee633737.aspx
http://msdn.microsoft.com/en-us/library/ee633759.aspx
http://msdn.microsoft.com/en-us/library/ee633745.aspx
http://msdn.microsoft.com/en-us/library/ee633724.aspx
http://msdn.microsoft.com/en-us/library/ee633733.aspx
http://msdn.microsoft.com/en-us/library/ff487058.aspx

QUESTION 60
You are developing a SQL Server Integration Services (SSIS) package that imports data into a data warehouse hosted on SQL Azure. The package uses a Foreach container to process text files found in a folder. The package must be deployed to a single server by using the Project Deployment model. Multiple SQL Server Agent jobs call the package. Each job is executed on a different schedule. Each job passes a different folder path to the package. You need to configure the package to accept the folder path from each job. Which package configuration should you use?

A. .dtsConfig file
B. Registry Entry
C. Parent Package Variable
D. XML Configuration File
E. Environment Variable

Correct Answer: E
Section: (none)
Explanation

Explanation/Reference:
-- BURGOS - YES
--\BURGOS

According to these references, this answer looks correct.

Reference:
http://msdn.microsoft.com/en-us/library/ms345179.aspx
http://msdn.microsoft.com/en-us/library/ms141708.aspx
http://msdn.microsoft.com/en-us/library/hh213214.aspx
http://msdn.microsoft.com/en-us/library/hh213296.aspx
http://msdn.microsoft.com/en-us/library/hh213293.aspx

QUESTION 61
You are adding a new capability to several dozen SQL Server Integration Services (SSIS) packages. The new capability is not available as an SSIS task. Each package must be extended with the same new capability. You need to add the new capability to all the packages without copying the code between packages. What should you do?

A. Use the Expression task.
B. Use the Script component.
C. Use the Script task.
D. Develop a custom task.
E. Develop a custom component.

Correct Answer: D
Section: (none)
Explanation

Explanation/Reference:
-- BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms135965.aspx
http://msdn.microsoft.com/en-us/library/ms345161.aspx

QUESTION 62
You are writing a SQL Server Integration Services (SSIS) package that transfers data from a legacy system.
Data integrity in the legacy system is very poor. Invalid rows are discarded by the package but must be logged
to a CSV file for auditing purposes. You need to establish the best technique to log these invalid rows while
minimizing the amount of development effort. What should you do?

A. Add a data tap on the output of a component in the package data flow.
B. Deploy the package by using an msi file.
C. Run the package by using the dtexecui.exe utility and the SQL Log provider.
D. Uses the dtutil /copy command.
E. Create an OnError event handler.
F. Use the Integration Services Deployment Wizard.
G. Use the gacutil command.
H. Create a reusable custom logging component.
I. Run the package by using the dtexec /rep /conn command.
J. Run the package by using the dtexec /dumperror /conn command.
K. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the
   configuration.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference:
http://www.rafael-salas.com/2012/01/ssis-2012-quick-peek-to-data-taps.html
http://msdn.microsoft.com/en-us/library/hh230989.aspx
http://msdn.microsoft.com/en-us/library/jj655339.aspx

QUESTION 63
You are deploying a project to the SQL Server Integration Services (SSIS) catalog. You need to ensure that the
project is deployed to the production environment with the least amount of administrative effort. What should
you do?

A. Create an On Error event handler.
B. Add a data tap on the output of a component in the package data flow.
C. Use the gacutil command.
D. Create a reusable custom logging component.
E. Run the package by using the dtexecui.exe utility and the SQL Log provider.
F. Run the package by using the dtexec /rep /conn command.
G. Use the Integration Services Deployment Wizard.
H. Run the package by using the dtexec /dumperror /conn command.
I. Use the dtutil /copy command.
J. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.
K. Deploy the package by using an msi file.

Correct Answer: G
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference:
http://msdn.microsoft.com/en-us/library/hh479588.aspx
http://msdn.microsoft.com/en-us/library/hh213290.aspx
http://msdn.microsoft.com/en-us/library/hh213373.aspx

QUESTION 64
You are the data steward for a Business Intelligence project. You must identify duplicate rows stored in a SQL Server table and output discoveries to a CSV file. A Data Quality Services (DQS) knowledge base has been created to support this project. You need to produce the CSV file with the least amount of development effort. What should you do?

A. Create an Integration Services package and use a Data Profiling transform.
B. Create a custom .NET application based on the Knowledgebase class.
C. Create a CLR stored procedure based on the Knowledgebase class.
D. Create a Master Data Services (MDS) business rule.
E. Create a data quality project.

Correct Answer: E
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/hh213052.aspx
http://msdn.microsoft.com/en-us/library/ff877917.aspx
http://msdn.microsoft.com/en-us/library/
microsoft.masterdataservices.services.datacontracts.knowledgebase.aspx
http://msdn.microsoft.com/en-us/library/bb895263.aspx

QUESTION 65
You are designing a data warehouse hosted on SQL Azure. The data warehouse currently includes the dimUser and dimDistrict dimension tables and the factSales fact table. The dimUser table contains records for

each user permitted to run reports against the warehouse, and the dimDistrict table contains information about sales districts. The system is accessed by users from certain districts, as well as by area supervisors and users from the corporate headquarters. You need to design a table structure to ensure that certain users can see sales data for only certain districts. Some users must be permitted to see sales data from multiple districts. What should you do?

A. Create a userDistrict table that contains primary key columns from the dimUser and dimDistrict tables.
B. Partition the factSales table on the district column.
C. Add a district column to the dimUser table.
D. For each district, create a view of the factSales table that includes a WHERE clause for the district.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
Looks Correct

QUESTION 66
You are editing a SQL Server Integration Services (SSIS) package that contains three Execute SQL tasks and no other tasks. The package and all three Execute SQL tasks have their TransactionOption property set to Supported. You need to ensure that if any of the Execute SQL tasks fail, all three tasks will roll back their changes. What should you do?

A. Change the TransactionOption property of the package to Required.
B. Change the TransactionOption property of all three Execute SQL tasks to Required.
C. Move the three Execute SQL tasks into a Sequence container.
D. Move the three Execute SQL tasks into a Foreach Loop container.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference: http://msdn.microsoft.com/en-us/library/ms137749.aspx
Reference: http://msdn.microsoft.com/en-us/library/microsoft.sqlserver.dts.runtime.dtstransactionoption.aspx

QUESTION 67
You are designing a SQL Server Integration Services (SSIS) package that uploads a file to a table named Orders in a SQL Azure database. The company's auditing policies have the following requirements:
▪ An entry must be written to a dedicated SQL Server log table named OrderLog.
▪ The entry must be written as soon as the file upload task completes.

You need to meet the company's policy requirements. Which event handler should you use?

A. OnWarning
B. OnVariableValueChanged
C. OnExecStatusChanged
D. OnPostExecute

Correct Answer: D
Section: (none)
Explanation

Explanation/Reference:
According to this reference, this answer looks correct.

Reference: http://msdn.microsoft.com/en-us/library/ms140223.aspx

QUESTION 68
You are designing a data warehouse with two fact tables. The first table contains sales per month and the second table contains orders per day. Referential integrity must be enforced declaratively. You need to design a solution that can join a single time dimension to both fact tables. What should you do?

A. Join the two fact tables.
B. Merge the fact tables.
C. Create a time dimension that can join to both fact tables at their respective granularity.
D. Create a surrogate key for the time dimension.

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms174537.aspx
http://technet.microsoft.com/en-us/library/ms174832.aspx
http://msdn.microsoft.com/en-us/library/ms174884.aspx
http://decipherinfosys.wordpress.com/2007/02/01/surrogate-keys-vs-natural-keys-for-primary-key/
http://www.agiledata.org/essays/keys.html
http://www.databasejournal.com/features/mssql/article.php/3922066/SQL-Server-Natural-Key-Verses-Surrogate-Key.htm
http://www.jamesserra.com/archive/2012/01/surrogate-keys/

QUESTION 69
You are completing the installation of the Data Quality Server component of SQL Server Data Quality Services (DQS). You need to complete the post-installation configuration. What should you do?

A. Run the dbimpexp.exe command.
B. Install the data providers that are used for data refresh.
C. Run the DQSInstaller.exe command.
D. Install the Analysis Services OLE DB Provider.

Correct Answer: C
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ff877917.aspx
http://msdn.microsoft.com/en-us/library/gg492277.aspx

QUESTION 70
You are using the Knowledge Discovery feature of the Data Quality Services (DQS) client application to modify an existing knowledge base. In the mapping configuration, two of the three columns are mapped to existing domains in the knowledge base. The third column, named Team Type, does not yet have a domain. You need to complete the mapping of the Team Type column. What should you do?

A. Add a column mapping for the Team Type column.
B. Add a domain for the Team Type column.
C. Map a composite domain to the source column.
D. Create a composite domain that includes the Team Type column.

Correct Answer: B
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ff877917.aspx
http://msdn.microsoft.com/en-us/library/hh213015.aspx
http://msdn.microsoft.com/en-us/library/gg524799.aspx

QUESTION 71
You are developing a SQL Server Integration Services (SSIS) project that copies a large amount of rows from a SQL Azure database. The project uses the Package Deployment Model. This project is deployed to SQL Server on a test server. You need to ensure that the project is deployed to the SSIS catalog on the production server. What should you do?

A. Add an OnError event handler to the SSIS project.
B. Use an msi file to deploy the package on the server.
C. Open a command prompt and run the gacutil command.
D. Open a command prompt and run the dtutil /copy command.
E. Open a command prompt and run the dtexec /rep /conn command.
F. Open a command prompt and run the dtexec /dumperror /conn command.
G. Run the package by using the dtexecui.exe utility and the SQL Log provider.
H. Create a reusable custom logging component and use it in the SSIS project.
I. Configure the SSIS solution to use the Project Deployment Model.
J. Configure the output of a component in the package data flow to use a data tap.
K. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL Server.

Correct Answer: I
Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/hh231102.aspx
http://msdn.microsoft.com/en-us/library/hh213290.aspx
http://msdn.microsoft.com/en-us/library/hh213373.aspx

Exam

QUESTION 1
DRAG AND DROP
You are designing a package control flow. The package moves sales order data from a SQL Azure transactional database to an on-premise reporting database. The package will run several times a day, while new sales orders are being added to the transactional database. The current design of the package control flow is shown in the answer area. (Click the Exhibit button.)
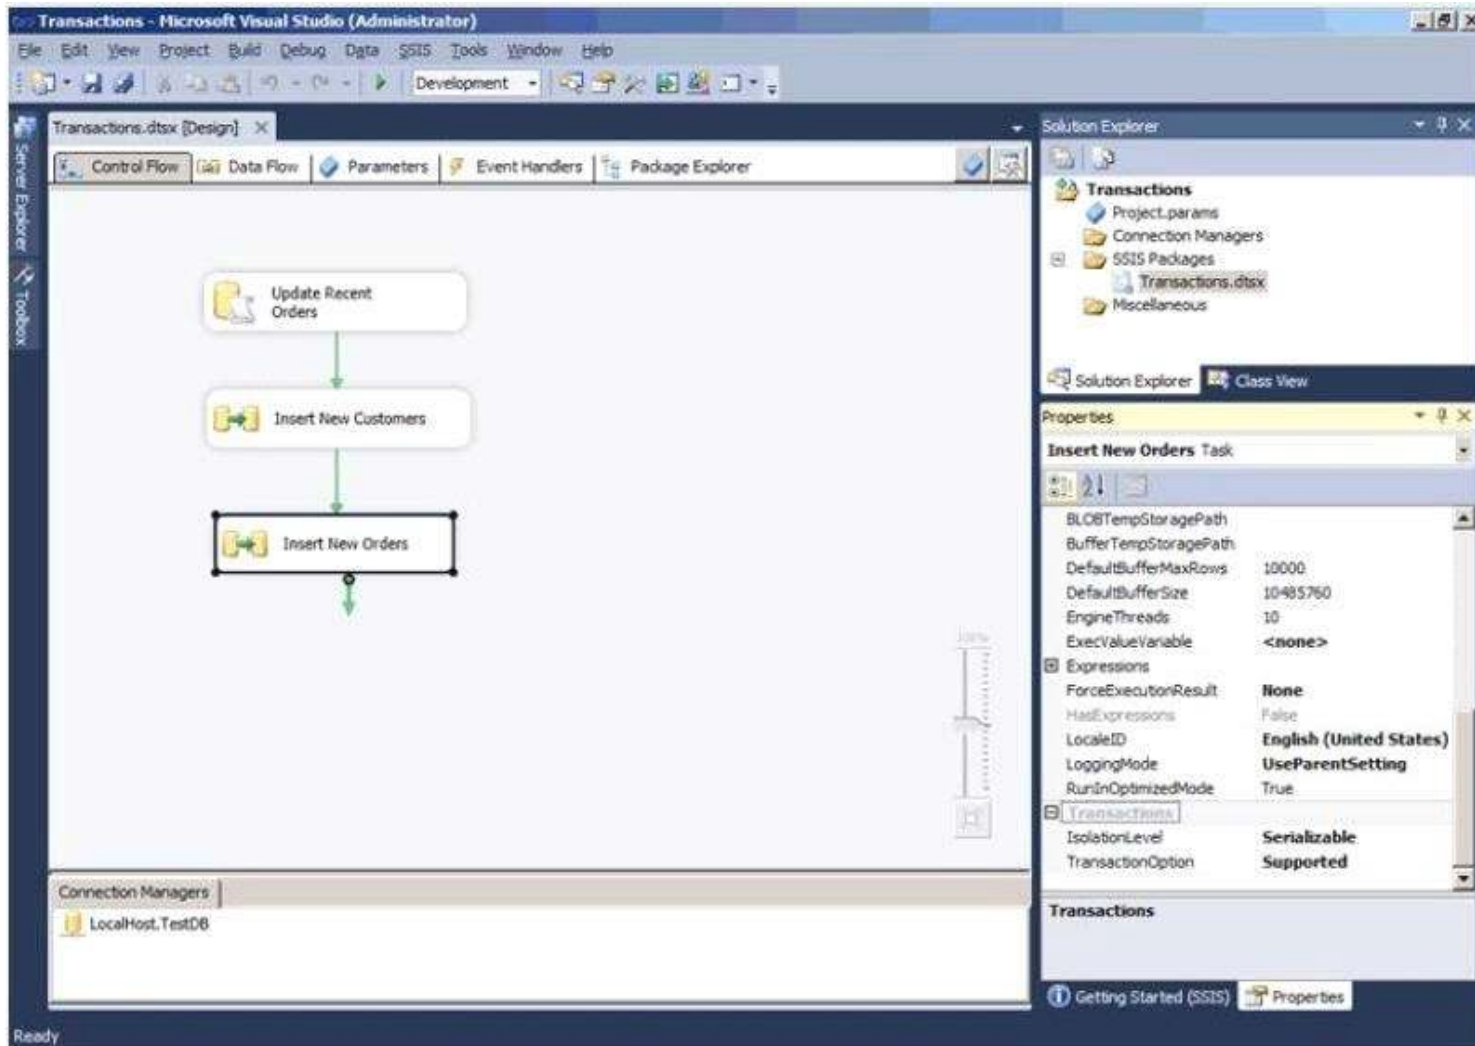
The Insert New Orders Data Flow task must meet the following requirements:
▪ Usage of the tempdb database should not be impacted.
▪ Concurrency should be maximized, while only reading committed transactions.
▪ If the task fails, only that task needs to be rolled back.

You need to configure the Insert New Orders Data Flow task to meet the requirements. How should you configure the transaction properties? (To answer, select the appropriate setting or settings in the answer area.)

Select and Place:

Chaos	
ReadCommitted	
ReadUncommitted	
RepeatableRead	
Required	
Serializable	
Snapshot	
Supported	
Unspecified	

Correct Answer:

Chaos

NotSupported

ReadCommitted

ReadUncommitted

RepeatableRead

Required

Serializable

Snapshot

Supported

Unspecified

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms137690.aspx
http://msdn.microsoft.com/en-us/library/ms137749.aspx
http://msdn.microsoft.com/en-us/library/microsoft.sqlserver.dts.runtime.dtscontainer.isolationlevel.aspx
http://msdn.microsoft.com/en-us/library/ms173763.aspx
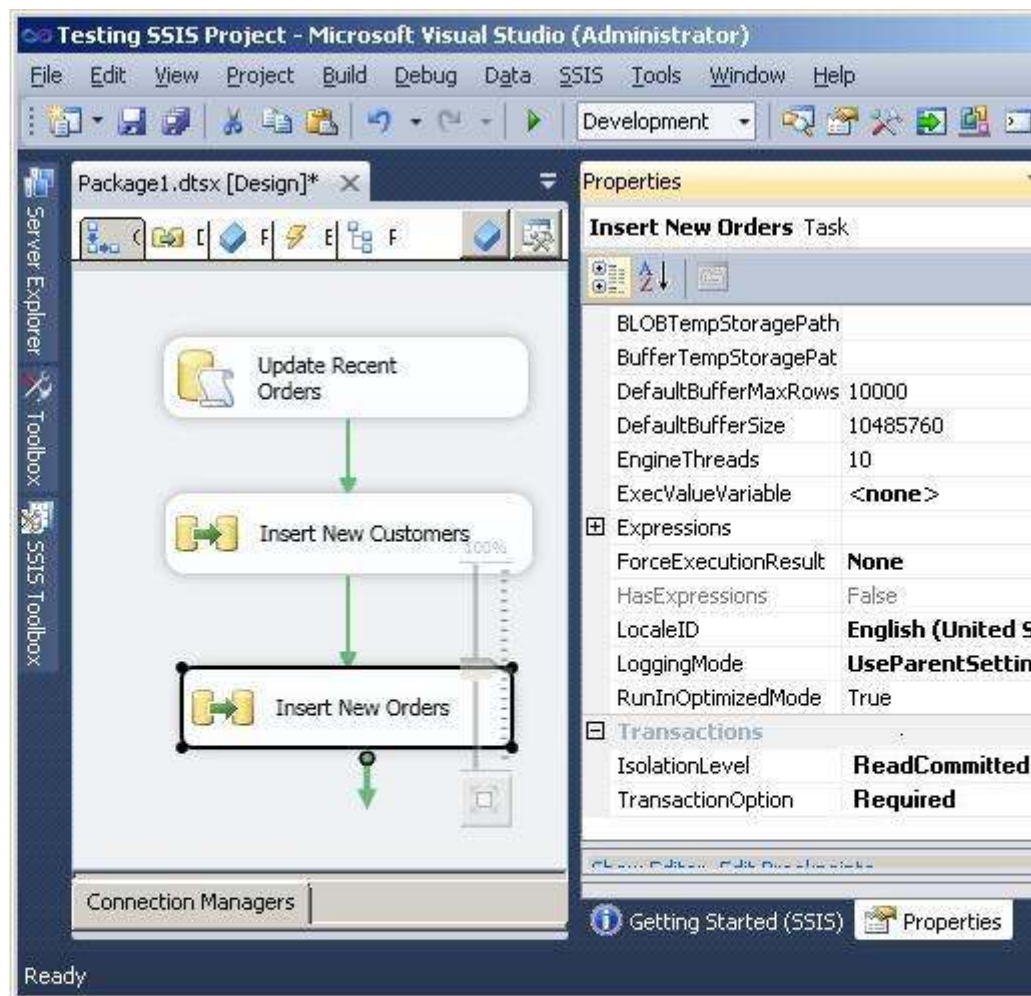
QUESTION 2
You are designing an extract, transform, load (ETL) process with SQL Server Integration Services (SSIS). Two packages, Package A and Package B, will be designed. Package A will execute Package B. Both packages must reference a file path corresponding to an input folder where files will be located for further processing. You

need to design a solution so that the file path can be easily configured with the least administrative and development effort. Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
	Create an SSIS project containing Package A and Package B.
	Add a project parameter named FilePath.
	Add a package parameter named FilePath to Package A and Package B.
	Deploy the project to the SSIS catalog.
<< Move	Configure the project to use an Environment that sets the FilePath parameter.
Remove >>	Add a Parent Package Variable configuration in Package B to set the FilePath variable.
	Add a configuration file in Package A to set the FilePath value.
	Create one SSIS project containing only Package A and another SSIS project containing only Package B.

Correct Answer:

Create an SSIS project containing Package A and Package B.
Add a project parameter named FilePath.
Deploy the project to the SSIS catalog.
Configure the project to use an Environment that sets the FilePath parameter.

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference:
http://msdn.microsoft.com/en-us/library/hh479588.aspx
http://msdn.microsoft.com/en-us/library/hh213290.aspx
http://msdn.microsoft.com/en-us/library/hh213373.aspx

QUESTION 3
A SQL Server Integration Services (SSIS) package named DataFeed interacts with an external vendor data feed. The package is executed several times a day, either as part of other packages' control flow or by itself. The external data feed is unreliable because network failures and slow response times are frequent. The package is currently deployed on the file system. To analyze the reliability of the external data feed, you must

collect execution data. Every time the DataFeed package is executed, the following information must be logged:
• Start Time
• End Time
• Execution Result
• Execution Duration
You need to design a logging solution that meets the requirements by using the least amount of administrative and development effort. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title		Answer Choices Title
▲ ▼		Add OnError and On Warning event handlers
		Query the catalog.executable_statistics view
		Query the ExecutionLog Table.
	<< Move	Deploy the project that contains the package to the SSIS Catalog and execute the package on the server.
	Remove >>	Deploy the package to the msdb database.
		Add an Execute SQL task to the event handlers.
		Query the catalog.execution_data_statistics view.
		Create a new Project and add the package to the project.

Correct Answer:

Create a new Project and add the package to the project.

Deploy the project that contains the package to the SSIS Catalog and execute the package on the server.

Query the catalog.executable_statistics view

Section: (none)
Explanation

Explanation/Reference:

QUESTION 4
You use SQL Server Integration Services (SSIS) for extract, transformation, load (ETL) processing.
Issues concerning addresses are discovered in the data warehouse that you manage. Validation must separate the addresses into three categories:
• Valid addresses
• Autocorrected addresses
• Invalid addresses
You need to enhance the SSIS packages to perform address validation by using an external service.

Which transformation should you use? (To answer, drag the appropriate transformation from the list of options to the correct location or locations in the answer area.)

Select and Place:

Lookup
Term Lookup
Fuzzy Lookup
Fuzzy Grouping
DQS Cleansing
Slowly Changing Dimension

OLE DB Source

Conditional Split

Invalid Addresses

Valid Addresses

Autocorrected Addresses
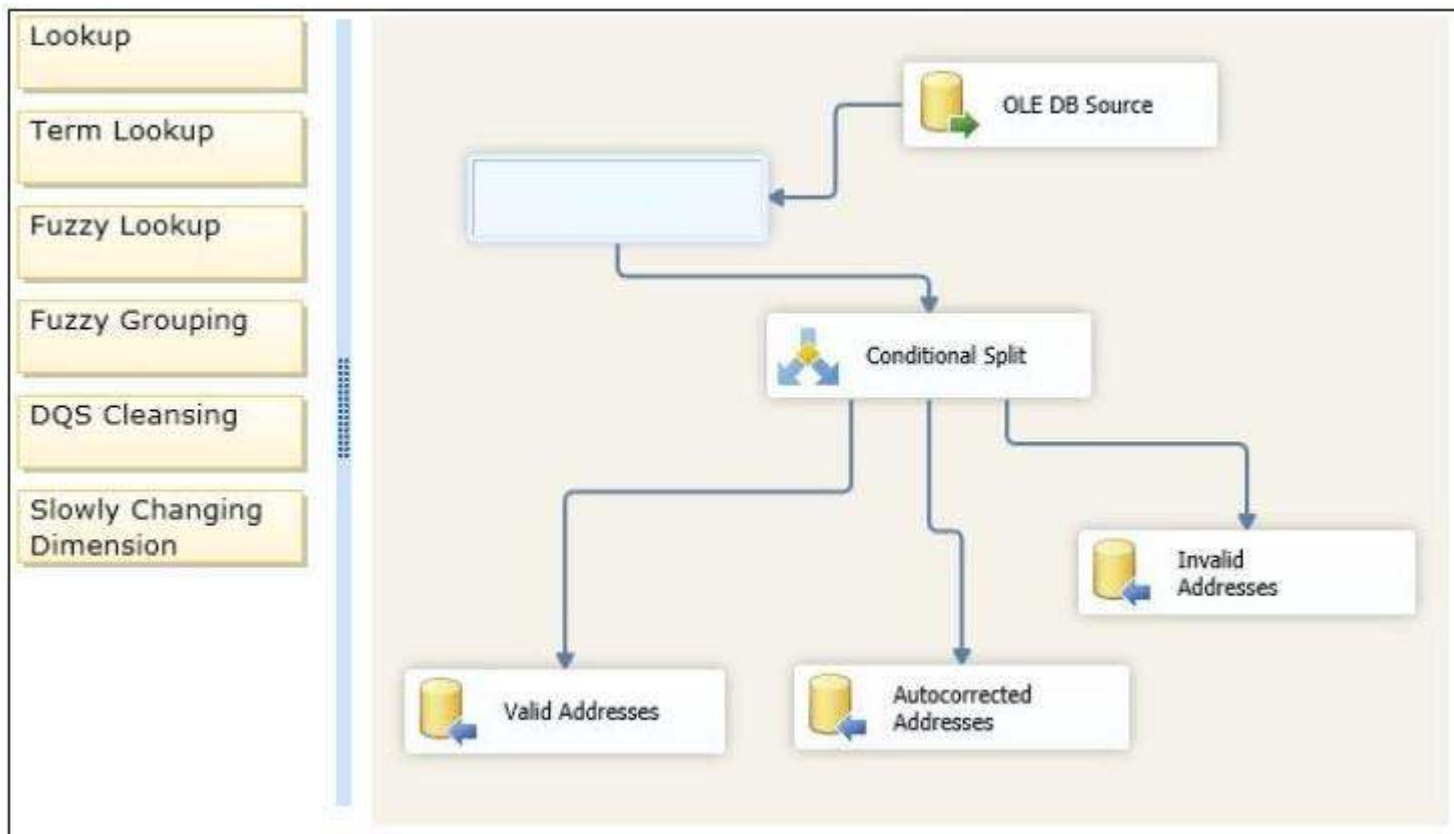
Correct Answer:
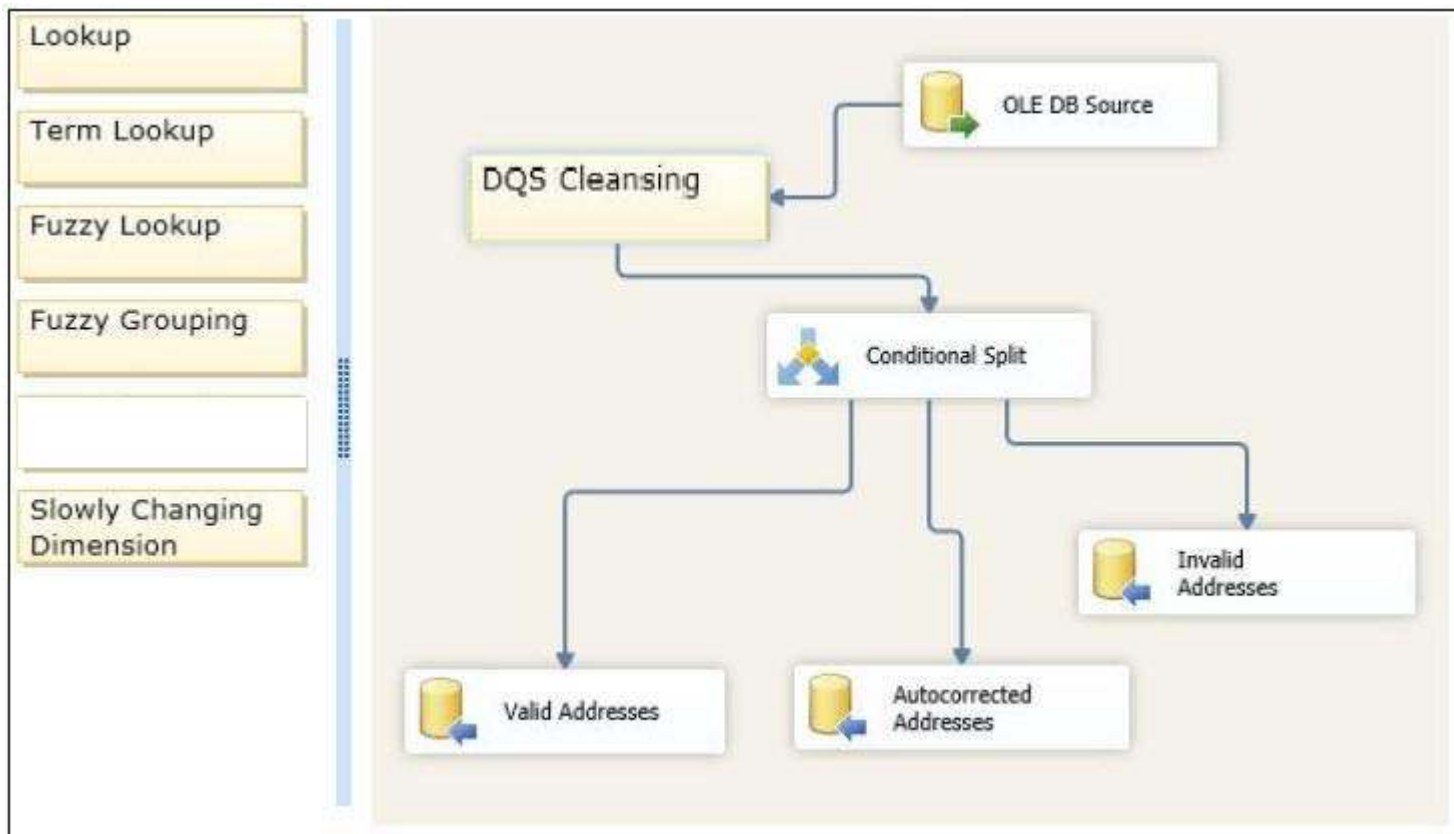
Section: (none)
Explanation

Explanation/Reference:

QUESTION 5
You are designing a SQL Server Integration Services (SSIS) package configuration strategy. The package configuration must meet the following requirements:
• Include multiple properties in a configuration.
• Force packages to load all settings in the configuration.
• Support Encrypting File System (EFS) formats.
You need to select the appropriate configuration. Which configuration type should you use?
To answer, select the appropriate option from the drop-down list in the dialog box.

Hot Area:

Correct Answer:

Section: (none)
Explanation

Explanation/Reference:

QUESTION 6
DRAG DROP
You are developing a SQL Server Integration Services (SSIS) package.

The package uses custom functionality that accesses a SQL Server database.

The custom functionality must be implemented by using Language Integrated Query (LINQ).

You need to ensure that the LINQ code can be debugged at design time.

What should you select from the SSIS Toolbox? (To answer, drag the appropriate item to the correct location in the answer area.)

Select and Place:

ADO NET Destination
DataReader Destination
Slowly Changing Dimension
Execute Process Task
Script Component
Data Conversion
Conditional Split

Package.dtsx [Design]* ✕

Control... | Data Flow | Parame... | Event... |

Data Flow Task: Data Flow Task

Correct Answer:

ADO NET Destination

DataReader Destination

Slowly Changing Dimension

Execute Process Task

Data Conversion

Conditional Split

Package.dtsx [Design]* ✕

Control... Data Flow Parame... Event...

Data Flow Task: Data Flow Task

Script Component
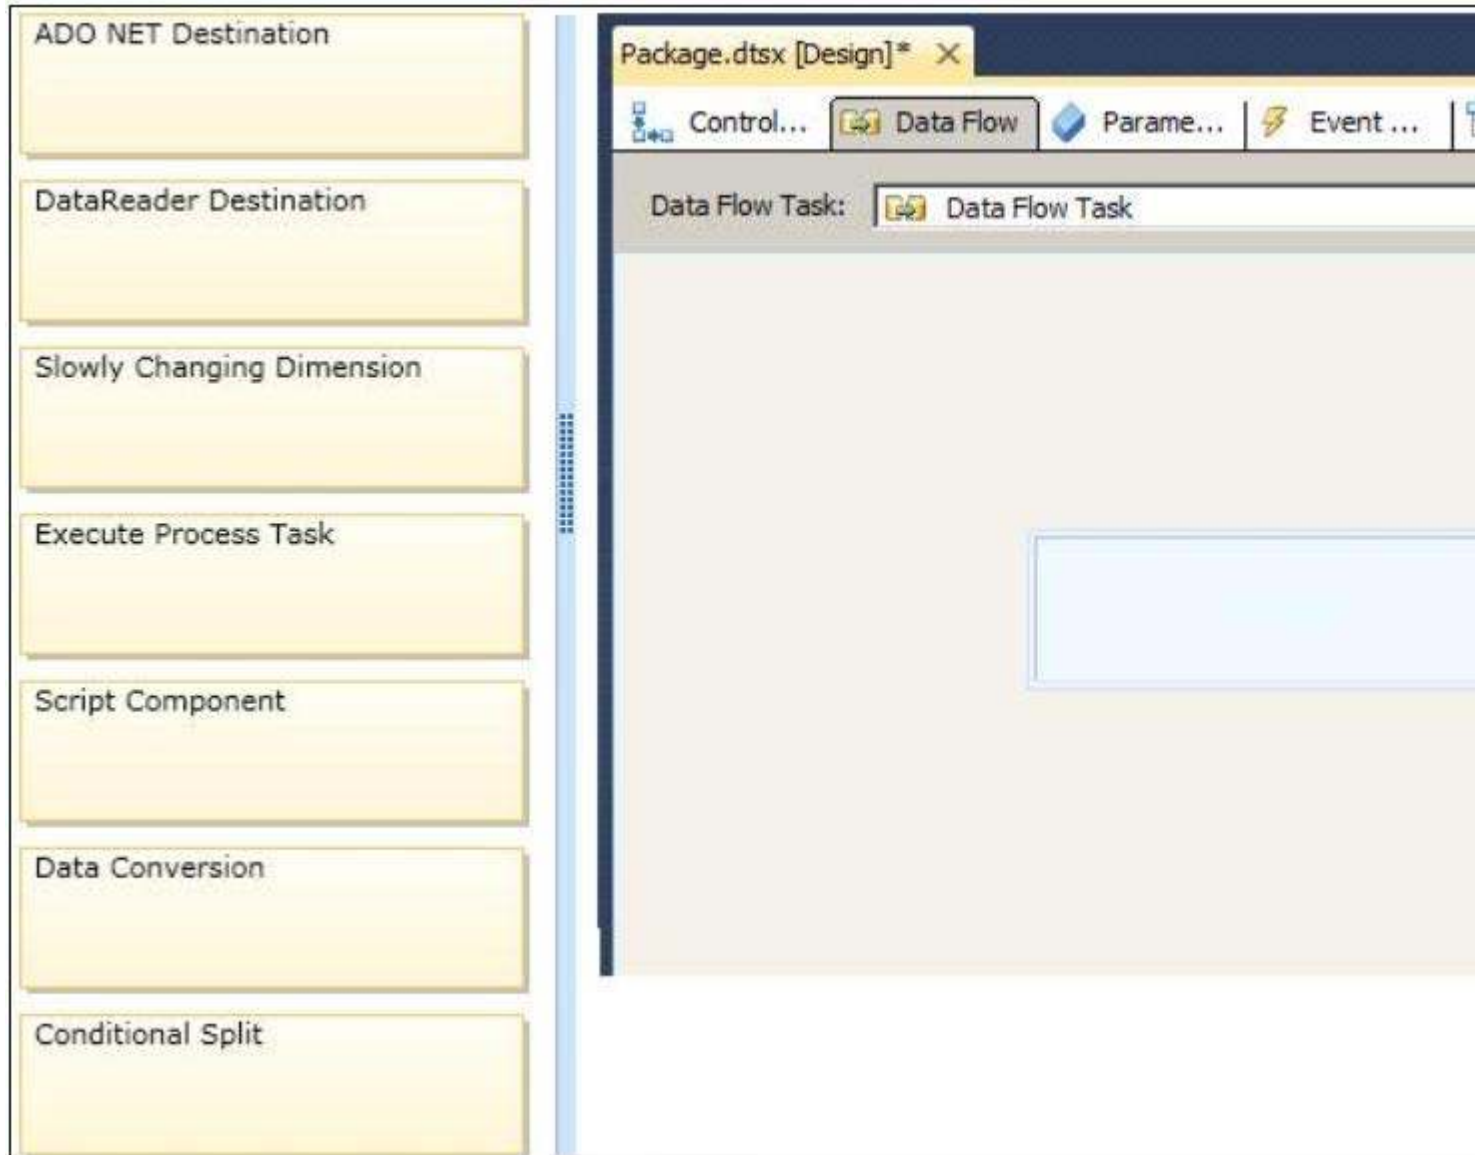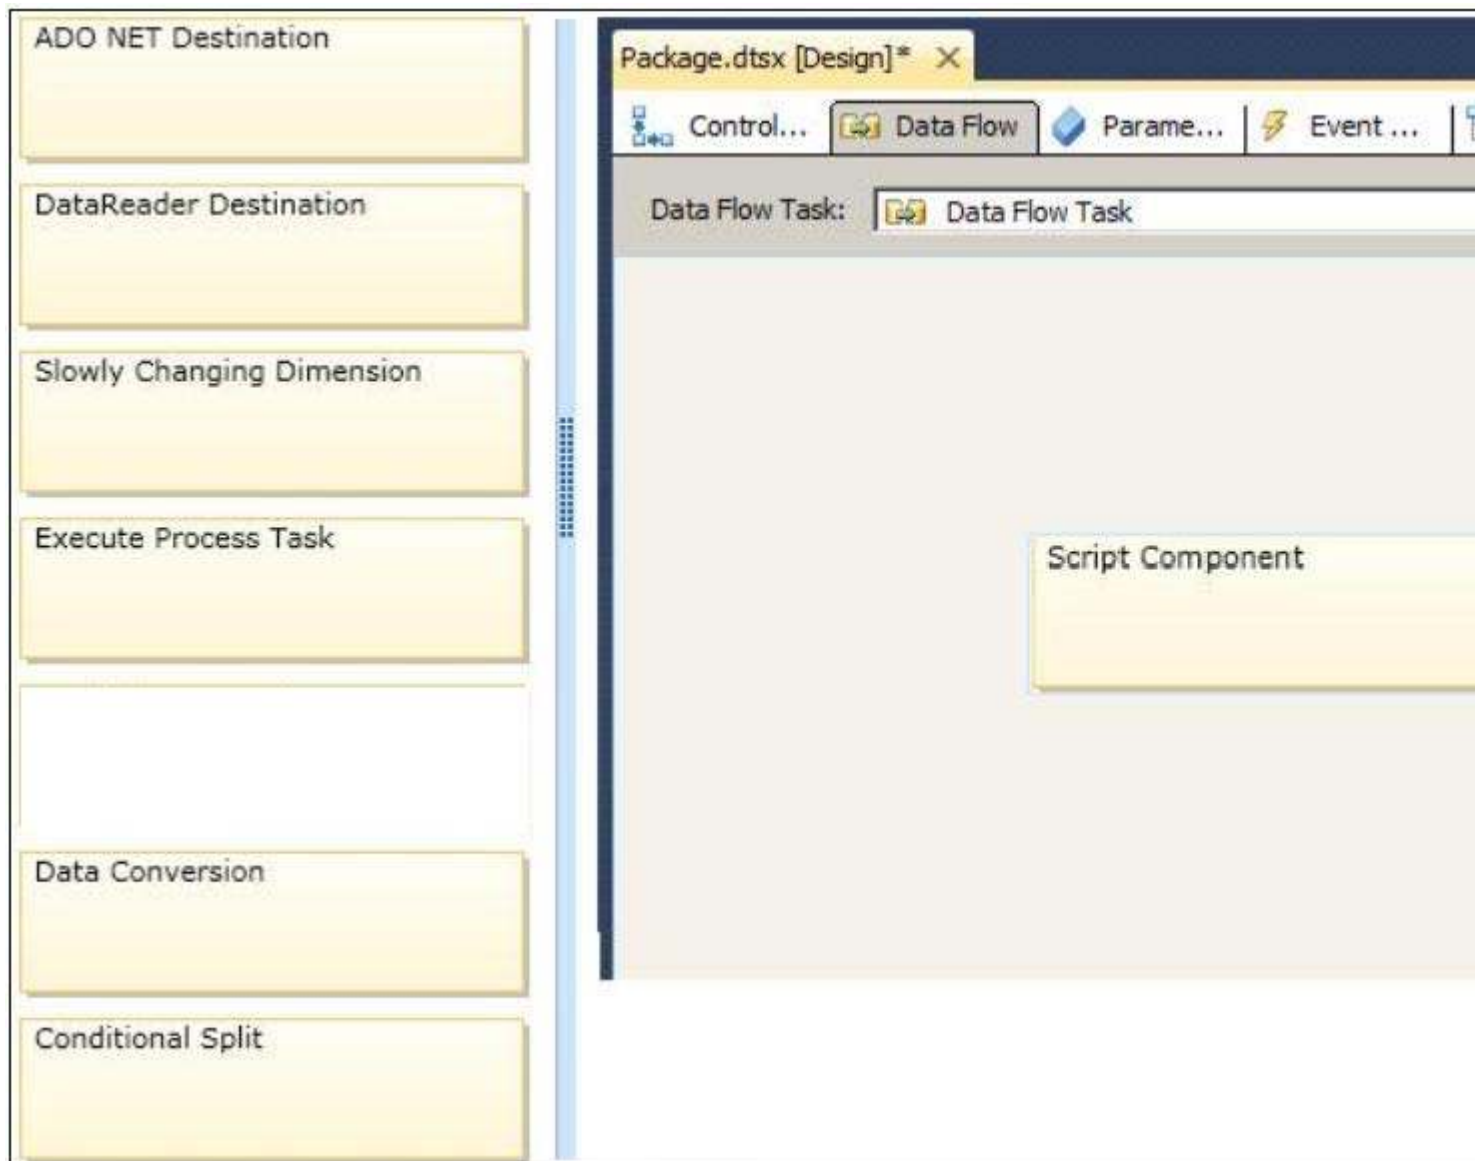
Section: (none)
Explanation

Explanation/Reference:

QUESTION 7
HOTSPOT
You are editing a SQL Server Integration Services (SSIS) package. You need to edit the Script task that has an expression defined. Which task should you select? To answer, select the appropriate setting or settings in the answer area.

Hot Area:
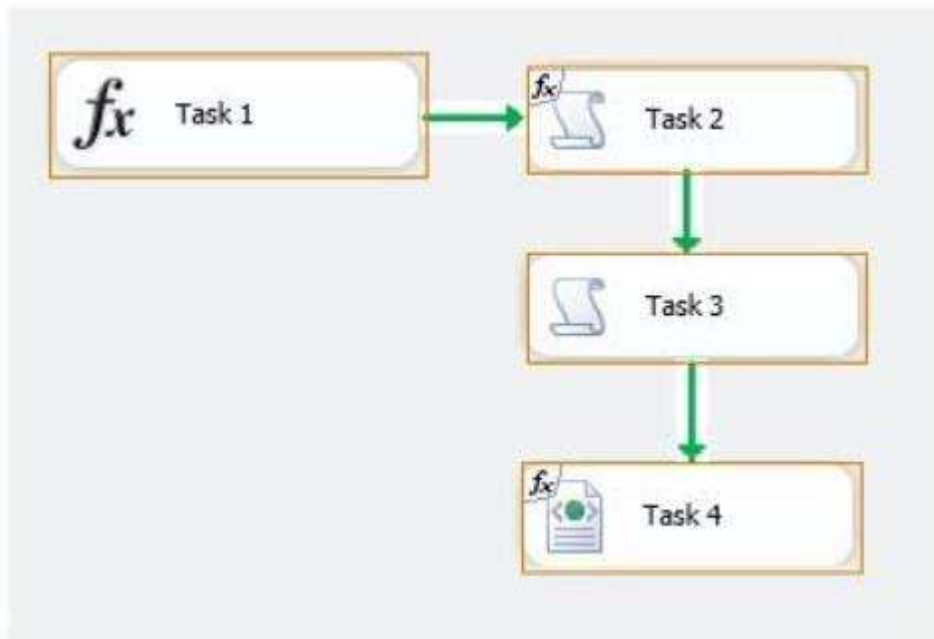
Correct Answer:
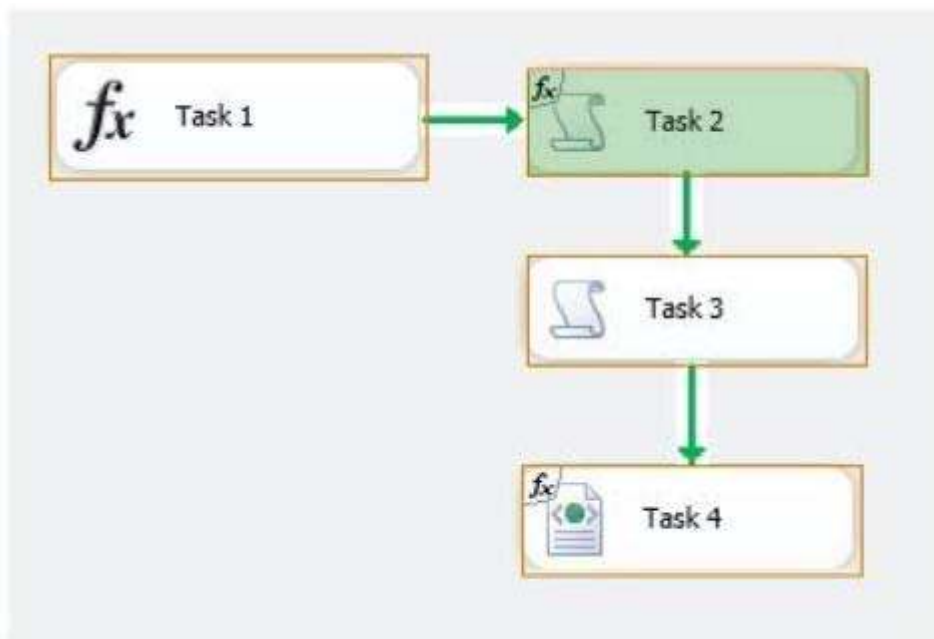
Section: (none)
Explanation

Explanation/Reference:

QUESTION 8
HOTSPOT
You are editing a SQL Server Integration Services (SSIS) project named Project1 in SQL Server Data Tools. A package Connection Manager has been parameterized with project scope. You need to display the parameters that have been generated for the Connection Manager. What should you use? To answer, select the appropriate setting or settings in the answer area.

Hot Area:

Solution Explorer

| Parameters | Event Handlers | Package Explorer |

Project1
- Project.params
- Connection Mana
- ☐ SSIS Packages
  - Package.dtsx
- Miscellaneous

Correct Answer:

Solution Explorer

| Parameters | Event Handlers | Package Explorer |

Project1
- Project.params
- Connection Mana
- ☐ SSIS Packages
  - Package.dtsx
- Miscellaneous

Section: (none)
Explanation

Explanation/Reference:

QUESTION 9
You administer a Microsoft SQL Server database. Service accounts for SQL Agent are configured to use a local user. A Microsoft SQL Server Integration Services (SSIS) job step has been created within a SQL Server Agent job. The SSIS package accesses a network share when exporting data from a SQL Server database. When you execute the SQL Server Agent job, it fails due to a permissions failure on a share on a remote server. You need to ensure that the SQL Server Agent job can execute the SSIS package. Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
	Add a proxy that references the local user.
	Add a proxy that references the credential.
<< Move	Create a local user account and grant local administrator on the SQL Server instance.
Remove >>	Create a credential that references the local user.
	Create a credential that references the domain user.
	Assign the proxy to the Operating System subsystem.
	Assign the proxy to the SSIS package execution subsystem.
	Create a domain user account and grant permissions to the domain user account to access the network share.

Correct Answer:

Create a domain user account and grant permissions to the domain user account to access the network share.

Create a credential that references the domain user.

Add a proxy that references the credential.

Assign the proxy to the SSIS package execution subsystem.

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms175834.aspx
http://msdn.microsoft.com/en-us/library/ms189522.aspx
http://msdn.microsoft.com/en-us/library/ms190703.aspx
http://msdn.microsoft.com/en-us/library/ms161950.aspx

QUESTION 10
You are creating a SQL Server Integration Services (SSIS) package to populate a fact table from a source table. The fact table and source table are located in a SQL Azure database. The source table has a price field and a tax field. The OLE DB source uses the data access mode of Table. You have the following requirements:
• The fact table must populate a column named TotalCost that computes the sum of the price and tax columns.
• Before the sum is calculated, any records that have a price of zero must be discarded.
You need to create the SSIS package in SQL Server Data Tools.

In what sequence should you order four of the listed components for the data flow task? (To answer, move the appropriate components from the list of components to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
	Conditional Split
	Data Conversion
	OLE DB Destination
	Multicast
	Derived Column
	Copy Column
	Data Profiling Task
	OLE DB Source

Correct Answer:

OLE DB Source
Conditional Split
Derived Column
OLE DB Destination

Section: (none)
Explanation

Explanation/Reference:

QUESTION 11
DRAG DROP
You are editing a SQL Server Integration Services (SSIS) package that uses checkpoints.
The package performs the following steps:
1. Download a sales transaction file by using FTP.
2. Truncate a staging table.
3. Load the contents of the file to the staging table.
4. Merge the data with another data source for loading to a data warehouse.

The checkpoints are currently working such that if any of the four steps fail, the package will restart from the failed step the next time it executes. You need to modify the package to ensure that if either the Truncate Staging Table or the Load Sales to Staging task fails, the package will always restart from the Truncate Staging Table task the next time the package runs. Which three steps should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Correct Answer:

Open the SSIS Toolbox

Move the Truncate Staging Table and Load
Sales to Staging tasks to add a new Sequence
container and set their FailParentOnFailure
properties to True

Set the FailPackageOnFailure property of the
container to True

Section: (none)
Explanation

Explanation/Reference:

QUESTION 12
A new SQL Server Integration Services (SSIS) project is deployed to the SSIS catalog. To troubleshoot some
data issues, you must output the data streaming through several data flows into text files for further analysis.
You have the list of data flow package paths and identification strings of the various task components that must
be analyzed. You need to create these output files with the least amount of administrative and development
effort. Which three stored procedures should you execute in sequence? (To answer, move the appropriate
actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
	Catalog.create_folder
	catalog.Create_Execution_dump
<< Move	Catalog.add_data_tap
Remove >>	Catalog.configure_catalog
	Catalog.add_data_tap_by_guid
	Catalog.create_Execution
	Catalog.Start_execution

Correct Answer:

Catalog.create_Execution
Catalog.add_data_tap
Catalog.Start_execution

Section: (none)
Explanation

Explanation/Reference:
correct

QUESTION 13
DRAG DROP
You are building a fact table in a data warehouse. The table must have a columnstore index. The table cannot
be partitioned. You need to design the fact table and load it with data. Which three actions should you perform
in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange

them in the correct order.)

Build List and Reorder:

Ordered List Title		Answer Choices Title
▲ ▼	<< Move Remove >>	Enable XVelocity on the server.
		Create the columnstore index.
		Load the data.
		Create the table
		Set the Allow Snapshot Isolation database property to True

Correct Answer:

Create the table
Load the data.
Create the columnstore index.

Section: (none)
Explanation

Explanation/Reference:

QUESTION 14
DRAG DROP
You are developing a SQL Server Integration Services (SSIS) package that is ready for deployment to a production server. The package contains sensitive information secured by using the EncryptSensitiveWithUserKey package protection level. You are preparing the package for deployment by the production operations team. You need to ensure that the production operations team can open and execute the package without re-entering the sensitive information. Which three steps should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title		Answer Choices Title
		Open the package in SQL Server Data Tools
		In the Properties window, change the ProtectionLevel property to EncryptSensitiveWithPassword and enter the PackagePassword
		Use the [Encrypt] option with the required protection level and password
	<< Move	Run the dtexecUI command prompt utility
	Remove >>	Click the task which contains the sensitive information
		Click the design surface of the Control Flow designer
		Run the dtexec command prompt utility
		When prompted, reset the package protection level for production

Correct Answer:

Open the package in SQL Server Data Tools

Click the design surface of the Control Flow designer

In the Properties window, change the ProtectionLevel property to EncryptSensitiveWithPassword and enter the PackagePassword

Section: (none)
Explanation

Explanation/Reference:

QUESTION 15
You are designing an extract, transform, load (ETL) process with SQL Server Integration Services (SSIS). Two packages, Package A and Package B, will be designed. Package A will execute Package B. Both packages must reference a file path corresponding to an input folder where files will be located for further processing. You need to design a solution so that the file path can be easily configured with the least administrative and development effort. Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title		Answer Choices Title
		Create an SSIS project containing Package A and Package B.
		Add a project parameter named FilePath.
		Add a package parameter named FilePath to Package A and Package B.
▲ ▼	<< Move / Remove >>	Deploy the project to the SSIS catalog.
		Configure the project to use an Environment that sets the FilePath parameter.
		Add a Parent Package Variable configuration in Package B to set the FilePath variable.
		Add a configuration file in Package A to set the FilePath value.
		Create one SSIS project containing only Package A and another SSIS project containing only Package B.

Correct Answer:

Create an SSIS project containing Package A and Package B.

Add a project parameter named FilePath.

Deploy the project to the SSIS catalog.

Configure the project to use an Environment that sets the FilePath parameter.

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

Reference:
http://msdn.microsoft.com/en-us/library/hh479588.aspx
http://msdn.microsoft.com/en-us/library/hh213290.aspx
http://msdn.microsoft.com/en-us/library/hh213373.aspx

QUESTION 16
You develop a SQL Server Integration Services (SSIS) project by using the Project Deployment model. The project contains many packages. It is deployed on a server named Development1. The project will be deployed to several servers that run SQL Server 2012. The project accepts one required parameter. The data type of the parameter is a string. A SQL Agent job is created that will call the master.dtsx package in the project. A job step is created for the SSIS package. The job must pass the value of an SSIS Environment Variable to the project parameter. The value of the Environment Variable must be configured differently on each server that runs SQL Server. The value of the Environment Variable must provide the server name to the project parameter. You need to configure SSIS on the Development1 server to pass the Environment Variable to the package. Which four actions should you perform in sequence by using SQL Server Management Studio? (To answer, move the

appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
	In the SSIS catalog, create a folder named DevelopmentEnv.
	In the SSIS catalog, create an Environment named DevelopmentEnv.
	In the DevelopmentEnv Environment, create an Environment Variable. Name the Environment Variable PackageEnvironment. Set the Environment Variable value to Development1.
<< Move	In the project configuration, create an Environment reference to DevelopmentEnv. Assign the value of the PackageEnvironment Environment Variable to the project parameter.
Remove >>	In the project configuration, edit the value of the project parameter to DevelopmentEnv.
	In the SQL Agent job, choose the DevelopmentEnv Environment in the Configuration tab of the SQL Server Integration Services job step.

Correct Answer:

In the SSIS catalog, create an Environment named DevelopmentEnv.

In the DevelopmentEnv Environment, create an Environment Variable. Name the Environment Variable PackageEnvironment. Set the Environment Variable value to Development1.

In the project configuration, create an Environment reference to DevelopmentEnv. Assign the value of the PackageEnvironment Environment Variable to the project parameter.

In the SQL Agent job, choose the DevelopmentEnv Environment in the Configuration tab of the SQL Server Integration Services job step.

Section: (none)
Explanation

Explanation/Reference:
Verified answer as correct.

References:
http://msdn.microsoft.com/en-us/library/hh479588.aspx
http://msdn.microsoft.com/en-us/library/hh213230.aspx
http://msdn.microsoft.com/en-us/library/hh213214.aspx

QUESTION 17
A SQL Server Integration Services (SSIS) project has been deployed to the SSIS catalog. The project
includes a project Connection Manager to connect to the data warehouse. The SSIS catalog includes
two Environments:
• Development
• QA
Each Environment defines a single Environment Variable named ConnectionString of type string. The
value of each variable consists of the connection string to the development or QA data warehouses.
You need to be able to execute deployed packages by using either of the defined Environments.
Which three actions should you perform in sequence? (To answer, move the appropriate actions
from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
	In the deployed project, set the ConnectionString property of the Connection Manager to use the ConnectionString Environment Variable
	In the deployed package, set the ConnectionString property of the connection Manager to use the ConnectionString Environment Variable
	Add the Environment references to the deployed project
	Deploy the project to the Development and QA Environments
	Add permissions to the Development and QA Environments
	Execute the deployed package by using the required Environment

`<< Move`
`Remove >>`

Correct Answer:

Deploy the project to the Development and QA
Environments

In the deployed project, set the
ConnectionString property of the Connection
Manager to use the ConnectionString
Environment Variable

Execute the deployed package by using the
required Environment

Section: (none)
Explanation

Explanation/Reference:

QUESTION 18
You are validating whether a SQL Server Integration Services (SSIS) package named Master.dtsx in the SSIS catalog is executing correctly. You need to display the number of rows in each buffer passed between each data flow component of the package. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
	Execute a SQL statement with a package name of Master.dtsx against the catalog.executions view and return its execution ID.
	Run the Master.dtsx package with its logging level set to Performance.
<< Move	Execute a SQL statement with the execution ID equal to the previous retrieved execution ID against the catalog.execution_data_statistics view and return the rows_sent column values for all the rows.
Remove >>	Run the Master.dtsx package with the logging level set to Verbose.
	Execute a SQL statement with a package name of Master.dtsx against the msdb..syssislog table and return its execution ID.

Correct Answer:

Run the Master.dtsx package with the logging level set to Verbose.

Execute a SQL statement with a package name of Master.dtsx against the catalog.executions view and return its execution ID.

Execute a SQL statement with the execution ID equal to the previous retrieved execution ID against the catalog.execution_data_statistics view and return the rows_sent column values for all the rows.

Section: (none)
Explanation

Explanation/Reference:

QUESTION 19
You plan to deploy a SQL Server Integration Services (SSIS) project by using the project deployment model. You need to monitor control flow tasks to determine whether any of them are running longer than usual. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of

actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title		Answer Choices Title
		Write a query against the catalog.operation_messages view. Add a calculation to the Query to compare durations to the catalog.executables view
		Execute the query
		Write a query against the catalog.execution_component_phases view. Add a calculation to the query to compare durations to the catalog.executables view.
	<< Move	Connect to the SSISDB database
	Remove >>	Connect to the msdb database
		Write a query against the catalog.execution_component_phases view. Add a calculation to the query to compare durations to the catalog.executions view.
		Write a query against the catalog.operation_messages view. Add a calculation to the Query to compare durations to the catalog.executions view

Correct Answer:

> Connect to the SSISDB database
>
> Write a query against the catalog.execution_component_phases view. Add a calculation to the query to compare durations to the catalog.executions view.
>
> Execute the query

Section: (none)
Explanation

Explanation/Reference:

QUESTION 20
DRAG DROP
A Data Flow task in a SQL Server Integration Services (SSIS) package produces run- time errors. You need to edit the package to log specific error messages. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title		Answer Choices Title
		Create an OnError event handler for the Data Flow task.
		Define an expression for the ActionAtEvent property
	<< Move	Add a Notify Operator task
	Remove >>	Add a WMI Event Watcher task
		In the ReadOnlyVariables list, select System::ErrorDescription and add conditional logic to invoke the FireError method
		Add a Script task

Correct Answer:

Create an OnError event handler for the Data
Flow task.

Add a Script task

In the ReadOnlyVariables list, select
System::ErrorDescription and add conditional
logic to invoke the FireError method

Section: (none)
Explanation

Explanation/Reference:

QUESTION 21
DRAG DROP
You are designing a SQL Server Integration Services (SSIS) package. The package moves order-related data to a staging table named Order. Every night the staging data is truncated and then all the recent orders from the online store database are inserted into the staging table. Your package must meet the following requirements:
▪ If the truncate operation fails, the package execution must stop and report an error.
▪ If the Data Flow task that moves the data to the staging table fails, the entire refresh operation must be rolled back.
▪ For auditing purposes, a log entry must be entered in a SQL log table after each execution of the Data Flow task.

The TransactionOption property for the package is set to Required. You need to design the package to meet the requirements. How should you design the control flow for the package? (To answer, drag the appropriate setting from the list of settings to the correct location or locations in the answer area.)
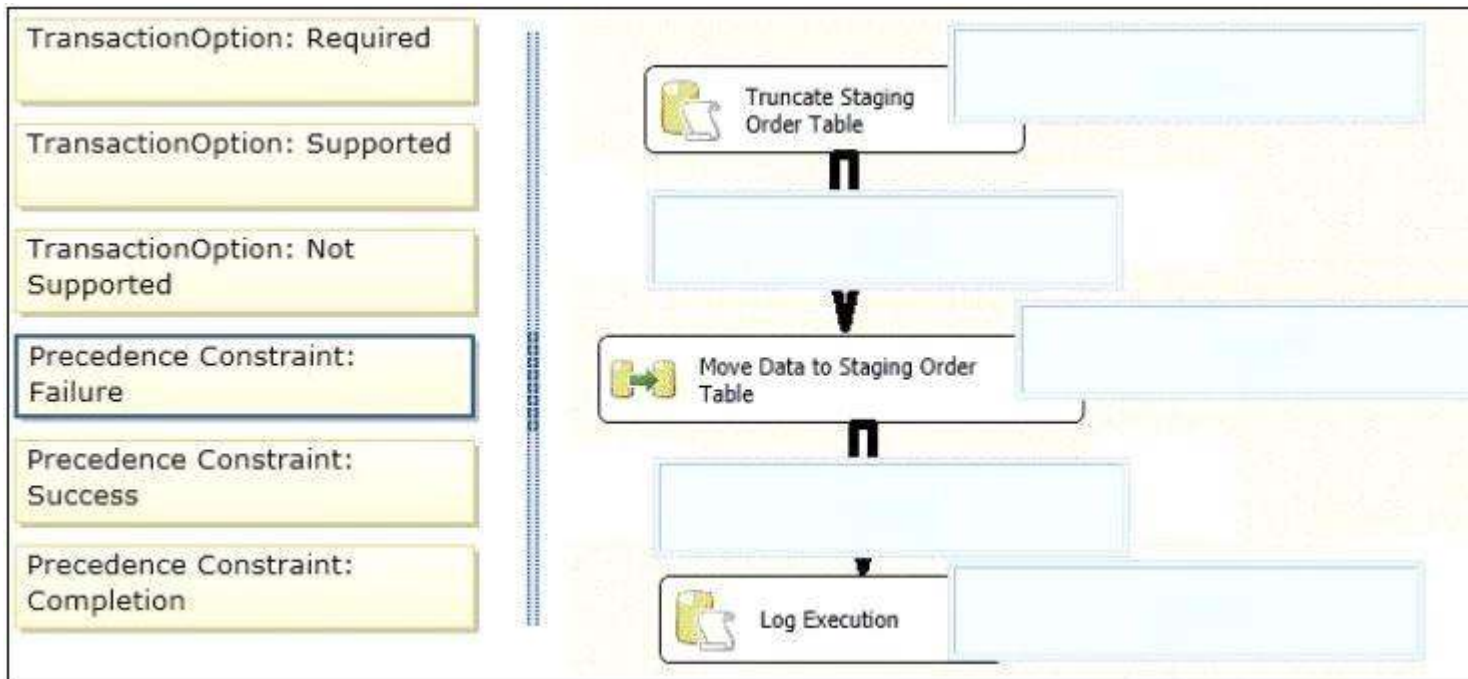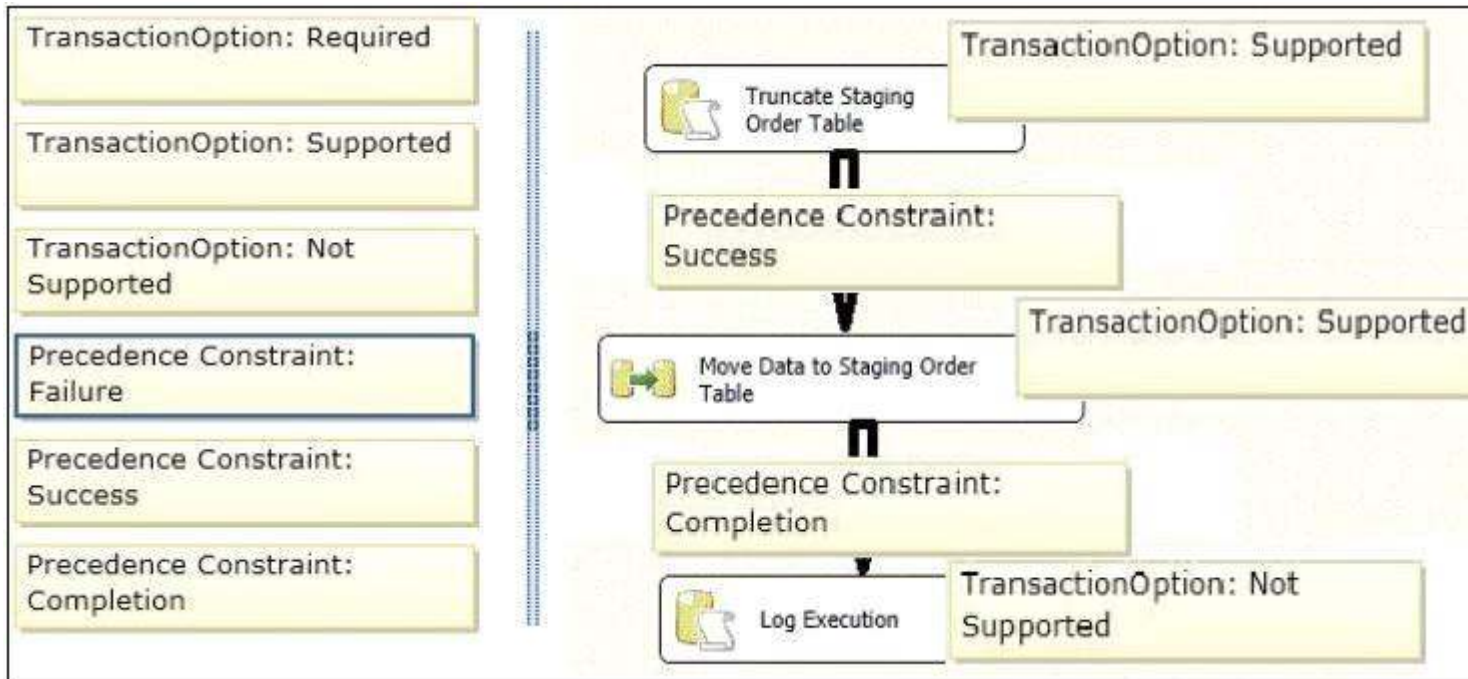
Select and Place:

Correct Answer:

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms137690.aspx
http://msdn.microsoft.com/en-us/library/ms141144.aspx

QUESTION 22
DRAG AND DROP
You are developing a SQL Server Integration Services (SSIS) package to insert new data into a data mart. The package uses a Lookup transformation to find matches between the source and destination. The data flow has the following requirements:

▪ New rows must be inserted.
▪ Lookup failures must be written to a flat file.

In the Lookup transformation, the setting for rows with no matching entries is set to Redirect rows to no match output. You need to configure the package to direct data into the correct destinations. How should you design the data flow outputs? (To answer, drag the appropriate transformation from the list of answer options to the correct location in the answer area.)
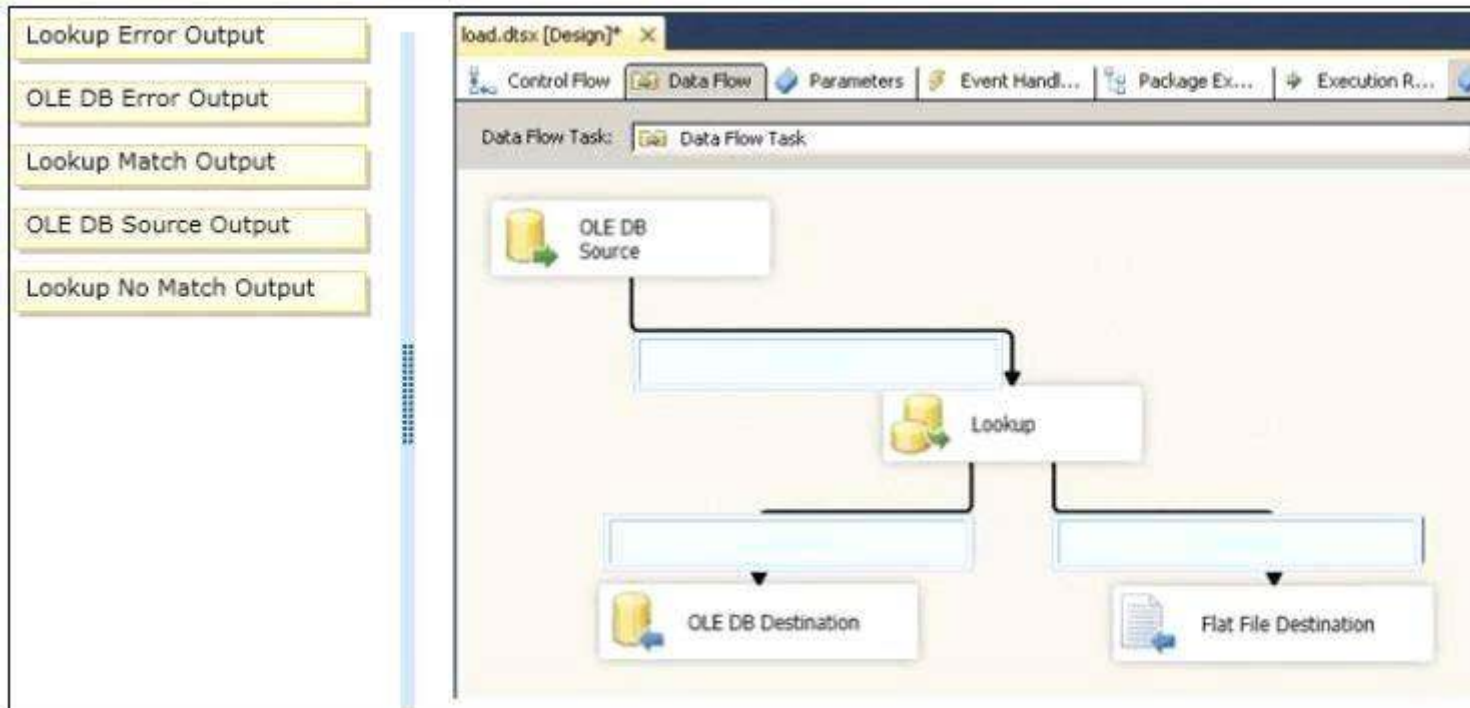
Select and Place:

Correct Answer:

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms141821.aspx
http://msdn.microsoft.com/en-us/library/bb895366.aspx

QUESTION 23
You are building a SQL Server Integration Services (SSIS) package to load product data sourced from a SQL
Azure database to a data warehouse. Before the product data is loaded, you create a batch record by using an
Execute SQL task named Create Batch. After successfully loading the product data, you use another Execute
SQL task named Set Batch Success to mark the batch as successful.

Package.dtsx [Design]* ✕

Control Flow | Data Flow | Parameters | Eve

Create Batch

Load Products

Set Batch Success

You need to create and execute an Execute SQL task to mark the batch as failed if either the Create Batch or Load Products task fails. Which three steps should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
	Set the Multiple Constraints option to Logical OR.
	Set the Evaluation Operation option to Expression OR Constraint.
	Connect the Create Batch and Load Products tasks to the Set Batch Failure task with failure constraints.
	Create an Execute SQL task named Set Batch Failure to mark the batch as failed.
	Set the Multiple Constraints option to Logical AND.
	Connect the Create Batch and Load Products tasks to the Set Batch Failure task with completion constraints.
	Set the ForceExecutionValue property of the Create Batch task to True.

<< Move
Remove >>

Correct Answer:

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms141003.aspx
http://msdn.microsoft.com/en-us/library/ms141261.aspx
http://msdn.microsoft.com/en-us/library/ms141722.aspx
http://msdn.microsoft.com/en-us/library/ms139895.aspx

QUESTION 24
You are developing a data flow transformation to merge two data sources. One source contains product data
and the other source contains data about the country in which the product was manufactured. Both data
sources contain a two-character CountryCode column and both use SQL Server. Both data sources contain an
ORDER BY clause to sort the data by the CountryCode column in ascending order. You use a Merge Join
transformation to join the data.

You need to ensure that the Merge Join transformation works correctly without additional transformations. What
should you do? (Each answer presents a part of the solution. Choose all that apply.)

A. Set the appropriate SortKeyPosition properties on the data sources.
B. Set the ValidateExternalMetaData property on the Merge Join transformation to True.
C. Set the IsSorted property on both data sources.
D. Set the MaxBuffersPerInput property on the Merge Join transformation to 2.

Correct Answer: AC
Section: (none)
Explanation

Explanation/Reference:

According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms141775.aspx
http://msdn.microsoft.com/en-us/library/ms137653.aspx
http://siddhumehta.blogspot.com/2009/05/validateexternalmetadata-property.html
http://msdn.microsoft.com/en-us/library/ms135950.aspx

QUESTION 25
HOTSPOT
You are developing a data flow to load sales data into a fact table. In the data flow, you configure a Lookup Transformation in full cache mode to look up the product data for the sale. The lookup source for the product data is contained in two tables. You need to set the data source for the lookup to be a query that combines the two tables. Which page of the Lookup Transformation Editor should you select to configure the query? To answer, select the appropriate page in the answer area.

Hot Area:

Correct Answer:

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms141821.aspx
http://msdn.microsoft.com/en-us/library/ms189697.aspx

QUESTION 26
DRAG DROP
You are creating a sales data warehouse. When a product exists in the product dimension, you update the product name. When a product does not exist, you insert a new record. In the current implementation, the DimProduct table must be scanned twice, once for the insert and again for the update. As a result, inserts and updates to the DimProduct table take longer than expected. You need to create a solution that uses a single command to perform an update and an insert. How should you use a MERGE T-SQL statement to accomplish this goal? (To answer, drag the appropriate answer choice from the list of options to the correct location or locations in the answer area. You may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Source

Target

Product

Inserted

MATCHED THEN

NOT MATCHED BY TARGET THEN

NOT MATCHED BY PRODUCT AND (Target.ProductID is NULL) THEN

MATCHED AND (Target.ProductID is NULL) THEN

MERGE dbo.DimProduct AS _____

USING (SELECT ProductID, ProductName, ProductColor, ProductCategory

   FROM dbo.StagingProduct) AS _____

ON (Target.ProductID = Source.ProductID)

WHEN _____

   UPDATE SET Target.ProductName = Source.ProductName

WHEN _____

   INSERT (ProductID, ProductName, ProductColor, ProductCategory)
   VALUES (Source.ProductID, Source.ProductName, Source.ProductColor, Source.ProductCategory)
OUTPUT $action, Inserted.*, Deleted.*;

Correct Answer:

Product

Inserted

NOT MATCHED BY PRODUCT AND (Target.ProductID is NULL) THEN

MATCHED AND (Target.ProductID is NULL) THEN

MERGE dbo.DimProduct AS   Target
USING (SELECT ProductID, ProductName, ProductColor, ProductCategory
     FROM dbo.StagingProduct) AS   Source
ON (Target.ProductID = Source.ProductID)
WHEN   MATCHED THEN
     UPDATE SET Target.ProductName = Source.ProductName
WHEN   NOT MATCHED BY TARGET THEN
     INSERT (ProductID, ProductName, ProductColor, ProductCategory)
     VALUES (Source.ProductID, Source.ProductName, Source.ProductColor, Source.ProductCategory)
OUTPUT $action, Inserted.*, Deleted.*;

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/bb510625.aspx
http://msdn.microsoft.com/en-us/library/cc280522.aspx
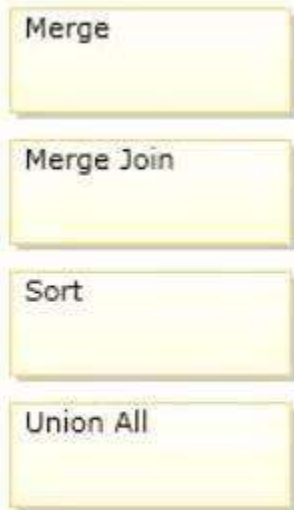
QUESTION 27
DRAG DROP
You are developing a SQL Server Integration Services (SSIS) package that imports unsorted data into a data warehouse hosted on SQL Azure. You have the following requirements:
▪   A destination table must contain all of the data in two source tables.
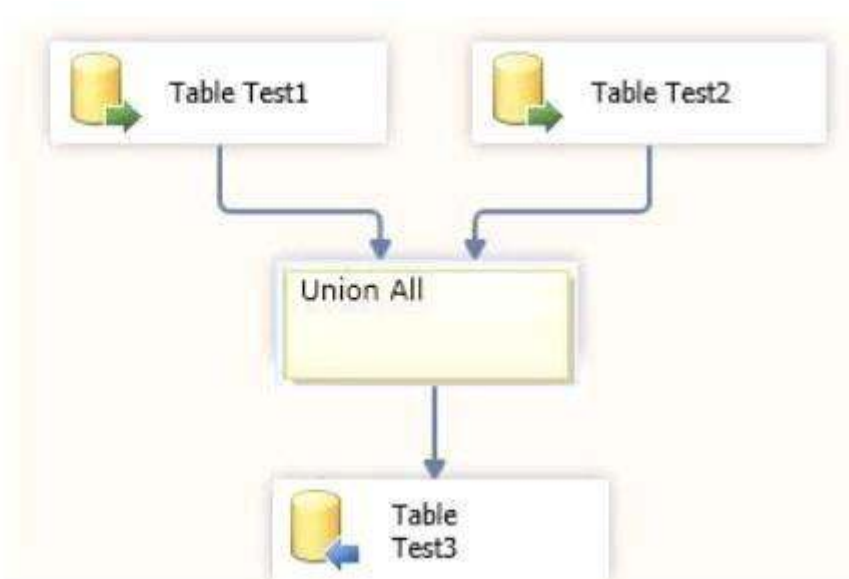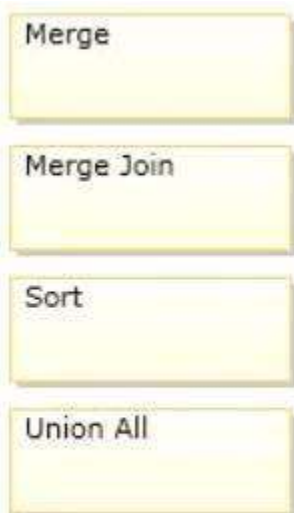▪   Duplicate records must be inserted into the destination table.

You need to develop a data flow that imports the data while meeting the requirements. How should you develop the data flow? (To answer, drag the appropriate transformation from the list of transformations to the correct location in the answer area.)

Select and Place:

Correct Answer:

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms141703.aspx
http://msdn.microsoft.com/en-us/library/ms141775.aspx
http://msdn.microsoft.com/en-us/library/ms141020.aspx
http://msdn.microsoft.com/en-us/library/ms140182.aspx

QUESTION 28
DRAG AND DROP
You are developing a SQL Server Integration Services (SSIS) package that imports data into a data

warehouse. You are developing the part of the SSIS package that populates the ProjectDates dimension table. The business key of the ProjectDates table is the ProjectName column. The business user has given you the dimensional attribute behavior for each of the four columns in the ProjectDates table:

▪ ExpectedStartDate - New values should be tracked over time.
▪ ActualStartDate - New values should not be accepted.
▪ ExpectedEndDate - New values should replace existing values.
▪ ActualEndDate - New values should be tracked over time.

You use the SSIS Slowly Changing Dimension Transformation. You must configure the Change Type value for each source column. Which settings should you select? (To answer, select the appropriate setting or settings in the answer area. Each Change Type may be used once, more than once, or not at all.)

Select and Place:

Dimension Columns	Change Type
ExpectedStartDate	
ActualStartDate	
ExpectedEndDate	
ActualEndDate	

Fixed Attribute

Changing Attribute

Historical Attribute

Correct Answer:

Dimension Columns	Change Type
ExpectedStartDate	Historical Attribute
ActualStartDate	Fixed Attribute
ExpectedEndDate	Changing Attribute
ActualEndDate	Historical Attribute

Fixed Attribute

Changing Attribute

Historical Attribute

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms141715.aspx
http://msdn.microsoft.com/en-us/library/ms141662.aspx

QUESTION 29
You administer a Microsoft SQL Server database. You want to import data from a text file to the database. You need to ensure that the following requirements are met:
▪ Data import is performed by using a stored procedure.

- Data is loaded as a unit and is minimally logged.

Which data import command and recovery model should you choose? (To answer, drag the appropriate data import command or recovery model to the appropriate location or locations in the answer area. Each data import command or recovery model may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Command/Recovery Model Name	Command/Recovery Model	
BCP	Data import command	
BULK INSERT	Recovery model	
Bulk-logged		
OPENDATASOURCE		
Full		

Correct Answer:

Command/Recovery Model Name	Command/Recovery Model	
BCP	Data import command	BULK INSERT
BULK INSERT	Recovery model	Bulk-logged
Bulk-logged		
OPENDATASOURCE		
Full		

Section: (none)
Explanation

Explanation/Reference:
According to these references, this answer looks correct.

References:
http://msdn.microsoft.com/en-us/library/ms162802.aspx
http://msdn.microsoft.com/en-us/library/ms188365.aspx
http://msdn.microsoft.com/en-us/library/ms175937.aspx
http://msdn.microsoft.com/en-us/library/aa337544.aspx