

Microsoft.Exactquestions.70-467.v2014-08-20.by.SONIA.50q

Number: 70-467
Passing Score: 800
Time Limit: 120 min
File Version: 24.5



MICROSOFT 70-467 EXAM QUESTIONS & ANSWERS

Exam Name: Designing Business Intelligence Solutions with Microsoft SQL Server 2012

Exact questions

QUESTION 1

DRAG DROP

You are designing a SQL Server Reporting Services (SSRS) solution. A report project must access multiple SQL Azure databases. Each database is on a different host. The databases have identical schema and security configurations.

You have the following requirements:

- The report must support subscriptions.
- Users must be able to select the host when running the report.

What should you do?

To answer, drag the appropriate phrase or phrases from the list to the correct location or locations in the answer area. (Answer choices may be used once, more than once, or not all.)

SQL Azure data.

SQL Azure hosts.

a shared dataset.

stored credentials.

integrated security.

data source in the report.

an expression-based connection string.

shared data source in the report

Create a

Create a report parameter that displays available values of

Create

Configure the data source to use

A.

SQL Azure data.

a shared dataset.

integrated security.

shared data source in the report

Create a

Create a report parameter that displays available values of

Create

Configure the data source to use

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* To include data in a report, you must first create data connections, also known as data sources, and then create datasets.

* A data connection includes the data source type, connection information, and the type of credentials to use. There are two types of data sources: embedded and shared. An embedded data source is defined in the report and used only by that report (fits this scenario). A shared data source is defined independently from a report

and can be used by multiple reports.

* Built-in data extensions include the following data connection types:

?Microsoft SQL Server

?Microsoft SQL Server Analysis Services

?Microsoft SharePoint List

?Windows Azure SQL Database

Etc.

* Expression-based connection strings are evaluated at run time. For example, you can specify the data source as a parameter, include the parameter reference in the connection string, and allow the user to choose a data source for the report.

* Credentials You provide the credentials that are needed to access the data. The data source owner must have granted you the appropriate permissions to access both the data source and the specific data on the data source.

Reference: Data Connections, Data Sources, and Connection Strings (SSRS)

QUESTION 2

DRAG DROP

You are designing a SQL Server Reporting Services (SSRS) solution. An existing report aggregates data from a SQL Server database in a chart. You need to use the chart in a new report and ensure that other users can use the chart in their reports. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

In Power View, open the report that contains the chart.

In Report Designer, insert the report part into a new report.

In Report Designer, open the report that contains the chart.

In Report Builder, insert the report part into a new report.

In Power View, insert the report part into a new report.

Select the chart for publication as a report part and publish the report.

<http://www.lead2pass.com/70-467.html>

A.

In Power View, open the report that contains the chart.

In Report Designer, insert the report part into a new report.

In Power View, insert the report part into a new report.

Select the chart for publication as a report part and publish the report.

In Report Builder, insert the report part into a new report.

In Report Designer, open the report that contains the chart.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* In Report Designer, after you create tables, charts, and other report items in a project, you can publish them as report parts to a report server or SharePoint site integrated with a report server so that you and others can reuse them in other reports.

* . By using Report Builder, you can customize and update reports that were created in SQL Server Data Tools (SSDT) Report Designer.

* In Report Builder, IT pros and power users can create powerful operational reports, and reusable report parts and shared datasets.

Incorrect:

* (incorrect) Power View, a feature of SQL Server 2012 Reporting Services Add-in for Microsoft SharePoint Server 2010 Enterprise Edition, is an interactive data exploration, visualization, and presentation experience. It provides intuitive ad-hoc reporting for business users such as data analysts, business decision makers, and information workers. They can easily create and interact with views of data from data models based on PowerPivot workbooks published in a PowerPivot Gallery, or tabular models deployed to SQL Server 2012 Analysis Services (SSAS) instances. Power View is a browser-based Silverlight application launched from SharePoint Server 2010 that enables users to present and share insights with others in their organization through interactive presentations.

Reference: Getting Started with Report Builder

Reference: Report Parts in Report Designer (SSRS)

QUESTION 3

You are designing a subscription strategy for a SQL Server Reporting Services (SSRS) report. You have an application that populates a table with user-specific subscription schedules and report formats. You need to ensure that users can receive reports by email according to their preferences. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create a standard subscription for each record in the table.
- B. Create a data-driven subscription for each record in the schedule table.
- C. Create one data-driven subscription. Schedule the subscription to frequently retrieve user preferences.
- D. Create a standard subscription for each subscription schedule.

Correct Answer: C

Section: (none)

Explanation

QUESTION 4

You are modifying a SQL Server Reporting Services (SSRS) report for a SQL Server Analysis Services (SSAS) cube. The report defines a report parameter of data type Date/Time with which users can filter the report by a single date. The parameter value cannot be directly used to filter the Multidimensional Expressions (MDX) query for the dataset. You need to ensure that the report displays data filtered by the user-entered value. You must achieve this goal by using the least amount of development effort. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Edit the dataset query parameter. Change the Value property of the report parameter to an expression that uses the same format as the date dimension member key value.
- B. Edit the dataset query parameter. Change the Name property of the dataset query parameter so that it points to a name value for each date dimension member.
- C. Edit the dataset query parameter. Create a subcube subquery that uses the StrToSet MDX function and accepts the report parameter value.
- D. Change the dataset query to Transact-SQL (T-SQL). Use the OPENROWSET function to query the cube. Output the cube results to the T-SQL query and use a Convert function to change the report parameter

value into the same format as the date dimension member.

Correct Answer: A

Section: (none)

Explanation

QUESTION 5

You are designing a strategy for an enterprise reporting solution that uses SQL Server Reporting Services (SSRS). Many of the SSRS reports will use common utilities and functions, including the following:

?Report utility functions and business logic in code
?Standardized report formatting properties such as fonts and colors for report branding
Formatting may change and new functions may be added as the reporting solution evolves. You need to create a strategy for deploying the formatting and code across the entire enterprise reporting solution. You must also ensure that reports can be easily updated to reflect formatting and function changes. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create a report as a template. Apply standardized formatting to the template. Store code in the Code section of the template.
- B. Build a web service that retrieves formatting properties and runs code. Call the web service through a report dataset.
- C. Store the formatting properties and code in database objects. Use stored procedures to populate a default value for report parameters and map each parameter to a corresponding formatting property.
- D. Create an assembly that contains formatting properties and code. Deploy the assembly on the Reporting Server and reference the assembly from each report.

Correct Answer: D


Section: (none)

Explanation

QUESTION 6

DRAG DROP

You are designing a dataset for a SQL Server Reporting Services (SSRS) report. The report includes the report items displayed in the following graphic.



The dataset is sourced from a commonly used stored procedure in an inventory data mart hosted in a SQL Azure database. It returns the status for all products across all storage sites. The report must display data for

the storage site that is selected by the Site report parameter. You cannot change the stored procedure code. You need to filter the dataset to use only data specific to the selected site.

How should you configure the filter?

To answer, drag the appropriate expression or expressions to the correct location or locations in the answer area. (Answer choices may be used once, more than once, or not all.)

Available Expressions:

- =DataSet!StorageSite.Value
- =Fields!StorageSite.Value
- =Parameters!Site.Value
- = "Site1"
- =SiteParameters.Value
- =StoredProc!StorageSite.Value

Change filters.

Include rows where the following conditions are true.

Expression: Text

Operator:

Value:

A.

Available Expressions:

- =DataSet!StorageSite.Value
-
-
- = "Site1"
- =SiteParameters.Value
- =StoredProc!StorageSite.Value

Change filters.

Include rows where the following conditions are true.

Expression: Text

Operator:

Value:

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

To set a filter on the dataset

?Open a report in Design view.

?Right-click a dataset in the Report Data pane and then click Dataset Properties. The Dataset Properties dialog box opens.

?Click Filters. This displays the current list of filter equations. By default, the list is empty.

?Click Add. A new blank filter equation appears.

?In Expression, type or select the expression for the field to filter. To edit the expression, click the expression (fx) button.

Box 1: Here we use the Fields expression.

?From the drop-down box, select the data type that matches the type of data in the expression you created in step 5.

?In the Operator box, select the operator that you want the filter to use to compare the values in the Expression box and the Value box. The operator you choose determines the number of values that are used from the next step.

Box 2: we test for equality.

?In the Value box, type the expression or value against which you want the filter to evaluate the value in Expression.

Box 3: we compare to the value of the Parameter named Site.

?Click OK.

Reference: How to: Add a Filter (Reporting Services)

QUESTION 7

You are designing a SQL Server Integration Services (SSIS) solution. The solution will contain an SSIS project that includes several SSIS packages. Each SSIS package will define the same connection managers and variables. You have the following requirements:

?Ensure that the deployment model supports changing the content of connection strings by using parameters at execution time.

?Ensure that the deployment model automatically starts from calls to the catalog.start_execution stored procedure in the SSISDB database.

?Maximize performance at execution time.

?Minimize development effort.

You need to design a solution that meets the requirements. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Use a project deployment model. Modify connection manager properties to use project parameters. Ensure that the SSISDB database is created.
- B. Use a project deployment model. Configure connections in an XML configuration file referenced by an environment variable that corresponds to the SQL Server environment of each SSIS package.
- C. Use a package deployment model. Use a SQL Server package configuration with a common filter. Change the contents of the SSIS Configurations table at runtime.
- D. Use a package deployment model. Save each SSIS package to a file share that can be accessed from all environments.

Correct Answer: A

Section: (none)

Explanation

QUESTION 8

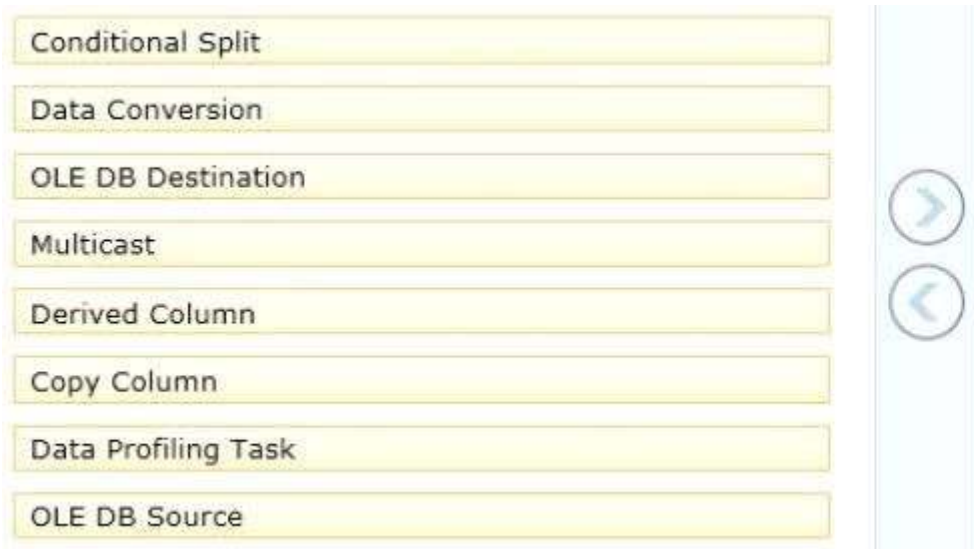
DRAG DROP

You are creating a SQL Server Integration Services (SSIS) package to populate a fact table from a source table. The fact table and source table are located in a SQL Azure database. The source table has a price field and a tax field. The OLE DB source uses the data access mode of Table. You have the following requirements:

?The fact table must populate a column named TotalCost that computes the sum of the price and tax columns.

?Before the sum is calculated, any records that have a price of zero must be discarded. You need to create the SSIS package in SQL Server Data Tools. In what sequence should you order four of the listed components for the data flow task? (To answer, move the appropriate components from the list of components to the answer area and arrange them in the correct order.)

<http://www.lead2pass.com/70-467.html>



A.



Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* You configure a Data Flow task by adding components to the Data Flow tab. SSIS supports three types of data flow components:

Sources: Where the data comes from

Transformations: How you can modify the data

Destinations: Where you want to put the data

* Creating a data flow includes the following steps:

/ Adding one or more sources to extract data from files and databases, and add connection managers to connect to the sources.

/ Adding the transformations that meet the business requirements of the package. A data flow is not required to include transformations.

Some transformations require a connection manager. For example, the Lookup transformation uses a connection manager to connect to the database that contains the lookup data. / Connecting data flow components by connecting the output of sources and transformations to the input of transformations and destinations.

/ Adding one or more destinations to load data into data stores such as files and databases, and adding connection managers to connect to the data sources. / Configuring error outputs on components to handle problems. At run time, row-level errors may occur when data flow components convert data, perform a lookup, or evaluate expressions. For example, a data column with a string value cannot be converted to an integer, or an expression tries to divide by zero. Both operations cause errors, and the rows that contain the errors can be

processed separately using an error flow. / Include annotations to make the data flow self-documenting.

* The capabilities of transformations vary broadly. Transformations can perform tasks such as updating, summarizing, cleaning, merging, and distributing data. You can modify values in columns, look up values in tables, clean data, and aggregate column values.

* The Data Flow task encapsulates the data flow engine that moves data between sources and destinations, and lets the user transform, clean, and modify data as it is moved. Addition of a Data Flow task to a package control flow makes it possible for the package to extract, transform, and load data.

A data flow consists of at least one data flow component, but it is typically a set of connected data flow components: sources that extract data; transformations that modify, route, or summarize data; and destinations that load data.

QUESTION 9

HOTSPOT

You are configuring the partition storage settings for a SQL Server Analysis Services (SSAS) cube. The partition storage must meet the following requirements:

?Optimize the storage of source data and aggregations in the cube.

?Use proactive caching.


?Drop cached data that is more than 30 minutes old. ?Update the cache when data changes, with a silence interval of 10 seconds.

You need to select the partition storage setting.

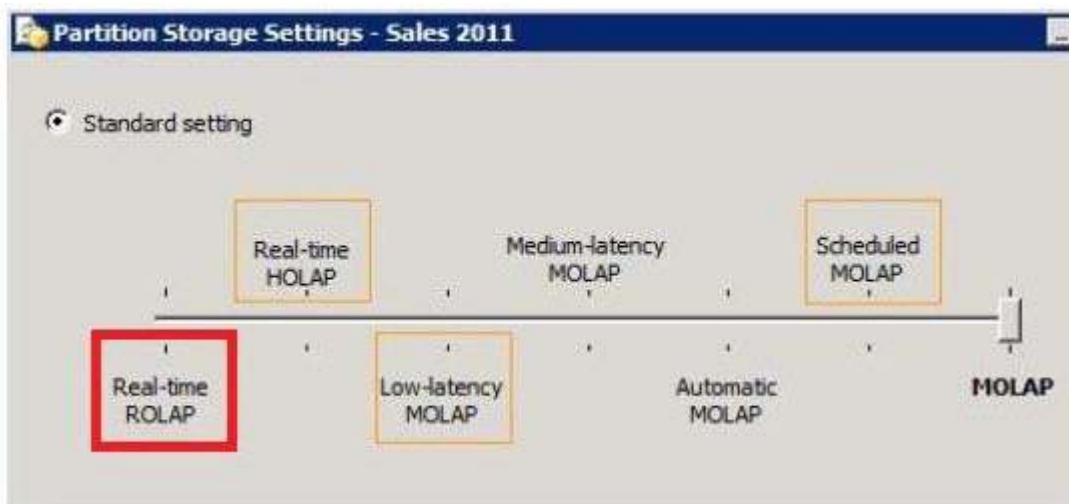
Which setting should you select?

To answer, select the appropriate setting in the answer area.

<http://www.lead2pass.com/70-467.html>



A.



Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* Real Time ROLAP

OLAP is in real time. Detail data and aggregations are stored in relational format. The server listens for notifications when data changes and all queries reflect the current state of the data (zero latency). This setting would typically be used for a data source with very frequent and continuous updates when the very latest data is always required by users. Depending on the types of queries generated by client applications, this method is liable to give the slowest response times.

Reference: Choosing a Standard Storage Setting

QUESTION 10

DRAG DROP

You administer a SQL Server Analysis Services (SSAS) instance. You need to capture a continuous log of detailed event and subevent durations and custom trace events from queries executed in the SSAS instance. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

<http://www.lead2pass.com/70-467.html>

Write a Multidimensional Expressions (MDX) script to query the DISCOVER_TRACES dynamic management view (DMV).

Write an XMLA script to log the extended events of the trace.

Launch SQL Server Profiler and connect to the instance.

Execute the script.

Launch SQL Server Management Studio and connect to the instance.

Configure the trace to save to a SQL Server database table.

A.

Write an XMLA script to log the extended events of the trace.

Launch SQL Server Profiler and connect to the instance.

Configure the trace to save to a SQL Server database table.

Launch SQL Server Management Studio and connect to the instance.

Write a Multidimensional Expressions (MDX) script to query the DISCOVER_TRACES dynamic management view (DMV).

Execute the script.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

mdx script "discover_traces"

Box 1: Launch SQL Server Profiler and connect to the instance. Box 2: Write an XMLA script to log the extended events of the trace.

Box 3: Execute the script.

Note:

* Auditing an instance of SQL Server or a SQL Server database involves tracking and logging events that occur on the system. The SQL Server Audit object collects a single instance of server- or database-level actions and groups of actions to monitor. The audit is at the SQL Server instance level. You can have multiple audits per SQL Server instance. The Server Audit Specification object belongs to an audit. You can create one server audit specification per audit, because both are created at the SQL Server instance scope.

* Trace events can be started and captured using SQL Server Profiler, or can be started from an XMLA command as SQL Server Extended Events and later analyzed.

* Extended Event tracing is enabled using a similar XMLA create object script.

Reference: Analysis Services Trace Events

QUESTION 11

DRAG DROP

You plan to deploy a SQL Server Integration Services (SSIS) project by using the project deployment model. You need to monitor control flow tasks to determine whether any of them are running longer than usual. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Write a query against the **catalog.operation_messages** view. Add a calculation to the query to compare durations to the **catalog.executables** view.

Execute the query.

Write a query against the **catalog.execution_component_phases** view. Add a calculation to the query to compare durations to the **catalog.executables** view.

Connect to the **SSISDB** database.

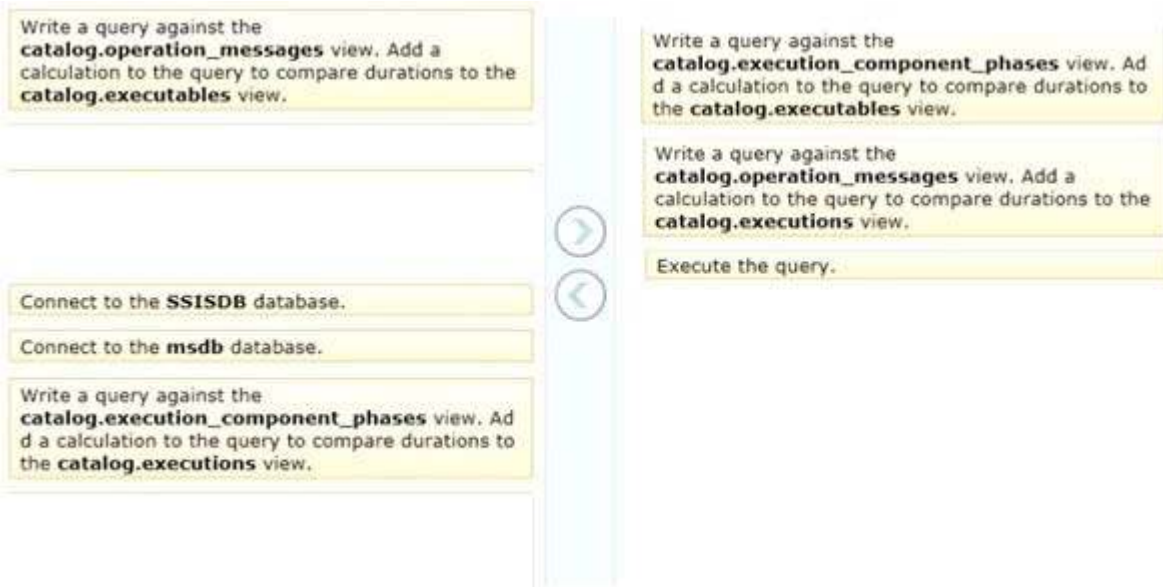
Connect to the **msdb** database.

Write a query against the **catalog.execution_component_phases** view. Add a calculation to the query to compare durations to the **catalog.executions** view.

Write a query against the **catalog.operation_messages** view. Add a calculation to the query to compare durations to the **catalog.executions** view.



A.



Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* execution_component_phases

Displays the time spent by a data flow component in each execution phase.

* The following example uses the catalog.execution_component_phases view to find the total amount of time that a specific package has spent executing in all phases (active_time), and the total elapsed time for the package (total_time).

use SSISDB

```

select package_name, task_name, subcomponent_name, execution_path, SUM(DATEDIFF
(ms,start_time,end_time)) as active_time, DATEDIFF(ms,min(start_time), max(end_time)) as total_time from
catalog.execution_component_phases
where execution_id = 1841
  
```

group by package_name, task_name, subcomponent_name, execution_path order by package_name, task_name, subcomponent_name, execution_path

* catalog.executables

This view displays a row for each executable in the specified execution. An executable is a task or container that you add to the control flow of a package.

*(incorrect) catalog.executions (SSISDB Database)

Displays the instances of package execution in the Integration Services catalog. Packages that are executed with the Execute Package task run in the same instance of execution as the parent package.

This view displays a row for each instance of execution in the catalog.

*(incorrect) catalog.operation_messages

Displays messages that are logged during operations in the Integration Services catalog. This view displays a row for each message that is logged during an operation in the catalog. The message can be generated by the server, by the package execution process, or by the execution engine.

Reference: catalog.execution_component_phases

Reference: catalog.executables

QUESTION 12

You are designing an extract, transform, load (ETL) process for loading data from a SQL Server database into a large fact table in a data warehouse each day with the prior day's sales data. The ETL process for the fact table must meet the following requirements:

?Load new data in the shortest possible time.

?Remove data that is more than 36 months old.

?Ensure that data loads correctly.

?Minimize record locking.

?Minimize impact on the transaction log.

You need to design an ETL process that meets the requirements. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Partition the destination fact table by date. Insert new data directly into the fact table and delete old data directly from the fact table.
- B. Partition the destination fact table by date. Use partition switching and staging tables both to remove old data and to load new data.
- C. Partition the destination fact table by customer. Use partition switching both to remove old data and to load new data into each partition.
- D. Partition the destination fact table by date. Use partition switching and a staging table to remove old data. Insert new data directly into the fact table.

Correct Answer: B

Section: (none)

Explanation

QUESTION 13

DRAG DROP

You are validating whether a SQL Server Integration Services (SSIS) package named Master.dtsx in the SSIS catalog is executing correctly. You need to display the number of rows in each buffer passed between each data flow component of the package. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Execute a SQL statement with a package name of Master.dtsx against the **catalog.executions** view and return its execution ID.

Run the Master.dtsx package with the logging level set to **Performance**.

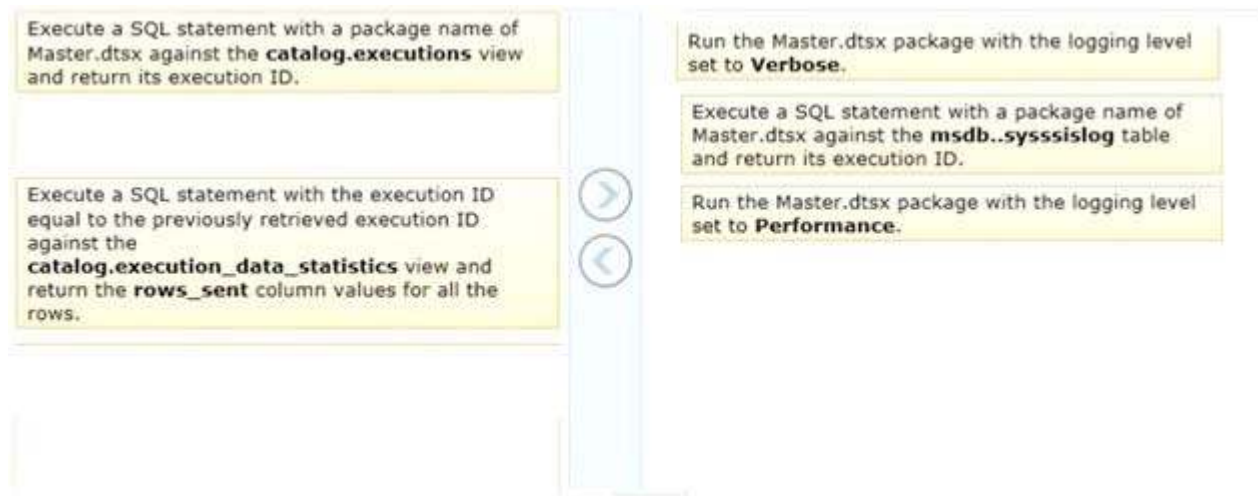
Execute a SQL statement with the execution ID equal to the previously retrieved execution ID against the **catalog.execution_data_statistics** view and return the **rows_sent** column values for all the rows.

Run the Master.dtsx package with the logging level set to **Verbose**.

Execute a SQL statement with a package name of Master.dtsx against the **msdb..sysssislog** table and return its execution ID.



A.



Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* You are going to become very very familiar indeed with [catalog].[executions]. It is a view that provides a record of all package executions on the server and, most importantly, it contains [execution_id] ?the identifier for each execution and the field to which all other objects herein will be related.

QUESTION 14

You are creating a Multidimensional Expressions (MDX) calculation for Projected Revenue in a cube. For Customer A, Projected Revenue is defined as 150 percent of the Total Sales for the customer. For all other customers, Projected Revenue is defined as 110 percent of the Total Sales for the customer. You need to calculate the Projected Revenue as efficiently as possible. Which calculation should you use? (More than one answer choice may achieve the goal. Select the BEST answer.)

- ☐ A.

```
CREATE MEMBER CurrentCube.[Measures].[Projected Revenue]
AS [Measures].[Total Sales];
SCOPE ([Customer].[Customer Name].MEMBERS, [Measures].[Projected Revenue]);
    [Measures].[Total Sales] * 1.1;
    IF [Customer].[Customer Name].CurrentMember.Name = "Customer A"
        THEN [Measures].[Total Sales] * 1.5
    END IF;
END SCOPE;
```
- ☐ B.

```
CREATE MEMBER CurrentCube.[Measures].[Projected Revenue]
AS CASE WHEN [Customer].[Customer Name].CurrentMember.Name = "Customer A"
THEN [Measures].[Total Sales] * 1.5
ELSE [Measures].[Total Sales] * 1.1 END
```
- ☐ C.

```
CREATE MEMBER CurrentCube.[Measures].[Projected Revenue]
AS [Measures].[Total Sales] * 1.1;
SCOPE ([Customer].[Customer Name].&[Customer A], [Measures].[Projected Revenue]);
    THIS = [Measures].[Total Sales] * 1.5;
END SCOPE;
```
- ☐ D.

```
CREATE MEMBER CurrentCube.[Measures].[Projected Revenue]
AS [Measures].[Total Sales] * 1.1;
SCOPE ([Customer].[Customer Name].MEMBERS, [Measures].[Projected Revenue]);
    [Customer].[Customer Name].&[Customer A] = [Measures].[Total Sales] * 1.5;
END SCOPE;
```

- A. Option A
B. Option B
C. Option C
D. Option D

Correct Answer: C


Section: (none)

Explanation

Explanation/Reference:

QUESTION 15

You are developing the database schema for a SQL Server Analysis Services (SSAS) BI Semantic Model (BISM). The BISM will be based on the schema displayed in the following graphic.



You have the following requirements:

?Ensure that queries of the data model correctly display average student age by class and average class level

by student.

?Minimize development effort.

You need to design the data model.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create a multidimensional project and define measures and a reference relationship.
- B. Create a tabular project and define calculated columns.
- C. Create a multidimensional project and define measures and a many-to-many dimensional relationship.
- D. Create a tabular project and define measures.

Correct Answer: C

Section: (none)

Explanation

QUESTION 16

You are modifying a star schema data mart that feeds order data from a SQL Azure database into a SQL Server Analysis Services (SSAS) cube. The data mart contains two large tables that include flags and indicators for some orders. There are 100 different flag columns, each with 10 different indicator values. Some flags reuse indicators. The tables both have a granularity that matches the fact table.

You have the following requirements:

?Allow users to slice data by all flags and indicators. ?Modify the date dimension table to include a surrogate key of a numeric data type and add the surrogate key to the fact table.

?Use the most efficient design strategy for cube processing and queries.

You need to modify the schema.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Define the surrogate key as an INT data type. Combine the distinct flag/indicator combinations into a single dimension.
- B. Define the surrogate key as an INT data type. Create a single fact dimension in each table for its flags and indicators.
- C. Define the surrogate key as a BIGINT data type. Combine the distinct flag/indicator combinations into a single dimension.
- D. Define the surrogate key as a BIGINT data type. Create a single fact dimension in each table for its flags and indicators.

Correct Answer: A

Section: (none)

Explanation

QUESTION 17

You are defining a named set by using Multidimensional Expressions (MDX) in a sales cube. The cube includes a Product dimension that contains a Category hierarchy and a Color attribute hierarchy. You need to return only the blue products in the Category hierarchy. Which set should you use? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A.
- B.
- C.
- D.

Correct Answer: C

Section: (none)

Explanation

QUESTION 18

An existing cube dimension that has 30 attribute hierarchies is performing very poorly. You have the following requirements:

?Implement drill-down browsing.

?Reduce the number of attribute hierarchies but ensure that the information contained within them is available to users on demand.

?Optimize performance.

You need to redesign the cube dimension to meet the requirements. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. set the AggregateFunction property to Sum on all measures. Use the SCOPE statement in a Multidimensional Expressions (MDX) calculation to tune the aggregation types.
- B. Set the AttributeHierarchyOptimizedState property to FullyOptimized on the attribute hierarchies.
- C. Create user-defined hierarchies. For the attributes sourced by the levels of the user-defined hierarchies, set the RelationshipType property to Rigid. Run incremental processing.
- D. Remove as many attribute hierarchies as possible from the dimension. Reintroduce the information in the attribute hierarchies as properties. Implement natural hierarchies and set the AttributeHierarchyVisible property to False for attributes used as levels in the natural hierarchies.

Correct Answer: D

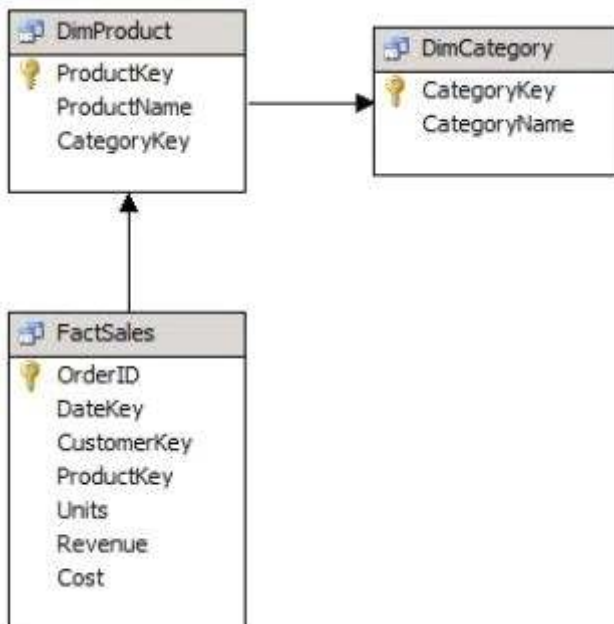
Section: (none)

Explanation

QUESTION 19

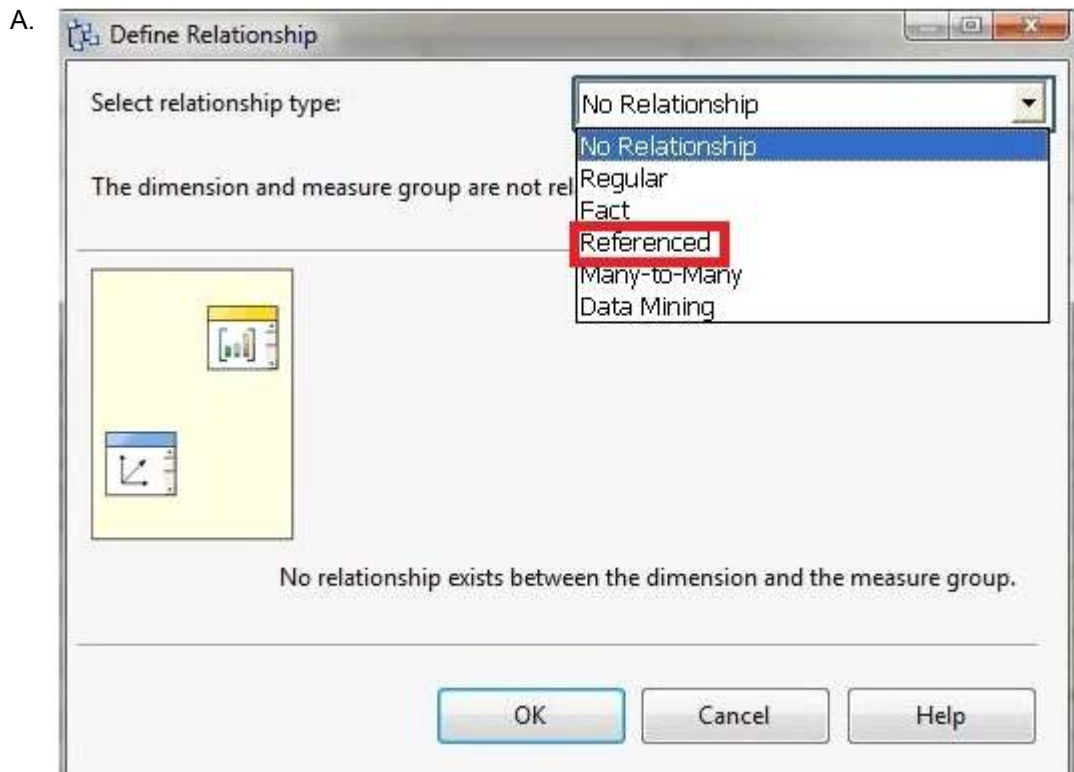
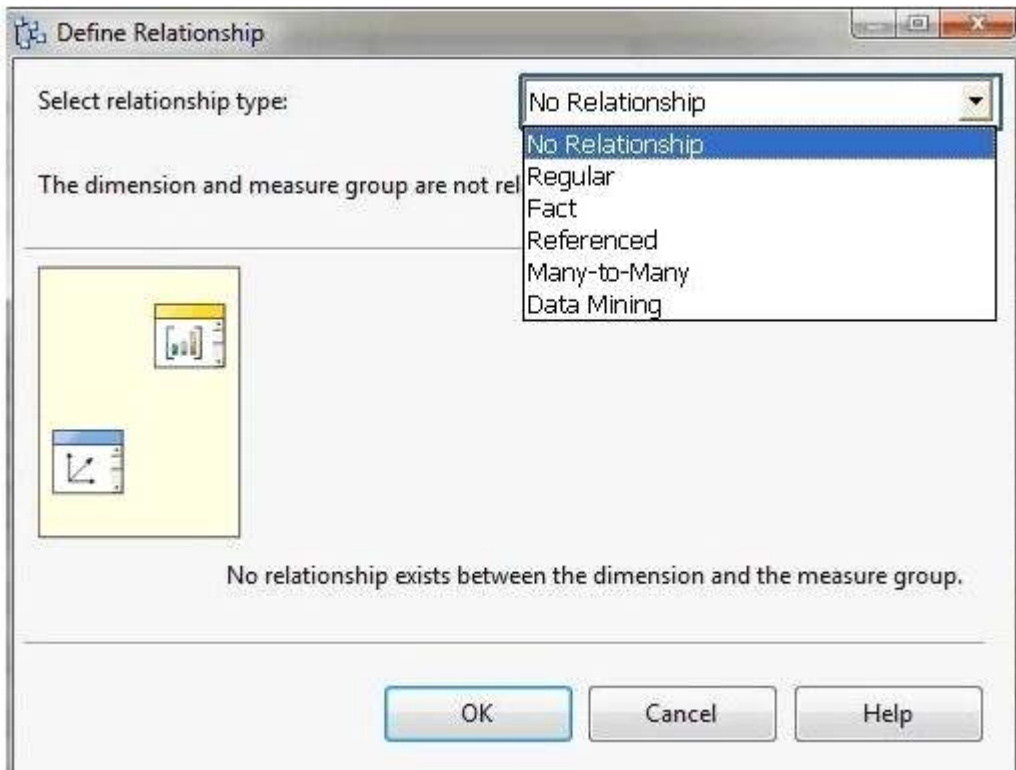
HOTSPOT

You are developing a SQL Server Analysis Services (SSAS) cube. A dimension named Category is based on the DimCategory table. A subset of the data source view is shown in the following graphic.



You need to relate the Category dimension to the Sales measure group. Which relationship type should you choose?

To answer, select the appropriate option from the drop-down list in the dialog box.



Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

Referenced Relationship

Link a dimension to a fact table indirectly through a dimension that is linked directly through a primary key to foreign key relationship.

QUESTION 20

You are designing a partitioning strategy for a large fact table in a Manufacturing data warehouse. Tens of millions of new inventory fact records are loaded into the data warehouse weekly, outside of business hours. Most queries against the database are generated by reports and by cube processing. Data is frequently queried at the day level and occasionally at the month level.

- A. Partition the inventory fact table by month, and compress each partition.
- B. Partition the inventory fact table by day, and compress each partition.
- C. Partition the inventory fact table by year.
- D. Partition the inventory fact table by week.

Correct Answer: B

Section: (none)

Explanation

QUESTION 21

DRAG DROP

You are administering a SQL Server Analysis Services (SSAS) database on a server. The database hosts a financial cube based on a SQL Azure database. You need to grant read access to the financial cube for all users in the group USA\PowerUsers. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Add the **USA\PowerUsers** group to the role. Set the cube access for the role to **Read**.

Add the **USA\PowerUsers** group to the role. Set the cube access for the role to **Select**.

Add the group **USA\PowerUsers** as a SQL Server login to the server.

In SQL Server Management Studio (SSMS), connect to the Database Engine instance on the server.

Add the **USA\PowerUsers** group to the role. Set the cube access for the role to **Read and Process**.

Create a new role for the database.

In SQL Server Management Studio (SSMS), connect to the SSAS instance on the server.

- A.
- B.
- C.
- D.

Correct Answer:

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Box 1: In the SQL Management Studio (SSMS), connect to the SSAS instance on the server.

Box 2: Create a new role for the database.

Box 3: Add the USA\PowerUsers group to the role. Set the cube access for the role to Read and Process.

Note:

* A member of the server role for Microsoft SQL Server Analysis Services, or a member of a database role that has Full Control (Administrator) permissions in a particular database, can create a database role that only has permission to process specified objects within the database. Giving a database role permission to process a database object lets an administrator delegate the task of processing certain objects, without also granting extraneous permissions to the user who is performing the processing.

* To give a database role permission to process a cube In SQL Server Management Studio, connect to the instance of Analysis Services, expand Roles for the appropriate database in Object Explorer, and then double-click a database role (or right-click Roles and select New Role to create a new database role). If this is a new role, make sure that you enter a name for the role in the Role namebox.

Click Cubes in the Select a Page pane, locate the cube in the Cube list, and then select the Process check box for the cube.

Click the OK button.

* There is no write permissions on a cube.

Reference: Grant Process Permissions on an Analysis Services Multidimensional Database

QUESTION 22

A SQL Server Analysis Services (SSAS) cube contains a large measure group. The fact table supporting the measure group is loaded with new data throughout the day.

You have the following requirements:

?Ensure that the cube displays current data as quickly as possible.

?Maximize availability of the cube.

?Maximize query performance for all aggregation levels. You need to choose a partitioning strategy that meets the requirements. Which partitioning strategy should you choose? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create one partition for the current day that uses multidimensional OLAP (MOLAP) with proactive caching as a storage mode.
- B. Create one partition for the current month that uses hybrid OLAP (HOLAP) as a storage mode.
- C. Create one partition for the current day that uses relational OLAP (ROLAP) as a storage mode.
- D. Create one partition for the current day that uses multidimensional OLAP (MOLAP) as a storage mode. Process the partition each night.

Correct Answer: A

Section: (none)

Explanation

QUESTION 23

You are designing an extract, transform, load (ETL) process for loading data from a SQL Azure database into a large fact table in a data warehouse each day with the prior day's sales data. The ETL process for the fact table must meet the following requirements:

?Load new data in the shortest possible time.

?Remove data that is more than 36 months old.

?Minimize record locking.

?Minimize impact on the transaction log.

You need to design an ETL process that meets the requirements. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Partition the fact table by date. Insert new data directly into the fact table and delete old data directly from the fact table.
- B. Partition the fact table by customer. Use partition switching both to remove old data and to load new data into each partition.
- C. Partition the fact table by date. Use partition switching and staging tables both to remove old data and to load new data.
- D. Partition the fact table by date. Use partition switching and a staging table to remove old data. Insert new data directly into the fact table.

Correct Answer: C

Section: (none)

Explanation

QUESTION 24

DRAG DROP

You are validating whether a SQL Server Integration Services (SSIS) package named Master.dtsx in the SSIS catalog is executing correctly. You need to display the number of rows in each buffer passed between each data flow component of the package. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Execute a SQL statement with a package name of Master.dtsx against the **catalog.executions** view and return its execution ID.

Run the Master.dtsx package with the logging level set to **Basic**.

Execute a SQL statement with the execution ID equal to the previously retrieved execution ID against the **catalog.execution_data_statistics** view and return the **rows_sent** column values for all the rows.

Run the Master.dtsx package with the logging level set to **Verbose**.

Execute a SQL statement with a package name of Master.dtsx against the **catalog.event_messages** view and return its execution ID.

A.

Execute a SQL statement with a package name of Master.dtsx against the **catalog.executions** view and return its execution ID.

Execute a SQL statement with the execution ID equal to the previously retrieved execution ID against the **catalog.execution_data_statistics** view and return the **rows_sent** column values for all the rows.

Run the Master.dtsx package with the logging level set to **Verbose**.

Execute a SQL statement with a package name of Master.dtsx against the **catalog.event_messages** view and return its execution ID.

Run the Master.dtsx package with the logging level set to **Basic**.

B.

C.

D.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* You are going to become very very familiar indeed with [catalog].[executions]. It is a view that provides a record of all package executions on the server and, most importantly, it contains [execution_id] ?the identifier for each execution and the field to which all other objects herein will be related.

QUESTION 25

DRAG DROP

You plan to deploy a SQL Server Integration Services (SSIS) project by using the project deployment model. You need to monitor control flow tasks to determine whether any of them are running longer than usual. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Write a query against the **catalog.execution_component_phases** view. Add a calculation to the query to compare durations to the **catalog.executions** view.

Write a query against the **catalog.execution_data_statistics** view. Add a calculation to the query to compare durations to the **catalog.executions** view.

Write a query against the **catalog.execution_data_statistics** view. Add a calculation to the query to compare durations to the **catalog.execution_data_taps** view.

Write a query against the **catalog.execution_component_phases** view. Add a calculation to the query to compare durations to the **catalog.execution_data_taps** view.

Execute the query.

Connect to the **SSISDB** database.

Connect to the **msdb** database.

A.

Write a query against the **catalog.execution_component_phases** view. Add a calculation to the query to compare durations to the **catalog.executions** view.

Execute the query.

Connect to the **msdb** database.

Write a query against the **catalog.execution_data_statistics** view. Add a calculation to the query to compare durations to the **catalog.executions** view.

Write a query against the **catalog.execution_data_statistics** view. Add a calculation to the query to compare durations to the **catalog.execution_data_taps** view.

Write a query against the **catalog.execution_component_phases** view. Add a calculation to the query to compare durations to the **catalog.execution_data_taps** view.

Connect to the **SSISDB** database.

B.

C.

D.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* execution_component_phases

Displays the time spent by a data flow component in each execution phase.

* The following example uses the catalog.execution_component_phases view to find the total amount of time that a specific package has spent executing in all phases (active_time), and the total elapsed time for the package (total_time).

use SSISDB

```
select package_name, task_name, subcomponent_name, execution_path, SUM(DATEDIFF
(ms,start_time,end_time)) as active_time, DATEDIFF(ms,min(start_time), max(end_time)) as total_time from
catalog.execution_component_phases
```

```
where execution_id = 1841
```

```
group by package_name, task_name, subcomponent_name, execution_path order by package_name,
task_name, subcomponent_name, execution_path
```

* catalog.executables

This view displays a row for each executable in the specified execution. An executable is a task or container that you add to the control flow of a package.

*(incorrect) catalog.executions (SSISDB Database)

Displays the instances of package execution in the Integration Services catalog. Packages that are executed with the Execute Package task run in the same instance of execution as the parent package.

This view displays a row for each instance of execution in the catalog.

*(incorrect) catalog.operation_messages

Displays messages that are logged during operations in the Integration Services catalog. This view displays a row for each message that is logged during an operation in the catalog. The message can be generated by the server, by the package execution process, or by the execution engine.

Reference: catalog.execution_component_phases

Reference: catalog.executables

QUESTION 26

DRAG DROP

You are designing a SQL Server Reporting Services (SSRS) solution. A report project must access multiple SQL Server databases. Each database is on a different instance. The databases have identical schema and security configurations. You have the following requirements:

?The report must support subscriptions.

?Users must be able to select the host when running the report.

What should you do?

To answer, drag the appropriate phrase or phrases from the list to the correct location or locations in the answer area. (Answer choices may be used once, more than once, or not all.)

a shared dataset.

stored credentials.

integrated security.

SQL Server data.

SQL Server instances.

data source in the report.

an expression-based connection string.

shared data source in the report

Create a

Create a report parameter that displays available values of

Create

Configure the data source to use

A.

a shared dataset.

integrated security.

SQL Server data.

shared data source in the report

Create a data source in the report.

Create a report parameter that displays available values of SQL Server instances.

Create an expression-based connection string.

Configure the data source to use stored credentials.

B.

C.

D.

Correct Answer: A
Section: (none)
Explanation

Explanation/Reference:

Explanation:

Note:

The report need a data source.

Through a report parameter the user can select among the available SQL Server instances. This selection is used through an expression-based connection string. Authentication is handled through stored credentials.

QUESTION 27

DRAG DROP

You are designing a SQL Server Reporting Services (SSRS) solution. An existing report aggregates data from a SQL Azure database in a chart. You need to use the chart in a new report and ensure that other users can use the chart in their reports. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

In Report Builder, insert the report part into a new report.

In Report Designer, open the report that contains the chart.

In Report Designer, insert the report part into a new report.

In Power View, open the report that contains the chart.

Select the chart for publication as a report part and publish the report.

A.

In Report Builder, insert the report part into a new report.

Select the chart for publication as a report part and publish the report.

In Report Designer, insert the report part into a new report.

In Power View, open the report that contains the chart.

In Report Designer, open the report that contains the chart.

B.

C.

D.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* In Report Designer, after you create tables, charts, and other report items in a project, you can publish them as report parts to a report server or SharePoint site integrated with a report server so that you and others can reuse them in other reports.

* . By using Report Builder, you can customize and update reports that were created in SQL Server Data Tools (SSDT) Report Designer.

* In Report Builder, IT pros and power users can create powerful operational reports, and reusable report parts and shared datasets.

Incorrect:

* (incorrect) Power View, a feature of SQL Server 2012 Reporting Services Add-in for Microsoft SharePoint Server 2010 Enterprise Edition, is an interactive data exploration, visualization, and presentation experience. It provides intuitive ad-hoc reporting for business users such as data analysts, business decision makers, and information workers. They can easily create and interact with views of data from data models based on PowerPivot workbooks published in a PowerPivot Gallery, or tabular models deployed to SQL Server 2012 Analysis Services (SSAS) instances. Power View is a browser-based Silverlight application launched from SharePoint Server 2010 that enables users to present and share insights with others in their organization through interactive presentations.

Reference: Getting Started with Report Builder

Reference: Report Parts in Report Designer (SSRS)

QUESTION 28

You are designing a multidimensional OLAP (MOLAP) cube. The MOLAP cube must meet the following requirements:

Ensure that workloads for aggregation tuning can be automatically collected. Require the least amount of effort to perform manual aggregation tuning. Minimize impact on the performance of previously tuned queries. You need to design a MOLAP cube that meets the requirements. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Enable SQL Server Analysis Services (SSAS) query logging. Run the Usage-Based Optimization Wizard to generate aggregations. Merge the wizard results with existing aggregation designs.
- B. Set up multiple partitions. Run the Aggregation Design Wizard periodically for each measure group. After the wizard finishes, discard the old aggregation design and accept the new one.
- C. Set up multiple partitions. Run the Aggregation Design Wizard on each partition. Schedule the aggregations by using an XMLA script in SQL Server Agent.
- D. Set the AggregationUsage property of all attributes based on natural keys to Full.

Correct Answer: A

Section: (none)

Explanation

QUESTION 29

You are designing a fact table in a SQL Server database. The fact table must meet the following requirements:

?Include a columnstore index.

?Allow users to choose up to 10 dimension tables and up to five facts at one time. ?Maximize performance of queries that aggregate measures by using any of the 10 dimensions.

?Support billions of rows.

?Use the most efficient design strategy.

You need to design the fact table to meet the requirements. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Design a fact table with 5 dimensional key columns and 10 measure columns. Place the columnstore index on the dimensional key columns.
- B. Design a fact table with 5 dimensional key columns and 10 measure columns. Place the columnstore index on the measure columns.
- C. Design a fact table with 10 dimensional key columns and 5 measure columns. Place the columnstore index on the dimensional key columns and the measure columns.
- D. Design a fact table with 10 dimensional key columns and 5 measure columns. Place the columnstore index on only the measure columns.

Correct Answer: C

Section: (none)

Explanation

QUESTION 30

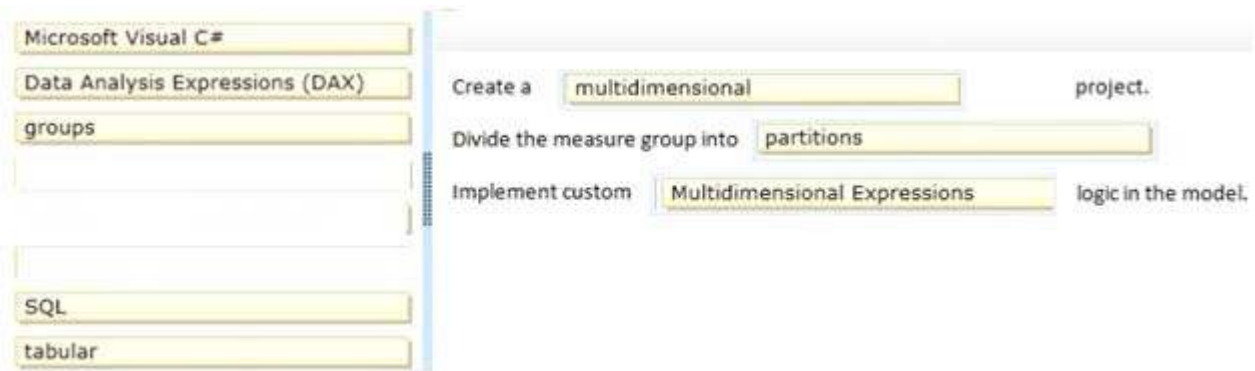
DRAG DROP

You are designing a SQL Server Analysis Services (SSAS) data model on a very large data warehouse. The fact tables in the data warehouse contain terabytes of data in tens of billions of rows. You must support the following features:

?Complex attribute/column relationships

?Advanced calculations in the data model definition ?Advanced calculations using logic deployed in a custom assembly You need to choose the correct SSAS design strategy. What should you do? To answer, drag the appropriate term or terms to the correct location or locations in the answer area. (Answer choices may be used once, more than once, or not all.)

A.



Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

Box 1:

The primary reason for building an Analysis Services multidimensional model is to achieve fast performance of ad hoc queries against business data. A multidimensional model is composed of cubes and dimensions that can be annotated and extended to support complex query constructions.

Box 2:

A partition is a container for a portion of the measure group data. Partitions are not seen from MDX queries; all queries reflect the whole content of the measure group, regardless of how many partitions are defined for the measure group. The data content of a partition is defined by the query bindings of the partition, and by the slicing expression.

Box 3:

Multidimensional Expressions (MDX) is the query language that you use to work with and retrieve multidimensional data in Microsoft SQL Server 2005 Analysis Services (SSAS).

QUESTION 31

You are creating a Multidimensional Expressions (MDX) calculation for Projected Revenue in a cube. For Product A, Projected Revenue is defined as 150 percent of the Total Sales of the product. For all other products, Projected Revenue is defined as 110 percent of the Total Sales of the product. You need to calculate the Projected Revenue as efficiently as possible. Which calculation should you use? (More than one answer choice may achieve the goal. Select the BEST answer.)

- ☐ A. `CREATE MEMBER CurrentCube.[Measures].[Projected Revenue]
AS CASE WHEN [Product].[Product Name].CurrentMember.Name = "Product
THEN [Measures].[Total Sales] * 1.5
ELSE [Measures].[Total Sales] * 1.1 END`
- ☐ B. `CREATE MEMBER CurrentCube.[Measures].[Projected Revenue]
AS [Measures].[Total Sales] * 1.1;
SCOPE ([Product].[Product Name].MEMBERS, [Measures].[Projected Revenue]
[Product].[Product Name].&[Product A] = [Measures].[Total Sales]
END SCOPE;`
- ☐ C. `CREATE MEMBER CurrentCube.[Measures].[Projected Revenue]
AS [Measures].[Total Sales] * 1.1;
SCOPE ([Product].[Product Name].&[Product A], [Measures].[Projected Revenue]
THIS = [Measures].[Total Sales] * 1.5;
END SCOPE;`
- ☐ D. `CREATE MEMBER CurrentCube.[Measures].[Projected Revenue]
AS [Measures].[Total Sales];
SCOPE ([Product].[Product Name].MEMBERS, [Measures].[Projected Revenue]
[Measures].[Total Sales] * 1.1;
IF [Product].[Product Name].CurrentMember.Name = "Product A"
THEN [Measures].[Total Sales] * 1.5
END IF;
END SCOPE;`

- A.
B.
C.
D.

Correct Answer: C


Section: (none)

Explanation

Explanation/Reference:

QUESTION 32

You are developing the database schema for a SQL Server Analysis Services (SSAS) BI Semantic Model (BISM). The BISM will be based on the schema displayed in the following graphic.



You have the following requirements:

?Ensure that queries of the data model correctly display average student age by class. ?Ensure that the solution supports role-based security and partitions.

?Minimize development effort.

You need to design the data model. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create a multidimensional project and define measures and a many-to-many dimensional relationship. Create partitions in SQL Server Management Studio (SSMS).
- B. Create a multidimensional project and define measures and a reference relationship. Create partitions in SQL Server Data Tools (SSDT).
- C. Create a tabular project and define measures. Create partitions in SQL Server Data Tools (SSDT).
- D. Create a tabular project and define calculated columns. Create partitions in SQL Server Management Studio (SSMS).

Correct Answer: A

Section: (none)

Explanation

QUESTION 33

You are defining a named set by using Multidimensional Expressions (MDX) in a sales cube. The cube includes a Customer dimension that contains a Geography hierarchy and a Gender attribute hierarchy. You need to return only the female customers in the Geography hierarchy. Which set should you use? (More than one answer choice may achieve the goal. Select the BEST answer.)

<http://www.lead2pass.com/70-467.html>

- ☐ A. Exists
- ```
(
 [Customer].[Customer Geography].[Customer Name].Members,
 [Customer].[Gender].[Female]
)
```
- ☐ B. Generate
- ```
(
    [Customer].[Gender].[Female],
    [Customer].[Model Name].[Model Name].Members, ALL
)
```
- ☐ C. Filter
- ```
(
 [Customer].[Customer Geography].[Customer Name].Members,
 ([Customer].[Gender].[Female], [Measures].[Sales Amount]) > 0
)
```
- ☐ D. CrossJoin
- ```
(
    [Customer].[Customer Geography].[Customer Name].Members,
    [Customer].[Gender].[Female]
)
```

- A.
B.
C.
D.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 34

HOTSPOT

You are designing a SQL Server Integration Services (SSIS) package configuration strategy. The package configuration must meet the following requirements:

- ?Include multiple properties in a configuration.
- ?Force packages to load all settings in the configuration.
- ?Support Encrypting File System (EFS) formats.

You need to select the appropriate configuration. Which configuration type should you use? To answer, select the appropriate option from the drop-down list in the dialog box.

Package Configuration Wizard

Select Configuration Type
You can select the type of configuration to create.

SSIS supports a variety of sources to use for setting the properties of objects.

Configuration type: Environment variable

☒ Specify configuration settings directly

Configuration file name:

☐ Configuration location is stored in an environment variable

Environment variable:

Environment variable dropdown menu:

- Environment variable
- XML configuration file
- Registry entry
- Parent package variable
- SQL Server

A.

Package Configuration Wizard

Select Configuration Type
You can select the type of configuration to create.

SSIS supports a variety of sources to use for setting the properties of objects.

Configuration type: Environment variable

☒ Specify configuration settings directly

Configuration file name:

☐ Configuration location is stored in an environment variable

Environment variable:

Environment variable dropdown menu:

- Environment variable
- XML configuration file
- Registry entry
- Parent package variable
- SQL Server

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

Package Configuration Types

The following table describes the package configuration types.

* SQL Server table

A table in a SQL Server database contains the configuration. The table can include multiple configurations.

* XML configuration file

An XML file contains the configurations. The XML file can include multiple configurations.

* Environment variable

An environment variable contains the configuration.

* Registry entry

A Registry entry contains the configuration.

* Parent package variable

A variable in the package contains the configuration. This configuration type is typically used to update properties in child packages.

Reference: Package Configurations

QUESTION 35

You are designing a SQL Server Integration Services (SSIS) solution that will load multiple Online Transactional Processing (OLTP) data sources into a SQL Server data mart. You have the following requirements:

?Ensure that the process supports the creation of an exception report that details possible duplicate key values, null ratios within columns, and column-length distributions of values. ?Ensure that users can generate the exception report in an XML format.

?Use the minimum development effort.

You need to design the SSIS solution to meet the requirements. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Use a Data Profiling task. Use a Data Flow task to extract the XML output of the Data Profiling task into a SQL Server table. Query the table to view the exceptions.
- B. Use Data Flow tasks to process the clean data.
- C. Use a Data Profiling task. Read the exceptions in Data Profile Viewer.
- D. Design a stored procedure that examines data for common dirty data patterns. Use an Execute SQL task.

Correct Answer: C

Section: (none)

Explanation

QUESTION 36

You are designing a SQL Server Integration Services (SSIS) solution. The solution will contain an SSIS project that includes several SSIS packages. Each SSIS package will define the same connection managers and variables. You have the following requirements:

?The deployment model must support changing the content of connection strings by using parameters at execution time.

?The deployment model must automatically log events to the SSISOB database.

?Maximize performance at execution time.<http://www.lead2pass.com/70-467.html>

You need to design a solution that meets the requirements. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Use a project deployment model. Modify connection manager properties to use project parameters.
- B. Use a package deployment model. Save each SSIS package to a file share that can be accessed from all environments.
- C. Use a package deployment model. Configure connections in an XML configuration file referenced by an environment variable that corresponds to the SQL Server environment of each SSIS package.
- D. Use a project deployment model. Modify connection manager properties to use package parameters.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Case Study: 1

Tailspin Toys

Background

You are the business intelligence (BI) solutions architect for Tailspin Toys. You produce solutions by using SQL Server 2012 Business Intelligence edition and Microsoft SharePoint Server 2010 Service Pack 1 (SP1) Enterprise edition.

Technical Background


Extract, transform, load (ETL) processes populate the data warehouse every 24 hours.

ETL Processes

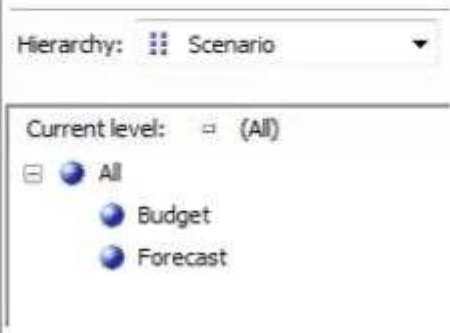
One SQL Server Integration Services (SSIS) package is designed and developed to populate each data warehouse table. The primary source of data is extracted from a SQL Azure database. Secondary data sources include a Microsoft Dynamics CRM 2011 on-premises database. ETL developers develop packages by using the SSIS project deployment model. The ETL developers are responsible for testing the packages and producing a deployment file. The deployment file is given to the ETL administrators. The ETL administrators belong to a Windows security group named SSISOwners that maps to a SQL Server login named SSISOwners.

Data Models

The IT department has developed and manages two SQL Server Analysis Services (SSAS) BI Semantic Model (BISM) projects: Sales Reporting and Sales Analysis. The Sales Reporting database has been developed as a tabular project. The Sales Analysis database has been developed as a multidimensional project. Business analysts use PowerPivot for Microsoft Excel to produce self-managed data models based directly on the data warehouse or the corporate data models, and publish the PowerPivot workbooks to a SharePoint site. The sole purpose of the Sales Reporting database is to support business user reporting and ad-hoc analysis by using Power View. The database is configured for DirectQuery mode and all model queries result in SSAS querying the data warehouse. The database is based on the entire data warehouse. The Sales Analysis database consists of a single SSAS cube named Sales. The Sales cube has been developed to support sales monitoring, analysts, and planning. The Sales cube metadata is shown in the following graphic.



Details of specific Sales cube dimensions are described in the following table.

Dimension	Hierarchies and levels	Additional information
Date	Calendar <ul style="list-style-type: none"> • Year • Quarter • Month • Date 	All attributes are hidden. The appropriate dimension and attribute Type properties have been configured.
Salesperson	Geography <ul style="list-style-type: none"> • Country • Region • Salesperson 	Based on the DimSalesperson and DimRegion tables. All attributes are hidden.
Scenario	Scenario (attribute hierarchy) <ul style="list-style-type: none"> • Scenario 	

The Sales measure group is based on the FactSales table. The Sales Plan measure group is based on the FactSalesPlan table. The Sales Plan measure group has been configured with a multidimensional OLAP (MOLAP) writeback partition. Both measure groups use MOLAP partitions, and aggregation designs are assigned to all partitions. Because the volumes of data in the data warehouse are large, an incremental processing strategy has been implemented. The Sales Variance calculated member is computed by subtracting the Sales Plan forecast amount from Sales. The Sales Variance % calculated member is computed by dividing Sales Variance by Sales. The cube's Multidimensional Expressions (MDX) script does not set any color properties.

Analysis and Reporting

SQL Server Reporting Services (SSRS) has been configured in SharePoint integrated mode. A business analyst has created a PowerPivot workbook named Manufacturing Performance that integrates data from the data warehouse and manufacturing data from an operational database hosted in SQL Azure. The workbook has been published in a PowerPivot Gallery library in SharePoint Server and does not contain any reports. The analyst has scheduled daily data refresh from the SQL Azure database. Several SSRS reports are based on the PowerPivot workbook, and all reports are configured with a report execution mode to run on demand. Recently users have noticed that data in the PowerPivot workbooks published to SharePoint Server is not being refreshed. The SharePoint administrator has identified that the Secure Store Service target application used by the PowerPivot unattended data refresh account has been deleted.

Business Requirements

ETL Processes

All ETL administrators must have full privileges to administer and monitor the SSIS catalog, and to import and manage projects.

Data Models

The budget and forecast values must never be accumulated when querying the Sales cube. Queries should return the forecast sales values by default. Business users have requested that a single field named SalespersonName be made available to report the full name of the salesperson in the Sales Reporting data model. Writeback is used to initialize the budget sales values for a future year and is based on a weighted

allocation of the sales achieved in the previous year.

Analysis and Reporting

Reports based on the Manufacturing Performance PowerPivot workbook must deliver data that is no more than one hour old. Management has requested a new report named Regional Sales. This report must be based on the Sales cube and must allow users to filter by a specific year and present a grid with every region on the columns and the Products hierarchy on the rows. The hierarchy must initially be collapsed and allow the user to drill down through the hierarchy to analyze sales. Additionally, sales values that are less than \$5000 must be highlighted in red.

Technical Requirements

Data Warehouse

Business logic in the form of calculations should be defined in the data warehouse to ensure consistency and availability to all data modeling experiences. The schema design should remain as denormalized as possible and should not include unnecessary columns. The schema design must be extended to include the product dimension data.

ETL Processes

Package executions must log only data flow component phases and errors.


Data Models

Processing time for all data models must be minimized. A key performance indicator (KPI) must be added to the Sales cube to monitor sales performance. The KPI trend must use the Standard Arrow indicator to display improving, static, or deteriorating Sales Variance % values compared to the previous time period.

Analysis and Reporting

IT developers must create a library of SSRS reports based on the Sales Reporting database. A shared SSRS data source named Sales Reporting must be created in a SharePoint data connections library.

Data Warehouse Schema



QUESTION 37

You need to fix the PowerPivot data refresh problem by using the least amount of administrative effort. What should you do?

- A. Use the PowerPivot Configuration Tool and select the Upgrade Features, Services, Applications and Solutions option.
- B. Use the PowerPivot Configuration Tool and select the Configure or Repair PowerPivot for SharePoint option.
- C. Reinstall SSAS in PowerPivot for SharePoint mode by using the SQL Server 2012 installation media.
- C. In SharePoint Central Administration, create a target application and configure the PowerPivot service application settings to use the target application.

Correct Answer: B

Section: (none)

Explanation

QUESTION 38

You need to configure the Scenario attribute to ensure that business users appropriately query the Sales Plan measure. What should you do? (Each correct answer presents part of the solution. Choose all that apply.)

- A. Set the AttributeHierarchyVisible property to False.
- B. Set the IsAggregatable property to False.
- C. Set the Usage property to Parent.
- D. set the DefaultMember property to the Forecast member.
- E. Set the AttributeHierarchyEnabled property to False.
- F. Set the RootMemberIf property to ParentIsMissing.

Correct Answer: CD

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* From scenario:

The Sales measure group is based on the FactSales table. The Sales Plan measure group is based on the FactSalesPlan table. The Sales Plan measure group has been configured with a multidimensional OLAP (MOLAP) writeback partition. Both measure groups use MOLAP partitions, and aggregation designs are assigned to all partitions.

QUESTION 39

DRAG DROP

You need to configure the attribute relationship types for the Salesperson dimension. Which configuration should you use?

To answer, drag the appropriate pair of attributes and attribute relationships from the list to the correct location or locations in the answer area. (Answer choices may be used once, more than once, or not all.)



- B.
- C.
- D.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:


Note:

* You connect a "higher-level" attribute to a "lower-level" attribute.

* Best practice design says relationships should be rigid (bold filled lines) when members aren't shifting around.

*

Example:



QUESTION 40

DRAG DROP

You need to extend the schema design to store the product dimension data. Which design should you use? To answer, drag the appropriate table or tables to the correct location or locations in the answer area. (Fill from left to right. Answer choices may be used once, more than once, or not all.)

DimCategory

CategoryKey
CategoryName

DimSubcategory

SubcategoryKey
SubcategoryName
CategoryName

DimProduct

ProductKey
ProductSKU
ProductName
Size
SubcategoryKey
SubcategoryName
CategoryName

DimProduct

ProductKey
ProductSKU
ProductName
Size
SubcategoryKey
SubcategoryName
CategoryName
StartDate
EndDate

DimSubcategory

SubcategoryKey
SubcategoryName
CategoryKey

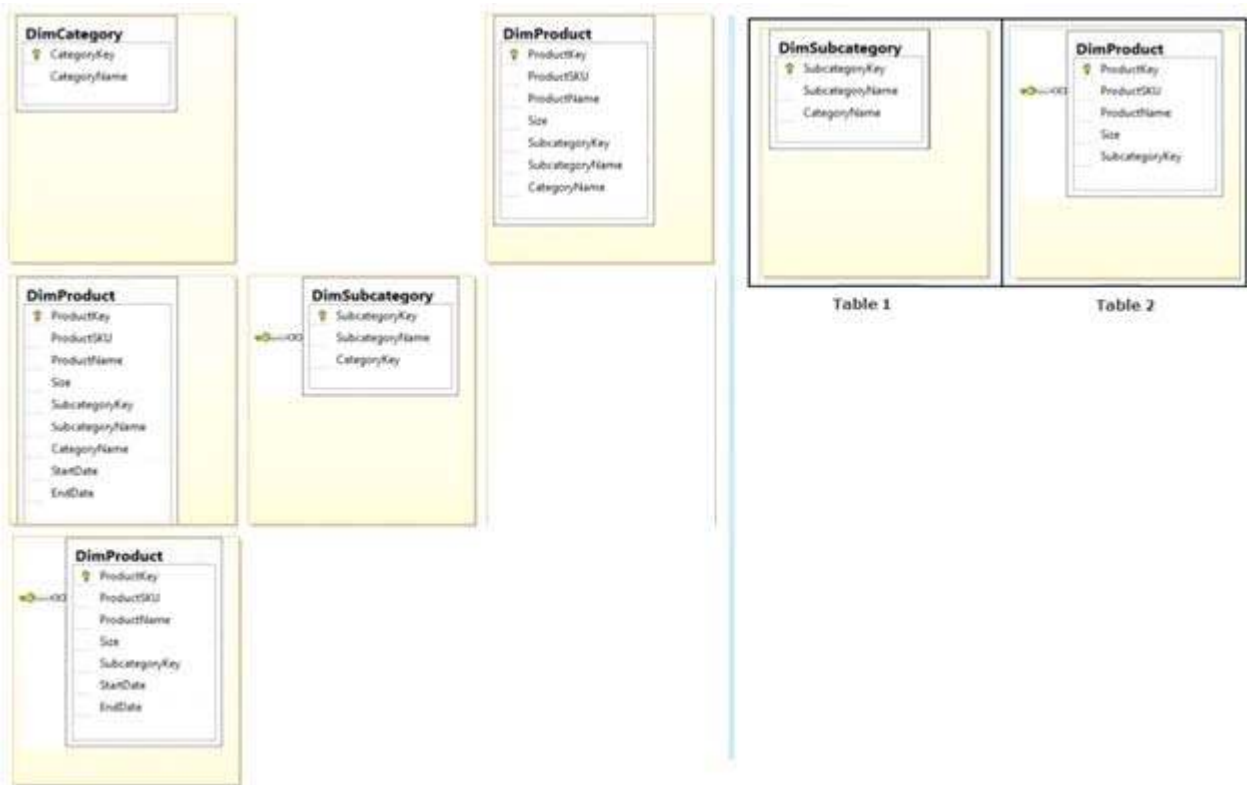
DimProduct

ProductKey
ProductSKU
ProductName
Size
SubcategoryKey

DimProduct

ProductKey
ProductSKU
ProductName
Size
SubcategoryKey
StartDate
EndDate

A.



B.

C.

D.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Explanation:

* From scenario:

/ The schema design must be extended to include the product dimension data.

* DimProduct table need to contain a foreign key to the DimSubCategory table. No further SubCategory data should be stored in the DimProduct table.

* No time related columns (StartDate, EndDate) should be included in the DimProduct table.

Case Study: 2

Contoso, Ltd

General Background

You are the SQL Server Administrator for Contoso, Ltd. You have been tasked with upgrading all existing SQL Server instances to SQL Server 2012.

Technical Background

The corporate environment includes an Active Directory Domain Services (AD DS) domain named contoso.com. The forest and domain levels are set to Windows Server 2008. All default containers are used for computer and user accounts. All servers run Windows Server 2008 R2 Service Pack 1 (SP1). All client computers run Windows 7 Professional SP1. All servers and client computers are members of the contoso.com domain. The current SQL Server environment consists of a single instance failover cluster of SQL Server 2008 R2 Analysis Services (SSAS). The virtual server name of the cluster is SSASCluster. The cluster includes two nodes: Node1 and Node2. Node1 is currently the active node. In anticipation of the upgrade, the prerequisites and shared components have been upgraded on both nodes of the cluster, and each node was rebooted during a weekly maintenance window. A single- server deployment of SQL Server 2008 R2 Reporting Services (SSRS) in native mode is installed on a server named SSRS01. The Reporting Server service is configured to

use a domain service account. SSRS01 hosts reports that access the SSAS databases for sales data as well as modeling data for the Research team. SSRS01 contains 94 reports used by the organization. These reports are generated continually during business hours. Users report that report subscriptions on SSRS01 are not being delivered. You run the reports on demand from Report Manager and find that the reports render as expected. A new server named SSRS02 has been joined to the domain, SSRS02 will host a single-server deployment of SSRS so that snapshots of critical reports are accessible during the upgrade. The server configuration is shown in the exhibit. (Click the Exhibit button.) The production system includes three SSAS databases that are described in the following table.

Database name	Size
Customer Sales	350 MB
Manufacturing	1.2 GB
Research	620 MB

All SSAS databases are backed up once a day, and backups are stored offsite.

Business Requirements

After the upgrade users must be able to perform the following tasks:

- ?Ad-hoc analysis of data in the SSAS databases by using the Microsoft Excel PivotTable client.
- ?Daily operational analysis by executing a custom application that uses ADOMD.NET and existing Multidimensional Expressions (MDX) queries. The detailed data must be stored in the model.

Technical Requirements

You need to minimize downtime during the SSASCluster upgrade. The upgrade must minimize user intervention and administrative effort. The upgrade to SQL Server 2012 must maximize the use of all existing servers, require the least amount of administrative effort, and ensure that the SSAS databases are operational as soon as possible. You must implement the highest level of domain security for client computers connecting to SSRS01. The SSRS instance on SSRS01 must use Kerberos delegation to connect to the SSAS databases. Email notification for SSRS01 has not been previously configured. Email notification must be configured to use the SMTP server SMTP01 with a From address of reports@contoso.com. Report distribution must be secured by using SSL and must be limited to the contoso.com domain. You have the following requirements for SSRS02:

- ?Replicate the SSRS01 configuration.
- ?Ensure that all current reports are available on SSRS02.
- ?Minimize the performance impact on SSRS01.

In preparation for the upgrade, the SSRS-related components have been installed on the new SSRS02 server by using the Reporting Services file-only installation mode. The Reporting Services databases have been restored from SSRS01 and configured appropriately. You must design a strategy to recover the SSRS instance on SSRS01 in the event of a system failure. The strategy must ensure that SSRS can be recovered in the minimal amount of time and that reports are available as soon as possible. Only functional components must be recovered.


SSRS02 is the recovery server and is running the same version of SSRS as SSRS01. A full backup of the SSRS databases on SSRS01 is performed nightly. The report server configuration files, custom assemblies, and extensions on SSRS02 are manually synchronized with SSRS01.

Prior to implementing the upgrade to SQL Server 2012, you must back up all existing SSAS databases. Databases on SSRS01 is performed nightly. The report server configuration files, custom assemblies, and extensions on SSRS02 are manually synchronized with SSRS01. Prior to implementing the upgrade to SQL Server 2012, you must back up all existing SSAS databases. The backup must include only the partitioning, metadata, and aggregations to minimize the processing time required when restoring the databases. You must minimize processing time and the amount of disk space used by the backups. Before upgrading SSAS on the SSASCluster, all existing databases must be moved to a temporary staging server named SSAS01 that hosts a default instance of SQL Server 2012 Analysis Services. This server will be used for testing client applications connecting to SSAS 2012, and as a disaster recovery platform during the upgrade. You must move the databases by using the least amount of administrative effort and minimize downtime. All SSAS databases other than the Research database must be converted to tabular BI Semantic Models (BISMs) as part of the upgrade to SSAS 2012. The Research team must have access to the Research database for modeling throughout the upgrade. To facilitate this, you detach the Research database and attach it to SSAS01. While testing the

Research database on SSAS01, you increase the compatibility level to 1100. You then discover a compatibility issue with the application. You must roll back the compatibility level of the database to 1050 and retest. After completing the upgrade, you must do the following:

1. Design a role and assign an MDX expression to the Allowed member set property of the Customer dimension to allow sales representatives to browse only members of the Customer dimension that are located in their sales regions. Use the sales representatives' logins and minimize impact on performance.
2. Deploy a data model to allow the ad-hoc analysis of data. The data model must be cached and source data from an OData feed.

Server Configuration



QUESTION 41

You need to configure security for the SSRS instance on SSRS01 to connect to SSAS and minimize downtime. What should you do? (Each correct answer presents part of the solution. Choose all that apply.)

- A. Register a service principal name for the Report Server service.
- B. Register a service principal name for the Analysis Services service.
- C. Restart the IIS service.
- D. Configure SSRS01 to use the Negotiate authentication type.
- E. Configure SSRS01 to use the Custom authentication type.

Correct Answer: AD

Section: (none)

Explanation

Explanation/Reference:

Explanation:

A (not B): If you are deploying Reporting Services in a network that uses the Kerberos protocol for mutual authentication, you must create a Service Principal Name (SPN) for the Report Server service if you configure it to run as a domain user account.

D (not E):

* See step 6 below.

To register an SPN for a Report Server service running as a domain user ?Install Reporting Services and configure the Report Server service to run as a domain user account. Note that users will not be able to connect to the report server until you complete the following steps.

?Log on to the domain controller as domain administrator.

?Open a Command Prompt window.

?Copy the following command, replacing placeholder values with actual values that are valid for your network:

?Setspn -a http/<computer-name>.<domain-name>:<port> <domain-user-account> ?Run the command.

?Open the RsReportServer.config file and locate the <AuthenticationTypes> section. Add

<RSWindowsNegotiate/> as the first entry in this section to enable NTLM.

* RSWindowsNegotiate. If you initially set the Windows service account for the report server to NetworkService or LocalSystem in Reporting Services Configuration Manager, RSWindowsNegotiate is added to the RsReportServer.config file as the default setting. With this setting, the report server can accept requests from client applications requesting Kerberos or NTLM authentication. If Kerberos is requested and the authentication fails, the report server switches to NTLM authentication and prompts the user for credentials unless the network is configured to manage authentication transparently.

Using RSWindowsNegotiate is your best option because it provides the greatest flexibility for multiple clients in an intranet environment.

Not C: IIS is not mention in this scenario.

Note:

* From scenario:

/ A single-server deployment of SQL Server 2008 R2 Reporting Services (SSRS) in native mode is installed on a server named SSRS01. The Reporting Server service is configured to use a domain service account.

Reference: Register a Service Principal Name (SPN) for a Report Server

QUESTION 42

You need to perform the pre-upgrade database backup operation by using SQL Server Management Studio (SSMS). How should you configure the backup options?

- A. Select the Apply compression check box. Select the Encrypt backup file check box and supply a password.
- B. Clear the Apply compression check box. Select the Encrypt backup file check box and supply a password.
- C. Clear the Apply compression check box. Clear the Encrypt backup file check box.
- D. Select the Apply compression check box. Clear the Encrypt backup file check box.

Correct Answer: A

Section: (none)

Explanation

QUESTION 43

You need to roll back the compatibility level of the Research database. What should you do?

- A. Restore a backup of the previous version of the database.<http://www.lead2pass.com/70-467.html>
- B. Use an ALTER DATABASE statement to set the compatibility option.
- C. Change the CompatibilityLevel property in the XMLA script, and then execute the script.
- D. In SQL Server Management Studio (SSMS), change the compatibility level in the database properties.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 44

You need to develop a BISM that meets the business requirements for ad-hoc and daily operational analysis. You must minimize development effort. Which development approach and mode should you use?

- A. Develop a tabular project and configure the model with the DirectQuery mode setting on and the project query mode set to DirectQuery.
- B. Develop a tabular project and configure the model with the DirectQuery mode setting on and the project query mode set to In-Memory with DirectQuery.
- C. Develop a multidimensional project and configure the model with the DirectQuery mode setting off.
- D. Develop a multidimensional project and configure the cube to use hybrid OLAP (HOLAP) storage mode.

Correct Answer: B

Section: (none)

Explanation

QUESTION 45

You need to configure SSRS to send the required notification messages. Which configuration settings should you use? (Each correct answer presents a partial solution. Choose all that apply.)

- A. <SendUsing>2</SendUsing>http://www.lead2pass.com/70-467.html
- B. <SendUsing>contoso.com</SendUsing>
- C. <SMTPServer>SMTP01/SMTPServer>
- D. <SMTPServerPort>110</SMTPServerPort>
- E. <SMTPServer>SSRS01/SMTPServer>
- F. <From>reports@contoso.com</From>
- G. <PermittedHosts>contoso.com</PermittedHosts>

Correct Answer: ACF

Section: (none)

Explanation

Explanation/Reference:

Explanation:

A:

* In the configuration file, the delivery method is set through the SendUsing configuration setting.

* SendUsing specifies a method for sending messages. You can choose between a network SMTP service or a local SMTP service pickup directory. To use a remote SMTP service, this value must be set to 2 in the RSReportServer.config file.

C, F:

* From scenario: Email notification for SSRS01 has not been previously configured. Email notification must be configured to use the SMTP server SMTP01 with a From address of reports@contoso.com.

* SMTPServer specifies the remote SMTP server or forwarder. This value is a required value if you are using a remote SMTP server or forwarder.

Note:

Configuration Options for Remote SMTP Service

The connection between the report server and an SMTP server or forwarder is determined by the following configuration settings:

* SendUsing specifies a method for sending messages. You can choose between a network SMTP service or a local SMTP service pickup directory. To use a remote SMTP service, this value must be set to 2 in the RSReportServer.config file.

* SMTPServer specifies the remote SMTP server or forwarder. This value is a required value if you are using a remote SMTP server or forwarder.

* From sets the value that appears in the From: line of an e-mail message. This value is a required value if you are using a remote SMTP server or forwarder. Other values that are used for remote SMTP service include the

following (note that you do not need to specify these values unless you want to override the default values).

* SMTPServerPort is configured for port 25.

* SMTPAuthenticate specifies how the report server connects to the remote SMTP server. Reference: Configure a Report Server for E-Mail Delivery (Reporting Services), Configuration Options for Remote SMTP Service

QUESTION 46

You need to create the Package Activity report. What should you do?

- A. Create a log table and use SSIS event handlers to write to the log table. Then create an SSRS report that uses the log table.
- B. use the SSIS log provider for SQL Server. Then create an SSRS report that uses the sysssislog table.
- C. Create a log table and build a custom log provider to write to the log table. Then create an SSRS report that uses the log table.
- D. Create an SSRS report that uses the catalog.executions and catalog.execution_data_statistics views.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

<http://www.lead2pass.com/70-467.html>

QUESTION 47

You need to design a cube partitioning strategy to be implemented as the cube size increases. What should you do?

- A. Use relational OLAP (ROLAP) on all local partitions.
- B. Implement monthly remote partitions.
- C. Use multidimensional OLAP (MOLAP) on all local partitions.
- D. Implement monthly local partitions.

Correct Answer: B

Section: (none)

Explanation

QUESTION 48

You need to choose the appropriate key to use when designing a dimension table based on the Customer table. What should you do?

- A. Use a surrogate key.
- B. Use a natural key.<http://www.lead2pass.com/70-467.html>
- C. Use the CustomerNumber column as the key.
D Concatenate the CustomerName and CustomerNumber columns and use the concatenated string as the key.
- D. Use the CustomerName column as the key.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 49

You need to scale out SSAS. What should you do?<http://www.lead2pass.com/70-467.html>

- A. Back up the cube on ServerB and restore it on ServerC each day.
- B. Create an empty cube on ServerC and link to the objects in the cube on ServerB.
- C. Process the cube on both ServerB and ServerC each day.
- D. Synchronize the cube from ServerB to ServerC each day.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 50

You need to implement the aggregation designs for the cube. What should you do?

- A. Use the Usage-Based Optimization Wizard.
- B. Use the Aggregation Design Wizard.
- C. Partition the cube by month.
- D. Implement cache warming in SSAS via an SSIS package.

Correct Answer: B

Section: (none)

Explanation