Using Automatic Index Management

Erin Stellato
PRINCIPAL CONSULTANT

@erinstellato   www.sqlskills.com/blogs/erin

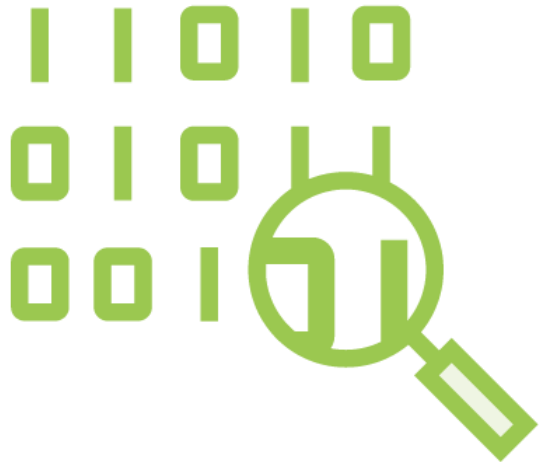# Module Overview

Automatic Index Management

Underlying data

Using the data

Creating indexes can help queries execute more efficiently

- Nonclustered indexes are a subset of columns from a table

Creating the right indexes for a table requires knowing your workload

SQL Server: Indexing for Performance

- http://bit.ly/2BCQFpX

Automatic Index Management

Azure SQL Database

Uses information from DMVs and an internal model
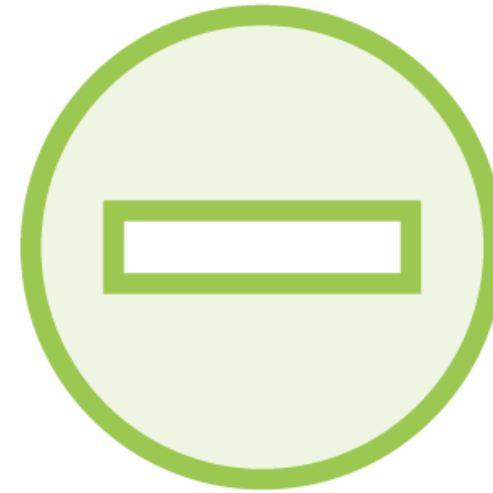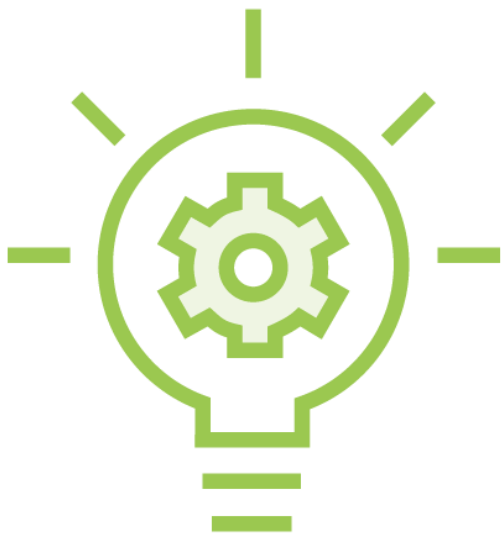
Disabled by default

Managing Indexes

Uses missing index data and an internal model to create new indexes

Relies on object access data to remove unused indexes

Indexes created and removed during times of non-peak workload

Verification of action (addition/removal) occurs and action can be reverted

Index recommendations are provided if you prefer manual intervention

Demo

Using Automatic Index Management

What We Covered

Automatic Index Management

Underlying data

Using the data

Next steps

Get the Latest Release

Download SQL Server 2017
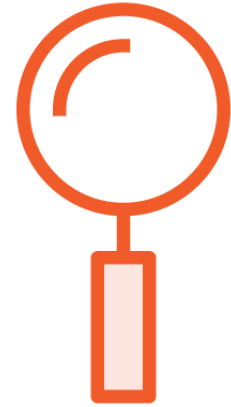
Azure SQL Database

Start Using New Features

Wait statistics collection enabled by default

Enable Automatic Plan Correction

Enable Automatic Index Management

Course Summary

Wait statistics information now captured in Query Store

Automatic Plan Correction forces plans to stabilize query performance

Automatic Index Management creates and drops indexes to improve query performance