

SQL Server: Detecting and Correcting Database Corruption

Module 5: DBCC CHECKDB and Related Commands

Paul S. Randal

<http://www.SQLskills.com/blogs/paul/>
Paul@SQLskills.com



pluralsight
hardcore developer training

Introduction

- Consistency checking is done using the DBCC CHECKDB command or by using some of the DBCC commands that run a subset of the DBCC CHECKDB functionality
- SQL Server Books Online for DBCC in general is at <http://bit.ly/c2pV9l>
- In this module we'll cover:
 - DBCC CHECKDB overview
 - Other consistency checking commands
 - Consistency checks that DBCC CHECKDB performs
 - DBCC CHECK* options and behaviors
 - Potential problems
 - How to break up consistency checks over time

DBCC CHECKDB (1)

- **This is the main consistency checking command**
- **Runs all consistency checks against a single database**
- **Very simple to execute:**
 - DBCC CHECKDB (N'database')
 - Or change context to the desired database and just do DBCC CHECKDB
- **The only way to read all allocated pages in the database**
 - Use to force page checksums to be checked
- **Possible to change some of the behavior using a variety of options**
 - We'll cover these later
- **Very resource intensive with CPU, memory, and I/O**
- **Will use multiple threads on Enterprise Edition**
 - Parallelism can be disabled using documented trace flag 2528

DBCC CHECKDB (2)

- **Many algorithms have been added to minimize runtime and run online without blocking locks since SQL Server 2000**
 - Compared with 6.5 or 7.0, which used locking all the time
 - Performance increases from version to version
- **Features worth mentioning from SQL Server 2005 onward:**
 - Progress reporting
 - Data purity checks
 - “Last known good” stored in the database boot page
 - No false failures like in SQL Server 2000
- **My comprehensive blog post series is at <http://bit.ly/TFvxmG>**
- **Detailed information (~90 pages) in the *SQL Server 2008 Internals* and *SQL Server 2012 Internals* books**

Other DBCC CHECK* Commands (1)

■ DBCC CHECKALLOC

- Checks consistency of allocation structures in a database
 - After first checking the consistency of critical system tables
- Examples:
 - Ensure allocation structures in the database are valid
 - Ensure no two tables have the same data file pages allocated
- Functionality is part of DBCC CHECKDB

■ DBCC CHECKTABLE

- Checks consistency of a table and its indexes
- Examples:
 - Ensure each data file page in the table and its indexes can be read and has a valid structure
 - Ensure each nonclustered index row has a matching table row
 - Ensure order of keys in indexes are correct
 - Ensure that any FILESTREAM data is correctly linked to the table
- Functionality is part of DBCC CHECKDB

Other DBCC CHECK* Commands (2)

- **DBCC CHECKCATALOG**

- Check relationships between system catalogs
 - E.g. column metadata must exist for all columns listed in a table's metadata, and vice-versa
- Functionality is part of DBCC CHECKDB

- **DBCC CHECKFILEGROUP**

- Perform consistency checks for a filegroup
- Runs allocation checks for the database, then logical consistency checks for table and index partitions in the specified filegroup
- Similar functionality to DBCC CHECKDB

Other DBCC CHECK* Commands (3)

- **DBCC CHECKIDENT**

- Checks and can reset the identity value for a table
- Functionality is not part of DBCC CHECKDB

- **DBCC CHECKCONSTRAINTS**

- Checks FOREIGN KEY and CHECK constraints for validity
- Functionality is not part of DBCC CHECKDB

What Exactly Does DBCC CHECKDB Do? (1)

- **Primitive checks of critical system tables**
 - If any problems are found, the only option is to restore from backups or try to extract data manually
- **Allocation checks**
- **Logical checks of critical system tables**
- **Logical checks of all other tables**
- **Metadata checks**
- **Service Broker data validation**
- **Indexed view, XML index, spatial index checks**
 - Only when the WITH EXTENDED_LOGICAL_CHECKS option is used, from SQL Server 2008 onward
- **If a repair option is specified, repairs are done at each stage if necessary and possible**

What Exactly Does DBCC CHECKDB Do? (2)

- **Another way to think of it is that DBCC CHECKDB does:**
 - Primitive system table checks, then
 - DBCC CHECKALLOC, then
 - DBCC CHECKTABLE of system tables, then
 - DBCC CHECKTABLE of user tables, then
 - DBCC CHECKCATALOG, then
 - Service Broker checks and other optional checks

- **It does not do:**
 - DBCC CHECKIDENT
 - DBCC CHECKCONSTRAINTS

Does DBCC CHECKDB Run During Startup?

- **When SQL Server starts, you may see messages in the error log about DBCC CHECKDB**
 - 2013-08-18 12:02:34.23 spid4s CHECKDB for database 'master' finished without errors on 2013-08-16 03:10:47.163 (local time). This is an informational message only; no user action is required.
- **This does not mean that DBCC CHECKDB was performed during instance startup**
- **It is reporting the “last known good” time for databases on the instance**
- **This is the time that DBCC CHECKDB last finished without finding any corruptions**
 - Stored in the boot page of the database (page 9 in file 1)
 - Not propagated to availability group replicas
 - Can be viewed using DBCC DBINFO (covered in the advanced course)

DBCC CHECK* Options (1)

- **NO_INFOMSGS**
 - Skip printing informational messages
 - Recommended to always use
- **ALL_ERRORMSGs**
 - Print all error messages
 - Not required from SQL Server 2008 SP1 onward
- **NOINDEX**
 - Skip checking nonclustered indexes
 - Not recommended
- **DATA_PURITY**
 - Perform column data validation
 - Default for all databases created on SQL Server 2005 or higher
 - Except master and model (it's a bug, see my blog at <http://bit.ly/17mAGYQ>)
 - Used only after upgrading from SQL Server 2000

DBCC CHECK* Options (2)

- **ESTIMATEONLY**

- Estimate how much *tempdb* space is required
 - This functionality is broken in current builds of SQL Server 2008 R2 and 2012

- **TABLOCK**

- Use locks instead of a database snapshot
- Not permitted on the *master* database
- Default for *tempdb*, where it causes allocation checks to be skipped

- **EXTENDED_LOGICAL_CHECKS**

- Perform extra validation of some data structures
 - Indexed views, XML indexes, spatial indexes

- **PHYSICAL_ONLY**

- Skips the majority of logical consistency checks
- Reduces the time necessary for completion
- Reduces CPU and memory resources necessary

- **Repair options will be discussed in Module 8**

How Are DBCC CHECK* Commands Online?

- SQL Server 2000 used transaction log analysis, which was slow
- Since SQL Server 2005, DBCC commands use a database snapshot, which is a transactionally-consistent, point-in-time view of the database
 - Not required if the database is read-only, single-user, or the target database is itself a database snapshot
- **The database snapshot is “hidden”**
 - Created as NTFS alternate streams of the existing data files
- **Sometimes problems can occur:**
 - The database snapshot runs out of space
 - No permissions to create the database snapshot
 - 3rd-party NTFS filter drivers that do not support alternate streams correctly
- **Solution: create your own database snapshot and check that**

How Long Will DBCC CHECKDB Take to Run?

- **Depends on many factors:**
 - Size of the database
 - Concurrent I/O load on the server
 - Concurrent CPU activity on the server
 - Concurrent update activity on the database
 - Throughput capabilities of the I/O subsystem
 - Number of CPUs on the server
 - Speed of the disks where *tempdb* is placed
 - Complexity of the database schema
 - Which options were specified
 - Number and type of corruptions that exist
- **See my blog post for more details at <http://bit.ly/RRL570>**

Potential Problems with DBCC CHECKDB (1)

- **Hidden database snapshots use NTFS alternate streams on existing database files**
 - Some 3rd-party software uses kernel filter drivers that do not cope with these correctly and return garbage data so DBCC CHECKDB reports corruption
- **The initial phase where the database snapshot is created cannot be interrupted due to limitations in the Engine**
 - This can make killing a DBCC CHECKDB impossible, and the SPID to appear as if it is in rollback
- **Database snapshots in general have an issue where they can run into NTFS limitations on size due to file-system fragmentation**
 - See <http://support.microsoft.com/kb/2002606> for more details

Potential Problems with DBCC CHECKDB (2)

- **Beware of unexpected very long run times**
 - Most likely a corruption has triggered an algorithm to look deeper into the database structure to determine exactly where the problem is
- **Consistency checks make use of *tempdb* for temporary data storage, so *tempdb* must be sized appropriately**
 - Use WITH ESTIMATEONLY to get an estimate of how much space is required
 - Beware that in some current builds that functionality is broken

Breaking Up Consistency Checks Over Time

- **Several options for doing the equivalent of DBCC CHECKDB if it's not possible to restore to another server and run DBCC CHECKDB there**
 - E.g. no other system to use or the database is too large
- **Option 1:**
 - Use DBCC CHECKFILEGROUP to perform per-filegroup checks of all filegroups over a number of days, plus DBCC CHECKCATALOG
 - This is another benefit of using multiple filegroups when your database size becomes very large (e.g. more than 100 GB)
- **Option 2:**
 - Use DBCC CHECKTABLE to perform consistency checks of all tables over a number of days, plus DBCC CHECKALLOC and DBCC CHECKCATALOG
 - Can be based on time, table size, or some other factor

Summary

- Most people just use DBCC CHECKDB but it is useful to understand the different consistency checking commands and their options
- Also useful to know what DBCC CHECKDB is doing and some of the problems that can occur with how it is implemented
- In the next module, we'll discuss:
 - Interpreting the output of DBCC CHECKDB
 - Simple examples of corruption