

SQL Server: Deadlock Analysis and Prevention

Module 2: Locking Overview

Jonathan M. Kehayias
Jonathan@SQLskills.com



Introduction


- **This module provides a brief overview of locking in SQL Server to provide the necessary background to troubleshoot deadlocks**
 - Full details are beyond the scope of this course, but will be covered in the upcoming course on SQL Server Transactions, Locking, and Isolation Levels
- **Locks are used by the Database Engine to synchronize access to a resource by multiple concurrent transactions**
- **Understanding how locking occurs in SQL Server is critical for being able to properly analyze and prevent deadlocks**
- **In this module we'll cover:**
 - Lock resource hierarchy
 - Lock modes
 - Lock compatibility
 - Isolation levels and locking hints
 - Lock escalation

Terminology

- **Transaction** – a unit of work performed within the database
- **Lock** - the synchronization mechanism on a resource that protects changes amongst multiple concurrent transactions
- **Lock mode** - defines the level of access that other transactions have while the resource is locked
- **Blocking** - when a transaction requests a lock mode that conflicts with a currently held lock and has to wait for that lock to be released
- **Deadlock** - when two transactions block each other trying to acquire locks on resources the other transaction holds in a conflicting mode

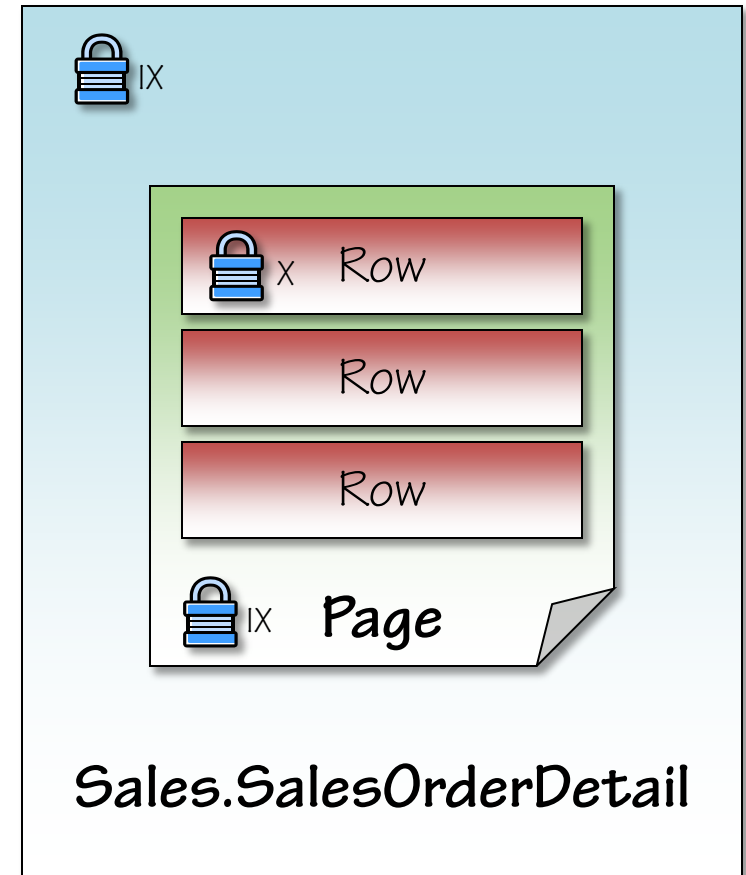
Lock Granularity

- **RID/KEY – a single row is locked**
 - RID – row identifier for a single row in a heap
 - KEY – index key for a single row in a index
- **PAGE - a single page in the database is locked**
- **HoBT – a heap or B-tree (index) partition is locked**
- **TABLE – the entire table is locked**
- **METADATA - the table schema definition is locked**
- **Locks are acquired at multiple levels of granularity to fully protect the lowest-level resource**
- **Locks are always acquired 'top-down', from the table level down to individual rows**
 - This forms the lock hierarchy



Lock Hierarchy

- DELETE statement begins executing that affects one row in a table
- A intent exclusive (IX) lock is acquired for the table
- A intent exclusive (IX) lock is acquired for the page containing the row
- A exclusive (X) lock is acquired for the row being modified



Shared and Update Locks

- **Shared (S) – acquired for read operations that do not modify data**
 - No transaction can modify data while the shared lock exists
 - Concurrent SELECT operations can read the data
 - Locks are released as soon as the read operation completes unless the isolation level is repeatable read or higher or hints are used
- **Update (U) – acquired for resources that will be updated**
 - Only one transaction can acquire an update (U) lock at a time
 - Prevents deadlocks caused by lock conversions from a shared (S) lock to an exclusive (X) lock
 - The update (U) lock is converted to an exclusive (X) lock to modify the data
 - Concurrent shared (S) locks are allowed

Exclusive and Intent Locks

- **Exclusive (X) – acquired for data modifications**
 - Prevents access to a resource from concurrent transactions
 - Ensures that multiple changes cannot be made to the same resource at the same time
 - Reads can only occur using NOLOCK or read uncommitted, read committed snapshot, or snapshot isolation levels
- **Intent (I) – acquired to establish the lock hierarchy**
 - Acquired on higher-level resources to protect locks on lower-level resources
 - E.g. a table-level intent exclusive (IX) lock is required before a page-level exclusive (X) or intent exclusive (IX) lock can be acquired
 - Prevents transactions from modifying/locking the higher-level resource with an incompatible lock for the lower-level lock being acquired
 - Intent shared (IS), intent exclusive (IX), shared with intent exclusive (SIX), intent update (IU), shared intent update (SIU) and update intent exclusive (UIX)

Schema and Key-Range Locks

- **Schema (Sch) – acquired to protect or modify an object's definition**
 - Schema modification (Sch-M) locks are acquired during DDL operations to prevent access to the object while its definition is changed
 - Schema stability (Sch-S) locks are acquired when compiling and executing queries to prevent modification of the object definition by DDL operations
- **Key-range - acquired to protect a range of rows when using the serializable isolation level**
 - Prevents other transactions from inserting rows into the range to ensure that the data returned by a query would remain the same if it were run again within the transaction

Lock Compatibility

- If a resource is already locked when a transaction requests a lock on it, the new lock can only be acquired if it is compatible with the existing lock on the resource
- The most common locks are shown here but a full compatibility matrix is available in Books Online (<http://bit.ly/SQLLockCompat>)

		Existing lock mode					
Requested mode		IS	S	U	IX	SIX	X
	Intent shared (IS)	Yes	Yes	Yes	Yes	Yes	No
	Shared (S)	Yes	Yes	Yes	No	No	No
	Update (U)	Yes	Yes	No	No	No	No
	Intent exclusive (IX)	Yes	No	No	Yes	No	No
	Shared with intent exclusive (SIX)	Yes	No	No	No	No	No
	Exclusive (X)	No	No	No	No	No	No

Isolation Levels and Locking Hints

- **SQL Server provides five different transaction isolation levels**
 - Read Uncommitted, Read Committed, Repeatable Read, Snapshot, Serializable
- **Locking hints can be specified for individual tables referenced in DML statements**
 - HOLDLOCK, NOLOCK, NOWAIT, PAGLOCK, READCOMMITTED, READCOMMITTEDLOCK, READPAST, READUNCOMMITTED, REPEATABLEREAD, ROWLOCK, SERIALIZABLE, TABLOCK, TABLOCKX, UPDLOCK, XLOCK
- **The isolation levels and locking hints change the locking behavior of transactions**
 - Full details are beyond the scope of this course, but will be covered in the upcoming course on SQL Server Transactions, Locking, and Isolation Levels
 - More details can be found in Books Online: (<http://bit.ly/TranIsoLevs>)

Lock Escalation

- **The lock granularity is chosen during query compilation**
 - Row, page, partition, or table
- **During query execution, the lock granularity may be escalated if the resources necessary for the lower-level locks are not available**
 - Lock escalation can occur from row-to-table or page-to-table
 - If partition-level escalation is enabled in SQL Server 2008 onwards, escalation can be row-to-partition or page-to-partition
 - Locks NEVER escalate row-to-page-to-partition/table
- **Lock escalation can be desirable under some circumstances, but the higher-level locking reduces concurrency and can lead to blocking**

Summary

- Locking is an integral part of relational database management systems and protects data during concurrent activity in the database
- Locks can be acquired at different levels of granularity and in different modes to control concurrent data access
- Blocking occurs when a transaction attempts to acquire a lock with a mode that conflicts with a lock that is held by another transaction
- Deadlocks occur when two transactions block each other trying to acquire locks on resources the other transaction holds in a conflicting mode
- The next module will look at:
 - Deadlock detection