

# SQL Server: Myths and Misconceptions

## Module 5: Locking

Paul S. Randal

Paul@SQLskills.com



# Introduction

- **Locking always occurs in SQL Server**
  - Lots of confusion about how locking works
  - Many misconceptions about what operations cause blocking
- **In this module:**
  - Six myths around locking and blocking

# Locking Myth #1

**UNTRUE!!**

- **Myth: Online index operations do not acquire locks**
- **There are two blocking locks held by online index operations**
  - 'Almost Online Index Operations' isn't a very good marketing term
- **Share table lock held at the start so the new index can be created**
- **Schema-Modification lock held at the end so the old index can be dropped**
- **Both locks can cause blocking while they're queued**
- **Table Intent-Share lock held for the duration of the operation**
  - Usually is not a blocking lock

# Locking Myth #2

**UNTRUE!!**

- **Myth: Backups cause blocking**
- **Backups do not generally take locks so do not cause blocking**
- **One exception: bulk-operation lock prevents minimally-logged operations from occurring concurrently with backups**
- **Backups are very I/O intensive so might cause I/O contention which *\*appears\** to be blocking as the workload slows down drastically**
- **To avoid this affect, consider:**
  - Taking large data backups at off-peak times
  - Taking more frequent transaction log backups
  - Separating database and backup file storage

# Locking Myth #3

**UNTRUE!!**

- **Myth: DBCC CHECKDB causes blocking**
- **DBCC CHECKDB has been online by default since SQL Server 2000**
  - In SQL Server 7.0 and before, it did take blocking locks
  - In SQL Server 2000 it performed online log analysis
  - In SQL Server 2005 onwards it uses a database snapshot
- **It is one of the most I/O intensive operations in SQL Server**
  - Its read-ahead mechanism is very aggressive
  - The I/O load can make it \*seem\* like the workload is blocked

## Locking Myth #4

**UNTRUE!!**

- **Myth: Lock escalation goes row to page, then page to table**
- **Lock escalation never does a two-step escalation**
- **Lock escalation usually escalates to a table lock**
- **SQL Server 2008 introduced a partition-level lock**
  - Lock escalation can be set to escalate to a partition
  - This allows concurrent partition-level operations with lock escalation
  - Default is still table-level lock escalation as some customers experienced deadlocks when partition-level escalation was tried as the default during early SQL Server 2008 CTPs

# Locking Myth #5

**YES AND NO!!**

- **Myth: Blocking and deadlocks can occur with unrelated table rows**
- **Before SQL Server 2008 R2 it was possible to have blocking and deadlocks involving unrelated rows in a table**
  - The lock hash algorithm before SQL Server 2008 R2 was prone to occasional invalid lock hash collisions
  - The rewritten algorithm is 100 million times less likely to have this occur
- **You can use the undocumented %%LOCKRES%% function to investigate suspected lock hash collisions**
  - See <http://bit.ly/JNDQWC> for an example

## Locking Myth #6

**UNTRUE!!**

- **Myth: NOLOCK / READ UNCOMMITTED means no locks are acquired**
- **NOLOCK and READ UNCOMMITTED do have to acquire some locks**
  - Schema-stability locks to prevent the structure of the table/index changing
  - BULK\_OPERATION locks on heaps to prevent reading of unformatted pages
- **NOLOCK and READ UNCOMMITTED are the same under the covers**
- **And they still have to take latches to access the physical page images in memory, so there's still some potential for blocking at the latch level**