

# SQL Server: Optimizing Ad Hoc Statement Performance

Module 3: Estimates and Selectivity

Kimberly L. Tripp

[Kimberly@SQLskills.com](mailto:Kimberly@SQLskills.com)

<http://www.SQLskills.com/blogs/Kimberly>

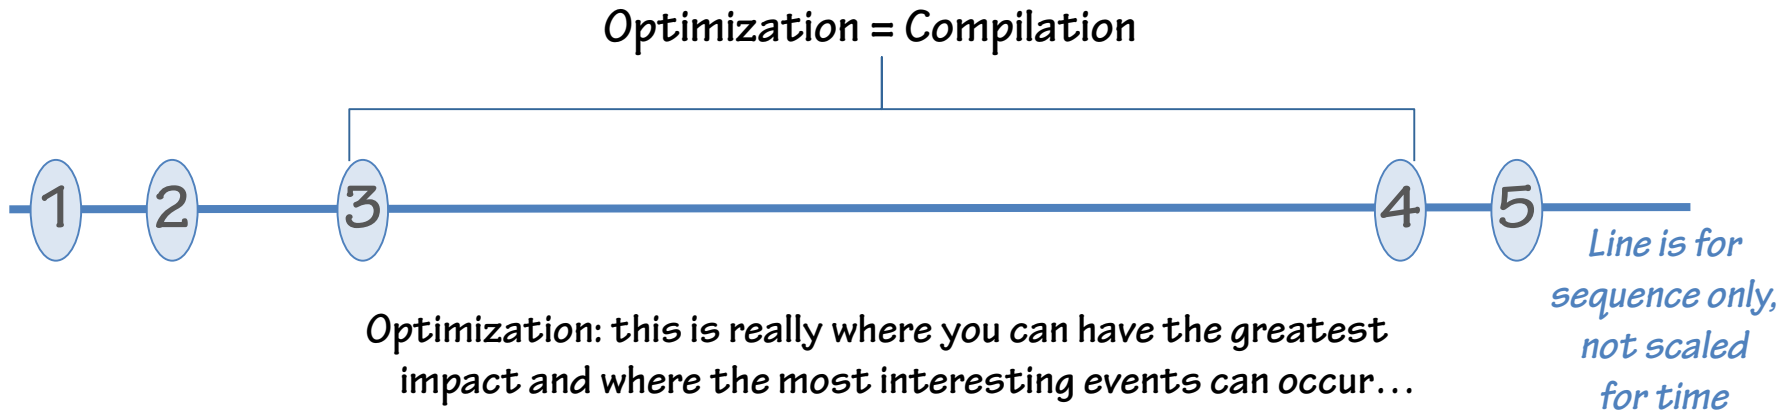


**pluralsight**   
hardcore developer training

# Course Overview

- **Statement execution methods**
- **Estimates and selectivity**
  - Statement execution simplified
  - Cost-based optimization
  - Understanding selectivity
  - Estimates, statistics, and heuristics
  - How do you see statistics?
  - What do statistics tell us about our data?
  - When and how does SQL Server use statistics?
- **Statement caching**
- **Plan cache pollution**
- **Statement execution summary**

# Statement Execution Simplified



1. Parse
2. Standardization/normalization/algebraization  
⇒ query tree (not Transact-SQL anymore)
3. Cost-based optimization (statistics are used to come up with an optimal plan, as well as other things)
4. Compilation
5. Execution

# Cost-Based Optimization

- Find a reasonable subset of possible algorithms to access data based on:
  - The query                      Sometimes a rewrite helps...
  - Any joins                      Sometimes a derived table (sub-query in the FROM clause)...
  - Any SARGs                      Your SARGs need to be well-defined...
  - Data selectivity
  - Join density
- The more information the optimizer has the better...
- How do you provide the BEST information?
- One of the best ways to “influence” your query plans is through effective statistics (and better indexes)

# Understanding Selectivity

- Imagine a table of employee data – for a Chicago company
- The table is clustered by *EmployeeID*
- Imagine executing this query:

```
SELECT [e].*  
FROM [dbo].[EmployeesAddresses] AS [e]  
WHERE [e].[city] = 'Chicago'    not selective enough  
WHERE [e].[city] = 'Glenview'  not an easy answer  
WHERE [e].[city] = 'Peoria'    selective enough
```

- When is an index on *city* useful?
  - When the data is selective ENOUGH...

More importantly,  
how does SQL Server know?

# Demo Summary: Estimates and Selectivity

- **Estimates come from:**

- Statistics – if they exist or if they can be (auto) created, using:
  - The histogram: when the value can be “sniffed” (parameters)
  - The density vector: when the value cannot be “sniffed” (variables)
- Heuristics – if there are no statistics available and SQL Server cannot auto create them
  - These are internal “magic” numbers (cannot be changed)
  - They often result in very poor plans (`LEAVE AUTO_CREATE_STATISTICS ON`)
  - Sometimes this is the only option when better estimations cannot occur (comparison between columns (e.g. `col1 > col2`))

- **Statistics have to be reasonably small to be fast/useful**

- **They’re just estimates**

- They’re not always guaranteed to be accurate
- They’re just meant to get us closer to the right value

# How Do You See Statistics?

- **DBCC SHOW\_STATISTICS** (**tname**, **statname**)
  - Gives you ALL the statistical details
    - Number of rows and number of rows on which the statistics were based
    - Densities for all LEFT-based subsets of the column, including the cluster key (last – if not already somewhere in the index)
    - Histogram for the high-order element
- **sp\_autostats** **tname**

Index Name	AUTOSTATS	Last Updated
[member_ident]	ON	2008-08-26 17:18:12.593
[member_corporation_link]	ON	2008-08-26 17:18:12.673
[member_region_link]	ON	2008-08-26 17:18:12.793
[MemberName]	ON	2008-10-29 11:13:29.220
[_WA_Sys_00000003_OCBAE877]	ON	2008-10-29 11:28:32.313

# What Do Statistics Tell Us About Our Data? (1)

1		2		Statistics Header			
Name	Updated	Rows	Rows Sampled	Steps	Density	Average key length	String Index
MemberName	Oct 10 2008 1:02AM	10000	10000	26	0	21.5526	YES

## Density Vector

3			All density	Average Length	Columns
			0.03846154	5.6154	Lastname
			0.0001	16.5526	Lastname, Firstname
			0.0001	17.5526	Lastname, Firstname, MiddleInitial
			0.0001	21.5526	Lastname, Firstname, MiddleInitial, member_no

## Histogram

4					
RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS	AVG_RANGE_ROWS	
ANDERSON	0	385	0	1	
BARR	0	385	0	1	
CHEN	0	385	0	1	
...	...	...	...	...	
ZUCKER	0	384	0	1	



# What Do Statistics Tell Us About Our Data? (2)

## ■ Statistics date (#1)

- This is the date that the statistics were last updated:
  - Through SQL Server's auto-updating mechanism
    - Database option: *AUTO\_UPDATE\_STATISTICS*
  - Manually, by executing one of the following:
    - *sp\_updatestats*
    - *UPDATE STATISTICS*
- Or, if they've never been updated then it represents date they were created:
  - Through SQL Server's auto-create mechanism
    - Database option: *AUTO\_CREATE\_STATISTICS*
  - Manually, by executing one of the following:
    - *sp\_createstats*
    - *CREATE STATISTICS*
- Can also get this information from the function: *STATS\_DATE()* and the new (new in SQL 2008 R2 SP2 / and SQL Server 2012 SP1) DMV:

**SELECT** \*

**FROM** [sys].[dm\_db\_stats\_properties](object\_id, index\_id)

# What Do Statistics Tell Us About Our Data? (3)

- **Data analyzed to build the statistics (#2)**
  - Rows – number of rows in the table at the time the statistics were built
  - Rows Sampled – the number of rows that were analyzed to generate the statistic
- **Sampling**
  - Does not directly indicate a problem with statistics
  - Could be a problem if your data is heavily skewed
- **Is it a problem?**
  - Using showplan tooltip – estimate vs. actual rows
    - If query performance is poor AND the actual is significantly OFF from the estimate then you might want to verify the statistics creation (rows v. rows sampled)
    - If statistics were based on a sampling and performance is improved after statistics have been updated, then you might want to turn off auto update for this index (using *STATISTICS\_NORECOMPUTE*) and schedule an *UPDATE STATISTICS WITH FULLSCAN*

# What Do Statistics Tell Us About Our Data? (4)

## ■ Density vector (#3)

- Shows the average distribution of data given the LEFT-based subsets of the entire key
- Rows \* all density = average number of rows returned
  - Based on that left-based subset of columns supplied
  - Density information
    - Density for *LastName*
      - $10,000 \text{ Rows} * 0.03846154 = 384.6154$
    - Density for *LastName, FirstName* combined
      - $10,000 \text{ Rows} * 0.0001 = 1$

## ■ Histogram (#4)

- Only stores data for the leading column of the index (sometimes referred to as the high-order element (e.g. *LastName*))
- Has actual values with details about that “step”
  - Anderson 385 rows
  - Barr 385 rows
  - ...

# When and How Does SQL Server Use Statistics?

- Estimation comes from “sniffing” the value
  - **Result:** estimate comes from *HISTOGRAM*
  - **Pro:** estimate is usually more accurate
  - **Pro:** \*that\* execution gets a plan designed for \*that\* value
  - **Con:** If/when this plan is saved – subsequent executions are prone to “parameter sniffing problems” (PSP)
- Value cannot be “sniffed”
  - **Result:** estimate comes from the *DENSITY\_VECTOR*
  - **Depends:** estimate is an “average”
  - **Depends:** the plan generated is designed for the “average” value – not \*that\* value
  - **Pro:** If/when this plan is saved – subsequent executions are NOT prone to PSP
  - **Con:** When your data is NOT [relatively] evenly distributed, this plan might not be good for anyone

# Summary: Estimates and Selectivity

- **Method: ad hoc statement**

- Can have literals
  - Can be “sniffed” and estimated using the *HISTOGRAM*
  - Can generate an optimal plan
- Can have variables
  - Cannot be “sniffed” (they are unknown during optimization/compilation)
  - Optimizes based on the average distribution of data (using the *DENSITY\_VECTOR*)
- Can be parameterized and cached but it’s extremely unlikely (only when safe)
  - Requires CPU/compilation on every execution

- **Method: *sp\_executesql***

- Can generate an optimal plan for the *first* execution
  - Saves CPU/compilations costs for subsequent executions
- Can be prone to parameter sniffing problems (PSP)
  - When the optimal plan varies (based on the parameters passed) then subsequent executions may suffer by using the plan chosen by the *first* execution’s parameters

- **Method: Dynamic String Execution through *EXEC (@string)***

- Turns the statement into an ad hoc statement
- It behaves exactly like an ad hoc statement