


SQL Server: Optimizing Ad Hoc Statement Performance

Module 4: Statement Caching

Kimberly L. Tripp

Kimberly@SQLskills.com

<http://www.SQLskills.com/blogs/Kimberly>



pluralsight 
hardcore developer training

Course Overview

- **Statement execution methods**
- **Estimates and selectivity**
- **Statement caching**
 - What affects ad hoc statement behavior?
 - Default ad hoc statement behavior
 - Ad hoc statement textual matching
 - Ad hoc statements – safe vs. unsafe
 - Ad hoc statement caching
 - Analyzing the plan cache
 - Changing ad hoc statement behavior
- **Plan cache pollution**
- **Statement execution summary**

What Affects Ad Hoc Statement Behavior?

- Default behavior for ad hoc statements affected by server and database settings
- In SQL Server 2005, behavior defined solely by the database option/setting:
 - Database option: *PARAMETERIZATION*
 - Default = *SIMPLE* (generally, recommended)
 - Optionally, *FORCED*
- In SQL Server 2008 onward, an additional server setting affects ad hoc statement behavior:
 - Server configuration setting for “*optimize for ad hoc workloads*”
 - Default = *OFF*
 - Optionally, *ON* (generally, recommended)

Default Ad Hoc Statement Behavior

- **On first execution (or, when SQL Server doesn't already have a plan in cache for a statement) then:**
 - Generate a plan, use it, and place it in the cache for future use
 - This is the cached plan type of "Adhoc" from *sys.dm_exec_cached_plans* (objtype) column
 - Analyze the statement to determine if it's "safe"
 - If it's safe – generate a plan, parameterize it, keep it in cache for subsequent executions ⇒ called auto-parameterization
 - This is the cached plan type of "Prepared" from *sys.dm_exec_cached_plans* (objtype)
 - If it's unsafe, only the ad hoc plan will be placed in the cache
- **On subsequent executions, every statement checks to see if there's a match of the statement in the plan cache**
 - If there's an **exact textual match** of the statement, use that plan
 - If there's a **parameterized plan**, use that plan

Ad Hoc Statement Textual Matching

- **Must be exact in every way**
 - No spaces, tabs, or hard-returns
 - No case differences (even when the database is not case-sensitive)
- **All of these statements have a different textual plan in cache**

```
select * from member
```

```
select * from dbo.member
```

```
select *      from member
```

```
select *      from      member
```

- Each has a different *sys.dm_exec_sql_text (sql_handle)*
- Each has a different *sys.dm_exec_query_plan (plan_handle)*

NOTE: Every statement above has the SAME *query_hash* (and, the same *query_plan_hash*)

Ad Hoc Statements – Safe vs. Unsafe

- **Statements are unsafe when the statement :**
 - Uses an *IN* clause
 - Has more than one table in the *FROM* clause
 - Uses expressions joined by *OR* in a *WHERE* clause
 - When a *SELECT* query contains a sub-query
- **OK, have I lost you yet? It's VERY restrictive (see whitepaper: *Plan Caching in SQL Server 2008*, Appendix A for a complete list)**
 - Available from <http://bit.ly/zt8bw>
- **A safe statement has a benefit – auto parameterization**
 - Idea: automatically parameterize search arguments and place the statement in cache for future re-use of the plan (saves CPU but not plan cache size)
 - Needs to be a VERY straightforward statement/plan
 - Parameters do not change plan choice
 - Doesn't happen very often and NOT something I would rely on and/or get excited about

Ad Hoc Statement Caching

- Take the following “unsafe” query (identical, except for the SARG):

```
SELECT ... WHERE [m].[lastname] = 'Tripp'  
SELECT ... WHERE [m].[lastname] = 'Tripps'  
SELECT ... WHERE [m].[lastname] = 'Tripped'  
SELECT ... WHERE [m].[lastname] = 'Falls'
```
- Each plan takes 16,384 bytes (or roughly 16KB)
- This “query class” is harder to track because each is listed in the cache
- All have the same *query_hash* but not necessarily the same *query_plan_hash*
- Query *sys.dm_exec_query_stats* for *query_hash* and aggregate over *query_hash*, *query_plan_hash* to see how many plans a particular query might have in the cache
- If there’s only one *query_plan_hash*, the statement may be stable
 - However, this is completely index and data dependent

Verifying Plans in Cache NOW

- **Use *syscacheobjects* (in SQL Server 2000 and SQL Server 2005)**
 - *SELECT * FROM master.dbo.syscacheobjects*
 - This lists plans in cache as well as parameterized plans
 - *usecounts* and *refcounts* are interesting in terms of frequency of execution
 - *pagesused* gives insight into the size of the plan
- **Use DMVs to see the same and more in SQL Server 2005 onward**
 - DMV: *sys.dm_exec_query_stats* – gives much of what you see from *syscacheobjects*
 - DMF: *sys.dm_exec_sql_text (sql_handle)* – returns the actual text but only when needed (if just query stats then it's faster not to include text)
 - DMF: *sys.dm_exec_query_plan (plan_handle)* – can give you the XML Showplan as well and there's no prior equivalent
- **SQL Server 2008 added *query_hash* and *query_plan_hash* to *sys.dm_exec_query_stats***

Analyzing the Plan Cache

- Every statement goes into the ad hoc plan cache for exact textual matching
 - Review *sys.dm_exec_cached_plans*
- SQL Server 2008 added *query_hash* and *query_plan_hash* to *sys.dm_exec_query_stats*
 - SQL templatises the parameters – similar to *sp_get_query_template*
 - Aggregate by *query_hash* to find similar queries
 - Aggregate by *query_hash*, *query_plan_hash* to find similar queries and their plans
 - Queries that have only one *query_plan_hash* are STABLE
 - Queries that have more than one (sometimes dozens) are UNSTABLE
- Check out Books Online topic “Finding and Tuning Similar Queries by Using Query and Query Plan Hashes” at <http://bit.ly/QfUmRY>

The “Cumulative Effect” of Queries

```
SELECT [Query Hash] = [qs2].[query_hash]
, [Query Plan Hash] = [qs2].[query_plan_hash]
, [Avg CPU Time] = SUM ([qs2].[total_worker_time]) /
    SUM ([qs2].[execution_count])
, [Example Statement Text] = MIN ([qs2].[statement_text])
FROM (SELECT [qs].*, [statement_text] = SUBSTRING ([st].[text],
    ([qs].[statement_start_offset] / 2) + 1
    , ((CASE [statement_end_offset]
        WHEN - 1 THEN DATALENGTH ([st].[text])
        ELSE [qs].[statement_end_offset] END
        - [qs].[statement_start_offset]) / 2) + 1)
FROM [sys].[dm_exec_query_stats] AS [qs]
CROSS APPLY [sys].[dm_exec_sql_text]
    ([qs].[sql_handle]) AS [st]) AS [qs2]
GROUP BY [qs2].[query_hash], [qs2].[query_plan_hash]
ORDER BY [Avg CPU Time] DESC;
```

Changing Ad Hoc Statement Behavior (1)

- **Server setting: *optimize for ad hoc workloads***
 - On first execution, only the *query_hash* will go into cache
 - On second execution (if), the plan will be placed in cache
- **Significantly reduces the amount of cache allocated**
 - Query plans are on 8KB boundaries; plans can be quite large
 - Minimum plan size is 8KB and some plans can be MB in size
 - Plans for ad hoc statements that only execute once have only a few hundred bytes (just the statement/*query_hash*) in cache (rather than the ad hoc plan)
 - Does not eliminate plan cache pollution but slows it down
- **Generally, highly recommended**

Changing Ad Hoc Statement Behavior (2)

- **Database option: parameterization *FORCED***
 - More statements are forced to be cached on first execution (using parameter sniffing)
 - **Pro:** if you have stable plans (only one *query_plan_hash* for a *query_hash* that has lots of executions) from a lot of ad hoc clients then this might help to reduce CPU
 - **Con:** If you have some statements that really aren't safe, you could end up executing bad plans (PSP)
 - Recommendation: BEFORE turning this on, investigate plan stability by using *query_hash* and *query_plan_hash*
 - Not all statements are cached/parameterized
 - A *WHERE* clause condition defined by *LIKE* is not parameterized
 - *INSERT EXEC*
 - See Books Online topic "Forced Parameterization" for more information at <http://bit.ly/1guvurN>
 - Some features cannot be leveraged or may generate suboptimal plans in *FORCED* parameterization:
 - Filtered indexes and filtered statistics
 - Indexed views and indexes on computed columns
 - Partitioned tables and partitioned views
- **Generally, not recommended but highly useful when appropriate**

Multiple Plans (Tipping/Covering)

- **High-priority queries and queries that are executed often need to reduce their cumulative effect on the server**
- **By covering a query you can reduce the number of possible plans that exist for a query**
 - Reduces I/Os required to access the data the query needs
 - Reducing [often, drastically] I/Os translates into:
 - Fewer pages in cache results in better cache efficiency; less time
 - Important note: It's NOT just about I/Os – there are cases where plans with fewer I/Os are actually MORE expensive because of other operations (like sorts, temp tables, etc.) so be careful not to get too wrapped up in just I/Os. Review the showplan and statistics time as well!)
- **In some cases when you cover a query and create a stable plan, SQL Server might recognize it as safe (and parameterize/cache it)**
- **In other cases, even if all of your plans are stable, SQL Server might not recognize them as safe**
 - It's then when you might consider FORCED parameterization to save CPU

Summary: Statement Caching

- **Ad hoc statements**

- Simple parameterization – almost all statements will be compiled just for that execution; they will not be parameterized/saved (see rules for parameterization in whitepaper)
- Forced parameterization – most statements will be parameterized/saved (you'll see this in decreased CPU/compilations and potential for parameter-sniffing problems)

- ***sp_executesql* (or, prepared statements)**

- A fantastic way to reduce the CPU/cache overhead that ad hoc has, but should only be used when a plan is stable and consistent
- This is a better way to force a statement into cache as opposed to using forced parameterization database-wide

- **Stored procedures**

- Can be “sniffed” but there are exceptions to what SQL Server will store in their plans