


SQL Server: Optimizing Ad Hoc Statement Performance

Module 5: Plan Cache Pollution

Kimberly L. Tripp

Kimberly@SQLskills.com

<http://www.SQLskills.com/blogs/Kimberly>



pluralsight
hardcore developer training

Course Overview

- **Statement execution methods**
- **Estimates and selectivity**
- **Statement caching**
- **Plan cache pollution**
 - Ad hoc plan cache pollution defined
 - Plan cache usage and limits
 - Plan cache stores
 - Verifying the state of plan cache
 - Balancing plan cache pollution, CPU, and parameter-sniffing problems
- **Statement execution summary**

Ad Hoc Plan Cache Pollution Defined

- **BECAUSE every statement goes into the ad hoc plan cache for exact textual matching**
 - Plans start to fill up cache
- **Plans are more expensive to put into cache than data, therefore SQL Server values them a bit more highly (sometimes stealing pages from the buffer pool to give to the plan cache)**
- **Plans – even if only used once – can sit in the cache for quite some time**

Plan Cache Usage and Limits

- The “plan cache” is also known as the “procedure cache”
- Uses “stolen” pages from the buffer pool (data pages)
- View cached plans: *sys.dm_exec_cached_plans*
- Plan cache memory limits:
 - SQL Server 2005 SP2 and higher
 - 75% of visible target memory from 0-4GB
 - + 10% of visible target memory from 4GB-64GB
 - + 5% of visible target memory > 64GB
 - SQL Server 2005 RTM/SP1
 - 75% of visible target memory from 0-8GB
 - + 50% of visible target memory from 8GB-64GB
 - + 25% of visible target memory > 64GB
 - SQL Server 2000
 - 4GB upper cap on the plan cache

Current Versions	
Memory	Plan Cache
4GB	3.0GB
8GB	3.5GB
16GB	4.2GB
32GB	5.8GB
64GB	9.0GB
128GB	12.2GB
256GB	30.0GB

Plan Cache Stores

- ***CACHESTORE_OBJCP* = "Object Plans"**
 - Stored procedures, functions, triggers...
 - Generally, it's desirable to have a higher value especially when reused
- ***CACHESTORE_SQLCP* = "SQL Plans"**
 - Ad hoc SQL statements (including parameterized ones)
 - Prepared statements
 - OK when highly reused, but often not reused
- **Lots of other cachestores but "SQL Plans" is what we're focused on and it can be the most problematic**

```
SELECT [mc].*
FROM [sys].[dm_os_memory_clerks] AS [mc]
ORDER BY [mc].[pages_kb] DESC
```

Verifying State of Plan Cache

-- How much of your cache is allocated to single-use plans?

```
SELECT [Cache Type] = [cp].[objtype]
, [Total Plans] = COUNT_BIG (*)
, [Total MBs] = SUM (CAST ([cp].[size_in_bytes]
    AS DECIMAL (18, 2))) / 1024 / 1024
, [Avg Use Count] = AVG ([cp].[usecounts])
, [Total MBs - USE Count 1] = SUM (CAST ((CASE
    WHEN [cp].[usecounts] = 1 THEN [cp].[size_in_bytes]
    ELSE 0 END) AS DECIMAL (18, 2))) / 1024 / 1024
, [Total Plans - USE Count 1]
    = SUM (CASE WHEN [cp].[usecounts] = 1 THEN 1 ELSE 0 END)
, [Percent Wasted] = [Total MBs - USE Count 1]/[Total MBs]*100
FROM [sys].[dm_exec_cached_plans] AS [cp]
GROUP BY [cp].[objtype]
ORDER BY [Total MBs - USE Count 1] DESC;
```

Balancing Plan Cache Pollution, CPU, and PSP (1)

- **Server setting: *optimize for ad hoc workloads***
 - On first execution, only the *query_hash* will go into cache
 - On second execution (if), the plan will be placed in cache
- **Create a single and more consistent plan with covering indexes – might make the plans more stable!**
 - SQL Server will pick up SOME stable statements IF and ONLY IF they're SAFE
 - (Back to the rules from the whitepaper)
 - Note: this only reduces compilation costs (e.g. CPU) but it does not reduce plan cache pollution because every ad hoc statement still goes into the ad hoc cache
 - If you create a bunch of stable plans that SQL doesn't see as SAFE but they essentially are (one *query_plan_hash* for each *query_hash*) then you can consider the database setting: forced parameterization
 - If you're finding A LOT of single-use statements that have the same *query_hash* and are executed frequently but with only one *query_plan_hash* then this is ideal!
 - But, remember, ad hoc statements are always placed in the ad hoc plan cache so you still need *optimize for ad hoc workloads*...set that FIRST!

Balancing Plan Cache Pollution, CPU, and PSP (2)

- Analyze the plan cache for the number of query plans per *query_hash* (as well as the number of executions)

```
SELECT [query_hash]
      , [# of Plans] = COUNT (DISTINCT [query_plan_hash])
      , [Execution Total]= SUM ([execution_count])
```

...

- Two primary scenarios to consider

Scenario 1

query_hash	# of Plans	# of Executions
Ox04BB791B589774AD	1	6456456
Ox1706E9EC3049A95B	6	276543
Ox5BD9FF487079B335	1	124345
Ox6604520C5200ABCO	1	78905
Ox77BA5A89C7EBE605	1	14342
OxA078B4BC8768A9A6	1	4567
OxB81E270A58A79D16	1	6

Mostly stable plans
(only 1 plan per *query_hash*)

Scenario 2

query_hash	# of Plans	# of Executions
Ox04BB791B589774AD	34	6456456
Ox1706E9EC3049A95B	1	276543
Ox5BD9FF487079B335	8	124345
Ox6604520C5200ABCO	3	78905
Ox77BA5A89C7EBE605	2	14342
OxA078B4BC8768A9A6	9	4567
OxB81E270A58A79D16	24	6

Mostly Unstable plans
(multiple plans per *query_hash*)

Verifying State of Plan Cache (2)

```
-- How much is each query_hash using and how many plans?
SELECT [qs].[query_hash]
      , [Distinct Plan Count]
        = COUNT (DISTINCT [qs].[query_plan_hash])
      , [Execution Total]
        = SUM ([qs].[execution_count])
      , [Total MB]
        = SUM (cp.size_in_bytes) / 1024.0 / 1024.0
FROM [sys].[dm_exec_query_stats] AS [qs]
     INNER JOIN [sys].[dm_exec_cached_plans] AS [cp]
       ON [qs].[plan_handle] = [cp].[plan_handle]
GROUP BY [qs].[query_hash]
ORDER BY [Execution Total] DESC;
GO
```

Balancing Plan Cache Pollution, CPU, and PSP (3)

- Scenario 1: Consider changing parameterization to *FORCED*
- Then, for all of the statements that have more than one plan, you'll need to make sure that they are not forced
 - Can you change the code?
 - Change the ad hoc statements to be executed with a recompilation option (see my next course: *Optimizing Stored Procedure Performance*)
 - What if you can't change the code?
 - Use "templated" plan guides to take the few statements that are NOT stable and make *SIMPLE* (recompiled)
 - Use *sp_get_query_template* to get the templated query and parameters

```
EXEC sp_create_plan_guide N'PlanGuideName'  
    , @StatementWithUnStableQueryPlan  
    , N'TEMPLATE '  
    , NULL  
    , @Parameters  
    , N'OPTION(PARAMETERIZATION SIMPLE)';
```

Balancing Plan Cache Pollution, CPU, and PSP (4)

- Scenario 2: Keep the default parameterization mode: *SIMPLE*
- Then, for all of the statements that have only one plan, force them individually
 - Can you change the code?
 - Change the ad hoc statements to be executed with *sp_executesql* or stored procedures (see my next course: *Optimizing Stored Procedure Performance*)
 - What if you can't change the code?
 - Use "templated" plan guides to take the few statements that are stable and make them *FORCED* (reduced compilation/CPU)
 - Use *sp_get_query_template* to get the templated query and parameters

```
EXEC sp_create_plan_guide N'PlanGuideName'  
    , @StatementWithStableQueryPlan  
    , N'TEMPLATE '  
    , NULL  
    , @Parameters  
    , N'OPTION(PARAMETERIZATION FORCED)';
```

Balancing Plan Cache Pollution, CPU, and PSP (5)

- Start by configuring '*optimize for ad hoc workloads*'
 - This helps, but it's not always enough
- Additionally, consider a SQL Agent Job to periodically wake up and analyze the plan cache for single-use plans
 - If more than a certain percentage of memory is being wasted...
 - If more than a certain amount of memory (GB) is being wasted...
- Clearing the plan cache
 - **DBCC FREEPROCCACHE**
[({ plan_handle | sql_handle | pool_name })]
[WITH NO_INFOMSGS]
 - All plans: *DBCC FREEPROCCACHE*
 - A single plan: *DBCC FREEPROCCACHE (plan_handle)*
 - All plans for a single database: *DBCC FLUSHPROCINDB (DBID)*
 - **DBCC FREESYSTEMCACHE** ('ALL' [, pool_name])
[WITH {[MARK_IN_USE_FOR_REMOVAL], [NO_INFOMSGS]}]
 - All ad hoc and prepared plans: **DBCC FREESYSTEMCACHE** (N'SQL Plans')

Alternatives to Ad Hoc Statements

- **When you can change the code and when the plan is stable:**
 - Use forced statement caching
 - *sp_executesql*
 - Stored procedures (more options/control)
 - Only one plan goes into the cache
 - Saves cache (an ad hoc plan is NOT saved, only the prepared statement)
 - Saves compilation time
 - This SOUNDS perfect...
 - When used with a statement whose plans are inconsistent – problems from parameter sniffing can be worse
- **When you can't change the code then you need to determine:**
 - Can you tune an ad hoc statement to make it safe (or, at least stable)?
 - Can you benefit from changing to forced parameterization? (test!)
 - Are there any statements that require a plan guide template to force them to the opposite of the database option chosen?

Summary: Plan Cache Pollution

- **My general recommendation:**
 - Reduce plan cache pollution
 - Better to use the cache allocated by reusing plans for stable plans
- **How can you better reuse plans?**
 - Better tuning to create safe or more stable plans
 - Use *sp_executesql* to force those that are stable
 - Use stored procedures
- **Reusing plans isn't always a good thing:**
 - Recompilation is not always bad
 - Compiled plans aren't always perfect for all executions
- **Ad hoc statements are going to happen in every environment**
 - Must do some form of maintenance
- **Plan cache pollution is also very likely – incredibly rare that we aren't setting the configuration option AND setting up regular clearing of the 'SQL Plans' cache**