

SQL Server: Troubleshooting Query Plan Quality Issues

Module 2: Why Query Plan Quality Matters

Joe Sack

Joe@SQLskills.com

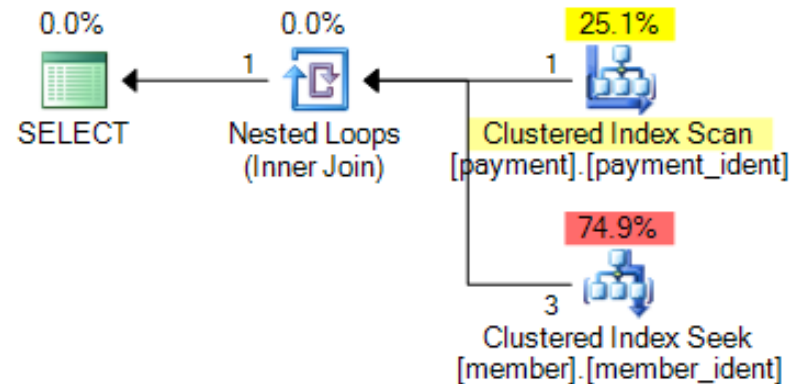
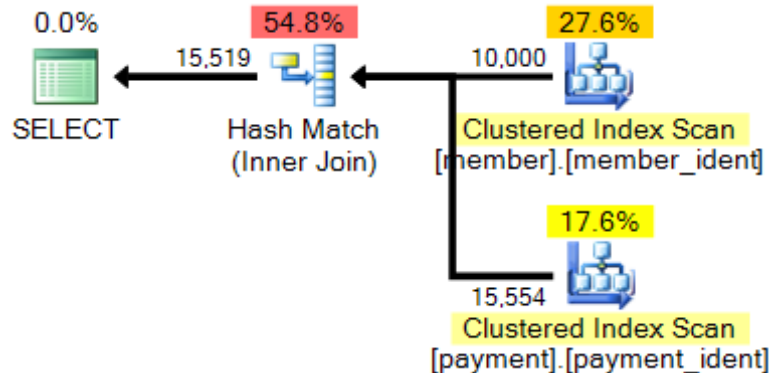


Module Introduction

- When a query execution plan is generated based on skewed data, the consequences can be significant
- It's common to troubleshoot the side-effects of query plan quality issues, but troubleshooting the root cause can be far more effective
- This module will cover the fundamentals of why you should care about query plan quality

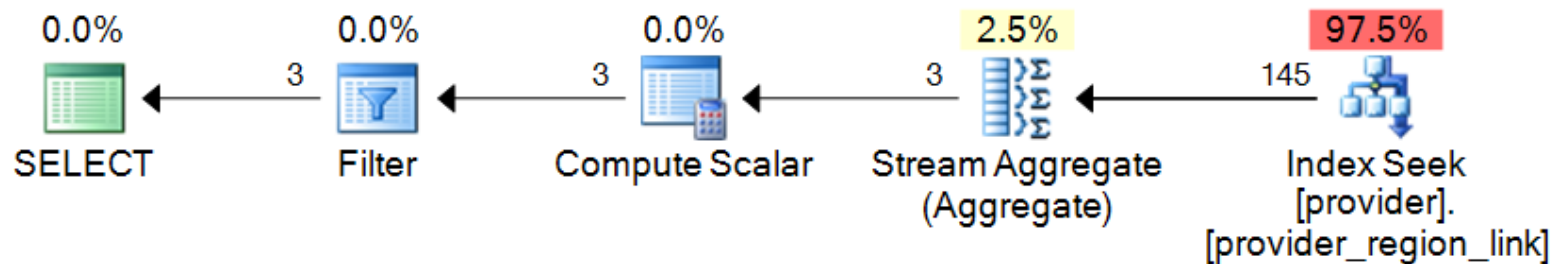
Which is the “Good” Plan?

```
SELECT
    [member].[member_no],
    [member].[lastname],
    [payment].[payment_no],
    [payment].[payment_dt],
    [payment].[payment_amt]
FROM [dbo].[member]
INNER JOIN [dbo].[payment] ON
    [member].[member_no] = [payment].[member_no];
```



Cardinality Estimates

- Estimate of rows processed per operator in a query plan
 - Row estimates calculated from leaf-level to root
 - Non-leaf level operators look at descendent operator estimates and apply additional estimation activities, for example, filtering



- Cardinality estimation relies on statistics, constraints, and heuristics

Costing and Plan Quality

- **Row estimates are fed into individual operator cost models**
 - Costing is a major factor in deciding which plan is chosen
- **What if row count assumptions are incorrect?**
 - Bad assumptions can lead to inefficient plan choices
 - A suboptimal plan can result in a slow-running query, excessive resource consumption, and reduced concurrency

Operator Cost ⁽¹⁾

- **Cost originally equated to elapsed time in seconds required to run on a specific Microsoft employee's machine (during SQL 7.0 time-frame)**
- **So really, "cost" today in the context of query plans is a unit-less measure**
 - Cost does not equate to time
- **Cost is used for relative comparison across plan operators and between plans**

Operator Cost ⁽²⁾

- **Operator cost = I/O cost + CPU cost**
- **Cost calculation varies by operator**
 - Some have I/O and CPU costs
 - Some have just CPU cost
- **Sub-tree cost = cost of specific operator + descendants**
 - Total cost found in root operator

Operator Cost ⁽³⁾

- **I/O cost assumptions (example):**
 - Data pages NOT in cache
 - Random I/O = 0.003125
 - Sequential I/O = 0.000740741
- **These aren't formally documented**

Operator Memory

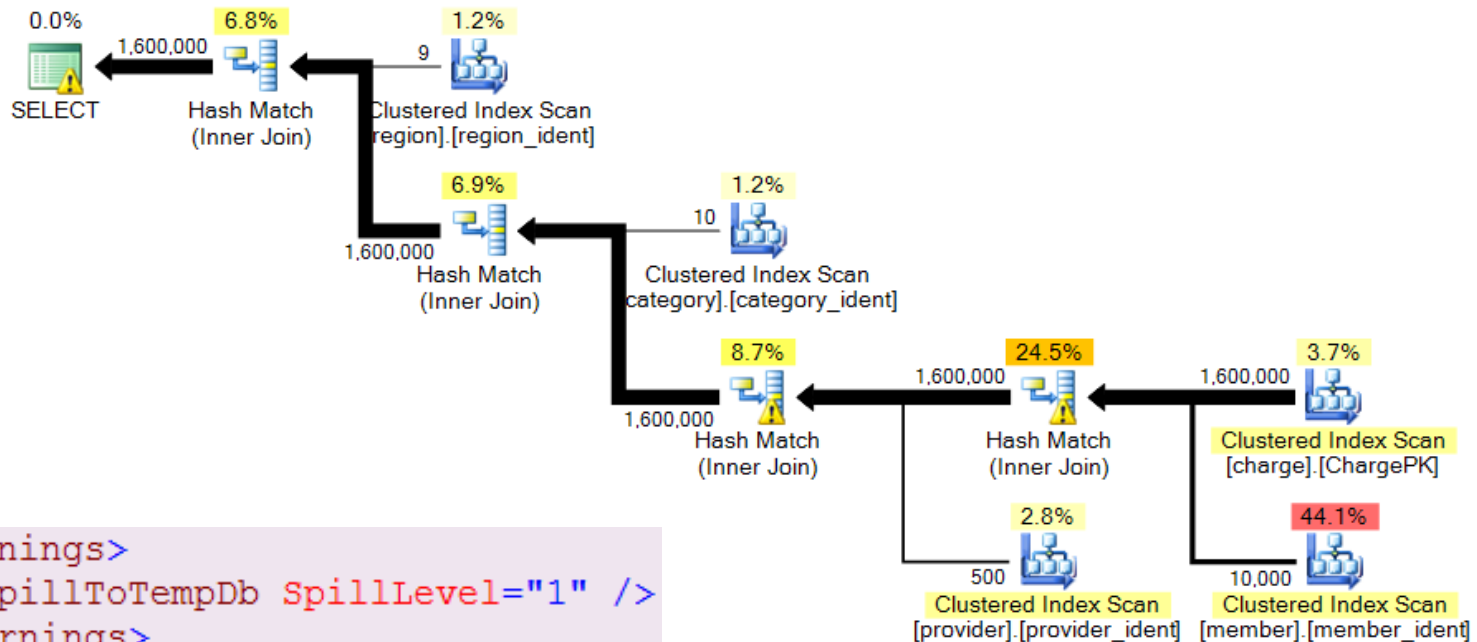
- Each type of operator requires varying amounts of memory in order to perform the associated operation
- Some operators require more memory because they cache rows
 - More rows = more memory required
 - SQL Server performs estimates of the required memory and tries to reserve the memory grant prior to execution
 - This is where cardinality estimation is critical

Memory Operators

- **Memory consuming operators include:**
 - Hash Match
 - Sort
- **When available memory is insufficient, queries that require a lot of memory may wait to execute**
- **Under-estimating memory (due to cardinality estimate issues) can cause memory spills to tempdb (I/O)**
- **Over-estimating memory can reduce concurrency!**

Under-estimates and Spills

- Significant under-estimates?
 - Hash and Sort operations example
 - Memory grants may be insufficient
 - Risk of spills to disk, excessive I/O



Over-estimates and Concurrency

- **Significant over-estimates?**
 - Hash and Sort operation impact example
 - Memory grants may be unnecessarily inflated
 - Concurrent requests wait for inflated memory grants
 - Concurrency throttled unnecessarily

	requested_memory_kb	granted_memory_kb	required_memory_kb
1	94512	NULL	36496
2	94608	94608	36496
3	95008	NULL	36496

sys.dm_exec_query_memory_grants

Impacted Query Optimizer Decisions

- **Decisions impacted by cardinality estimates?**
 - Index selection
 - Seek or scan
 - Parallel or serial plan execution
 - Join algorithms
 - Inner or outer table selection
 - Spool generation
 - Bookmark lookups
 - Stream or hash aggregates
 - Wide or narrow updates

Excessive Resource Consumption

- **Bad plans can be responsible for**
 - Excessive I/O
 - Inflated CPU
 - Memory pressure
 - Decreased throughput
 - Reduced concurrency
- **The takeaway?**
 - Don't ignore cardinality estimate issues