

SQL Server: Troubleshooting Query Plan Quality Issues

Module 4: Query Plan Quality Patterns and Resolutions

Joe Sack

Joe@SQLskills.com



pluralsight
hardcore developer training

Module Introduction

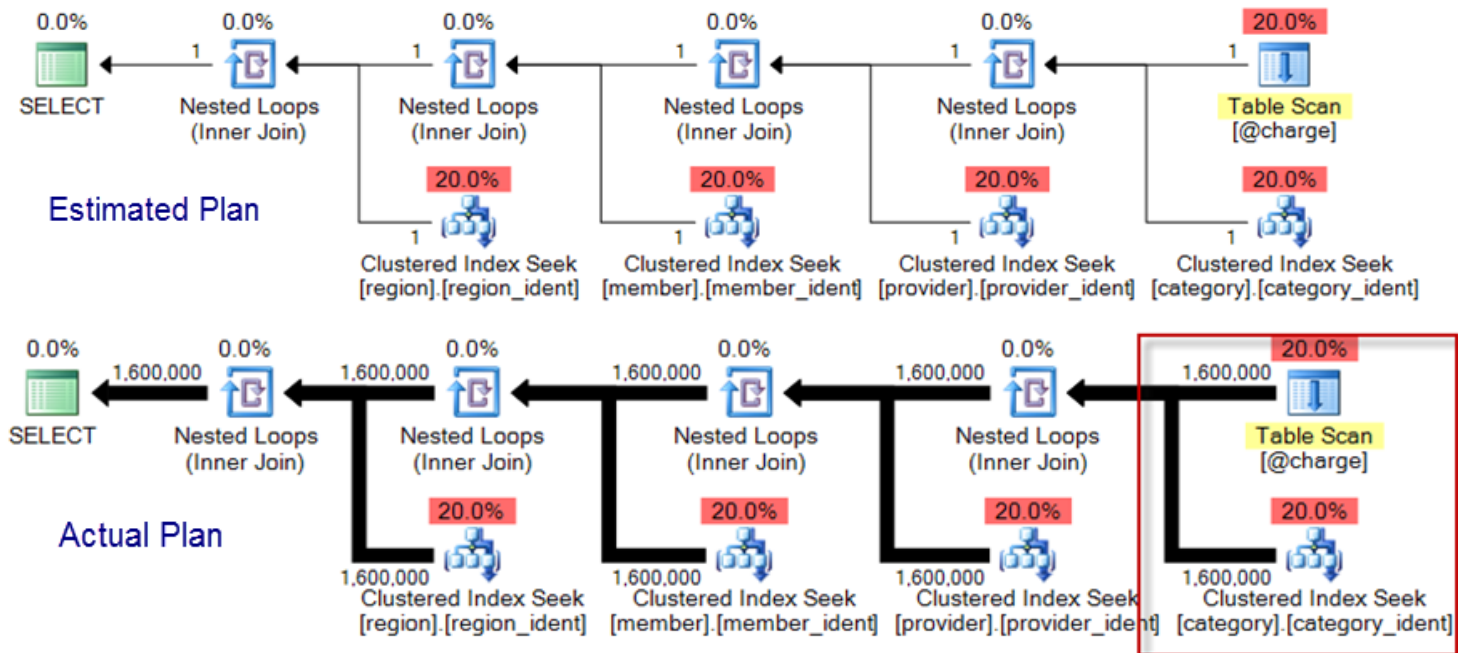
- **There are many reasons why query plan quality may be poor**
- **This module will review the more common patterns**
- **For each pattern we'll also cover the recommended next steps and associated resolutions**

Before Jumping In...

- **Is performance or concurrency degraded?**
 - Not all skews translate to poor performance
 - E.g. a hash join being used due to a skew, but ultimately the hash join operation is more efficient for that particular query
 - No magic number for defining when a skew is problematic
- **Sometimes the resolution involves significant refactoring of the T-SQL code and associated schema**
 - Keep the cost of effort versus the desired benefit in mind

Issue Prioritization

- Poor assumptions propagate “up” the tree (right-to-left)



- Troubleshoot the root cause, not the side effects

Missing or Stale Statistics (1)

- **Missing or out-of-date statistics**
 - Of cardinality estimate issues, this is often the easiest to address
- **Next steps?**
 - Check database options
 - `is_auto_create_stats_on`
 - `is_auto_update_stats_on`
- **Check “no recompute” settings**
 - `sys.stats` catalog view with `no_recompute = 1`
 - Can be changed via `sp_autostats`, `UPDATE STATISTICS`, `CREATE STATISTICS`, `CREATE INDEX`, `ALTER INDEX`

Missing or Stale Statistics (2)

- **Next steps?**

- Check STATS_DATE to validate last update
- Manual statistics updates may be needed where auto-updates are inadequate
 - For example, for very large tables where the automatic update threshold is significantly higher, less frequent

- **TF 2371 for changes to automatic statistics updates, decreasing dynamic threshold as table rows grow**

- Trade-off is increased recompilation, new plans
- See “Changes to automatic update statistics in SQL Server – traceflag 2371” from the Microsoft SAP on SQL Server Blog, <http://bit.ly/qOAlqs>

Statistics Sampling Issue

- **Precision of statistics histogram may be inadequate**
 - Very large table or significant data skews

- **Next steps?**
 - Compare DBCC SHOW_STATISTICS “rows sampled” versus “rows”
 - Can also use the new sys.dm_db_stats_properties
 - See Erin Stellato’s blog post for more details -> <http://bit.ly/NyCzWf>
 - Evaluate impact based on using higher percent sampling or FULLSCAN during manual statistics update
 - Overhead trade-offs for updating statistics on very large tables
 - Explore filtered statistics for very large tables
 - Beware of update threshold, which is based on overall table thresholds and NOT the filter itself

Hidden Column Correlation

- **Are columns correlated?**
 - Query Optimizer assumption is that columns are independent
- **How to resolve?**
 - Create multi-column statistics
 - If appropriate, create multi-column indexes

Comparison of Intra-Table Columns

- **Cardinality estimate issues can occur when comparing intra-table columns**
- **How to resolve?**
 - Computed columns and ensuring that associated statistics are generated
 - Self-joins, common table expressions
 - Normalization (e.g. order header vs. detail)

Table Variable Usage

- **Small number of rows?**
 - May be a non-issue
- **Large or volatile result set**
 - Can significantly impact query plan quality
- **Next steps**
 - Compare cardinality using a temporary table or permanent table instead

Misc	
Actual Execution Mode	Row
Actual Number of Batches	0
Actual Number of Rows	1600000
Actual Rebinds	0
Actual Rewinds	0
Defined Values	@charge.charge_no, @charge.member_no, @charge...
Description	Scan rows from a table.
Estimated CPU Cost	0.0001581
Estimated Execution Mode	Row
Estimated I/O Cost	0.003125
Estimated Number of Executions	1
Estimated Number of Rows	1
Estimated Operator Cost	0.0032831 (20%)
Estimated Rebinds	0
Estimated Rewinds	0
Estimated Row Size	23 B
Estimated Subtree Cost	0.0032831
Forced Index	False
ForceScan	False
Logical Operation	Table Scan
Node ID	4
NoExpandHint	False
Number of Executions	1
Object	[@charge] [charge]
Ordered	False
Output List	@charge.charge_no, @charge.member_no, @charge...
Parallel	False
Physical Operation	Table Scan
TableCardinality	0

Scalar and MSTV UDFs

- **Black-box from a cardinality estimate perspective**
 - Multi-statement table-valued function
 - Scalar UDFs
- **Resolutions**
 - Inline versus multi-statement?
 - Is a function even necessary?

Parameter Sniffing

- **Initial compilation of a plan based on the first value passed**

- Can be a good thing, but with wildly varying plan shapes, can be bad

- **If suspected, validate:**

- ParameterCompiledValue
- ParameterRuntimeValue

Parameter List	@member_no
Column	@member_no
Parameter Compiled Value	(1)
Parameter Runtime Value	(2)

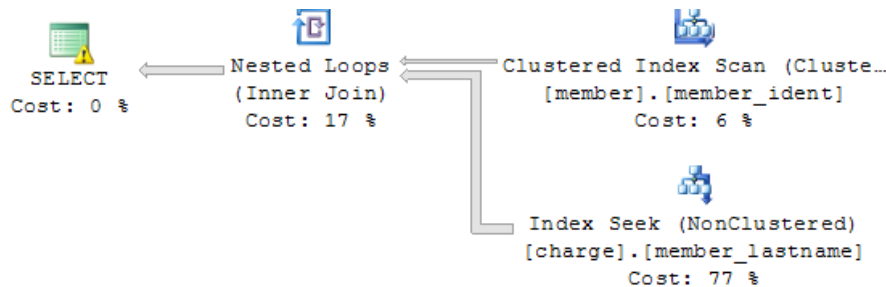
- **Several options to explore...**

- Procedure branching and isolation based on skew parameters
- OPTIMIZE FOR, local variable method
 - But beware since local variables may introduce cardinality estimation issues of their own
- Various recompile options at different scopes

Implicit Data-Type Conversion Issues

- Implicit data-type conversions

- Watch for data types of columns in search and join conditions



Warnings

Type conversion in expression `(CONVERT_IMPLICIT(nvarchar(15), [Credit].[dbo].[member].[lastname],0))` may affect "CardinalityEstimate" in query plan choice, Type conversion in expression `(CONVERT_IMPLICIT(nvarchar(15),[Credit].[dbo].[member].[lastname],0)=N'XAVIER')` may affect "SeekPlan" in query plan choice

- How to resolve?

- Use identical data types
- Use consistent naming conventions so that you can use catalog views to monitor any inconsistencies across column references

Complex Predicates

- **Wrapping column references in functions and complex expressions can hamper accurate cardinality estimation**
 - Not all are problematic, e.g. LOWER, UPPER, GETDATE
- **How to resolve?**
 - Keep your expression column references “clean”
 - Move complex expressions off of the columns
 - Prepare the value before passing to the predicate

Query Complexity

- **Sometimes the complexity of a query can make it almost impossible to have accurate cardinality estimates**
 - Views referencing views referencing views referencing functions within other views which reference the same tables referenced within views...
- **Resolution?**
 - If such queries are not meeting performance SLAs, one valid technique is to break apart a monolithic query into smaller chunks and intermediate result sets
 - Gives the optimizer fewer factors to consider
 - Multiple optimal plans may outperform one large suboptimal one
 - Trade-off of intermediate result set overhead must be weighed against plan quality benefits

Hints

- **Hint usage can cause cardinality estimate issues**
 - Various hints, including join hints, FAST number_rows
- **How to resolve?**
 - Evaluate whether hints are still needed and remove accordingly

Distributed Queries

- **Permissions associated with the linked-server definition and associated distributed queries can impact whether proper cardinality estimates can be generated**

- **Resolution?**
 - Ensure db_ddladmin fixed database role for the linked server account (minimum required permissions)
 - SQL Server 2012 Service Pack 1 has a fix that alleviates this requirement
 - See blog post “Distributed Query Plan Quality and SQL Server 2012 SP1”
<http://bit.ly/Tapozh>
 - Localization via:
 - ETL?
 - Replication?
 - Consolidation?

Query Optimizer Bugs

- Some query plan quality issues may be “bugs” or regressions
- Do you think it’s a bug? If so:
 - Check <http://connect.microsoft.com>
 - For existing entries, add your comments
 - When there is no match, take the time to enter in the reproduction of the bug with as much detail as possible

Course Summary

- Cardinality estimates are critical for proper operator costing and overall query execution plan quality
- Poorly performing query? Pay attention to skews
- Problem manifests in a variety of way, so your solutions will vary
- Resolve operator leaf-level skews first before addressing intermediate-level operator cardinality estimation skews
- Improving cardinality estimates, improves query plan quality
- Thanks for watching!