

SQL Server: Performance Troubleshooting Using Wait Statistics

Module 2: SQL Server Threading Model

Paul S. Randal

Paul@SQLskills.com



Introduction

- **SQL Server is really a mini operating system**
- **It performs and controls its own:**
 - Memory management
 - I/O
 - Scheduling
- **In this module we'll cover:**
 - Threads
 - Scheduling
 - Scheduling states and state transitions
 - Special cases


What Are Threads?


- **A thread is unit of execution within a process**
- **Multiple threads can exist within a process**
 - Each is given small amount of processor time, then waits while other threads execute (called scheduling)
 - This gives the impression of operations happening concurrently
 - See this Wikipedia link for more background: <http://bit.ly/Mx94C4>
- **SQL Server uses operating system threads to perform its work**
 - Called worker threads
 - Some are dedicated threads, such as for performing checkpoints, deadlock monitoring, and ghost record cleanup
 - Most are in a pool of threads that SQL Server uses to process user requests

Parallelism

- **The Query Processor may determine that a particular operation can be performed more efficiently by using multiple threads**
 - Called running in parallel
 - Multiple threads execute a small portion of the operation at the same time
- **In that case, the Query Optimizer will produce a query plan that can be parallelized**
 - Otherwise it generates plan that can only run serially (i.e. single-threaded)
- **At execution time, the Query Execution portion of the Query Processor decide what parallelism to use**
 - If the plan is serial only, no parallelism is used
 - If the plan can be run in parallel, the decision is made on how far to parallelize based on available resources, and a variety of configuration options

Parallel Threads Example

- As part of a query plan, you may see the  operator, for example
- This is a Repartition Streams operation
 - Uses producer and consumer threads, plus a control thread
- For a degree-of-parallelism = 4 operation, the threads would look like:



Controlling Parallelism


- **By default, the server-wide limit on parallelism is disabled**
 - The sp_configure option 'max degree of parallelism' is set to '0'
 - SQL Server decides how far to parallelize each query
- **If configured to anything except '0', this limit can be over-ridden**
 - Use the OPTION MAXDOP (x) syntax, where 'x' is the degree of parallelism requested (from 0 to 64)
 - Any user with any privilege can do this
- **In Enterprise Edition of SQL Server 2008 onwards, Resource Governor can be used to provide a hard limit to the degree of parallelism a query can utilize**
 - Using the MAX_DOP option for a workload group
 - This cannot be over-ridden in any way
- **Specifying a degree of parallelism does not guarantee it will be used**

Thread Scheduling

- **SQL Server performs its own thread scheduling**
 - Called non-preemptive scheduling
 - More efficient (for SQL Server) than relying on Windows scheduling
 - Performed by the SQLOS layer of the Storage Engine
- **Each processor core (whether logical or physical) has a scheduler**
 - A scheduler is responsible for managing the execution of work by threads
 - Schedulers exist for user threads and for internal operations
 - Use the `sys.dm_os_schedulers` DMV to view schedulers
- **When SQL Server has to call out to the OS, it must switch the calling thread to preemptive mode so the OS can interrupt it if necessary**
 - More information on this in Module 5

Components of a Scheduler

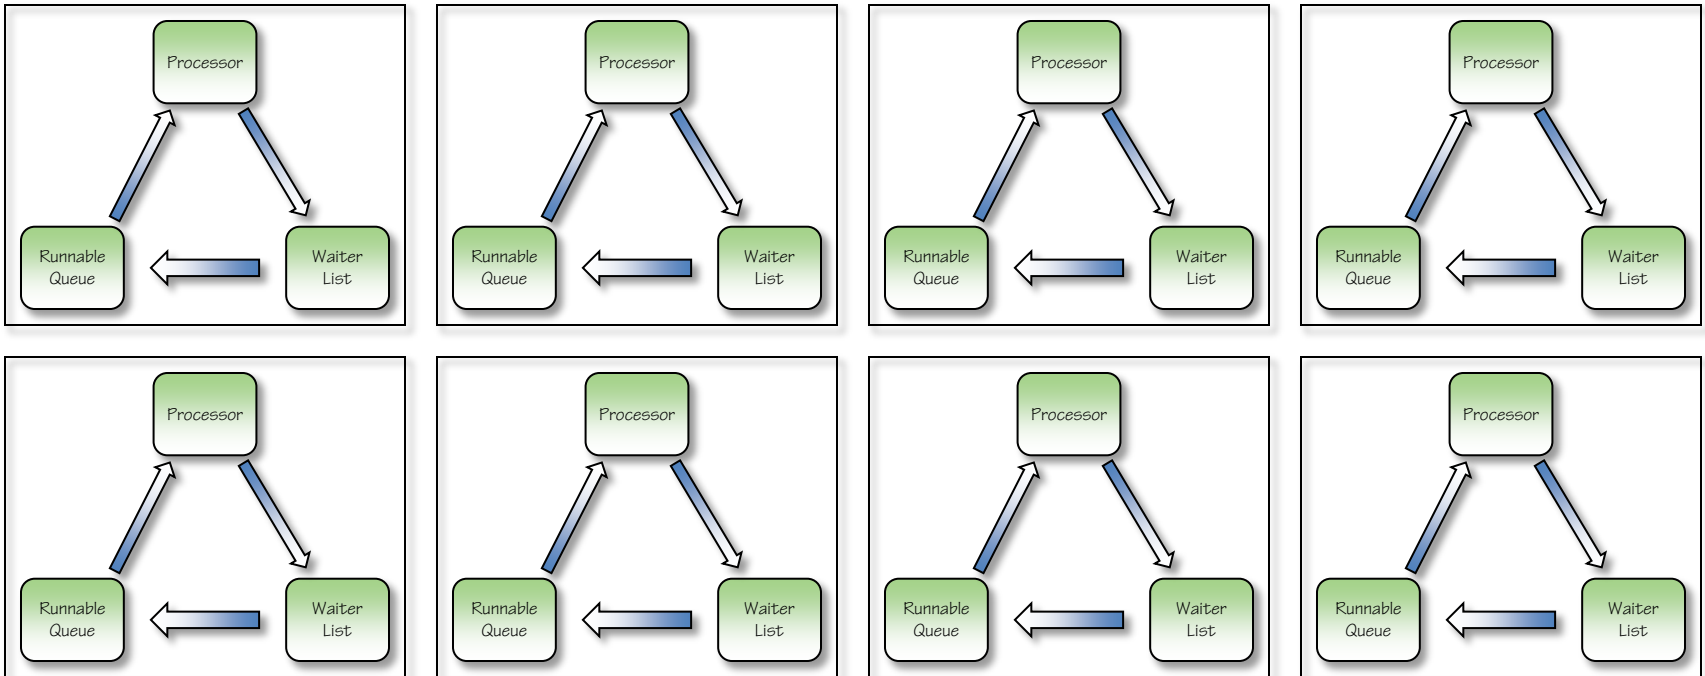
- All schedulers are composed of three 'parts'



- Threads transition around these parts until their work is complete

Schedulers in SQL Server

- **One scheduler per logical or physical processor core**
 - Plus some extra ones for internal tasks and the Dedicated Admin Connection
- **For example, for a server with four physical processor cores, with hyper-threading enabled, there will be eight user schedulers**

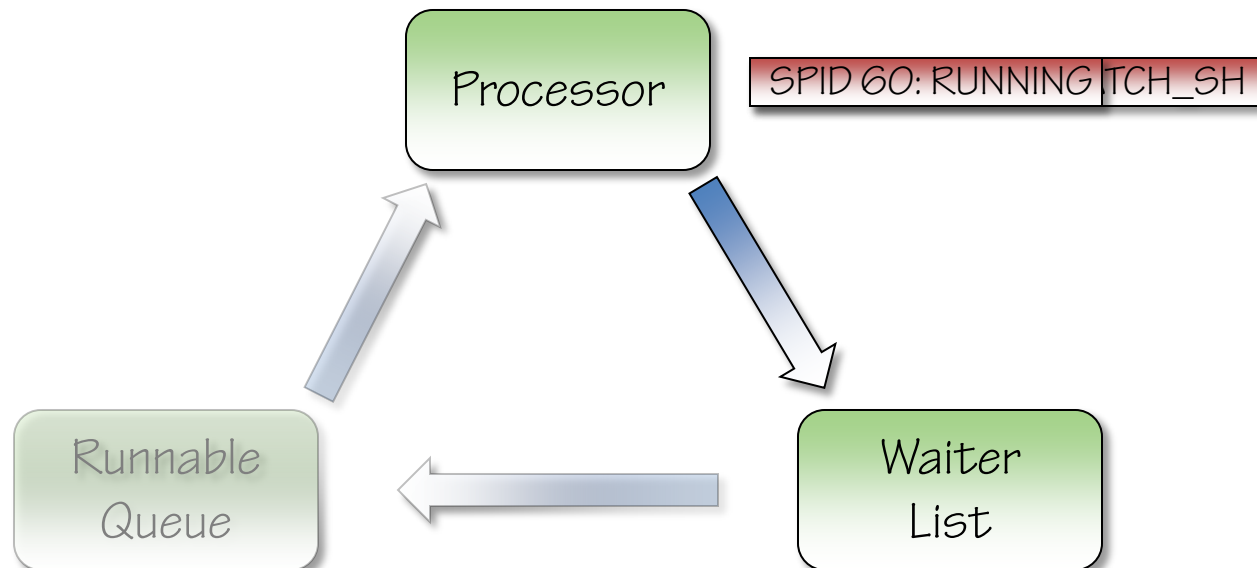


Thread States

- A thread can be in one of three states when being actively used as part of processing a query
- **RUNNING**
 - The thread is currently executing on the processor
- **SUSPENDED**
 - The thread is currently on the Waiter List waiting for a resource
- **RUNNABLE**
 - The thread is currently on the Runnable Queue waiting to execute on the processor
- Threads transition between these states until their work is complete

Transition from RUNNING to SUSPENDED

- A thread continues executing on the processor until it must wait for a resource to become available
 - The thread's state changes from RUNNING to SUSPENDED
 - The thread moves to the Waiter List
 - This process is called being 'suspended'

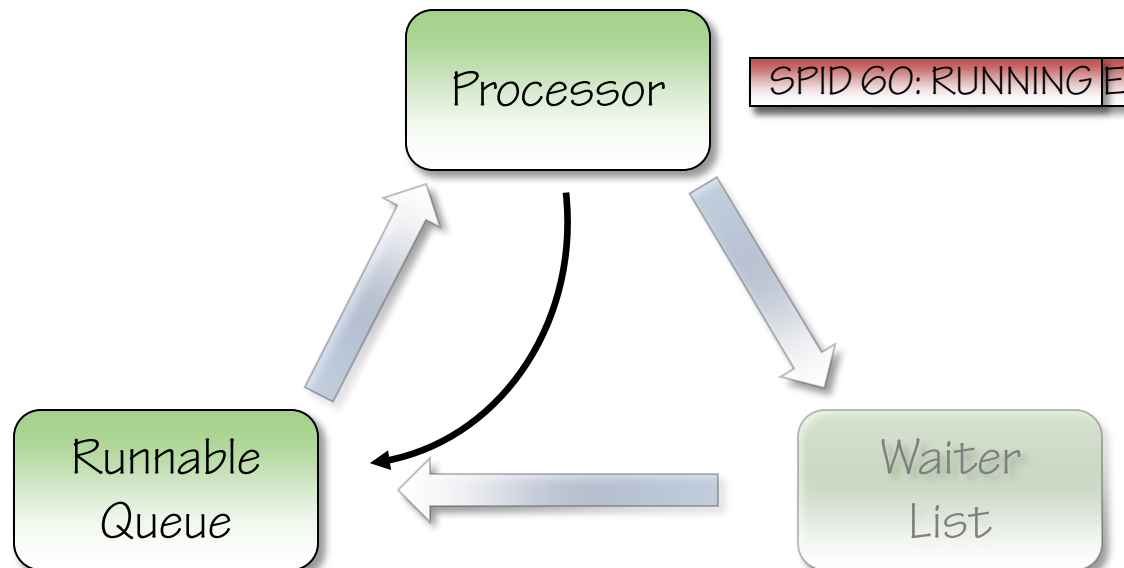


The Waiter List

- The Waiter List is an unordered list of threads that are suspended
- Any thread can be notified at any time that the resource it is waiting for is now available
 - Again, absolutely no ordering
- There is no limit to how long a thread remains on the waiter list
 - Although execution timeouts or lock timeouts may take effect
- There is no practical limit to how many threads can be on a scheduler's waiter list at any time
- The `sys.dm_os_waiting_tasks` DMV shows which threads are currently waiting and what they are waiting for
 - More details on this in Module 3

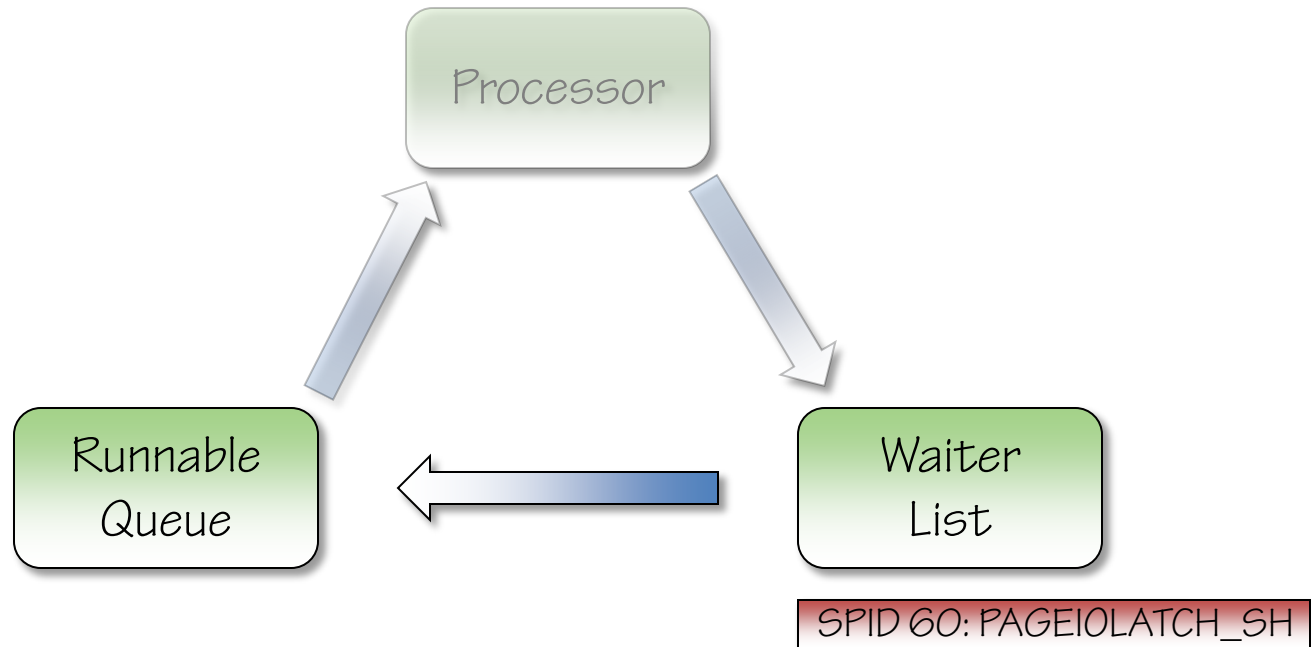
Special Case: Quantum Exhaustion

- If a thread does not need to wait for a resource, it will continue executing until its quantum is exhausted
 - Thread quantum is fixed at 4 milliseconds and cannot be changed
- If this occurs the thread moves to the bottom of the Runnable Queue
 - The thread's state changes from RUNNING to RUNNABLE
 - More information on this in Module 5



Transition from SUSPENDED to RUNNABLE

- A thread continues to wait until it is told that the resource is available
 - The thread's state changes from SUSPENDED to RUNNABLE
 - The thread moves to the bottom of the Runnable Queue
 - This process is called being 'signaled'



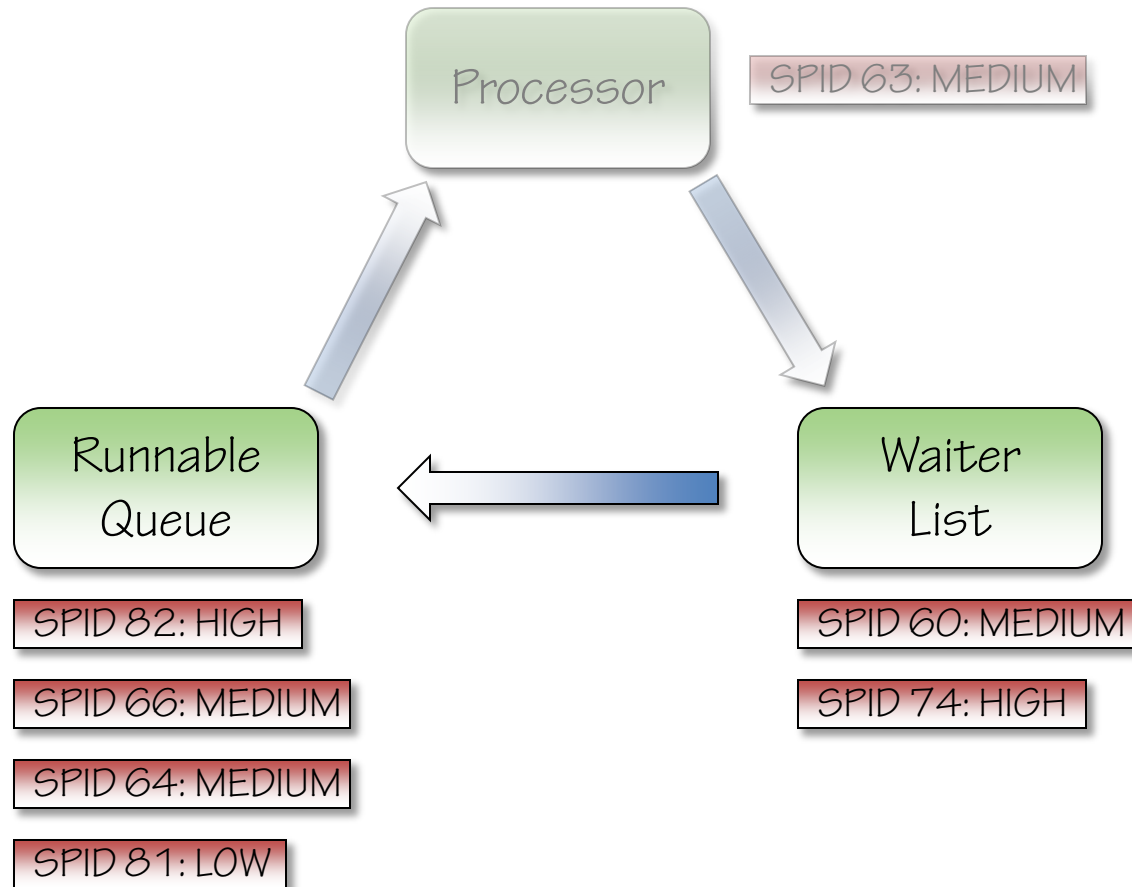
The Runnable Queue

- **The Runnable Queue is a strict First-In-First-Out (FIFO) queue**
 - There is a special case that will be discussed on the next slide
- **Threads enter the queue at the bottom and progress to the top**
- **The thread at the top of the queue is the one that will execute on the processor when the processor becomes free**
 - When the currently executing thread is suspended or exhausts its quantum
- **The size of the Runnable Queue can be seen from the `runnable_tasks_count` column in `sys.dm_os_schedulers`**

Special Case: Resource Governor

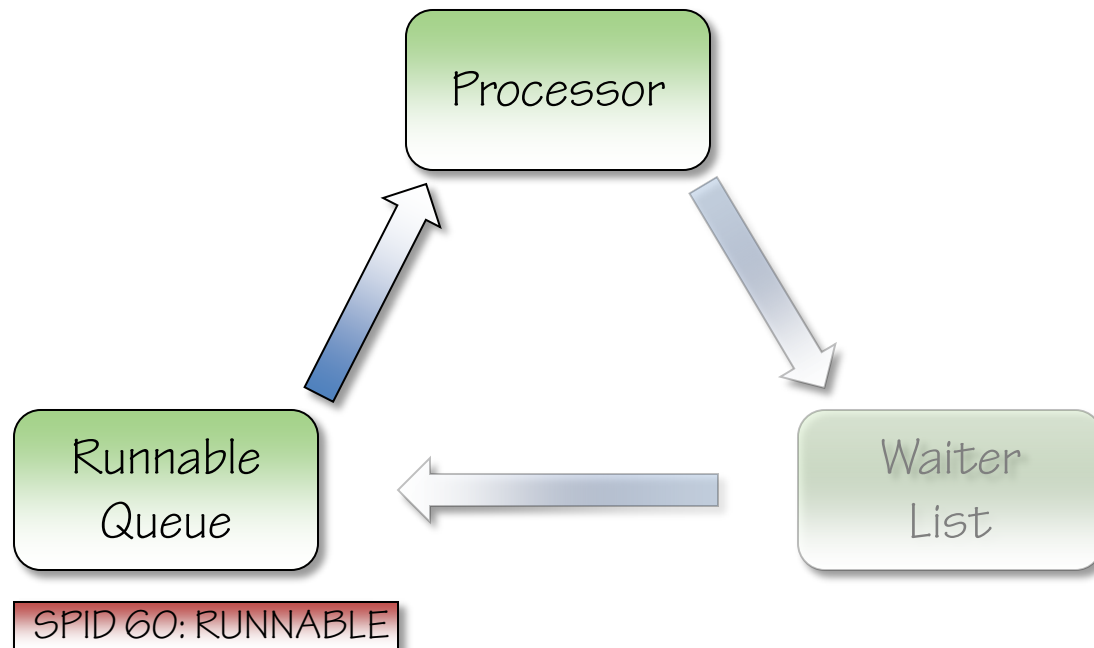
- **Using Resource Governor (in Enterprise Edition of SQL Server 2008 onwards) it is possible to define multiple workload groups that share the same resource pool**
 - A workload group can be thought of as a bucket of connections
 - A resource pool can be thought of as a set of CPU and memory limits, but more details are beyond the scope of this course
- **Multiple workload groups for a resource pool can be assigned relative priorities**
 - The default priority is medium
 - Possible values are high, medium, and low
- **The relative priorities change how the Runnable Queue works**
 - High to medium to low equates to 9 to 3 to 1 in terms of the priority of the threads that are permitted to execute
 - For example, for each low-priority thread on the runnable queue, nine high-priority threads will be allowed to jump over it and execute first

Resource Governor Example

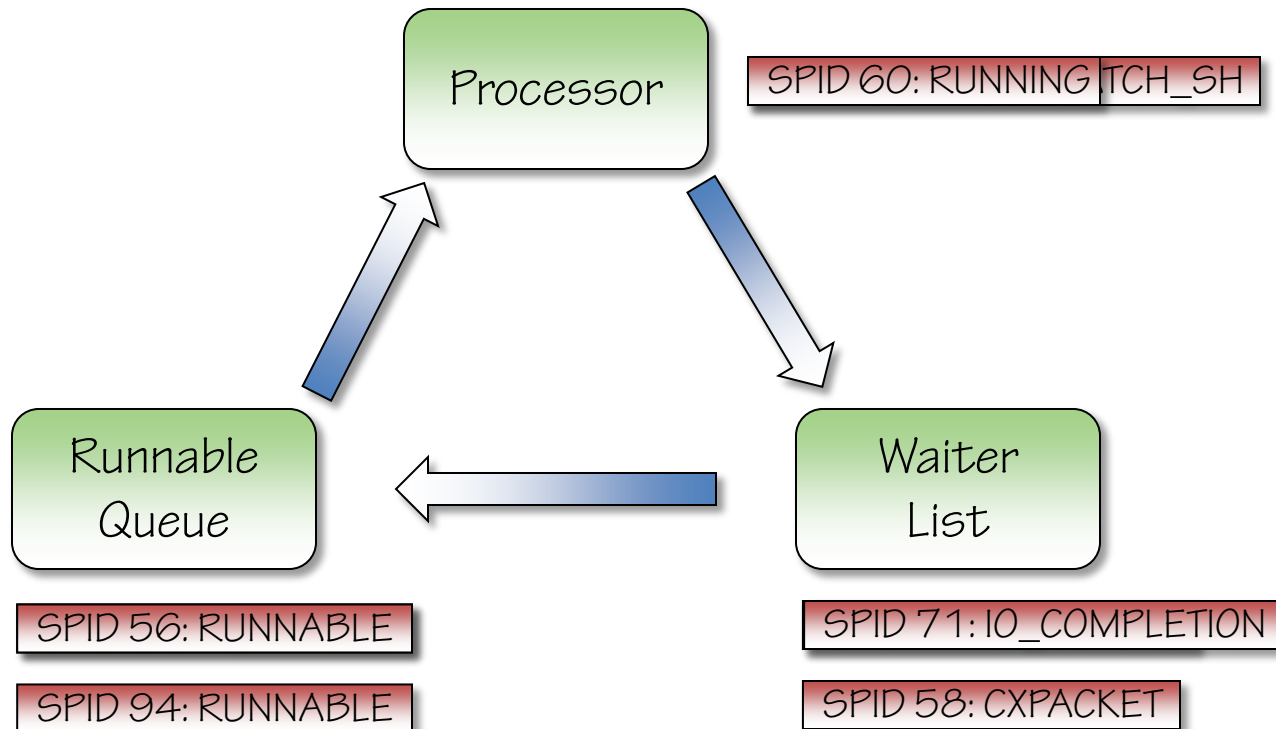


Transition from RUNNABLE to RUNNING

- The thread waits on the Runnable Queue until it reaches the top and the processor becomes available
 - The thread's state changes from RUNNABLE to RUNNING



Pulling It All Together



Summary

- SQL Server performs its own thread scheduling
- Threads move from RUNNING to SUSPENDED to RUNNABLE and back to RUNNING again until their work is completed
 - This can be altered somewhat using Resource Governor
- Analysis of the reasons for a thread going from RUNNING to SUSPENDED and how long a thread remains SUSPENDED is the basis of the waits and queues methodology
- The next module will define the various wait times and how to extract them from SQL Server