

SQL Server: Why Physical Database Design Matters

Module 3: Data Types and Index Size

Kimberly L. Tripp

Kimberly@SQLskills.com

<http://www.SQLskills.com/blogs/Kimberly>



pluralsight 
hardcore developer training

Introduction: Does Data Type Choice Impact Indexes?

- **Again, you might be thinking: disk space is cheap – who cares?...**
- **Once you understand the basic indexes structures, you'll see how profound the effect of your choice can be**
- **We'll discuss the impact of data type choice by index type:**
 - Clustered index
 - Nonclustered index
 - Columnstore index
- **Throughout this module, I'll be discussing/demonstrating:**
 - Key considerations around index structures
 - What the physical structures look like and how to analyze them
 - How does SQL Server access data (based on index usage)
 - What is the effect on performance

What Structures Exist for a Table?

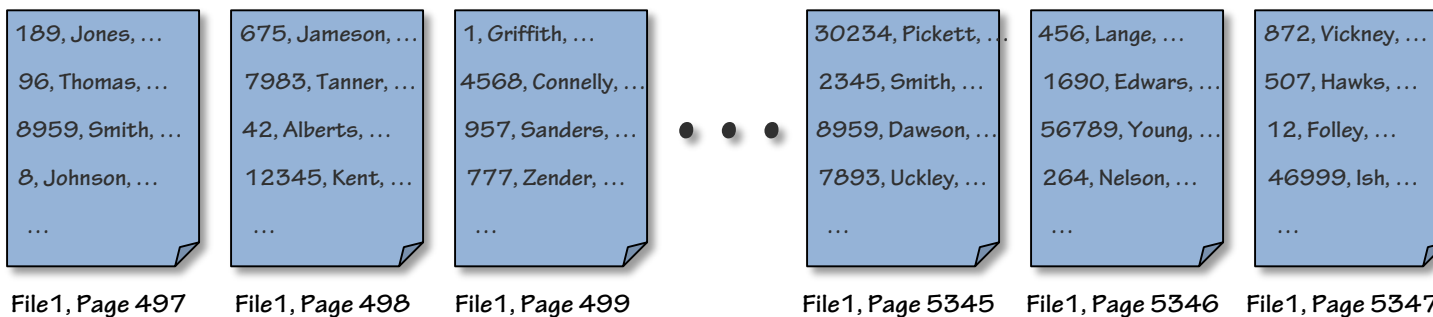
- **Table structure is either:**
 - Unordered: the table is called a heap
 - Ordered: through creation of a clustered index and the table is called a clustered table (not to be confused with other RDBMS' clustered table)
- **Indexes**
 - Clustered: only one can exist per table as this defines the data's order
 - Not *required*, but highly recommended
 - The data is physically ordered at create/rebuild
 - The data is logically ordered through a doubly-linked list
 - Cluster key choice is CRITICAL!
 - Nonclustered: not required, can have up to 999 of these (249 in 2000/2005)
 - DO NOT affect the base table's structure
 - Are affected by whether or not the table is clustered
 - **Hint: The nonclustered index dependency on the clustered index should impact your choice for the clustering key!**

What About Columnstore Indexes?

- Traditional clustered and nonclustered indexes are also known as “row-based” indexes
- Columnstore indexes are new in SQL Server 2012 and have a completely different internal structure as values for a single column are stored together
 - Data type has an impact but data distribution is even more interesting (in terms of possible column-level compression)
- These are beyond the scope of this course and have very specific/limited uses
 - SQL Server 2012: max of one nonclustered columnstore index per table and once created, the table is read-only
 - SQL Server 2014: columnstore indexes can be clustered and read/write
- Columnstore indexes will not be discussed here, check out Joe Sack’s course *SQL Server 2012: Nonclustered Columnstore Indexes*

Table Structure: Heap

- A table without a clustered index
- Records are not ordered and there is no doubly-linked list
- Accessed via allocation structures only so if no indexes exist then a full table scan is required for any SELECT query
- Imagine 80,000 records at 20 rows/per page = 4,000 pages
- Table scan costs at least 4,000 I/Os
 - Why “at least”?



*4,000 pages
of Employees
in no specific
order*

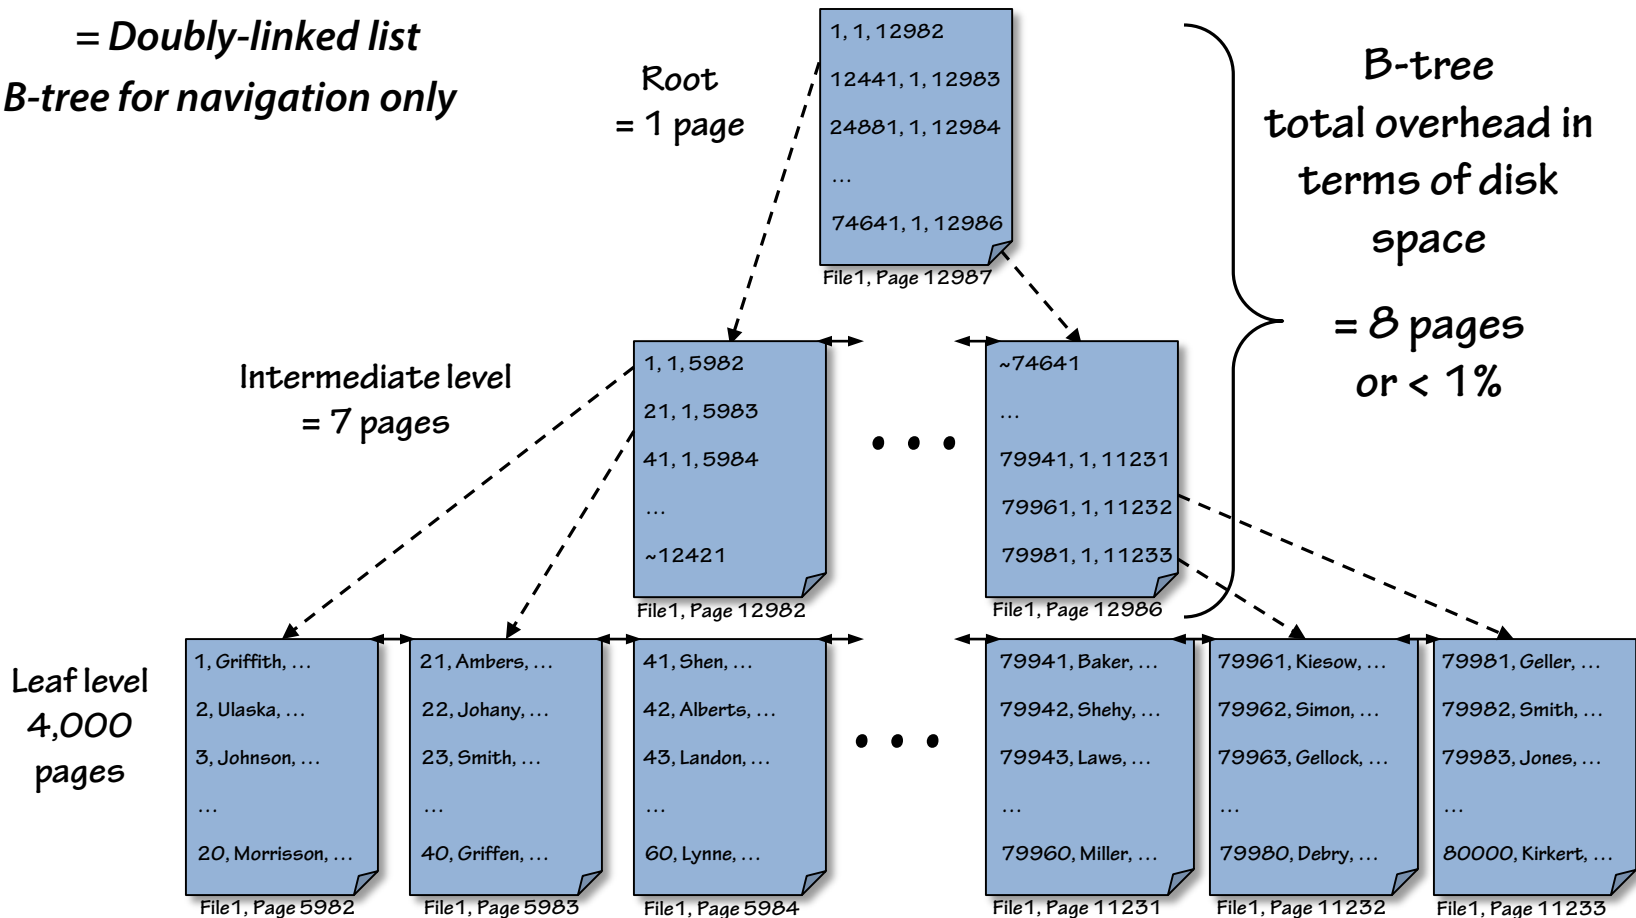
Table Structure: Clustered Table

Row Data

= Ordered leaf level

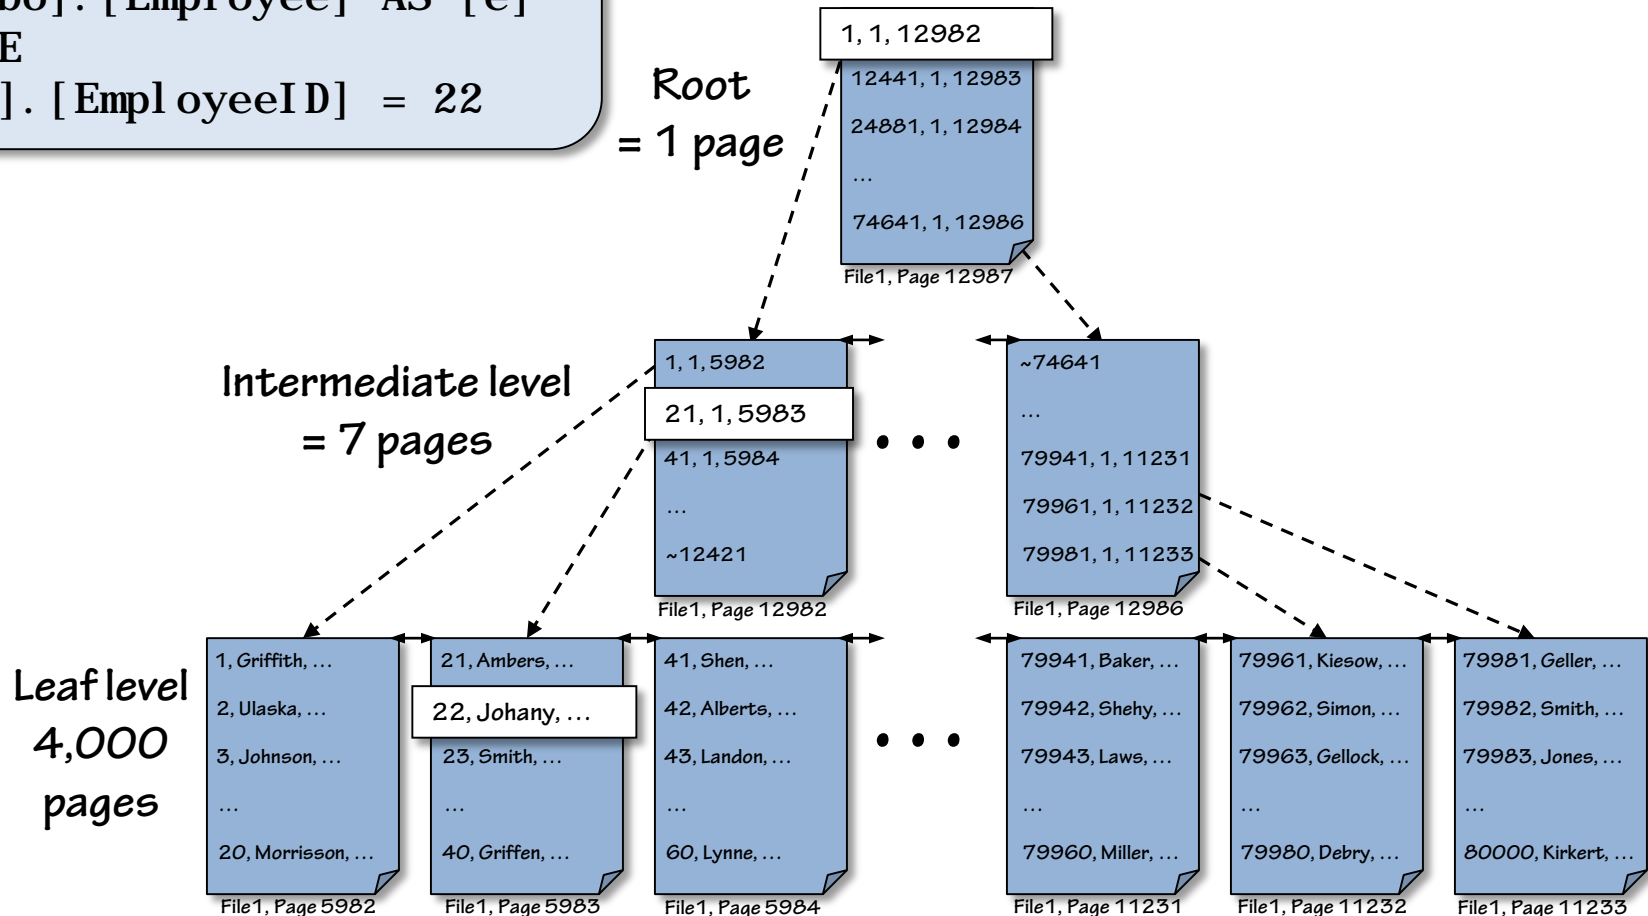
= Doubly-linked list

B-tree for navigation only



Accessing Data Using a Clustered Index

```
SELECT [e]. *  
FROM  
    [dbo].[Employee] AS [e]  
WHERE  
    [e].[EmployeeID] = 22
```




Nonclustered Indexes: The Book Analogy

- Think of a book with indexes in the back
- The book has one form of logical ordering
- To lookup data, you use the indexes in the back
- Using a “Common Name” you look up that value in the index
- Once you find that index value then, you need to lookup the actual data based on its page in the book... i.e. a “bookmark” lookup
- The bookmark always depends on the book’s content order

Index – Species Common Name

Index – Animal by Type, Name
Bird, Mammal, Reptile, etc...

Index – Animal by
Country, Name



Index – Animals by Habitat, Name
Air, Land, Water

Index – Species Scientific Name

Index – Animal by
Continent, Country, Name

Accessing Data Using a Nonclustered Index

- **Nonclustered indexes work a lot like an index in the back of a book:**
 - Nonclustered index leaf level is the data defined by the index, in sorted order (exactly the same as a book)
 - Nonclustered index b-tree doesn't exist in a book but does in SQL Server.
 - Its use is just like the b-tree in a clustered index, i.e. navigational only
- **Nonclustered indexes use the clustering key as the lookup value**
 - Instead of a page number (or, physical locator – which could change), SQL Server uses the clustering key as the lookup ID in a nonclustered index
 - Can be both good OR bad, depending on the key
- **Remember: the nonclustered index dependency on the clustered index should impact your choice for the clustering key!**
 - This is why...

Clustering Key Usage in Nonclustered Indexes

Imagine the internals of a nonclustered index on SocialSecurityNumber on three different versions of the Employee table each with a different clustering key

SSN	Lookup	Uniquifier
000-00-0184	Smith	0 (0 bytes)
000-00-0236	Jones	1 (4 bytes)
000-00-0395	Smith	1 (4 bytes)
000-00-0418	Jones	0 (0 bytes)

The lookup value is non-unique
(and wide as the column type of
nvarchar(40)).
Also, what if there are multiple
rows with the same lastname
(Smiths/Jones/Anderson)?

CL: Lastname

SSN	Lookup
000-00-0184	92CF41D7-17BF-49F7-B5C8-D3246C19B302
000-00-0236	2F87EEBB-FBA1-4C06-B7F1-BE63285B5935
000-00-0395	2EF09CA4-6E48-47AA-A688-3D9FDEA220EO
⋮	⋮

The lookup value
is a GUID = 16 bytes

CL: GUID

SSN	Lookup
000-00-0184	31101
000-00-0236	22669
000-00-0395	18705
⋮	⋮

The lookup value
is an int = 4 bytes

CL: EmployeeID

Each table starts at 80,000 rows over 4,000 pages (due to the average row size of 400 bytes/row and therefore 20 rows/page). Then EACH/EVERY index must include the (entire) lookup value.

Clustering Key Widens Nonclustered Indexes

- **Imagine a real-world scenario**
 - Table has 8 nonclustered indexes and 10 million rows
- **What's the overhead required (and total space) for the bookmark lookups in the nonclustered indexes:**
 - With a clustering key of an int (4 bytes)
 - With a clustering key of a GUID (16 bytes)
 - With a natural key (6 columns and ~64 bytes)
 - NOTE: This is just the overhead of the data type without factoring in nullable/non-unique.

Simple calculations for <u>overhead</u> in the LEAF level of the nonclustered indexes based on CL key columns defined		
CL Key Column(s)	Width of CL key (bytes)	MB
int	4	305.18
datetime	8	610.35
datetime, int	12	915.53
guid	16	1,220.70
composite	32	2,441.41
composite	64	4,882.81

Nonclustered Index Overhead

- Table has 8 nonclustered indexes and 10 million rows
- What is the required disk space for placing the clustering key in each and every nonclustered index
- Add required overhead for nullability as well as whether the column is unique vs. non-unique
- Keys of: int / bigint / datetime, int / GUID are likely to be unique and non-nullable (marked with *)

NOTE: Did not factor in additional overhead for composite keys and the number of variable-width columns they might have.

CL Key Column(s)	Bytes	MB
int *	4	305.18
int, nullable	7	534.06
int, non-unique (min)	4	305.18
int, non-unique (max)	12	915.53
int, non-unique (min), nullable	7	534.06
int, non-unique (max), nullable	15	1,144.41
bigint *	8	610.35
bigint, nullable	11	839.23
datetime, int *	12	915.53
datetime, int, nullable	15	1,144.41
guid *	16	1,220.70
guid, nullable	19	1,449.58
composite 32 bytes (comp32) *	32	2,441.41
comp32, nullable	35	2,670.29
comp32, non-unique (min)	32	2,441.41
comp32, non-unique (max)	40	3,051.76
comp32, non-unique (min), nullable	35	2,670.29
comp32, non-unique (max), nullable	43	3,280.64
composite 64 bytes (comp64) *	64	4,882.81
comp64, nullable	67	5,111.69
comp64, non-unique (min)	64	4,882.81
comp64, non-unique (max)	72	5,493.16
comp64, non-unique (min), nullable	67	5,111.69
comp64, non-unique (max), nullable	75	5,722.05
composite 128 bytes (comp128) *	128	9,765.63
comp128, nullable	131	9,994.51
comp128, non-unique (min)	128	9,765.63
comp128, non-unique (max)	136	10,375.98
comp128, non-unique (min), nullable	131	9,994.51
comp128, non-unique (max), nullable	139	10,604.86

Is it Really That Much Space?

- What about 100 million rows with 12 nonclustered indexes?

Simple calculations for <u>overhead</u> in the LEAF level of the nonclustered indexes based on CL key columns defined		
CL Key Column(s)	Width of CL key (bytes)	MB
int	4	4,577.64
bigint	8	9,155.27
datetime, int	12	13,732.91
guid	16	18,310.55
composite32, nullable	35	40,054.32
composite64, nullable	67	76,675.42
composite128, nullable	131	149,917.60

- You're looking at gigabytes of storage, memory, backups
- Insert/update performance (logging)
- Maintenance requirements
- My point is that it really does add up
- It is something you need to strategize/analyze and DESIGN!

Clustered Index Criteria

- **How do you keep your clustering key as streamlined as possible?**
 - **Unique**
 - Yes: No extra time/space overhead, data takes care of this criteria
 - NO: SQL Server must “uniquify” the rows on INSERT
 - **Static**
 - Yes: Reduces overhead
 - NO: Costly to maintain during updates to the key
 - **Narrow**
 - Yes: Keeps the nonclustered indexes narrow
 - NO: Unnecessarily wastes space
 - **Non-nullable/fixed-width**
 - Yes: Reduces overhead
 - NO: Adds overhead to ALL nonclustered indexes
 - **Ever-increasing key value**
 - Yes: Reduces index fragmentation
 - NO: Inserts/updates might cause significant index fragmentation

Choose a GOOD Clustering Key

- **Identity column**

- Adding this column and clustering on it can be extremely beneficial, even when you don't "use" this data

- **DateCol, identity**

- Composite key defined in that order
 - Do not use date alone as that would need to be "uniquified"
- Great clustering key for partitioned tables
- Ideal where you have a lot of data-related queries (even if not partitioned)

- **GUID**

- NO: if populated by client-side call to .NET client to generate the GUID
 - OK as the primary key but not as the clustering key
- NO: if populated by server-side NEWID() function
 - OK as the primary key but not as the clustering key
- Maybe: if populated by the server-side NEWSEQUENTIALID() function as it creates a more sequential pattern (and therefore less fragmentation)
 - But, this isn't really why you chose to use a GUID...

Primary Key does NOT have to be the Clustering Key

- **Primary key: relational integrity**
- **Clustering key: internal mechanism for looking up rows (bookmark lookup)**
- **SQL Server enforces uniqueness of a primary key through an index and defaults to clustered**
 - 1 clustered index per table
 - 1 primary key per table
- **If the primary key is a natural key then you probably want to enforce it with a nonclustered index**
 - Might be very wide
 - Very expensive to duplicate in each and every nonclustered index
- **If the table doesn't have a column (or small set of columns) that meets these criteria then consider adding a surrogate [identity] key and then create the clustered index on it**

Scenario: What is the Real Cost?

- AdventureWorksDW.dbo.FactInternetSales
- Clustered index (composite index of two seemingly narrow columns):
 - SalesOrderNumber type: **nvarchar(20)**
 - SalesOrderLineNumber type: tinyint
- Nonclustered indexes (all, single-column nonclustered):
 - IX_FactInternetSales_ShipDateKey: ShipDateKey
 - IX_FactInternetSales_CurrencyKey: CurrencyKey
 - IX_FactInternetSales_CustomerKey: CustomerKey
 - IX_FactInternetSales_DueDateKey: DueDateKey
 - IX_FactInternetSales_OrderDateKey: OrderDateKey
 - IX_FactInternetSales_ProductKey: ProductKey
 - IX_FactInternetSales_PromotionKey: PromotionKey
- Everything seems narrow and somewhat optimal, until we talk about the data types of the clustering key columns

Scenario: What's in That Key?

- **What? What does the data look like?**
 - SalesOrderNumber = 7 characters (SO12345) which is 14 bytes of data
 - EVERY row is SO + 5-digit number – why?
 - Then, add variable-width overhead (each column has 2 bytes in the variable block as an offset)
 - When this is the first (or only) variable-width column then the addition of a variable block within the row adds 2 more bytes
 - Each of these 7 character “numbers” requires 18 bytes (14 + 2 + 2)
- **If the clustering key requires it then EVERY nonclustered index requires it**
- **7 nonclustered indexes:**
 - Ironically, ALL columns (of all nonclustered indexes) are:
 - Non-nullable and fixed-width
 - They do not require a variable-block on their own, but when you add the clustering key

Scenario: What's the Total Cost?

- What's the physical cost of this poorly defined column: SalesOrderNumber vs. an int (and ditching the type since there's only one type in all of the data)
- 14 bytes wasted per row, per index
- 7 nonclustered indexes x 14 = 98 bytes (completely wasted) per row
- What if the table were larger?
 - Imagine 10 million rows and 10 nonclustered indexes:
 - $10,000,000 \times 140 / 1024 / 1024 = \mathbf{1.335\ GB}$ of nonsense
 - Imagine 100 million rows and 10 nonclustered indexes:
 - $100,000,000 \times 140 / 1024 / 1024 = \mathbf{13.35\ GB}$ of nonsense
 - Imagine 1 billion rows and 12 nonclustered indexes:
 - $1,000,000,000 \times 154 / 1024 / 1024 = \mathbf{143.42\ GB}$ of nonsense

Summary: The Effect of Data Type Choice

- Data type choice can have a profound affect on:
 - Column size
 - Row size
 - Index size
- ***Everything* is less efficient/effective when design is poor**
 - Waste disk space
 - Waste cache (and this is still very costly)
 - Larger backups, more time to backup
 - Logging is more costly with wider rows (DML is negatively affected)
 - Maintenance is more costly (and possibly required more)
 - Queries can be less efficient
 - Required to put data into memory than really needed
 - Returning more data than what's necessary