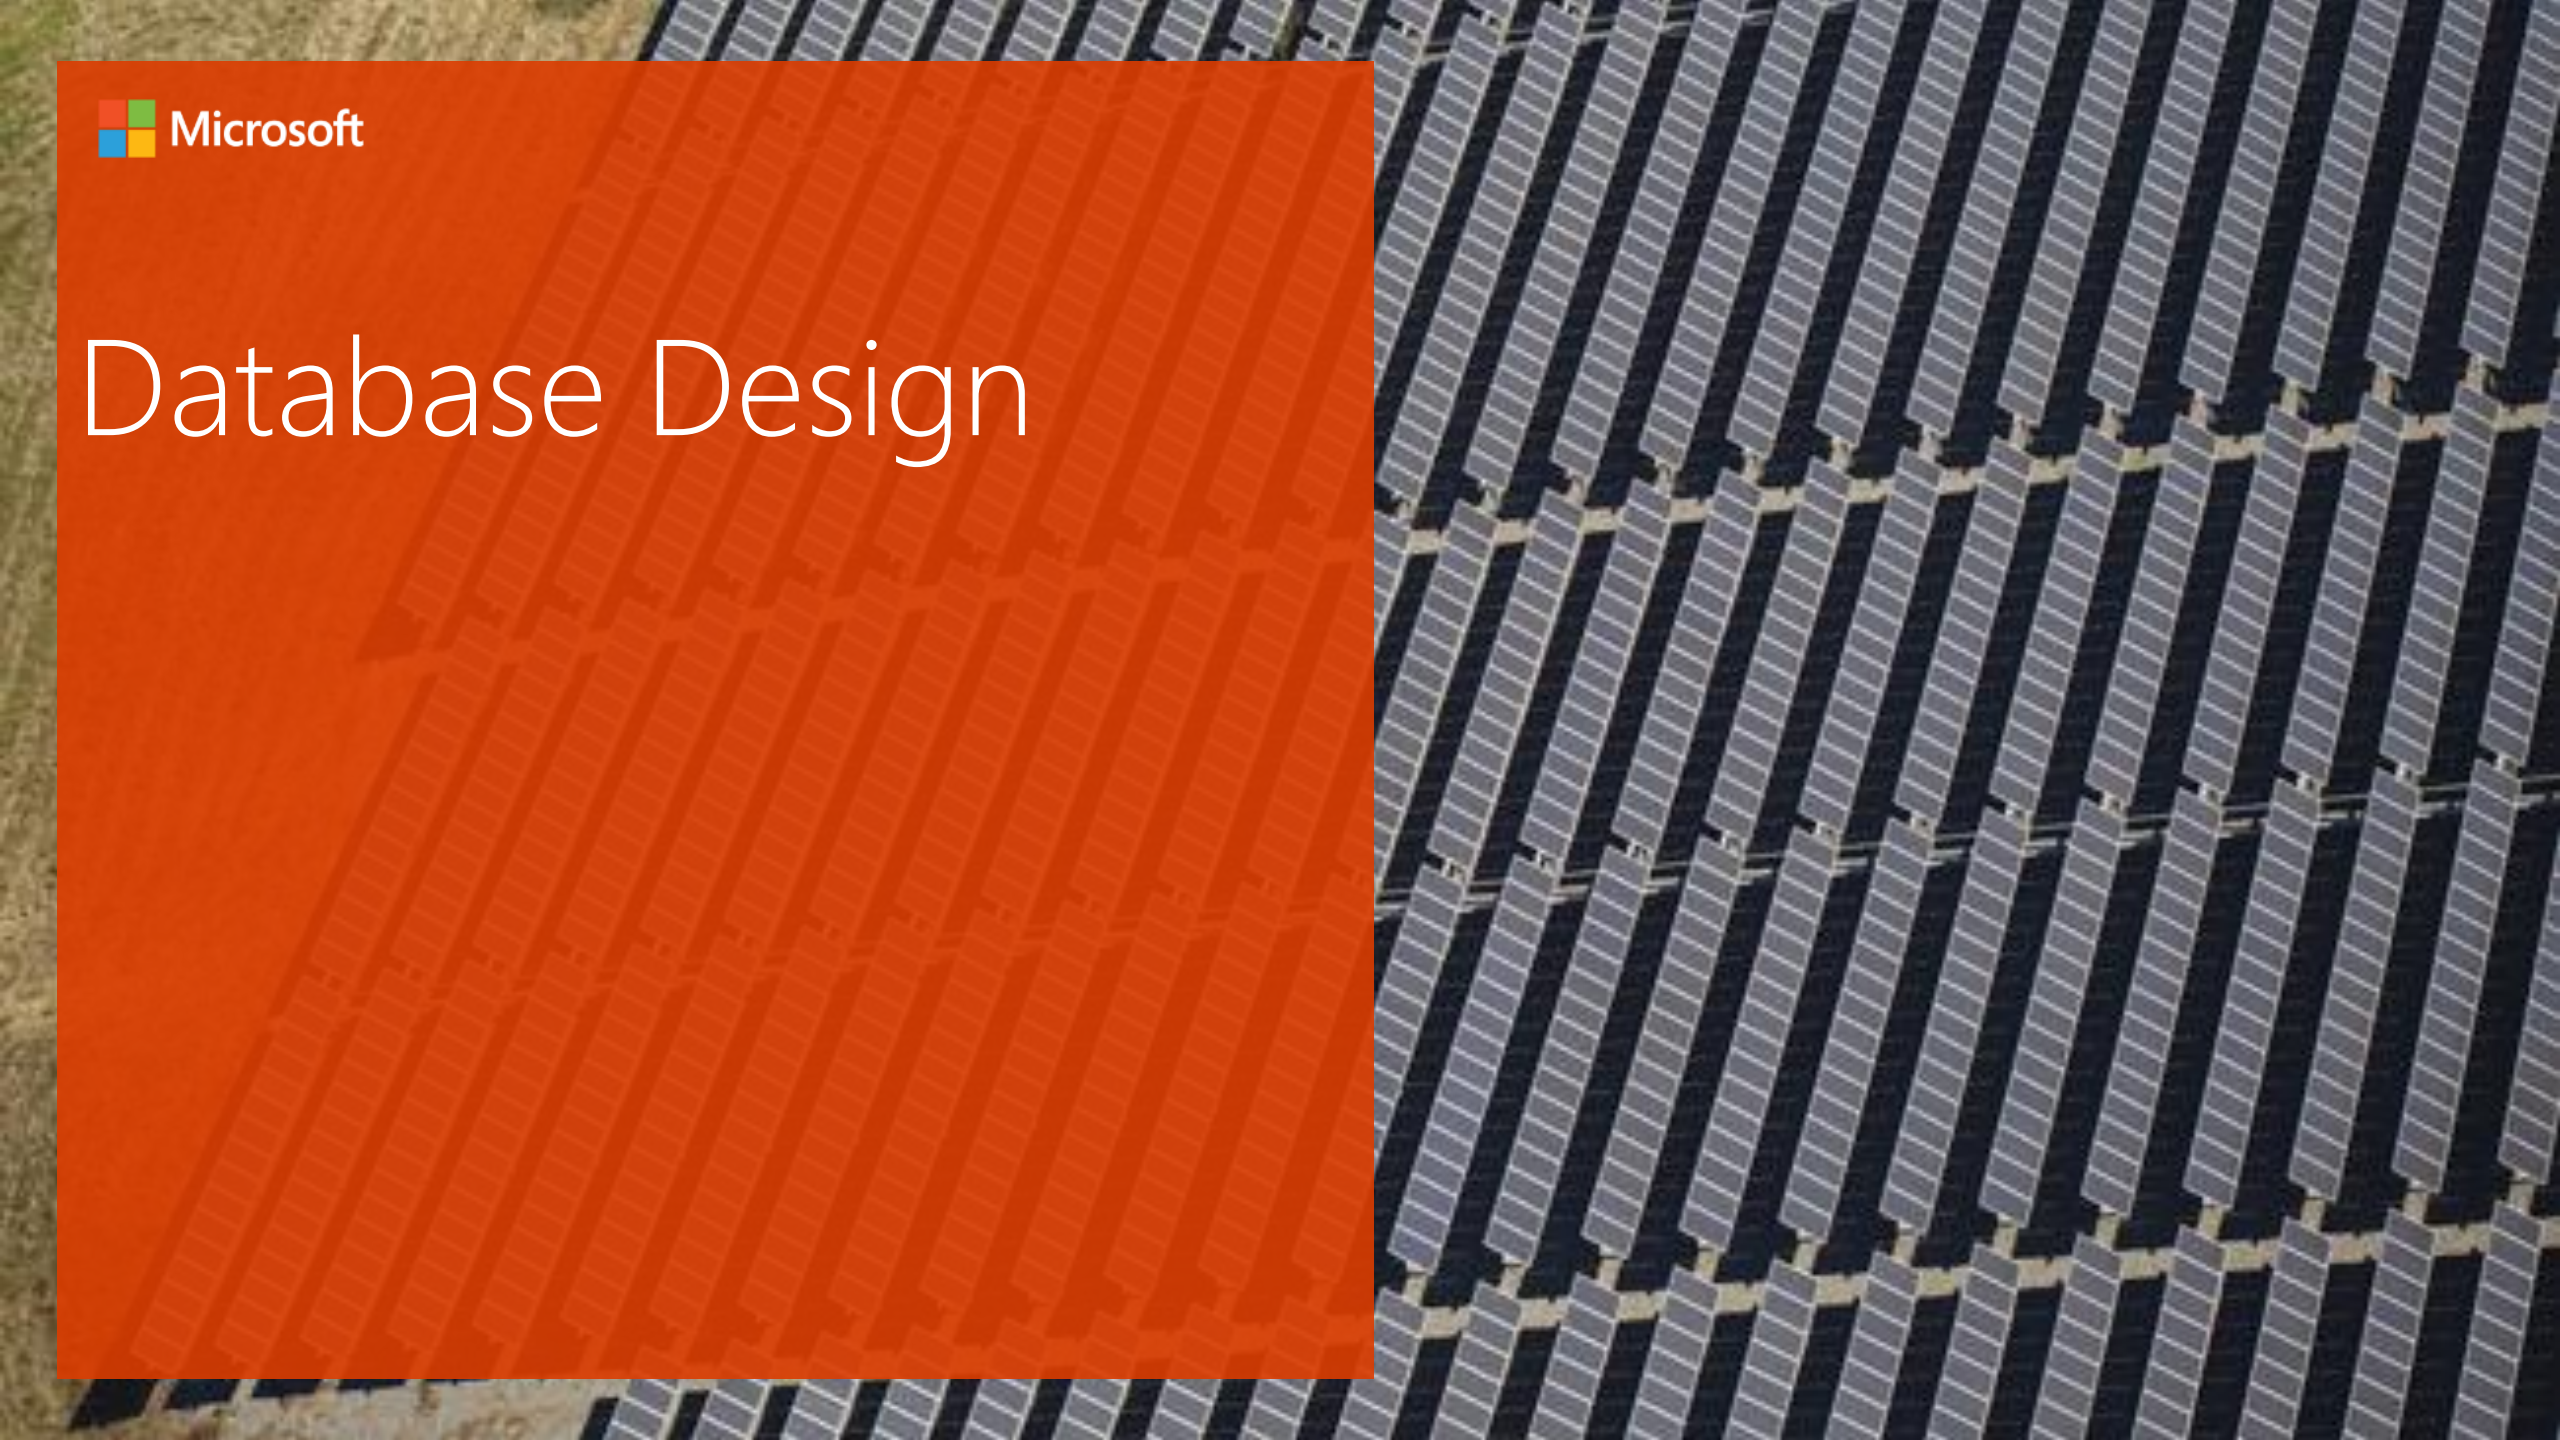
Microsoft
Database Design

Agenda

- DDL
- Creating a Database
- Understanding AUTOGROW
- Database Sizing
- Tempdb
- Transferring Ownership

Database DDL

Sample DDL

```
CREATE DATABASE PDW_2012
WITH
( AUTOGROW = ON
, DISTRIBUTED_SIZE = 1024.0GB
, REPLICATED_SIZE = 20GB
, LOG_SIZE = 10GB
);
DROP DATABASE PDW_2012;
```

DDL Points to Note

- File sizes specified as a decimal
- File size unit of measure is GB only
- Collation changes at database level not allowed
- Appliance default collation: Latin1_General_100_CI_AS_KS_WS
- AUTOGROW setting & allocation units cannot be seen via metadata

DDL Points to Note

Handled by PDW:

- Filegroup specifications
- File placement

Creating a Database

Creating a Database

User Creates Database

- PDW creates a logical shell database on the Control node
- PDW creates a physical application database on each Compute node

Total number of databases created
= # of Compute nodes + 1

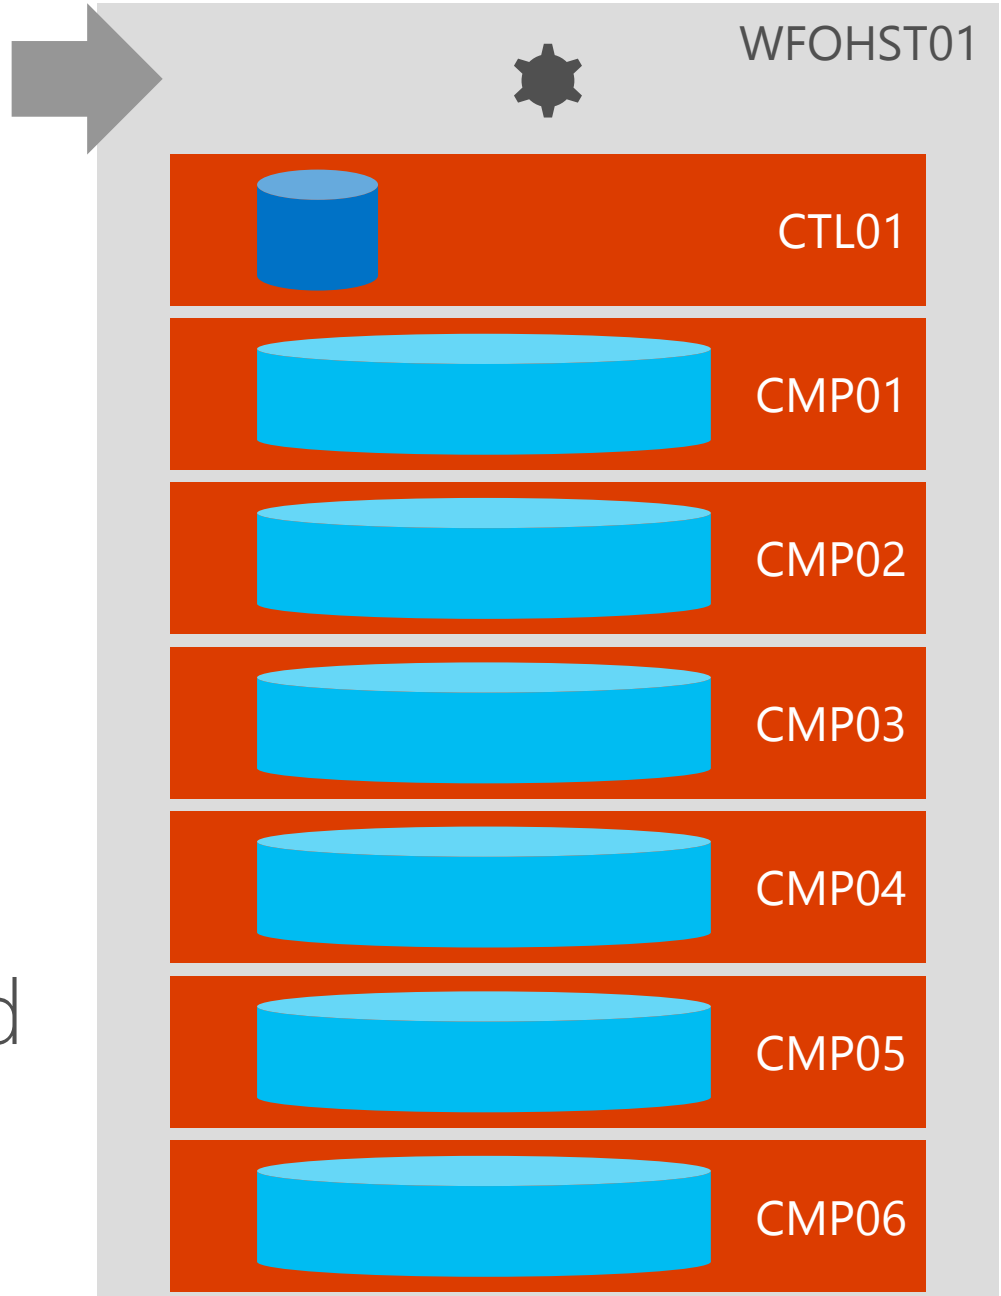

Control Node Database

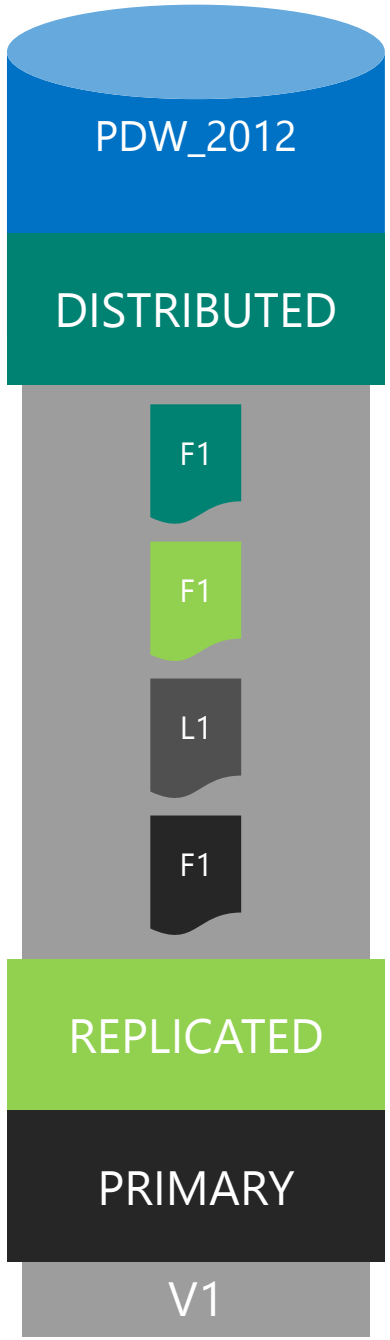

Database Extended Properties

Name	Value
pdw_distributed_size	Value specified in DDL
pdw_is_autogrow	True or False
pdw_is_valid_full_backup	True or False
pdw_log_size	Value specified in DDL
pdw_physical_name	DB_<GUID>
pdw_replicated_size	Value specified in DDL

Compute Node Database

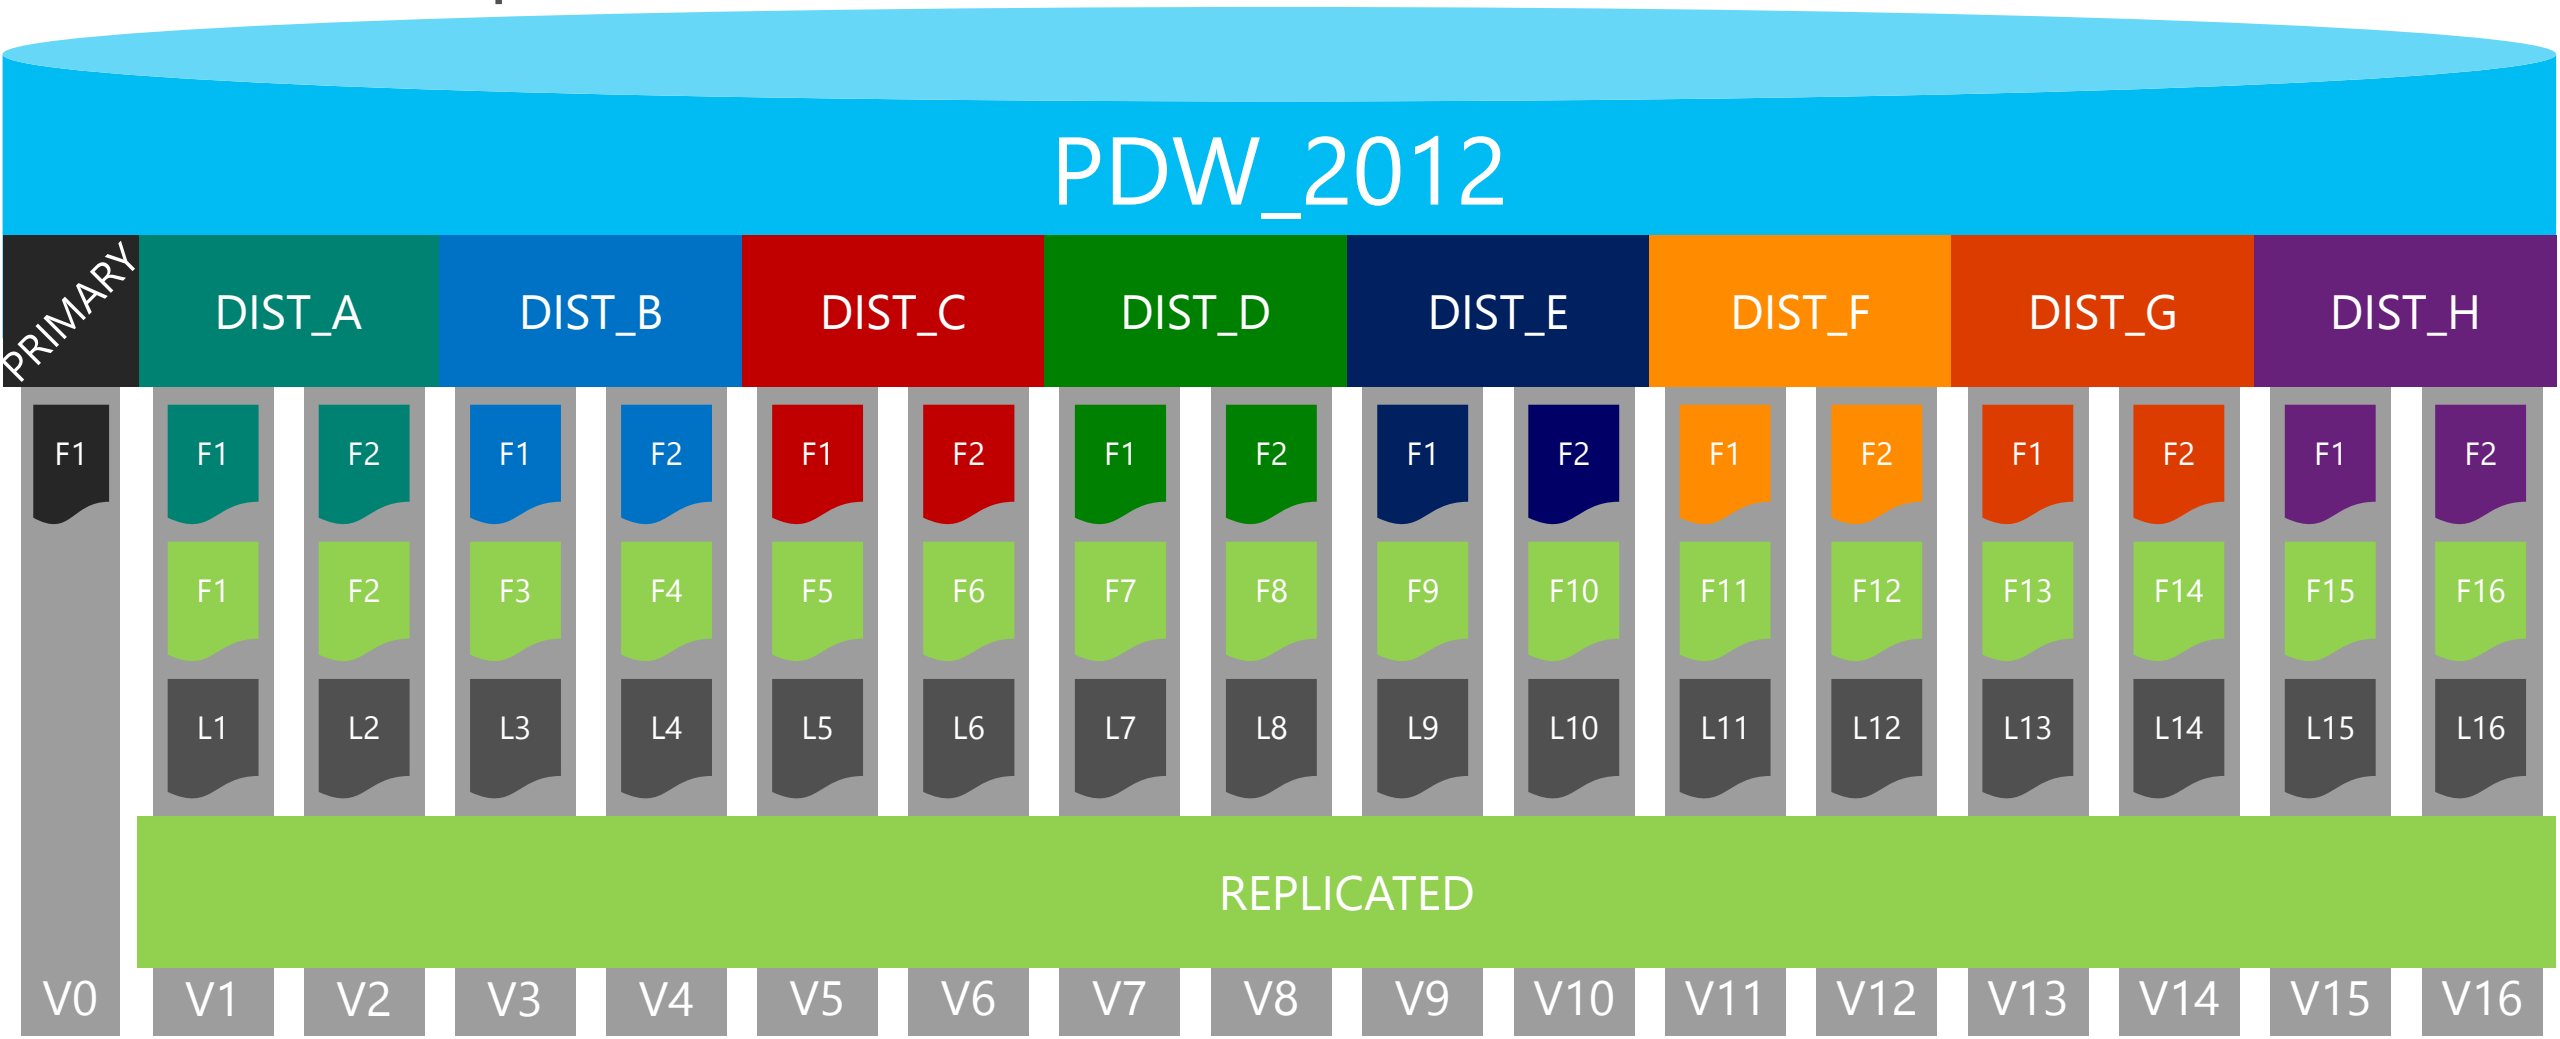

Physical Database Configuration

- Simple Recovery Model
- MAXDOP 4

Database Autogrow

AUTOGROW

- Is not discoverable from PDW catalog views
- sys.database_files shows configuration of the shell database on the Control node
- There is no sys.pdw_nodes_database_files catalog view
- Primary filegroup handled differently to other filegroups: always ON for primary

Compute Node AUTOGROW - ON

FILE	TYPE	VALUE	COMMENT
ROWS	max_size	-1	Each data file is not limited by size
ROWS	growth	512	Each data file grows by a fixed amount of 512 8KB pages (4MB)
ROWS	is_percent_growth	0	Table grows in fixed amounts not by percent
LOG	max_size	268435456	Each log file will grow to a max of 2TB
LOG	growth	10	Each log file will grow by 10%
LOG	is_percent_growth	1	Log file grows by percentage increments

Compute Node AUTOGROW - OFF

FILE	TYPE	VALUE	COMMENT
ROWS	max_size	-1	Not relevant as growth = 0
ROWS	growth	0	Data files do not grow
ROWS	is_percent_growth	0	Not relevant as growth = 0
LOG	max_size	268435456	Not relevant as growth = 0
LOG	growth	0	Log files do not grow
LOG	is_percent_growth	0	Not relevant as growth = 0

Sizing

Allocating Space – REPLICATED_SIZE

Base on a single copy of all replicated tables

- Each Compute node

Allocating Space – REPLICATED_SIZE

Allocation Amounts

Total storage

- Value of REPLICATED_SIZE* # of Compute nodes

Per replicated data file:

- Value of REPLICATED_SIZE/ 16

Allocating Space – DISTRIBUTED_SIZE

Divided evenly across the appliance

- Across Compute nodes
- Across distributions
- Across data files allocated to the distribution filegroups

Allocating Space – DISTRIBUTED_SIZE

Allocation Amounts

Per Compute node:

- Value of DISTRIBUTED_SIZE / # of Compute nodes

Per distribution:

- (Value of DISTRIBUTED_SIZE / # Compute nodes)/ 8

Per distributed data file:

- (Value of DISTRIBUTED_SIZE / # Compute nodes)/ 8/ 2

Allocating Size – LOG_SIZE

Divided evenly across the appliance

- Across Compute nodes
- Across data volumes

Allocating Size – LOG_SIZE

Amount allocated

Per Compute node:

- Value of LOG_SIZE/ # of Compute nodes

Per data volume:

- (Value of LOG_SIZE/ # of Compute nodes)/ 16

Re-sizing User Databases

ALTER DATABASE

- Sets one property at a time
- Can be smaller than original value
- Also used for switching AUTOGROW ON | OFF

DBCC SHRINKLOG(SIZE <target> | DEFAULT)

- Target combined size on all nodes

General Sizing Advice

Pre-size your Database!

- Enable Instant File Initialization
- Allocate sufficient capacity for 6-12 months
- Keep sufficient spare capacity to rebuild your largest table
- Create log twice the size of your largest transaction
- AUTOGROW ON – avoids 3a.m. alarm call

Example

```
CREATE DATABASE PDW_2012
WITH
( AUTOGROW = ON
, DISTRIBUTED_SIZE = 1024.0GB
, REPLICATED_SIZE = 20GB
, LOG_SIZE = 10GB
);
```

Q. How much storage has been allocated for distributed tables?

A. 1024GB

Per	Capacity (6 Nodes)
Compute Node	1.67GB
Log File	0.104GB

Guidance

Insufficient Capacity Consequences

Restricts ability to rebuild tables / indexes

Prevents automatic data re-distribution

- Extend appliance
- Restoring a database

Files using AUTOGROW get heavily fragmented

- Impacts performance

Remember Skew!

- File layout is important for understanding skew
- When data files for a distribution are full then the database is full
- When volume is full appliance is full

Data Loading & Sizing

- When loading data internal tables are created
- Internal tables hold a complete copy of the data being loaded
- You can nominate a specific database for these internal tables to isolate them from user database
- Failure to nominate a database will result in the internal tables being created inside target database
- Data load pattern can impact your database sizing!

General Loading Recommendations

- Create (and use) a staging database
- Size database to twice the size of largest load
- Size log to twice the size of largest load or batch
- If parallel loading load size includes all loads
- Staging database size likely to be different for initial/historic load and incremental load

FASTAPPEND load mode does not use internal tables.
Remember this for Data Loading Patterns segment!

Tempdb

Temporary Databases

Two Temporary Databases

- tempdb
- pdwtempdb

pdwtempdb

- User database
- Centralized resource for the appliance
- Holds user-defined temporary tables
- Isolates temporary tables into separate db
- Helps with sizing and fragmentation
- Does not have same optimisations as tempdb

pdwtempdb Database Sizing

Database Pre-sized at

- 1280MB per distributed data file
- 128MB per replicated file
- 64MB per log file
- Autogrows at 10% (unlimited)
- Re-created when appliance is re-started

Cannot be shrunk by DBCC SHRINKLOG or ALTER DATABASE

Contents visible in SSDT (tempdb)

pdwtempdb on 6 Compute Node PDW

6 Compute Nodes = 96 Data Volumes

pdwtempdb CREATE DATABASE size

- Distributed_Size = 122880MB or 120GB
- Replicated_Size = 12288MB or 12GB
- Log_Size = 6144MB or 6GB

tempdb

- System database created by SQL Server
- Exists on Compute and Control nodes
- Spread across all data volumes
- Used for data movement
- Tables created in tempdb are called Q tables
- Also used by normal SQL Server operations such as sorts and spills

tempdb on Control Node

One data file created
One log file created
Database Pre-sized at

- Data file 1000MB
  - autogrows at 1000MB (unlimited)
- Log file 100MB
  - autogrows at 100MB (unlimited)
- Re-created when appliance is re-started

Cannot be shrunk by DBCC SHRINKLOG or ALTER DATABASE

Contents visible in SSDT (tempdb)

tempdb on Compute Nodes

One file created per data volume (16 total)

Database Pre-sized at

- 1000MB per data file
  - Data file AUTOGROWs at 1000MB (unlimited)
- 100MB per log file
  - Log file AUTOGROWs at 100MB (limited 2097152MB)
- Re-created when appliance is re-started

tempdb on 6 Compute Node PDW

Control Node

- Data file = 1000MB
- Log file = 100MB

Compute Nodes (96 Volumes)

- Data files = 96000MB or 93.75GB
- Log files = 9600MB or 9.375GB

SSDT and tempdb

- tempdb shown in SSDT is a logical view
- Representation for the appliance
- Q Tables – tempdb
- #tables - pdwtempdb

Ownership

Database Ownership

- Login used to create the database is the owner of that database
- If login needs to be dropped then ownership must first be transferred

The sa account cannot be dropped or disabled in PDW. It is therefore a good option for database ownership

Transferring Database Ownership

- Use the ALTER AUTHORIZATION command

```
ALTER AUTHORIZATION ON DATABASE::
PDW_2012 TO sa;
```