



# Resolving Queries

Understanding Data Movement

# Agenda

- Understanding D-SQL Plans
- Understanding Data Movement
- Data Movement Service (DMS)
- Optimising DMS

# Understanding D-SQL plans

# Accessing the D-SQL plan

Before execution

- Use EXPLAIN

After execution

- DMV: sys.dm\_pdw\_request\_steps
- DMV: sys.dm\_pdw\_sql\_requests
- Management Console: Queries windows

# EXPLAIN

- PDW Specific syntax
- Returns the D-SQL plan (XML)
- Like an estimated execution plan but for MPP
- Sometimes called the MPP Plan
- Details the operations or steps required to resolve the query

# Example usage

EXPLAIN

SELECT p.\*

FROM dbo.FactInternetSales fis

LEFT JOIN dbo.DimCustomer p on fis.CustomerKey = p.CustomerKey

WHERE fis.OrderDateKey = 20040101

OPTION (LABEL = 'Shuffle : Join Cost')

;

EXPLAIN

CREATE TABLE dbo.DimEmployee

WITH (DISTRIBUTION = Hash(EmployeeKey))

AS

SELECT\*

FROM[AdventureWorksPDW2012].dbo.DimEmployee

OPTION (LABEL = 'Trim')

;

# Usage Guidance

Explain can be used before

- Select
- Insert
- Update
- Delete
- CTAS
- CRTAS
- CETAS

Explain cannot be used in conjunction with

- Variables
- Stored Procedures
- DDL

# D-SQL Operations

## SQL Operations

- Rnd\_ID
- On
- Return
- RemoteOn
- MetadataCreate

## DMS Operations

- Shuffle
- Broadcast
- Partition
- Move
- Trim
- DistributeReplicated
- Copy



# SQL Operations

## Rnd\_ID / Random ID

- Used to create unique names
- Most common usage : create names for temporary tables used in data movement

## On

- Used to perform an action on a database or an object
- Typical use case is to create a temporary table for data movement

## Return

- Used to return final result set back to end user
- Max of one return step in any MPP plan
- Return may also be used to perform the final aggregation

# DMS Operations

- Data Movement Service implements DMS operations
- Equate to data movement strategies
- Transparent to the user
- Optional steps – data doesn't have to move
- Fulfil data movement requests

# Typical Orchestration Steps

1. Random ID Operation
  - Create random name for temporary table
2. ON Operation
  - Create Temporary table using randomly generated name
3. DMS Operation
  - Move data into position to satisfy query request
4. Return Operation
  - Return results back to user
5. ON Operation
  - Drop Temporary Table

# Understanding Data Movement

# Why Data Moves

- Incompatible Join
- Incompatible Aggregation
- Re-distribute data
- Data Consistency
- Query syntax

# Joining Data

- Data is spread across an appliance
- Before a join can take place data needs to be co-located
- One or more sets of data may need to be re-distributed to enable the join
- Not all joins require data re-distribution

# Join Compatibility Matrix

Left Table	Right Table	Inner	Left	Right	Full	Cross
Replicated	Replicated					
Replicated	Distributed					
Distributed	Replicated					
Distributed	Distributed					

Conditions!

For a Distributed – Distributed join to be compatible (green) join must

- Contain distribution key of both columns
- Match data types on distribution keys
- Be an equality join

# Incompatible Join Example

EXPLAIN

```
SELECT      p.*  
FROM        dbo.FactInternetSales fis  
LEFT JOIN   dbo.DimCustomer p  
ON          fis.CustomerKey = p.CustomerKey  
WHERE       fis.OrderDateKey = 20040101  
OPTION      (LABEL = 'Shuffle : Join Cost');
```

- FactInternetSales distributed on OrderDateKey
- DimCustomer is distributed on Customer
- Therefore join incompatible
- Resolution: Move FactInternetSales. DimCustomer is distributed on the joining key
- N.B. Filter is applied to FactInternetSales prior to the move



# Aggregation Incompatibility

Data is aggregation incompatible when

- Data needs to be moved for a full aggregation to take place

Two approaches to resolve the incompatibility:

- Re-distribute data by a column in the group by
  - Keeps data down on the compute nodes
- Push data to a central point for aggregation
  - Uses the control node

# Incompatible Aggregation example

EXPLAIN

SELECT COUNT(\*)

FROM dbo.FactInternetSales fis

GROUP BY ProductKey

OPTION (LABEL = 'Shuffle : Aggregate');

- FactInternetSales distributed by OrderDateKey
- Query groups by product
- Therefore aggregation incompatible
- Resolution: Move FactInternetSales and re-distribute data on ProductKey
- N.B. Data is pre-aggregated by ProductKey prior to movement

# Re-distributing Data

Typically found when data is being persisted rather than returned to the user

You can move:

- From distributed to replicated
- From replicated to distributed
- From distributed (a) to distributed (b)

# Re-distribution Example

```
EXPLAIN
CREATE TABLE dbo.DimEmployee_dist
WITH (DISTRIBUTION = Hash(EmployeeKey))
AS
SELECT*
FROM[AdventureWorksPDW2012].dbo.DimEmployee
OPTION (LABEL = 'Trim')
;
```



Distributed



Replicated

- DimEmployee is a replicated table
- CTAS requests create DimEmployee\_dist hashed on EmployeeKey
- Therefore data needs to be re-hashed by Employee key
- Data also needs to be persisted in the new table
- N.B. Only data for this compute node needs to be kept

# Query Syntax

Depending on the distribution key

- OVER clause
- DISTINCT counts

... may trigger data movement...

Any expression on the distribution key will also trigger data movement

# What Data Moves

As little as is possible!

- Remove columns
  - Retain columns required for query resolution
- Remove rows
  - Apply where clause predicates
- Pre-aggregate Data
  - Group by the distribution key for partial aggregation
- Transport remote rows
  - Only send rows to other nodes that need to be stored remotely

# How Data Moves

Via Data Movement Service (DMS)

- Component of PDW
- Exists on Control and Compute nodes
- Responsible for all load & query data movement
- Controlled by the PDW Engine (DMSManager.dll)

# Data Movement Service



# DMS Functionality

## Primary functionality

- Re-distribute data across the compute nodes
- Centralize data to the control node
- De-centralizes data from control to compute
- Import and export data
  - Polybase
  - Remote Tables (export only)
- Load data

# DMS Components

- DMS Manager
- DMS Core Service

# DMS Manager

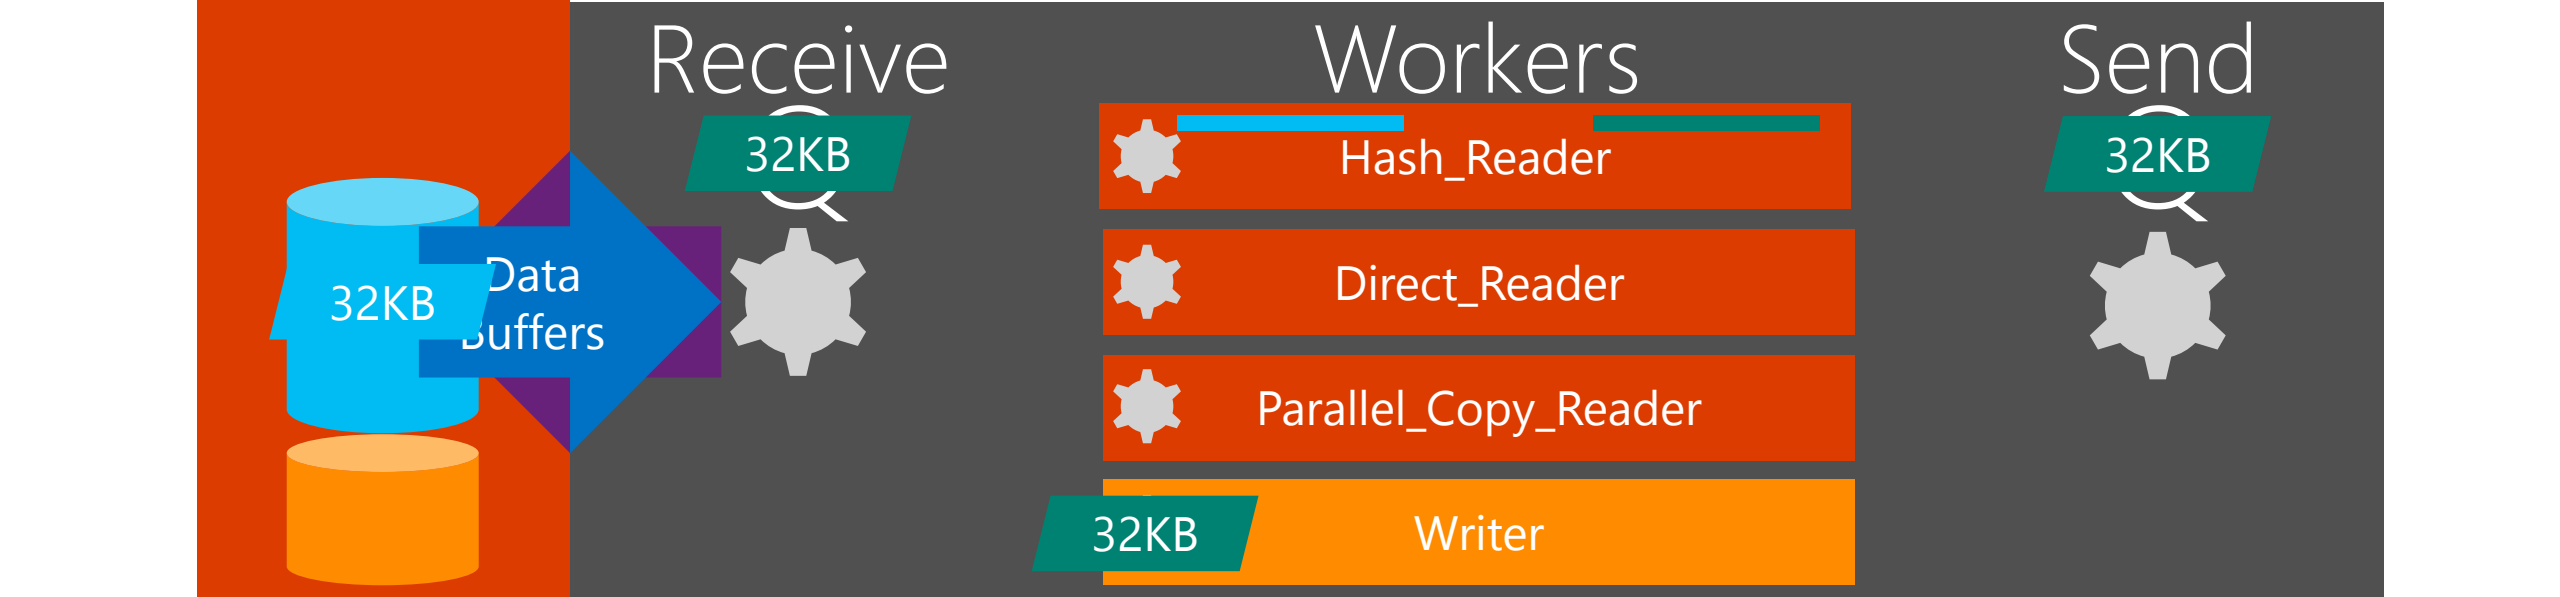
- Part of the PDWEngine service
- Exists only on the control node
- Takes instruction from the PDW Engine
- Communicates instructions to the DMS Core Services

# DMS Core Service

- Runs on control node and all compute nodes
- Receives instructions from the DMS manager
- Responsible for physically moving data
  - Reads the data from the tables in PDW
  - Moves the data to the new location
  - Writes the data to the target

# DMS Concepts

- Buffers
- Data Channel
- Queues
- Workers

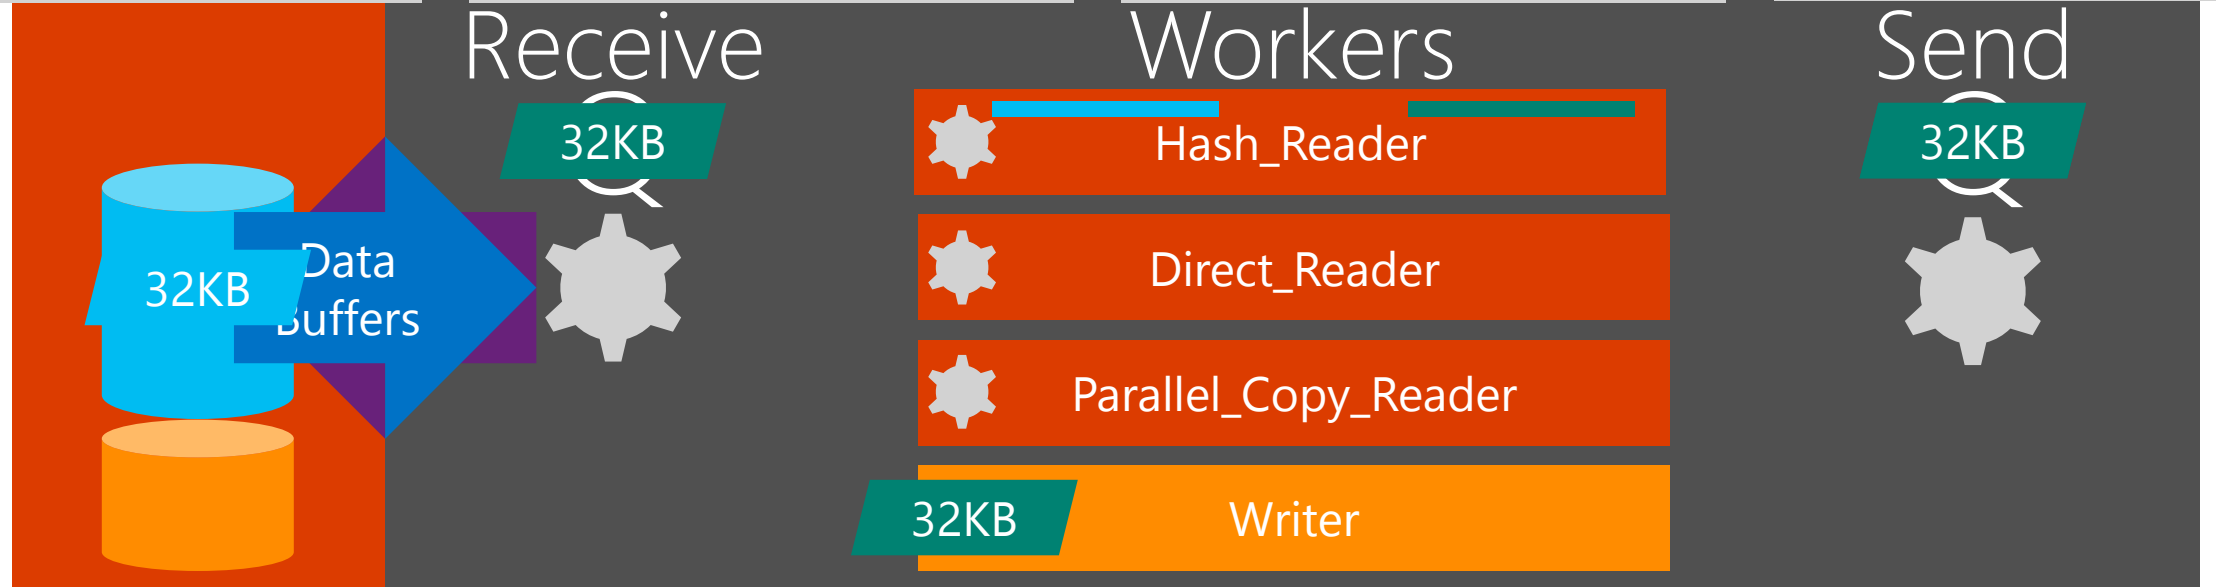


Step 1.  
Execute query &  
receive buffers

Step 2.  
DMS Reader pulls  
buffer & processes

Step 3.  
Sender ships new  
buffer to target node

Step 4.  
Writer bulk loads  
data into target table



# DMS Buffers

- Transports data between nodes using buffers
- Buffers are capped in size and are formatted to native ODBC types
- When loading DMS uses 256KB buffers to ingest data
- Majority of the time DMS uses 32KB buffers
- A row must be able to fit inside a 32KB buffer

# Data Channels

- DMS transports buffers over data channels
- Data channels instantiated when PDW started
- Shared channels are created between each of the DMS Core processes
- Binary & Reject Channels are created for load
- Default data channel for movement
- Sender threads transmit buffers
- Receiver threads receive buffers delivering them to appropriate workers

## Example

- 6 Compute Nodes = 7 DMS Core Services
- $6 \times 7 = 42$  connections in each channel



# Queues

Each DMS Core Service operates two queues

- Receive Queue
- Send Queue

Why? De-coupling

- Reader worker tasks de-coupled from data channel sender tasks
- Data channel receiver tasks de-coupled from the writer worker tasks

# Workers

- Hash\_Reader
  - Used by Shuffle, Trim
- Direct\_Reader
  - Used by Broadcast, Move, DistributeReplicated
- Parallel\_Copy\_Reader\*
  - Used by Partition, Remote Table Copy
- Writer
  - Used for all writes

# Reader Workers

- Acquires data from the source
- Process the data inside the buffer row by row
- Hash the distribution key if required
- Write the row to the appropriate send buffer

# Writer Workers

- One writer worker per target table
- If the table is a heap then write can also use BU locks and invoke multiple threads
- Takes work from the appropriate Receive Queue
- Invokes ODBC Bulk API to load data into target

# Data Type Handling

- Pre-dominantly handled via SQL Native Client
  - ODBC and ODBC BCP api
  - Fast with minimal type conversion
- Some legacy code paths use ADO.NET
  - Results in Native to managed datatype overheads
  - Slower with greater overheads
  - Used by SSIS destination adaptor

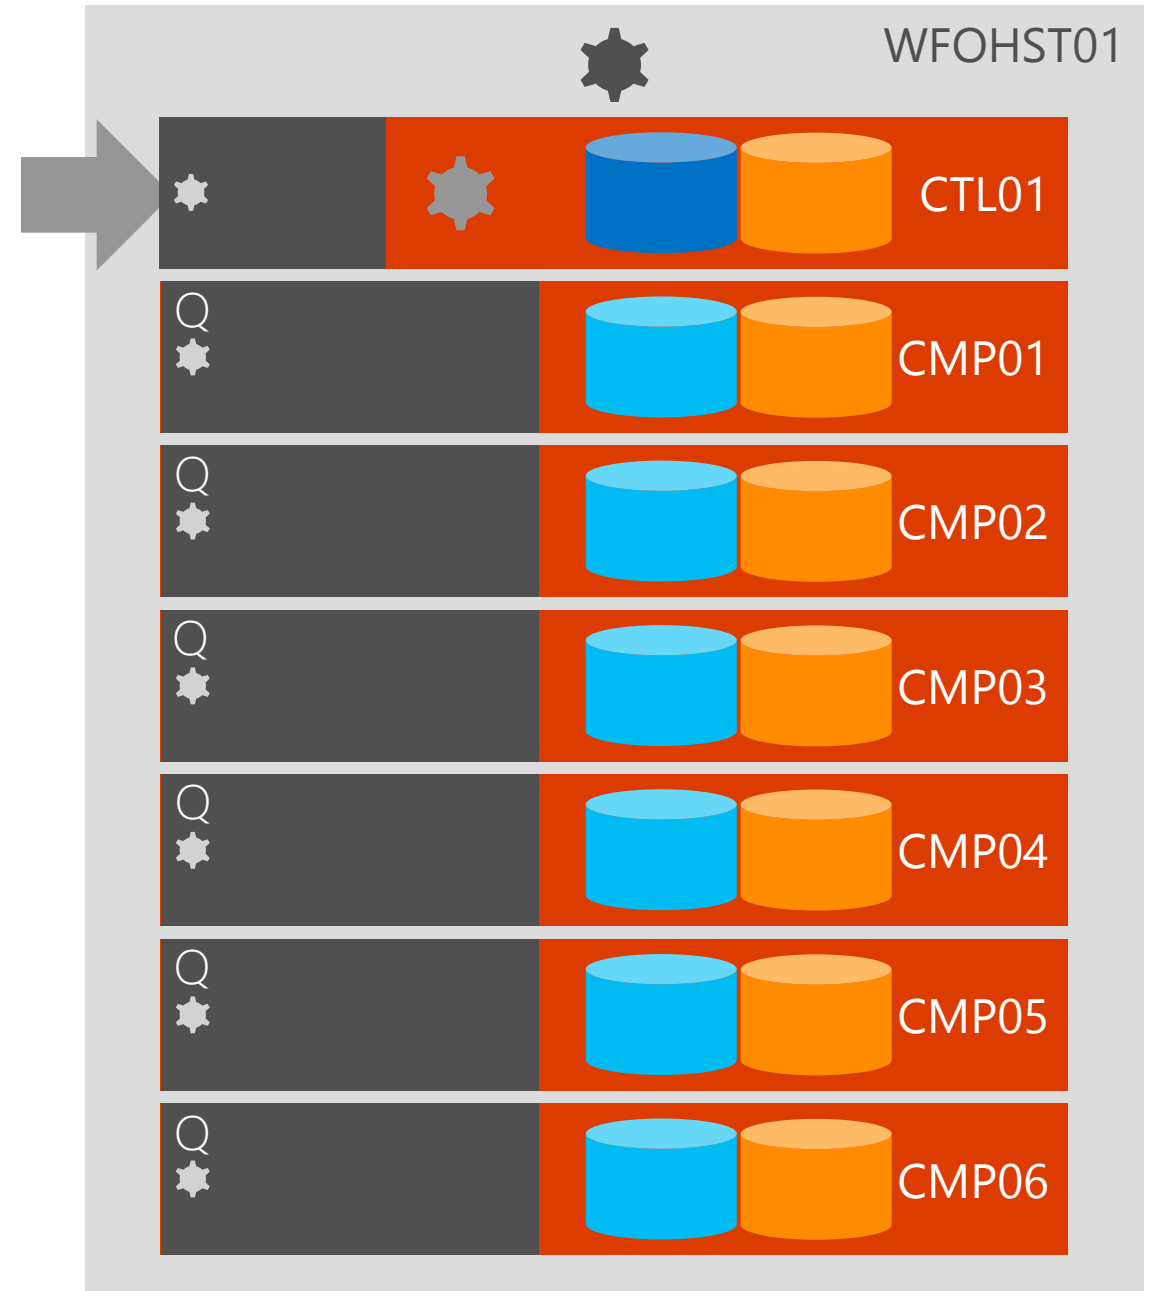
# DMS Operations

DMS moves data via the following operations

- Shuffle
- Broadcast
- Partition Move
- Master Move
- Trim
- DistributeReplicated

# Key

-  Tempdb
-  Shell database
-  Compute database
-  PDW Engine Service
-  DMS Service
-  PDW Cluster



# Shuffle Move Scenario



# Introducing the Shuffle

- Re-distribute data on a different key than the current distribution
- Where possible data will be filtered and aggregated prior to movement
- Join must equal true for a Shuffle to be used
  - Cannot have  $A=B$  OR  $B=C$

# Explain Plan

```
<?xml version="1.0" encoding="utf-8"?>
```

```
<dsql_query>
```

```
<sql>SELECT COUNT(*)
```

```
FROM dbo.FactInternetSales fis
```

```
GROUP BY ProductKey
```

```
OPTION (LABEL = 'Shuffle : Aggregate')</sql>
```

```
<dsql_operations total_cost="0.20019264" total_number_operations="5">
```

```
<dsql_operation operation_type="RND_ID">
```

```
<identifier>TEMP_ID_90858</identifier>
```

```
</dsql_operation>
```

```
<dsql_operation operation_type="ON">
```

```
<location permanent="false" distribution="AllDistributions" />
```

```
<sql_operations>
```

```
<sql_operation type="statement">CREATE TABLE [tempdb].[dbo].[TEMP_ID_90858]
```

```
</sql_operations>
```

```
</dsql_operation>
```

```
<dsql_operation operation_type="SHUFFLE_MOVE">
```

```
<operation_cost cost="0.20019264" accumulative_cost="0.20019264" average_rowsize="12" output_rows="521.335" />
```

```
<source_statement>SELECT [T1_1].[ProductKey] AS [ProductKey],
```

```
[T1_1].[col] AS [col]
```

```
FROM (SELECT COUNT_BIG(CAST ((0) AS INT)) AS [col], <shuffle_columns>ProductKey;</shuffle_columns>
```

```
[T2_1].[ProductKey] AS [ProductKey]
```

```
FROM [Instructor].[dbo].[FactInternetSales] AS T2_1
```

```
GROUP BY [T2_1].[ProductKey]) AS T1_1</source_statement>
```

```
<destination_table>[TEMP_ID_90858]</destination_table>
```

```
<shuffle_columns>ProductKey;</shuffle_columns>
```

```
</dsql_operation>
```

```
<dsql_operation operation_type="RETURN">...</dsql_operation>
```

```
<dsql_operation operation_type="ON">...</dsql_operation>
```

```
</dsql_operations>
```

```
</dsq
```

Random  
Table name

Temp Table

Shuffle

Distribution  
Key

Shuffle  
Target

# Explain Plan

```
<?xml version="1.0" encoding="utf-8"?>
<dsql_query>
  <sql>SELECT COUNT(*)
FROM dbo.FactInternetSales fis
GROUP BY ProductKey
OPTION (LABEL = 'Shuffle : Aggregate')</sql>
  <dsql_operations total_cost="0.20019264" total_number_operations="5">
    <dsql_operation operation_type="RND_ID">...</dsql_operation>
    <dsql_operation operation_type="ON">...</dsql_operation>
    <dsql_operation operation_type="SHUFFLE_MOVE">...</dsql_operation>
    <dsql_operation operation_type="RETURN">
      <location distribution="AllDistributions" />
      <select>SELECT [T1_1].[col] AS [col]
FROM   (SELECT CONVERT (INT, [T2_1].[col], 0) AS [col]
      FROM   (SELECT ISNULL([T3_1].[col], CONVERT (BIGINT, 0, 0)) AS [col]
            FROM   (SELECT SUM([T4_1].[col]) AS [col]
                  FROM   [tempdb].[dbo].[TEMP_ID_90858] AS T4_1
                  GROUP BY [T4_1].[ProductKey]) AS T3_1) AS T2_1) AS T1_1</select>
    </dsql_operation>
    <dsql_operation operation_type="ON">
      <location permanent="false" distribution="AllDistributions" />
      <sql>DROP TABLE [tempdb].[dbo].[TEMP_ID_90858]</sql>
    </dsql_operation>
  </dsql_operations>
</dsql_query>
```

The diagram illustrates the execution flow of the SQL query using the XML plan. Arrows point to specific operations:

- Return** (Red arrow) points to the `<dsql_operation operation_type="RETURN">` tag.
- Final Aggregation** (Grey arrow) points to the `<select>` block within the `<dsql_operation operation_type="RETURN">` tag.
- On** (Blue arrow) points to the `<dsql_operation operation_type="ON">` tag.
- Shuffle Target** (Orange arrow) points to the `<location permanent="false" distribution="AllDistributions" />` tag.
- Drop Q Table** (Green arrow) points to the `<sql>DROP TABLE [tempdb].[dbo].[TEMP_ID_90858]</sql>` tag.

# Shuffle – Part 1

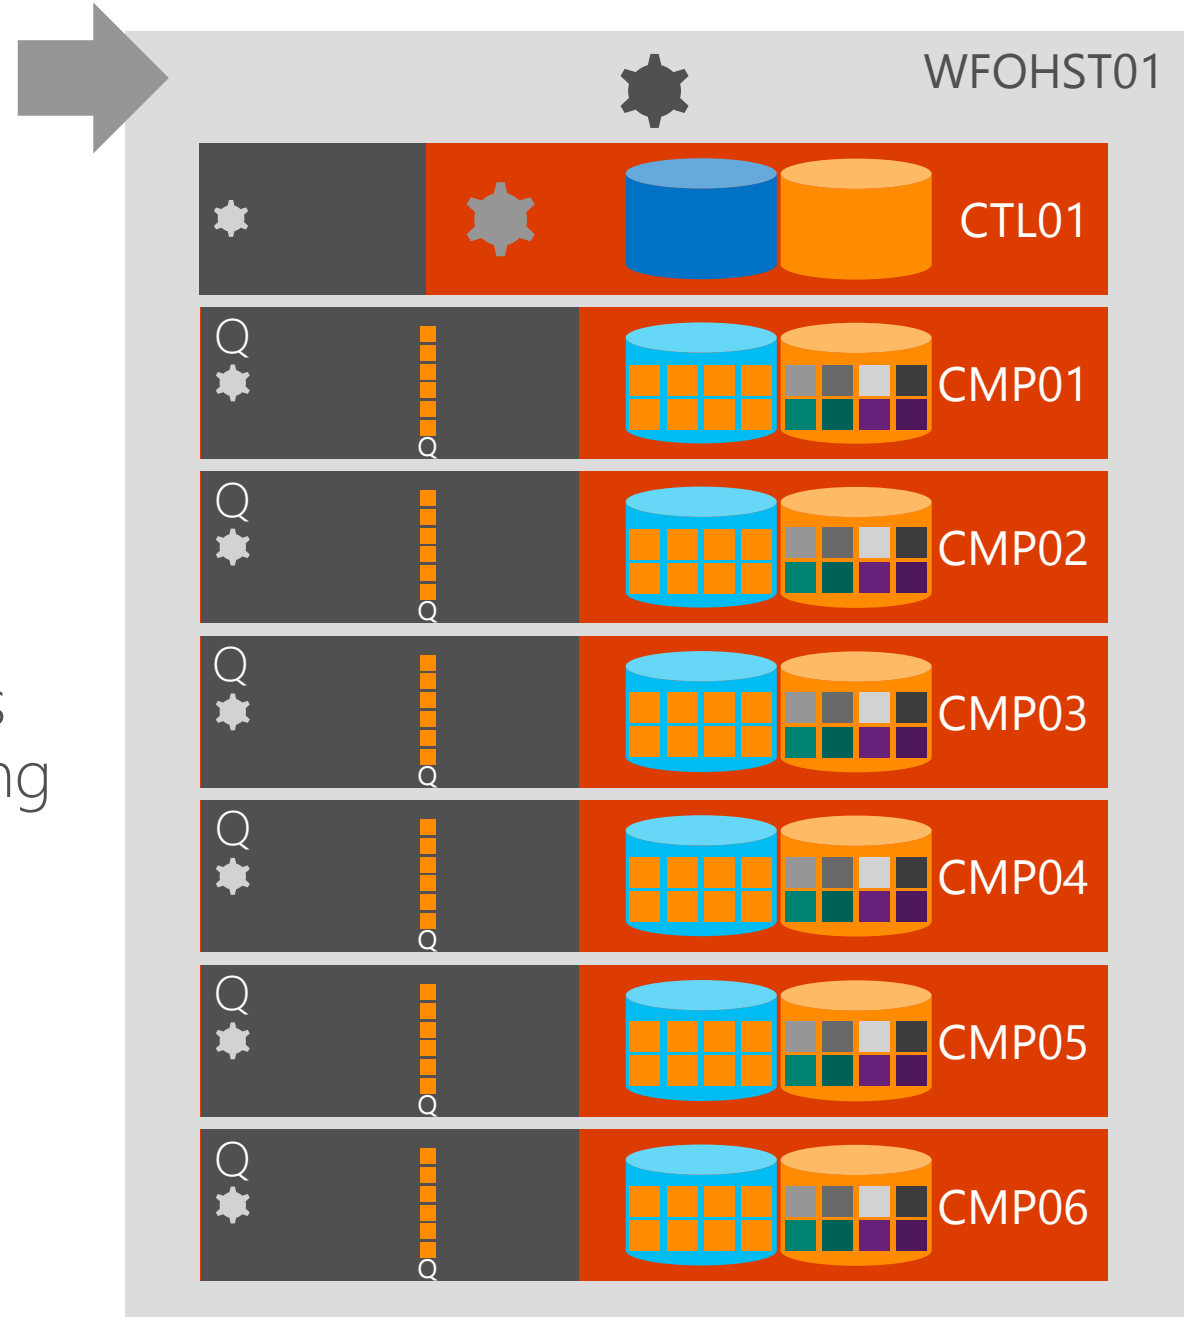
## User submits query

### PDWEngine

- Creates a DSQL plan, by querying the shell database
- Generates Random ID
- Creates geometry aligned temp tables
- **FROM** `dbo.FactInternetSales`
- **GROUP BY** `ProductKey`;

### DMS

- Executes query from PDWEngine
- Query filters and pre-aggregates data
- Receives data in 32KB buffers
- Places buffers in a read queue for processing



# Shuffle – Part 2

## The DMS Hash\_Reader Worker

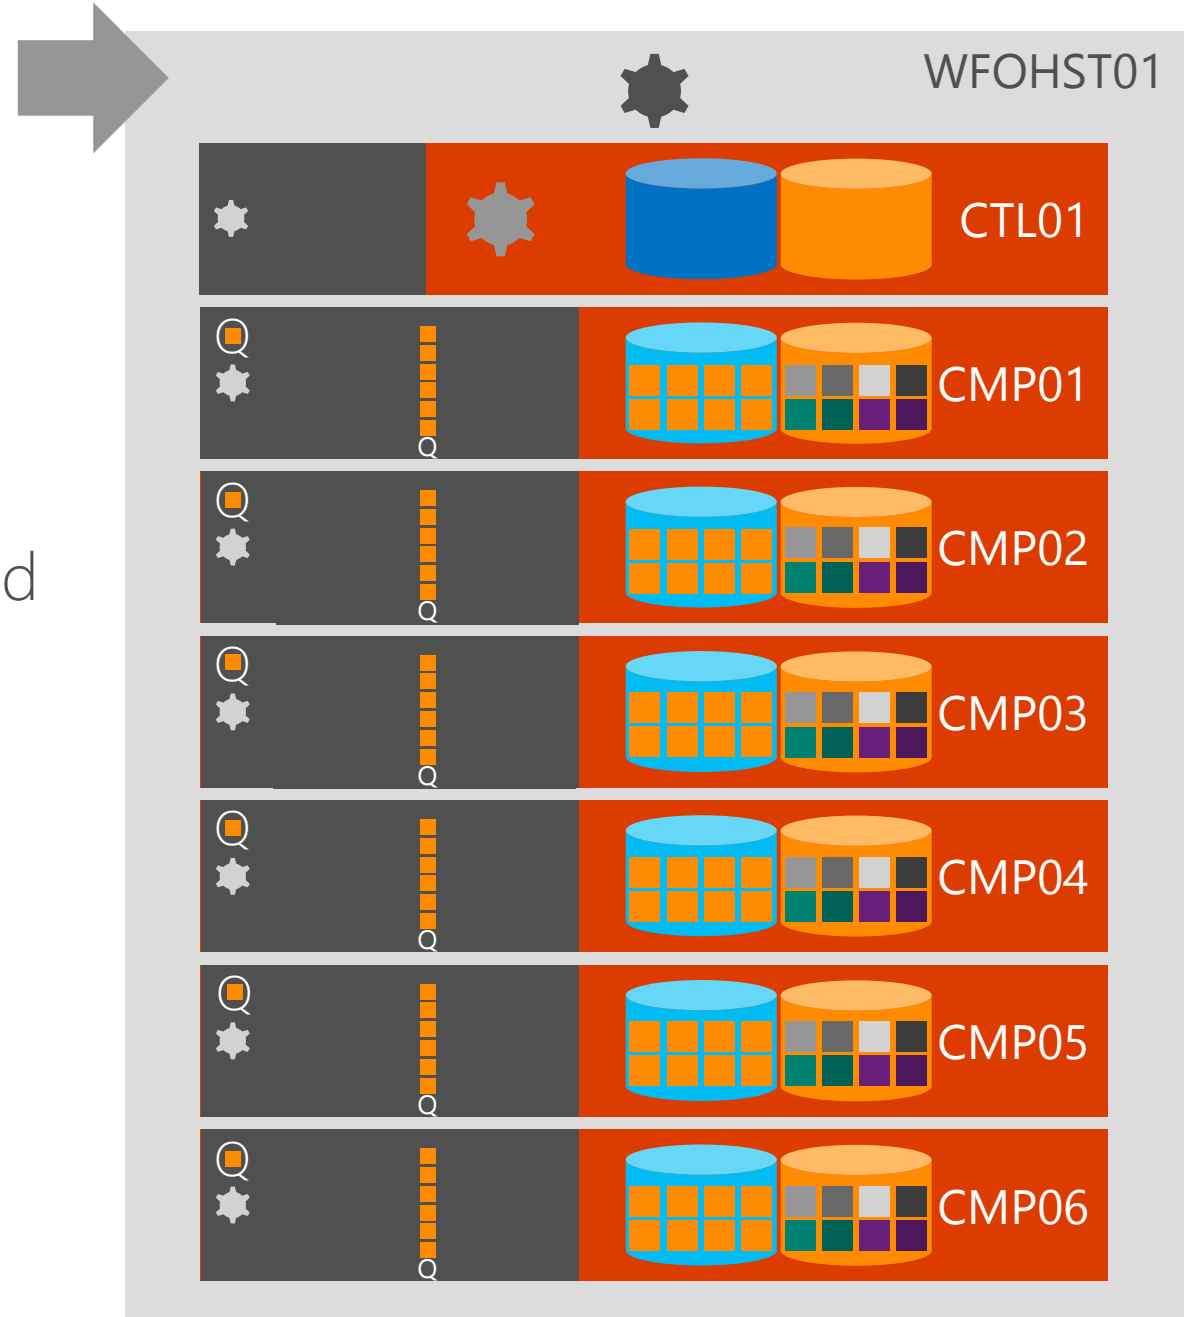
- Takes buffers from the read queue
- Read the buffer row by row
- Hash new distribution key column
- Write row into new distribution aligned buffer

## Sender

- Transmits full buffers to target distribution

## Receiver

- Places buffer on distribution aligned write queue



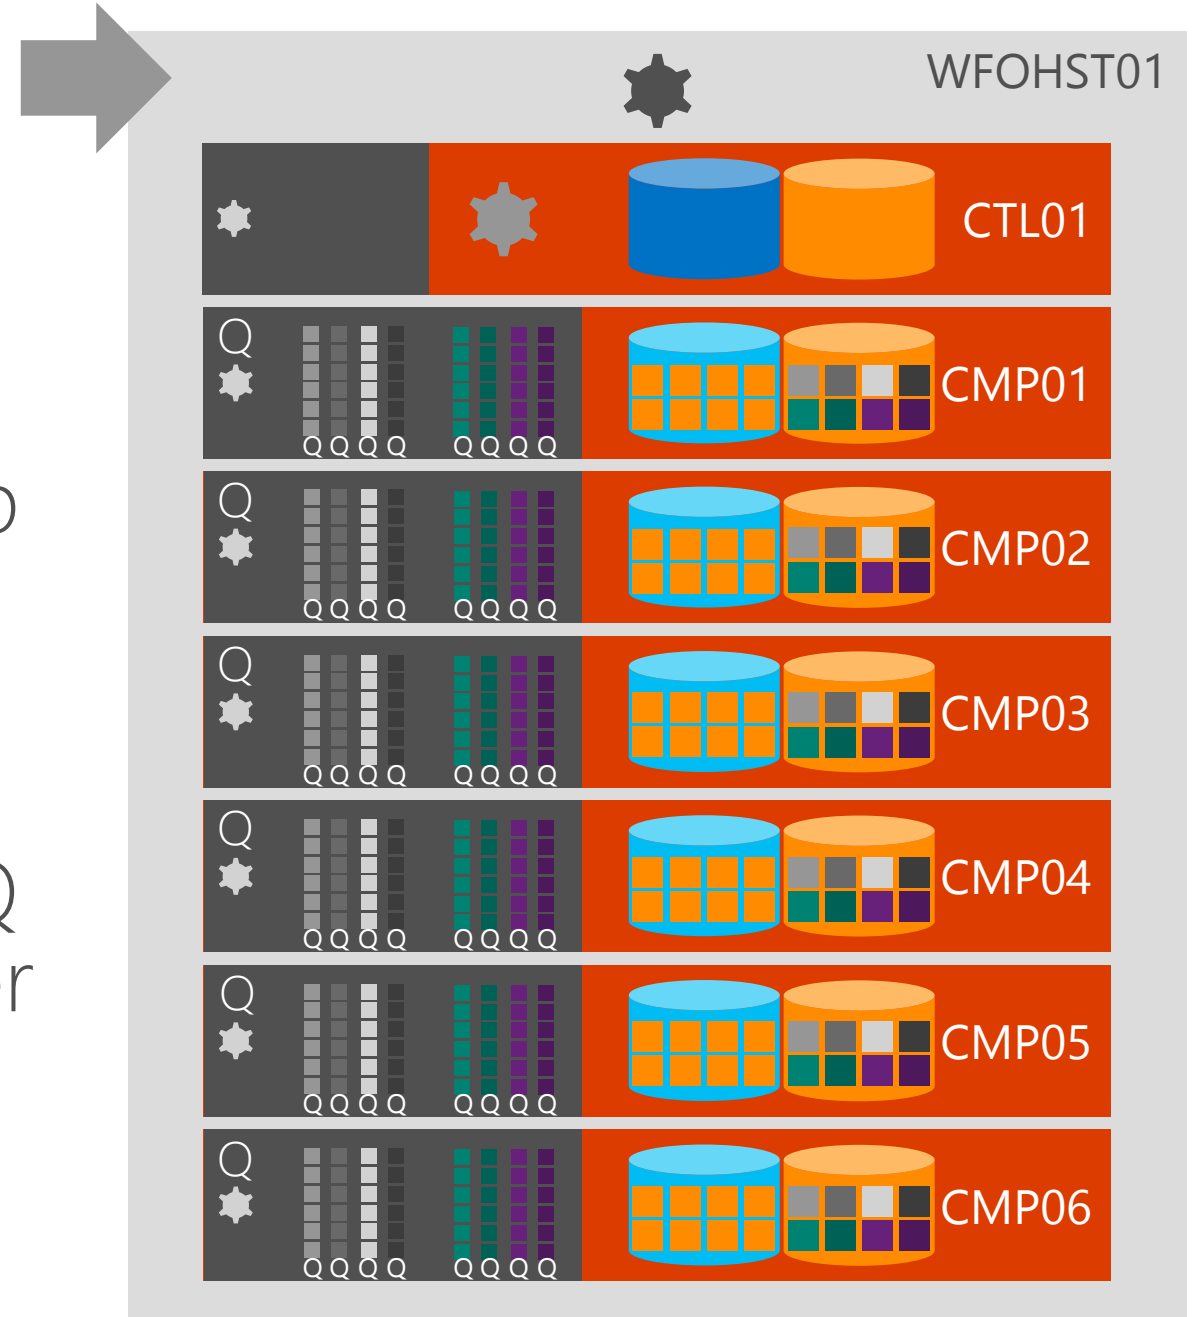
# Shuffle – Part 3

## DMS Writer Worker

- Bulk load buffers into distribution aligned Tempdb “Q” tables

## PDWEngine

- Submits distribution compatible query against Q tables returning data to user
- Drops Q Tables



# Broadcast Move Scenario

# Understanding Broadcast

- Distributed data > Replicated data
- Copies qualifying data from distributions and replicates it to all nodes
- Typically used when the cost estimation for the movement is low
- No hashing of the data is required
- Worker type is DIRECT\_READER
- Broadcast is the failsafe movement type



# Explain

```
<?xml version="1.0" encoding="utf-8"?>
<dsql_query>
  <sql>SELECT c.FirstName + ' ' + c.LastName
FROM [AdventureWorksPDW2012].dbo.FactInternetSales fis
LEFT JOIN [AdventureWorksPDW2012].dbo.DimCustomer c on fis.CustomerKey = c.CustomerKey
WHERE c.LastName = 'Smith'
OPTION (LABEL = 'Broadcast : Dim')</sql>
  <dsql_operations total_cost="1.548252288" total_number_operations="5">
    <dsql_operation operation_type="RND_ID">...</dsql_operation>
    <dsql_operation operation_type="ON">...</dsql_operation>
    <dsql_operation operation_type="BROADCAST_MOVE">
      <operation_cost cost="1.548252288" accumulative_cost="1.548252288" average_rowsize="204" output_rows="31.6228" />
      <source_statement>SELECT [T1_1].[CustomerKey] AS [CustomerKey],
        [T1_1].[FirstName] AS [FirstName],
        [T1_1].[LastName] AS [LastName]
FROM (SELECT [T2_1].[CustomerKey] AS [CustomerKey],
        [T2_1].[FirstName] AS [FirstName],
        [T2_1].[LastName] AS [LastName]
FROM [AdventureWorksPDW2012].[dbo].[DimCustomer] AS T2_1
WHERE ([T2_1].[LastName] = (N'Smith')) AS T1_1</source_statement>
      <destination_table>[TEMP_ID_90857]</destination_table>
    </dsql_operation>
    <dsql_operation operation_type="RN">...</dsql_operation>
    <dsql_operation operation_type="ON">
      <location permanent="1" on="AllComputeNodes" />
      <sql_operations>
        <sql_operation>TABLE [tempdb].[dbo].[TEMP_ID_90857]</sql_operation>
      </sql_operations>
    </dsql_operation>
  </dsql_operations>
</dsql_query>
```

The diagram illustrates the execution plan for the provided SQL query. It features three main components represented by arrows pointing to specific parts of the XML output:

- Broadcast**: An orange arrow points to the `<dsql_operation operation_type="BROADCAST_MOVE">` element, which represents the broadcast operation.
- Costing & Estimates**: An orange arrow points to the `<operation_cost cost="1.548252288" accumulative_cost="1.548252288" average_rowsize="204" output_rows="31.6228" />` element, which provides cost and estimate data for the broadcast operation.
- Destination**: A dark gray arrow points to the `<destination_table>[TEMP_ID_90857]</destination_table>` element, which specifies the destination table for the broadcast operation.

```
<?xml version="1.0" encoding="utf-8"?>
```

```
<dsql_query>
```

```
  <sql>SELECT c.FirstName + ' ' + c.LastName
```

```
FROM [AdventureWorksPDW2012].dbo.FactInternetSales fis
```

```
LEFT JOIN [AdventureWorksPDW2012].dbo.DimCustomer c on fis.CustomerKey = c.CustomerKey
```

```
WHERE c.LastName = 'Smith'
```

```
OPTION (LABEL = 'Broadcast : Dim')</sql>
```

```
  <dsql_operations total_cost="1.548252288" total_number_operations="5">
```

```
    <dsql_operation operation_type="RND_ID">...</dsql_operation>
```

```
    <dsql_operation operation_type="ON">...</dsql_operation>
```

```
    <dsql_operation operation_type="BROADCAST_MOVE" operation>
```

```
    <dsql_operation operation_type="RETURN">
```

```
      <location distribution="AllDistributions" />
```

```
      <select>SELECT [T1_1].[col] AS [col]
```

```
FROM      (SELECT [T2_1].[col] AS [col]
```

```
          FROM      (SELECT ([T3_1].[FirstName] + N' ' + [T3_1].[LastName]) AS [col],
```

```
                        [T3_1].[CustomerKey] AS [CustomerKey]
```

```
                      FROM      [tempdb].[dbo].[TEMP_ID_90857] AS T3_1) AS T2_1
```

```
          INNER JOIN
```

```
            [AdventureWorksPDW2012].[dbo].[FactInternetSales] AS T2_2
```

```
          ON ([T2_1].[CustomerKey] = [T2_2].[CustomerKey])) AS T1_1</select>
```

```
    </dsql_operation>
```

```
    <dsql_operation operation_type="ON">
```

```
      <location permanent="false" distribution="AllComputeNodes" />
```

```
      <sql_operations>
```

```
        <sql_operation type="statement">DROP TABLE [tempdb].[dbo].[TEMP_ID_90857]</sql_operation>
```

```
      </sql_operations>
```

```
    </dsql_operation>
```

```
  </dsql_operations>
```

```
</dsql_query>
```



Return

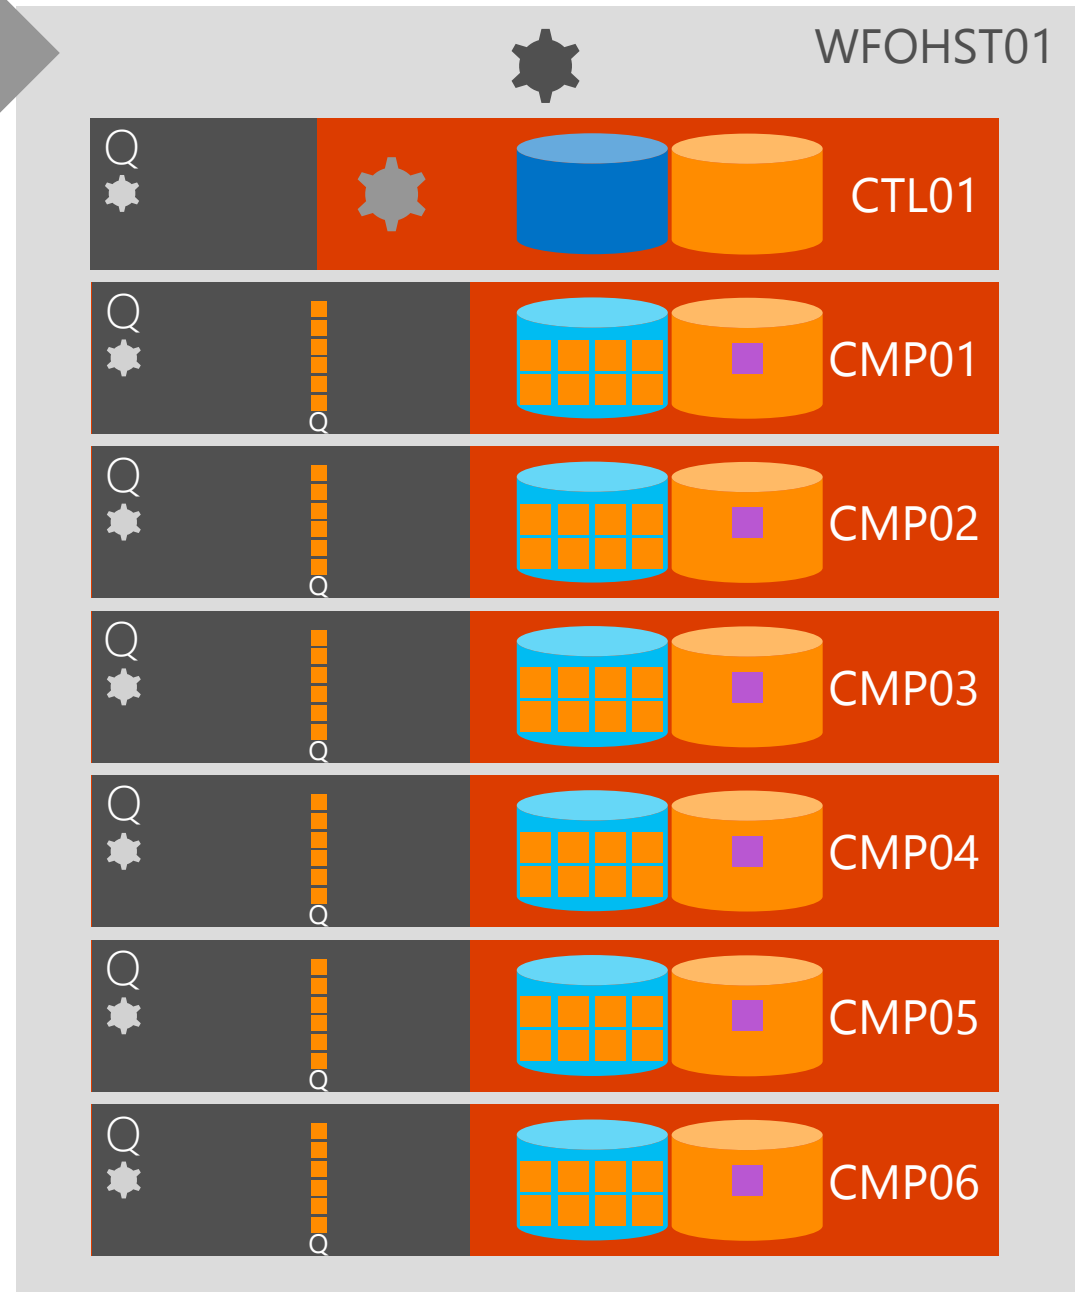
# Broadcast – Part 1

## User submits query

### PDWEngine

- Creates a DSQL plan, by querying the shell database
  - Generates Random ID
  - Creates geometry aligned temp tables
  - Instructs DMS to read data by executing a dsql query against the distributions
- ```
SELECT c.FirstName + ' ' + c.LastName  
FROM dbo.FactInternetSales fis  
LEFT JOIN dbo.DimCustomer c  
ON fis.CustomerKey=c.CustomerKey  
WHERE c.LastName = 'Smith'  
OPTION (LABEL = 'Broadcast')  
DMS
```

- Executes query from PDWEngine
- Query filters and pre-aggregates data
- Receives data in 32KB buffers
- Places buffers in a read queue for processing



# Broadcast – Part 2

DMS Direct\_Reader worker

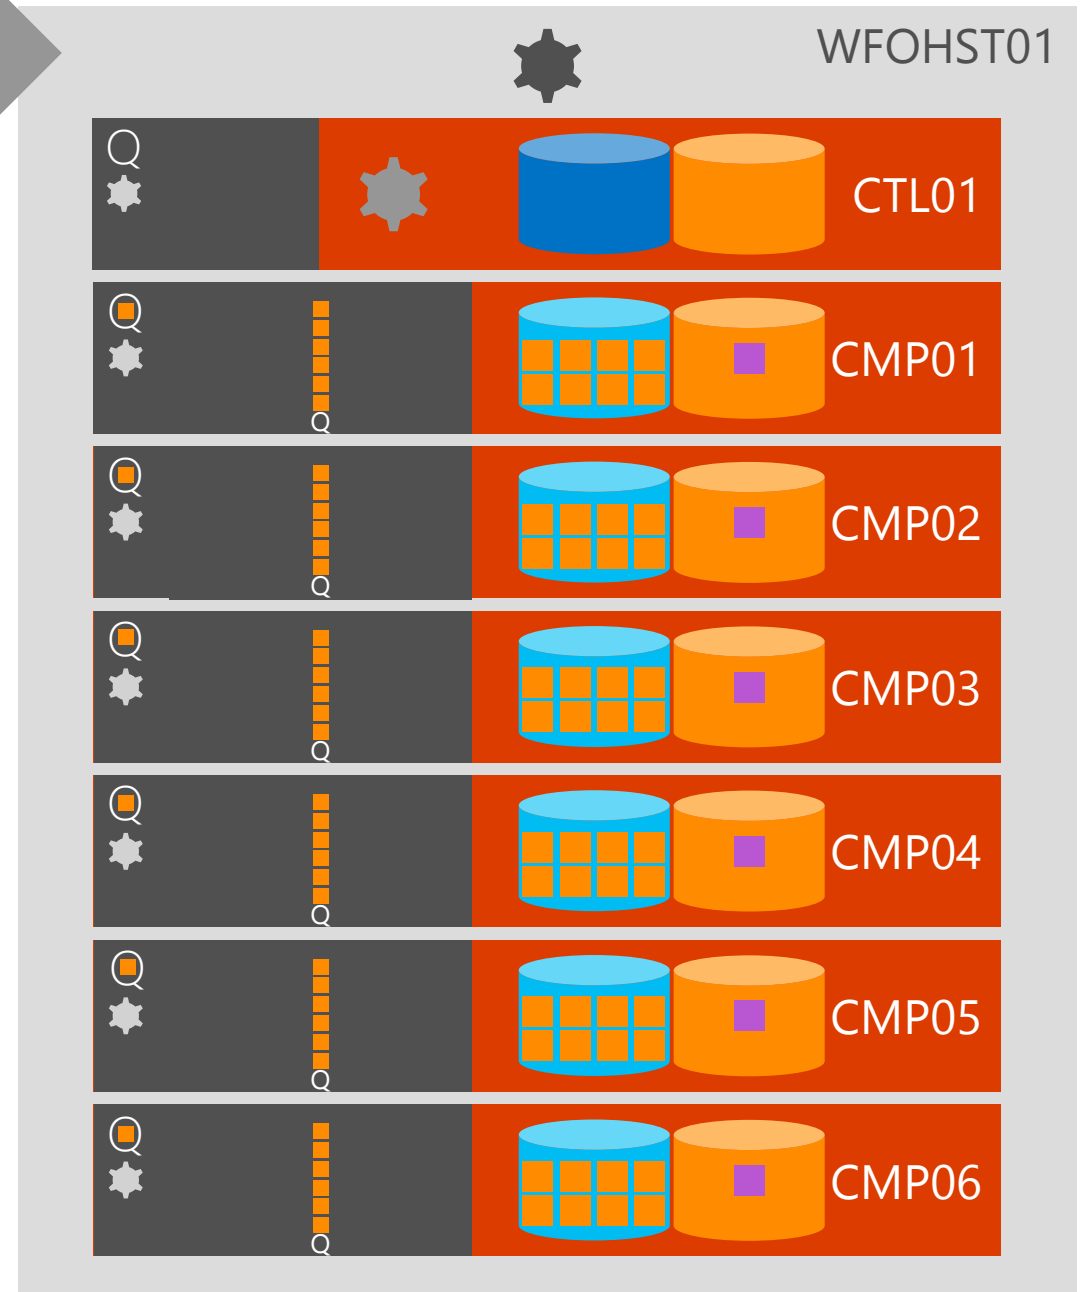
- Take buffer from queue
- Read buffer row by row
- Write row into new geometry aligned buffer

Sender

- Transmits full buffers to target node

Receiver

- Place buffer on geometry aligned write queue



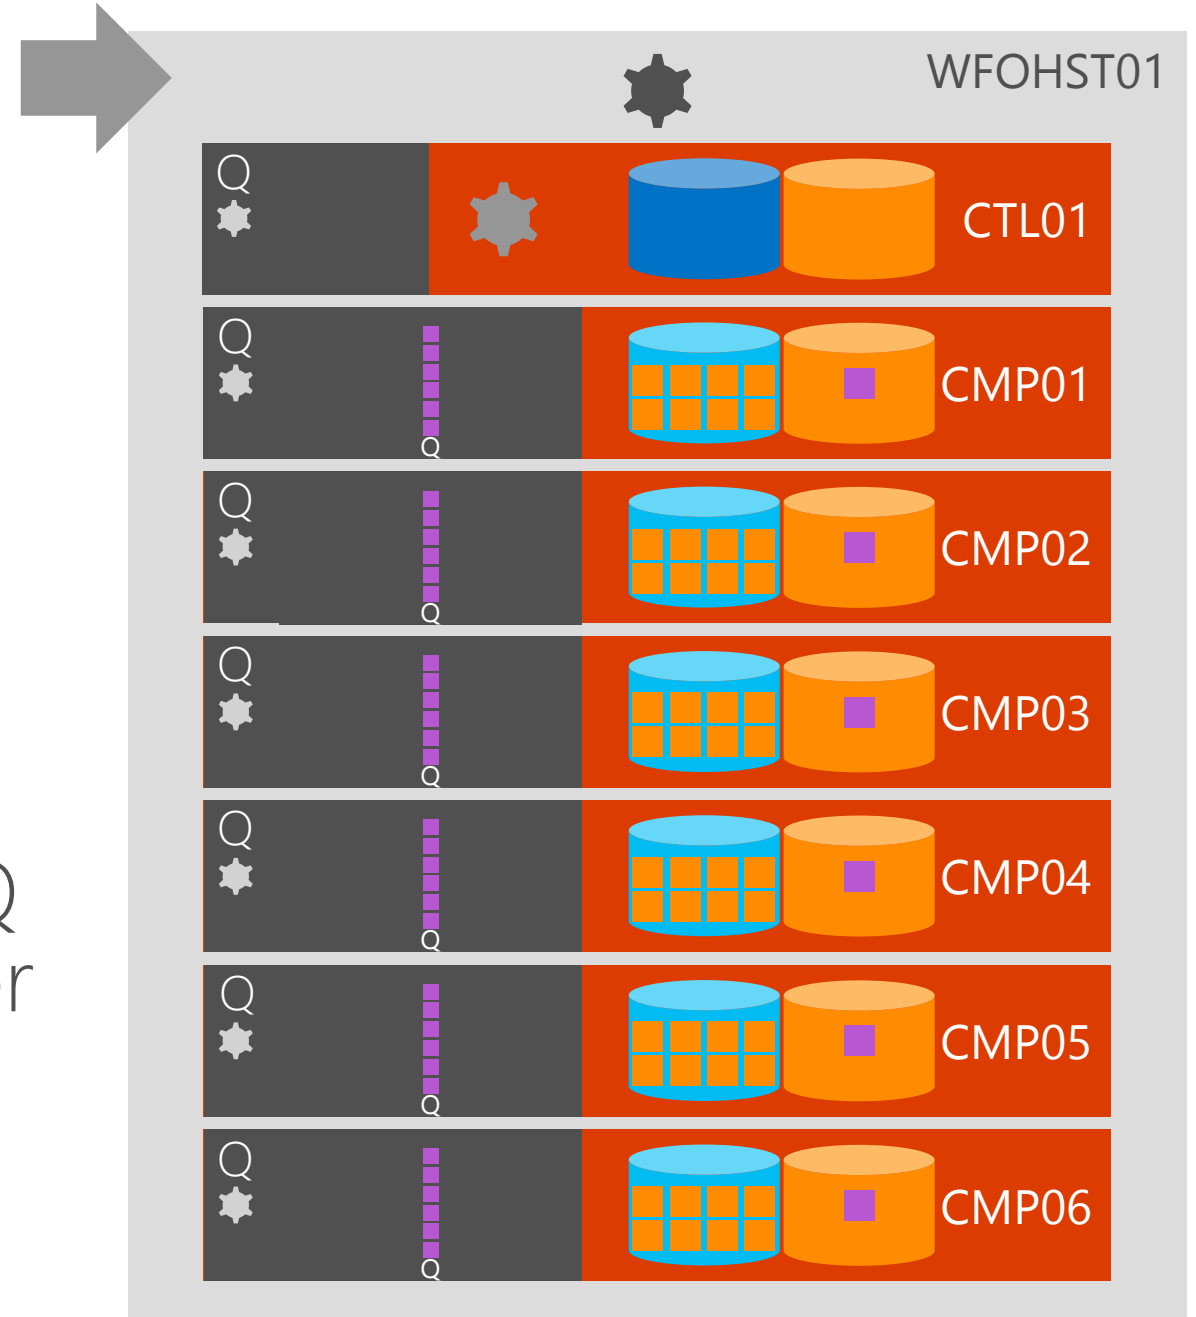
# Broadcast – Part 3

## DMS Writer Worker

- Bulk load buffers into geometry aligned tempdb "Q" tables

## PDW Engine

- Submits distribution compatible query against Q tables returning data to user
- Drops Q Tables



# Partition Move Scenario

# Understanding the Partition Move

- Centralise data to the control node
- Used for final aggregation
- Also used for temporary storage of an aggregation prior to another move operation
- No hashing of data required for Partition Move
- Worker type is PARALLEL\_COPY\_READER
- Can be an expensive operation
- Control node can become the bottleneck





# Partition – Part 1

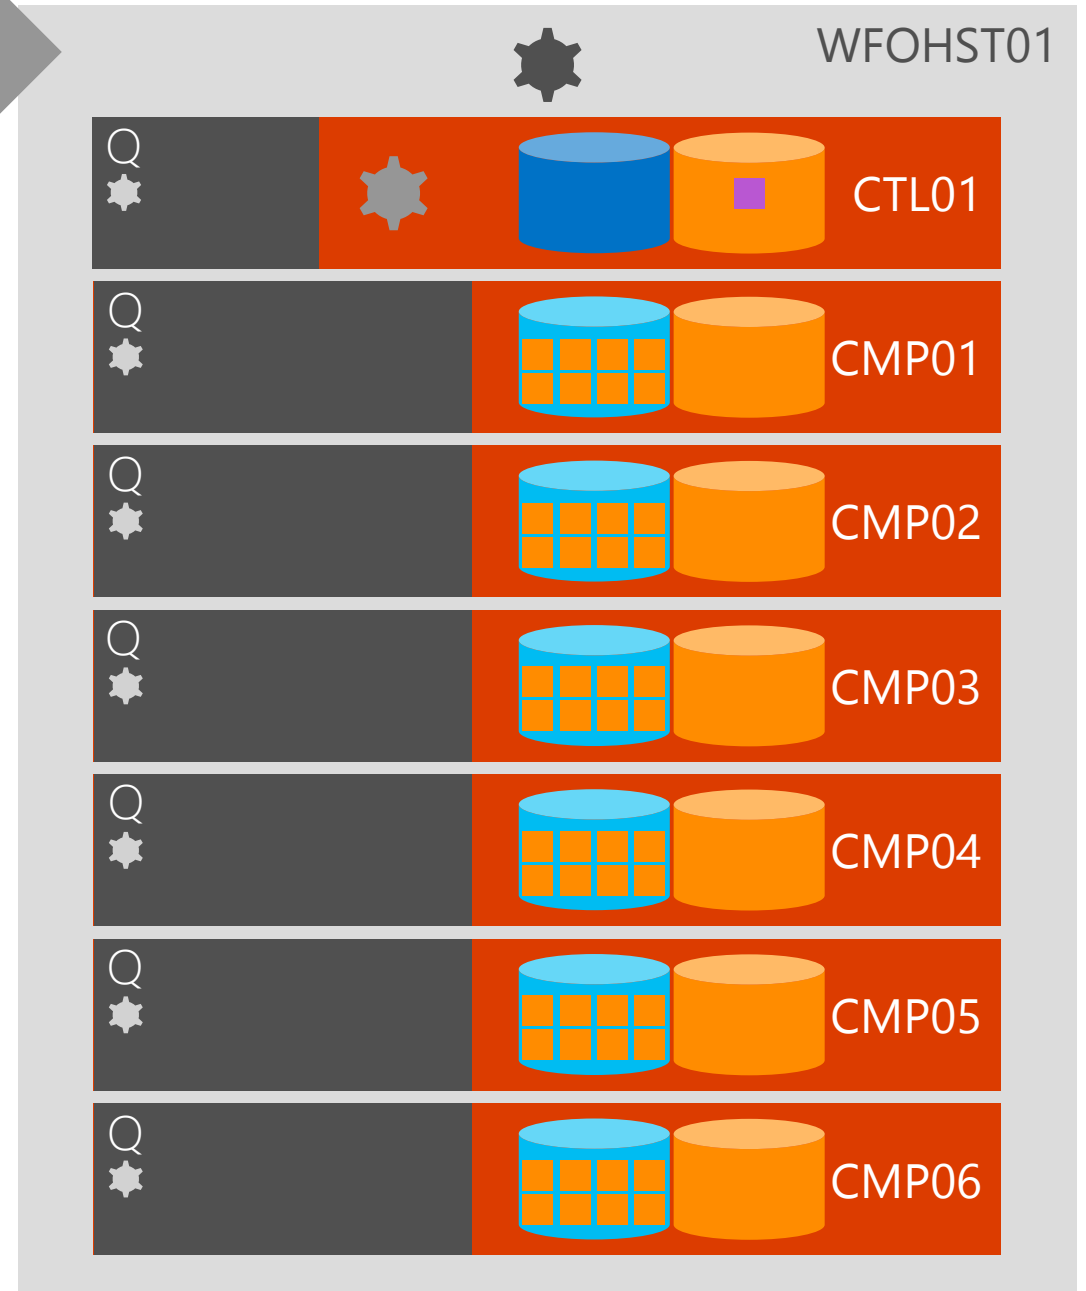
## User submits query

### PDWEngine

- Creates a DSQL plan, by querying the shell database
- Generates Random ID
- **SELECT COUNT(\*)**  
Creates temp table on Control Node
- **FROM dbo.FactInternetSales fis**  
Instructs DMS to read data by executing a dsq query against the distributions
- **OPTION (LABEL = 'Partition')**

### DMS

- Executes query from PDWEngine
- Query filters and pre-aggregates data
- Process data using Parallel\_Copy\_Reader worker
- Bypass data channel writing the row directly to Tempdb on control node

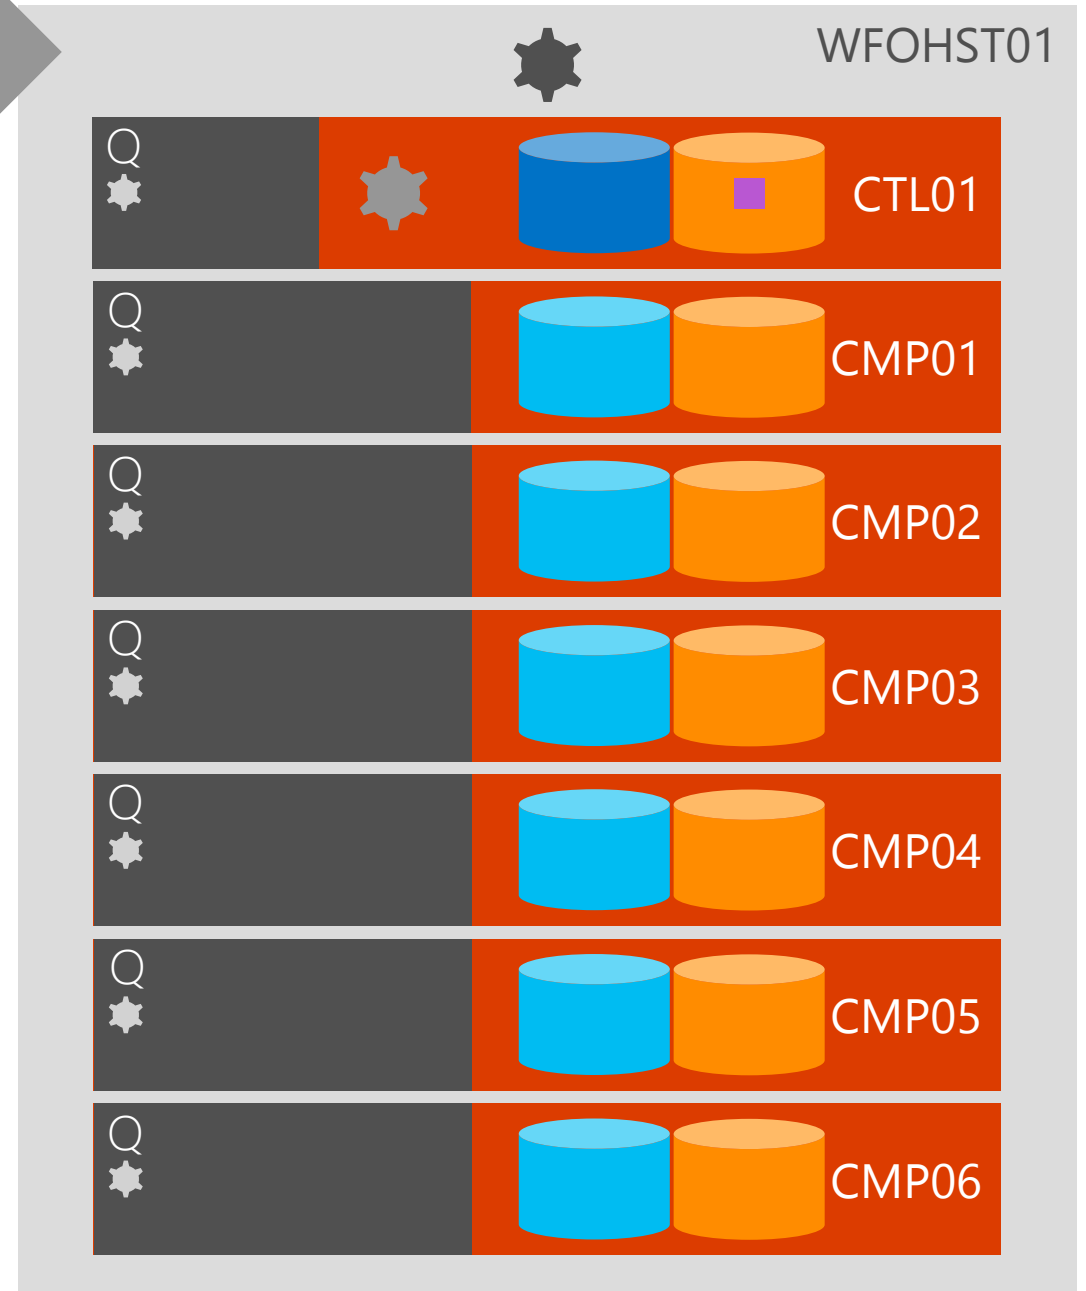


# Partition – Part 2



## PDWEngine

- Submits distribution compatible query against the Q table returning data to user
- Drops Q table



# Master Move Scenario

# Understanding The Master Move

- No user data is allowed to persist on the control node
- A mechanism is required to push the data back to the compute nodes
- A move operation is typically preceded by a partition move

# Explain Plan

```
<?xml version="1.0" encoding="utf-8"?>
<dsql_query>
  <sql>CREATE TABLE dbo.Top_5_Customers
WITH (DISTRIBUTION = REPLICATE)
AS
SELECT TOP 5 CustomerKey,FirstName,LastName
FROM [AdventureWorksPDW2012].dbo.DimCustomer
OPTION (LABEL = 'Move')</sql>
  <dsql_operations total_cost="0.4896" total_number_operations="10">
    <dsql_operation operation_type="ON">
      <location permanent="true" distribution="AllComputeNodes" />
      <sql_operations>
        <sql_operation type="statement">CREATE TABLE [Instructor].[dbo].[Top_5_Customers]
        </sql_operation>
      </sql_operations>
    </dsql_operation>
    <dsql_operation operation_type="RND_ID">
      <identifier>TEMP_ID_90853</identifier>
    </dsql_operation>
    <dsql_operation operation_type="ON">
      <location permanent="false" distribution="Control" />
      <sql_operations>
        <sql_operation type="statement">CREATE TABLE [tempdb].[dbo].[TEMP_ID_90853]
        </sql_operation>
      </sql_operations>
    </dsql_operation>
    <dsql_operation operation_type="PARTITION_MOVE">
      <operation_cost cost="0.2448" accumulative_cost="0.2448" output_rows="204" output_rows="5" />
      <location distribution="AllDistributions" />
      <source_statement>SELECT [T1_1].[CustomerKey] AS [CustomerKey],
        [T1_1].[FirstName] AS [FirstName],
        [T1_1].[LastName] AS [LastName]
FROM (SELECT TOP (CAST ((5) AS BIGINT)) [T2_1].[CustomerKey] AS [CustomerKey],
        [T2_1].[FirstName] AS [FirstName],
        [T2_1].[LastName] AS [LastName]
FROM [AdventureWorksPDW2012].[dbo].[DimCustomer] AS T2_1) AS T1_1</source_statement>
      <destination>Control</destination>
      <destination_table>[TEMP_ID_90853]</destination_table>
    </dsql_operation>
  </dsql_operations>
</dsql_query>
```



Partition  
Move

# Explain Plan

```
<?xml version="1.0" encoding="utf-8"?>
<dsql_query>
  <sql>CREATE TABLE dbo.Top_5_Customers
WITH (DISTRIBUTION = REPLICATE)
AS
SELECT TOP 5 CustomerKey,FirstName,LastName
FROM [AdventureWorksPDW2012].dbo.DimCustomer
OPTION (LABEL = 'Move')</sql>
  <dsql_operations total_cost="0.4896" total_number_operations="10">
    <dsql_operation operation_type="ON">...</dsql_operation>
    <dsql_operation operation_type="RND_ID">...</dsql_operation>
    <dsql_operation operation_type="ON">...</dsql_operation>
    <dsql_operation operation_type="PARTITION_MOVE">...</dsql_operation>
    <dsql_operation operation_type="RND_ID">
      <identifier>TEMP_ID_90854</identifier>
    </dsql_operation>
    <dsql_operation operation_type="ON">
      <location permanent="false" distribution="AllComputeNodes" />
    </dsql_operation>
  </dsql_operations>
  <sql_operation_type="statement">CREATE TABLE [tempdb].[dbo].[TEMP_ID_90854]
  </sql_operation_type>
  <destination>Compute</destination>
  <source_statement>
    <sql_operation_type="MASTER_TABLE_MOVE">
      <cost>0.2448</cost> <accumulative_cost>0.2448</accumulative_cost> <output_rows>5</output_rows> />
    <sql_operation_type="SELECT">
      <source_statement>
        SELECT [T1_1].[CustomerKey] AS [CustomerKey],
        [T1_1].[FirstName] AS [FirstName],
        [T1_1].[LastName] AS [LastName]
      FROM [tempdb].[dbo].[TEMP_ID_90853] AS T2_1) AS T1_1</source_statement>
      <destination>Compute</destination>
      <destination_table>[TEMP_ID_90854]</destination_table>
    </sql_operation>
  </sql_operation_type>
  <dsql_operation operation_type="ON">...</dsql_operation>
  <dsql_operation operation_type="ON">...</dsql_operation>
  <dsql_operation operation_type="ON">...</dsql_operation>
</dsql_operations>
<meta-data>
  <full />
</meta-data>
</dsql_query>
```

Target Node

Move

Target table

<destination\_table>[TEMP\_ID\_90854]</destination\_table>

```

<?xml version="1.0" encoding="utf-8"?>
<dsql_query>
  <sql>CREATE TABLE dbo.Top_5_Customers
WITH (DISTRIBUTION = REPLICATE)
AS
SELECT TOP 5 CustomerKey,FirstName,LastName
FROM [AdventureWorksPDW2012].dbo.DimCustomer
OPTION (LABEL = 'Move')</sql>
  <dsql_operations total_cost="0.4896" total_number_operations="10">
    <dsql_operation operation_type="ON">...</dsql_operation>
    <dsql_operation operation_type="RND_ID">...</dsql_operation>
    <dsql_operation operation_type="ON">...</dsql_operation>
    <dsql_operation operation_type="PARTITION_MOVE">...</dsql_operation>
    <dsql_operation operation_type="RND_ID">...</dsql_operation>
    <dsql_operation operation_type="ON">...</dsql_operation>
    <dsql_operation operation_type="MASTER_TABLE_MO">...</dsql_operation>
    <dsql_operation operation_type="ON">
      <location permanent="true" distribution="AllComputeNodes" />
      <sql_operations>
        <sql_operation type="statement">INSERT INTO [Instructor].[dbo].[Top_5_Customers] WITH (TABLOCK)
SELECT [T1_1].[CustomerKey],
      [T1_1].[FirstName],
      [T1_1].[LastName]
FROM   [tempdb].[dbo].[TEMP_ID_90854] AS T1_1
OPTION (MAXDOP 1)</sql_operation>
      </sql_operations>
    </dsql_operation>
    <dsql_operation operation_type="ON">...</dsql_operation>
    <dsql_operation operation_type="ON">...</dsql_operation>
  </dsql_operations>
</meta-data>...</meta-data>
</dsql_query>

```



MAXDOP 1  
Write

# Move – Part 1

## PDW Engine User submits query

- Creates a DSQL plan, by querying the shell database

```
CREATE TABLE dbo.Top_5_Customers  
WITH (DISTRIBUTION = D REPLICATE)
```

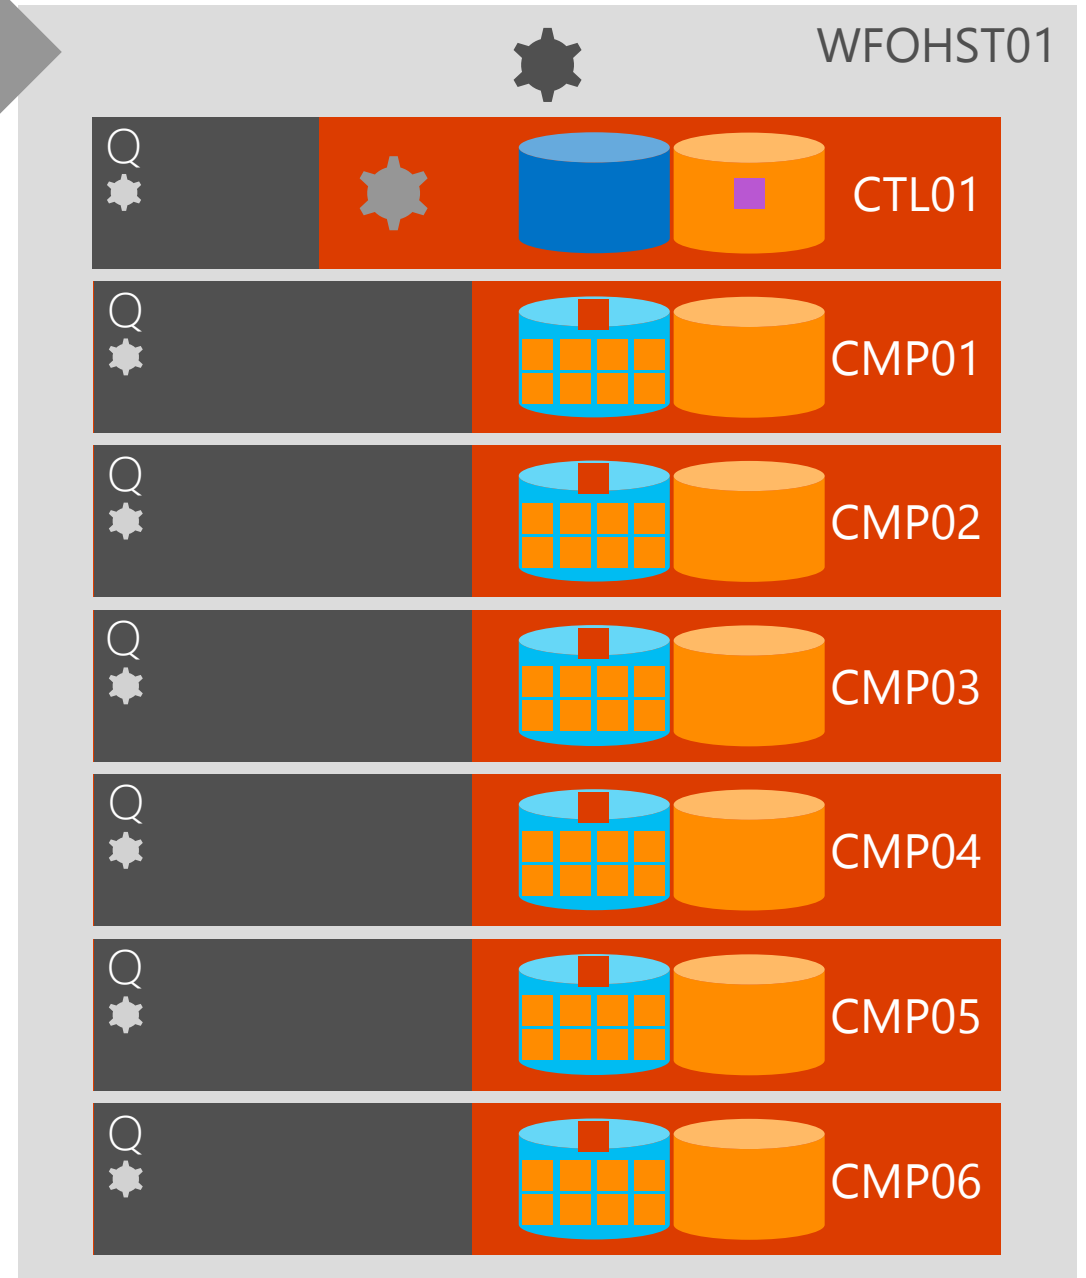
AS Creates temp table on Control Node

```
SELECT TOP 5  
CustomerKey, FirstName, LastName
```

```
FROM dbo.DimCustomer
```

```
OPTION (LABEL = 'Move')
```

- Executes query from PDW Engine
- Query filters and pre-aggregates data
- Process data using Parallel\_Copy\_Reader worker
- Bypass data channel writing the row directly to Tempdb on control node





# Move – Part 2

## PDWEngine

- Generates Random ID
- Instruct DMS to read data held on control node

## DMS

- Executes query from PDWEngine
- Receives data in 32KB buffers
- Places buffers in a read queue for processing
- Reads buffers using Direct\_Reader worker
- Writes row to new geometry aligned buffer
- Places buffer on send queue



# Move – Part 3

## Sender

- Moves buffers to compute nodes via DMS data channel

## Receiver

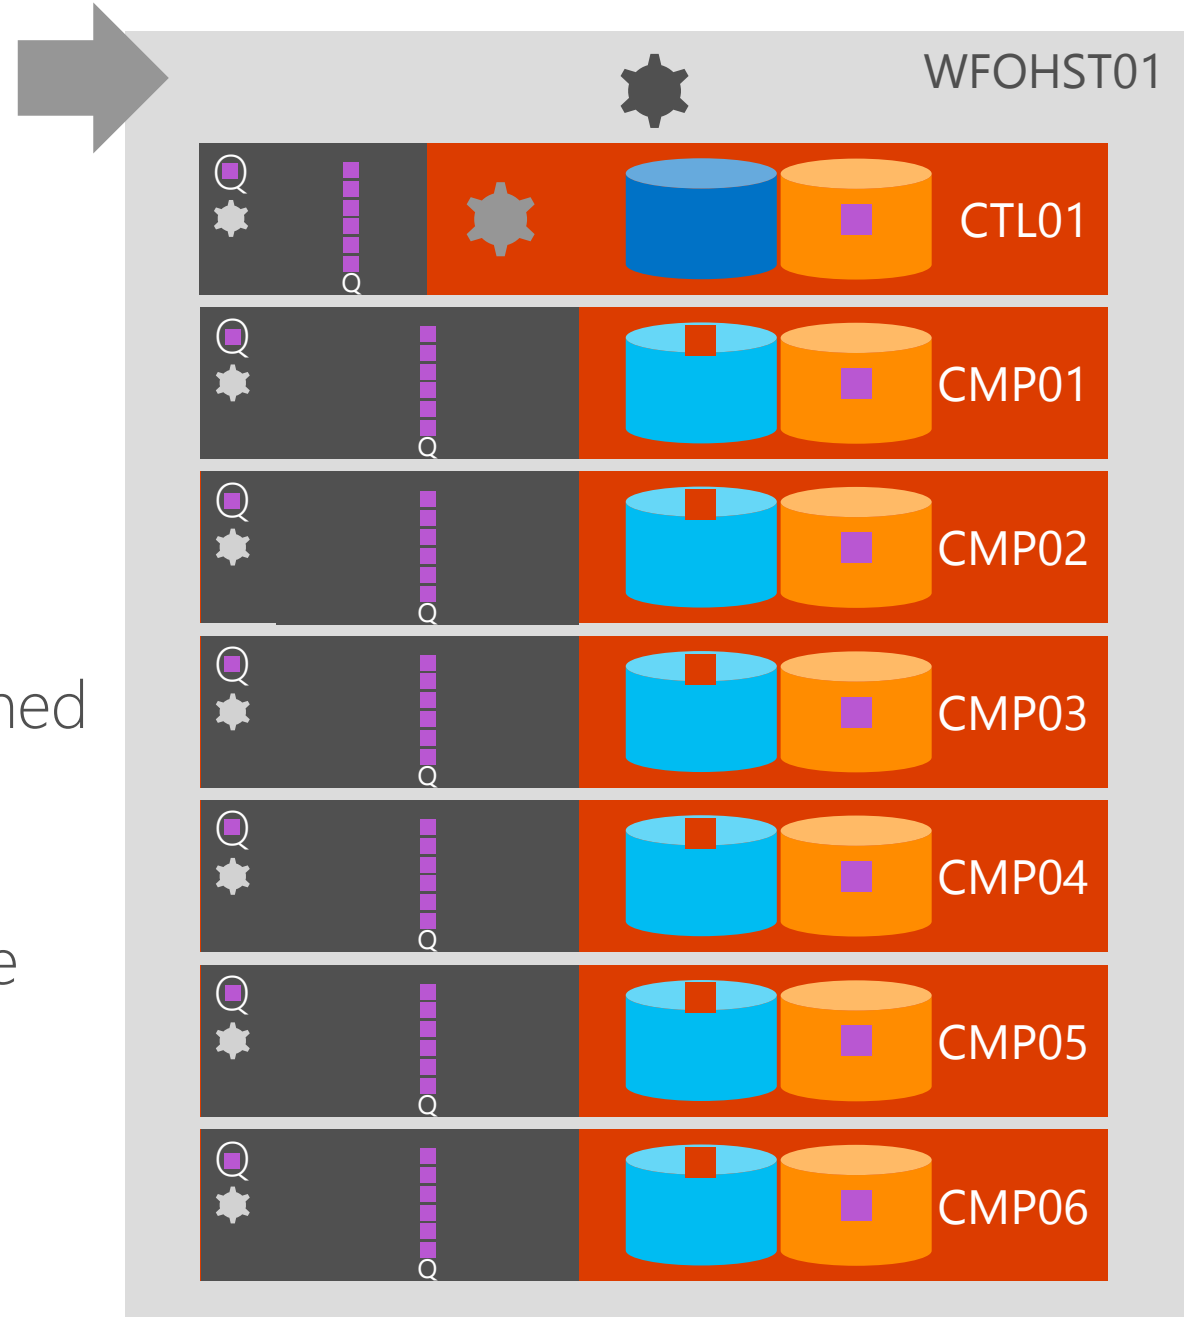
- Receives buffers and queues them

## Writer worker

- Bulk loads buffers into geometry aligned tempdb "Q" tables

## PDWEngine

- Submits MAXDOP(1) write queries to populate target table in user database
- Drops Q Tables



# Trim Move Scenario

# Introducing Trim Move

- Changing table geometry
- From a replicated table to distributed table
- All data exists on the compute node already
- No data leaves the compute node

# Explain Plan

```
<?xml version="1.0" encoding="utf-8"?>
<dsql_query>
  <sql>CREATE TABLE dbo.DimEmployee
WITH (DISTRIBUTION = Hash(EmployeeKey))
AS
SELECT *
FROM [AdventureWorksPDW2012].dbo.DimEmployee
OPTION (LABEL = 'Trim')</sql>
  <dsql_operations total_cost="24.272" total_number_operations="2">
    <trim_operation type="statement">CREATE TABLE [Instructor].[dbo].[DimEmployee]
    <trim_columns>EmployeeKey;</trim_columns>
    <trim_columns>EmployeeKey;</trim_columns>
    <destination_table>[DimEmployee]</destination_table>
  </dsql_operation>
</dsql_operations>
  <meta-data>
    <partitioned>
      <partitioning_column index="1">EmployeeKey</partitioning_column>
    </partitioned>
  </meta-data>
</dsql_query>
```

Distribution Key

Trim

Target

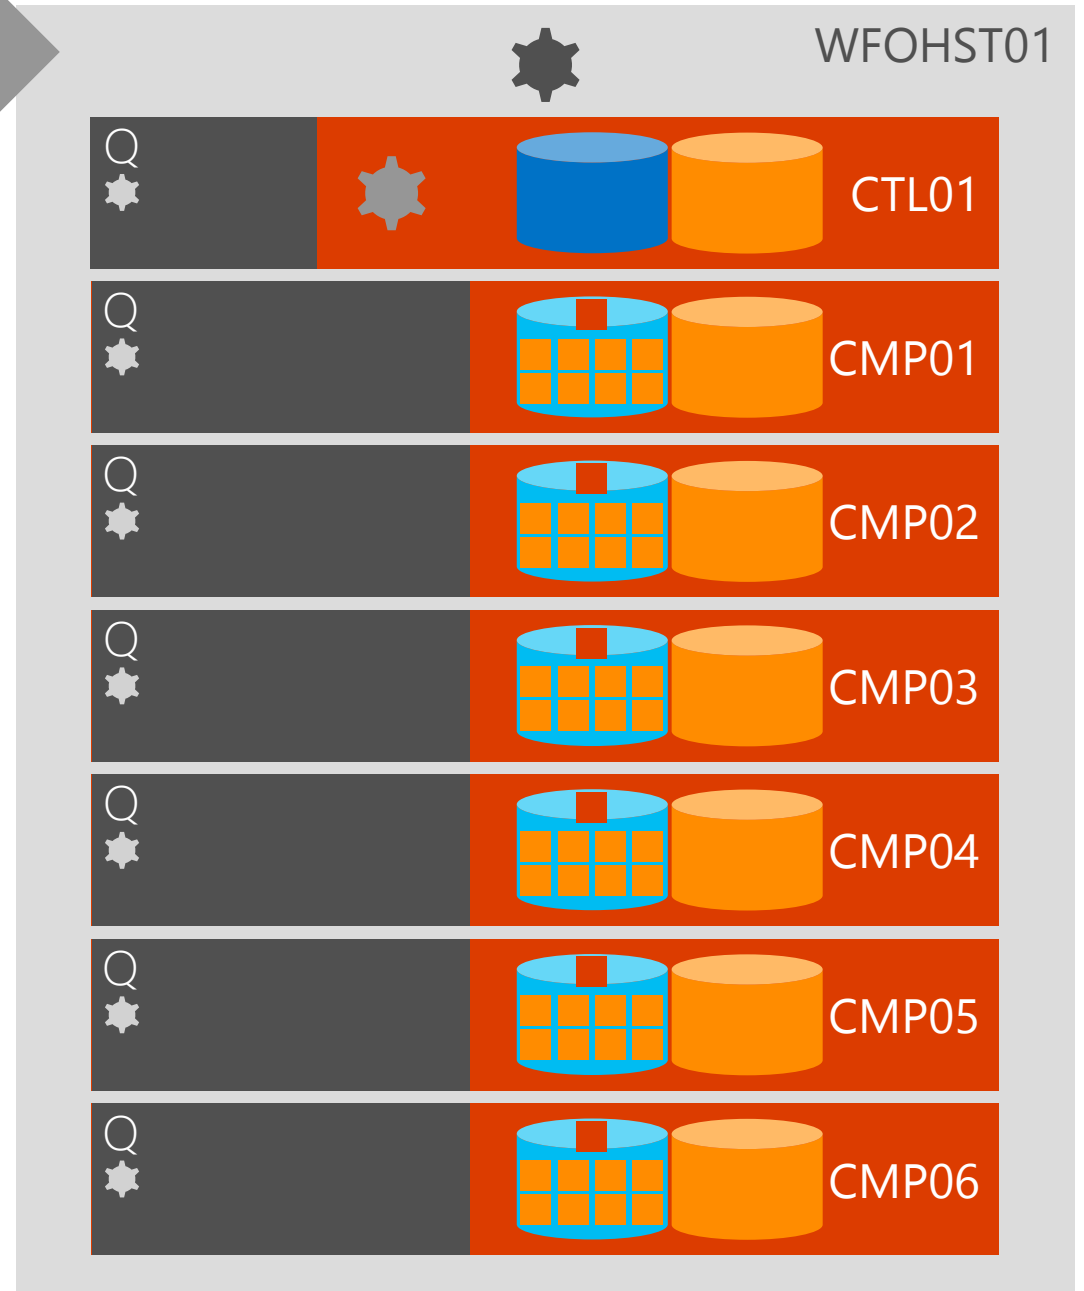
<destination\_table>[DimEmployee]</destination\_table>

# Trim – Part 1



PDW Engine User submits query

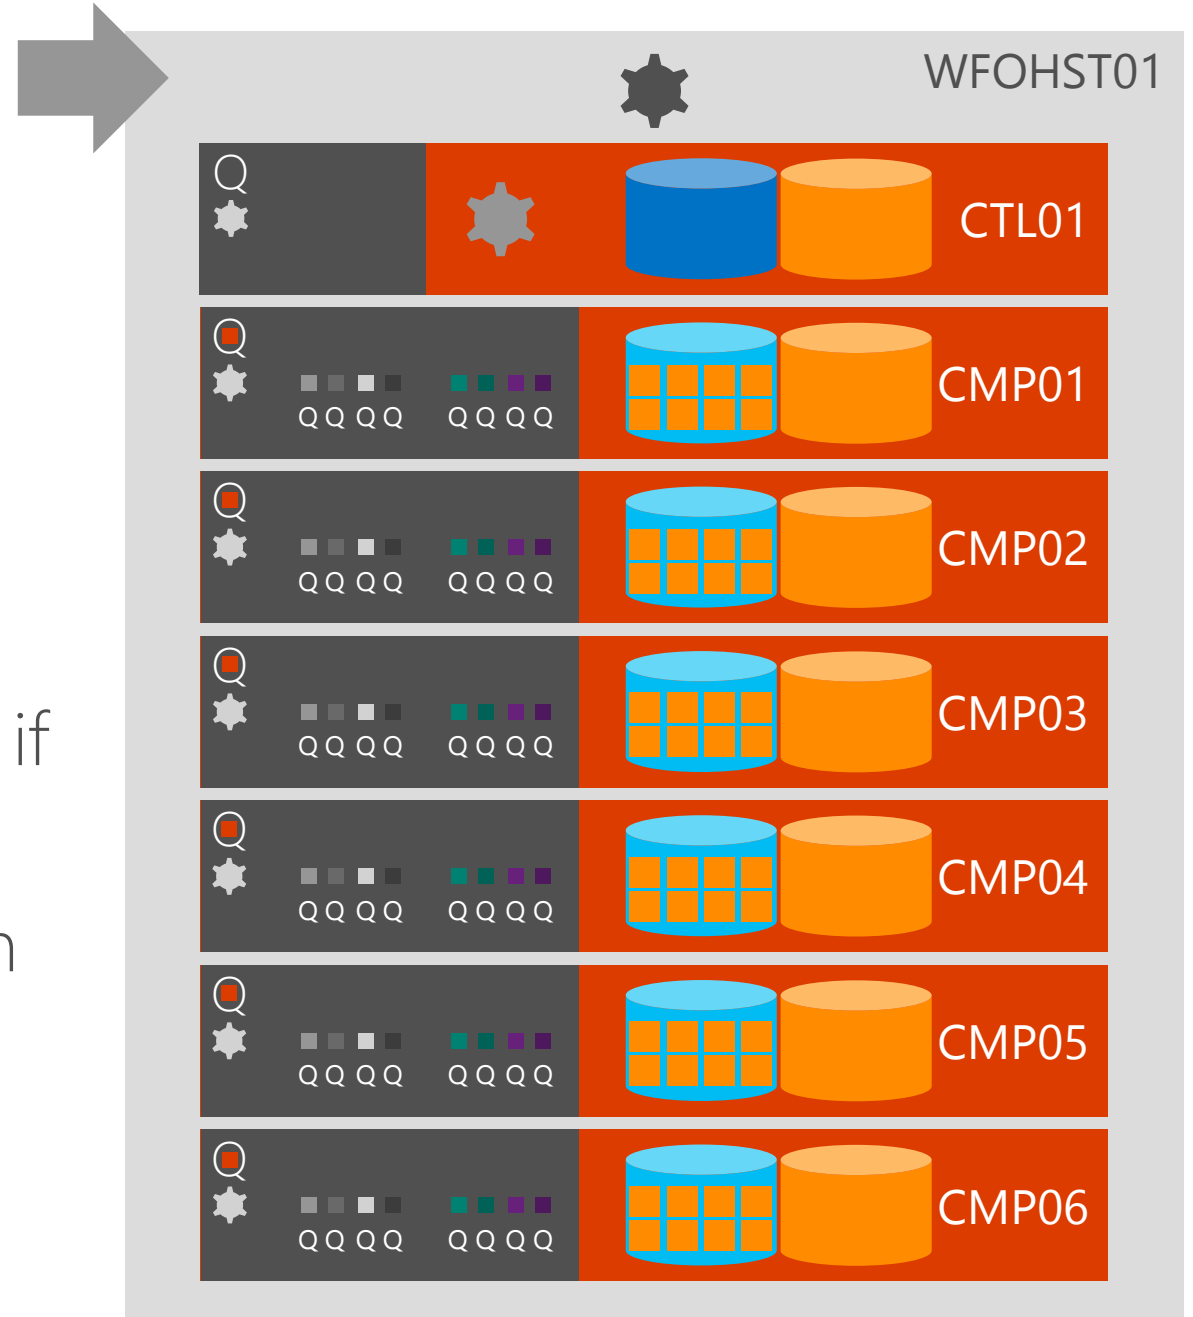
- Creates a DSQL plan, by querying the shell database
- `CREATE TABLE dbo.DimEmployee`  
`WITH (DISTRIBUTION = Hash(EmployeeKey))`  
`AS`
- `SELECT*`  
`FROM`  
`[AdventureWorksPDW2012].dbo.DimEmployee`  
`OPTION (LABEL = 'Trim')`
- Executes query from PDW Engine
- Receives data in 32KB buffers
- Places buffers a read queue for processing



# Trim – Part 2

The DMS performs the Trim move

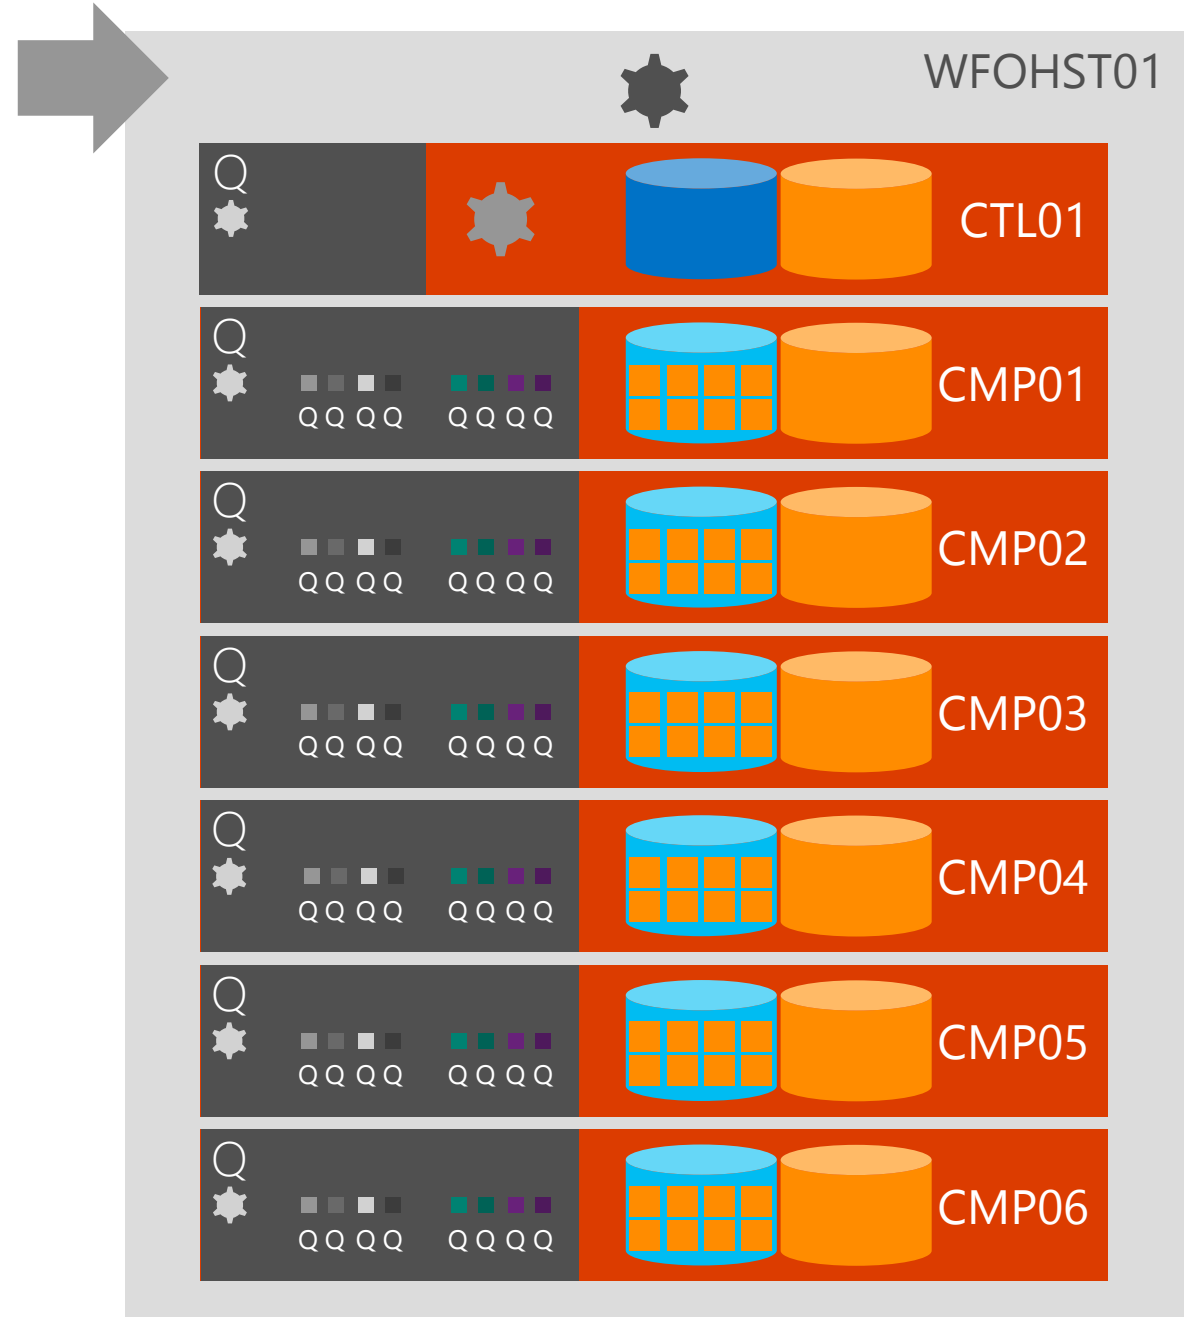
- Pull a buffer from queue
- Reads buffer row by row using Hash\_Reader worker
- Hash on new distribution key
- Only writes rows to the new distribution aligned output buffer if the row belongs on this compute node
- When buffer is full place buffer on distribution aligned write queue ready for writing



# Trim – Part 3

## The DMS Writer Workers

- Pull buffers from the write queue
- Bulk loads the data directly into the target tables





# DistributeReplicatedTable Move Scenario

# Introducing DistributeReplicatedTable Move

- Required to provide consistency across replicated tables
- Write is first persisted on one compute node in Tempdb
- Persisted data is then shared with other nodes by bulk inserting the values into their Tempdb
- DistributeReplicatedTable DMS movement does not insert data into the target
  - Performed by PDWEngine as a follow up step

# Explain Plan

Source Compute  
node

DistributeReplicated

Target

```
<?xml version="1.0" encoding="utf-8"?>
<dsql_query>
  <source_node>0</source_node>
  <sql>INSERT INTO [Instructor].[dbo].[TimeLog] WITH (TABLOCK) ([DateNow])
  SELECT SY
  OPTION (L
  <dsql_operation operation_type="DISTRIBUTE_REPLICATED_TABLE_MOVE">
    <source_node>0</source_node>
    <source_statement>SELECT [T1_1].[PDWExpr1001] AS [PDWExpr1001]
  FROM   (VALUES (ISNULL(CONVERT (DATETIME2 (7), N'2014-03-23 06:56:17.9700886', 0), CONVERT (DATETIME2 (7), N'2014-03-23 06:56:17.9700886', 0)))
    <destination_table>[TEMP_ID_90843]</destination_table>
  </dsql_operation>
  <dsql_operation operation_type="ON">
    <location permanent="AllComputeNodes" />
    <sql_operations>
      <sql_operation type="INSERT">INSERT INTO [Instructor].[dbo].[TimeLog] WITH (TABLOCK) ([DateNow])
  SELECT [T1_1].[PDWExpr1001]
  FROM   [tempdb].[dbo].[TEMP_ID_90843] AS T1_1
  OPTION (MAXDOP 1)</sql_operation>
    <destination_table>[TEMP_ID_90843]</destination_table>
    </sql_operations>
  </dsql_operation>
  <dsql_operation operation_type="ON">...</dsql_operation>
</dsql_operations>
</dsql_query>
```

# DistributeReplicated – Part 1

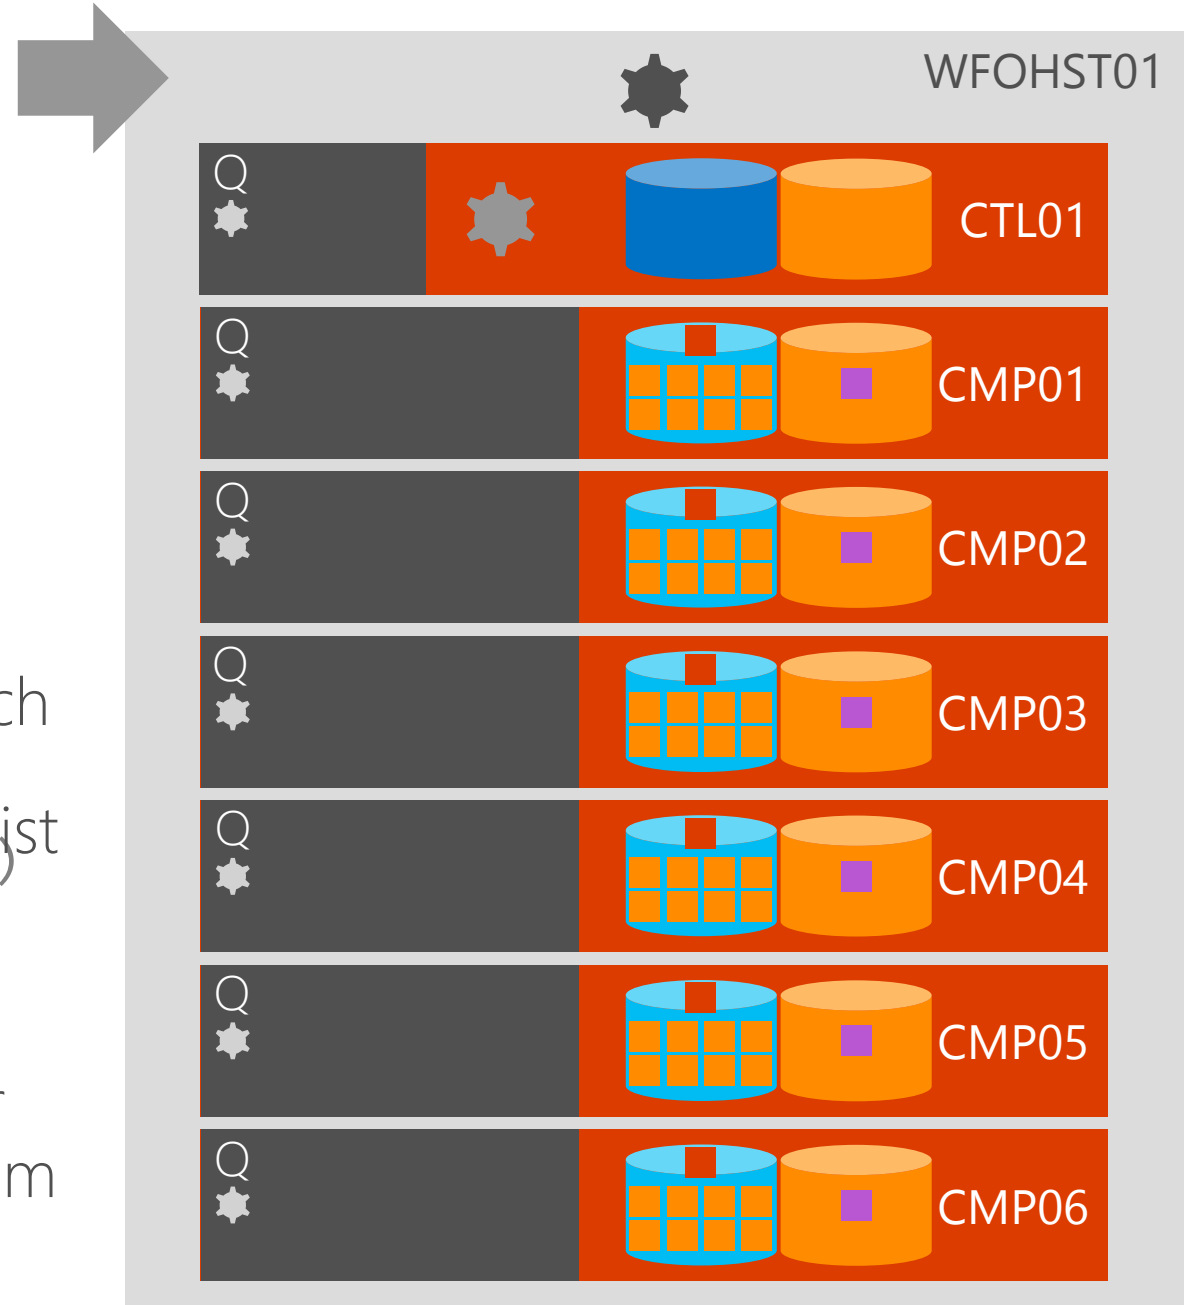
User submits write query

PDWEngine

- Creates a DSQL plan, by querying the shell database  
`CREATE TABLE TimeLog (DateNow datetime2)`
- Generates a Random ID
- Creates a replicated temp table on each node  
`INSERT INTO TimeLog`  
`SELECT SYSDATETIME()`
- Writes data into one node first to persist the values and act as the data source  
`OPTION (LABEL = 'ReplicatedDistributed')`

DMS

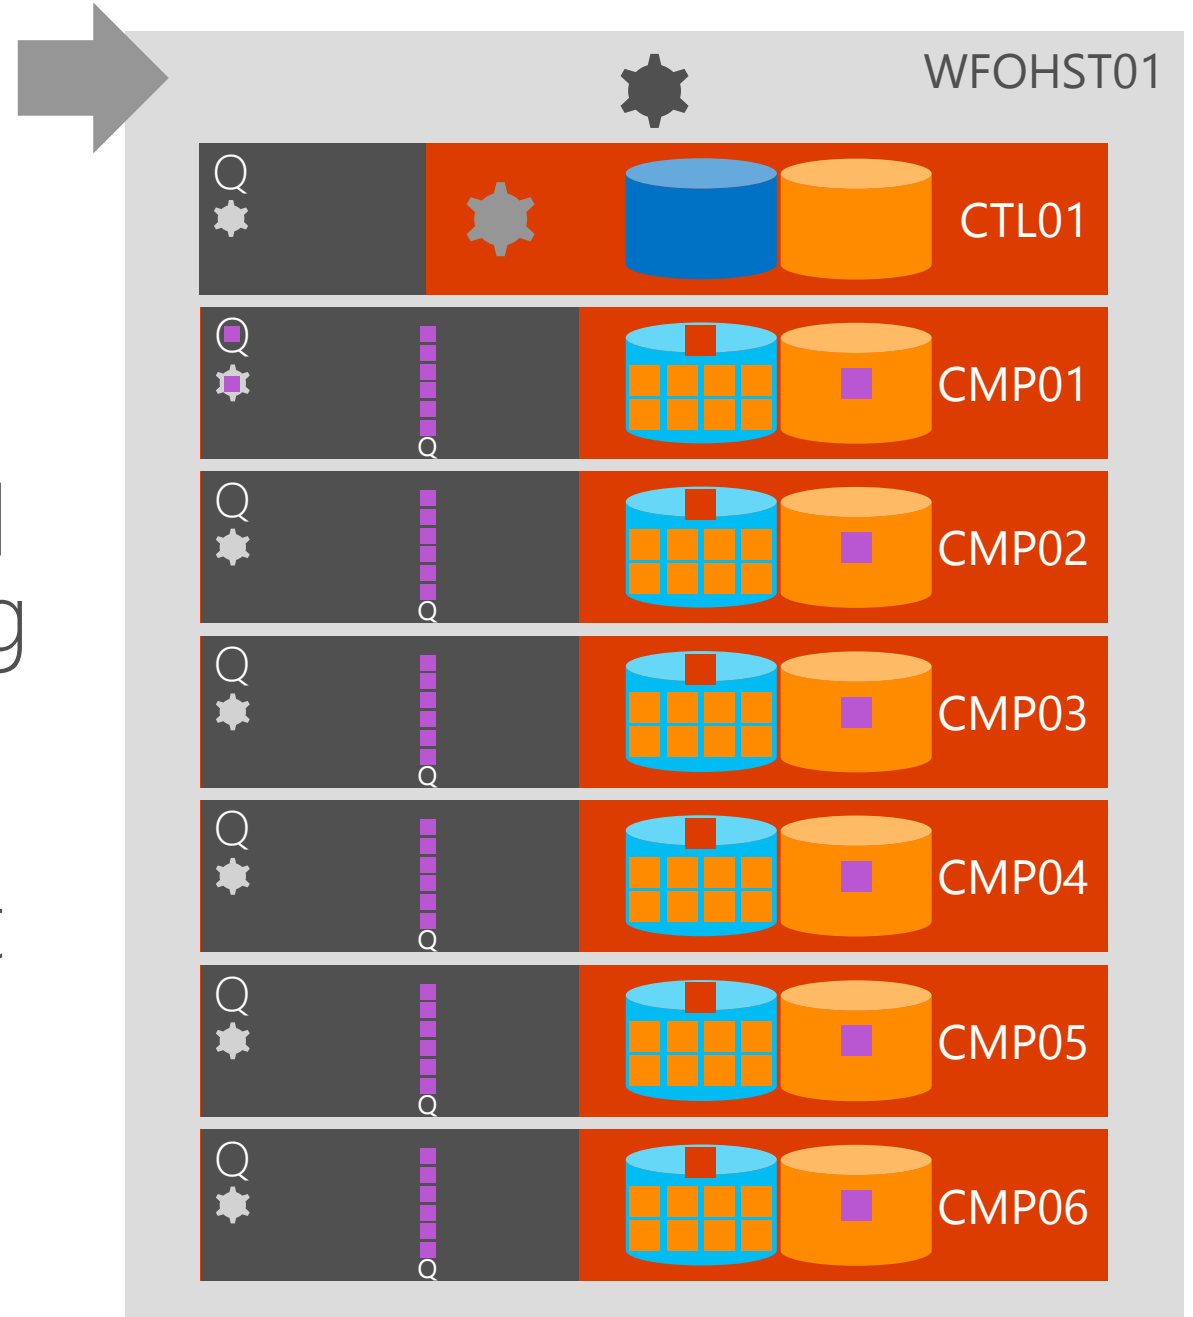
- Executes query from PDWEngine
- Reads data using Hash\_Reader worker
- Holds data in 32KB buffers placing them in a read queue for processing



# DistributeReplicated – Part 2

The DMS performs the DistributeReplicated move

- Pull buffers from the read queue on the node acting as the source
- Distributes the data to all compute nodes placing it onto the read queue



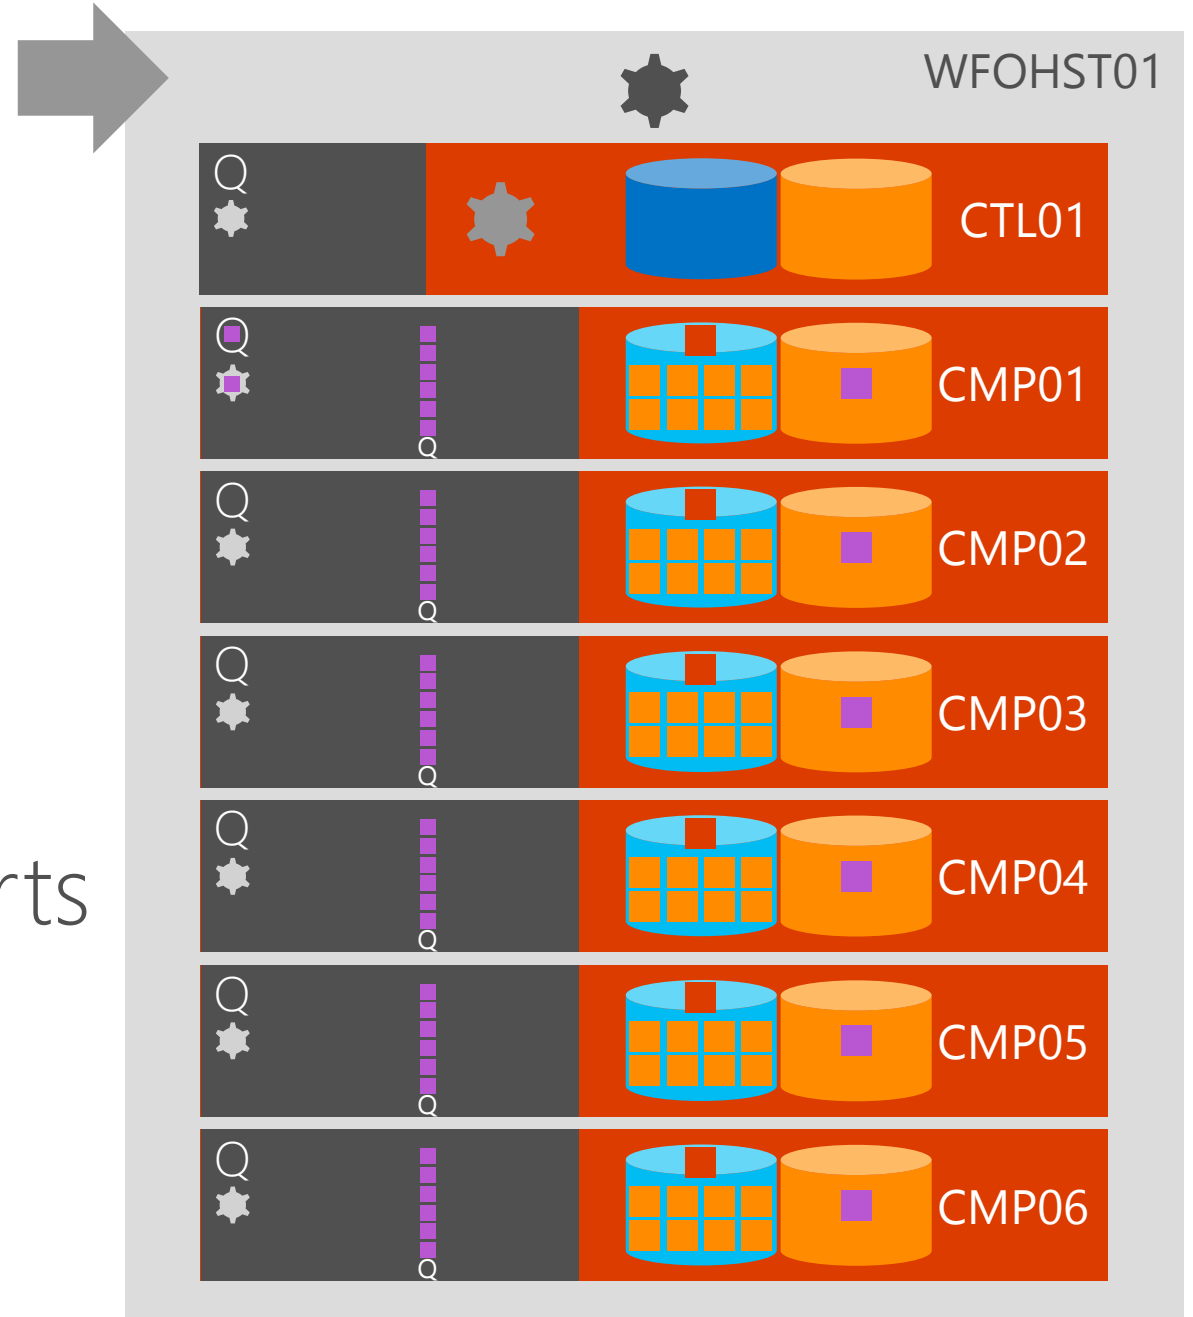
# DistributeReplicated – Part 3

## The DMS Writer Worker

- Pulls buffers from write queue
- Bulk load buffers into Tempdb “Q” table

## PDWEngine

- Submits MAXDOP(1) inserts populating target table
- Drops tables in Tempdb



# Optimising DMS

# Buffer Density

- Buffers used by query processing max 32KB
- Whole rows must fit into each buffer
  - No overflow
- All columns exploded to max size in buffer
  - i.e. VARCHAR(4000) = 4000 bytes in the buffer
- Nullable columns add eight bytes to column

Wide variable length  
character columns really  
hurt buffer density

Rationalise data types  
where possible avoiding  
nullable columns

Large numbers of nullable  
columns add up quickly.  
Akin to death by a  
thousand cuts!



# Calculating Density

```
WITH T AS
(
SELECT t.name          AS [Table_Name]
,      c.name          AS [Column_Name]
,      c.is_nullable   AS [NullValue]
,      CASE WHEN ty.name IN ('char', 'varchar', 'nchar', 'nvarchar') THEN 1
            ELSE 0
      END AS [VariableLength]
,      CASE WHEN ty.name IN('char', 'varchar')      THEN c.max_length + 1
            WHEN ty.name IN('nchar', 'nvarchar')    THEN c.max_length + 2
            WHEN ty.name IN('binary', 'varbinary')  THEN c.max_length
            ELSE ty.max_length
      END AS [DataLength]
FROM sys.tables t
JOIN sys.columns c ON t.object_id = c.object_id
JOIN ( SELECT      name
        ,          system_type_id
        ,          user_type_id
        ,          CASE WHEN name = 'time' THEN 12
                        WHEN name IN ('date', 'datetime', 'datetime2', 'datetimeoffset') THEN max_length * 2
                        WHEN name = 'smalldatetime' THEN 16
                        WHEN name = 'decimal' THEN 19
                        ELSE max_length
                      END as max_length
      ) as ty
FROM sys.types t
      ) ty ON c.system_type_id = ty.system_type_id
      AND c.user_type_id = ty.user_type_id
)
```

# Calculating Density Part 2

```
SELECT [Table_Name]
,      CAST(32768. / SUM(CASE WHEN NullValue = 1 OR VariableLength = 1 THEN 8
                           ELSE 0
                           END +[DataLength]) AS INT) AS RowsPerBuffer
,      SUM(CASE WHEN NullValue = 1 OR VariableLength = 1 THEN 8
            ELSE 0
            END + [DataLength]) AS RowSize
,      32768. % SUM(CASE WHEN NullValue = 1 OR VariableLength = 1 THEN 8
                   ELSE 0
                   END +[DataLength]) AS BufferFreeBytes
,      CAST(((32768. % SUM(CASE WHEN NullValue = 1 OR VariableLength = 1 THEN 8
                           ELSE 0
                           END +[DataLength]))) / 32768) * 100 AS DECIMAL(8,5)) AS [BufferFreePct]
FROM T
GROUP BY T.[Table_Name]
```

# Monitoring DMS

- Use performance monitor to track send and receive queue depth
- Use `sys.dm_pdw_dms_workers` to monitor throughput of the workers

```

SELECT r.request_id      AS Request_request_id      ,      w.dms_step_index      AS Worker_dms_step_index
,      r.session_id      AS Request_session_id      ,      w.pdw_node_id      AS Worker_pdw_node_id
,      r.status          AS Request_status          ,      w.distribution_id     AS Worker_distribution_id
,      r.submit_time     AS Request_submit_time     ,      w.type              AS Worker_type
,      r.start_time      AS Request_start_time      ,      w.status           AS Worker_status
,      r.end_compile_time AS Request_end_compile_time ,      w.bytes_per_sec     AS Worker_bytes_per_sec
,      r.end_time        AS Request_end_time        ,      w.bytes_processed   AS Worker_bytes_processed
,      r.total_elapsed_time AS Request_total_elapsed_time ,      w.rows_processed   AS Worker_rows_processed
,      r.[label]         AS Request_label           ,      w.start_time        AS Worker_start_time
,      r.error_id        AS Request_error_id        ,      w.end_time          AS Worker_end_time
,      d.name            AS Request_name            ,      w.total_elapsed_time AS Worker_total_elapsed_time
,      r.command         AS Request_command         ,      w.cpu_time          AS Worker_cpu_time
,      r.resource_class  AS Request_resource_class  ,      w.query_time        AS Worker_query_time
,      s.step_index      AS Step_step_index        ,      w.buffers_available AS Worker_buffers_available
,      s.operation_type  AS Step_operation_type    ,      w.dms_cpid          AS Worker_dms_cpid
,      s.distribution_type AS Step_distribution_type ,      w.sql_spid          AS Worker_sql_spid
,      s.location_type   AS Step_location_type     ,      w.error_id          AS Worker_error_id
,      s.status          AS Step_status            ,      w.source_info       AS Worker_source_info
,      s.error_id        AS Step_error_id          ,      w.destination_info  AS Worker_destination_info
,      s.start_time      AS Step_start_time
,      s.end_time        AS Step_end_time
,      s.total_elapsed_time AS Step_total_elapsed_time
,      s.row_count       AS Step_row_count
,      s.command         AS Step_command
FROM sys.dm_pdw_exec_requests r
JOIN sys.databases d          ON r.database_id = d.database_id
JOIN sys.dm_pdw_request_steps s ON r.request_id = s.request_id
JOIN sys.dm_pdw_dms_workers w ON s.request_id = w.request_id
                                AND s.step_index = w.step_index
WHERE [Label] = 'Move';

```

# Single Row Inserts & replicated tables

```
CREATE TABLE T1(col1 INT)  
WITH (DISTRIBUTION=REPLICATE)
```

```
INSERT INTO T1  
SELECT 1  
OPTION (LABEL = 'DistributeReplicatedTableMove : Triggers Movement');
```

# Movement

```
<?xml version="1.0" encoding="utf-8"?>
<dsql_query>
  <sql>INSERT INTO T1
  SELECT 1
  OPTION (LABEL = 'DistributeReplicatedTableMove : Triggers Movement')</sql>
  <dsql_operations total_cost="0" total_number_operations="5">
    <dsql_operation operation_type="RND_ID">...</dsql_operation>
    <dsql_operation operation_type="ON">...</dsql_operation>
    <dsql_operation operation_type="DISTRIBUTE_REPLICATED_TABLE_MOVE">
      <source_node>0</source_node>
      <source_statement>SELECT [T1_1].[PDWExpr1001] AS [PDWExpr1001]
FROM   (VALUES (CAST ((1) AS INT))) AS T1_1(PDWExpr1001)</source_statement>
      <destination_table>[TEMP_ID_90621]</destination_table>
    </dsql_operation>
    <dsql_operation operation_type="ON">...</dsql_operation>
    <dsql_operation operation_type="ON">...</dsql_operation>
  </dsql_operations>
</dsql_query>
```

# No Movement

```
INSERT INTO T1  
VALUES (1)  
OPTION (LABEL = 'DistributeReplicatedTableMove : No Movement');
```

```
<?xml version="1.0" encoding="utf-8"?>  
=<dsql_query>  
=<sql>EXPLAIN  
  INSERT INTO T1  
  VALUES (1)  
  OPTION (LABEL = 'DistributeReplicatedTableMove : No Movement')</sql>  
=<dsql_operations total_cost="0" total_number_operations="1">  
=<dsql_operation operation_type="ON">  
  <location permanent="true" distribution="AllComputeNodes" />  
=<sql_operations>  
  <sql_operation type="statement">INSERT INTO [Instructor].[dbo].[T1] VALUES(1)</sql_operation>  
  </sql_operations>  
=</dsql_operation>  
=</dsql_operations>  
=</dsql_query>
```

