



*The World's Largest Community  
of SQL Server Professionals*

# Inside the SQL Server Transaction Log

Brad M. McGehee

Director of DBA Education

Red Gate Software

[www.bradmcgehee.com/presentations](http://www.bradmcgehee.com/presentations)

# My Assumptions About You

- You are probably a DBA (production or developer) who has at least one year's experience.
- You have a basic understanding of how the SQL Server data cache works.
- You understand the ACID properties of databases.
- You understand the differences between the full, bulk-insert, and simple recovery models.
- You understand how to perform full, differential, and transaction log backups.

# What We Are Going to Learn Today

- Why Does the Transaction Log Exist
- How the Transaction Log Works
- How Are Log Records Written to the Transaction Log
- How the Transaction Log Can Become an IO Bottleneck
- How to Determine if the Transaction Log is a Bottleneck
- How to Deal with Transaction Log Bottlenecks and Boost Performance

# Why Does The Transaction Log Exist

- The transaction log stores a **record of all the data modifications** performed in a database.
- In the event of an unexpected shut-down, the data in the transaction log can be used during **recovery** to **roll forward any transactions completed**, but not written to the database file on disk at the time of the shut-down.
- In addition, any **uncompleted transactions** that were partially written to the database file on disk before the failure can be **rolled back** during recovery.
- Both of these actions ensure **database integrity** and the ACID properties of transactions. **This is the reason the Transaction Log exists.**

# How the Transaction Log Works

- In the following slides, I am going to start out with the **big picture** of how the transaction log works.
- For the purposes of my explanation, I am going to assume that the database uses the **full recovery model** and that **transaction log backups are performed** periodically.
- After discussing the big picture, I will provide a **specific example**, with more detail, so you better understand what happens under the covers.
- **I won't be covering every detail** of how the transaction log works, as there is not enough time. See *Inside Microsoft SQL Server 2008* for more details.

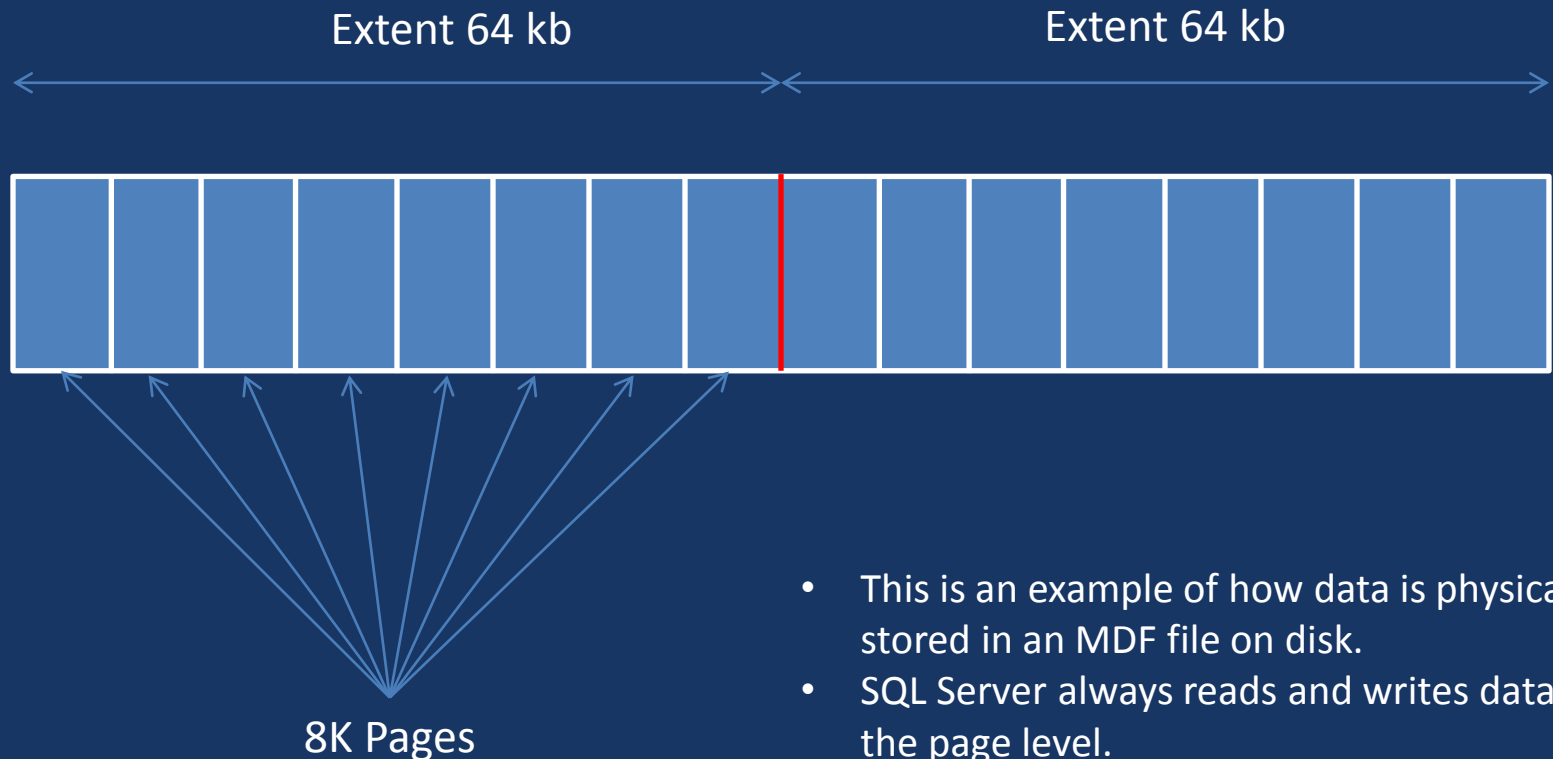
# How We Often Think of Database and T-Log Files

MDF File (let's ignore NDF files for now)

LDF File

This is a logical view, not a physical view of the MDF and LDF files.

# Looking at a MDF File in Detail



- This is an example of how data is physically stored in an MDF file on disk.
- SQL Server always reads and writes data at the page level.
- Some data pages may be full, others partially full, and others empty.
- Data may be written randomly or sequentially.

# Looking at a LDF File in Detail

**LDF files are not physically divided into pages like MDF files.** Instead, they are divided into **virtual log files (VLFs)**, often of uneven sizes from small to large. They are used to determine which parts of a log can be **set to reusable** after a transaction log backup. SQL Server creates the VLFs and determines their sizes. **LDF files are circular, and don't grow if they don't have to.**



Unlike database pages, which hold rows of data, **VLFs store log records**. Log records are **written sequentially** in the log file, into available VLFs as needed, in a circular fashion. In other words, multiple log records for a single transaction can overflow from one VLF to another. (DEMO DBCC LOGININFO)




# What are Log Records

- **Every modification** to the database creates multiple log records.
- Log records describe **what is being changed** and **what the change is**.
- A log record's **size will vary**, depending on the nature of the change.
- Log records are **written sequentially** (although mixed with other log records) to the log file and stored in VLFs, as previously discussed.
- Besides log records, **extra space in the log is reserved** in case a roll back occurs and **compensation log records** need to be written to the log. This is one reason why logging takes up more space than the actual modification.
- Log records are assigned **LSNs** (Log Sequence Number), an **ever-increasing, three-part number** that **uniquely defines the position of a log record within the transaction log**, and they are critical as they are used to identify which log records need to be rolled back or rolled forward during recovery.(DEMO DBCC LOG).

# VLFs Can Be in One of Four States

Start of Physical Log File

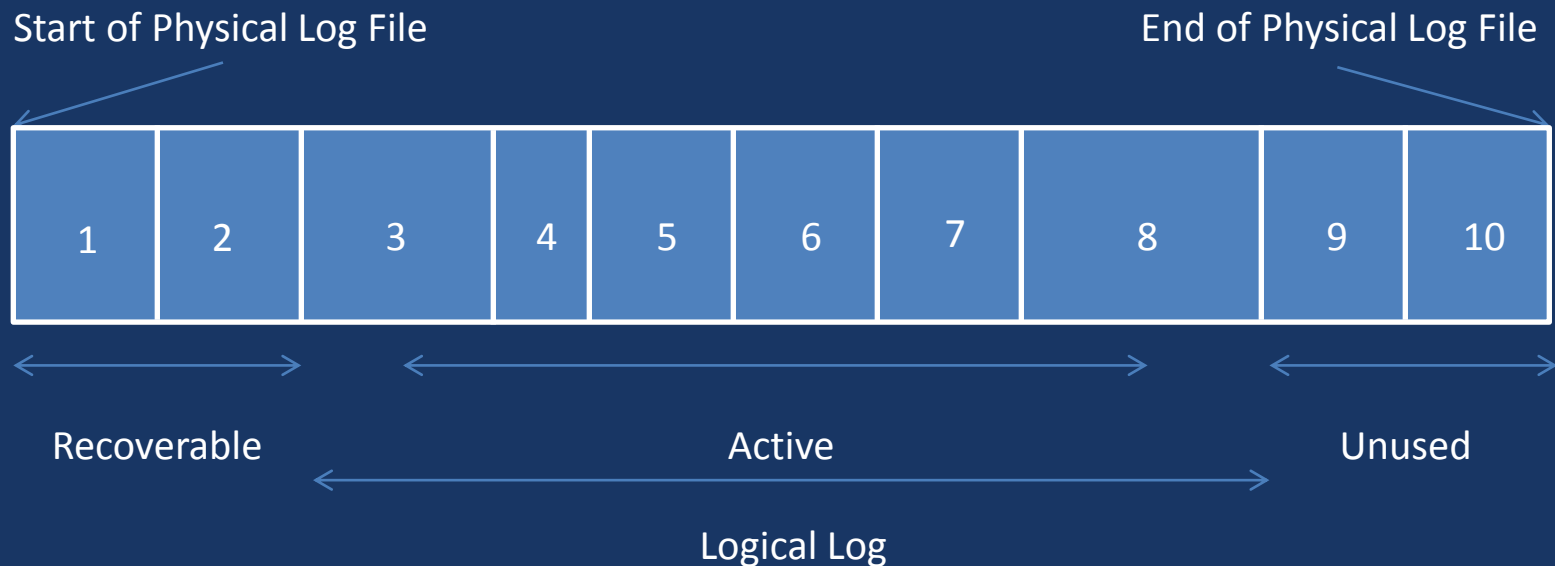
End of Physical Log File



A VLF can be in any one of four different states:

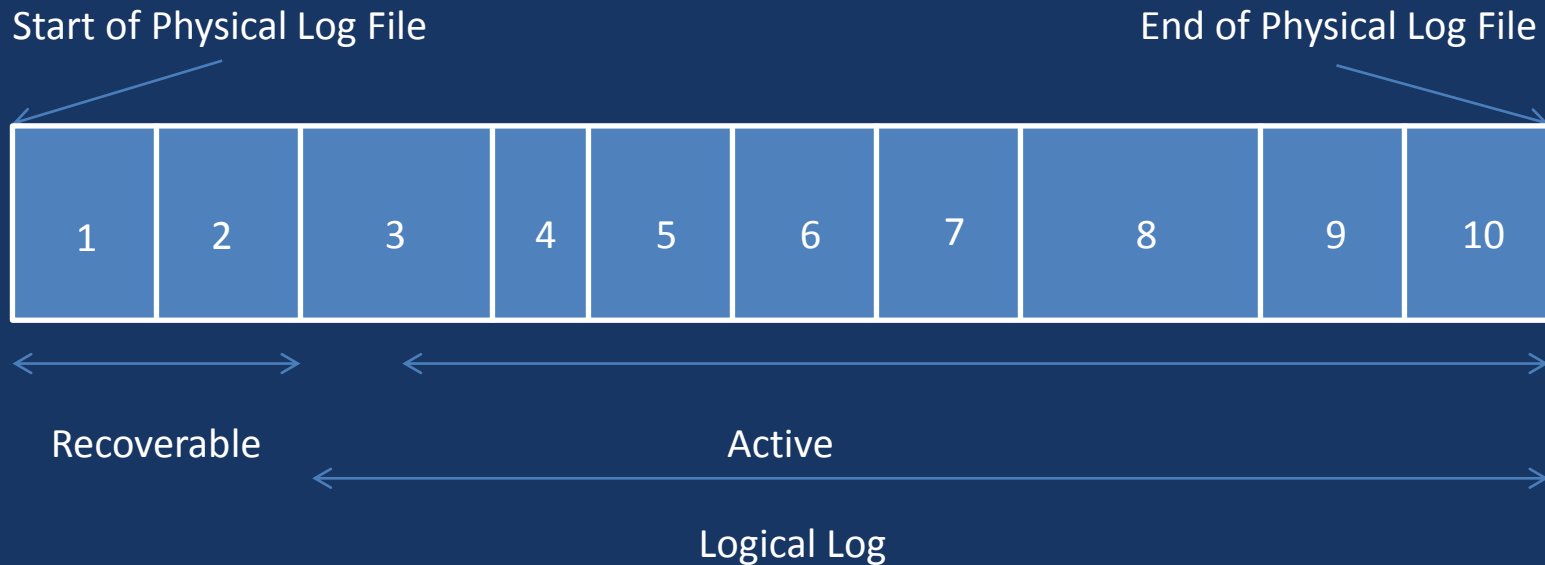
- **Active:** Includes **at least one active log record** that is part of the active log. Active records still need to be used for some purpose. (e.g. active transaction, replication, mirroring). Log backup leaves it unaffected.
- **Recoverable:** No active log records, but **not available**. Space cannot be reused until transaction log backup occurs.
- **Reusable:** Inactive log records. **Space can be reused**.
- **Unused:** **No log records** have ever been recorded in it. VLF is zeroed out.

# Simple Transaction Log Example




- Let's assume the transaction log is in the above state.
- The **logical log** is the part of the log that includes the first active VLF through the last active VLF.

# Log Records Are Added



- Let's assume that **data has been added** and the transaction log has grown to the full size of the physical LDF file.

# Transaction Log Backup is Performed



- Let's **assume that a transaction log backup** was performed.
- Notice how **recoverable VLFs (1-2) change to reusable** VLFs.
- Note that **VLF 3 is still active**, as it still has active log records.

# Log Records Added After Log Backup

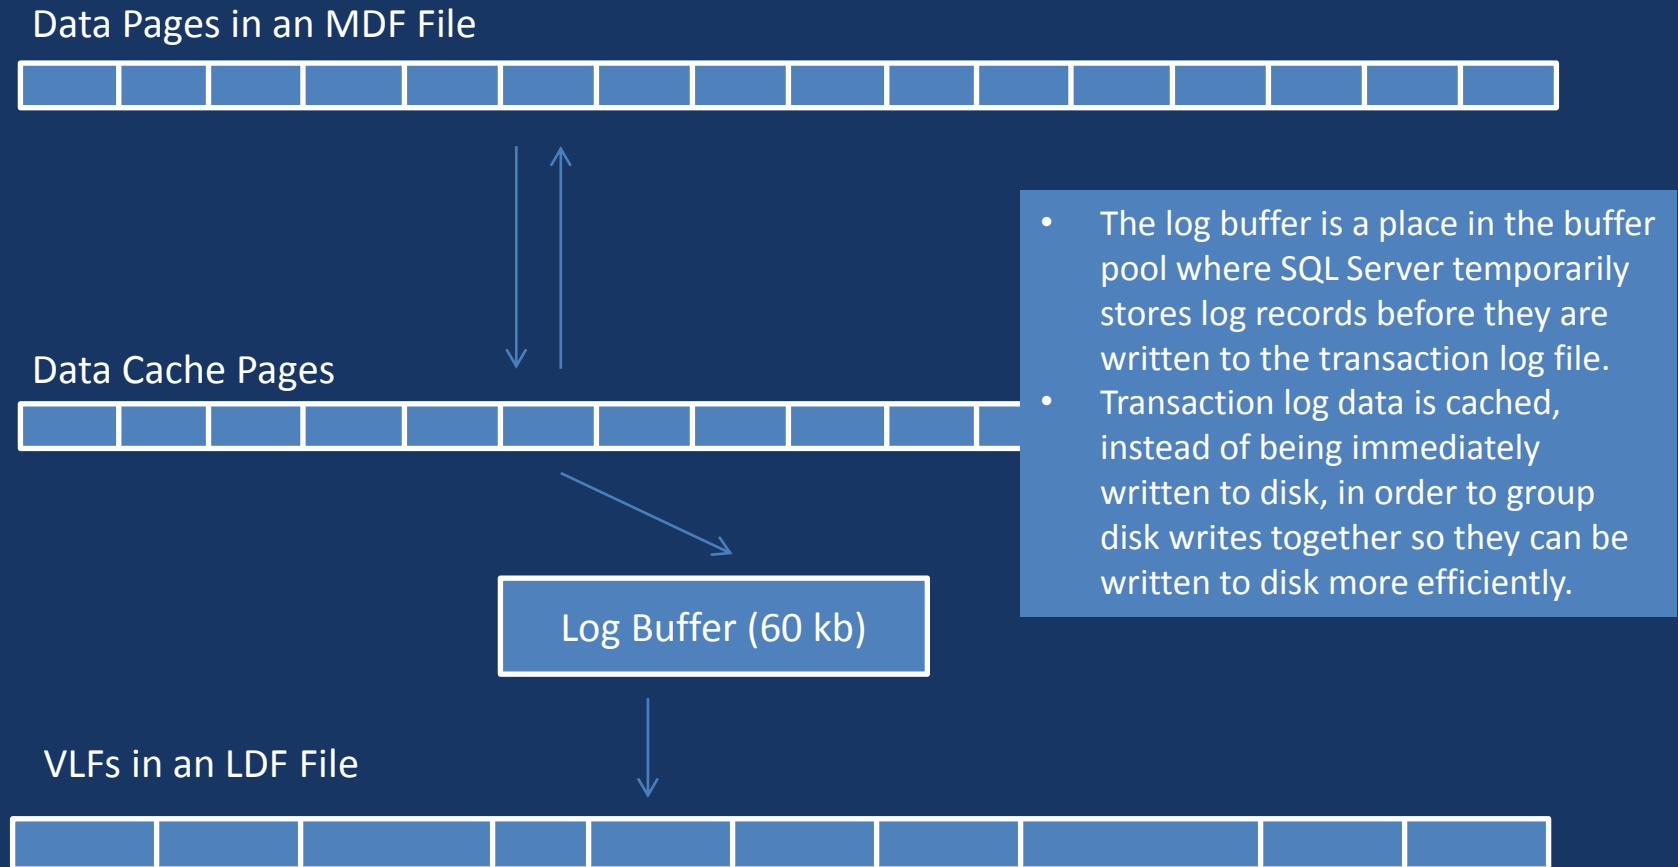


- Let's assume that more **log records have been added**.
- If reusable VLFs exist, then the **transaction log wraps** around from end to front.
- What would happen if **more data were added** now?

# How Are Log Records Written to the Transaction Log

- Now that we know a little about how the transaction log works internally, we need to **learn how transaction log records are written to the transaction log.**
- It is important to understand how log records are written to the LDF file, as how this works can directly **affect the performance** of a SQL Server instance.

# Introducing the Log Buffer





# Following the Life of a Transaction Through to the Transaction Log

- Now that we have a basic understanding of the architecture, **let's look at a specific example.**
- Let's assume that we have a **transaction that wants to INSERT a single row** into a table.
- To keep things simple, let's assume there is **no other activity** going on and that the **data cache and log buffer are empty.**

# Our Example #1

- When a transaction begins, here's what happens:
  - A **BeginTran log record** is written to the log buffer.
  - A **data page is retrieved** from disk and put into a page of the data cache.
  - SQL Server then *creates an* **INSERT log record**, which is then added to the log buffer.
  - The **new row is then INSERTed into the data page** stored in the data cache and the page is marked as being dirty.

## Our Example #2

- Now that the row has been inserted, a **Commit log record** is created and written to the log buffer, and all of the log records associated with the transaction are now **flushed (written) to the transaction log file** from the log buffer.
- This process is known as **write-ahead logging**. In other words, **all data modifications must first be made to the transaction log** before it is considered a complete transaction.
- Now that the **transaction is complete**, the client is notified that the batch completed successfully.

# A Visual Example

Data Pages in an MDF File



Page is dirty once transaction completes. Eventually a checkpoint or the lazy writer process will write (harden) the page to disk.

Data Cache Pages



Log Buffer (60 kb)



VLF Files in an LDF File



# How the Transaction Log Can Become an IO Bottleneck

- As **all data modifications** have to be **written** to the transaction log, this contributes heavily to disk IO, potentially causing **performance bottlenecks**.
- What you may not know is that not only is the **transaction log written to, it is also read**.
- Let's take a more detailed look at what data causes **writes and reads** to the transaction log, and see how it affects performance.

# Common Types of Log Write Activity

- DML activity
- Index maintenance as a result of DML activity
- Page splitting as a result of DML activity
- Auto Stats creation and updating
- DDL activity
- Database maintenance, such as rebuilding or reorganizing indexes, updating statistics, etc.

# Read Log Activity #1

- Every time a transaction has to be **rolled back**, log records have to be read so the roll back can occur.
- Creating a **database snapshot** requires crash recovery to be run, which reads the transaction log.
- Running **DBCC CHECKDB** creates a database snapshot as part of its internal process, and as we just noted, creating a database snapshot requires crash recovery to be run, and the transaction log read.
- Any kind of **backup**--full, differential, or log--all require the transaction log to be read.

# Read Log Activity #2

- If a database is in the simple recovery mode, every time a **checkpoint** occurs, the transaction log has to be read.
- **Transactional replication** reads the transaction log in order to move changes from the publisher to the distributor.
- Using **Change Data Capture** uses the transactional replication log reader to track changes, which in turn reads the transaction log.
- **Database mirroring** reads the transaction log in order to move changes from the primary to the mirror.



# How to Determine if the Transaction Log is a Bottleneck

- A variety of **tools and metrics** can be used to help diagnose if the transaction log has become an IO bottleneck.
- I will mention several, but will focus on just **two Performance Monitor counters**, as I feel these two are the best guide to determining if your transaction logs are incurring IO bottlenecks.

# These Tools Can Be Used

- DBCC OPENTRAN
- DBCC SQLPERF(LOGSPACE)
- fn\_virtualfilestats
- sys.dm\_os\_wait\_stats
- sys.dm\_io\_virtual\_file\_stats
- sys.dm\_io\_pending\_io\_requests
- Performance Monitor Counters:
  - **SQL Server:Databases** (Log Bytes Flushed/sec, Log Flushes/sec, Log Flush Wait Time)
  - **Average Disk sec/Read** (The average time, in milliseconds, of a read of data from disk; read latency)
  - **Average Disk sec/Write** (The average time, in milliseconds, of a write of data to disk; write latency)

<http://sqlcat.com/technicalnotes/archive/2008/12/09/diagnosing-transaction-log-performance-issues-and-limits-of-the-log-manager.aspx>

# When Does an IO Bottleneck Occur

- The **Average Disk sec/Read** and the **Average Disk sec/Write** counters provide you with the disk latency for the array where your transaction log(s) are located. Below is a chart I use to help me determine if IO is a problem.
  - Less than 5 ms = excellent (**this should be your goal**)
  - Between 5- 10 ms = good
  - Between 10-20 ms = okay
  - Between 20-50 ms = slow
  - Greater than 50-100 ms = potentially serious IO bottleneck
  - Greater than 100 ms = definite IO bottleneck
- Like any recommendations, the **numbers above are generic** and may not fit your particular environment.

# How to Deal with Transaction Log Bottlenecks

- There are two ways to deal with transaction log bottlenecks:
  - **Prevent them in the first place.** Implementing best practices before problems arise is much easier than after the fact.
  - If they have been identified **after the fact**, implement as many transaction log best practices as you can, but your options are often more limited, and reboots required.
- Let's discuss some best practices that you might want to consider implementing in order to prevent or reduce transaction log IO bottlenecks.

# Managing Virtual Log Files: Too Many

- When the number of VLFs in a transaction log become too many, called VLF fragmentation, **performance can suffer greatly**. This includes DML, backup, and recovery performance.
- At a point of about **100-200 VLFs** or so, log performance can begin to degrade.
- To find out how many VLFs you have in a log file, run this code: “**DBCC LOGINFO** (database\_name)”. The number of rows returned is the number of VLFs in the database’s log file.

# Managing Virtual Log Files: Too Many

- If the number of VLFs **exceeds 100-200** (depending on how large the log file is), consider manually shrinking the log file, and then manually growing it to create the “optimal” number of VLFs for your environment.
- The main reason there are **too many VLFs** is because **autogrowth is used** to grow the LDF file. A big mistake.
- By default, the **number of VLFs created** is:
  - Amounts less than 64 MB = 4 VLFs
  - Amounts of 64 MB to less than 1 GB = 8 VLFs
  - Chunks of 1 GB or larger = 16 VLFs
- For more information, see:
  - <http://sqlskills.com/blogs/kimberly/post/8-Steps-to-better-Transaction-Log-throughput.aspx>
  - <http://www.sqlskills.com/BLOGS/KIMBERLY/post/Transaction-Log-VLFs-too-many-or-too-few.aspx>

# Managing Virtual Log Files: Too Few

- If the log file will be 16 GB or larger, it is recommended that they are **manually grown 4 GB** at a time to ensure that the VLFs are not so large that it takes a long time for a transaction log to be cleared.
- In other words, **you can have too few VLFs**, which can prevent VLFs from being freed up on a timely basis for transaction log backups.
- In large databases, try to **keep VLFs no larger than 500 MB** or so.
  - EX: 16 GB created in one shot equals 16 VLFs of 1 GB each
  - EX: 16 GB created in two shots equals 32 VLFs of 500 MB each

# Remove Physical File Fragmentation

- Use the Windows **defrag.exe** command to see if the array with your transaction logs has physical file fragmentation.
- If it does, then use the same tool, or other defragging tools, to **remove it**.
- Unfortunately, you have to **take SQL Server down** to safely defrag the files.
- The best way to deal with physical file fragmentation is to **prevent it** from happening in the first place, which is discussed next.



# Preallocate Transaction Log File Size

- **Pre-sizing your log files** to their expected size offers several benefits as it prevents autogrowth from kicking in, which can cause problems. Some of the benefits include:
  - Reduces physical file fragmentation.
  - Prevents too many VLFs from being created. In other words, you control how many VLFs are created.
  - Prevents unexpected time delays (blocking) as the transaction log is growing because of autogrowth.

# Disk Partition Alignment

- For optimum performance, the partitions on your disk need to be **properly aligned**, or you risk losing a lot of your IO subsystem's performance.
- If you are not familiar with this topic, read the Microsoft white paper *Disk Partition Alignment Best Practices for SQL Server* by Jimmy May and Denny Lee. It is available at: [http://msdn.microsoft.com/en-us/library/dd758814\(v=sql.100\).aspx](http://msdn.microsoft.com/en-us/library/dd758814(v=sql.100).aspx).
- Most commonly a problem running on **Windows 2003**, or hardware that has been upgraded from **Windows 2003 to Windows 2008**.

# Select a Fast IO Subsystem

- Since we know the potential negative effect of transaction logs on the overall performance of SQL Server, we want to ensure that they run on the **fastest IO subsystem** possible, given your budget restrictions.
- Ideally, the transaction log should be located on a **RAID 10** array (RAID 1 if you can't afford RAID 10).
- In some very high-end OLTP systems, some companies have been using **RAID 1** SSDs (Solid State Drives) or Fusion-IO devices for maximum performance.

# Separate Data and Log Files

- You have probably heard this advice a hundred times, but it is important to **separate your data (MDF, NDF) and log (LDF) files on separate arrays** in order to reduce IO contention. This is one of the easiest things you can do, and is effective in reducing IO contention between data and log files.
- In a perfect world with an unlimited budget, **each log file should be on its own array**, as mixing multiple log files on the same array can introduce a lot of random reads and writes.
- The **more log files there are** on the same shared array, the bigger this problem becomes.
- At the very least, put high-performance, mission critical logs on their **own array**.

# Use Standard Performance Tuning Techniques

- Only perform as much DML or DDL as required.
- Keep transactions short.
- Minimize the possibility that a transaction has to be rolled back. Rollback sooner than later.
- Eliminate redundant or unused indexes (including indexed views), as they all need to be maintained, and maintenance requires log activity.
- Take steps to minimize page splitting. For example, use a monotonically increasing clustered index key for each of your tables, and select a fill factor that minimizes internal page splitting.

# Perform Transaction Log Backups Often

- Every time a **transaction log backup is made, unused portions of it are cleared out** and can be used for new logging.
- The more often you perform transaction log backups, **the quicker they take**, helping to reduce overhead on the transaction log.
- If you don't back up the transaction logs often, and they grow large, there will be a **larger IO hit** during the transaction log backup operation.
- So, **smaller and more frequent transaction log backups are preferred** over larger and less frequent transaction log backups.

# Use Minimally Logged Operations If Appropriate

- While all database modifications are logged in SQL Server, some database modifications can be **minimally logged**, which incur less IO overhead than fully logged operations.
- These include **BCP, BULK INSERT, INSERT INTO, SELECT INTO**, among several others.
- For these operations to be minimally logged, the database has to be using the **bulk-logged or simple recovery** models, among other requirements.

# Transactional Replication & Mirroring

- If you use either **transactional replication** or **mirroring**, the transaction log is used to move transactions between databases.
- If transactional replication or mirroring **can't keep up**, log records that haven't been moved remain active, taking up space in the transaction log.
- If you use either of these, **plan for a large transaction log**, and try to ensure that replication or mirroring don't get behind. For example, ensure that network connection between the servers is fast enough to keep up.



# Schedule Database Maintenance During Slow Times

- Full and differential database backups, index rebuilds, index reorganization, statistics updates, and DBCC CHECKDB, all can **contribute to transaction log activity**.
- Consider performing these operations at times of the day **when your server is not so busy**, and be sure you **don't overlap** running them at the same time.
- This will help to **spread out transaction log IO** throughout the day, helping to minimize potential bottlenecks.

# Take Aways From This Session

- Transaction logs play a **critical role** in SQL Server, and because of this, you need to pay attention to them.
- Transaction logs can become an **IO bottleneck** for busy databases because all activity is logged.
- By **following best practices**, you can help reduce or eliminate any transaction log bottlenecks, helping to boost the performance of your servers.
- Your SQL Servers may or may not be experiencing a transaction log bottleneck, but **you won't know until you check for yourself**.
- If you don't have a transaction log IO bottleneck, don't be complacent. **Be a proactive DBA** and implement as many of the above strategies to mitigate any potential future transaction log bottlenecks.

# Find Out More

- Free E-books on SQL Server:
  - [www.sqlservercentral.com/Books](http://www.sqlservercentral.com/Books)
- Check these websites out:
  - [www.SQLServerCentral.com](http://www.SQLServerCentral.com)
  - [www.Simple-Talk.com](http://www.Simple-Talk.com)
- Blogs:
  - [www.bradmcgehee.com](http://www.bradmcgehee.com)
  - [www.twitter.com/bradmcgehee](http://www.twitter.com/bradmcgehee)
- Contact me at:
  - [bradmcgehee@hotmail.com](mailto:bradmcgehee@hotmail.com)