

Braindump2go.70-458

Number: 70-458
Passing Score: 700
Time Limit: 160 min
File Version: 12.39



Vendor: Microsoft

Exam Code: 70-458

Exam Name: Transition Your MCTS on SQL Server 2008 to MCSA: SQL Server 2012,
Part 2 Exam

Awesome dump, pass the exam today, all questions are from this dump!!!!!!1

Exam A

QUESTION 1

You administer a Microsoft SQL Server 2012 instance. The instance contains a database that supports a retail sales application. The application generates hundreds of transactions per second and is online 24 hours per day and 7 days per week. You plan to define a backup strategy for the database. You need to ensure that the following requirements are met:

- No more than 5 minutes worth of transactions are lost.
- Data can be recovered by using the minimum amount of administrative effort.

What should you do? Choose all that apply.

- A. Configure the database to use the SIMPLE recovery model.
- B. Create a DIFFERENTIAL database backup every 4 hours.
- C. Create a LOG backup every 5 minutes.
- D. Configure the database to use the FULL recovery model.
- E. Create a FULL database backup every 24 hours.
- F. Create a DIFFERENTIAL database backup every 24 hours.

Correct Answer: BCDE

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms187048.aspx>

<http://msdn.microsoft.com/en-us/library/ms186289.aspx>

<http://msdn.microsoft.com/en-us/library/ms175477.aspx>

<http://msdn.microsoft.com/en-us/library/ms189860.aspx>

<http://msdn.microsoft.com/en-us/library/ms179314.aspx>

<http://msdn.microsoft.com/en-us/library/ms175526.aspx>

<http://msdn.microsoft.com/en-us/library/ms191539.aspx>

<http://msdn.microsoft.com/en-us/library/ms191429.aspx>

<http://msdn.microsoft.com/en-us/library/ms191253.aspx>

QUESTION 2

You administer a Microsoft SQL Server 2012 instance named SQL2012 that hosts an OLTP database of 1 terabyte in size. The database is modified by users only from Monday through Friday from 09:00 hours to 17:00 hours. Users modify more than 30 percent of the data in the database during the week. Backups are performed as shown in the following schedule:

Type	Frequency
Full	Sunday at 20:00 hours
Differential	Monday through Friday at 20:00 hours
Log	Monday through Friday between 08:00 hours and 18:00 hours

The Finance department plans to execute a batch process every Saturday at 09:00 hours. This batch process will take a maximum of 8 hours to complete. The batch process will update three tables that are 10 GB in size.

The batch process will update these tables multiple times. When the batch process completes, the Finance department runs a report to find out whether the batch process has completed correctly. You need to ensure that if the Finance department disapproves the batch process, the batch operation can be rolled back in the minimum amount of time. What should you do on Saturday?

- A. Perform a differential backup at 08:59 hours.
- B. Record the LSN of the transaction log at 08:59 hours. Perform a transaction log backup at 17:01 hours.
- C. Create a database snapshot at 08:59 hours.
- D. Record the LSN of the transaction log at 08:59 hours. Perform a transaction log backup at 08:59 hours.
- E. Create a marked transaction in the transaction log at 08:59 hours. Perform a transaction log backup at 17:01 hours.
- F. Create a marked transaction in the transaction log at 08:59 hours. Perform a transaction log backup at 08:59 hours.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

--Burgos. NO

A table about backup strategy is missing. I recover a table from equal question from 70-857. SNAPSHOT sounds to me as correct answer.

--\Burgos.

QUESTION 3

You are designing a data warehouse hosted on SQL Azure. The data warehouse currently includes the dimUser and dimDistrict dimension tables and the factSales fact table. The dimUser table contains records for each user permitted to run reports against the warehouse, and the dimDistrict table contains information about sales districts. The system is accessed by users from certain districts, as well as by area supervisors and users from the corporate headquarters. You need to design a table structure to ensure that certain users can see sales data for only certain districts. Some users must be permitted to see sales data from multiple districts. What should you do?

- A. Create a userDistrict table that contains primary key columns from the dimUser and dimDistrict tables.
- B. Partition the factSales table on the district column.
- C. Add a district column to the dimUser table.
- D. For each district, create a view of the factSales table that includes a WHERE clause for the district.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

QUESTION 4

You develop three Microsoft SQL Server 2012 databases named Database1, Database2, and Database3. You have permissions on both Database1 and Database2. You plan to write and deploy a stored procedure named dbo.usp_InsertEvent in Database3. dbo.usp_InsertEvent must execute other stored procedures in the other databases. You need to ensure that callers that do not have permissions on Database1 or Database2 can execute the stored procedure. Which Transact-SQL statement should you use?

- A. USE Database2

- B. EXECUTE AS OWNER
- C. USE Database1
- D. EXECUTE AS CALLER

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

Verified answer as correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188354.aspx>

<http://blog.sqlauthority.com/2007/10/06/sql-server-executing-remote-stored-procedure-calling-stored-procedure-on-linked-server/>

QUESTION 5

You administer a Microsoft SQL Server 2012 failover cluster that contains two nodes named Node A and Node B. A single instance of SQL Server is installed on the cluster. An additional node named Node C has been added to the existing cluster. You need to ensure that the SQL Server instance can use all nodes of the cluster. What should you do?

- A. Create a ConfigurationFile.ini file from Node B, and then run the AddNode command-line tool on Node A.
- B. Use Node A to install SQL Server on Node C.
- C. Run the Add Node to SQL Server Failover Cluster Wizard on Node C.
- D. Use Cluster Administrator to add a new Resource Group to Node B.
- E. Run the New SQL Server stand-alone installation Wizard on Node C.
- F. Use Node B to install SQL Server on Node C.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

--Burgos. NO

This question appears two times in original VCE file with different set of choices, but two choices appears in both questions (the correct and the shorter "Use Node B..." choice). I discharge one of questions and put all different choices in this questions.

--\Burgos.

Verified answer as correct.

References:

<http://msdn.microsoft.com/en-us/library/hh231721.aspx>

<http://msdn.microsoft.com/en-us/library/ms189134.aspx>

QUESTION 6

You administer all the deployments of Microsoft SQL Server 2012 in your company. You need to ensure that an OLTP database that uses a storage area network (SAN) remains available if any of the servers fail. You also need to minimize the amount of storage used by the database. Which configuration should you use?

- A. Two servers configured in different data centers
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
One server configured as an Active Secondary

- B. SQL Server that includes an application database configured to perform transactional replication
- C. Two servers configured in the same data center
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
One server configured as an Active Secondary
- D. Two servers configured in different data centers
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
- E. Two servers configured in the same data center
A primary server configured to perform log-shipping every 10 minutes
A backup server configured as a warm standby
- F. Two servers configured on the same subnet
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
- G. SQL Server that includes an application database configured to perform snapshot replication
- H. Two servers configured in a Windows Failover Cluster in the same data center
SQL Server configured as a clustered instance

Correct Answer: H

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ff650328.aspx>

<http://msdn.microsoft.com/en-us/library/ms189134>

<http://msdn.microsoft.com/en-us/library/hh750283.aspx>

<http://msdn.microsoft.com/en-us/library/ff878716.aspx>

QUESTION 7

You administer all the deployments of Microsoft SQL Server 2012 in your company. A database contains a large product catalog that is updated periodically. You need to be able to send the entire product catalog to all branch offices on a monthly basis. Which configuration should you use?

- A. Two servers configured in the same data center
A primary server configured to perform log-shipping every 10 minutes
A backup server configured as a warm standby
- B. Two servers configured in the same data center
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
One server configured as an Active Secondary
- C. SQL Server that includes an application database configured to perform snapshot replication
- D. Two servers configured in different data centers
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
One server configured as an Active Secondary
- E. Two servers configured on the same subnet
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
- F. Two servers configured in different data centers
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
- G. SQL Server that includes an application database configured to perform transactional replication
- H. Two servers configured in a Windows Failover Cluster in the same data center
SQL Server configured as a clustered instance

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms151832.aspx>

QUESTION 8

You administer all the deployments of Microsoft SQL Server 2012 in your company. You need to ensure that data changes are sent to a non-SQL Server database server in near real time. You also need to ensure that data on the primary server is unaffected. Which configuration should you use?

- A. SQL Server that includes an application database configured to perform transactional replication
- B. Two servers configured in different data centers
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
- C. Two servers configured in different data centers
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
One server configured as an Active Secondary
- D. SQL Server that includes an application database configured to perform snapshot replication
- E. Two servers configured in the same data center
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
One server configured as an Active Secondary
- F. Two servers configured on the same subnet
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
- G. Two servers configured in a Windows Failover Cluster in the same data center
SQL Server configured as a clustered instance
- H. Two servers configured in the same data center
A primary server configured to perform log-shipping every 10 minutes
A backup server configured as a warm standby

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://technet.microsoft.com/en-us/library/ms151738.aspx>

QUESTION 9

You administer all the deployments of Microsoft SQL Server 2012 in your company. You need to ensure that an OLTP database that includes up-to-the-minute reporting requirements can be off-loaded from the primary database to another server. You also need to be able to add indexes to the secondary database. Which configuration should you use?

- A. Two servers configured in different data centers
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
One server configured as an Active Secondary
- B. Two servers configured in the same data center

- SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
One server configured as an Active Secondary
- C. Two servers configured in the same data center
A primary server configured to perform log-shipping every 10 minutes
A backup server configured as a warm standby
 - D. Two servers configured in different data centers
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
 - E. Two servers configured on the same subnet
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
 - F. SQL Server that includes an application database configured to perform transactional replication
 - G. SQL Server that includes an application database configured to perform snapshot replication
 - H. Two servers configured in a Windows Failover Cluster in the same data center
SQL Server configured as a clustered instance

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to this reference, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ff878253.aspx>

QUESTION 10

You administer all the deployments of Microsoft SQL Server 2012 in your company. You have two servers in the same data center that hosts your production database. You need to ensure that the database remains available if a catastrophic server failure or a disk failure occurs. You also need to maintain transactional consistency of the data across both servers. You need to achieve these goals without manual intervention. Which configuration should you use?

- A. Two servers configured in a Windows Failover Cluster in the same data center
SQL Server configured as a clustered instance
- B. SQL Server that includes an application database configured to perform transactional replication
- C. Two servers configured in the same data center
A primary server configured to perform log-shipping every 10 minutes
A backup server configured as a warm standby
- D. Two servers configured in different data centers
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
One server configured as an Active Secondary
- E. Two servers configured in the same data center
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
One server configured as an Active Secondary
- F. Two servers configured in different data centers
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
- G. SQL Server that includes an application database configured to perform snapshot replication
- H. Two servers configured on the same subnet
SQL Server Availability Group configured in Synchronous-Commit Availability Mode

Correct Answer: H

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ff877931.aspx>

QUESTION 11

You are developing a data flow transformation to merge two data sources. One source contains product data and the other source contains data about the country in which the product was manufactured. Both data sources contain a two-character CountryCode column and both use SQL Server. Both data sources contain an ORDER BY clause to sort the data by the CountryCode column in ascending order. You use a Merge Join transformation to join the data.



You need to ensure that the Merge Join transformation works correctly without additional transformations. What should you do? (Each answer presents a part of the solution. Choose all that apply.)

- A. Set the appropriate SortKeyPosition properties on the data sources.
- B. Set the ValidateExternalMetaData property on the Merge Join transformation to True.
- C. Set the IsSorted property on both data sources.
- D. Set the MaxBuffersPerInput property on the Merge Join transformation to 2.

Correct Answer: AC

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms141775.aspx>

<http://msdn.microsoft.com/en-us/library/ms137653.aspx>

<http://siddhumehta.blogspot.com/2009/05/validateexternalmetadata-property.html>

<http://msdn.microsoft.com/en-us/library/ms135950.aspx>

QUESTION 12

You are developing a SQL Server Integration Services (SSIS) package to load data into a data warehouse.



You need to establish the correct order for loading each table to maximize parallel processing. Which order should you use?

- A. 1. DimCurrency, DimScenario, DimAccount in parallel
2. DimOrganization
3. FactFinance
- B. 1. DimCurrency, DimOrganization in parallel
2. DimScenario, DimAccount in parallel
3. FactFinance
- C. 1. DimCurrency, FactFinance in parallel
2. DimOrganization, DimScenario, DimAccount in parallel
- D. 1. FactFinance
2. DimOrganization, DimScenario, DimAccount in parallel
3. DimCurrency
- E. 1. DimCurrency
2. DimOrganization
3. DimScenario, DimAccount in parallel
4. FactFinance

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms139892.aspx>

<http://msdn.microsoft.com/en-us/library/ms141261.aspx>

QUESTION 13

You are developing a SQL Server Integration Services (SSIS) package. You need to design a package to change a variable value during package execution by using the least amount of development effort. What should you use?

- A. Expression task

- B. Script task
- C. Execute SQL task
- D. Execute Process task
- E. Term Extraction transformation

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/hh213137.aspx>

QUESTION 14

You are editing a SQL Server Integration Services (SSIS) package that contains three Execute SQL tasks and no other tasks. The three Execute SQL tasks modify products in staging tables in preparation for a data warehouse load. The package and all three Execute SQL product tasks have their TransactionOption property set to Supported. You need to ensure that if any of the three Execute SQL product tasks fail, all three tasks will roll back their changes. What should you do?

- A. Change the TransactionOption property of the package to Required.
- B. Change the TransactionOption property of all three Execute SQL product tasks to Required.
- C. Move the three Execute SQL product tasks into a Foreach Loop container.
- D. Move the three Execute SQL product tasks into a Sequence container.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms137690.aspx>

<http://msdn.microsoft.com/en-us/library/ms141144.aspx>

QUESTION 15

You administer a Microsoft SQL Server 2012 server that has SQL Server Integration Services (SSIS) installed. You plan to deploy new SSIS packages to the server. The SSIS packages use the Project Deployment Model together with parameters and Integration Services environment variables. You need to configure the SQL Server environment to support these packages. What should you do?

- A. Create SSIS configuration files for the packages.
- B. Create an Integration Services catalog.
- C. Install Data Quality Services.
- D. Install Master Data services.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES NO

--\BURGOS

According to these references, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/hh479588.aspx>

<http://msdn.microsoft.com/en-us/library/hh213290.aspx>

<http://msdn.microsoft.com/en-us/library/hh213373.aspx>

QUESTION 16

You are developing a SQL Server Integration Services (SSIS) package that imports data into a data warehouse. You add an Execute SQL task to the control flow. The task must execute a simple INSERT statement. The task has the following requirements:

- The INSERT statement must use the value of a string package variable. The variable name is StringVar.
- The Execute SQL task must use an OLE DB Connection Manager.

In the Parameter Mapping tab of the Execute SQL task, StringVar has been added as the only parameter. You must configure the SQLStatement property of the Execute SQL task. Which SQL statement should you use?

- A. `INSERT INTO dbo.Table (variablevalue) VALUES ($Project::StringVar)`
- B. `INSERT INTO dbo.Table (variablevalue) VALUES (@StringVar)`
- C. `INSERT INTO dbo.Table (variablevalue) VALUES ($Package::StringVar)`
- D. `INSERT INTO dbo.Table (variablevalue) VALUES (?)`

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms141003.aspx>

<http://msdn.microsoft.com/en-us/library/ms140355.aspx>

<http://msdn.microsoft.com/en-us/library/cc280502.aspx>

QUESTION 17

You are developing a SQL Server Integration Services (SSIS) package to load data into a SQL Server table on Server A. The package includes a data flow and is executed on ServerB. The destination table has its own identity column. The destination data load has the following requirements:

- The identity values from the source table must be used.
- Default constraints on the destination table must be ignored.
- Batch size must be 100,000 rows.

You need to add a destination and configure it to meet the requirements. Which destination should you use?

- A. ADO NET Destination with Bulk Insert
- B. SQL Server Destination
- C. OLE DB Destination with Fast Load
- D. ADO NET Destination without Bulk Insert

E. OLE DB Destination without Fast Load

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms141237.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms139821.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms141095.aspx>

QUESTION 18

You are creating a SQL Server Master Data Services (MDS) model for a company. The source data for the company is stored in a single table that contains the manager-to-subordinate relationships. You need to create a hierarchy representing the organizational structure of the company. Which hierarchy type should you use?

- A. Natural
- B. Explicit
- C. Parent
- D. Recursive

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

Reference: <http://technet.microsoft.com/en-us/library/ff487006.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ee633747.aspx>

Reference: <http://technet.microsoft.com/en-us/library/ee633759.aspx>

QUESTION 19

You develop a SQL Server Integration Services (SSIS) package that imports SQL Azure data into a data warehouse every night. The SQL Azure data contains many misspellings and variations of abbreviations. To import the data, a developer used the Fuzzy Lookup transformation to choose the closest-matching string from a reference table of allowed values. The number of rows in the reference table is very large. If no acceptable match is found, the Fuzzy Lookup transformation passes a null value. The current setting for the Fuzzy Lookup similarity threshold is 0.50. Many values are incorrectly matched. You need to ensure that more accurate matches are made by the Fuzzy Lookup transformation without degrading performance. What should you do?

- A. Change the Exhaustive property to True.
- B. Change the similarity threshold to 0.85.
- C. Change the similarity threshold to 0.40.
- D. Increase the maximum number of matches per lookup.
- E. Decrease the maximum number of matches per lookup.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

--Burgos. YES, (same differences, four choices correct was .85 instead .55)
I join two equal questions that had different choices.
--\Burgos.

According to this reference, this answer looks correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms137786.aspx>

QUESTION 20

You are designing a SQL Server Integration Services (SSIS) package that uploads a file to a table named Orders in a SQL Azure database. The company's auditing policies have the following requirements:

- An entry must be written to a dedicated SQL Server log table named OrderLog.
- The entry must be written as soon as the file upload task completes.

You need to meet the company's policy requirements. Which event handler should you use?

- A. OnWarning
- B. OnVariableValueChanged
- C. OnExecStatusChanged
- D. OnPostExecute
- E. OnProgress
- F. OnInformation
- G. OnComplete

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (four last choices)
--\BURGOS

--Burgos.
I join two equal questions that had different choices.
--\Burgos.

According to this reference, this answer looks correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms140223.aspx>

QUESTION 21

You are developing a SQL Server Integration Services (SSIS) package that imports data from a relational database to a data warehouse. You are importing data from a relational table named Projects. The table has change data capture enabled on all columns. You need to process only the most recent values from rows that have been inserted or updated since the previous execution of the package. Which query should you use as the data source?

- A. `SELECT * FROM cdc.fn_cdc_get_all_changes_Projects (@from_lsn, @to_lsn, N' all update old')`
- B. `SELECT * FROM cdc.fn_cdc_get_all_changes_Projects (@from_lsn, @to_lsn, N' all')`
- C. `SELECT * FROM cdc.fn_cdc_get_net_changes_Projects (@from_lsn, @to_lsn)`
- D. `SELECT * FROM cdc.Projects_CT WHERE @from_lsn >= _$start_lsn AND @to_lsn < _$start_lsn`

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

Reference: <http://msdn.microsoft.com/en-us/library/bb522511.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/bb510627.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/cc645937.aspx>

QUESTION 22

You are designing a SQL Server Integration Services (SSIS) data flow to load sales transactions from a source system into a data warehouse hosted on SQL Azure. One of the columns in the data source is named ProductCode. Some of the data to be loaded will reference products that need special processing logic in the data flow. You need to enable separate processing streams for a subset of rows based on the source product code. Which data flow transformation should you use?

- A. Source Assistant
- B. Audit
- C. Destination Assistant
- D. Script Component

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms137640.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms141150.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ff929138.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ff929116.aspx>

QUESTION 23

You are designing a SQL Server Integration Services (SSIS) package that uses the Fuzzy Lookup transformation. The reference data to be used in the transformation does not change. You need to reuse the Fuzzy Lookup match index to increase performance and reduce maintenance. What should you do?

- A. Select the GenerateAndPersistNewIndex option in the Fuzzy Lookup Transformation Editor.
- B. Select the GenerateNewIndex option in the Fuzzy Lookup Transformation Editor.
- C. Select the DropExistingMatchIndex option in the Fuzzy Lookup Transformation Editor.
- D. Execute the sp_FuzzyLookupTableMaintenanceUninstall stored procedure.
- E. Execute the sp_FuzzyLookupTableMaintenanceInvoke stored procedure.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES
--\BURGOS

According to this reference, this answer looks correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms137786.aspx>

QUESTION 24

You are creating a SQL Server Integration Services (SSIS) package to retrieve product data from two different sources. One source is hosted in a SQL Azure database. Each source contains products for different distributors. Products for each distributor source must be combined for insertion into a single product table destination. You need to select the appropriate data flow transformation to meet this requirement. Which transformation types should you use? (Each correct answer presents a complete solution. Choose all that apply.)

- A. Multicast
- B. Merge Join
- C. Term Extraction
- D. Union All
- E. Merge

Correct Answer: DE

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

The question is confused, I'm Not sure if the products were the same products (from different distributors, MERGE) or different products from different distributors (UNION ALL). I answered equal this DUMP.

--\BURGOS

According to these references, this answer looks correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms141703.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms141775.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms141020.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms141809.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms137701.aspx>

QUESTION 25

You are implementing the indexing strategy for a fact table in a data warehouse. The fact table is named Quotes. The table has no indexes and consists of seven columns:

- [ID]
- [QuoteDate]
- [Open]
- [Close]
- [High]
- [Low]
- [Volume]

Each of the following queries must be able to use a columnstore index:

```
SELECT AVG ([Close]) AS [AverageClose]
FROM Quotes
WHERE [QuoteDate] BETWEEN '20100101' AND '20101231'.
```

```
SELECT AVG([High] - [Low]) AS [AverageRange]
FROM Quotes
WHERE [QuoteDate] BETWEEN '20100101' AND '20101231'.
```

```
SELECT SUM([Volume]) AS [SumVolume]
FROM Quotes
WHERE [QuoteDate] BETWEEN '20100101' AND '20101231'
```

You need to ensure that the indexing strategy meets the requirements. The strategy must also minimize the number and size of the indexes. What should you do?

- A. Create one columnstore index that contains [ID], [Close], [High], [Low], [Volume], and [QuoteDate].
- B. Create three columnstore indexes:
 - One containing [QuoteDate] and [Close]
 - One containing [QuoteDate], [High], and [Low]
 - One containing [QuoteDate] and [Volume]
- C. Create one columnstore index that contains [QuoteDate], [Close], [High], [Low], and [Volume].
- D. Create two columnstore indexes:
 - One containing [ID], [QuoteDate], [Volume], and [Close]
 - One containing [ID], [QuoteDate], [High], and [Low]

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

Reference: <http://msdn.microsoft.com/en-us/library/gg492088.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/gg492153.aspx>

QUESTION 26

You are developing a SQL Server Integration Services (SSIS) package that imports data into a data warehouse hosted on SQL Azure. The package uses a Foreach container to process text files found in a folder. The package must be deployed to a single server by using the Project Deployment model. Multiple SQL Server Agent jobs call the package. Each job is executed on a different schedule. Each job passes a different folder path to the package. You need to configure the package to accept the folder path from each job. Which package configuration should you use?

- A. .dtsConfig file
- B. Registry Entry
- C. Parent Package Variable
- D. XML Configuration File
- E. Environment Variable

Correct Answer: E

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms345179.aspx>

<http://msdn.microsoft.com/en-us/library/ms141708.aspx>
<http://msdn.microsoft.com/en-us/library/hh213214.aspx>
<http://msdn.microsoft.com/en-us/library/hh213296.aspx>
<http://msdn.microsoft.com/en-us/library/hh213293.aspx>

QUESTION 27

You are creating a SQL Server Master Data Services (MDS) model. This model is used to store a master list of products. An attribute must be added to the Product entity to define the sales manager responsible for each product. You need to create an attribute in the Product entity that prevents users from entering invalid sales manager values. Which type of attribute should you create?

- A. Derived
- B. Domain-based
- C. User-defined
- D. Parent
- E. Explicit
- F. Recursive

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/bb190163.aspx>
<http://msdn.microsoft.com/en-us/library/ee633737.aspx>
<http://msdn.microsoft.com/en-us/library/ee633759.aspx>
<http://msdn.microsoft.com/en-us/library/ee633745.aspx>
<http://msdn.microsoft.com/en-us/library/ee633724.aspx>
<http://msdn.microsoft.com/en-us/library/ee633733.aspx>
<http://msdn.microsoft.com/en-us/library/ff487058.aspx>

QUESTION 28

You are designing a data warehouse with two fact tables. The first table contains sales per month and the second table contains orders per day. Referential integrity must be enforced declaratively. You need to design a solution that can join a single time dimension to both fact tables. What should you do?

- A. Join the two fact tables.
- B. Merge the fact tables.
- C. Create a time dimension that can join to both fact tables at their respective granularity.
- D. Create a surrogate key for the time dimension.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms174537.aspx>
<http://technet.microsoft.com/en-us/library/ms174832.aspx>
<http://msdn.microsoft.com/en-us/library/ms174884.aspx>
<http://decipherinfosys.wordpress.com/2007/02/01/surrogate-keys-vs-natural-keys-for-primary-key/>
<http://www.agiledata.org/essays/keys.html>
<http://www.databasejournal.com/features/mssql/article.php/3922066/SQL-Server-Natural-Key-Verses-Surrogate-Key.htm>
<http://www.jamesserra.com/archive/2012/01/surrogate-keys/>

QUESTION 29

You are completing the installation of the Data Quality Server component of SQL Server Data Quality Services (DQS). You need to complete the post-installation configuration. What should you do?

- A. Run the dbimpexp.exe command.
- B. Install the data providers that are used for data refresh.
- C. Run the DQSInstaller.exe command.
- D. Install the Analysis Services OLE DB Provider.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (another version of this question)
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ff877917.aspx>
<http://msdn.microsoft.com/en-us/library/gg492277.aspx>

QUESTION 30

You are completing the installation of the Data Quality Server component of SQL Server Data Quality Services (DQS). You need to complete the post-installation configuration. What should you do?

- A. Run the Configuration component in the Data Quality Client.
- B. Install ADOMD.NET.
- C. Run the Data Quality Server Installer.
- D. Make the data available for DQS operations.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (this version of this question)
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ff877917.aspx>
<http://msdn.microsoft.com/en-us/library/gg492277.aspx>

QUESTION 31

You are using the Knowledge Discovery feature of the Data Quality Services (DQS) client application to modify an existing knowledge base. In the mapping configuration, two of the three columns are mapped to existing domains in the knowledge base. The third column, named Team Type, does not yet have a domain. You need to complete the mapping of the Team Type column. What should you do?

- A. Add a column mapping for the Team Type column.
- B. Add a domain for the Team Type column.
- C. Map a composite domain to the source column.
- D. Create a composite domain that includes the Team Type column.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ff877917.aspx>

<http://msdn.microsoft.com/en-us/library/hh213015.aspx>

<http://msdn.microsoft.com/en-us/library/gg524799.aspx>

QUESTION 32

You are developing a SQL Server Integration Services (SSIS) project that copies a large amount of rows from a SQL Azure database. The project uses the Package Deployment Model. This project is deployed to SQL Server on a test server. You need to ensure that the project is deployed to the SSIS catalog on the production server. What should you do?

- A. Add an OnError event handler to the SSIS project.
- B. Use an msi file to deploy the package on the server.
- C. Open a command prompt and run the gacutil command.
- D. Open a command prompt and run the dtutil /copy command.
- E. Open a command prompt and run the dtexec /rep /conn command.
- F. Open a command prompt and run the dtexec /dumperror /conn command.
- G. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- H. Create a reusable custom logging component and use it in the SSIS project.
- I. Configure the SSIS solution to use the Project Deployment Model.
- J. Configure the output of a component in the package data flow to use a data tap.
- K. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL Server.

Correct Answer: I

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (same context, five questions, one of each answer: "data tap", "event handler", "dtexecui.exe/SQLLogProvider", "gacutil" and "wizard")

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/hh231102.aspx>
<http://msdn.microsoft.com/en-us/library/hh213290.aspx>
<http://msdn.microsoft.com/en-us/library/hh213373.aspx>

QUESTION 33

You are developing a SQL Server Integration Services (SSIS) package. To process complex scientific data originating from a SQL Azure database, a custom task component is added to the project. You need to ensure that the custom component is deployed on a test environment correctly. What should you do?

- A. Add an OnError event handler to the SSIS project.
- B. Use an msi file to deploy the package on the server.
- C. Open a command prompt and run the gacutil command.
- D. Open a command prompt and run the dtutil /copy command.
- E. Open a command prompt and run the dtexec /rep /conn command.
- F. Open a command prompt and run the dtexec /dumperror /conn command.
- G. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- H. Create a reusable custom logging component and use it in the SSIS project.
- I. Configure the SSIS solution to use the Project Deployment Model.
- J. Configure the output of a component in the package data flow to use a data tap.
- K. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL Server.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (same context, five questions, one of each answer: "data tap", "event handler", "dtexecui/SQLLogProvider", "gacutil" and "wizard")
--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms403356.aspx>

QUESTION 34

You are implementing a SQL Server Integration Services (SSIS) package that loads data hosted in a SQL Azure database into a data warehouse. The source system contains redundant or inconsistent data. When the package finds invalid data, the row containing the invalid data must be omitted but it must also be written to a text file for further analysis. You need to establish the best technique to log these invalid rows while keeping the amount of development effort to a minimum. What should you do?

- A. Add an OnError event handler to the SSIS project.
- B. Use an msi file to deploy the package on the server.
- C. Open a command prompt and run the gacutil command.
- D. Open a command prompt and run the dtutil /copy command.
- E. Open a command prompt and run the dtexec /rep /conn command.
- F. Open a command prompt and run the dtexec /dumperror /conn command.
- G. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- H. Create a reusable custom logging component and use it in the SSIS project.
- I. Configure the SSIS solution to use the Project Deployment Model.
- J. Configure the output of a component in the package data flow to use a data tap.
- K. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL

Server.

Correct Answer: J

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (same context, five questions, one of each answer: "data tap", "event handler", "dtexec/SQLLogProvider", "gacutil" and "wizard")
--\BURGOS

According to these references, this answer looks correct.

References:

<http://technet.microsoft.com/en-us/library/hh230989.aspx>
<http://www.rafael-salas.com/2012/01/ssis-2012-quick-peek-to-data-taps.html>
<http://msdn.microsoft.com/en-us/library/ms162820.aspx>
<http://msdn.microsoft.com/en-us/library/hh231187.aspx>
<http://technet.microsoft.com/en-us/library/ms140223.aspx>
<http://msdn.microsoft.com/en-us/library/jj655339.aspx>

QUESTION 35

You are troubleshooting an existing SQL Server Integration Services (SSIS) package. On several occasions, the package execution does not finish and no data seems to have been transferred. You need to ensure that package logging occurs. Your solution must minimize deployment and development efforts. What should you do?

- A. Add an OnError event handler to the SSIS project.
- B. Use an msi file to deploy the package on the server.
- C. Open a command prompt and run the gacutil command.
- D. Open a command prompt and run the dtutil /copy command.
- E. Open a command prompt and run the dtexec /rep /conn command.
- F. Open a command prompt and run the dtexec /dumperror /conn command.
- G. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- H. Create a reusable custom logging component and use it in the SSIS project.
- I. Configure the SSIS solution to use the Project Deployment Model.
- J. Configure the output of a component in the package data flow to use a data tap.
- K. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL Server.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (same context, five questions, one of each answer: "data tap", "event handler", "dtexec/SQLLogProvider", "gacutil" and "wizard")
Two similar questions with different suggested answers. I opted by "On Error Event Handler" for "You need to ensure that package logging occurs"
--\Burgos.

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms141212.aspx>
<http://www.mssqltips.com/sqlservertip/2450/ssis-package-deployment-model-in-sql-server-2012-part-1-of-2/>
<http://www.mssqltips.com/sqlservertip/2450/ssis-package-deployment-model-in-sql-server-2012-part-2-of-2/>

QUESTION 36

To facilitate the troubleshooting of SQL Server Integration Services (SSIS) packages, a logging methodology is put in place. The methodology has the following requirements:

- The deployment process must be simplified.
- All the logs must be centralized in SQL Server.
- Log data must be available via reports or T-SQL.
- Log archival must be automated.

You need to configure a logging methodology that meets the requirements while minimizing the amount of deployment and development effort. What should you do?

- A. Open a command prompt and run the gacutil command.
- B. Open a command prompt and execute the package by using the SQL Log provider and running the dtexecui.exe utility.
- C. Add an OnError event handler to the SSIS project.
- D. Use an msi file to deploy the package on the server.
- E. Configure the output of a component in the package data flow to use a data tap.
- F. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL Server.
- G. Open a command prompt and run the dtexec /rep /conn command.
- H. Open a command prompt and run the dtutil /copy command.
- I. Open a command prompt and run the dtexec /dumperror /conn command.
- J. Configure the SSIS solution to use the Project Deployment Model.
- K. Create a reusable custom logging component and use it in the SSIS project.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (same context, five questions, one of each answer: "data tap", "event handler", "dtexecui.exe/SQLLogProvider", "gacutil" and "wizard")
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms140246.aspx>

<http://msdn.microsoft.com/en-us/library/hh231187.aspx>

QUESTION 37

To ease the debugging of packages, you standardize the SQL Server Integration Services (SSIS) package logging methodology. The methodology has the following requirements:

- Centralized logging in SQL Server
- Simple deployment
- Availability of log information through reports or T-SQL
- Automatic purge of older log entries
- Configurable log details

You need to configure a logging methodology that meets the requirements while minimizing the amount of deployment and development effort. What should you do?

- A. Use the gacutil command.
- B. Use the dtutil /copy command.
- C. Use the Integration Services Deployment Wizard.

- D. Create an OnError event handler.
- E. Create a reusable custom logging component.
- F. Run the package by using the dtexec /rep /conn command.
- G. Run the package by using the dtexec /dumperror /conn command.
- H. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- I. Add a data tap on the output of a component in the package data flow.
- J. Deploy the package by using an msi file.
- K. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.

Correct Answer: H

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (same context, five questions, one of each answer: "data tap", "event handler", "dtexecui.exe/SQLLogProvider", "gacutil" and "wizard")
 --\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms140246.aspx>

<http://msdn.microsoft.com/en-us/library/hh231187.aspx>

QUESTION 38

You are deploying a new SQL Server Integration Services (SSIS) project to the test environment. A package in the project uses a custom task component. You need to ensure that the custom object is deployed on the test environment correctly. What should you do?

- A. Run the package by using the dtexec /rep /conn command.
- B. Create a reusable custom logging component.
- C. Create an OnError event handler.
- D. Use the gacutil command.
- E. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.
- F. Run the package by using the dtexec /dumperror /conn command.
- G. Use the Integration Services Deployment Wizard.
- H. Deploy the package by using an msi file.
- I. Add a data tap on the output of a component in the package data flow.
- J. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- K. Use the dtutil /copy command.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (same context, five questions, one of each answer: "data tap", "event handler", "dtexecui.exe/SQLLogProvider", "gacutil" and "wizard")
 --\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms403356.aspx>

QUESTION 39

You are writing a SQL Server Integration Services (SSIS) package that transfers data from a legacy system. Data integrity in the legacy system is very poor. Invalid rows are discarded by the package but must be logged to a CSV file for auditing purposes. You need to establish the best technique to log these invalid rows while minimizing the amount of development effort. What should you do?

- A. Add a data tap on the output of a component in the package data flow.
- B. Deploy the package by using an msi file.
- C. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- D. Uses the dtutil /copy command.
- E. Create an OnError event handler.
- F. Use the Integration Services Deployment Wizard.
- G. Use the gacutil command.
- H. Create a reusable custom logging component.
- I. Run the package by using the dtexec /rep /conn command.
- J. Run the package by using the dtexec /dumperror /conn command.
- K. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (same context, five questions, one of each answer: "data tap", "event handler", "dtuiexec/SQLLogProvider", "gacutil" and "wizard")
--\BURGOS

According to these references, this answer looks correct.

Reference:

<http://www.rafael-salas.com/2012/01/ssis-2012-quick-peek-to-data-taps.html>

<http://msdn.microsoft.com/en-us/library/hh230989.aspx>

<http://msdn.microsoft.com/en-us/library/jj655339.aspx>

QUESTION 40

You are deploying a project to the SQL Server Integration Services (SSIS) catalog. You need to ensure that the project is deployed to the production environment with the least amount of administrative effort. What should you do?

- A. Create an On Error event handler.
- B. Add a data tap on the output of a component in the package data flow.
- C. Use the gacutil command.
- D. Create a reusable custom logging component.
- E. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- F. Run the package by using the dtexec /rep /conn command.
- G. Use the Integration Services Deployment Wizard.
- H. Run the package by using the dtexec /dumperror /conn command.
- I. Use the dtutil /copy command.
- J. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.
- K. Deploy the package by using an msi file.

Correct Answer: G

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (same context, five questions, one of each answer: "data tap", "event handler", "dtexec/SQLLogProvider", "gacutil" and "wizard")

--\BURGOS

According to these references, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/aa479588.aspx>

<http://msdn.microsoft.com/en-us/library/aa213290.aspx>

<http://msdn.microsoft.com/en-us/library/aa213373.aspx>

QUESTION 41

Occasionally a job that executes an existing SQL Server Integration Services (SSIS) package does not complete and nothing is processed. You need to ensure that package logging occurs. Your solution must minimize deployment and development efforts. What should you do?

- A. Create a reusable custom logging component.
- B. Use the gacutil command.
- C. Use the Integration Services Deployment Wizard.
- D. Run the package by using the dtexec /rep /conn command.
- E. Add a data tap on the output of a component in the package data flow.
- F. Create an OnError event handler.
- G. Use the dtutil /copy command.
- H. Deploy the package by using an msi file.
- I. Run the package by using the dtexec /dumperror /conn command.
- J. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- K. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (same context, five questions, one of each answer: "data tap", "event handler", "dtexec/SQLLogProvider", "gacutil" and "wizard")

Two similar questions with different suggested answers. I opted by "On Error Event Handler" for "You need to ensure that package logging occurs"

--\Burgos.

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms140246.aspx>

<http://msdn.microsoft.com/en-us/library/aa231187.aspx>

QUESTION 42

You are adding a new capability to several dozen SQL Server Integration Services (SSIS) packages. The new capability is not available as an SSIS task. Each package must be extended with the same new capability. You need to add the new capability to all the packages without copying the code between packages. What should you do?

- A. Use the Expression task.
- B. Use the Script component.
- C. Use the Script task.
- D. Develop a custom task.
- E. Develop a custom component.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms135965.aspx>

<http://msdn.microsoft.com/en-us/library/ms345161.aspx>

QUESTION 43

You are the data steward for a Business Intelligence project. You must identify duplicate rows stored in a SQL Server table and output discoveries to a CSV file. A Data Quality Services (DQS) knowledge base has been created to support this project. You need to produce the CSV file with the least amount of development effort. What should you do?

- A. Create an Integration Services package and use a Data Profiling transform.
- B. Create a custom .NET application based on the Knowledgebase class.
- C. Create a CLR stored procedure based on the Knowledgebase class.
- D. Create a Master Data Services (MDS) business rule.
- E. Create a data quality project.

Correct Answer: E

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/hh213052.aspx>

<http://msdn.microsoft.com/en-us/library/ff877917.aspx>

<http://msdn.microsoft.com/en-us/library/microsoft.masterdataservices.services.datacontracts.knowledgebase.aspx>

<http://msdn.microsoft.com/en-us/library/bb895263.aspx>

QUESTION 44

You administer a Microsoft SQL Server database named Sales. The database is 3 terabytes in size. The Sales database is configured as shown in the following table:

Filegroup	File
PRIMARY	<ul style="list-style-type: none"> Sales.mdf
SALES	<ul style="list-style-type: none"> Sales_1.ndf Sales_2.ndf Sales_3.ndf
ARCHIVES	<ul style="list-style-type: none"> SalesArch_1.ndf SalesArch_2.ndf

You discover that Sales_2.ndf is corrupt. You need to recover the corrupted data in the minimum amount of time. What should you do?

- A. Perform a file restore.
- B. Perform a transaction log restore.
- C. Perform a restore from a full backup.
- D. Perform a filegroup restore.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (different order of choices)

I put a reference ms186858 below. Is possible to restore FILE or FILEGROUP in SQL 2012 and this choice sounds better to me. In original VCE file appears another equal question that suggest to use a FULL restore, I delete this question.

--\Burgos.

According to these references, this answer looks correct.

Reference: <http://technet.microsoft.com/en-us/library/ms187048.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/aa337540.aspx>

Reference: http://msdn.microsoft.com/en-us/library/ms186858.aspx#restoring_using_FILE_n_FG

QUESTION 45

You administer a Microsoft SQL Server 2012 database. You create an availability group named haContosoDbs. Your primary replica is available at Server01\Contoso01. You need to configure the availability group to have the highest availability. You also need to ensure that no data is lost. Which Transact-SQL statement should you use?

- A. ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = ASYNCHRONOUS_COMMIT, FAILOVER_MODE = AUTOMATIC)
- B. ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = SYNCHRONOUS_COMMIT, FAILOVER_MODE = MANUAL)
- C. ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = SYNCHRONOUS_COMMIT, FAILOVER_MODE = AUTOMATIC)
- D. ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = ASYNCHRONOUS_COMMIT, FAILOVER_MODE = MANUAL)

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO (I had another one, with "low latency" (ASYNCHRONOUS))
--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ff877931.aspx>

QUESTION 46

You administer a Microsoft SQL Server 2012 database. You create an availability group named haContosoDbs. Your primary replica is available at Server01\Contoso01. You need to configure the availability group to minimize transaction latency on any available secondary databases. In the event of a database failure, the designated secondary database should come online automatically. Which Transact-SQL statement should you use?

- A. ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = ASYNCHRONOUS_COMMIT, FAILOVER_MODE = AUTOMATIC)
- B. ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = SYNCHRONOUS_COMMIT, FAILOVER_MODE = MANUAL)
- C. ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = SYNCHRONOUS_COMMIT, FAILOVER_MODE = AUTOMATIC)
- D. ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = ASYNCHRONOUS_COMMIT, FAILOVER_MODE = MANUAL)

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES
--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ff877931.aspx>

QUESTION 47

You administer two Microsoft SQL Server 2012 servers. Each server resides in a different, untrusted domain. You plan to configure database mirroring. You need to be able to create database mirroring endpoints on both servers. What should you do?

- A. Configure the SQL Server service account to use Network Service.
- B. Use a server certificate.
- C. Use a database certificate.
- D. Configure the SQL Server service account to use Local System.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (same differences)
--\BURGOS

According to these references, this answer looks correct.

References:

<http://technet.microsoft.com/en-us/library/ms179511.aspx>

<http://technet.microsoft.com/en-us/library/ms186384.aspx>

QUESTION 48

You administer a Microsoft SQL Server 2012 database. You configure Transparent Data Encryption (TDE) on the Orders database by using the following statements:

```
CREATE MASTER KEY ENCRYPTION BY PASSWORD = 'MyPassword1!'
CREATE CERTIFICATE TDE_Certificate WITH SUBJECT = 'TDE Certificate';

BACKUP CERTIFICATE TDE_Certificate TO FILE = 'd:\TDE_Certificate.cer'
WITH PRIVATE KEY (FILE = 'd:\TDE_Certificate.key', ENCRYPTION BY PASSWORD =
'MyPassword1!');

CREATE DATABASE ENCRYPTION KEY
WITH ALGORITHM = AES_256
ENCRYPTION BY SERVER CERTIFICATE TDE_Certificate;

ALTER DATABASE Orders SET ENCRYPTION ON;
```

You attempt to restore the Orders database and the restore fails. You copy the encryption file to the original location. A hardware failure occurs and so a new server must be installed and configured. After installing SQL Server to the new server, you restore the Orders database and copy the encryption files to their original location. However, you are unable to access the database. You need to be able to restore the database. Which Transact-SQL statement should you use before attempting the restore?

- A. ALTER DATABASE Master SET ENCRYPTION OFF;
- B. CREATE CERTIFICATE TDE_Certificate FROM FILE = 'd:\TDE_Certificate.cer' WITH PRIVATE KEY (FILE = 'D:\TDE_Certificate.key', DECRYPTION BY PASSWORD = 'MyPassword1!');
- C. CREATE CERTIFICATE TDE_Certificate WITH SUBJECT = 'TDE Certificate';
USE Orders;
CREATE DATABASE ENCRYPTION KEY
WITH ALGORITHM = AES_256
ENCRYPTION BY SERVER CERTIFICATE TDE_Certificate;
- D. CREATE CERTIFICATE TDE_Certificate FROM FILE = 'd:\TDE_Certificate.cer';
- E. DECLARE @startdate date
SET @startdate = GETDATE()

CREATE CERTIFICATE TDE_Certificate FROM FILE = 'd:\TDE_Certificate.cer'
WITH START_DATE = @startdate;
- F. CREATE ASSEMBLY TDE_Assembly FROM 'd:\TDE_Certificate.cer'
WITH PERMISSION_SET = SAFE;
GO
CREATE CERTIFICATE TDE_Certificate FROM ASSEMBLY TDE_Assembly;

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

--Burgos. YES (with four choices)

This question appears two times in original VCE file with different set of choices, but two choices appears in both questions (the correct and the shorter "CREATE CERTIFICATE" choice). I discharge one of questions and put all different choices in this questions.

--\Burgos.

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ff773063.aspx>

<http://msdn.microsoft.com/en-us/library/bb510663.aspx>

<http://msdn.microsoft.com/en-us/library/bb934049.aspx>

QUESTION 49

You administer a Microsoft SQL Server 2012 database that has Trustworthy set to On. You create a stored procedure that returns database-level information from Dynamic Management Views. You grant User1 access to execute the stored procedure. You need to ensure that the stored procedure returns the required information when User1 executes the stored procedure. You need to achieve this goal by granting the minimum permissions required. What should you do? (Each correct answer presents a complete solution. Choose all that apply.)

- A. Create a SQL Server login that has VIEW SERVER STATE permissions.
Create an application role and a secured password for the role.
- B. Modify the stored procedure to include the EXECUTE AS OWNER statement.
Grant VIEW SERVER STATE permissions to the owner of the stored procedure.
- C. Create a SQL Server login that has VIEW SERVER STATE permissions.
Modify the stored procedure to include the EXECUTE AS {newlogin} statement.
- D. Grant the db_owner role on the database to User1.
- E. Grant the sysadmin role on the database to User1.

Correct Answer: BC

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms187861.aspx>

<http://msdn.microsoft.com/en-us/library/ms191291.aspx>

QUESTION 50

You use a contained database named ContosoDb within a domain. You need to create a user who can log on to the ContosoDb database. You also need to ensure that you can port the database to different database servers within the domain without additional user account configurations. Which type of user should you create?

- A. SQL user with login
- B. SQL user without login
- C. Domain user
- D. User mapped to an asymmetric key

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to this reference, this answer looks correct.

Reference: <http://msdn.microsoft.com/en-us/library/ff929071.aspx>

QUESTION 51

You are the lead database administrator (DBA) of a Microsoft SQL Server 2012 environment. All DBAs are members of the DOMAIN\JrDBAs Active Directory group. You grant DOMAIN\JrDBAs access to the SQL Server. You need to create a server role named SpecialDBARole that can perform the following functions:

- View all databases.
- View the server state.
- Assign GRANT, DENY, and REVOKE permissions on logins.

You need to add DOMAIN\JrDBAs to the server role. You also need to provide the least level of privileges necessary. Which SQL statement or statements should you use? Choose all that apply.

- A. `CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION setupadmin;`
- B. `ALTER SERVER ROLE [SpecialDBARole] ADD MEMBER [DOMAIN\JrDBAs];`
- C. `CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION securityadmin;`
- D. `GRANT VIEW DEFINITION TO [SpecialDBARole];`
- E. `CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION serveradmin;`
- F. `GRANT VIEW SERVER STATE, VIEW ANY DATABASE TO [SpecialDBARole];`

Correct Answer: BCF

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

Verified answer as correct.

References:

<http://msdn.microsoft.com/en-us/library/ee677634.aspx>

<http://msdn.microsoft.com/en-us/library/ms188659.aspx>

<http://msdn.microsoft.com/en-us/library/ms191291.aspx>

QUESTION 52

You administer a Microsoft SQL Server 2012 server. You plan to deploy new features to an application. You need to evaluate existing and potential clustered and non-clustered indexes that will improve performance. What should you do?

- A. Query the sys.dm_db_index_usage_stats DMV.
- B. Query the sys.dm_db_missing_index_details DMV.
- C. Use the Database Engine Tuning Advisor.
- D. Query the sys.dm_db_missing_index_columns DMV.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms174202.aspx>

QUESTION 53

You administer a SQL Server 2012 server that contains a database named SalesDb. SalesDb contains a schema named Customers that has a table named Regions. A user named UserA is a member of a role named Sales. UserA is granted the Select permission on the Regions table. The Sales role is granted the Select permission on the Customers schema. You need to ensure that UserA is disallowed to select from the Regions table. Which Transact-SQL statement should you use?

- A. DENY SELECT ON Object::Regions FROM Sales
- B. DENY SELECT ON Schema::Customers FROM Sales
- C. REVOKE SELECT ON Object::Regions FROM Sales
- D. REVOKE SELECT ON Schema::Customers FROM Sales
- E. DENY SELECT ON Object::Regions FROM UserA
- F. DENY SELECT ON Schema::Customers FROM UserA
- G. REVOKE SELECT ON Object::Regions FROM UserA
- H. REVOKE SELECT ON Schema::Customers FROM UserA
- I. EXEC sp_addrolemember 'Sales', 'UserA'
- J. EXEC sp_droprolemember 'Sales', 'UserA'

Correct Answer: E

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (four questions with same context, two REVOKE and two DENY, I did not remember exactly)
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188369.aspx>

<http://msdn.microsoft.com/en-us/library/ms187750.aspx>

<http://msdn.microsoft.com/en-us/library/ff848791.aspx>

QUESTION 54

You administer a SQL Server 2012 server that contains a database named SalesDb. SalesDb contains a schema named Customers that has a table named Regions. A user named UserA is a member of a role named Sales. UserA is granted the Select permission on the Regions table. The Sales role is granted the Select permission on the Customers schema. You need to ensure that the following requirements are met:

- The Sales role does not have the Select permission on the Customers schema.
- UserA has the Select permission on the Regions table.

Which Transact-SQL statement should you use?

- A. DENY SELECT ON Object::Regions FROM Sales
- B. DENY SELECT ON Schema::Customers FROM Sales
- C. REVOKE SELECT ON Object::Regions FROM Sales
- D. REVOKE SELECT ON Schema::Customers FROM Sales
- E. DENY SELECT ON Object::Regions FROM UserA
- F. DENY SELECT ON Schema::Customers FROM UserA
- G. REVOKE SELECT ON Object::Regions FROM UserA
- H. REVOKE SELECT ON Schema::Customers FOR UserA
- I. EXEC sp_addrolemember 'Sales', 'UserA'

J. EXEC sp_droprolemember 'Sales', 'UserA'

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (four questions with same context, two REVOKE and two DENY, I did not remember exactly)
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188369.aspx>

<http://msdn.microsoft.com/en-us/library/ms187750.aspx>

<http://msdn.microsoft.com/en-us/library/ff848791.aspx>

QUESTION 55

You administer a SQL Server 2012 server that contains a database named SalesDb. SalesDb contains a schema named Customers that has a table named Regions. A user named UserA is a member of a role named Sales. UserA is granted the Select permission on the Regions table. The Sales role is granted the Select permission on the Customers schema. You need to remove the Select permission for UserA on the Regions table. You also need to ensure that UserA can still access all the tables in the Customers schema, including the Regions table, through the Sales role permissions. Which Transact-SQL statement should you use?

- A. DENY SELECT ON Object::Regions FROM UserA
- B. DENY SELECT ON Schema::Customers FROM UserA
- C. EXEC sp_addrolemember 'Sales', 'UserA'
- D. REVOKE SELECT ON Object::Regions FROM UserA
- E. REVOKE SELECT ON Object::Regions FROM Sales
- F. EXEC sp_droprolemember 'Sales', 'UserA'
- G. REVOKE SELECT ON Schema::Customers FROM UserA
- H. DENY SELECT ON Object::Regions FROM Sales
- I. DENY SELECT ON Schema::Customers FROM Sales
- J. REVOKE SELECT ON Schema::Customers FROM Sales

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (four questions with same context, two REVOKE and two DENY, I did not remember exactly)
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188369.aspx>

<http://msdn.microsoft.com/en-us/library/ms187750.aspx>

<http://msdn.microsoft.com/en-us/library/ff848791.aspx>

QUESTION 56

You administer a SQL Server 2012 server that contains a database named SalesDb. SalesDb contains a schema named Customers that has a table named Regions. A user named UserA is a member of a role named Sales. UserA is granted the Select permission on the Regions table and the Sales role is granted the Select

permission on the Customers schema. You need to ensure that the Sales role, including UserA, is disallowed to select from the Regions table. Which Transact-SQL statement should you use?

- A. REVOKE SELECT ON Schema::Customers FROM UserA
- B. REVOKE SELECT ON Object::Regions FROM UserA
- C. EXEC sp_addrolemember 'Sales', 'UserA'
- D. DENY SELECT ON Schema::Customers FROM Sales
- E. EXEC sp_droprolemember 'Sales', 'UserA'
- F. REVOKE SELECT ON Schema::Customers FROM Sales
- G. DENY SELECT ON Object::Regions FROM UserA
- H. REVOKE SELECT ON Object::Regions FROM Sales
- I. DENY SELECT ON Schema::Customers FROM UserA
- J. DENY SELECT ON Object::Regions FROM Sales

Correct Answer: J

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (four questions with same context, two REVOKE and two DENY, I did not remember exactly)
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188369.aspx>

<http://msdn.microsoft.com/en-us/library/ms187750.aspx>

<http://msdn.microsoft.com/en-us/library/ff848791.aspx>

QUESTION 57

You administer a SQL Server 2012 server that contains a database named SalesDb. SalesDb contains a schema named Customers that has a table named Regions. A user named UserA is a member of a role named Sales. UserA is granted the Select permission on the Regions table. The Sales role is granted the Select permission on the Customers schema. You need to ensure that UserA is disallowed to select from any of the tables in the Customers schema. Which Transact-SQL statement should you use?

- A. DENY SELECT ON Object::Regions FROM UserA
- B. DENY SELECT ON Object::Regions FROM Sales
- C. REVOKE SELECT ON Schema::Customers FROM Sales
- D. REVOKE SELECT ON Schema::Customers FROM UserA
- E. REVOKE SELECT ON Object::Regions FROM Sales
- F. REVOKE SELECT ON Object::Regions FROM UserA
- G. DENY SELECT ON Schema::Customers FROM Sales
- H. DENY SELECT ON Schema::Customers FROM UserA
- I. EXEC sp_addrolemember 'Sales', 'UserA'
- J. EXEC sp_droprolemember 'Sales', 'UserA'

Correct Answer: H

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (four questions with same context, two REVOKE and two DENY, I did not remember

exactly)
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188369.aspx>

<http://msdn.microsoft.com/en-us/library/ms187750.aspx>

<http://msdn.microsoft.com/en-us/library/ff848791.aspx>

QUESTION 58

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN). The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours. Five percent of the existing data is modified each day. The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands. Each data load adds 3 GB of data to the database. These data load operations must occur in the minimum amount of time. A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours. You need to ensure that your backup will continue if any invalid checksum is encountered. Which backup option should you use?

- A. STANDBY
- B. Differential
- C. FULL
- D. CHECKSUM
- E. BULK_LOGGED
- F. CONTINUE_AFTER_ERROR
- G. SIMPLE
- H. DBO_ONLY
- I. COPY_ONLY
- J. SKIP
- K. RESTART
- L. Transaction log
- M. NO_CHECKSUM
- N. NORECOVERY

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (ALL five questions with same context)

--\BURGOS

Verified answer as correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms186865.aspx>

<http://msdn.microsoft.com/en-us/library/microsoft.sqlserver.management.smo.backuprestorebase.continueaftererror.aspx>

[microsoft.sqlserver.management.smo.backuprestorebase.continueaftererror.aspx](http://msdn.microsoft.com/en-us/library/microsoft.sqlserver.management.smo.backuprestorebase.continueaftererror.aspx)

QUESTION 59

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN). The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours. Five percent of the existing data is modified each day. The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands. Each data load adds 3 GB of data to the database. These data load operations must occur in the minimum amount of time. A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours. On Wednesday at 10:00 hours, the development team requests you to refresh the database on a development server by using the most recent version. You need to perform a full database backup that will be restored on the development server. Which backup option should you use?

- A. NORECOVERY
- B. FULL
- C. NO_CHECKSUM
- D. CHECKSUM
- E. Differential
- F. BULK_LOGGED
- G. STANDBY
- H. RESTART
- I. SKIP
- J. Transaction log
- K. DBO ONLY
- L. COPY_ONLY
- M. SIMPLE
- N. CONTINUE AFTER ERROR

Correct Answer: L

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (ALL five questions with same context)

--\BURGOS

Verified answer as correct.

References:

<http://msdn.microsoft.com/en-us/library/ms191495.aspx>

<http://msdn.microsoft.com/en-us/library/ms186858.aspx>

QUESTION 60

You are migrating a database named Orders to a new server that runs Microsoft SQL Server 2012. You attempt to add the [Corpnet\User1] login to the database. However, you receive the following error message: "User already exists in current database."

You need to configure the [Corpnet\User1] login to be able to access the Orders database and retain the original permissions. You need to achieve this goal by using the minimum required permissions. Which Transact-SQL statement should you use?

- A. DROP USER [User1];
CREATE USER [Corpnet\User1] FOR LOGIN [Corpnet\User1];
ALTER ROLE [db_owner] ADD MEMBER [Corpnet\User1];
- B. ALTER SERVER ROLE [sysadmin] ADD MEMBER [Corpnet\User1];
- C. ALTER USER [Corpnet\User1] WITH LOGIN [Corpnet\User1];
- D. ALTER ROLE [db_owner] ADD MEMBER [Corpnet\User1];

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (version "User1")
--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms176060.aspx>

QUESTION 61

You create an availability group that has replicas named HA/Server01 and HA/Server02. Currently, HA/Server01 is the primary replica. You have multiple queries that read data and produce reports from the database. You need to offload the reporting workload to the secondary replica when HA/Server01 is the primary replica. What should you do?

- A. Set the Availability Mode property of HA/Server02 to Asynchronous commit.
- B. Set the Readable Secondary property of HA/Server02 to Read-intent only.
- C. Set the Connections in Primary Role property of HA/Server01 to Allow read/write connections.
- D. Set the Availability Mode property of HA/Server01 to Asynchronous commit.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (version "User1")
--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/jj542414.aspx>

QUESTION 62

You administer a Microsoft SQL Server 2012 server that hosts a transactional database and a reporting database. The transactional database is updated through a web application and is operational throughout the day. The reporting database is only updated from the transactional database. The recovery model and backup schedule are configured as shown in the following table:

Database	Description
Transactional database	Recovery model: <ul style="list-style-type: none"> Full Backup schedule: <ul style="list-style-type: none"> Full database backup: midnight, daily Differential database backup: on the hour, every two hours starting at 02:00 hours except at 00:00 hours Log backup: every half hour, except at the times of full and differential backups
Reporting database	Recovery model: <ul style="list-style-type: none"> Simple Backup schedule: <ul style="list-style-type: none"> Full database backup: 01:00 hours daily Differential database backup: 13:00 hours daily Data updates: <ul style="list-style-type: none"> Changes in data are updated from the transactional database to the reporting database at 00:30 hours and at 12:30 hours The update takes 15 minutes

The differential backup of the reporting database fails. Then, the reporting database fails at 14:00 hours. You need to ensure that the reporting database is restored. You also need to ensure that data loss is minimal. What should you do?

- A. Restore the latest full backup, and restore the latest differential backup. Then, restore the latest log backup.
- B. Perform a point-in-time restore.
- C. Restore the latest full backup.
- D. Restore the latest full backup, and restore the latest differential backup. Then, restore each log backup taken before the time of failure from the most recent differential backup.
- E. Restore the latest full backup. Then, restore the latest differential backup.
- F. Restore the latest full backup. Then, restore each differential backup taken before the time of failure from the most recent full backup.
- G. Perform a page restore.
- H. Perform a partial restore.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

--BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms187048.aspx>
<http://msdn.microsoft.com/en-us/library/ms186289.aspx>
<http://msdn.microsoft.com/en-us/library/ms175477.aspx>
<http://msdn.microsoft.com/en-us/library/ms189860.aspx>
<http://msdn.microsoft.com/en-us/library/ms179314.aspx>
<http://msdn.microsoft.com/en-us/library/ms175526.aspx>
<http://msdn.microsoft.com/en-us/library/ms191539.aspx>
<http://msdn.microsoft.com/en-us/library/ms191429.aspx>
<http://msdn.microsoft.com/en-us/library/ms191253.aspx>

QUESTION 63

You administer a Microsoft SQL Server 2012 server that hosts a transactional database and a reporting database. The transactional database is updated through a web application and is operational throughout the day. The reporting database is only updated from the transactional database. The recovery model and backup schedule are configured as shown in the following table:

Database	Description
Transactional database	<p>Recovery model:</p> <ul style="list-style-type: none">• Full <p>Backup schedule:</p> <ul style="list-style-type: none">• Full database backup: midnight, daily• Differential database backup: on the hour, every two hours starting at 02:00 hours except at 00:00 hours• Log backup: every half hour, except at the times of full and differential backups
Reporting database	<p>Recovery model:</p> <ul style="list-style-type: none">• Simple <p>Backup schedule:</p> <ul style="list-style-type: none">• Full database backup: 01:00 hours daily• Differential database backup: 13:00 hours daily <p>Data updates:</p> <ul style="list-style-type: none">• Changes in data are updated from the transactional database to the reporting database at 00:30 hours and at 12:30 hours• The update takes 15 minutes

One of the hard disk drives that stores the reporting database fails at 16:40 hours. You need to ensure that the reporting database is restored. You also need to ensure that data loss is minimal. What should you do?

- A. Restore the latest full backup. Then, restore each differential backup taken before the time of failure from the most recent full backup.
- B. Perform a partial restore.

- C. Restore the latest full backup, and restore the latest differential backup. Then, restore the latest log backup.
- D. Restore the latest full backup.
- E. Perform a page restore.
- F. Restore the latest full backup, and restore the latest differential backup. Then, restore each log backup taken before the time of failure from the most recent differential backup.
- G. Restore the latest full backup. Then, restore the latest differential backup.
- H. Perform a point-in-time restore.

Correct Answer: G

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms187048.aspx>

<http://msdn.microsoft.com/en-us/library/ms186289.aspx>

<http://msdn.microsoft.com/en-us/library/ms175477.aspx>

<http://msdn.microsoft.com/en-us/library/ms189860.aspx>

<http://msdn.microsoft.com/en-us/library/ms179314.aspx>

<http://msdn.microsoft.com/en-us/library/ms175526.aspx>

<http://msdn.microsoft.com/en-us/library/ms191539.aspx>

<http://msdn.microsoft.com/en-us/library/ms191429.aspx>

<http://msdn.microsoft.com/en-us/library/ms191253.aspx>

QUESTION 64

You administer a Microsoft SQL Server 2012 server that hosts a transactional database and a reporting database. The transactional database is updated through a web application and is operational throughout the day. The reporting database is only updated from the transactional database. The recovery model and backup schedule are configured as shown in the following table:

Database	Description
Transactional database	Recovery model: <ul style="list-style-type: none"> • Full Backup schedule: <ul style="list-style-type: none"> • Full database backup: midnight, daily • Differential database backup: on the hour, every two hours starting at 02:00 hours except at 00:00 hours • Log backup: every half hour, except at the times of full and differential backups
Reporting database	Recovery model: <ul style="list-style-type: none"> • Simple Backup schedule: <ul style="list-style-type: none"> • Full database backup: 01:00 hours daily • Differential database backup: 13:00 hours daily Data updates: <ul style="list-style-type: none"> • Changes in data are updated from the transactional database to the reporting database at 00:30 hours and at 12:30 hours • The update takes 15 minutes

At 16:20 hours, you discover that pages 17, 137, and 205 on one of the database files are corrupted on the transactional database. You need to ensure that the transactional database is restored. You also need to ensure that data loss is minimal. What should you do?

- A. Perform a partial restore.
- B. Restore the latest full backup, and restore the latest differential backup. Then, restore each log backup taken before the time of failure from the most recent differential backup.
- C. Perform a point-in-time restore.
- D. Restore the latest full backup.
- E. Restore the latest full backup, and restore the latest differential backup. Then, restore the latest log backup.
- F. Perform a page restore.
- G. Restore the latest full backup. Then, restore each differential backup taken before the time of failure from the most recent full backup.
- H. Restore the latest full backup. Then, restore the latest differential backup.

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms187048.aspx>
<http://msdn.microsoft.com/en-us/library/ms186289.aspx>
<http://msdn.microsoft.com/en-us/library/ms175477.aspx>
<http://msdn.microsoft.com/en-us/library/ms189860.aspx>
<http://msdn.microsoft.com/en-us/library/ms179314.aspx>
<http://msdn.microsoft.com/en-us/library/ms175526.aspx>
<http://msdn.microsoft.com/en-us/library/ms191539.aspx>
<http://msdn.microsoft.com/en-us/library/ms191429.aspx>
<http://msdn.microsoft.com/en-us/library/ms191253.aspx>

QUESTION 65

You administer a Microsoft SQL Server 2012 server that hosts a transactional database and a reporting database. The transactional database is updated through a web application and is operational throughout the day. The reporting database is only updated from the transactional database. The recovery model and backup schedule are configured as shown in the following table:

Database	Description
Transactional database	<p>Recovery model:</p> <ul style="list-style-type: none">• Full <p>Backup schedule:</p> <ul style="list-style-type: none">• Full database backup: midnight, daily• Differential database backup: on the hour, every two hours starting at 02:00 hours except at 00:00 hours• Log backup: every half hour, except at the times of full and differential backups
Reporting database	<p>Recovery model:</p> <ul style="list-style-type: none">• Simple <p>Backup schedule:</p> <ul style="list-style-type: none">• Full database backup: 01:00 hours daily• Differential database backup: 13:00 hours daily <p>Data updates:</p> <ul style="list-style-type: none">• Changes in data are updated from the transactional database to the reporting database at 00:30 hours and at 12:30 hours• The update takes 15 minutes

At 14:00 hours, you discover that pages 71, 520, and 713 on one of the database files are corrupted on the reporting database. You need to ensure that the databases are restored. You also need to ensure that data loss is minimal. What should you do?

A. Perform a partial restore.

- B. Restore the latest full backup, and restore the latest differential backup. Then, restore each log backup taken before the time of failure from the most recent differential backup.
- C. Restore the latest full backup.
- D. Restore the latest full backup, and restore the latest differential backup. Then, restore the latest log backup.
- E. Perform a page restore.
- F. Restore the latest full backup. Then, restore each differential backup taken before the time of failure from the most recent full backup.
- G. Perform a point-in-time restore.
- H. Restore the latest full backup. Then, restore the latest differential backup.

Correct Answer: H

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms187048.aspx>

<http://msdn.microsoft.com/en-us/library/ms186289.aspx>

<http://msdn.microsoft.com/en-us/library/ms175477.aspx>

<http://msdn.microsoft.com/en-us/library/ms189860.aspx>

<http://msdn.microsoft.com/en-us/library/ms179314.aspx>

<http://msdn.microsoft.com/en-us/library/ms175526.aspx>

<http://msdn.microsoft.com/en-us/library/ms191539.aspx>

<http://msdn.microsoft.com/en-us/library/ms191429.aspx>

<http://msdn.microsoft.com/en-us/library/ms191253.aspx>

QUESTION 66

You administer a Microsoft SQL Server 2012 database named ContosoDb. The database contains a table named Suppliers and a column named IsActive in the Purchases schema. You create a new user named ContosoUser in ContosoDb. ContosoUser has no permissions to the Suppliers table. You need to ensure that ContosoUser can delete rows that are not active from Suppliers. You also need to grant ContosoUser only the minimum required permissions. Which Transact-SQL statement should you use?

- A. `GRANT DELETE ON Purchases.Suppliers TO Contoso User`
- B. `CREATE PROCEDURE Purchases.PurgeInactiveSuppliers
WITH EXECUTE AS USER = 'dbo'
AS
DELETE FROM Purchases.Suppliers WHERE IsActive = 0
GO
GRANT EXECUTE ON Purchases.PurgeInactiveSuppliers TO ContosoUser`
- C. `GRANT SELECT ON Purchases.Suppliers TO ContosoUser`
- D. `CREATE PROCEDURE Purchases.PurgeInactiveSuppliers
AS
DELETE FROM Purchases.Suppliers WHERE IsActive = 0
GO
GRANT EXECUTE ON Purchases.PurgeInactiveSuppliers TO ContosoUser`

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188354.aspx>

<http://msdn.microsoft.com/en-us/library/ms187926.aspx>

QUESTION 67

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN). The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours. Five percent of the existing data is modified each day. The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands. Each data load adds 3 GB of data to the database. These data load operations must occur in the minimum amount of time. A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours. You implement log shipping of the financial database to another SQL Server 2012 instance. You decide to failover to this secondary database. You need to ensure that all transactions will be replicated to the secondary database. Which backup option should you use?

- A. Differential
- B. Transaction Log
- C. FULL
- D. SIMPLE
- E. SKIP
- F. RESTART
- G. STANDBY
- H. CHECKSUM
- I. DBO_ONLY
- J. COPY_ONLY
- K. NORECOVERY
- L. NO_CHECKSUM
- M. CONTINUE_AFTER_ERROR
- N. BULK_LOGGED

Correct Answer: K

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (ALL five questions with same context)

--\BURGOS

According to these references, this answer looks correct.

References:

<http://technet.microsoft.com/en-us/library/ms187103.aspx>

<http://msdn.microsoft.com/en-us/library/ms191233.aspx>

<http://msdn.microsoft.com/en-us/library/ms178117.aspx>

QUESTION 68

You are migrating a database named Orders to a new server that runs Microsoft SQL Server 2012. You attempt to add a SQL login, [User1], to the database. However, you receive the following error message: "User already exists in current database."

You need to configure the [User1] login to be able to access the Orders database and retain the original permissions. You need to achieve this goal by using the minimum required permissions. Which Transact-SQL statement should you use?

- A. `DROP USER [User1];`
`CREATE USER [User1] FOR LOGIN [User1];`
`ALTER ROLE [db_owner] ADD MEMBER [User1];`
- B. `ALTER ROLE [db_owner] ADD MEMBER [User1];`
- C. `ALTER USER [User1] WITH LOGIN [User1];`
- D. `ALTER SERVER ROLE [sysadmin] ADD MEMBER [User1];`

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

```
-- BURGOS - YES (version "User1")
--\BURGOS
```

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms176060.aspx>

QUESTION 69

You are creating a SQL Server Master Data Services (MDS) model. This model is used to store a master list of products. An attribute must be added to the Product entity to define the sales manager responsible for each product. You need to create an attribute in the Product entity that prevents users from entering invalid sales manager values. Which type of attribute should you create?

- A. Derived
- B. Domain-based
- C. User-defined
- D. Parent
- E. Explicit
- F. Recursive

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

```
-- BURGOS - NO
--\BURGOS
```

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/bb190163.aspx>

<http://msdn.microsoft.com/en-us/library/ee633737.aspx>

<http://msdn.microsoft.com/en-us/library/ee633759.aspx>

<http://msdn.microsoft.com/en-us/library/ee633745.aspx>
<http://msdn.microsoft.com/en-us/library/ee633724.aspx>
<http://msdn.microsoft.com/en-us/library/ee633733.aspx>
<http://msdn.microsoft.com/en-us/library/ff487058.aspx>

QUESTION 70

You administer a SQL Server 2012 server that contains a database named SalesDb. SalesDb contains a schema named Customers that has a table named Regions. A user named UserA is a member of a role named Sales. UserA is granted the Select permission on the Regions table. The Sales role is granted the Select permission on the Customers schema. You need to ensure that the Sales role, including UserA, is disallowed to select from any of the tables in the Customers schema. Which Transact-SQL statement should you use?

- A. REVOKE SELECT ON Schema::Customers FROM UserA
- B. DENY SELECT ON Object::Regions FROM UserA
- C. EXEC sp_addrolemember 'Sales', 'UserA'
- D. DENY SELECT ON Object::Regions FROM Sales
- E. REVOKE SELECT ON Object::Regions FROM UserA
- F. DENY SELECT ON Schema::Customers FROM Sales
- G. DENY SELECT ON Schema::Customers FROM UserA
- H. EXEC sp_droprolemember 'Sales', 'UserA'
- I. REVOKE SELECT ON Object::Regions FROM Sales
- J. REVOKE SELECT ON Schema::Customers FROM Sales

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (four questions with same context, two REVOKE and two DENY, I did not remember exactly)
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188369.aspx>
<http://msdn.microsoft.com/en-us/library/ms187750.aspx>
<http://msdn.microsoft.com/en-us/library/ff848791.aspx>

QUESTION 71

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN). The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours. Five percent of the existing data is modified each day. The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands. Each data load adds 3 GB of data to the database. These data load operations must occur in the minimum amount of time. A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours. You need to ensure that the minimum amount of data is lost. Which recovery model should the database use?

- A. FULL

- B. DBO_ONLY
- C. CONTINUE_AFTER_ERROR
- D. CHECKSUM
- E. NO_CHECKSUM
- F. SIMPLE
- G. Transaction log
- H. SKIP
- I. RESTART
- J. COPY_ONLY
- K. NORECOVERY
- L. BULK_LOGGED
- M. Differential
- N. STANDBY

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (ALL five questions with same context)

--\BURGOS

This seems to be a trick question. In order to minimize the amount of data loss, we use Full recovery model. BulkLogged model has the potential for some data loss. According to this logic (and the below reference), this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms189275.aspx>

QUESTION 72

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN). The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours. Five percent of the existing data is modified each day. The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands. Each data load adds 3 GB of data to the database. These data load operations must occur in the minimum amount of time. A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours. You need to ensure that the backup size is as small as possible. Which backup should you perform every two hours?

- A. BULK_LOGGED
- B. NO_CHECKSUM
- C. FULL
- D. RESTART
- E. CHECKSUM
- F. STANDBY
- G. DBO.ONLY
- H. NORECOVERY

- I. SIMPLE
- J. Transaction log
- K. Differential
- L. CONTINUE_AFTER_ERROR
- M. COPY_ONLY
- N. SKIP

Correct Answer: J

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (ALL five questions with same context)

--\BURGOS

Verified answer as correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms186865.aspx>

<http://msdn.microsoft.com/en-us/library/ms191429.aspx>

<http://msdn.microsoft.com/en-us/library/ms179478.aspx>

QUESTION 73

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN). The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours. Five percent of the existing data is modified each day. The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands. Each data load adds 3 GB of data to the database. These data load operations must occur in the minimum amount of time. A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours. You implement log shipping of the financial database to another SQL Server 2012 instance. You decide to failover to this secondary database. You need to ensure that all transactions will be replicated to the secondary database. Which backup option should you use?

- A. Differential
- B. Transaction Log
- C. FULL
- D. SIMPLE
- E. SKIP
- F. RESTART
- G. STANDBY
- H. CHECKSUM
- I. DBO_ONLY
- J. COPY_ONLY
- K. NORECOVERY
- L. NO_CHECKSUM
- M. CONTINUE_AFTER_ERROR
- N. BULK_LOGGED

Correct Answer: K

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (ALL five questions with same context)

--\BURGOS

According to these references, this answer looks correct.

References:

<http://technet.microsoft.com/en-us/library/ms187103.aspx>

<http://msdn.microsoft.com/en-us/library/ms191233.aspx>

<http://msdn.microsoft.com/en-us/library/ms178117.aspx>

QUESTION 74

You administer a Microsoft SQL Server 2012 database that contains a table named OrderDetail. You discover that the NCI_OrderDetail_CustomerID non-clustered index is fragmented. You need to reduce fragmentation. You need to achieve this goal without taking the index offline. Which Transact-SQL batch should you use?

- A. `CREATE INDEX NCI_OrderDetail_CustomerID ON OrderDetail.CustomerID WITH DROP EXISTING`
- B. `ALTER INDEX NCI_OrderDetail_CustomerID ON OrderDetail.CustomerID REORGANIZE`
- C. `ALTER INDEX ALL ON OrderDetail REBUILD`
- D. `ALTER INDEX NCI_OrderDetail_CustomerID ON OrderDetail.CustomerID REBUILD`

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms188388.aspx>

QUESTION 75

You administer a Microsoft SQL Server 2012 database that has multiple tables in the Sales schema. Some users must be prevented from deleting records in any of the tables in the Sales schema. You need to manage users who are prevented from deleting records in the Sales schema. You need to achieve this goal by using the minimum amount of administrative effort. What should you do?

- A. Create a custom database role that includes the users. Deny Delete permissions on the Sales schema for the custom database role.
- B. Include the Sales schema as an owned schema for the db_denydatawriter role. Add the users to the db_denydatawriter role.
- C. Deny Delete permissions on each table in the Sales schema for each user.
- D. Create a custom database role that includes the users. Deny Delete permissions on each table in the Sales schema for the custom database role.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms181127.aspx>
<http://msdn.microsoft.com/en-us/library/ms189121.aspx>
<http://msdn.microsoft.com/en-us/library/ff848791.aspx>
<http://msdn.microsoft.com/en-us/library/ee677610.aspx>
<http://msdn.microsoft.com/en-us/library/ms187936.aspx>

QUESTION 76

You administer a Microsoft SQL Server 2012 database. The database has a table named Customers owned by UserA and another table named Orders owned by UserB. You also have a stored procedure named GetCustomerOrderInfo owned by UserB. GetCustomerOrderInfo selects data from both tables.

You create a new user named UserC. You need to ensure that UserC can call the GetCustomerOrderInfo stored procedure. You also need to assign only the minimum required permissions to UserC. Which permission or permissions should you assign to UserC? Choose all that apply.

- A. The Select permission on Customers
- B. The Execute permission on GetCustomerOrderInfo
- C. The Take Ownership permission on Customers
- D. The Control permission on GetCustomerOrderInfo
- E. The Take Ownership permission on Orders
- F. The Select permission on Orders

Correct Answer: AB

Section: (none)

Explanation**Explanation/Reference:**

-- BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms188676.aspx>
<http://stackoverflow.com/questions/2212044/sql-server-how-to-permission-schemas>
http://sqlservercentral.com/blogs/steve_jones/2012/03/14/ownership-chains-in-sql-server

QUESTION 77

You create an availability group named HaContoso that has replicas named Server01/HA, Server02/HA, and Server03/HA. Currently, Server01/HA is the primary replica. You need to ensure that the following requirements are met:

- Backup operations occur on Server02/HA.
- If Server02/HA is unavailable, backup operations occur on Server03/HA.
- Backup operations do not occur on Server01/HA.

How should you configure HaContoso?

- A.
 - Set the backup preference of HaContoso to Prefer Secondary.
 - Set the backup priority of Server02/HA to 20.
 - Set the backup priority of Server03/HA to 10.
- B.
 - Set the backup preference of HaContoso to Secondary only.

- Set the backup priority of Server02/HA to 20.
- Set the backup priority of Server03/HA to 10.
- C. ▪ Set the backup preference of HaContoso to Secondary only.
- Set the backup priority of Server02/HA to 10.
- Set the backup priority of Server03/HA to 20.
- D. ▪ Set the exclude replica of Server01/HA to true.
- Set the backup priority of Server02/HA to 10.
- Set the backup priority of Server03/HA to 20.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ff877884.aspx>

<http://msdn.microsoft.com/en-us/library/hh245119.aspx>

<http://msdn.microsoft.com/en-us/library/hh710053.aspx>

Exam B

QUESTION 1

You administer a Microsoft SQL Server database. You want to import data from a text file to the database. You need to ensure that the following requirements are met:

- Data import is performed by using a stored procedure.
- Data is loaded as a unit and is minimally logged.

Which data import command and recovery model should you choose? (To answer, drag the appropriate data import command or recovery model to the appropriate location or locations in the answer area. Each data import command or recovery model may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Command/Recovery Model Name	Command/Recovery Model
BCP	Data import command
BULK INSERT	Recovery model
Bulk-logged	
OPENDATASOURCE	
Full	

Correct Answer:

Command/Recovery Model Name	Command/Recovery Model
BCP	Data import command
BULK INSERT	Recovery model
Bulk-logged	
OPENDATASOURCE	
Full	

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES
--\BURGOS

According to these references, this answer looks correct.

References:

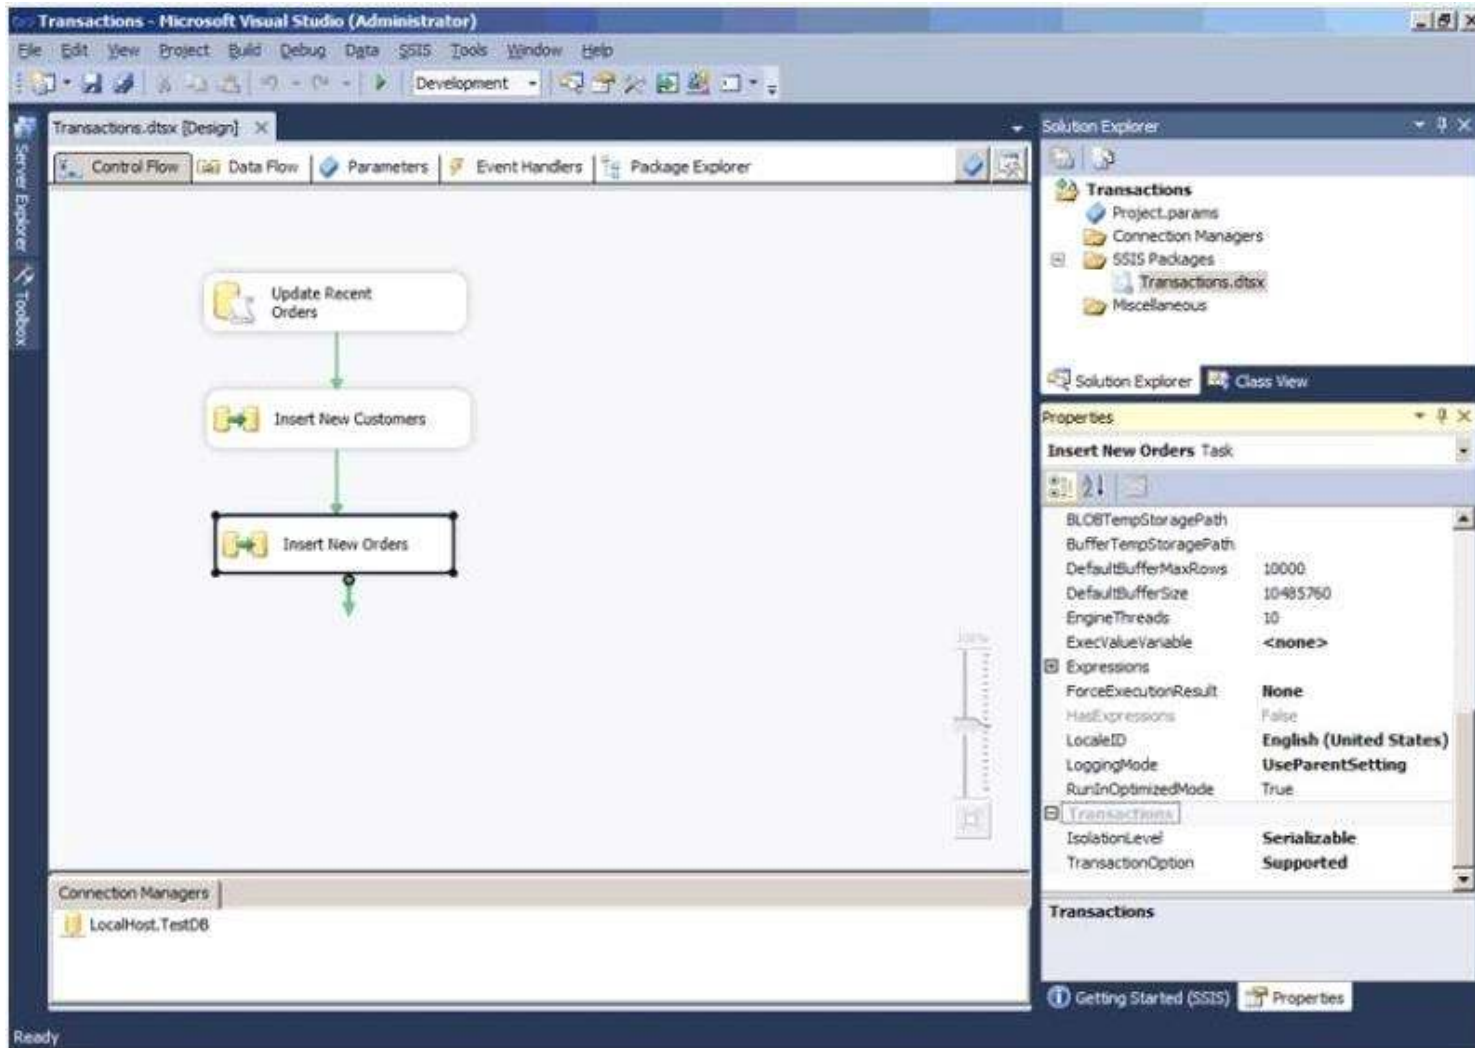
<http://msdn.microsoft.com/en-us/library/ms162802.aspx>

<http://msdn.microsoft.com/en-us/library/ms188365.aspx>
<http://msdn.microsoft.com/en-us/library/ms175937.aspx>
<http://msdn.microsoft.com/en-us/library/aa337544.aspx>

QUESTION 2

DRAW AND DROP

You are designing a package control flow. The package moves sales order data from a SQL Azure transactional database to an on-premise reporting database. The package will run several times a day, while new sales orders are being added to the transactional database. The current design of the package control flow is shown in the answer area. (Click the Exhibit button.)



The Insert New Orders Data Flow task must meet the following requirements:

- Usage of the tempdb database should not be impacted.
- Concurrency should be maximized, while only reading committed transactions.
- If the task fails, only that task needs to be rolled back.

You need to configure the Insert New Orders Data Flow task to meet the requirements. How should you configure the transaction properties? (To answer, select the appropriate setting or settings in the answer area.)

Select and Place:

Chaos	
NotSupported	
ReadCommitted	
ReadUncommitted	
RepeatableRead	
Required	
Serializable	
Snapshot	
Supported	
Unspecified	

Testing SSIS Project - Microsoft Visual Studio (Administrator)

File Edit View Project Build Debug Data SSIS Tools Window Help

Development

Package1.dtsx [Design]*

Server Explorer Toolbox SSIS Toolbox

Update Recent Orders

Insert New Customers

Insert New Orders

Connection Managers

Ready

Properties

Insert New Orders Task

BLOBTempStoragePath

BufferTempStoragePat

DefaultBufferMaxRows 10000

DefaultBufferSize 10485760

EngineThreads 10

ExecValueVariable <none>

Expressions

ForceExecutionResult None

HasExpressions False

LocaleID English (United S

LoggingMode UseParentSetting

RunInOptimizedMode True

Transactions


IsolationLevel

TransactionOption

Getting Started (SSIS) Properties

Correct Answer:

Chaos	
NotSupported	
ReadCommitted	
ReadUncommitted	
RepeatableRead	
Required	
Serializable	
Snapshot	
Supported	
Unspecified	



Section: (none)
Explanation

Explanation/Reference:

-- BURGOS - YES (Not Drag-and-drop, this question used four combo-box)
 --\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms137690.aspx>
<http://msdn.microsoft.com/en-us/library/ms137749.aspx>
<http://msdn.microsoft.com/en-us/library/microsoft.sqlserver.dts.runtime.dtscontainer.isolationlevel.aspx>
<http://msdn.microsoft.com/en-us/library/ms173763.aspx>

QUESTION 3

You are designing an extract, transform, load (ETL) process with SQL Server Integration Services (SSIS). Two packages, Package A and Package B, will be designed. Package A will execute Package B. Both packages must reference a file path corresponding to an input folder where files will be located for further processing. You need to design a solution so that the file path can be easily configured with the least administrative and development effort. Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲▼</div><div></div></div>	<div>Create an SSIS project containing Package A and Package B.</div> <div>Add a project parameter named FilePath.</div> <div>Add a package parameter named FilePath to Package A and Package B.</div> <div>Deploy the project to the SSIS catalog.</div> <div>Configure the project to use an Environment that sets the FilePath parameter.</div> <div>Add a Parent Package Variable configuration in Package B to set the FilePath variable.</div> <div>Add a configuration file in Package A to set the FilePath value.</div> <div>Create one SSIS project containing only Package A and another SSIS project containing only Package B.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Create an SSIS project containing Package A and Package B.

Add a project parameter named FilePath.

Deploy the project to the SSIS catalog.

Configure the project to use an Environment that sets the FilePath parameter.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

Reference:

<http://msdn.microsoft.com/en-us/library/hh479588.aspx>
<http://msdn.microsoft.com/en-us/library/hh213290.aspx>
<http://msdn.microsoft.com/en-us/library/hh213373.aspx>

QUESTION 4

You develop a SQL Server Integration Services (SSIS) project by using the Project Deployment model. The project contains many packages. It is deployed on a server named Development1. The project will be deployed to several servers that run SQL Server 2012. The project accepts one required parameter. The data type of the parameter is a string. A SQL Agent job is created that will call the master.dtsx package in the project. A job step is created for the SSIS package. The job must pass the value of an SSIS Environment Variable to the project parameter. The value of the Environment Variable must be configured differently on each server that runs SQL Server. The value of the Environment Variable must provide the server name to the project parameter. You need to configure SSIS on the Development1 server to pass the Environment Variable to the package. Which four actions should you perform in sequence by using SQL Server Management Studio? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲</div><div>▼</div></div>	<div>In the SSIS catalog, create a folder named DevelopmentEnv.</div> <div>In the SSIS catalog, create an Environment named DevelopmentEnv.</div> <div>In the DevelopmentEnv Environment, create an Environment Variable. Name the Environment Variable PackageEnvironment. Set the Environment Variable value to Development1.</div> <div>In the project configuration, create an Environment reference to DevelopmentEnv. Assign the value of the PackageEnvironment Environment Variable to the project parameter.</div> <div>In the project configuration, edit the value of the project parameter to DevelopmentEnv.</div> <div>In the SQL Agent job, choose the DevelopmentEnv Environment in the Configuration tab of the SQL Server Integration Services job step.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

In the SSIS catalog, create an Environment named DevelopmentEnv.

In the DevelopmentEnv Environment, create an Environment Variable. Name the Environment Variable PackageEnvironment. Set the Environment Variable value to Development1.

In the project configuration, create an Environment reference to DevelopmentEnv. Assign the value of the PackageEnvironment Environment Variable to the project parameter.

In the SQL Agent job, choose the DevelopmentEnv Environment in the Configuration tab of the SQL Server Integration Services job step.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

Verified answer as correct.

References:

<http://msdn.microsoft.com/en-us/library/hh479588.aspx>

<http://msdn.microsoft.com/en-us/library/hh213230.aspx>

<http://msdn.microsoft.com/en-us/library/hh213214.aspx>

QUESTION 5

You administer several Microsoft SQL Server 2012 database servers. Merge replication has been configured for an application that is distributed across offices throughout a wide area network (WAN). Many of the tables involved in replication use the XML and varchar(max) data types. Occasionally, merge replication fails due to timeout errors. You need to reduce the occurrence of these timeout errors. What should you do?

- A. Set the Merge agent on the problem subscribers to use the slow link agent profile.
- B. Create a snapshot publication, and reconfigure the problem subscribers to use the snapshot publication.
- C. Change the Merge agent on the problem subscribers to run continuously.
- D. Set the Remote Connection Timeout on the Publisher to 0.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms151223.aspx>

<http://msdn.microsoft.com/en-us/library/ms147839.aspx>

QUESTION 6

You administer three Microsoft SQL Server 2012 servers named ServerA, ServerB, and ServerC. ServerA is

the acting principal and ServerB is the mirror. You need to add ServerC as a witness to the existing mirroring session between ServerA and ServerB. You need to achieve this goal without delaying synchronization. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲▼</div><div></div></div>	<div>On ServerC, create an endpoint for use by the witness.</div> <div>Ensure that the same Windows Login exists on each server and grant Connect permissions to each server's endpoint.</div> <div>On ServerA, alter the principal database to use the endpoint on ServerC as the witness.</div> <div>On ServerA, pause the mirroring session between ServerA and ServerB.</div> <div>On ServerB, alter the principal database to use te endpoint on ServerC as the witness.</div> <div>Ensure that the same Proxy exists on each server and grant Connect permissions to each server's endpoint.</div> <div>On ServerA, resume the mirroring session between ServerA and ServerB.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

On ServerC, create an endpoint for use by the witness.

Ensure that the same Windows Login exists on each server and grant Connect permissions to each server's endpoint.

On ServerA, alter the principal database to use the endpoint on ServerC as the witness.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

Verified answer as correct.

Reference:

<http://msdn.microsoft.com/en-us/library/ms190430.aspx>

QUESTION 7

You administer three Microsoft SQL Server 2008 R2 instances. Database mirroring is configured in High-Safety mode with Automatic Failover between the following three servers:

- SQL1 is the Principal server.
- SQL2 is the mirror server.
- SQL3 is the witness server.

You need to upgrade SQL1 and SQL2 to SQL Server 2012. You need to ensure that downtime is minimized during the upgrade. Which six actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲</div><div>▼</div></div>	<div>Configure log shipping between SQL1 and SQL2.</div> <div>Upgrade SQL1 to SQL Server 2012.</div> <div>Upgrade SQL2 to SQL Server 2012.</div> <div>Disable log shipping between SQL1 and SQL2.</div> <div>Manually failover the database from SQL1 to SQL2.</div> <div>Manually failover the database from SQL2 to SQL1.</div> <div>Add SQL3 back to the database mirroring solution.</div> <div>Remove SQL3 from the database mirroring solution.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Remove SQL3 from the database mirroring solution.

Upgrade SQL2 to SQL Server 2012.

Manually failover the database from SQL1 to SQL2.

Upgrade SQL1 to SQL Server 2012.

Manually failover the database from SQL2 to SQL1.

Add SQL3 back to the database mirroring solution.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

Verified answer as correct.

Reference:

<http://msdn.microsoft.com/en-us/library/bb677181.aspx>

QUESTION 8

You administer a Microsoft SQL Server 2012 environment that contains a production SQL Server 2005 instance named SQL2005 and a development SQL Server 2012 instance named SQL2012. The development team develops a new application that uses the SQL Server 2012 functionality. You are planning to migrate a database from SQL2005 to SQL2012 so that the development team can test their new application. You need to migrate the database without affecting the production environment. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲</div><div>▼</div></div>	Perform a transaction log backup on SQL2005.
	Perform a full database backup on SQL2005.
	Perform a VSS backup on the database on SQL2005.
	Restore the VSS backup on SQL2012.
	Restore the full database backup on SQL 2012.
	Restore the database backup and transaction log backup on SQL 2012.
	Change the compatibility level for the database to 120 on SQL2012.
	Change the compatibility level for the database to 110 on SQL2012.

<< Move

Remove >>

Correct Answer:

Perform a full database backup on SQL2005.

Restore the full database backup on SQL 2012.

Change the compatibility level for the database to 110 on SQL2012.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - (Not Sure)
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms187048.aspx>
<http://msdn.microsoft.com/en-us/library/ms186289.aspx>
<http://msdn.microsoft.com/en-us/library/ms175477.aspx>
<http://msdn.microsoft.com/en-us/library/ms189860.aspx>
<http://msdn.microsoft.com/en-us/library/ms179314.aspx>
<http://msdn.microsoft.com/en-us/library/ms175526.aspx>

<http://msdn.microsoft.com/en-us/library/ms191539.aspx>
<http://msdn.microsoft.com/en-us/library/ms191429.aspx>
<http://msdn.microsoft.com/en-us/library/ms191253.aspx>

QUESTION 9

You administer a Microsoft SQL Server 2012 database. The database is backed up according to the following schedule:

- Daily full backup at 23:00 hours.
- Differential backups on the hour, except at 23:00 hours.
- Log backups every 10 minutes from the hour, except on the hour.

The database uses the Full recovery model. A developer accidentally drops a number of tables and stored procedures from the database between 22:40 hours and 23:10 hours. You perform a database restore at 23:30 hours to recover the dropped table. You need to restore the database by using the minimum amount of administrative effort. You also need to ensure minimal data loss. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲▼</div><div></div></div>	<div>Restore the most recent full backup.</div> <div>Restore the full backup taken the previous night.</div> <div>Restore the differential backup taken at 22:00 hours.</div> <div>Restore the transaction log backup taken at 22:40 hours.</div> <div>Restore each transaction log backup taken from 22:00 till 22:40 hours.</div> <div>Restore each transaction log backup taken from the most recent full backup.</div> <div>Restore each differential database backup taken from the previous night's full backup.</div> <div>Restore each transaction log backup taken from the previous night's full backup till 22:40 hours.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Restore the full backup taken the previous night.
Restore the differential backup taken at 22:00 hours.
Restore each transaction log backup taken from 22:00 till 22:40 hours.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms187048.aspx>
<http://msdn.microsoft.com/en-us/library/ms186289.aspx>
<http://msdn.microsoft.com/en-us/library/ms175477.aspx>
<http://msdn.microsoft.com/en-us/library/ms189860.aspx>
<http://msdn.microsoft.com/en-us/library/ms179314.aspx>
<http://msdn.microsoft.com/en-us/library/ms175526.aspx>
<http://msdn.microsoft.com/en-us/library/ms191539.aspx>
<http://msdn.microsoft.com/en-us/library/ms191429.aspx>
<http://msdn.microsoft.com/en-us/library/ms191253.aspx>

QUESTION 10

You administer a Microsoft SQL Server 2012 database. All database traffic to the SQL Server must be encrypted by using secure socket layer (SSL) certificates or the connection must be refused. Network administrators have deployed server certificates to the Windows store of all Windows servers on the network from a trusted Certificate Authority. This is the only Certificate Authority allowed to distribute certificates on the network.

You enable the Force Encryption flag for the MSSQLServer protocols, but client computers are unable to connect. They receive the following error message:

"A connection was successfully established with the server, but then an error occurred during the pre-login handshake, (provider: SSL Provider, error: 0 - The certificate chain was issued by an authority that is not trusted.) (Microsoft SQL Server)"

You notice the following entry in the SQL Server log:

"A self-generated certificate was successfully loaded for encryption."

You need to configure SQL Server to encrypt all client traffic across the network. You also need to ensure that client computers are able to connect to the server by using a trusted certificate. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> <div></div> </div>	<div>Restart the SQL Server.</div> <div>Leave the certificate blank in the drop-down list on the CERTIFICATES tab.</div> <div>Choose the new root-level certificate from the drop-down list on the CERTIFICATES tab.</div> <div>Install Certificate Services on the SQL Server, and create a new root-level certificate.</div> <div>From the SQL Configuration Manager on the SQL Server, open the PROTOCOLS properties for the SQL instance.</div> <div>Choose the server certificate provided by the network administrators from the drop-down list on the CERTIFICATES tab.</div> <div>From the SQL Configuration Manager on every client computer that will be connecting to SQL Server, open the PROTOCOLS properties for the SQL instance.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

From the SQL Configuration Manager on the SQL Server, open the PROTOCOLS properties for the SQL instance.

Choose the server certificate provided by the network administrators from the drop-down list on the CERTIFICATES tab.

Restart the SQL Server.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

Reference:

<http://thesqldude.com/2012/04/21/setting-up-ssl-encryption-for-sql-server-using-certificates-issues-tips-tricks/>

QUESTION 11

You administer a Microsoft SQL Server 2012 clustered instance that has two nodes named Node 1 and Node 2. Node 1 fails and the cluster fails over to Node 2. You need to replace Node 1 and add it to the cluster. Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> <div></div> </div>	<div>Evict Node 1 from the Windows Failover Cluster.</div> <div>Install Windows on a new server to replace Node 1.</div> <div>Run SQL Server Setup to add Node 1 to the failover cluster.</div> <div>Run Cluster Administrator Setup to add Node 1 to the failover cluster.</div> <div>Add Node 1 to the existing cluster by using SQL Server Configuration Manager.</div> <div>Add Node 1 to the existing cluster by using the Windows Failover Cluster Manager.</div> <div>Register the secondary instance with the Cluster Manager by using SQL Server Management Studio.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Evict Node 1 from the Windows Failover Cluster.

Install Windows on a new server to replace Node 1.

Add Node 1 to the existing cluster by using the Windows Failover Cluster Manager.

Run SQL Server Setup to add Node 1 to the failover cluster.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to this reference, this answer looks correct.

Reference:

<http://technet.microsoft.com/en-us/library/ms181075.aspx>

QUESTION 12

You administer a single Microsoft SQL Server instance on a two-node failover cluster that has nodes named Node A and Node B. The instance is currently running on Node A. You want to patch both Node A and Node B by using the most recent SQL Server Service Pack. You need to ensure that the following requirements are met:

- Both nodes receive the update.
- Downtime is minimized.
- No data is lost.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲▼</div><div></div></div>	<div>Pause Node A.</div> <div>Pause Node B.</div> <div>Failover from Node A to Node B.</div> <div>Start the SQL Server service on both nodes.</div> <div>Install the service pack on Node A.</div> <div>Install the service pack on Node B.</div> <div>Stop the SQL Server service on both nodes.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Install the service pack on Node B.

Failover from Node A to Node B.

Install the service pack on Node A.

Section: (none)**Explanation****Explanation/Reference:**

-- BURGOS - NO
--\BURGOS

-- Burgos.

Take care, exists two forms of this question, one with Node A as active running, another with Node B.

--\Burgos.

According to these references, this answer looks correct.

References:

<http://technet.microsoft.com/en-us/library/ms191295.aspx>

<http://msdn.microsoft.com/en-us/library/ms191009.aspx>

QUESTION 13

You administer a Microsoft SQL Server instance. You use a two-node SQL Server failover cluster. Node B is primary, and Node A is secondary. You need to install a security patch on both nodes. You need to ensure that the following requirements are met:

- Both nodes receive the update.
- Downtime is minimized.
- No data is lost.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> <div></div> </div>	<div> <div>Pause Node B.</div> <div>Pause Node A.</div> <div>Failover from Node B to Node A.</div> <div>Failover from Node A to Node B.</div> <div>Install the security patch on Node B.</div> <div>Install the security patch on Node A.</div> <div>Stop the SQL Server services on both nodes.</div> </div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Install the security patch on Node A.

Failover from Node B to Node A.

Install the security patch on Node B.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

-- Burgos.

Take care, exists two forms of this question, one with Node A as active running, another with Node B.

--\Burgos.

According to these references, this answer looks correct.

References:

<http://technet.microsoft.com/en-us/library/ms191295.aspx>

<http://msdn.microsoft.com/en-us/library/ms191009.aspx>

QUESTION 14

DRAG DROP

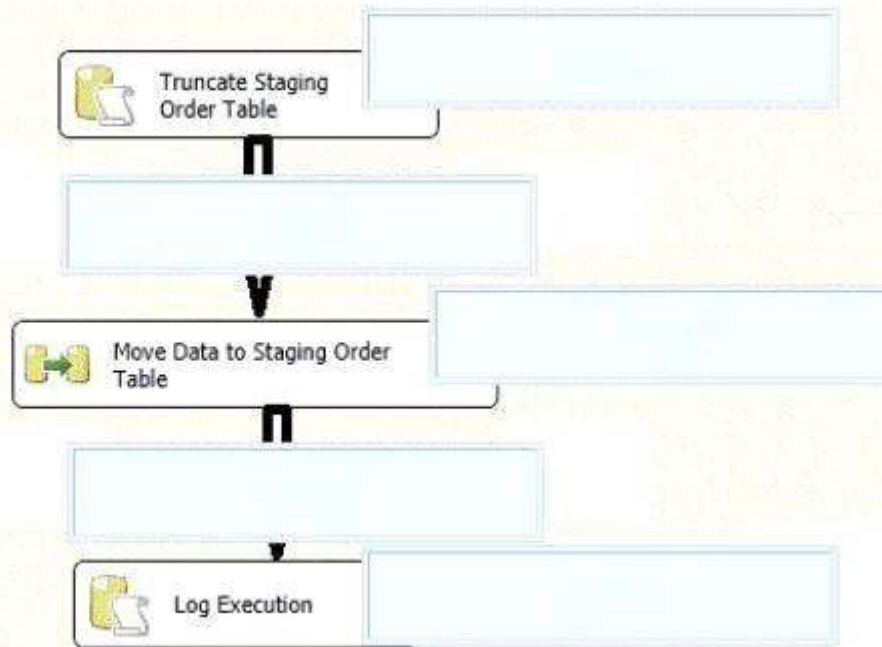
You are designing a SQL Server Integration Services (SSIS) package. The package moves order-related data to a staging table named Order. Every night the staging data is truncated and then all the recent orders from the online store database are inserted into the staging table. Your package must meet the following requirements:

- If the truncate operation fails, the package execution must stop and report an error.
- If the Data Flow task that moves the data to the staging table fails, the entire refresh operation must be rolled back.
- For auditing purposes, a log entry must be entered in a SQL log table after each execution of the Data Flow task.

The TransactionOption property for the package is set to Required. You need to design the package to meet the requirements. How should you design the control flow for the package? (To answer, drag the appropriate setting from the list of settings to the correct location or locations in the answer area.)

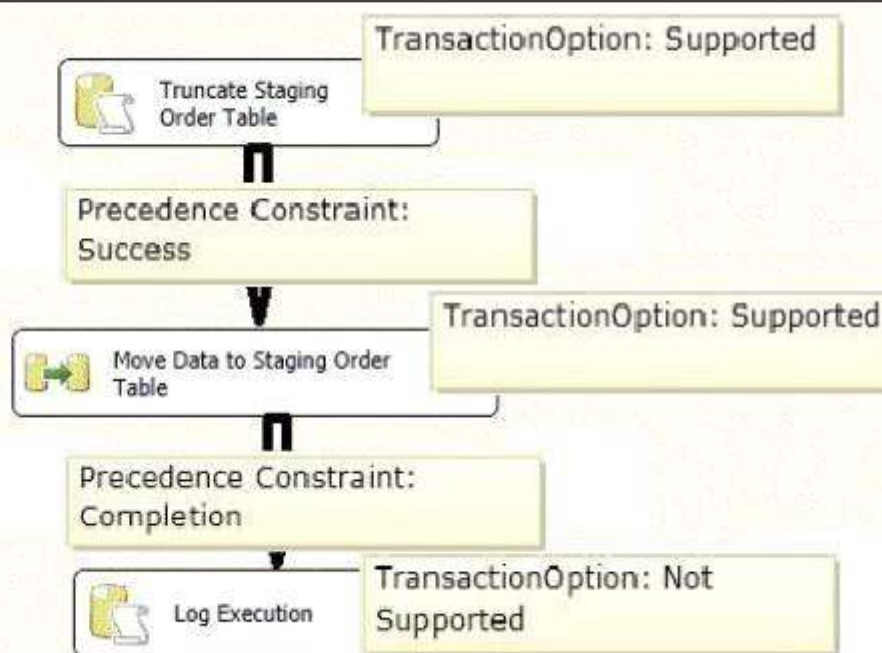
Select and Place:

- TransactionOption: Required
- TransactionOption: Supported
- TransactionOption: Not Supported
- Precedence Constraint: Failure
- Precedence Constraint: Success
- Precedence Constraint: Completion



Correct Answer:

- TransactionOption: Required
- TransactionOption: Supported
- TransactionOption: Not Supported
- Precedence Constraint: Failure
- Precedence Constraint: Success
- Precedence Constraint: Completion



Section: (none)

Explanation

Explanation/Reference:

- BURGOS - YES
- \BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms137690.aspx>
<http://msdn.microsoft.com/en-us/library/ms141144.aspx>

QUESTION 15

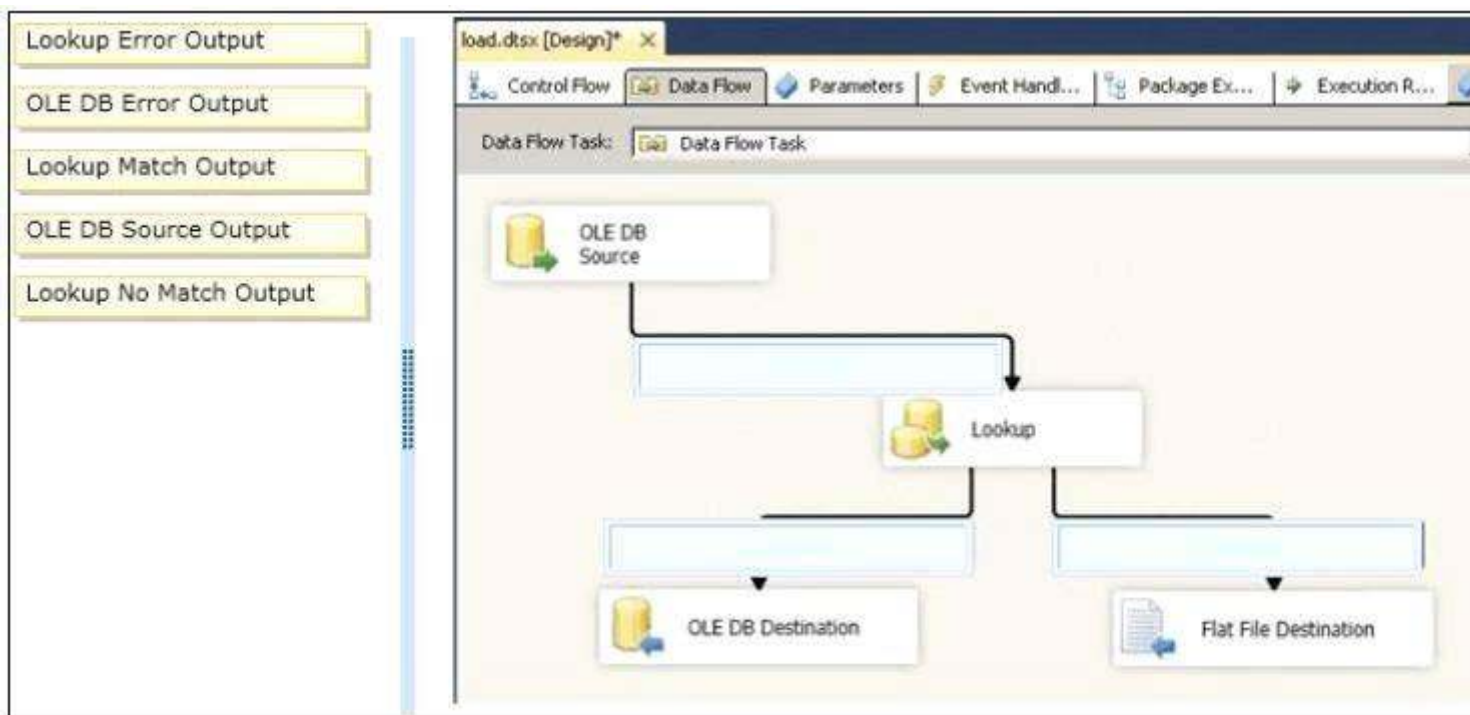
DRAG AND DROP

You are developing a SQL Server Integration Services (SSIS) package to insert new data into a data mart. The package uses a Lookup transformation to find matches between the source and destination. The data flow has the following requirements:

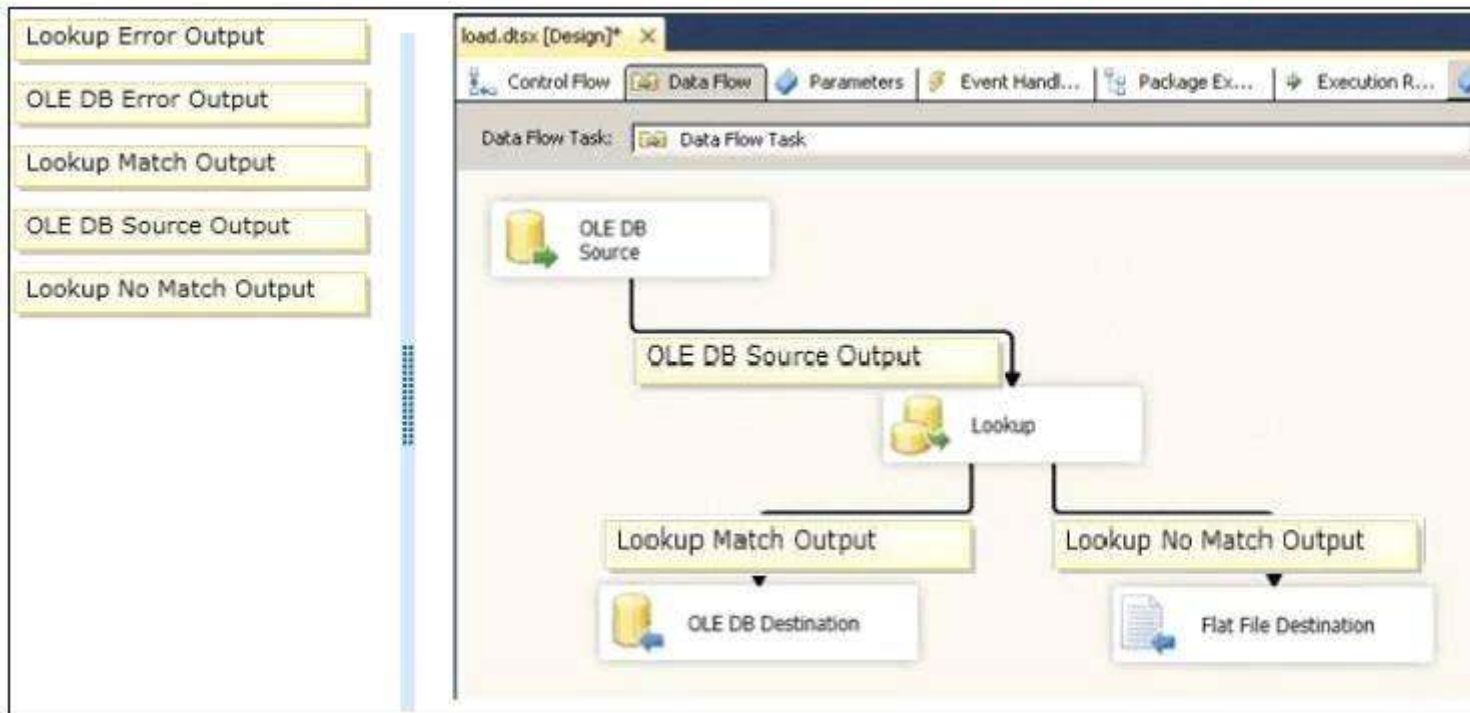
- New rows must be inserted.
- Lookup failures must be written to a flat file.

In the Lookup transformation, the setting for rows with no matching entries is set to Redirect rows to no match output. You need to configure the package to direct data into the correct destinations. How should you design the data flow outputs? (To answer, drag the appropriate transformation from the list of answer options to the correct location in the answer area.)

Select and Place:



Correct Answer:



Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms141821.aspx>

<http://msdn.microsoft.com/en-us/library/bb895366.aspx>

QUESTION 16

You are building a SQL Server Integration Services (SSIS) package to load product data sourced from a SQL Azure database to a data warehouse. Before the product data is loaded, you create a batch record by using an Execute SQL task named Create Batch. After successfully loading the product data, you use another Execute SQL task named Set Batch Success to mark the batch as successful.



You need to create and execute an Execute SQL task to mark the batch as failed if either the Create Batch or Load Products task fails. Which three steps should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div style="border: 1px solid gray; height: 300px; width: 100%;"></div>	<div style="border: 1px solid gray; padding: 5px;"> <p>Set the Multiple Constraints option to Logical OR.</p> <p>Set the Evaluation Operation option to Expression OR Constraint.</p> <p>Connect the Create Batch and Load Products tasks to the Set Batch Failure task with failure constraints.</p> <p>Create an Execute SQL task named Set Batch Failure to mark the batch as failed.</p> <p>Set the Multiple Constraints option to Logical AND.</p> <p>Connect the Create Batch and Load Products tasks to the Set Batch Failure task with completion constraints.</p> <p>Set the ForceExecutionValue property of the Create Batch task to True.</p> </div>
<div style="display: inline-block; border: 1px solid gray; padding: 2px 10px; margin: 2px;"> << Move </div> <div style="display: inline-block; border: 1px solid gray; padding: 2px 10px; margin: 2px;"> Remove >> </div>	

Correct Answer:

Create an Execute SQL task named Set Batch Failure to mark the batch as failed.

Connect the Create Batch and Load Products tasks to the Set Batch Failure task with failure constraints.

Set the Multiple Constraints option to Logical OR.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms141003.aspx>

<http://msdn.microsoft.com/en-us/library/ms141261.aspx>

<http://msdn.microsoft.com/en-us/library/ms141722.aspx>

<http://msdn.microsoft.com/en-us/library/ms139895.aspx>

QUESTION 17

HOTSPOT

You are developing a data flow to load sales data into a fact table. In the data flow, you configure a Lookup Transformation in full cache mode to look up the product data for the sale. The lookup source for the product data is contained in two tables. You need to set the data source for the lookup to be a query that combines the two tables. Which page of the Lookup Transformation Editor should you select to configure the query? To answer, select the appropriate page in the answer area.

Hot Area:

Lookup Transformation Editor

This transform enables the performance of simple equi-joins between the input and a reference data set.

General
Connection
Columns
Advanced
Error Output

Cache mode

☒ Full cache
☐ Partial cache
☐ No cache


Connection type

☐ Cache connection manager
☒ OLE DB connection manager

Specify how to handle rows with no matching entries

Fail component

Correct Answer:



Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms141821.aspx>

<http://msdn.microsoft.com/en-us/library/ms189697.aspx>

QUESTION 18

You are building a SQL Server Integration Services (SSIS) package to load data from all files that are automatically copied to a directory each night through an external FTP process. You need to load data from all copied files to a destination table in SQL Server. Which three steps should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div> <div></div> <div></div> </div> <div></div> </div>	<div>On the Collection page, select the Foreach File Enumerator option and configure the directory and files to process.</div> <div>Open the Control Flow page of the package.</div> <div>Create a For Loop container.</div> <div>Create a File System task.</div> <div>On the Collection page, select the Foreach ADO Enumerator option and configure the ADO object source variable.</div> <div>Create a Foreach Loop container.</div> <div>Open the Data Flow page of the package.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Open the Control Flow page of the package.

Create a Foreach Loop container.

On the Collection page, select the Foreach File Enumerator option and configure the directory and files to process.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms141724.aspx>

<http://msdn.microsoft.com/en-us/library/ms137728.aspx>

QUESTION 19

DRAG DROP

You are creating a sales data warehouse. When a product exists in the product dimension, you update the product name. When a product does not exist, you insert a new record. In the current implementation, the DimProduct table must be scanned twice, once for the insert and again for the update. As a result, inserts and updates to the DimProduct table take longer than expected. You need to create a solution that uses a single command to perform an update and an insert. How should you use a MERGE T-SQL statement to accomplish this goal? (To answer, drag the appropriate answer choice from the list of options to the correct location or locations in the answer area. You may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Source

Target

Product

Inserted

MATCHED THEN

NOT MATCHED BY TARGET THEN

NOT MATCHED BY PRODUCT AND (Target.ProductID is NULL) THEN

MATCHED AND (Target.ProductID is NULL) THEN

MERGE dbo.DimProduct AS

USING (SELECT ProductID, ProductName, ProductColor, ProductCategory

FROM dbo.StagingProduct) AS

ON (Target.ProductID = Source.ProductID)

WHEN

UPDATE SET Target.ProductName = Source.ProductName

WHEN

INSERT (ProductID, ProductName, ProductColor, ProductCategory)

VALUES (Source.ProductID, Source.ProductName, Source.ProductColor, Source.ProductCategory)

OUTPUT \$action, Inserted.*, Deleted.*;

Correct Answer:

Product

Inserted

NOT MATCHED BY PRODUCT AND (Target.ProductID is NULL) THEN

MATCHED AND (Target.ProductID is NULL) THEN

MERGE dbo.DimProduct AS Target

USING (SELECT ProductID, ProductName, ProductColor, ProductCategory

FROM dbo.StagingProduct) AS Source

ON (Target.ProductID = Source.ProductID)

WHEN MATCHED THEN

UPDATE SET Target.ProductName = Source.ProductName

WHEN NOT MATCHED BY TARGET THEN

INSERT (ProductID, ProductName, ProductColor, ProductCategory)

VALUES (Source.ProductID, Source.ProductName, Source.ProductColor, Source.ProductCategory)

OUTPUT \$action, Inserted.*, Deleted.*;

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/bb510625.aspx>

<http://msdn.microsoft.com/en-us/library/cc280522.aspx>

QUESTION 20

DRAG DROP

You are developing a SQL Server Integration Services (SSIS) package that imports unsorted data into a data warehouse hosted on SQL Azure. You have the following requirements:


- A destination table must contain all of the data in two source tables.
- Duplicate records must be inserted into the destination table.

You need to develop a data flow that imports the data while meeting the requirements. How should you develop

the data flow? (To answer, drag the appropriate transformation from the list of transformations to the correct location in the answer area.)

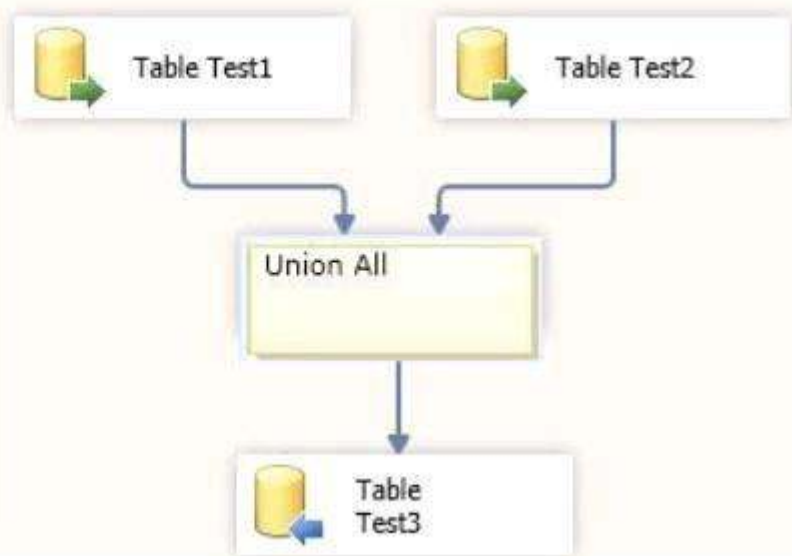
Select and Place:

- Merge
- Merge Join
- Sort
- Union All



Correct Answer:

- Merge
- Merge Join
- Sort
- Union All



Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - NO
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms141703.aspx>

<http://msdn.microsoft.com/en-us/library/ms141775.aspx>

<http://msdn.microsoft.com/en-us/library/ms141020.aspx>
<http://msdn.microsoft.com/en-us/library/ms140182.aspx>

QUESTION 21

DRAW AND DROP

You are developing a SQL Server Integration Services (SSIS) package that imports data into a data warehouse. You are developing the part of the SSIS package that populates the ProjectDates dimension table. The business key of the ProjectDates table is the ProjectName column. The business user has given you the dimensional attribute behavior for each of the four columns in the ProjectDates table:

- ExpectedStartDate - New values should be tracked over time.
- ActualStartDate - New values should not be accepted.
- ExpectedEndDate - New values should replace existing values.
- ActualEndDate - New values should be tracked over time.

You use the SSIS Slowly Changing Dimension Transformation. You must configure the Change Type value for each source column. Which settings should you select? (To answer, select the appropriate setting or settings in the answer area. Each Change Type may be used once, more than once, or not at all.)

Select and Place:

Fixed Attribute	▼
Changing Attribute	▼
Historical Attribute	▼

Dimension Columns	Change Type
ExpectedStartDate	▼
ActualStartDate	▼
ExpectedEndDate	▼
ActualEndDate	▼

Correct Answer:

Fixed Attribute	▼
Changing Attribute	▼
Historical Attribute	▼

Dimension Columns	Change Type
ExpectedStartDate	Historical Attribute ▼
ActualStartDate	Fixed Attribute ▼
ExpectedEndDate	Changing Attribute ▼
ActualEndDate	Historical Attribute ▼

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES (Not Drag-and-drop, this question used four combo-box)
--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms141715.aspx>

<http://msdn.microsoft.com/en-us/library/ms141662.aspx>

QUESTION 22

You administer a Microsoft SQL Server database. Service accounts for SQL Agent are configured to use a local user. A Microsoft SQL Server Integration Services (SSIS) job step has been created within a SQL Server Agent job. The SSIS package accesses a network share when exporting data from a SQL Server database. When you execute the SQL Server Agent job, it fails due to a permissions failure on a share on a remote server. You need to ensure that the SQL Server Agent job can execute the SSIS package. Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲▼</div><div></div></div>	<div>Add a proxy that references the local user.</div> <div>Add a proxy that references the credential.</div> <div>Create a local user account and grant local administrator on the SQL Server instance.</div> <div>Create a credential that references the local user.</div> <div>Create a credential that references the domain user.</div> <div>Assign the proxy to the Operating System subsystem.</div> <div>Assign the proxy to the SSIS package execution subsystem.</div> <div>Create a domain user account and grant permissions to the domain user account to access the network share.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Create a domain user account and grant permissions to the domain user account to access the network share.

Create a credential that references the domain user.

Add a proxy that references the credential.

Assign the proxy to the SSIS package execution subsystem.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms175834.aspx>

<http://msdn.microsoft.com/en-us/library/ms189522.aspx>

<http://msdn.microsoft.com/en-us/library/ms190703.aspx>

<http://msdn.microsoft.com/en-us/library/ms161950.aspx>

QUESTION 23

DRAG AND DROP

You administer several Microsoft SQL Server 2012 servers. Your company has a number of offices across the world connected by using a wide area network (WAN). Connections between offices vary significantly in both bandwidth and reliability. You need to identify the correct replication method for each scenario. What should you do? (To answer, drag the appropriate replication method or methods to the correct location or locations in the answer area. Each replication method may be used once, more than once, or not at all.)

Select and Place:

Replication Method	Scenario
Transactional Replication	Multiple databases on the same low-latency subnet must allow applications to write changes locally, and these changes must be replicated to all related databases.
Peer-to-Peer Replication	An order summary table is repopulated once a week. This table must be replicated to all databases.
Snapshot Replication	Field offices using unreliable connections keep a local copy of the product catalog and process orders locally. These orders must be periodically replicated to all other offices.
Merge Replication	Information in an order-tracking database must be replicated across a low-latency connection as changes occur to multiple reporting databases.

Correct Answer:

Replication Method	Scenario
Transactional Replication	Multiple databases on the same low-latency subnet must allow applications to write changes locally, and these changes must be replicated to all related databases.
Peer-to-Peer Replication	An order summary table is repopulated once a week. This table must be replicated to all databases.
Snapshot Replication	Field offices using unreliable connections keep a local copy of the product catalog and process orders locally. These orders must be periodically replicated to all other offices.
Merge Replication	Information in an order-tracking database must be replicated across a low-latency connection as changes occur to multiple reporting databases.

Section: (none)

Explanation

Explanation/Reference:

--BURGOS - YES (a really copy, word by word, including "related databas" without "E")

--\BURGOS

According to this reference, this answer looks correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms151198.aspx>

QUESTION 24

You administer two Microsoft SQL Server 2012 servers named ServerA and ServerB. You use a database named AdventureWorks. You need to prepare the AdventureWorks database for database mirroring. ServerB will act as the mirror in a mirroring partnership along with ServerA. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> <div></div> </div>	<div>Backup AdventureWorks on ServerA by using a full backup.</div> <div>Backup AdventureWorks on ServerA by using a full backup followed by a transaction log backup by using the NORECOVERY option.</div> <div>Backup AdventureWorks on ServerA by using a transaction log backup. Restore the transaction log backup by using the RECOVERY option on ServerB.</div> <div>Backup AdventureWorks on ServerA by using a transaction log backup. Restore the transaction log backup by using the NORECOVERY option on ServerB.</div> <div>Restore the full database backup of AdventureWorks by using the NORECOVERY option on ServerB as AdventureWorks.</div> <div>Restore the full database backup of AdventureWorks by using the NORECOVERY option on ServerB as AdventureWorks_Mirror.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Backup AdventureWorks on ServerA by using a full backup.

Restore the full database backup of AdventureWorks by using the NORECOVERY option on ServerB as AdventureWorks.

Backup AdventureWorks on ServerA by using a transaction log backup. Restore the transaction log backup by using the NORECOVERY option on ServerB.

Section: (none)

Explanation

Explanation/Reference:

--Burgos. YES (Contoso Version)

This question appears 2 times in original VCE file, one with AdventureWorks, another with Contoso. I discharge the Contoso version.

--\Burgos.

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms190941.aspx>

<http://msdn.microsoft.com/en-us/library/ms189852.aspx>

QUESTION 25

You administer a Microsoft SQL Server 2012 instance that contains a database of confidential data. You need to encrypt the database files at the page level. You also need to encrypt the transaction log files. Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲</div><div>▼</div></div>	<div>Create a master key.</div> <div>Create a certificate in the user database protected by the master key.</div> <div>Create a certificate in the master database protected by the master key.</div> <div>Create a database encryption key in the user database and protect it by a password.</div> <div>Create a database encryption key in the master database and protect it by a password.</div> <div>Create a database encryption key in the user database and protect it by the certificate.</div> <div>Create a database encryption key in the master database and protect it by the certificate.</div> <div>Set the database option to enable encryption.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Create a master key.

Create a certificate in the master database protected by the master key.

Create a database encryption key in the user database and protect it by the certificate.

Set the database option to enable encryption.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

--\BURGOS

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/bb510663.aspx>

<http://msdn.microsoft.com/en-us/library/bb934049.aspx>

QUESTION 26

You administer a Microsoft SQL Server 2012 database. The database uses SQL Server Agent jobs to perform regular FULL and LOG backups. The database uses the FULL recovery model. You plan to perform a bulk import of a very large text file. You need to ensure that the following requirements are met during the bulk operation:

- The database transaction log is minimally affected.
- The database is online and all user transactions are recoverable.
- All transactions are fully recoverable prior to the import.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div style="display: flex; align-items: center;"> <div style="margin-right: 10px;"> <input type="button" value="Up"/> <input type="button" value="Down"/> </div> <div style="border: 1px solid black; width: 280px; height: 120px;"></div> </div>	<div style="border: 1px solid black; padding: 5px;"> <p>Execute the BCP tool.</p> <p>Perform a FULL database backup.</p> <p>Perform a database LOG backup.</p> <p>Configure the database to use the FULL recovery model.</p> <p>Configure the database to use the BULK-LOGGED recovery model.</p> </div>
<div style="display: flex; justify-content: space-around; margin-top: 10px;"> <input type="button" value=" << Move"/> <input type="button" value=" Remove >>"/> </div>	

Correct Answer:

Perform a database LOG backup.

Configure the database to use the BULK-LOGGED recovery model.

Execute the BCP tool.

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES
--\BURGOS

According to these references, this answer looks correct.

Reference: <http://technet.microsoft.com/en-us/library/ms187048.aspx>

Reference: <http://technet.microsoft.com/en-us/library/ms189272.aspx>

QUESTION 27

HOTSPOT

You administer a Microsoft SQL Server 2012 database. The database contains a table that has the following definition:

```
CREATE TABLE [Sales].[Customer] (
    [CustomerID] int NOT NULL,
    [CustomerName] nvarchar(50) NOT NULL,
    [TerritoryID] int NULL,
    [LastContactDate] datetimeoffset NULL,
    [CustomerType] nchar(1) NOT NULL,
    [Notes] varchar(250) NULL
)
```

You want to export data from the table to a flat file by using the SQL Server Import and Export Wizard. You need to ensure that the following requirements are met:

- The first row of the file contains the first row of data.
- Each record is of the same length.
- The date follows the U.S. date format.
- The file supports international characters.

What should you do? (To answer, simply select the option or options in the answer area that you would configure.)

Hot Area:

The screenshot shows the 'SQL Server Import and Export Wizard' window, specifically the 'Choose a Destination' step. The title bar reads 'SQL Server Import and Export Wizard'. Below the title bar, the text 'Choose a Destination' is displayed, followed by the instruction 'Specify where to copy data to.'.

The 'Destination' dropdown menu is set to 'Flat File Destination'. Below this, the text 'Select a file and specify the file properties and the file format.' is shown.

The 'File name' field contains 'C:\Employee.csv', with a 'Browse...' button to its right. The 'Locale' dropdown menu is set to 'English (United States)', and the 'Unicode' checkbox is unchecked. The 'Code page' dropdown menu is set to '1252 (ANSI - Latin I)'. The 'Format' dropdown menu is empty. The 'Text qualifier' field contains '<none>'. The 'Column names in the first data row' checkbox is unchecked.

At the bottom of the window, there are five buttons: 'Help', '< Back', 'Next >', 'Finish >>', and 'Cancel'.

Correct Answer:

SQL Server Import and Export Wizard

Choose a Destination
Specify where to copy data to.

Destination: Flat File Destination

Select a file and specify the file properties and the file format.

File name: C:\Employee.csv Browse...

Locale: English (United States) Unicode

Code page: 1252 (ANSI - Latin I)

Format:

Text qualifier: <none>

☐ Column names in the first data row

Help < Back Next > Finish >> Cancel

Section: (none)

Explanation

Explanation/Reference:

-- BURGOS - YES

I changed retiring a check box of "Column names in the first data row" because "The first row of the file contains the first row of data."

--\BURGOS

Verified answer as correct.

References:

<http://msdn.microsoft.com/en-us/library/ms178804.aspx>

<http://msdn.microsoft.com/en-us/library/ms187828.aspx>

QUESTION 28

You administer a Microsoft SQL Server 2012 database. You use an OrderDetail table that has the following definition:

```
CREATE TABLE [dbo].[OrderDetail]
([SalesOrderID] [int] NOT NULL,
[SalesOrderDetailID] [int] IDENTITY(1,1) NOT NULL,
[CarrierTrackingNumber] [nvarchar](25) NULL,
[OrderQty] [smallint] NOT NULL,
[ProductID] [int] NOT NULL,
[SpecialOfferID] [int] NULL,
[UnitPrice] [money] NOT NULL);
```

You need to create a non-clustered index on the SalesOrderID column in the OrderDetail table to include only rows that contain a value in the SpecialOfferID column. Which four Transact-SQL statements should you use? (To answer, move the appropriate statements from the list of statements to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div style="border: 1px solid black; height: 150px; width: 100%;"></div>	<div style="border: 1px solid black; padding: 5px;"> WHERE FILTER ON SpecialOfferID IS NOT NULL; ON dbo.OrderDetail(SalesOrderID) ON dbo.OrderDetail(SalesOrderID) AS FILTERED_INDEX CREATE NONCLUSTERED INDEX FIdx_SpecialOfferID CREATE NONCLUSTERED FILTERED INDEX FIdx_SpecialOfferID </div>
<div style="display: flex; justify-content: space-around; align-items: center;"> <div style="border: 1px solid black; padding: 2px 10px;"><< Move</div> <div style="border: 1px solid black; padding: 2px 10px;">Remove >></div> </div>	

Correct Answer:

```
CREATE NONCLUSTERED INDEX
FIdx_SpecialOfferID
ON dbo.OrderDetail(SalesOrderID)
WHERE
SpecialOfferID IS NOT NULL;
```

Section: (none)

Explanation

Explanation/Reference:

--Burgos. YES

This question appears 2 times in original VCE file, one with SpecialOfferID column, another with CarrierTrackingNumber column. I discharge the CarrierTrackingNumber version.

--\Burgos.

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188783.aspx>
<http://msdn.microsoft.com/en-us/library/ms189280.aspx>