

MICROSOFT 70-458 EXAM QUESTIONS & ANSWERS

Number: 70-458
Passing Score: 700
Time Limit: 180 min
File Version: 56.4



MICROSOFT 70-458 EXAM QUESTIONS & ANSWERS

Exam Name: Transition Your MCTS on SQL Server 2008 to MCSA: SQL Server 2012, Part 2

Test-Papaers

QUESTION 1

You administer several Microsoft SQL Server 2012 servers. Your company has a number of offices across the world connected by using a wide area network (WAN). Connections between offices vary significantly in both bandwidth and reliability. You need to identify the correct replication method for each scenario. What should you do? (To answer, drag the appropriate replication method or methods to the correct location or locations in the answer area. Each replication method may be used once, more than once, or not at all.)

Replication Method	Scenario
Transactional Replication	Multiple databases on the same low-latency subnet must allow applications to write changes locally, and these changes must be replicated to all related databa
Peer-to-Peer Replication	An order summary table is repopulated once a week. This table must be replicated to all databases.
Snapshot Replication	Field offices using unreliable connections keep a local copy of the product catalog and process orders locally. These orders must be periodically replicated to all other
Merge Replication	Information in an order-tracking database must be replicated across a low-latency connection as changes occur to multiple reporting databases.

A.

Correct Answer:

Section: (none)

Explanation

Explanation/Reference:

Replication Method	Scenario
	Multiple databases on the same low-latency subnet must allow applications to write changes locally, and these changes must be replicated to all related databas
	An order summary table is repopulated once a week. This table must be replicated to all databases.
	Field offices using unreliable connections keep a local copy of the product catalog and process orders locally. These orders must be periodically replicated to all other
	Information in an order-tracking database must be replicated across a low-latency connection as changes occur to multiple reporting databases.

QUESTION 2

DRAG DROP

You are a database administrator of a Microsoft SQL Server 2012 environment. The environment contains two servers named SQLServer01 and SQLServer02. The database Contoso exists on SQLServer01. You plan to mirror the Contoso database between SQLServer01 and SQLServer02 by using database mirroring. You need to prepare the Contoso database for database mirroring. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Back up Contoso on SQLServer01 by using a full backup.	
Back up Contoso on SQLServer01 by using a full backup followed by a transaction log backup by using the NORECOVERY option.	
Back up Contoso on SQLServer01 by using a transaction log backup. Restore the transaction log backup by using the RECOVERY option on SQLServer02.	
Back up Contoso on SQLServer01 by using a transaction log backup. Restore the transaction log backup by using the NORECOVERY option on SQLServer02.	
Restore the full database backup of Contoso by using the NORECOVERY option on SQLServer02 as Contoso.	
Restore the full database backup of Contoso by using the RECOVERY option on SQLServer02 as Contoso_Mirror.	

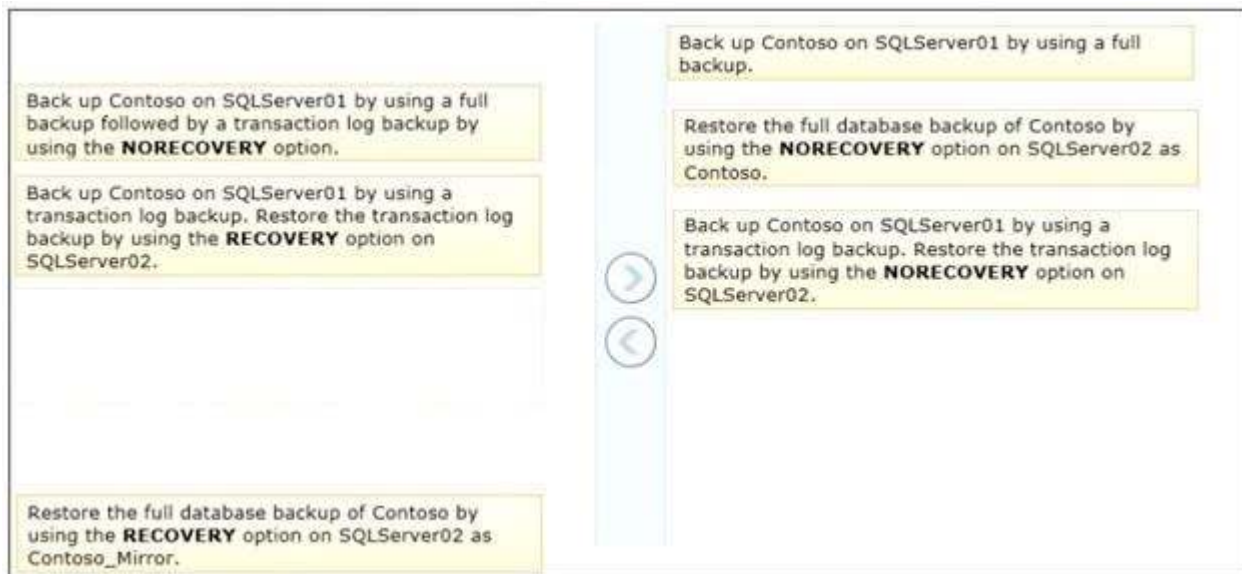
A.

Correct Answer:

Section: (none)

Explanation

Explanation/Reference:



QUESTION 3

You create an availability group that has replicas named HA/Server01 and HA/Server02. Currently, HA/Server01 is the primary replica. You have multiple queries that read data and produce reports from the database. You need to offload the reporting workload to the secondary replica when HA/Server01 is the primary replica. What should you do?

- A. Set the Availability Mode property of HA/Server02 to Asynchronous commit.
- B. Set the Readable Secondary property of HA/Server02 to Read-intent only.
- C. Set the Connections in Primary Role property of HA/Server01 to Allow read/write connections.
- D. Set the Availability Mode property of HA/Server01 to Asynchronous commit.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 4

You administer several Microsoft SQL Server 2012 database servers. Merge replication has been configured for an application that is distributed across offices throughout a wide area network (WAN). Many of the tables involved in replication use the XML and varchar (max) data types. Occasionally, merge replication fails due to timeout errors. You need to reduce the occurrence of these timeout errors. What should you do?

- A. Set the Merge agent on the problem subscribers to use the slow link agent profile.
- B. Create a snapshot publication, and reconfigure the problem subscribers to use the snapshot publication.
- C. Change the Merge agent on the problem subscribers to run continuously.
- D. Set the Remote Connection Timeout on the Publisher to 0.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 5

You administer two Microsoft SQL Server 2012 servers. Each server resides in a different, untrusted domain. You plan to configure database mirroring. You need to be able to create database mirroring endpoints on both servers. What should you do?

- A. Configure the SQL Server service account to use Network Service.
- B. Use a server certificate.
- C. Use a database certificate.
- D. Configure the SQL Server service account to use Local System.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 6

You are migrating a database named Orders to a new server that runs Microsoft SQL Server 2012. You attempt to add the [Corpnet\User1] login to the database. However, you receive the following error message: "User already exists in current database."

You need to configure the [Corpnet\User1] login to be able to access the Orders database and retain the original permissions. You need to achieve this goal by using the minimum required permissions. Which Transact-SQL statement should you use?

- A. DROP USER [User1];
CREATE USER [Corpnet\User1] FOR LOGIN [Corpnet\User1]; ALTER ROLE [db_owner] ADD MEMBER [Corpnet\User1];
- B. ALTER SERVER ROLS sysadmin ADD MEMBER [Corpnet\User1];
- C. ALTER USER [Corpnet\User1] WITH LOGIN [Corpnet\User1];
- D. ALTER ROLE [db_owner] ADD MEMBER [Corpnet\User1];

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 7

You use a contained database named ContosoDb within a domain. You need to create a user who can log on to the ContosoDb database. You also need to ensure that you can port the database to different database servers within the domain without additional user account configurations. Which type of user should you create?

- A. SQL user without login
- B. SQL user with a custom SID
- C. SQL user with login
- D. Domain user

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 8

You are the lead database administrator (DBA) of a Microsoft SQL Server 2012 environment. All DBAs are members of the DOMAIN\JrDBAs Active Directory group. You grant DOMAIN\JrDBAs access to the SQL Server. You need to create a server role named SpecialDBARole that can perform the following functions:

- View all databases.
- View the server state.
- Assign GRANT, DENY, and REVOKE permissions on logins. You need to add DOMAIN\JrDBAs to the server role. You also need to provide the least level of privileges necessary. Which SQL statement or statements should you use? Choose all that apply.

- A. CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION setupadmin;
- B. ALTER SERVER ROLE [SpecialDBARole] ADD MEMBER [DOMAIN\JrDBAs];
- C. CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION securityadmin;
- D. GRANT VIEW DEFINITION TO [SpecialDBARole];
- E. CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION serveradmin;
- F. GRANT VIEW SERVER STATE, VIEW ANY DATABASE TO [SpecialDBARole];

Correct Answer: CEF

Section: (none)

Explanation

Explanation/Reference:

QUESTION 9

You administer a Microsoft SQL Server database named Sales. The database is 3 terabytes in size. The Sales database is configured as shown in the following table.

Filegroup	File
PRIMARY	• Sales.mdf
XACTIONS	• Sales_1.ndf • Sales_2.ndf • Sales_3.ndf
ARCHIVES	• SalesArch_1.ndf • SalesArch_2.ndf

You discover that Sales_2.ndf is corrupt. You need to recover the corrupted data in the minimum amount of time. What should you do?

- A. Perform a file restore.
- B. Perform a transaction log restore.
- C. Perform a restore from a full backup.
- D. Perform a filegroup restore.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 10

You administer a SQL Server 2012 server that contains a database named SalesDb. SalesDb contains a schema named Customers that has a table named Regions. A user named UserA is a member of a role named Sales. UserA is granted the Select permission on the Regions table. The Sales role is granted the Select permission on the Customers schema. You need to ensure that the Sales role, including UserA, is disallowed to select from any of the tables in the Customers schema. Which Transact-SQL statement should you use?

- A. REVOKE SELECT ON Schema::Customers FROM UserA
- B. DENY SELECT ON Object::Regions FROM UserA
- C. EXEC sp_addrolemember 'Sales', 'UserA'
- D. DENY SELECT ON Object::Regions FROM Sales
- E. REVOKE SELECT ON Object::Regions FROM UserA
- F. DENY SELECT ON Schema::Customers FROM Sales
- G. DENY SELECT ON Schema::Customers FROM UserA
- H. EXEC sp_droprolemember 'Sales', 'UserA'
- I. REVOKE SELECT ON Object::Regions FROM Sales
- J. REVOKE SELECT ON Schema::Customers FROM Sales

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

QUESTION 11

You administer a SQL 2012 server that contains a database named SalesDb. SalesDb contains a schema named Customers that has a table named Regions. A user named UserA is a member of a role named Sales. UserA is granted the Select permission on the Regions table. The Sales role is granted the Select permission on the Customers schema. You need to remove the Select permission for UserA on the Regions table. You also need to ensure that UserA can still access all the tables in the Customers schema, including the Regions table, through the Sales role permissions. Which Transact-SQL statement should you use?

- A. DENY SELECT ON Object::Regions FROM UserA
- B. DENY SELECT ON Schema::Customers FROM UserA
- C. EXEC sp_addrolemember 'Sales', 'UserA'
- D. REVOKE SELECT ON Object::Regions FROM UserA
- E. REVOKE SELECT ON Object::Regions FROM Sales
- F. EXEC sp_droprolemember 'Sales', 'UserA'
- G. REVOKE SELECT ON Schema::Customers FROM UserA
- H. DENY SELECT ON Object::Regions FROM Sales
- I. DENY SELECT ON Schema::Customers FROM Sales
- J. REVOKE SELECT ON Schema::Customers FROM Sales

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

Get Complete Collection of 70-458 Exam's Question and Answers.
<http://www.Test4Prep.com>

Test4Prep

QUESTION 12

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN). The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D). · A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours. Five percent of the existing data is modified each day. The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands. Each data load adds 3 GB of data to the database. These data load operations must occur in the minimum amount of time. A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours. On Wednesday at 10:00 hours, the development team requests you to refresh the database on a development server by using the most recent version. You need to perform a full database backup that will be restored on the development server. Which backup option should you use?

- A. NORECOVERY
- B. FULL
- C. NO_CHECKSUM
- D. CHECKSUM
- E. Differential
- F. BULK_LOGGED
- G. STANDBY
- H. RESTART
- I. SKIP
- J. Transaction log
- K. DBO ONLY
- L. COPY_ONLY
- M. SIMPLE
- N. CONTINUE_AFTER_ERROR

Correct Answer: K

Section: (none)

Explanation

Explanation/Reference:

QUESTION 13

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN). The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D). · A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours. Five percent of the existing data is modified each day. The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands. Each data load adds 3 GB of data to the database. These data load operations must occur in the minimum amount of time. A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours. You need to ensure that the backup size is as small as possible. Which backup should you perform every two hours?

- A. BULK_LOGGED
- B. NO_CHECKSUM
- C. FULL
- D. RESTART
- E. CHECKSUM
- F. STANDBY
- G. DBO.ONLY
- H. NORECOVERY
- I. SIMPLE
- J. SKIP
- K. Transaction log
- L. COPY_ONLY
- M. Differential
- N. CONTINUE_AFTER_ERROR

Correct Answer: B

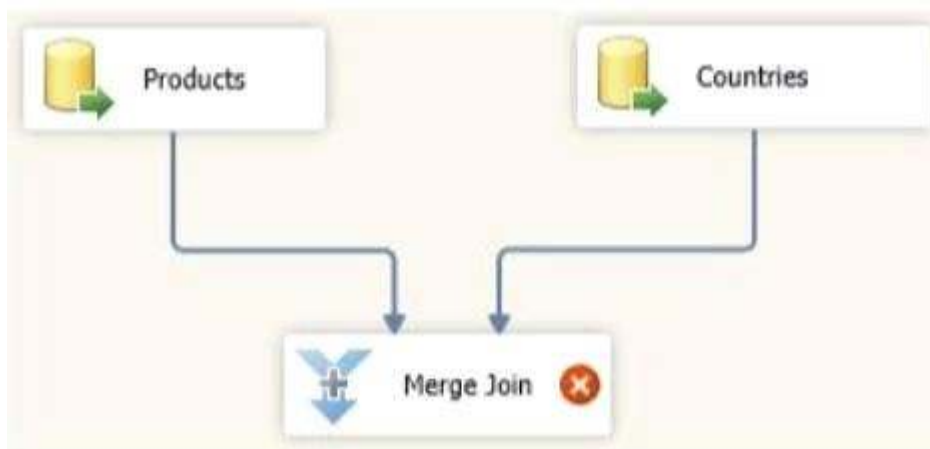
Section: (none)

Explanation

Explanation/Reference:

QUESTION 14

You are developing a data flow transformation to merge two data sources. One source contains product data and the other source contains data about the country in which the product was manufactured. Both data sources contain a two-character CountryCode column and both use SQL Server. Both data sources contain an ORDER BY clause to sort the data by the CountryCode column in ascending order. You use a Merge Join transformation to join the data.



You need to ensure that the Merge Join transformation works correctly without additional transformations. What should you do? (Each answer presents a part of the solution. Choose all that apply.)

- A. set the appropriate SortKeyPosition properties on the data sources.
- B. set the ValidateExternalMetaData property on the Merge Join transformation to True.
- C. set the IsSorted property on both data sources.
- D. Set the MaxBuffersPerInput property on the Merge Join transformation to 2.

Correct Answer: AB

Section: (none)

Explanation

Explanation/Reference:

QUESTION 15

You are creating a SQL Server Integration Services (SSIS) package to retrieve product data from two different sources. One source is hosted in a SQL Azure database. Each source contains products for different distributors. Products for each distributor source must be combined for insertion into a single product table destination. You need to select the appropriate data flow transformation to meet this requirement. Which transformation types should you use? (Each correct answer presents a complete solution. Choose all that apply.)

- A. Multicast
- B. Merge Join
- C. Term Extraction
- D. union All
- E. Merge

Correct Answer: AE

Section: (none)

Explanation

Explanation/Reference:

QUESTION 16

You are designing a SQL Server Integration Services (SSIS) package that uses the Fuzzy Lookup transformation. The reference data to be used in the transformation does not change. You need to reuse the Fuzzy Lookup match index to increase performance and reduce maintenance. What should you do?

- A. Select the GenerateAndPersistNewIndex option in the Fuzzy Lookup Transformation Editor.
- B. Select the GenerateNewIndex option in the Fuzzy Lookup Transformation Editor.
- C. Select the DropExistingMatchIndex option in the Fuzzy Lookup Transformation Editor.
- D. Execute the sp_FuzzyLookupTableMaintenanceUninstall stored procedure.
- E Execute the sp_FuzzyLookupTableMaintenanceInvoke stored procedure.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 17

You are creating a SQL Server Master Data Services (MDS) model for a company. The source data for the company is stored in a single table that contains the manager-to-subordinate relationships. You need to create a hierarchy representing the organizational structure of the company. Which hierarchy type should you use?

- A. Natural
- B. Explicit
- C. Parent
- D. Recursive

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 18

You are completing the installation of the Data Quality Server component of SQL Server Data Quality Services (DQS). You need to complete the post-installation configuration. What should you do?

- A. Run the Configuration component in the Data Quality Client.
- B. Install ADOMD.NET.
- C. Run the Data Quality Server Installer.
- D. Make the data available for DQS operations.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 19

You are the data steward for a Business Intelligence project. You must identify duplicate rows stored in a SQL Server table and output discoveries to a CSV file. A Data Quality Services (DQS) knowledge base has been created to support this project. You need to produce the CSV file with the least amount of development effort. What should you do?

- A. Create an Integration Services package and use a Data Profiling transform.
- B. Create a custom .NET application based on the Knowledgebase class.
C Create a data quality project.
- C. Create a CLR stored procedure based on the Knowledgebase class.
- D. Create a Master Data Services (MDS) business rule.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 20

You are designing a data warehouse with two fact tables. The first table contains sales per month and the second table contains orders per day. Referential integrity must be enforced declaratively. You need to design a solution that can join a single time dimension to both fact tables. What should you do?

- A. Join the two fact tables.
- B. Merge the fact tables.

- C. Create a time dimension that can join to both fact tables at their respective granularity.
- D. Create a surrogate key for the time dimension.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 21

DRAG DROP

You develop a SQL Server Integration Services (SSIS) project by using the Project Deployment model. The project contains many packages. It is deployed on a server named Development1. The project will be deployed to several servers that run SQL Server 2012. The project accepts one required parameter. The data type of the parameter is a string. A SQL Agent job is created that will call the master.dtsx package in the project. A job step is created for the SSIS package. The job must pass the value of an SSIS Environment Variable to the project parameter. The value of the Environment Variable must be configured differently on each server that runs SQL Server. The value of the Environment Variable must provide the server name to the project parameter. You need to configure SSIS on the Development1 server to pass the Environment Variable to the package. Which four actions should you perform in sequence by using SQL Server Management Studio? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

In the SSIS catalog, create an Environment named DevelopmentEnv .	
In the DevelopmentEnv Environment, create an Environment Variable. Name the Environment Variable PackageEnvironment . Set the Environment Variable value to Development1 .	
In the SQL Agent job, choose the DevelopmentEnv Environment in the Configuration tab of the SQL Server Integration Services job step.	
In the project configuration, create an Environment reference to DevelopmentEnv . Assign the value of the PackageEnvironment Environment Variable to the project parameter.	
In the project configuration, edit the value of the project parameter to DevelopmentEnv .	
In the SSIS catalog, create a folder named DevelopmentEnv .	

A.

Correct Answer:

Section: (none)

Explanation

Explanation/Reference:



QUESTION 22

You are editing a SQL Server Integration Services (SSIS) package that contains three Execute SQL tasks and no other tasks. The three Execute SQL tasks modify products in staging tables in preparation for a data warehouse load. The package and all three Execute SQL product tasks have their TransactionOption property set to Supported. You need to ensure that if any of the three Execute SQL product tasks fail, all three tasks will roll back their changes. What should you do?

- A. Change the TransactionOption property of the package to Required.
- B. Change the TransactionOption property of all three Execute SQL product tasks to Required.
- C. Move the three Execute SQL product tasks into a Foreach Loop container.
- D. Move the three Execute SQL product tasks into a Sequence container.

Correct Answer: A

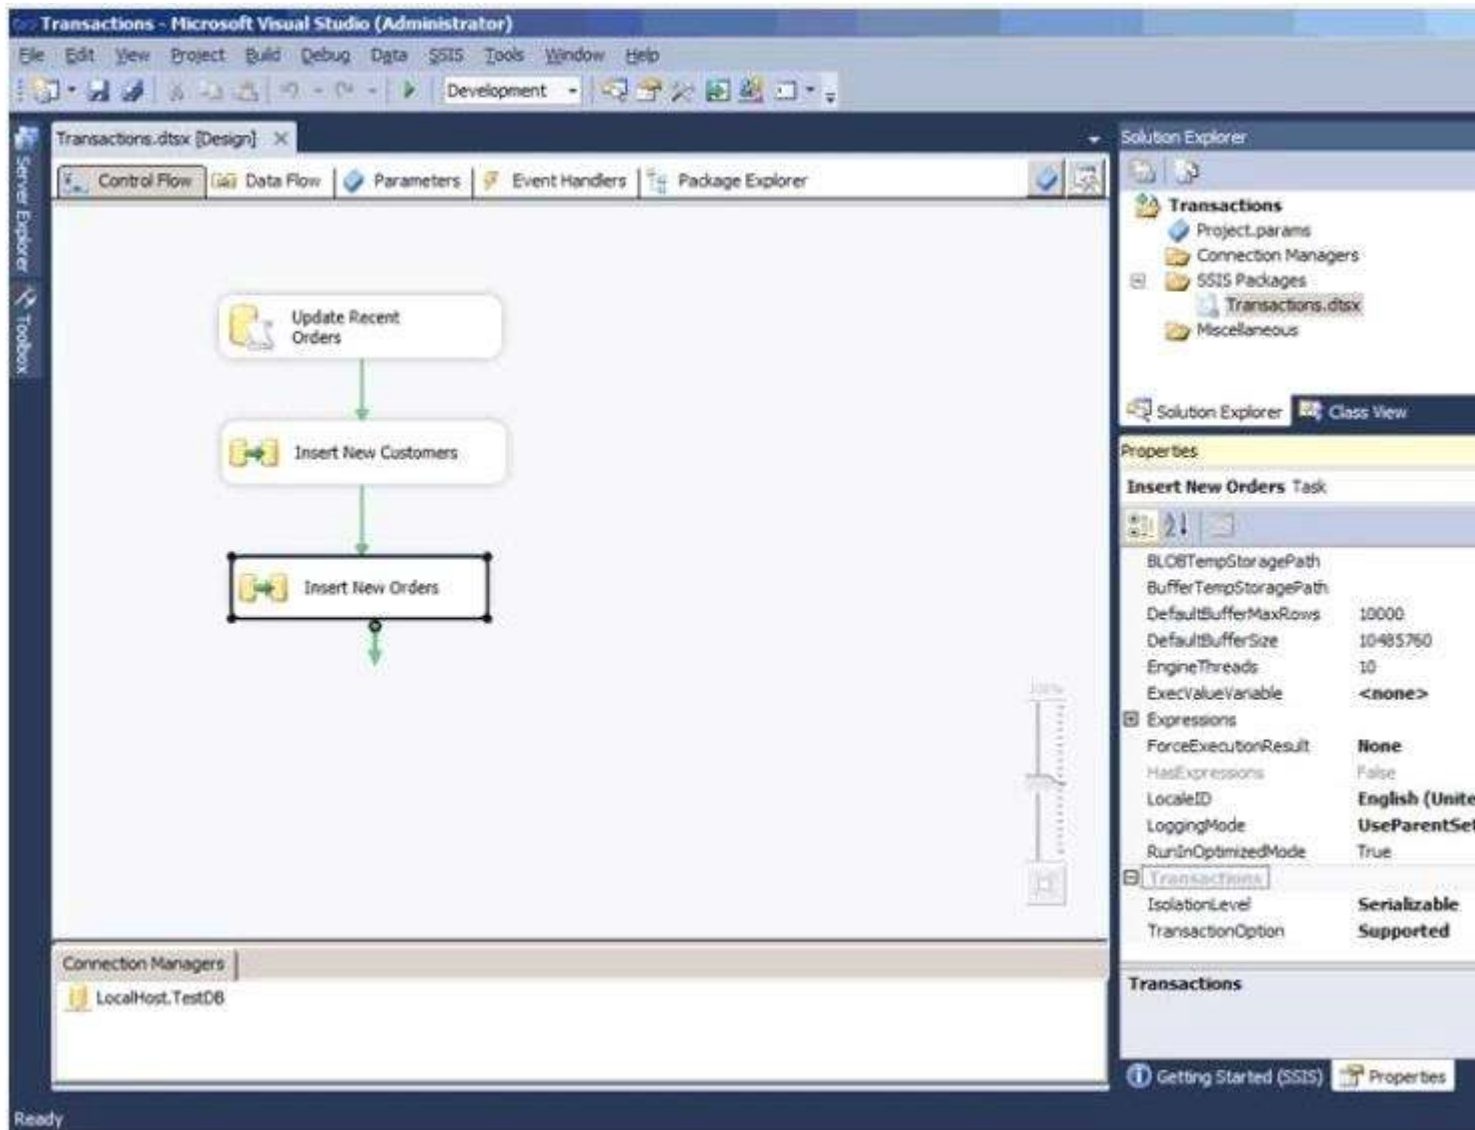
Section: (none)

Explanation

Explanation/Reference:

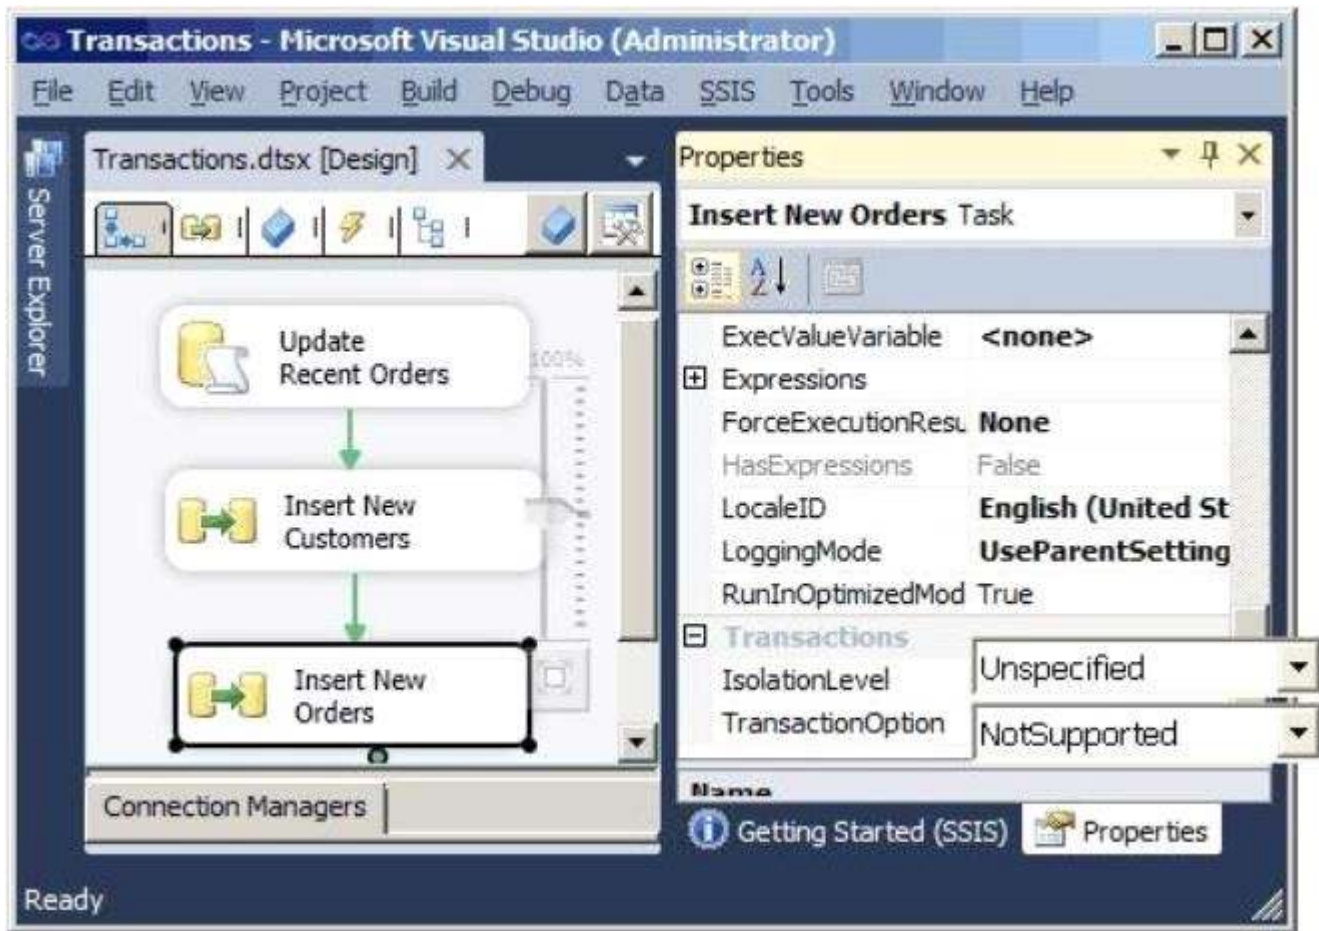
QUESTION 23

You are designing a package control flow. The package moves sales order data from a SQL Azure transactional database to an on-premise reporting database. The package will run several times a day, while new sales orders are being added to the transactional database. The current design of the package control flow is shown in the answer area. (Click the Exhibit button.)



The Insert New Orders Data Flow task must meet the following requirements:

*Usage of the tempdb database should not be impacted. *Concurrency should be maximized, while only reading committed transactions. *If the task fails, only that task needs to be rolled back. You need to configure the Insert New Orders Data Flow task to meet the requirements. How should you configure the transaction properties? To answer, select the appropriate setting or settings in the answer area.



A.

Correct Answer:

Section: (none)

Explanation

Explanation/Reference:

Answer: Select
"IsolationLevel as
Choas"

QUESTION 24

You are developing a SQL Server Integration Services (SSIS) package to load data into a data warehouse.



You need to establish the correct order for loading each table to maximize parallel processing. Which order should you use?

- A. 1. DimCurrency, DimScenario, DimAccount in parallel
2. DimOrganization
3. FactFinance
- B. 1. DimCurrency, DimOrganization in parallel
2. DimScenario, DimAccount in parallel
3. FactFinance
- C. 1. DimCurrency, FactFinance in parallel
2. DimOrganization, DimScenario, DimAccount in parallel
Get Complete Collection of 70-458 Exam's Question and Answers.
- D. 1. FactFinance
2. DimOrganization, DimScenario, DimAccount in parallel
3. DimCurrency
- E. 1. DimCurrency
2. DimOrganization
3. DimScenario, DimAccount in parallel
4. FactFinance

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 25

You are deploying a new SQL Server Integration Services (SSIS) project to the test environment. A package in the project uses a custom task component. You need to ensure that the custom object is deployed on the test environment correctly. What should you do?

- A. Run the package by using the dtexec /rep /conn command.
- B. Create a reusable custom logging component.
- C. Create an OnError event handler.
- D. Use the gacutil command.
E Use the dtutil /copy command.
- E. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.

- F. Run the package by using the dtexec /dumperror /conn command.
- G. Use the Project Deployment Wizard.
- H. Deploy the package by using an msi file.
- I. Add a data tap on the output of a component in the package data flow.
- J. Run the package by using the dtexecui.exe utility and the SQL Log provider.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 26

To ease the debugging of packages, you standardize the SQL Server Integration Services (SSIS) package logging methodology. The methodology has the following requirements:

- Centralized logging in SQL Server
- Simple deployment
- Availability of log information through reports or T-SQL
- Automatic purge of older log entries
- Configurable log details

You need to configure a logging methodology that meets the requirements while minimizing the amount of deployment and development effort. What should you do?

- A. Use the gacutil command.
- B. Use the dtutil /copy command.
- C. Use the Project Deployment Wizard.
- D. Create an OnError event handler.
- E. Create a reusable custom logging component.
- F. Run the package by using the dtexec /rep /conn command.
- G. Run the package by using the dtexec /dumperror /conn command.
- H. Run the package by using the dtexecui.exe utility and the SQL Log provider.
- I. Add a data tap on the output of a component in the package data flow.
- J. Deploy the package by using an msi file.
- K. Deploy the package to the Integration Services catalog by using dtutil and use SQL Server to store the configuration.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 27

You are troubleshooting an existing SQL Server Integration Services (SSIS) package. On several occasions, the package execution does not finish and no data seems to have been transferred. You need to ensure that package logging occurs. Your solution must minimize deployment and development efforts. What should you do?

- A. Add an OnError event handler to the SSIS project.
- B. Use an msi file to deploy the package on the server.
- C. Open a command prompt and run the gacutil command.
- D. Open a command prompt and run the dtutil /copy command.
- E. Open a command prompt and run the dtexec /rep /conn command.

- F. Open a command prompt and run the dtexec /dumperror /conn command.
- G. Open a command prompt and execute the package by using the SQL Log provider and running the dtexecui.exe utility.
- H. Create a reusable custom logging component and use it in the SSIS project.
- I. Configure the SSIS solution to use the Project Deployment Model.
- J. Configure the output of a component in the package data flow to use a data tap.
- K. Run the dtutil command to deploy the package to the SSIS catalog and store the configuration in SQL Server.

Correct Answer: I

Section: (none)

Explanation

Explanation/Reference:

QUESTION 28

You are developing a SQL Server Integration Services (SSIS) package that imports data into a data warehouse. You add an Execute SQL task to the control flow. The task must execute a simple INSERT statement. The task has the following requirements:

- The INSERT statement must use the value of a string package variable. The variable name is StringVar.
- The Execute SQL task must use an OLE DB Connection Manager. In the Parameter Mapping tab of the Execute SQL task, StringVar has been added as the only parameter. You must configure the SQLStatement property of the Execute SQL task. Which SQL statement should you use?

- A. INSERT INTO dbo.Table (variablevalue) VALUES (\$Project::StringVar)
- B. INSERT INTO dbo.Table (variablevalue) VALUES (@StringVar)
- C. INSERT INTO dbo.Table (variablevalue) VALUES (\$Package::StringVar)
- D. INSERT INTO dbo.Table (variablevalue) VALUES (?)

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 29

You are developing a SQL Server Integration Services (SSIS) package to load data into a SQL Server table on Server A. The package includes a data flow and is executed on ServerB. The destination table has its own identity column. The destination data load has the following requirements:

- The identity values from the source table must be used.
- Default constraints on the destination table must be ignored.
- Batch size must be 100,000 rows.

You need to add a destination and configure it to meet the requirements. Which destination should you use?

- A. ADO NET Destination with Bulk Insert
- B. SQL Server Destination
- C. OLE DB Destination with Fast Load
- D. ADO NET Destination without Bulk Insert
- E. OLE DB Destination without Fast Load

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 30

You are designing a SQL Server Integration Services (SSIS) data flow to load sales transactions from a source system into a data warehouse hosted on SQL Azure. One of the columns in the data source is named ProductCode. Some of the data to be loaded will reference products that need special processing logic in the data flow. You need to enable separate processing streams for a subset of rows based on the source product code. Which data flow transformation should you use?

- A. Source Assistant
- B. Audit
- C. Destination Assistant
- D. Script Component

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference: