

Microsoft.Certkiller.70-462.v2014-06-18.by.LARA.169q

Number: 70-462
Passing Score: 700
Time Limit: 120 min
File Version: 24.5

Exam Code: 70-462

Exam Name: Administering Microsoft SQL Server 2012 Databases



CERTKILLER
slay your exams

Exam A

QUESTION 1

DRAG DROP

You administer three Microsoft SQL Server 2012 servers named ServerA, ServerB, and ServerC.

ServerA is the acting principal and ServerB is the mirror.

You need to add ServerC as a witness to the existing mirroring session between ServerA and ServerB.

You need to achieve this goal without delaying synchronization.

Which **three** actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲</div><div>▼</div></div> <div></div>	<div>On Server C, Create an Endpoint for use by the witness.</div> <div>Ensure that the same Windows Login exists on each server and grant Connect permissions to each servers endpoint.</div> <div>On Server A, alter the principal database to use the endpoint on server C as the witness.</div> <div>On Server A, pause the mirroring session between Server A and Server B</div> <div>On Server B, alter the principle database to use the endpoint on Server C as the witness</div> <div>Ensure that the same Proxy exists on each server and grant Connect permissions to each server's endpoint.</div> <div>On Server A, resume the mirroring session between Server A and Server B.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

On Server C, Create an Endpoint for use by the witness.

Ensure that the same Windows Login exists on each server and grant Connect permissions to each servers endpoint.

On Server A, alter the principal database to use the endpoint on server C as the witness.

Section: (none)

Explanation

Explanation/Reference:

QUESTION 2

DRAG AND DROP

You administer several Microsoft SQL Server 2012 servers.

Your company has a number of offices across the world connected by using a wide area network (WAN).

Connections between offices vary significantly in both bandwidth and reliability.

You need to identify the correct replication method for each scenario.

What should you do? (To answer, drag the appropriate replication method or methods to the correct location or locations in the answer area. Each replication method may be used once, more than once, or not at all.)

Select and Place:

Replication Method	Scenario
Transactional Replication	Multiple databases on the same low-latency subnet must allow applications to write changes locally, and these changes must be replicated to all related databases.
Peer-to-Peer Replication	An order summary table is repopulated once a week. This table must be replicated to all databases.
Snapshot Replication	Field offices using unreliable connections keep a local copy of the product catalog and process orders locally. These orders must be periodically replicated to all other offices.
Merge Replication	Information in an order-tracking database must be replicated across a low-latency connection as changes occur to multiple reporting databases.

Correct Answer:

Replication Method	Scenario
	Multiple databases on the same low-latency subnet must allow applications to write changes locally, and these changes must be replicated to all related databases.
	An order summary table is repopulated once a week. This table must be replicated to all databases.
	Field offices using unreliable connections keep a local copy of the product catalog and process orders locally. These orders must be periodically replicated to all other offices.
	Information in an order-tracking database must be replicated across a low-latency connection as changes occur to multiple reporting databases.

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms151198.aspx>

QUESTION 3

You are a database administrator of a Microsoft SQL Server 2012 environment. The environment contains two servers named SQLServer01 and SQLServer02. The database Contoso exists on SQLServer01. You plan to mirror the Contoso database between SQLServer01 and SQLServer02 by using database mirroring. You need to prepare the Contoso database for database mirroring. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> </div> <div></div>	<div>Backup Contoso on SQLServer01 by using a full backup.</div> <div>Backup Contoso on SQLServer01 by using a full backup followed by a transaction log backup by using the NORECOVERY option.</div> <div>Backup Contoso on SQLServer01 by using a transaction log backup. Restore the transaction log backup by using the RECOVERY option on SQLServer02.</div> <div>Backup Contoso on SQLServer01 by using a transaction log backup. Restore the transaction log backup by using the NORECOVERY option on SQLServer02.</div> <div>Restore the full database backup of Contoso by using the NORECOVERY option on SQLServer02 as Contoso.</div> <div>Restore the full database backup of Contoso by using the NORECOVERY option on SQLServer02 as Contoso_Mirror.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Backup Contoso on SQLServer01 by using a full backup.

Restore the full database backup of Contoso by using the NORECOVERY option on SQLServer02 as Contoso.

Backup Contoso on SQLServer01 by using a transaction log backup. Restore the transaction log backup by using the NORECOVERY option on SQLServer02.

Section: (none)

Explanation

Explanation/Reference:

QUESTION 4

You administer a Microsoft SQL Server 2012 environment that contains a production SQL Server 2005 instance named SQL2005 and a development SQL Server 2012 instance named SQL2012.

The development team develops a new application that uses the SQL Server 2012 functionality.

You are planning to migrate a database from SQL2005 to SQL2012 so that the development team can test their new application.

You need to migrate the database without affecting the production environment.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲▼</div><div></div></div>	<div>Perform a transaction log backup on SQL2005.</div> <div>Perform a full database backup on SQL2005.</div> <div>Perform a VSS backup on the database on SQL2005.</div> <div>Restore the VSS backup on SQL2012.</div> <div>Restore the full database backup on SQL 2012.</div> <div>Restore the database backup and transaction log backup on SQL 2012.</div> <div>Change the compatibility level for the database to 120 on SQL2012.</div> <div>Change the compatibility level for the database to 110 on SQL2012.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Perform a full database backup on SQL2005.

Restore the full database backup on SQL 2012.

Change the compatibility level for the database to 110 on SQL2012.

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms177429.aspx>

QUESTION 5

You administer a Microsoft SQL Server 2012 database. You use an OrderDetail table that has the following definition:

```
CREATE TABLE [dbo].[OrderDetail]
([SalesOrderID] [int] NOT NULL,
[SalesOrderDetailID] [int] IDENTITY(1,1) NOT NULL,
[CarrierTrackingNumber] [nvarchar](25) NULL,
[OrderQty] [smallint] NOT NULL,
[ProductID] [int] NOT NULL,
[SpecialOfferID] [int] NULL,
[UnitPrice] [money] NOT NULL);
```

You need to create a non-clustered index on the SalesOrderID column in the OrderDetail table to include only rows that contain a value in the CarrierTrackingNumber column. Which four Transact-SQL statements should you use? (To answer, move the appropriate statements from the list of statements to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div style="display: flex; align-items: center;"> <div style="margin-right: 10px;"> <input type="button" value="▲"/> <input type="button" value="▼"/> </div> <div style="border: 1px solid black; width: 280px; height: 180px;"></div> </div>	<div style="border: 1px solid black; padding: 5px;"> WHERE FILTER ON CarrierTrackingNumber IS NOT NULL; ON dbo.OrderDetail(SalesOrderID) ON dbo.OrderDetail(SalesOrderID) AS FILTERED_INDEX CREATE NONCLUSTERED INDEX FIdx_CarrierTrackingNumber CREATE NONCLUSTERED FILTERED INDEX FIdx_CarrierTrackingNumber </div>
<div style="display: flex; justify-content: center; gap: 20px;"> <div style="border: 1px solid black; padding: 2px 10px;"><< Move</div> <div style="border: 1px solid black; padding: 2px 10px;">Remove >></div> </div>	

Correct Answer:

```
CREATE NONCLUSTERED INDEX
FIdx_CarrierTrackingNumber
ON dbo.OrderDetail(SalesOrderID)
WHERE
CarrierTrackingNumber IS NOT NULL;
```

Section: (none)

Explanation

Explanation/Reference:

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188783.aspx>

<http://msdn.microsoft.com/en-us/library/ms189280.aspx>

QUESTION 6

You administer a Microsoft SQL Server 2012 database. You use an OrderDetail table that has the following definition:

```
CREATE TABLE [dbo].[OrderDetail]
([SalesOrderID] [int] NOT NULL,
[SalesOrderDetailID] [int] IDENTITY(1,1) NOT NULL,
[CarrierTrackingNumber] [nvarchar](25) NULL,
[OrderQty] [smallint] NOT NULL,
[ProductID] [int] NOT NULL,
[SpecialOfferID] [int] NULL,
[UnitPrice] [money] NOT NULL);
```

You need to create a non-clustered index on the SalesOrderID column in the OrderDetail table to include only rows that contain a value in the SpecialOfferID column. Which four Transact-SQL statements should you use? (To answer, move the appropriate statements from the list of statements to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div style="border: 1px solid black; height: 150px; width: 100%;"></div>	<div style="border: 1px solid black; padding: 5px;"> WHERE FILTER ON SpecialOfferID IS NOT NULL; ON dbo.OrderDetail(SalesOrderID) ON dbo.OrderDetail(SalesOrderID) AS FILTERED_INDEX CREATE NONCLUSTERED INDEX FIdx_SpecialOfferID CREATE NONCLUSTERED FILTERED INDEX FIdx_SpecialOfferID </div>
<div style="display: flex; justify-content: space-around; align-items: center;"> << Move Remove >> </div>	

Correct Answer:

```
CREATE NONCLUSTERED INDEX
FIdx_SpecialOfferID
ON dbo.OrderDetail(SalesOrderID)
WHERE
SpecialOfferID IS NOT NULL;
```

Section: (none)

Explanation

Explanation/Reference:

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188783.aspx>

<http://msdn.microsoft.com/en-us/library/ms189280.aspx>

QUESTION 7

You administer a Microsoft SQL Server 2012 database. All database traffic to the SQL Server must be encrypted by using secure socket layer (SSL) certificates or the connection must be refused. Network

administrators have deployed server certificates to the Windows store of all Windows servers on the network from a trusted Certificate Authority. This is the only Certificate Authority allowed to distribute certificates on the network.

You enable the Force Encryption flag for the MSSQLServer protocols, but client computers are unable to connect. They receive the following error message:

"A connection was successfully established with the server, but then an error occurred during the pre-login handshake, (provider: SSL Provider, error: 0 - The certificate chain was issued by an authority that is not trusted.) (Microsoft SQL Server)"

You notice the following entry in the SQL Server log:

"A self-generated certificate was successfully loaded for encryption."

You need to configure SQL Server to encrypt all client traffic across the network. You also need to ensure that client computers are able to connect to the server by using a trusted certificate. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲▼</div><div></div></div>	<div>Restart the SQL Server.</div> <div>Leave the certificate blank in the drop-down list on the CERTIFICATES tab.</div> <div>Choose the new root-level certificate from the drop-down list on the CERTIFICATES tab.</div> <div>Install Certificate Services on the SQL Server, and create a new root-level certificate.</div> <div>From the SQL Configuration Manager on the SQL Server, open the PROTOCOLS properties for the SQL instance.</div> <div>Choose the server certificate provided by the network administrators from the drop-down list on the CERTIFICATES tab.</div> <div>From the SQL Configuration Manager on every client computer that will be connecting to SQL Server, open the PROTOCOLS properties for the SQL instance.</div>

<< Move

Remove >>

Correct Answer:

From the SQL Configuration Manager on the SQL Server, open the PROTOCOLS properties for the SQL instance.

Choose the server certificate provided by the network administrators from the drop-down list on the CERTIFICATES tab.

Restart the SQL Server.

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://thesqldude.com/2012/04/21/setting-up-ssl-encryption-for-sql-server-using-certificates-issues-tips-tricks/>

QUESTION 8

You administer a Microsoft SQL Server 2012 server that has multiple databases.

You need to ensure that users are **unable** to create stored procedures that begin with sp_.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲</div><div>▼</div></div>	<div>Enable StoredProcNamingPolicy.</div> <div>Evaluate StoredProcNamingPolicy.</div> <div>Create a Database Audit named StoredProcNamingConvention. Set the Filter to '@Name LIKE 'sp[_]%'.</div> <div>Create a Policy named StoredProcNamingPolicy. Set the Check condition to StoredProcNamingConvention and Evaluation Mode to On Demand.</div> <div>Create a Policy named StoredProcNamingPolicy. Set the Check condition to StoredProcNamingConvention and Evaluation Mode to On Change: Prevent.</div> <div>Create a Condition named StoredProcNamingConvention by using the Stored Procedure facet that has a single expression. Set the Field to @Name, Operator to NOT LIKE, and Value to 'sp[_]%'.</div> <div>Create a Condition named StoredProcNamingConvention by using the Stored Procedure facet that has a single expression. Set the Field to @Name, Operator to LIKE, and Value to 'sp[_]%'.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Create a Condition named StoredProcNamingConvention by using the Stored Procedure facet that has a single expression.
Set the Field to @Name, Operator to LIKE, and Value to 'sp[_]%'.
Create a Policy named StoredProcNamingPolicy.
Set the Check condition to StoredProcNamingConvention and Evaluation Mode to On Change: Prevent.
Enable StoredProcNamingPolicy.

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/bb510667.aspx>

Policies are created and managed by using Management Studio. The process includes the following steps:

1. Select a Policy-Based Management facet that contains the properties to be configured.
2. Define a condition that specifies the state of a management facet.
3. Define a policy that contains the condition, additional conditions that filter the target sets, and the evaluation mode.
4. Check whether an instance of SQL Server is in compliance with the policy.

Evaluation modes

There are four evaluation modes, three of which can be automated:

- On demand. This mode evaluates the policy when directly specified by the user.
- On change: prevent. This automated mode uses DDL triggers to prevent policy violations.

Important

If the nested triggers server configuration option is disabled, On change: prevent will not work correctly. Policy-Based Management relies on DDL triggers to detect and roll back DDL operations that do not comply with policies that use this evaluation mode. Removing the Policy-Based Management DDL triggers or disabling nest triggers, will cause this evaluation mode to fail or perform unexpectedly.

- On change: log only. This automated mode uses event notification to evaluate a policy when a relevant change is made.
- On schedule. This automated mode uses a SQL Server Agent job to periodically evaluate a policy.

QUESTION 9

You administer a Microsoft SQL Server database.

You want to import data from a text file to the database.

You need to ensure that the following requirements are met:

- Data import is performed by using a stored procedure.
- Data is loaded as a unit and is minimally logged.

Which data import command and recovery model should you choose? (To answer, drag the appropriate data import command or recovery model to the appropriate location or locations in the answer area. Each data import command or recovery model may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Command/Recovery Model Name	Command/Recovery Model
BCP	Data import command
BULK INSERT	Recovery model
Bulk-logged	
OPENDATASOURCE	
Full	

Correct Answer:

Command/Recovery Model Name	Command/Recovery Model
BCP	Data import command
BULK INSERT	Recovery model
Bulk-logged	
OPENDATASOURCE	
Full	

Section: (none)

Explanation

Explanation/Reference:

QUESTION 10

You administer a Microsoft SQL Server 2012 database.

The database is backed up according to the following schedule:

- Daily full backup at 23:00 hours.
- Differential backups on the hour, except at 23:00 hours.
- Log backups every 10 minutes from the hour, except on the hour.

The database uses the Full recovery model.

A developer accidentally drops a number of tables and stored procedures from the database between 22:40 hours and 23:10 hours.

You perform a database restore at 23:30 hours to recover the dropped table.

You need to restore the database by using the minimum amount of administrative effort.

You also need to ensure minimal data loss.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲</div><div>▼</div></div>	<div>Restore the most recent full backup.</div> <div>Restore the full backup taken the previous night.</div> <div>Restore the differential backup taken at 22:00 hours.</div> <div>Restore the transaction log backup taken at 22:40 hours.</div> <div>Restore each transaction log backup taken from 22:00 till 22:40 hours.</div> <div>Restore each transaction log backup taken from the most recent full backup.</div> <div>Restore each differential database backup taken from the previous night's full backup.</div> <div>Restore each transaction log backup taken from the previous night's full backup till 22:40 hours.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Restore the full backup taken the previous night.

Restore the differential backup taken at 22:00 hours.

Restore each transaction log backup taken from 22:00 till 22:40 hours.

Section: (none)

Explanation

Explanation/Reference:

QUESTION 11

You administer a Microsoft SQL Server 2012 instance that contains a database of confidential data.

You need to encrypt the database files at the page level.

You also need to encrypt the transaction log files.

Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> </div> <div></div>	<div>Create a master key.</div> <div>Create a certificate in the user database protected by the master key.</div> <div>Create a certificate in the master database protected by the master key.</div> <div>Create a database encryption key in the user database and protect it by a password.</div> <div>Create a database encryption key in the master database and protect it by a password.</div> <div>Create a database encryption key in the user database and protect it by the certificate.</div> <div>Create a database encryption key in the master database and protect it by the certificate.</div> <div>Set the database option to enable encryption.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Create a master key.

Create a certificate in the master database protected by the master key.

Create a database encryption key in the user database and protect it by the certificate.

Set the database option to enable encryption.

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/bb510663.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/bb934049.aspx>

QUESTION 12

You administer a Microsoft SQL Server database.

Service accounts for SQL Agent are configured to use a local user.

A Microsoft SQL Server Integration Services (SSIS) job step has been created within a SQL Server Agent job.

The SSIS package accesses a network share when exporting data from a SQL Server database.

When you execute the SQL Server Agent job, it fails due to a permissions failure on a share on a remote server.

You need to ensure that the SQL Server Agent job can execute the SSIS package.

Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲▼</div><div></div></div>	<div>Add a proxy that references the local user.</div> <div>Add a proxy that references the credential.</div> <div>Create a local user account and grant local administrator on the SQL Server instance.</div> <div>Create a credential that references the local user.</div> <div>Create a credential that references the domain user.</div> <div>Assign the proxy to the Operating System subsystem.</div> <div>Assign the proxy to the SSIS package execution subsystem.</div> <div>Create a domain user account and grant permissions to the domain user account to access the network share.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Create a domain user account and grant permissions to the domain user account to access the network share.

Create a credential that references the domain user.

Add a proxy that references the credential.

Assign the proxy to the SSIS package execution subsystem.

Section: (none)

Explanation

Explanation/Reference:

QUESTION 13

DRAG AND DROP

You administer a Microsoft SQL Server 2012 server.

A variety of issues occur from time to time in the production environment.

You need to identify the appropriate tool for each issue.

Which tool or tools should you use? (To answer, drag the appropriate tool or tools to the correct issue or

issues in the answer area. Each tool may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Tool	Issue	
DBCC CHECKDB	You want to verify network utilization.	
Performance Monitor	You suspect that a process is being blocked.	
sys.dm_exec_requests DMV	You need to validate the integrity of the database.	
SQL Server error log	A SQL Agent job fails on a specific step, and you need the details of that step.	
Job History	SQL Server will not start.	

Correct Answer:

Tool	Issue	
	You want to verify network utilization.	Performance Monitor
	You suspect that a process is being blocked.	sys.dm_exec_requests DMV
	You need to validate the integrity of the database.	DBCC CHECKDB
	A SQL Agent job fails on a specific step, and you need the details of that step.	Job History
	SQL Server will not start.	SQL Server error log

Section: (none)

Explanation

Explanation/Reference:

QUESTION 14

You administer a single Microsoft SQL Server instance on a two-node failover cluster that has nodes named Node A and Node B.

The instance is currently running on Node A.

You want to patch both Node A and Node B by using the most recent SQL Server Service Pack.

You need to ensure that the following requirements are met:

- Both nodes receive the update.
- Downtime is minimized.
- No data is lost.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲▼</div><div></div></div>	<div>Pause Node A.</div> <div>Pause Node B.</div> <div>Failover from Node A to Node B.</div> <div>Start the SQL Server service on both nodes.</div> <div>Install the service pack on Node A.</div> <div>Install the service pack on Node B.</div> <div>Stop the SQL Server service on both nodes.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Install the service pack on Node B.

Failover from Node A to Node B.

Install the service pack on Node A.

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://technet.microsoft.com/en-us/library/ms191009.aspx>

Reference: <http://technet.microsoft.com/en-us/library/ms191295.aspx>

QUESTION 15

You administer a Microsoft SQL Server 2012 database. The database uses SQL Server Agent jobs to perform regular FULL and LOG backups. The database uses the FULL recovery model. You plan to perform a bulk import of a very large text file. You need to ensure that the following requirements are met during the bulk operation:

- The database transaction log is minimally affected.
- The database is online and all user transactions are recoverable.
- All transactions are fully recoverable prior to the import.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> <div></div> </div>	<div> <div>Execute the BCP tool.</div> <div>Perform a FULL database backup.</div> <div>Perform a database LOG backup.</div> <div>Configure the database to use the FULL recovery model.</div> <div>Configure the database to use the BULK-LOGGED recovery model.</div> </div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Perform a database LOG backup.

Configure the database to use the BULK-LOGGED recovery model.

Execute the BCP tool.

Section: (none)

Explanation

Explanation/Reference:

QUESTION 16

You administer a Microsoft SQL Server 2012 clustered instance that has two nodes named Node 1 and Node 2.

Node 1 fails and the cluster fails over to Node 2.

You need to replace Node 1 and add it to the cluster.

Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> </div> <div></div>	<div>Evict Node 1 from the Windows Failover Cluster.</div> <div>Install Windows on a new server to replace Node 1.</div> <div>Run SQL Server Setup to add Node 1 to the failover cluster.</div> <div>Run Cluster Administrator Setup to add Node 1 to the failover cluster.</div> <div>Add Node 1 to the existing cluster by using SQL Server Configuration Manager.</div> <div>Add Node 1 to the existing cluster by using the Windows Failover Cluster Manager.</div> <div>Register the secondary instance with the Cluster Manager by using SQL Server Management Studio.</div>
<div><< Move</div> <div>Remove >></div>	

Correct Answer:

Evict Node 1 from the Windows Failover Cluster.

Install Windows on a new server to replace Node 1.

Add Node 1 to the existing cluster by using the Windows Failover Cluster Manager.

Run SQL Server Setup to add Node 1 to the failover cluster.

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://technet.microsoft.com/en-us/library/ms181075.aspx>

QUESTION 17

HOTSPOT

You administer a Microsoft SQL Server 2012 database. The database contains a table that has the following definition:

```
CREATE TABLE [Sales].[Customer] (
    [CustomerID] int NOT NULL,
    [CustomerName] nvarchar(50) NOT NULL,
    [TerritoryID] int NULL,
    [LastContactDate] datetimeoffset NULL,
    [CustomerType] nchar(1) NOT NULL,
    [Notes] varchar(250) NULL
)
```

You want to export data from the table to a flat file by using the SQL Server Import and Export Wizard. You need to ensure that the following requirements are met:

- The first row of the file contains the first row of data.
- Each record is of the same length.
- The date follows the U.S. date format.
- The file supports international characters.

What should you do? (To answer, simply select the option or options in the answer area that you would configure.)

Hot Area:



The screenshot shows the 'SQL Server Import and Export Wizard' window, specifically the 'Choose a Destination' step. The window title is 'SQL Server Import and Export Wizard'. The subtitle is 'Choose a Destination' with the instruction 'Specify where to copy data to.' Below this, the 'Destination' dropdown menu is set to 'Flat File Destination'. The instruction 'Select a file and specify the file properties and the file format.' is displayed. The 'File name' field contains 'C:\Employee.csv' with a 'Browse...' button to its right. The 'Locale' dropdown is set to 'English (United States)' and the 'Unicode' checkbox is unchecked. The 'Code page' dropdown is set to '1252 (ANSI - Latin I)'. The 'Format' dropdown is empty. The 'Text qualifier' field contains '<none>'. The 'Column names in the first data row' checkbox is unchecked. At the bottom, there are buttons for 'Help', '< Back', 'Next >', 'Finish >>', and 'Cancel'.

Correct Answer:



Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

References:

<http://msdn.microsoft.com/en-us/library/ms178804.aspx>

<http://msdn.microsoft.com/en-us/library/ms187828.aspx>

QUESTION 18

You create an availability group that has replicas named HA/Server01 and HA/Server02. Currently, HA/Server01 is the primary replica.

You have multiple queries that read data and produce reports from the database.

You need to offload the reporting workload to the secondary replica when HA/Server01 is the primary replica.

What should you do?

- A. Set the Availability Mode property of HA/Server02 to Asynchronous commit.
- B. Set the Readable Secondary property of HA/Server02 to Read-intent only.
- C. Set the Connections in Primary Role property of HA/Server01 to Allow read/write connections.
- D. Set the Availability Mode property of HA/Server01 to Asynchronous commit.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/jj542414.aspx>

QUESTION 19

You administer a Microsoft SQL Server 2012 instance.

You need to stop a blocking process that has an SPID of 64 without stopping other processes.

What should you do?

- A. Execute the following Transact-SQL statement: ALTER SESSION KILL '64'
- B. Execute the following Transact-SQL statement: KILL 64
- C. Restart the SQL Server service.
- D. Execute the following Transact-SQL statement: EXECUTE sp_KillSPID 64

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 20

You administer a Windows Azure SQL Database database named Orders

You need to create a copy of Orders named Orders_Reporting.

Which Transact-SQL command should you use?

- A. BACKUP DATABASE Orders TO DISK = 'D:\Orders.bak'RESTORE DATABASE Orders_Reporting FROM DISK = 'D:\Orders.bak'
- B. BACKUP DATABASE Orders TO DISK = 'D:\Orders.bak'CREATE DATABASE Orders_Reporting FROM DISK = 'D:\Orders.bak'
- C. CREATEDATABASE Orders_ReportingAS COPY OF Orders
- D. BACKUP DATABASE Orders TO DISK = 'D:\Orders.bak'MIRROR TO DISK = 'Orders_Reporting'

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 21

You administer all the deployments of Microsoft SQL Server 2012 in your company.

You need to ensure that an OLTP database that includes up-to-the-minute reporting requirements can be off-loaded from the primary database to another server.

You also need to be able to **add indexes to the secondary database**.

Which configuration should you use?

- A. Two servers configured in the same data center
 - primary server configured to perform log-shipping every 10 minutes
 - backup server configured as a warm standby

- B. Two servers configured on the same subnet
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
- C. Two servers configured in different data centers
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
- D. Two servers configured in the same data center
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
One server configured as an Active Secondary
- E. Two servers configured in different data centers
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
One server configured as an Active Secondary
- F. Two servers configured in a Windows Failover Cluster in the same data center
SQL Server configured as a clustered instance
- G. SQL Server that includes an application database configured to perform transactional replication
- H. SQL Server that includes an application database configured to perform snapshot replication

Correct Answer: G

Section: (none)

Explanation

Explanation/Reference:

QUESTION 22

You administer a Microsoft SQL Server 2012 default instance.

The instance is hosted by a server that has a local firewall configured.

The firewall only allows inbound connections on port 1433.

The server only hosts a single instance of SQL Server.

You need to ensure that the instance is configured to allow remote connections even if the SQL Server is unresponsive to client connections. What should you do? Choose all that apply.

- A. Enable inbound connections on TCP port 1434 in the Windows Firewall on the server.
- B. Execute the following Transact-SQL command:
sp_configure 'remote admin connections',
- C. Execute the Reconfigure command.
- D. Execute the following Transact-SQL command:
sp_configure 'remote access', 1
- E. Restart the SQL Server Agent Service.
- F. Enable inbound connections on TCP port 135 in the Windows Firewall on the server.

Correct Answer: ABC

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms191464.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms190468.aspx>

QUESTION 23

You administer all the deployments of Microsoft SQL Server 2012 in your company.

You have two servers in the same data center that hosts your production database.

You need to ensure that the database remains available if a catastrophic server failure or a disk failure occurs.
You also need to maintain transactional consistency of the data across both servers.
You need to achieve these goals without manual intervention.

Which configuration should you use?

- A. Two servers configured in a Windows Failover Cluster in the same data center
SQL Server configured as a clustered instance
- B. SQL Server that includes an application database configured to perform transactional replication
- C. Two servers configured in the same data center
A primary server configured to perform log-shipping every 10 minutes
A backup server configured as a warm standby
- D. Two servers configured in different data centers
SQL Server Availability Group configured in Synchronous-Commit Availability Mode
One server configured as an Active Secondary
- E. Two servers configured in the same data center
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
One server configured as an Active Secondary
- F. Two servers configured in different data centers
SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
- G. SQL Server that includes an application database configured to perform snapshot replication
- H. Two servers configured on the same subnet
SQL Server Availability Group configured in Synchronous-Commit Availability Mode

Correct Answer: H

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ff877931.aspx>

QUESTION 24

You create an availability group named HaContoso that has replicas named Server01/HA, Server02/HA, and Server03/HA.

Currently, Server01/HA is the primary replica.

You need to ensure that the following requirements are met:

- **Backup operations occur on Server02/HA.**
- **If Server02/HA is unavailable, backup operations occur on Server03/HA.**
- **Backup operations do not occur on Server01/HA.**

How should you configure HaContoso?

- A.
 - Set the backup preference of HaContoso to Prefer Secondary.
 - Set the backup priority of Server02/HA to 20.
 - Set the backup priority of Server03/HA to 10.
- B.
 - Set the backup preference of HaContoso to Secondary only.
 - Set the backup priority of Server02/HA to 20.
 - Set the backup priority of Server03/HA to 10.
- C.
 - Set the backup preference of HaContoso to Secondary only.
 - Set the backup priority of Server02/HA to 10.
 - Set the backup priority of Server03/HA to 20.
- D.
 - Set the exclude replica of Server01/HA to true.
 - Set the backup priority of Server02/HA to 10.
 - Set the backup priority of Server03/HA to 20.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ff877884.aspx>

QUESTION 25

You administer a Microsoft SQL Server 2012 server. When transaction logs grow, SQL Server must send an email message to the database administrators. You need to configure SQL Server to **send the email messages**. What should you configure?

- A. SQL Mail
- B. An Extended Events session
- C. Alerts and operators in SQL Server Agent
- D. Policies under Policy-Based Management

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 26

You administer a Microsoft SQL Server 2012 database.

The database has a table named Customers owned by UserA and another table named Orders owned by UserB.

You also have a stored procedure named GetCustomerOrderInfo owned by UserB.

GetCustomerOrderInfo selects data from both tables.

You create a new user named UserC. You need to ensure that UserC can call the GetCustomerOrderInfo stored procedure.

You also need to assign only the minimum required permissions to UserC.

Which permission or permissions should you assign to UserC? Choose all that apply.

- A. The Select permission on Customers
- B. The Execute permission on GetCustomerOrderInfo
- C. The Take Ownership permission on Customers
- D. The Control permission on GetCustomerOrderInfo
- E. The Take Ownership permission on Orders
- F. The Select permission on Orders

Correct Answer: AB

Section: (none)

Explanation

Explanation/Reference:

The question seems to be missing something. Or the original answer is incorrect. I've changed it to what I believe to be the correct answer. The original answer included "The Select permission on Orders.", but due to ownership chaining, you would only need to give Execute permissions to UserC to access the Orders table

since UserB is the owner.

(BF) - need to test this

Reference:

<http://msdn.microsoft.com/en-us/library/ms188676.aspx>

<http://stackoverflow.com/questions/2212044/sql-server-how-to-permission-schemas>

http://sqlservercentral.com/blogs/steve_jones/2012/03/14/ownership-chains-in-sql-server

QUESTION 27

You administer a Microsoft SQL Server 2012 database that has Trustworthy set to On.

You create a stored procedure that returns database-level information from Dynamic Management Views.

You grant User1 access to execute the stored procedure.

You need to ensure that the stored procedure returns the required information when User1 executes the stored procedure.

You need to achieve this goal by granting the minimum permissions required. What should you do? (Each correct answer presents a complete solution. Choose all that apply.)

- A. Create a SQL Server login that has VIEW SERVER STATE permissions.
Create an application role and a secured password for the role.
- B. Modify the stored procedure to include the EXECUTE AS OWNER statement.
Grant VIEW SERVER STATE permissions to the owner of the stored procedure.
- C. Create a SQL Server login that has VIEW SERVER STATE permissions.
Modify the stored procedure to include the EXECUTE AS {newlogin} statement.
- D. Grant the db_owner role on the database to User1.
- E. Grant the sysadmin role on the database to User1.

Correct Answer: BC

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms187861.aspx>

QUESTION 28

You are migrating a database named Orders to a new server that runs Microsoft SQL Server 2012.

You attempt to add the [Corpnet\User1] login to the database. However, you receive the following error message:

"User already exists in current database."

You need to configure the [Corpnet\User1] login to be able to access the Orders database and retain the original permissions.

You need to achieve this goal by using the minimum required permissions. Which Transact-SQL statement should you use?

- A. DROP USER [User1];
CREATE USER [Corpnet\User1] FOR LOGIN [Corpnet\User1];
ALTER ROLE [db_owner] ADD MEMBER [Corpnet\User1];
- B. ALTER SERVER ROLE [sysadmin] ADD MEMBER [Corpnet\User1];
- C. ALTER USER [Corpnet\User1] WITH LOGIN [Corpnet\User1];
- D. ALTER ROLE [db_owner] ADD MEMBER [Corpnet\User1];

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms176060.aspx>

QUESTION 29

You administer a Microsoft SQL Server 2012 database.

You need to ensure that the size of the transaction log file does not exceed 2 GB.

What should you do?

- A. Execute `sp_configure 'max log size', 2G`.
- B. use the `ALTER DATABASE...SET LOGFILE` command along with the `maxsize` parameter.
- C. In SQL Server Management Studio, right-click the instance and select Database Settings. Set the maximum size of the file for the transaction log.
- D. In SQL Server Management Studio, right-click the database, select Properties, and then click Files. Open the Transaction log Autogrowth window and set the maximum size of the file.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

QUESTION 30

You administer a Microsoft SQL Server 2012 server.

The MSSQLSERVER service uses a domain account named `CONTOSO\SQLService`.

You plan to configure Instant File Initialization.

You need to ensure that Data File **Autogrow** operations use Instant File Initialization.

What should you do? Choose all that apply.

- A. Restart the SQL Server Agent Service.
- B. Disable snapshot isolation.
- C. Restart the SQL Server Service.
- D. Add the `CONTOSO\SQLService` account to the Perform Volume Maintenance Tasks local security policy.
- E. Add the `CONTOSO\SQLService` account to the Server Operators fixed server role.
- F. Enable snapshot isolation.

Correct Answer: CD

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms175935.aspx>

Reference: <http://www.mssqltips.com/sqlservertip/2752/effect-of-instant-file-initialization-within-sql-server/>

How To Enable Instant File Initialization

1. Open Local Security Policy and go to Local Policies --> User Rights Assignment.



2. Double click Perform Volume Maintenance Tasks and add your SQL Server database engine service

account.

3. Restart the SQL Server service using SQL Server Configuration Manager and this setting should now be enabled.

QUESTION 31

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN).

The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours.

Five percent of the existing data is modified each day.

The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands.

Each data load adds 3 GB of data to the database.

These data load operations must occur in the minimum amount of time.

A full database backup is performed every Sunday at 10:00 hours.

Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours.

You need to ensure that your backup will continue if any invalid checksum is encountered.

Which backup option should you use?

- A. STANDBY
- B. Differential
- C. FULL
- D. CHECKSUM
- E. BULK_LOGGED
- F. CONTINUE_AFTER_ERROR
- G. SIMPLE
- H. DBO_ONLY
- I. COPY_ONLY
- J. SKIP
- K. RESTART
- L. Transaction log
- M. NO_CHECKSUM
- N. NORECOVERY

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms186865.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/microsoft.sqlserver.management.smo.backuprestorebase.continueaftererror.aspx>

QUESTION 32

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN).

The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours.

Five percent of the existing data is modified each day.

The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands.

Each data load adds 3 GB of data to the database.

These data load operations must occur in the minimum amount of time.

A full database backup is performed every Sunday at 10:00 hours.

Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours.

On Wednesday at 10:00 hours, the development team requests you to refresh the database on a development server by using the most recent version.

You need to perform a full database backup that will be restored on the development server.

Which backup option should you use?

- A. NORECOVERY
- B. FULL
- C. NO_CHECKSUM
- D. CHECKSUM
- E. Differential
- F. BULK_LOGGED
- G. STANDBY
- H. RESTART
- I. SKIP
- J. Transaction log
- K. DBO ONLY
- L. COPY_ONLY
- M. SIMPLE
- N. CONTINUE AFTER ERROR

Correct Answer: J

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms191495.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms186858.aspx>

QUESTION 33

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN).

The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours.

Five percent of the existing data is modified each day.

The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands.

Each data load adds 3 GB of data to the database.

These data load operations must occur in the minimum amount of time.

A full database backup is performed every Sunday at 10:00 hours.

Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours.

You need to ensure that the minimum **amount of data is lost**.

Which recovery model should the database use?

- A. FULL
- B. DBO_ONLY
- C. CONTINUE_AFTER_ERROR
- D. CHECKSUM
- E. NO_CHECKSUM
- F. SIMPLE
- G. Transaction log
- H. SKIP
- I. RESTART
- J. COPY_ONLY
- K. NORECOVERY
- L. BULK_LOGGED
- M. Differential
- N. STANDBY

Correct Answer: L

Section: (none)

Explanation

Explanation/Reference:

I'd still prefer bulk logged

Reference: <http://msdn.microsoft.com/en-us/library/ms189275.aspx>

QUESTION 34

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN).

The financial database has the following characteristics:

- A data file of 2 terabytes is located on a dedicated LUN (drive D).
- A transaction log of 10 GB is located on a dedicated LUN (drive E).
- Drive D has 1 terabyte of free disk space.
- Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours.

Five percent of the existing data is modified each day.

The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and 15:15 hours by using the BCP or BULK INSERT commands.

Each data load adds 3 GB of data to the database.

These data load operations must occur in the minimum amount of time.

A full database backup is performed every Sunday at 10:00 hours. Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours.

You need to ensure that the backup **size** is as small as possible.

Which backup should you perform every two hours?

- A. BULK_LOGGED
- B. NO_CHECKSUM
- C. FULL
- D. RESTART
- E. CHECKSUM
- F. STANDBY
- G. DBO.ONLY
- H. NORECOVERY
- I. SIMPLE
- J. Transaction log
- K. Differential
- L. CONTINUE_AFTER_ERROR
- M. COPY_ONLY
- N. SKIP

Correct Answer: J

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms186865.aspx>
Reference: <http://msdn.microsoft.com/en-us/library/ms191429.aspx>
Reference: <http://msdn.microsoft.com/en-us/library/ms179478.aspx>

QUESTION 35

You administer a Microsoft SQL Server 2012 database named Contoso on a server named Server01.

You need to **write messages to the Application Log when users are added to or removed from a fixed server role in Server01.**

What should you create?

- A. a Database Audit Specification
- B. a Policy
- C. an Alert
- D. a SQL Profiler Trace
- E. a Resource Pool
- F. an Extended Event session
- G. a Server Audit Specification

Correct Answer: G

Section: (none)

Explanation

Explanation/Reference:

QUESTION 36

You administer a Microsoft SQL Server 2012 database named Contoso on a server named Server01.

You need to be **notified immediately when fatal errors** occur on Server01.

What should you create?

- A. an Alert
- B. a Server Audit Specification
- C. an Extended Event session
- D. a Resource Pool
- E. a Policy
- F. a SQL Profiler Trace
- G. a Database Audit Specification

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 37

You administer a Microsoft SQL Server 2012 database named Contoso on a server named Server01.

You need to diagnose deadlocks that happen when executing a specific set of stored procedures by **recording events and playing them back on a different test server.**

What should you create?

- A. an Extended Event session
- B. a Policy
- C. a Database Audit Specification
- D. an Alert
- E. a Server Audit Specification
- F. a SQL Profiler Trace
- G. a Resource Pool

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

QUESTION 38

You administer a Microsoft SQL Server 2012 database named Contoso on a server named Server01.

You need to **prevent users from disabling server audits in Server01**.

What should you create?

- A. an Alert
- B. a Resource Pool
- C. an Extended Event session
- D. a Database Audit Specification
- E. a SQL Profiler Trace
- F. a Server Audit Specification
- G. a Policy

Correct Answer: G

Section: (none)

Explanation

Explanation/Reference:

QUESTION 39

You administer a SQL Server 2012 server that contains a database named SalesDb.

SalesDb contains a schema named Customers that has a table named Regions.

A user named UserA is a member of a role named Sales.

UserA is granted the Select permission on the Regions table.

The Sales role is granted the Select permission on the Customers schema.

You need to ensure that the **Sales role**, including UserA, is disallowed to select from any of the tables in the **Customers schema**.

Which Transact-SQL statement should you use?

- A. REVOKE SELECT ON Schema::Customers FROM UserA
- B. DENY SELECT ON Object::Regions FROM UserA
- C. EXEC sp_addrolemember 'Sales', 'UserA'
- D. DENY SELECT ON Object::Regions FROM Sales
- E. REVOKE SELECT ON Object::Regions FROM UserA
- F. DENY SELECT ON Schema::Customers FROM **Sales**
- G. DENY SELECT ON Schema::Customers FROM UserA
- H. EXEC sp_droprolemember 'Sales', 'UserA'
- I. REVOKE SELECT ON Object::Regions FROM Sales
- J. REVOKE SELECT ON Schema::Customers FROM Sales

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

References:

<http://msdn.microsoft.com/en-us/library/ms188369.aspx>

<http://msdn.microsoft.com/en-us/library/ms187750.aspx>

<http://msdn.microsoft.com/en-us/library/ff848791.aspx>

QUESTION 40

You administer a SQL Server 2012 server that contains a database named SalesDb.

SalesDb contains a schema named Customers that has a table named Regions.

A user named UserA is a member of a role named Sales.

UserA is granted the Select permission on the Regions table.

The Sales role is granted the Select permission on the Customers schema.

You need to **ensure that UserA is** disallowed to select from any of the tables in **the Customers schema**.

Which Transact-SQL statement should you use?

- A. DENY SELECT ON Object::Regions FROM UserA
- B. DENY SELECT ON Object::Regions FROM Sales
- C. REVOKE SELECT ON Schema::Customers FROM Sales
- D. REVOKE SELECT ON Schema::Customers FROM UserA
- E. REVOKE SELECT ON Object::Regions FROM Sales
- F. REVOKE SELECT ON Object::Regions FROM UserA
- G. DENY SELECT ON Schema::Customers FROM Sales
- H. DENY SELECT ON Schema::Customers FROM UserA
- I. EXEC sp_addrolemember 'Sales', 'UserA'
- J. EXEC sp_droprolemember 'Sales', 'UserA'

Correct Answer: H

Section: (none)

Explanation

Explanation/Reference:

<http://msdn.microsoft.com/en-us/library/ms188369.aspx>

<http://msdn.microsoft.com/en-us/library/ms187750.aspx>

<http://msdn.microsoft.com/en-us/library/ff848791.aspx>

QUESTION 41

You administer a SQL 2012 server that contains a database named SalesDb.

SalesDb contains a schema named Customers that has a table named Regions.

A user named UserA is a member of a role named Sales.

UserA is granted the Select permission on the Regions table.

The Sales role is granted the Select permission on the Customers schema.

You need to remove the Select permission for UserA on the Regions table.

You also need to ensure that **UserA** can **still access** all the tables in the Customers schema, including **the Regions table**, through the Sales role permissions.

Which Transact-SQL statement should you use?

- A. DENY SELECT ON Object::Regions FROM UserA
- B. DENY SELECT ON Schema::Customers FROM UserA
- C. EXEC sp_addrolemember 'Sales', 'UserA'
- D. REVOKE SELECT ON Object::Regions FROM UserA
- E. REVOKE SELECT ON Object::Regions FROM Sales
- F. EXEC sp_droprolemember 'Sales', 'UserA'
- G. REVOKE SELECT ON Schema::Customers FROM UserA
- H. DENY SELECT ON Object::Regions FROM Sales
- I. DENY SELECT ON Schema::Customers FROM Sales
- J. REVOKE SELECT ON Schema::Customers FROM Sales

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

<http://msdn.microsoft.com/en-us/library/ms188369.aspx>

<http://msdn.microsoft.com/en-us/library/ms187750.aspx>

<http://msdn.microsoft.com/en-us/library/ff848791.aspx>

QUESTION 42

You administer a SQL Server 2012 server that contains a database named SalesDb.

SalesDb contains a schema named Customers that has a table named Regions.

A user named UserA is a member of a role named Sales.

UserA is granted the Select permission on the Regions table and the Sales role is granted the Select permission on the Customers schema.

You need to ensure that the **Sales role**, including UserA, is disallowed to select from the **Regions table**.

Which Transact-SQL statement should you use?

- A. REVOKE SELECT ON Schema::Customers FROM UserA

- B. REVOKE SELECT ON Object::Regions FROM UserA
- C. EXEC sp_addrolemember 'Sales', 'UserA'
- D. DENY SELECT ON Schema::Customers FROM Sales
- E. EXEC sp_droprolemember 'Sales', 'UserA'
- F. REVOKE SELECT ON Schema::Customers FROM Sales
- G. DENY SELECT ON Object::Regions FROM UserA
- H. REVOKE SELECT ON Object::Regions FROM Sales
- I. DENY SELECT ON Schema::Customers FROM UserA
- J. DENY SELECT ON Object::Regions FROM Sales

Correct Answer: J

Section: (none)

Explanation

Explanation/Reference:

<http://msdn.microsoft.com/en-us/library/ms188369.aspx>

<http://msdn.microsoft.com/en-us/library/ms187750.aspx>

<http://msdn.microsoft.com/en-us/library/ff848791.aspx>

QUESTION 43

You administer all the deployments of Microsoft SQL Server 2012 in your company.

You need to ensure that an OLTP database that includes up-to-the-minute reporting requirements can be off-loaded from the primary database to another server.

You also Need to be able to **add indexes to the secondary database**.

Which configuration should you use?

- A. • Two servers configured in different data centers
 - SQL Server Availability Group configured in Synchronous-Commit Availability Mode
 - One server configured as an Active Secondary
- B. • Two servers configured in the same data center
 - SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
 - One server configured as an Active Secondary
- C. • Two servers configured in the same data center
 - A primary server configured to perform log-shipping every 10 minutes
 - A backup server configured as a warm standby
- D. • Two servers configured in different data centers
 - SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
- E. • Two servers configured on the same subnet
 - SQL Server Availability Group configured in Synchronous-Commit Availability Mode
- F. • SQL Server that includes an application database configured to perform transactional replication
- G. • SQL Server that includes an application database configured to perform snapshot replication

- H. • Two servers configured in a Windows Failover Cluster in the same data center
• SQL Server configured as a clustered instance

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

I'm not answering as I don't have any idea (original is F)

QUESTION 44

You administer all the deployments of Microsoft SQL Server 2012 in your company.

You need to ensure that data changes are sent to a non-SQL Server database server in near real time.

You also need to ensure that data on the primary server is unaffected.

Which configuration should you use?

- A. • SQL Server that includes an application database configured to perform transactional replication
- B. • Two servers configured in different data centers
• SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
- C. • Two servers configured in different data centers
• SQL Server Availability Group configured in Synchronous-Commit Availability Mode
• One server configured as an Active Secondary
- D. • SQL Server that includes an application database configured to perform snapshot replication
- E. • Two servers configured in the same data center
• SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
• One server configured as an Active Secondary
- F. • Two servers configured on the same subnet
• SQL Server Availability Group configured in Synchronous-Commit Availability Mode
- G. • Two servers configured in a Windows Failover Cluster in the same data center
• SQL Server configured as a clustered instance
- H. • Two servers configured in the same data center
• A primary server configured to perform log-shipping every 10 minutes
• A backup server configured as a warm standby

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Original Answer is A and I'll agree as it's the only one using Transactional Replication.

Must be to do with "You need to ensure that data changes are sent to a **non-SQL Server** database server in **near real time**."

QUESTION 45

You administer all the deployments of Microsoft SQL Server 2012 in your company.

A database contains a large product catalog that is updated periodically.

You need to be able to send the **entire product catalog to all branch offices on a monthly basis**.

Which configuration should you use?

- A.
 - Two servers configured in the same data center
 - A primary server configured to perform log-shipping every 10 minutes
 - A backup server configured as a warm standby
- B.
 - SQL Server that includes an application database configured to perform transactional replication
- C.
 - Two servers configured in the same data center
 - SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
 - One server configured as an Active Secondary
- D.
 - Two servers configured in a Windows Failover Cluster in the same data center
 - SQL Server configured as a clustered instance
- E.
 - SQL Server that includes an application database configured to perform snapshot replication
- F.
 - Two servers configured in different data centers
 - SQL Server Availability Group configured in Synchronous-Commit Availability Mode
 - One server configured as an Active Secondary
- G.
 - Two servers configured on the same subnet
 - SQL Server Availability Group configured in Synchronous-Commit Availability Mode
- H.
 - Two servers configured in different data centers
 - SQL Server Availability Group configured in Asynchronous-Commit Availability Mode

Correct Answer: E

Section: (none)

Explanation

Explanation/Reference:

I think that I'll stick with "E" as that was the original answer and "You need to be able to send the **entire product catalog to all branch offices on a monthly basis**." definitely sounds like it wants snapshot

QUESTION 46

You administer all the deployments of Microsoft SQL Server 2012 in your company.

You need to ensure that an OLTP database that uses a storage area network (SAN) remains available if any of the servers fail.

You also need to **minimize the amount of storage used by the database**.

Which configuration should you use?

- A. • Two servers configured in different data centers
 - SQL Server Availability Group configured in Synchronous-Commit Availability Mode
 - One server configured as an Active Secondary
- B. • SQL Server that includes an application database configured to perform transactional replication
- C. • Two servers configured in the same data center
 - SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
 - One server configured as an Active Secondary
- D. • Two servers configured in different data centers
 - SQL Server Availability Group configured in Asynchronous-Commit Availability Mode
- E. • Two servers configured in the same data center
 - A primary server configured to perform log-shipping every 10 minutes
 - A backup server configured as a warm standby
- F. • Two servers configured on the same subnet
 - SQL Server Availability Group configured in Synchronous-Commit Availability Mode
- G. • SQL Server that includes an application database configured to perform snapshot replication
- H. • Two servers configured in a Windows Failover Cluster in the same data center
 - SQL Server configured as a clustered instance

Correct Answer: H

Section: (none)

Explanation

Explanation/Reference:

I'll stick with "H" as that seems the logical choice.

QUESTION 47

You administer a Microsoft SQL Server 2012 instance.

After a routine shutdown, the drive that contains **tempdb** fails.

You need to be able to start the SQL Server. What should you do?

- A. Modify tempdb location in startup parameters.
- B. Start SQL Server in minimal configuration mode.

- C. Start SQL Server in single-user mode.
- D. Configure SQL Server to bypass Windows application logging.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms186400.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms345408.aspx>

QUESTION 48

You administer a single server that contains a Microsoft SQL Server 2012 default instance.

You plan to **install a new** application that requires the deployment of a database on the server.

The application login requires sysadmin permissions.

You need to ensure that the application login is unable to access other production databases. What should you do?

- A. Use the SQL Server default instance and configure an affinity mask.
- B. Install a new named SQL Server instance on the server.
- C. Use the SQL Server default instance and enable Contained Databases.
- D. Install a new default SQL Server instance on the server.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

I would have gone with Contained Databases, but the application requires sysadmin permissions.

QUESTION 49

You administer a Microsoft SQL Server 2012 instance named SQL2012 that hosts an OLTP database of 1 terabyte in size.

The database is modified by users only from Monday through Friday from 09:00 hours to 17:00 hours.

Users modify more than 30 percent of the data in the database during the week. Backups are performed as shown in the following schedule:

Type	Frequency
Full	Sunday at 20:00 hours
Differential	Monday through Friday at 20:00 hours
Log	Monday through Friday between 08:00 hours and 18:00 hours

The Finance department plans to execute a batch process every Saturday at 09:00 hours.

This batch process will take a maximum of 8 hours to complete.

The batch process will update three tables that are 10 GB in size.

The batch process will update these tables multiple times.

When the batch process completes, the Finance department runs a report to find out whether the batch process has completed correctly.

You need to ensure that if the Finance department disapproves the batch process, the batch operation can be rolled back in the minimum amount of time.

What should you do on Saturday?

- A. Perform a differential backup at 08:59 hours.
- B. Record the LSN of the transaction log at 08:59 hours. Perform a transaction log backup at 17:01 hours.
- C. Create a database snapshot at 08:59 hours.
- D. Record the LSN of the transaction log at 08:59 hours. Perform a transaction log backup at 08:59 hours.
- E. Create a marked transaction in the transaction log at 08:59 hours. Perform a transaction log backup at 17:01 hours.
- F. Create a marked transaction in the transaction log at 08:59 hours. Perform a transaction log backup at 08:59 hours.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 50

You administer a Microsoft SQL Server 2012 instance.

The instance contains a database that supports a retail sales application.

The application generates hundreds of transactions per second and is online 24 hours per day and 7 days per week.

You plan to define a backup strategy for the database.

You need to ensure that the following requirements are met:

- No more than 5 minutes worth of transactions are lost.
- Data can be recovered by using the minimum amount of administrative effort.

What should you do? Choose all that apply.

- A. Configure the database to use the SIMPLE recovery model.
- B. Create a DIFFERENTIAL database backup every 4 hours.
- C. Create a LOG backup every 5 minutes.
- D. Configure the database to use the FULL recovery model.
- E. Create a FULL database backup every 24 hours.
- F. Create a DIFFERENTIAL database backup every 24 hours.

Correct Answer: BCDE

Section: (none)

Explanation

Explanation/Reference:

QUESTION 51

You administer a Microsoft SQL Server 2012 database that contains a table named OrderDetail.

You discover that the NCI_OrderDetail_CustomerID non-clustered index is fragmented.

You need to reduce fragmentation.

You need to achieve this goal without taking the index offline.

Which Transact-SQL batch should you use?

- A. `CREATE INDEX NCI_OrderDetail_CustomerID ON OrderDetail.CustomerID WITH DROP EXISTING`
- B. `ALTER INDEX NCI_OrderDetail_CustomerID ON OrderDetail.CustomerID REORGANIZE`
- C. `ALTER INDEX ALL ON OrderDetail REBUILD`
- D. `ALTER INDEX NCI_OrderDetail_CustomerID ON OrderDetail.CustomerID REBUILD`

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms188388.aspx>

QUESTION 52

You administer a Microsoft SQL Server 2012 server.

You plan to deploy new features to an application.

You need to evaluate existing and potential clustered and non-clustered indexes that will improve performance.

What should you do?

- A. Query the sys.dm_db_index_usage_stats DMV.
- B. Query the sys.dm_db_missing_index_details DMV.
- C. Use the Database Engine Tuning Advisor.
- D. Query the sys.dm_db_missing_index_columns DMV.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms174202.aspx>

QUESTION 53

You administer a Microsoft SQL Server 2012 instance. You need to configure a new database to support FILETABLES. What should you do? Choose all that apply.

- A. Disable FILESTREAM on the Database.
- B. Enable FILESTREAM on the Server Instance.
- C. Configure the Database for Partial Containment.
- D. Create a non-empty FILESTREAM file group.
- E. Enable Contained Databases on the Server Instance.
- F. Set the FILESTREAM directory name on the Database.

Correct Answer: BDF

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/gg509097.aspx>

QUESTION 54

You administer two instances of Microsoft SQL Server 2012. You deploy an application that uses a database on the named instance. The application is unable to connect to the database on the named instance. You need to ensure that the application can connect to the named instance. What should you do?

- A. Configure the application as data-tiered.
- B. Open port 1433 on the Windows firewall on the server.
- C. Configure the named SQL Server instance to use an account that is a member of the Domain Admins group.
- D. Start the SQL Server Browser Service.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 55

You administer a Microsoft SQL Server 2012 database. You configure Transparent Data Encryption (TDE) on the Orders database by using the following statements:

```
CREATE MASTER KEY ENCRYPTION BY PASSWORD = 'MyPassword1!'
CREATE CERTIFICATE TDE_Certificate WITH SUBJECT = 'TDE Certificate';

BACKUP CERTIFICATE TDE_Certificate TO FILE = 'd:\TDE_Certificate.cer'
WITH PRIVATE KEY (FILE = 'D:\TDE_Certificate.key', ENCRYPTION BY PASSWORD =
'MyPassword1!');

CREATE DATABASE ENCRYPTION KEY
WITH ALGORITHM = AES_256
ENCRYPTION BY SERVER CERTIFICATE TDE_Certificate;

ALTER DATABASE Orders SET ENCRYPTION ON;
```

You attempt to restore the Orders database and the restore fails. You copy the encryption file to the original location. A hardware failure occurs and so a new server must be installed and configured. After installing SQL Server to the new server, you restore the Orders database and copy the encryption files to their original location. However, you are unable to access the database. You need to be able to restore the database. Which Transact-SQL statement should you use before attempting the restore?

- A. ALTER DATABASE Master SET ENCRYPTION OFF;
- B. CREATE CERTIFICATE TDE_Certificate FROM FILE = 'd:\TDE_Certificate.cer' WITH PRIVATE KEY (FILE = 'D:\TDE_Certificate.key', DECRYPTION BY PASSWORD = 'MyPassword1!');
- C. CREATE CERTIFICATE TDE_Certificate WITH SUBJECT = 'TDE Certificate';
USE Orders;
CREATE DATABASE ENCRYPTION KEY
WITH ALGORITHM = AES_256
ENCRYPTION BY SERVER CERTIFICATE TDE_Certificate;

D. CREATE CERTIFICATE TDE_Certificate FROM FILE = 'd:\TDE_Certificate.cer';

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 56

You administer a Microsoft SQL Server 2012 database.

The database contains a customer table created by using the following definition:

```
CREATE TABLE dbo.Customer
(
  CustomerID INT PRIMARY KEY,
  CustomerName VARCHAR(100) NOT NULL,
  CustomerAddress1 CHAR(200) NOT NULL,
  CustomerAddress2 CHAR(200) NULL,
  CustomerCity VARCHAR(100) NOT NULL,
  CustomerPostalCode CHAR(5) NOT NULL);
```

You need to ensure that the minimum amount of disk space is used to store the data in the customer table.

What should you do?

- A. Implement row-level compression.
- B. Implement page-level compression.
- C. Convert all indexes to Column Store indexes.
- D. Implement Unicode compression.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 57

You administer a Microsoft SQL Server 2012 database named ContosoDb.

The database contains a table named Suppliers and a column named IsActive in the Purchases schema.

You create a new user named ContosoUser in ContosoDb. ContosoUser has no permissions to the Suppliers table.

You need to ensure that ContosoUser can delete rows that are not active from Suppliers.

You also need to grant ContosoUser only the minimum required permissions. Which Transact-SQL statement should you use?

- A. GRANT DELETE ON Purchases.Suppliers TO ContosoUser
- B. CREATE PROCEDURE Purchases.PurgeInactiveSuppliers
WITH EXECUTE AS USER = 'dbo'
AS

```
DELETE FROM Purchases.Suppliers WHERE IsActive = 0
GO
GRANT EXECUTE ON Purchases.PurgeInactiveSuppliers TO ContosoUser
C. GRANT SELECT ON Purchases.Suppliers TO ContosoUser
D. CREATE PROCEDURE Purchases.PurgeInactiveSuppliers
AS
DELETE FROM Purchases.Suppliers WHERE IsActive = 0
GO
GRANT EXECUTE ON Purchases.PurgeInactiveSuppliers TO ContosoUser
```

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

Need to check these

Reference: <http://msdn.microsoft.com/en-us/library/ms188354.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms187926.aspx>

Exam B

QUESTION 1

You administer a Microsoft SQL Server 2012 instance that has multiple databases , You have a two-node SQL Server failover cluster . the cluster uses a storage area network (SAN)

You discover I/O issues. The SAN is at capacity and additional disks cannot be added

You need to reduce the I/O workload on the SAN at a minimal cost

what should you do ?

- A. Move user databases to a local disk
- B. modify application code to use table variable
- C. expand the tempdb data and log files
- D. move the tempdb files to a local disk

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 2

You administer a Microsoft SQL Server 2012 instance named SQL2012. You are in the process of migrating a database from a SQL Server 2008 instance named SQL2008 to the SQL2012 instance

You have upgraded a database from the SQL2008 instance by using the side-by-side migration technique.

You need to migrate the SQL Server logins from the **SQL2008 instance to the SQL2012** instance

what should you do ?

- A. Use sp_grantlogin.
- B. Back up the master database on the SQL2008 instance. Restore the master database on the SQL2012 instance
- C. use BCP.EXE
- D. use the transfer logins task in a Microsoft SQL integrated Services package

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 3

You plan to install a Microsoft SQL server 2012 instance

the instance will support a database that has the following requirements :

- Store **Excel** workbooks on the file system
- Access the workbooks through Transact-SQL
- Include the workbooks in database backups

During installation , you need to ensure that the requirements will be met

which feature should you use ?

- A. FILESTREAM
- B. SQL Server Integration Services (SSIS)
- C. Excel Services
- D. OpenXML

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 4

You administer a Microsoft SQL Server 2012 database

Users report that a billing application becomes unresponsive during busy time of the day

while investigating, you notice large number of processes taking or waiting for table locks

You suspect that SQL Server is assigning stronger locks to queries

You start a SQL Profiler Trace

Which event should you select ?

- A. Lock: Deadlock
- B. Deadlock graph
- C. Lock: Timeout
- D. Lock: Escalation

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 5

You administer a Microsoft SQL Server 2012 database instance, You create a New user named UserA

You need to ensure that UserA is able to create SQL Server Agent jobs and execute SQL Server agent jobs owned by UserA

To which role should you add UserA?

- A. ServerGroupAdministratorGroup
- B. DatabaseMailUserRole
- C. SQLAgentUserRole
- D. ServerAdmin

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 6

You Plan to install Microsoft SQL Server 2012 for a web hosting company

your company plans to host multiple web sites, each supported by a SQL Server database

You need to select an edition of SQL Server that features backup compression of databases, basic data integration features, and low total test of ownership

which edition should you choose ?

- A. Express Edition with Tools
- B. Web Edition
- C. Express Edition with Advanced Service
- D. Standard Edition

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 7

You administer a Microsoft SQL Server 2012 Database,

You have a SQL Server Agent Job instance that runs using the service account, You have a job step within the job that privileges.

You need to ensure that the job step can run using a different user account.

what should you use ?

- A. a schedule
- B. an operator
- C. an alert
- D. a proxy

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 8

You administer a Microsoft SQL Server environment. You purchase a new server and plan to migrate your database from SQL server 2008 to SQL Server 2012

You want to evaluate to prepare for possible conflicts and issues that may arise during or after the migration

which SQL Server tool should you use ?

- A. Data Tools

- B. Distributed Replay
- C. Upgrade Advisor
- D. Migration Assistant

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 9

You administer a Microsoft SQL Server 2012 Enterprise Edition server that uses 64 cores

You discover performance issues when complex calculations are performed on large amount of data under heavy system

You need to limit the number of cores that handle the processing.

what should you configure ?

- A. Lightweight pooling
- B. processor affinity
- C. Max worker threads
- D. I/O affinity

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 10

You are implementing a SQL Server 2012 **five-node failover cluster**

You need to choose a quorum configuration

which configuration should you use ?

- A. Distributed file system (DFS)
- B. Cluster Shared Volume (CSV)
- C. Node and Disk Majority
- D. Node Majority

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

Node majority option is recommended for cluster with **odd number** of nodes

Node and Disk Majority: This option is recommended for cluster with **even number of nodes**

QUESTION 11

You are creating an application that will connect to the AgentPortal database by using a SQL login named AgentPortalUser ..Procedures in the database will **SP_send_dbmail** to send email messages.

You create a user account in the msdb database for the AgentPortalUser login. You use the Database Mail Configuration Database Mail Profile, Security has not been configured for the Database Mail profile

You need to ensure that AgentPortalUser can send email messages.

What should you do ?

- A. Enable the guest user in the msdb database.
- B. in the Database Mail Configuration Wizard, create an email account for each recipient's email address in the Database profile
- C. Disable the guest user in the msdb database
- D. configure the AgentportalUser user as a member of DatabaseMailUser
- E. n the Database Mail Configuration Wizard, configure the Database Mail profile as a private profile for the AgentPortalUser account.

Correct Answer: E

Section: (none)

Explanation

Explanation/Reference:

if ans E not found IN ANSWERS ---> choose B

Reference: <http://msdn.microsoft.com/en-us/library/ms189635.aspx>

QUESTION 12

You are a database administrator for a Microsoft SQL Server 2012 instance.

You need to ensure that data can be migrated from a production server to two reporting servers with minimal data to ensure that data on the reporting server is always accessible

which solution should you use ?

- A. Database snapshot
- B. Log Shipping
- C. Change Data Capture
- D. Availability Groups

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 13

You install Microsoft SQL Server 2012 on a new server.

After setup is complete, you attempt to start the SQL Server service.
After being in a starting state for a few moments,
the service goes back to a stopped state.

You need to determine the cause of the **failure**.

Which file should you use?

- A. %programfiles%\Microsoft SQL Server\MSSQL11.MSSQLSERVER\MSSQL\Log\Errorlog
- B. %programfiles%\Microsoft SQL Server\110\Setup Bootstrap\Log\Summary.txt

- C. %programfiles%\Microsoft SQL Server\110\Shared>ErrorDumps\SQLDmpr[XXXX].mdmp
- D. %programfiles%\Microsoft SQL Server\MSSQL11.MSSQLSERVER\MSSQL\DATA\mastlog.ldf

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 14

You administer a Microsoft SQL Server 2012

A process that normally runs in less than 10 seconds has been running for more than an hour. you exam discover that the process is using session ID 60.

You need to find out whether the process is being blocked.

which Transact-SQL statement should you use?

- A. SELECT * FROM sys.dm_exec_sessions WHERE session_id = 60
- B. DBCC INPUTBUFFER (60)
- C. EXEC sp_helpdb 60
- D. SELECT * FROM sys.dm_exec_request WHERE session_id = 60

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 15

You administer three Microsoft SQL Server 2012 servers

You need to install the Power View components on a stand-alone server.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the an.... area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> <div></div> </div>	<div>Install SQL Server 2012 Analysis Services.</div> <div>Configure Microsoft SQL Server Reporting Services in Native mode.</div> <div>Install SQL Server 2012 Reporting Services</div> <div>Install SQL Server 2012 Database Engine and PowerPivot for SharePoint</div> <div>Configure Microsoft SQL Server Reporting Services Add-in in SharePoint mode.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Install SQL Server 2012 Reporting Services

Install SQL Server 2012 Database Engine and PowerPivot for SharePoint

Configure Microsoft SQL Server Reporting Services Add-in in SharePoint mode.

Section: (none)

Explanation

Explanation/Reference:

QUESTION 16

You administer a Microsoft SQL Server 2012 database.

You need to convert the database to a contained database.

You also need to ensure that all users are converted to contained users.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> </div> <div></div>	<div>Execute the ALTER DATABASE statement along with CONTAINMENT=PARTIAL.</div> <div>Execute the ALTER DATABASE statement along with CONTAINMENT=TRUE.</div> <div>Execute sp_configure 'cross db ownership chaining', 1; RECONFIGURE.</div> <div>Execute sp_configure 'contained database authentication', 1; RECONFIGURE.</div> <div>Execute sp_migrate_user_to_contained for the database.</div> <div>Execute sp_migrate_user_to_contained for each user.</div>
<div><< Move</div> <div>Remove >></div>	

Correct Answer:

Execute sp_configure 'contained database authentication', 1; RECONFIGURE.

Execute the ALTER DATABASE statement along with CONTAINMENT=PARTIAL.

Execute sp_migrate_user_to_contained for each user.

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ff929071.aspx>

QUESTION 17

You are implementing a SQL Server 2012 four-node failover cluster

You need to choose a quorum configuration

which configuration should you use ?

- A. Distributed file system (DFS)
- B. Cluster Shared Volume (CSV)
- C. Node and Disk Majority
- D. Node Majority

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

Node and Disk Majority: This option is recommended for cluster with **even number of nodes**

QUESTION 18

You administer a Microsoft SQL Server 2012 database named orders.

Users report that during peak usage periods, certain operations are taking more time than expected. Your initial analysis blocking is the cause.

You need to gather more data to be able to determine which processes are being **blocked** and to identify the root cause.

What should you do ?

- A. Schedule A SQL Agent job to run every 60 Seconds and insert the results of executing the SP_who2 stored procedure table
- B. Use System Monitor to catch the Lock Wait Time event
- C. Use Sp_Configure to set the blocked process threshold. Start a Trace using SQL Server Profiler to catch the Blocking Pro-Report event.
- D. Start a Trace using SQL Server Profiler to catch the Lock: Timeout event

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 19

You administer a Microsoft SQL Server 2012 server.

One of the databases on the server supports a highly active OLTP application.

Users report abnormally long wait times when they submit data into the application.

You need to identify which queries are taking **longer than 1** second to run over an extended period of time. What should you do?

- A. Use SQL Profiler to trace all queries that are processing on the server. Filter queries that have a Duration value of more than 1,000.
- B. Use sp_configure to set a value for blocked process threshold. Create an extended event session.
- C. Run the sp_who command from a query window.
- D. Run the DBCC TRACEON 1222 command from a query window and review the SQL Server event log.
- E. Use the Job Activity monitor to review all processes that are actively running. Review the Job History to find out the duration of each step.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

Verified the SQL Profiler and DBCC answers as correct. However, while Profiler will show this information, the best practice with Profiler is to use it short-term. The question specifically states "over an extended period of time". That means Profiler wouldn't be the best tool for this scenario. Therefore, DBCC would be the best answer.

Reference: <http://www.mssqltips.com/sqlservertip/2130/finding-sql-server-deadlocks-using-trace-flag-1222/>

Reference: <http://msdn.microsoft.com/en-us/library/ms188396.aspx>

QUESTION 20

You administer a Microsoft SQL Server 2012 environment. One of the SQL Server 2012 instance contain a database named sales

You plan to migrate Sales to windows Azure SQL Database

To do so , you need to implement a contained database.

What should you do ? (Each correct answer presents part of the solution. Choose all that apply)

- A. Enable server property Contained Database Authentication
- B. Set database containment to partial
- C. Set database containment to full
- D. Disable server property cross db ownership chaining
- E. Set database containment to AZURE
- F. Disable server property Contained Database Authentication

Correct Answer: ABE

Section: (none)

Explanation

Explanation/Reference:

QUESTION 21

You are the lead database administrator (DBA) of a Microsoft SQL Server 2012 environment.

All DBAs are members of the DOMAIN\JrDBAs Active Directory group.

You grant DOMAIN\JrDBAs access to the SQL Server.

You need to create a server role named SpecialDBARole that can perform the following functions:

- View all databases.
- View the server state.
- Assign GRANT, DENY, and REVOKE permissions on logins.

You need to add DOMAIN\JrDBAs to the server role.

You also need to provide the least level of privileges necessary.

Which SQL statement or statements should you use? Choose all that apply.

- A. `CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION setupadmin;`
- B. `ALTER SERVER ROLE [SpecialDBARole] ADD MEMBER [DOMAIN\JrDBAs];`
- C. `CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION securityadmin;`
- D. `GRANT VIEW DEFINITION TO [SpecialDBARole];`
- E. `CREATE SERVER ROLE [SpecialDBARole] AUTHORIZATION serveradmin;`
- F. `GRANT VIEW SERVER STATE, VIEW ANY DATABASE TO [SpecialDBARole];`

Correct Answer: BCF

Section: (none)

Explanation

Explanation/Reference:

QUESTION 22

You install a Microsoft SQL Server 2012 instance

the instance will store data extracted from two database running on windows Azure SQL Database

You hire a data steward to perform interactive data cleaning and ad hoc querying and updating of the data

You need to ensure that the data steward is given the correct client tools to perform these tasks

Which set of tools should you install?

- A. SQL Server Management Studio and Distributed Reply Client
- B. Data Quality Client and SQL Server Data Tools
- C. Data Quality Client and Distributed Replay Client
- D. SQL Server Management Studio and Master Data Services

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 23

You administer three Microsoft SQL Server 2008 R2 instances. Database mirroring is configured in High-Safety mode with Automatic Failover between the following three servers:

- SQL1 is the Principal server.
- SQL2 is the mirror server.
- SQL3 is the witness server.

You need to upgrade SQL1 and SQL2 to SQL Server 2012. You need to ensure that downtime is minimized during the upgrade. Which six actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> <div></div> </div>	<div>Configure log shipping between SQL1 and SQL2.</div> <div>Upgrade SQL1 to SQL Server 2012.</div> <div>Upgrade SQL2 to SQL Server 2012.</div> <div>Disable log shipping between SQL1 and SQL2.</div> <div>Manually failover the database from SQL1 to SQL2.</div> <div>Manually failover the database from SQL2 to SQL1.</div> <div>Add SQL3 back to the database mirroring solution.</div> <div>Remove SQL3 from the database mirroring solution.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Remove SQL3 from the database mirroring solution.

Upgrade SQL2 to SQL Server 2012.

Manually failover the database from SQL1 to SQL2.

Upgrade SQL1 to SQL Server 2012.

Manually failover the database from SQL2 to SQL1.

Add SQL3 back to the database mirroring solution.

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

Reference: <http://msdn.microsoft.com/en-us/library/bb677181.aspx>

QUESTION 24

You administer a Microsoft SQL Server 2012 database.

Your database is experiencing deadlock issues.

You need to be able to monitor deadlocks.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div> <div></div> <div></div> </div> <div></div> </div>	<div>Start Microsoft SQL Server Management Studio.</div> <div>Start SQL Server Configuration Manager and locate the SQL Server service.</div> <div>Restart the SQL Server service for that particular instance.</div> <div>Run the DBCC TRACEON (1221, -1) Transact-SQL query.</div> <div>From the SQL Server Properties page, click the Startup parameters tab and add Trace Flag -T1222 to the start-up parameters list.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Start SQL Server Configuration Manager and locate the SQL Server service.

From the SQL Server Properties page, click the Startup parameters tab and add Trace Flag -T1222 to the start-up parameters list.

Restart the SQL Server service for that particular instance.

Section: (none)

Explanation

Explanation/Reference:

QUESTION 25

You administer a Microsoft SQL Server 2012 server that has a database named Contoso.

The Contoso database has a table named ProductPrices in a schema named Sales.

You need to create a script that writes audit events into the application log whenever data in the ProductPrices table is updated.

Which four Transact-SQL statements should you use? (To answer, move the appropriate statements from the list of statements to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div> <div>▲</div> <div>▼</div> </div> <div></div>	<div>CREATE DATABASE AUDIT SPECIFICATION C_AuditSpec FOR SERVER AUDIT C_Audit ADD (UPDATE ON Sales.ProductPrices BY dbo)</div> <div>ALTER DATABASE AUDIT SPECIFICATION C_AuditSpec WITH (STATE=ON)</div> <div>USE Master</div> <div>CREATE SERVER AUDIT C_Audit TO FILE (FILEPATH = 'ApplicationLog')</div> <div>ALTER SERVER AUDIT C_Audit WITH (STATE = ON)</div> <div>CREATE SERVER AUDIT C_Audit TO APPLICATION_LOG</div> <div>ALTER SERVER AUDIT C_Audit WITH (STATE = ON)</div> <div>USE Contoso</div> <div>CREATE SERVER AUDIT SPECIFICATION C_AuditSpec FOR SERVER AUDIT C_Audit ADD (SCHEMA_OBJECT_ACCESS_GROUP)</div> <div>ALTER SERVER AUDIT SPECIFICATION C_AuditSpec WITH (STATE = ON)</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

```
USE Master
CREATE SERVER AUDIT C_Audit
TO APPLICATION_LOG

ALTER SERVER AUDIT C_Audit
WITH (STATE = ON)

USE Contoso
CREATE DATABASE AUDIT SPECIFICATION
C_AuditSpec
FOR SERVER AUDIT C_Audit
ADD (UPDATE ON Sales.ProductPrices BY
dbo)

ALTER DATABASE AUDIT SPECIFICATION
C_AuditSpec WITH (STATE=ON)
```

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/cc280386.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/cc280448.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/cc280404.aspx>

QUESTION 26

You administer a Microsoft SQL Server 2012 database named Contoso that contains a Single user-defined database role BillingUsers

All objects in contoso are in the dbo schema

You need to grant EXECUTE permission for all stored procedures in contoso to BillingUsers

Which Transact-SQL statement should you see?

- A.

```
CREATE ROLE proc_caller
GRANT EXECUTE ON Schema : : dbo TO proc_caller
ALTER ROLE proc_caller ADD MEMBER BillingUsers
```
- B.

```
GRANT EXECUTE ON INFORMATION_SCHEMA.ROUTINES TO BillingUsers
```
- C.

```
EXEC sp_addrolemember 'executor', 'BillingUsers'
```
- D.

```
CREATE ROLE proc_caller
GRANT EXECUTE ON ALL PROCEDURES TO proc_caller
ALTER MEMBER BillingUsers ADD TO ROLE proc_caller
```

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 27

You administer a SQL Server 2012 database instance

You need to configure the SQL Server Database Engine service on a Failover cluster

Which user account should you see ?

- A. the BUILTIN\LocalService account
- B. a domain user
- C. a local administrative user
- D. the BUILTIN\NetworkService account

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 28

You administer a Microsoft SQL Server 2012 database instance

You plan to Migration the database to windows Azure SQL Database. You verify all object contained in the database are compatible with Windows Azure SQL Database

You need to ensure that database users and required server logins are migrated to windows Azure SQL Database

What should you do ?

- A. use the copy database wizard
- B. use the Database Transfer wizard
- C. use SQL Server Management Studio to deploy the database to Windows Azure SQL Database
- D. Backup the database from the local server and restore it to Windows Azure SQL Database

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 29

You Administer a Microsoft SQL Server 2012 failover cluster

You need to ensure that a failover occurs when the server diagnostics returns query_processing error

Which server configuration property should you set?

- A. FailureConditionLevel
- B. SqlDumperDumpPath
- C. SqlDumperDumpFlags
- D. HealthCheckTimeout

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

<http://technet.microsoft.com/en-us/library/ff878664.aspx>

QUESTION 30

You the administrator of a Microsoft SQL Server 2012 server

Your application consume significant resources. You need to manage the server workload by restricting resource-intensive

You need to dynamically limit **resource consumption**

what should you do ?

- A. Configure Resource Pools, Workload Groups, and Classifier Function, and then enable the Resource Governor
- B. Set up Service Broker to ensure that application are not allowed to consume more than the specified amount of resource
- C. create a new rule for each application that sets the resource limit allowed
- D. create a new plan Guide with a Scope Type of sql and define the resource limits for each application

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 31

You have been hired as a Database Consultant by ABC.com to design a SQL Server 2012 database solution.

You are tasked with designing a scale-out and high-availability SQL Server 2012 Online Transaction Processing (OLTP) database solution that will maintain copies of data across two server instances.

Your solution must provide scale-out of read operations by distributing the reads from clients across two SQL Server 2012 nodes. The data in both SQL Server nodes needs to be indexed.

What should you include in your solution?

- A. You should include two servers configured in an Active-Active SQL Server 2012 Cluster
- B. You should include a primary SQL Server 2012 database that uses transactional replication to replicate data to a secondary database.
- C. You should include two servers configured in an Active-Passive SQL Server 2012 Cluster.
- D. You should include two servers in an Asynchronous-Commit Availability Mode Availability Group.
- E. You should include two servers in a Synchronous-Commit Availability Mode Availability Group

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 32

You have been hired as a Database Consultant by ABC.com to design a database solution for a new application.

You are tasked with designing a high-availability database solution that uses SQL Server 2012 to host a primary database. The solution should maintain a near real-time copy of the data on a second non-Microsoft database.

What should you include in your solution?

- A. You should include a primary database with scheduled log shipping to the secondary database configured.
- B. You should include two servers configured in an Active-Passive SQL Server 2012 Cluster.
- C. You should include a primary SQL Server 2012 database that uses transactional replication to replicate data to the secondary database.

- D. You should include two servers in an Asynchronous-Commit Availability Mode Availability Group.
- E. You should include two servers in a Synchronous-Commit Availability Mode Availability Group.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 33

You work as the Senior Database Administrator (DBA) at ABC.com. The company has a main office and 10 branch offices. Each branch office contains a single database server running Microsoft SQL Server 2012. The main office has multiple clustered servers running Microsoft SQL Server 2012.

Your role includes the management of the entire Microsoft SQL Server 2012 infrastructure. The company runs a custom application that stores data in a large Microsoft SQL Server 2012 database. The primary database is hosted in the main office. Each branch office SQL Server hosts a copy of the database.

You need to configure a solution that will replicate the entire primary database from the main office SQL Server every weekend.

What should you include in your solution?

- A. Transactional Replication
- B. Log Shipping.
- C. Snapshot Replication.
- D. SQL Server Availability Group.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 34

You work as a Database Administrator (DBA) at ABC.com. The infrastructure includes servers running Microsoft SQL Server 2012. All databases are hosted on a SAN (Storage Area Network).

You need to design a database solution for a new application. You are tasked with designing a high-availability database solution. The solution must include a single copy of the database to save disk space and the database must remain online in the event of a SQL Server failure.

What should you include in your solution?

- A. You should include two servers configured as a failover cluster.
- B. You should include two servers and database mirroring.
- C. You should include two servers and log shipping.
- D. You should include two servers configure as a SQL Server Availability Group.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 35

You work as a Database Administrator (DBA) at ABC.com. The infrastructure includes servers running Windows Server 2008 R2 and Microsoft SQL Server 2012.

The company uses several custom applications that store data in databases on the Microsoft SQL Server 2012 servers.

A full backup of all databases is taken every night at midnight. A differential backup of all databases is taken on the hour every hour starting at 3am until the last backup at 11pm. A log backup is taken every 15 minutes for databases configured with the Full Recovery Model.

One application named ABCApp1 stores data in a database named ABCApp1DB. ABCApp1DB is configured with the Simple Recovery Model. ABCApp1DB fails at 3:25am. You discover that the last differential backup of ABCApp1DB failed.

You need to restore ABCApp1DB from backup as quickly as possible and minimize data loss.

Which of the following steps should you perform to restore ABCApp1DB? (Choose one or more answers).

- A. Restore the latest full backup.
- B. Restore the latest differential backup.
- C. Restore the latest log backup.
- D. Restore each differential backup taken since the last full backup.
- E. Restore each log backup since the last full backup.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 36

You are a database developer of a Microsoft SQL Server 2012 database. You are designing a table that will store Customer data from different sources. The table will include a column that contains the CustomerID from the source system and a column that contains the SourceID. A sample of this data is as shown in the following table.

SourceID	CustomerID	Customer Name
1	234	John Smith
3	7345	Jason Warren
3	4402	Susan Burk
2	866	Michael Allen

You need to ensure that the table has no duplicate CustomerID within a SourceID. You also need to ensure that the data in the table is in the order of SourceID and then CustomerID. Which Transact- SQL statement should you use?

- A.

```
CREATE TABLE Customer
(SourceID int NOT NULL IDENTITY,
CustomerID int NOT NULL IDENTITY,
CustomerName varchar(255) NOT NULL);
```
- B.

```
CREATE TABLE Customer
(SourceID int NOT NULL,
CustomerID int NOT NULL PRIMARY KEY CLUSTERED,
CustomerName varchar(255) NOT NULL);
```
- C.

```
CREATE TABLE Customer
(SourceID int NOT NULL PRIMARY KEY CLUSTERED,
CustomerID int NOT NULL UNIQUE,
CustomerName varchar(255) NOT NULL);
```
- D.

```
CREATE TABLE Customer
(SourceID int NOT NULL,
CustomerID int NOT NULL,
CustomerName varchar(255) NOT NULL,
CONSTRAINT PK_Customer PRIMARY KEY CLUSTERED
(SourceID, CustomerID));
```

Correct Answer: D

Section: (none)

Explanation

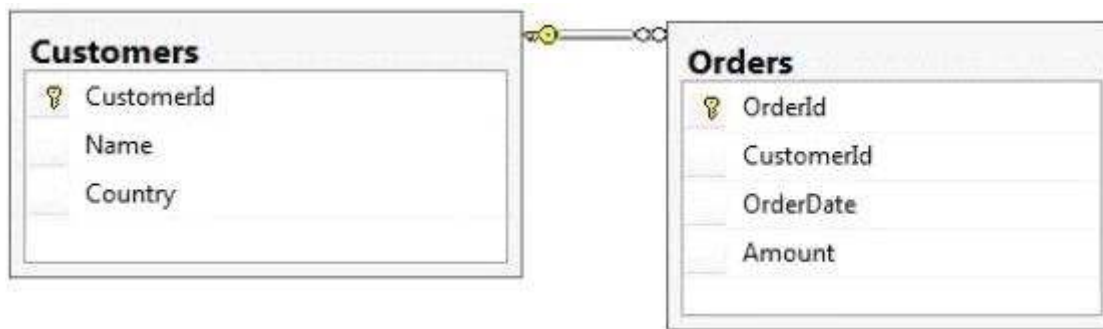
Explanation/Reference:

Verified the answer as correct.

D option, I met in test, goes with unique key on 2 combined columns. Still choose this.

QUESTION 37

You administer a Microsoft SQL Server 2012 database named ContosoDb. Tables are defined as shown in the exhibit. (Click the Exhibit button.)



You need to display rows from the Orders table for the Customers row having the CustomerId value set to 1 in the following XML format.

```

<Customers>
  <Name>Customer A</Name>
  <Country>Australia</Country>
  <Orders>
    <OrderId>1</OrderId>
    <OrderDate>2000-01-01T00:00:00</OrderDate>
    <Amount>3400.00</Amount>
  </Orders>
  <Orders>
    <OrderId>2</OrderId>
    <OrderDate>2001-01-01T00:00:00</OrderDate>
    <Amount>4300.00</Amount>
  </Orders>
</Customers>
  
```

Which Transact-SQL query should you use?

- A. `SELECT OrderId, OrderDate, Amount, Name, Country FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId WHERE Customers.CustomerId = 1 FOR XML RAW`
- B. `SELECT OrderId, OrderDate, Amount, Name, Country FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId WHERE Customers.CustomerId = 1 FOR XML RAW, ELEMENTS`
- C. `SELECT OrderId, OrderDate, Amount, Name, Country FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId WHERE Customers.CustomerId = 1 FOR XML AUTO`
- D. `SELECT OrderId, OrderDate, Amount, Name, Country FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId WHERE Customers.CustomerId = 1 FOR XML AUTO, ELEMENTS`
- E. `SELECT Name, Country, OrderId, OrderDate, Amount FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId WHERE Customers.CustomerId = 1 FOR XML AUTO`
- F. `SELECT Name, Country, OrderId, OrderDate, Amount FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId WHERE`

```
Customers.CustomerId= 1
FOR XML AUTO, ELEMENTS
```

- G. `SELECT Name AS '@Name', Country AS '@Country', OrderId, OrderDate, Amount FROM Orders INNER JOIN Customers ON Orders.CustomerId= Customers.CustomerId WHERE Customers.CustomerId= 1 FOR XML PATH ('Customers')`
- H. `SELECT Name AS 'Customers/Name', Country AS 'Customers/Country', OrderId, OrderDate, Amount FROM Orders INNER JOIN Customers ON Orders.CustomerId= Customers.CustomerId WHERE Customers.CustomerId= 1 FOR XML PATH ('Customers')`

Correct Answer: F

Section: (none)

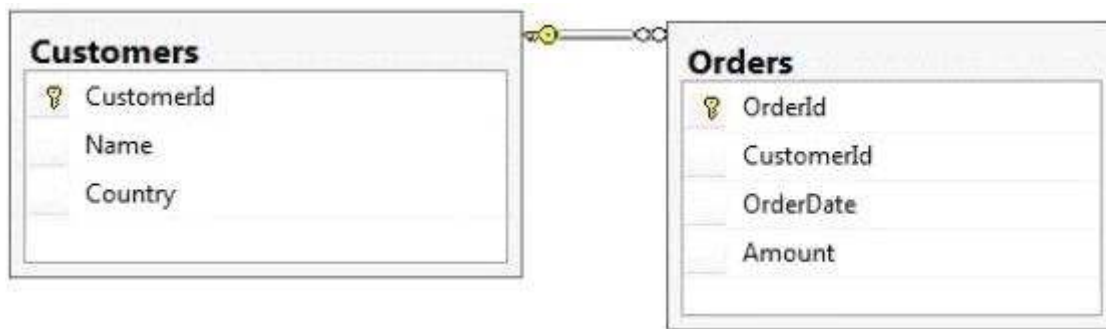
Explanation

Explanation/Reference:

Verified answer as correct.

QUESTION 38

You administer a Microsoft SQL Server 2012 database named ContosoDb. Tables are defined as shown in the exhibit. (Click the Exhibit button.)



You need to display rows from the Orders table for the Customers row having the CustomerId value set to 1 in the following XML format.

```
<Customers Name="Customer A" Country="Australia">
  <OrderId>1</OrderId>
  <OrderDate>2000-01-01T00:00:00</OrderDate>
  <Amount>3400.00</Amount>
</Customers>
<Customers Name="Customer A" Country="Australia">
  <OrderId>2</OrderId>
  <OrderDate>2001-01-01T00:00:00</OrderDate>
  <Amount>4300.00</Amount>
</Customers>
```

Which Transact-SQL query should you use?

- A. `SELECT OrderId, OrderDate, Amount, Name, Country FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId WHERE Customers.CustomerId = 1 FOR XML RAW`

- B. `SELECT OrderId, OrderDate, Amount, Name, Country FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId WHERE Customers.CustomerId = 1 FOR XML RAW, ELEMENTS`
- C. `SELECT OrderId, OrderDate, Amount, Name, Country FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId WHERE Customers.CustomerId = 1 FOR XML AUTO`
- D. `SELECT OrderId, OrderDate, Amount, Name, Country FROM Orders INNER JOIN Customers ON Orders.CustomerId = Customers.CustomerId WHERE Customers.CustomerId= 1 FOR XML AUTO, ELEMENTS`
- E. `SELECT Name, Country, OrderId, OrderDate, Amount FROM Orders INNER JOIN Customers ON Orders.CustomerId= Customers.CustomerId WHERE Customers.CustomerId= 1 FOR XML AUTO`
- F. `SELECT Name, Country, CrderId, OrderDate, Amount FROM Orders INNER JOIN Customers ON Orders.CustomerId= Customers.CustomerId WHERE Customers.CustomerId= 1 FOR XML AUTO, ELEMENTS`
- G. `SELECT Name AS '@Name', Country AS '@Country', OrderId, OrderDate, Amount FROM Orders INNER JOIN Customers ON Orders.CustomerId= Customers.CustomerId WHERE Customers.CustomerId= 1 FOR XML PATH ('Customers')`
- H. `SELECT Name AS 'Customers/Name', Country AS 'Customers/Country', OrderId, OrderDate, Amount FROM Orders INNER JOIN Customers ON Orders.CustomerId= Customers.CustomerId WHERE Customers.CustomerId= 1 FOR XML PATH ('Customers')`

Correct Answer: G

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

QUESTION 39

You use Microsoft SQL Server 2012 to develop a database application. You need to create an object that meets the following requirements:

- Takes an input variable
- Returns a table of values
- Cannot be referenced within a view

Which object should you use?

- A. Scalar-valued function
- B. Inline function
- C. User-defined data type
- D. Stored procedure

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 40

You administer two Microsoft SQL Server 2012 servers named ServerA and ServerB. You use a database named AdventureWorks. You need to prepare the AdventureWorks database for database mirroring. ServerB will act as the mirror in a mirroring partnership along with ServerA. Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲</div><div>▼</div></div>	<div>Backup AdventureWorks on ServerA by using a full backup.</div> <div>Backup AdventureWorks on ServerA by using a full backup followed by a transaction log backup by using the NORECOVERY option.</div> <div>Backup AdventureWorks on ServerA by using a transaction log backup. Restore the transaction log backup by using the RECOVERY option on ServerB.</div> <div>Backup AdventureWorks on ServerA by using a transaction log backup. Restore the transaction log backup by using the NORECOVERY option on ServerB.</div> <div>Restore the full database backup of AdventureWorks by using the NORECOVERY option on ServerB as AdventureWorks.</div> <div>Restore the full database backup of AdventureWorks by using the NORECOVERY option on ServerB as AdventureWorks_Mirror.</div>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

Backup AdventureWorks on ServerA by using a full backup.

Restore the full database backup of AdventureWorks by using the NORECOVERY option on ServerB as AdventureWorks.

Backup AdventureWorks on ServerA by using a transaction log backup. Restore the transaction log backup by using the NORECOVERY option on ServerB.

Section: (none)

Explanation

Explanation/Reference:

I don't think this question will ever come up. The MSDN link says it is being deprecated and to use Availability Groups instead:

<http://msdn.microsoft.com/en-us/library/ms189852.aspx>

However, the answer is correct.

<http://msdn.microsoft.com/en-us/library/ms190941.aspx>

Exam C

QUESTION 1

You administer two Microsoft SQL Server 2012 servers.
Each server resides in a different, untrusted domain.
You plan to configure database mirroring.
You need to be able to create database mirroring endpoints on both servers.
What should you do?

- A. Configure the SQL Server service account to use Network Service.
- B. Use a server certificate.
- C. Use a database certificate.
- D. Configure the SQL Server service account to use Local System.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

Explanation/Reference: valid

QUESTION 2

You administer a Microsoft SQL Server database named Sales. The database is 3 terabytes in size. The Sales database is configured as shown in the following table:

Filegroup	File
PRIMARY	<ul style="list-style-type: none">• Sales.mdf
XACTIONS	<ul style="list-style-type: none">• Sales_1.ndf• Sales_2.ndf• Sales_3.ndf
ARCHIVES	<ul style="list-style-type: none">• SalesArch_1.ndf• SalesArch_2.ndf

You discover that **Sales_2.ndf is corrupt**. You need to recover the corrupted data in the minimum amount of time. What should you do?

- A. Perform a file restore.
- B. Perform a transaction log restore.
- C. Perform a restore from a full backup.
- D. Perform a filegroup restore.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

VALID

According to these references, this answer looks correct.

Reference: <http://technet.microsoft.com/en-us/library/ms187048.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/aa337540.aspx>

Under the simple recovery model, the file must belong to a read-only filegroup.

Under the full or bulk-logged recovery model, before you can restore files, you must back up the active transaction log (known as the tail of the log). For more information, see [Back Up a Transaction Log \(SQL Server\)](#).

To restore a database that is encrypted, you must have access to the certificate or asymmetric key that was used to encrypt the database. Without the certificate or asymmetric key, the database cannot be restored. As a result, the certificate that is used to encrypt the database encryption key must be retained as long as the backup is needed. For more information, see [SQL Server Certificates and Asymmetric Keys](#).

QUESTION 3

You administer a Microsoft SQL Server database named Sales.

The database is 3 terabytes in size.

The Sales database is configured as shown in the following table.

Filegroup	File
PRIMARY	<ul style="list-style-type: none">• Sales.mdf
XACTIONS	<ul style="list-style-type: none">• Sales_1.ndf• Sales_2.ndf• Sales_3.ndf
ARCHIVES	<ul style="list-style-type: none">• SalesArch_1.ndf• SalesArch_2.ndf

You discover that **all files except Sales_2.ndf are corrupt**. You need to recover the corrupted data in the minimum amount of time. What should you do?

- A. Perform a restore from a full backup.
- B. Perform a transaction log restore.
- C. Perform a file restore.
- D. Perform a filegroup restore.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

valid

QUESTION 4

You administer a Microsoft SQL Server 2012 instance that contains a financial database hosted on a storage area network (SAN).

The financial database has the following characteristics:

A data file of 2 terabytes is located on a dedicated LUN (drive D).

A transaction log of 10 GB is located on a dedicated LUN (drive E).

Drive D has 1 terabyte of free disk space.

Drive E has 5 GB of free disk space.

The database is continually modified by users during business hours from Monday through Friday between 09:00 hours and 17:00 hours.

Five percent of the existing data is modified each day.

The Finance department loads large CSV files into a number of tables each business day at 11:15 hours and

15:15 hours by using the BCP or BULK INSERT commands.

Each data load adds 3 GB of data to the database.

These data load operations must occur in the minimum amount of time.

A full database backup is performed every Sunday at 10:00 hours.

Backup operations will be performed every two hours (11:00, 13:00, 15:00, and 17:00) during business hours.

You implement log shipping of the financial database to another SQL Server 2012 instance.

You decide to failover to this secondary database. You need to ensure that all transactions will be replicated to the secondary database.

Which backup option should you use?

- A. Differential
- B. Transaction Log
- C. FULL
- D. SIMPLE
- E. SKIP
- F. RESTART
- G. STANDBY
- H. CHECKSUM
- I. DBO_ONLY
- J. COPY_ONLY
- K. NORECOVERY
- L. NO_CHECKSUM
- M. CONTINUE_AFTER_ERROR
- N. BULK_LOGGED

Correct Answer: K

Section: (none)

Explanation

Explanation/Reference:

VALID

According to these references, this answer looks correct.

References:

<http://technet.microsoft.com/en-us/library/ms187103.aspx>

<http://msdn.microsoft.com/en-us/library/ms191233.aspx>

<http://msdn.microsoft.com/en-us/library/ms178117.aspx>

QUESTION 5

You administer several Microsoft SQL Server 2012 database servers.

Merge replication has been configured for an application that is distributed across offices throughout a wide area network (WAN).

Many of the tables involved in replication use the XML and varchar(max) data types. Occasionally, merge replication fails due to timeout errors.

You need to reduce the occurrence of these timeout errors. What should you do?

- A. Set the Merge agent on the problem subscribers to use the slow link agent profile.
- B. Create a snapshot publication, and reconfigure the problem subscribers to use the snapshot publication.
- C. Change the Merge agent on the problem subscribers to run continuously.
- D. Set the Remote Connection Timeout on the Publisher to 0.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

valid

When replication is configured, a set of agent profiles is installed on the Distributor. An agent profile contains a set of parameters that are used each time an agent runs: each agent logs in to the Distributor during its startup process and queries for the parameters in its profile. For merge subscriptions that use Web synchronization, profiles are downloaded and stored at the Subscriber. If the profile is changed, the profile at the Subscriber is updated the next time the Merge Agent runs. For more information about Web synchronization, see Web Synchronization for Merge Replication.

Replication provides a default profile for each agent and additional predefined profiles for the Log Reader Agent, Distribution Agent, and Merge Agent. In addition to the profiles provided, you can create profiles suited to your application requirements. An agent profile allows you to change key parameters easily for all agents associated with that profile. For example, if you have 20 Snapshot Agents and need to change the query timeout value (the **-QueryTimeout** parameter), you can update the profile used by the Snapshot Agents and all agents of that type will begin using the new value automatically the next time they run. You might also have different profiles for different instances of an agent. For example, a Merge Agent that connects to the Publisher and Distributor over a dialup connection could use a set of parameters that are better suited to the slower communications link by using the **slow link** profile.

QUESTION 6

You administer a Microsoft SQL Server 2012 instance that has several SQL Server Agent jobs configured. When SQL Server Agent jobs fail, the error messages returned by the job steps do not provide the required detail.

The following error message is an example error message:

"The job failed. The Job was invoked by User CONTOSO\ServiceAccount. The last step to run was step 1 (Subplan_1)."

You need to ensure that all available details of the job step failures for SQL Server Agent jobs are retained. What should you do?

- A. Configure output files.
- B. Expand agent logging to include information from all events.
- C. Disable the Limit size of job history log feature.
- D. Configure event forwarding.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

VALID

Reference: <http://msdn.microsoft.com/en-us/library/ms175488.aspx>

QUESTION 7

You administer a Microsoft SQL Server 2012 instance that has several SQL Server Agent jobs configured. When SQL Server Agent jobs fail, the error messages returned by the job steps are truncated.

The following error message is an example of the truncated error message:

"Executed as user CONTOSO\ServiceAccount. ...0.4035.00 for 64-bit Copyright (C) Microsoft Corp 1984-2011. All rights reserved. Started 63513 PM Error 2012-06-23 183536.87 Code 0XC001000E Source UserImport Description Code 0x00000000 Source Log Import Activity Descript... The package execution fa... The step failed."

You need to ensure that all the details of the job step failures are retained for SQL Server Agent jobs. What should you do?

- A. Expand agent logging to include information from all events.
- B. Disable the Limit size of job history log feature.
- C. Configure event forwarding.
- D. Configure output files.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

VALID

QUESTION 8

You administer a Microsoft SQL Server 2012 server that has SQL Server Integration Services (SSIS) installed.

You plan to deploy new SSIS packages to the server.

The SSIS packages use the Project Deployment Model together with parameters and Integration Services environment variables.

You need to configure the SQL Server environment to support these packages.

What should you do?

- A. Create SSIS configuration files for the packages.
- B. Create an Integration Services catalog.
- C. Install Data Quality Services.
- D. Install Master Data services.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

valid

QUESTION 9

You administer a Microsoft SQL Server 2012 database that has multiple tables in the Sales schema.

Some users must be prevented from deleting records in any of the tables in the Sales schema.

You need to manage users who are prevented from deleting records in the Sales schema.

You need to achieve this goal by using the minimum amount of administrative effort. What should you do?

- A. Create a custom database role that includes the users. Deny Delete permissions on the Sales schema for the custom database role.
- B. Include the Sales schema as an owned schema for the db_denydatawriter role. Add the users to the db_denydatawriter role.
- C. Deny Delete permissions on each table in the Sales schema for each user.
- D. Create a custom database role that includes the users. Deny Delete permissions on each table in the Sales schema for the custom database role.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Looks good. valid

QUESTION 10

You administer a Microsoft SQL Server 2012 database.

You provide temporary securityadmin access to User1 to the database server.

You need to know if User1 adds logins to securityadmin. Which server-level audit action group should you use?

- A. SERVER_STATE_CHANGE_GROUP
- B. SERVER_PRINCIPAL_IMPERSONATION_GROUP
- C. SUCCESSFUL_LOGIN_GROUP
- D. SERVER_ROLE_MEMBER_CHANGE_GROUP

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

valid

Verified answer as correct.

Reference: <http://technet.microsoft.com/en-us/library/cc280663.aspx>

SERVER_STATE_CHANGE_GROUP

This event is raised when the SQL Server service state is modified. Equivalent to the Audit Server Starts and Stops Event Class.

SERVER_PRINCIPAL_IMPERSONATION_GROUP

This event is raised when there is an impersonation within server scope, such as EXECUTE AS <login>. Equivalent to the Audit Server Principal Impersonation Event Class.

SUCCESSFUL_LOGIN_GROUP

Indicates that a principal has successfully logged in to SQL Server. Events in this class are raised by new connections or by connections that are reused from a connection pool. Equivalent to the Audit Login Event Class.

SERVER_ROLE_MEMBER_CHANGE_GROUP

This event is raised whenever a login is added or removed from a fixed server role. This event is raised for the sp_addsrvrolemember and sp_dropsrvrolemember stored procedures. Equivalent to the Audit Add Login to Server Role Event Class.

QUESTION 11

You administer a Microsoft SQL Server 2012 database.

Users report that an application that accesses the database displays an error, but the error does not provide meaningful information.

No entries are found in the SQL Server log or Windows event logs related to the error.

You need to identify the root cause of the issue by retrieving the error message.

What should you do?

- A. Create an Extended Events session by using the sqlserver.error_reported event.
- B. Create a SQL Profiler session to capture all ErrorLog and EventLog events.
- C. Flag all stored procedures for recompilation by using sp_recompile.
- D. Execute sp_who.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

valid

Reference: <http://msdn.microsoft.com/en-us/library/bb630282.aspx>

Event **sqlserver.error_reported**: This event gets fired every time that an error happens in the server

QUESTION 12

You administer a Microsoft SQL Server 2012 failover cluster that contains two nodes named Node A and Node B.

A single instance of SQL Server is installed on the cluster.

An additional node named Node C has been added to the existing cluster.

You need to ensure that the SQL Server instance can use all nodes of the cluster. What should you do?

- A. Run the New SQL Server stand-alone installation Wizard on Node C.
- B. Run the Add Node to SQL Server Failover Cluster Wizard on Node C.
- C. Use Node B to install SQL Server on Node C.
- D. Use Node A to install SQL Server on Node C.

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

valid

Verified answer as correct.

Reference: <http://technet.microsoft.com/en-us/library/ms191545.aspx>

To add a node to an existing SQL Server failover cluster, you must run SQL Server Setup on the node that is to be added to the SQL Server failover cluster instance. Do not run Setup on the active node.

The Installation Wizard will launch the SQL Server Installation Center. To add a node to an existing failover cluster instance, click Installation in the left-hand pane. Then, select **Add node to a SQL Server failover cluster**.

QUESTION 13

You maintain several databases on a 32-bit Microsoft SQL Server 2005 instance on a Windows Server 2008 R2 64-bit server.

You need to migrate the databases to a 64-bit SQL Server 2012 instance on the same server. You also need to ensure that the new Transact-SQL functionality in SQL Server 2012 can be used in the database after the migration.

What should you do? (Each correct answer presents part of the solution. Choose all that apply.)

- A. Perform a side-by-side installation of a 32-bit SQL Server 2012 instance.
- B. Perform a side-by-side installation of a 64-bit SQL Server 2012 instance.
- C. Perform an in-place upgrade to 64-bit SQL Server 2012.
- D. Detach the database from the old instance and attach it to the new instance.
- E. Change the compatibility level of the database.

Correct Answer: BDE

Section: (none)

Explanation

Explanation/Reference:

VALID

* use detach and attach operations to upgrade a SQL Server 2005, SQL Server 2008 or SQL Server 2008 R2 database in SQL Server 2012. After being attached to SQL Server 2012, the database is available immediately and is automatically upgraded.

QUESTION 14

You administer a Microsoft SQL Server 2012 instance.

You discover that the SQL Agent Error Log file is rapidly growing in size.

You need to ensure that the SQL Agent Error Log file does not grow rapidly when SQL Server agent jobs execute.

What should you do?

- A. Execute the sp_cycle_agent_errorlog stored procedure.
- B. Configure event forwarding.
- C. Enable the Auto Shrink option on the master database.
- D. Enable the Auto Shrink option on the msdb database.
- E. Disable the Include execution trace messages feature.

Correct Answer: E

Section: (none)

Explanation

Explanation/Reference:

QUESTION 15

You administer a Microsoft SQL Server 2012 database. You create an availability group named haContosoDbs. Your primary replica is available at Server01\Contoso01. You need to configure the availability group to have the highest availability. You also need to ensure that no data is lost. Which Transact-SQL statement should you use?

- ☐ A. `ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = ASYNCHRONOUS_COMMIT,`
- ☐ B. `ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = SYNCHRONOUS_COMMIT,`
- ☐ C. `ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = SYNCHRONOUS_COMMIT,`
- ☐ D. `ALTER AVAILABILITY GROUP haContosoDbs MODIFY REPLICA ON 'Server01\Contoso01' WITH (AVAILABILITY_MODE = ASYNCHRONOUS_COMMIT,`

- A. Option A
- B. Option B
- C. Option C
- D. Option D

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

VALID

QUESTION 16

You use Microsoft SQL Server 2012 to write code for a transaction that contains several statements. There is high contention between readers and writers on several tables used by your transaction. You need to minimize the use of the tempdb space. You also need to prevent reading queries from blocking writing queries. Which isolation level should you use?

- ☐ A. SERIALIZABLE
- ☐ B. SNAPSHOT
- ☐ C. READ COMMITTED SNAPSHOT
- ☐ D. REPEATABLE READ

- A.
- B.
- C.
- D.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 17

You use a contained database named ContosoDb within a domain. You need to create a user who can log on to the ContosoDb database. You also need to ensure that you can port the database to different database servers within the domain without additional user account configurations. Which type of user should you create?

- A. SQL user without login
- B. User mapped to an asymmetric key
- C. Domain user
- D. login mapped to a virtual account

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

VALID

QUESTION 18

You administer a Microsoft SQL Server 2012 server that hosts a transactional database and a reporting database.

The transactional database is updated through a web application and is operational throughout the day.

The reporting database is only updated from the transactional database.

The recovery model and backup schedule are configured as shown in the following table:

Database	Description
Transactional database	Recovery model: <ul style="list-style-type: none"> • Full Backup schedule: <ul style="list-style-type: none"> • Full database backup: midnight, daily • Differential database backup: on the hour, every two hours starting at 02:00 hours except at 00:00 hours • Log backup: every half hour, except at the times of full and differential backups
Reporting database	Recovery model: <ul style="list-style-type: none"> • Simple Backup schedule: <ul style="list-style-type: none"> • Full database backup: 01:00 hours daily • Differential database backup: 13:00 hours daily Data updates: <ul style="list-style-type: none"> • Changes in data are updated from the transactional database to the reporting database at 00:30 hours and at 12:30 hours • The update takes 15 minutes

At 16:20 hours, you discover that pages 17, 137, and 205 on one of the database files are corrupted on the **transactional** database.

You need to ensure that the transactional database is restored.

You also need to ensure that data loss is minimal.

What should you do?

- A. Perform a partial restore.
- B. Restore the latest full backup, and restore the latest differential backup. Then, restore each log backup taken before the time of failure from the most recent differential backup.
- C. Perform a point-in-time restore.
- D. Restore the latest full backup.
- E. Restore the latest full backup, and restore the latest differential backup. Then, restore the latest log backup.
- F. Perform a page restore.
- G. Restore the latest full backup. Then, restore each differential backup taken before the time of failure from the most recent full backup.
- H. Restore the latest full backup. Then, restore the latest differential backup.

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

valid

Requirements for Restoring Pages

A page restore is subject to the following requirements:

The databases must be using the full or bulk-logged recovery model. Some issues exist if you are using the bulk-logged model. For more information, see the following section.

Pages in read-only filegroups cannot be restored. Trying to make a filegroup read-only will fail if there is a page restore going on at the same time in the filegroup.

The restore sequence must start with a full, file, or filegroup backup.

A page restore requires an unbroken chain of log backups up to the current log file, and they must all be applied so that the page is brought up to date with the current log file.

As in a file-restore sequence, in each restore step, you can add more pages to the roll forward set.

A database backup and page restore cannot be run at the same time.

Bulk-logged Recovery Model and Page Restore

For a database that uses the bulk-logged recovery model, page restore has the following additional conditions:

Backing up while filegroup or page data is offline is problematic for bulk-logged data, because the offline data is not recorded in the log. Any offline page can prevent backing up the log. In this cases, consider using DBCC REPAIR, because this might cause less data loss than restoring to the most recent backup.

If a log backup of a bulk-logged database encounters a bad page, it fails unless WITH

CONTINUE_AFTER_ERROR is specified.

Page restore generally does not work with bulk-logged recovery.

A best practice for performing page restore is to set the database to the full recovery model, and try a log backup. If the log backup works, you can continue with the page restore. If the log backup fails, you either have to lose work since the previous log backup or you have to try running DBCC must be run with the REPAIR_ALLOW_DATA_LOSS option.

QUESTION 19

You administer a Microsoft SQL Server 2012 server that hosts a transactional database and a reporting database. The transactional database is updated through a web application and is operational throughout the day. The reporting database is only updated from the transactional database. The recovery model and backup schedule are configured as shown in the following table:

Database	Description
Transactional database	Recovery model: <ul style="list-style-type: none"> Full Backup schedule: <ul style="list-style-type: none"> Full database backup: midnight, daily Differential database backup: on the hour, every two hours starting at 02:00 hours except at 00:00 hours Log backup: every half hour, except at the times of full and differential backups
Reporting database	Recovery model: <ul style="list-style-type: none"> Simple Backup schedule: <ul style="list-style-type: none"> Full database backup: 01:00 hours daily Differential database backup: 13:00 hours daily Data updates: <ul style="list-style-type: none"> Changes in data are updated from the transactional database to the reporting database at 00:30 hours and at 12:30 hours The update takes 15 minutes

One of the hard disk drives that stores the reporting database fails at 16:40 hours. You need to ensure that the reporting database is restored. You also need to ensure that data loss is minimal. What should you do?

- Restore the latest full backup. Then, restore each differential backup taken before the time of failure from the most recent full backup.
- Perform a partial restore.
- Restore the latest full backup, and restore the latest differential backup. Then, restore the latest log backup.
- Perform a point-in-time restore.
- Restore the latest full backup.
- Perform a page restore.
- Restore the latest full backup, and restore the latest differential backup. Then, restore each log backup taken before the time of failure from the most recent differential backup.
- Restore the latest full backup. Then, restore the latest differential backup.

Correct Answer: H

Section: (none)

Explanation

Explanation/Reference:

VALID

QUESTION 20

You administer a Microsoft SQL Server 2012 server that hosts a transactional database and a reporting

database. The transactional database is updated through a web application and is operational throughout the day. The reporting database is only updated from the transactional database. The recovery model and backup schedule are configured as shown in the following table:

Database	Description
Transactional database	<p>Recovery model:</p> <ul style="list-style-type: none">• Full <p>Backup schedule:</p> <ul style="list-style-type: none">• Full database backup: midnight, daily• Differential database backup: on the hour, every two hours starting at 02:00 hours except at 00:00 hours• Log backup: every half hour, except at the times of full and differential backups
Reporting database	<p>Recovery model:</p> <ul style="list-style-type: none">• Simple <p>Backup schedule:</p> <ul style="list-style-type: none">• Full database backup: 01:00 hours daily• Differential database backup: 13:00 hours daily <p>Data updates:</p> <ul style="list-style-type: none">• Changes in data are updated from the transactional database to the reporting database at 00:30 hours and at 12:30 hours• The update takes 15 minutes

At 14:00 hours, you discover that pages 71, 520, and 713 on one of the database files are corrupted on the reporting database. You need to ensure that the databases are restored. You also need to ensure that data loss is minimal. What should you do?

- A. Perform a partial restore.
- B. Restore the latest full backup, and restore the latest differential backup. Then, restore each log backup taken before the time of failure from the most recent differential backup.
- C. Restore the latest full backup.
- D. Restore the latest full backup, and restore the latest differential backup. Then, restore the latest log backup.
- E. Perform a page restore.
- F. Restore the latest full backup. Then, restore each differential backup taken before the time of failure from the most recent full backup.
- G. Perform a point-in-time restore.
- H. Restore the latest full backup. Then, restore the latest differential backup.

Correct Answer: H

Section: (none)

Explanation

Explanation/Reference:
VALID

QUESTION 21
<http://www.ipass4sure.com>

You administer a Microsoft SQL Server 2012 server that hosts a transactional database and a reporting database. The transactional database is updated through a web application and is operational throughout the day. The reporting database is only updated from the transactional database. The recovery model and backup schedule are configured as shown in the following table:

4 <http://www.certarea.com>
70-462

Database	Description
Transactional database	Recovery model: <ul style="list-style-type: none">• Full Backup schedule: <ul style="list-style-type: none">• Full database backup: midnight, daily• Differential database backup: on the hour, every two hours starting at 02:00 hours except at 00:00 hours• Log backup: every half hour, except at the times of full and differential backups
Reporting database	Recovery model: <ul style="list-style-type: none">• Simple Backup schedule: <ul style="list-style-type: none">• Full database backup: 01:00 hours daily• Differential database backup: 13:00 hours daily Data updates: <ul style="list-style-type: none">• Changes in data are updated from the transactional database to the reporting database at 00:30 hours and at 12:30 hours• The update takes 15 minutes

The differential backup of the reporting database fails. Then, the reporting database fails at 14:00 hours. You need to ensure that the reporting database is restored. You also need to ensure that data loss is minimal. What should you do?

- A. Restore the latest full backup, and restore the latest differential backup. Then, restore the latest log backup.
- B. Perform a point-in-time restore.
- C. Restore the latest full backup, and restore the latest differential backup. Then, restore each log backup taken before the time of failure from the most recent differential backup.
- D. Restore the latest full backup. Then, restore the latest differential backup.
- E. Restore the latest full backup. Then, restore each differential backup taken before the time of failure from

- the most recent full backup.
- F. Perform a page restore.
 - G. Perform a partial restore.
 - H. Restore the latest full backup.

Correct Answer: H

Section: (none)

Explanation

Explanation/Reference:

valid

5 <http://www.certarea.com>
70-462

QUESTION 22

DRAG AND DROP

You administer a Microsoft SQL Server database that is used by an application.

Users of the application report performance issues.

You need to choose the appropriate tool for performance-tuning of SQL Server databases.

Which tool or tools should you use? (To answer, drag the appropriate tool or tools to their corresponding task or tasks in the answer area. Each tool may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Tool	Task
SQL Profiler	Generating alerts
System Monitor	Capturing and replaying trace activity
XEvents	Identifying cause of high page splits
	Troubleshooting cause of high page_io latch

Correct Answer:

Tool	Task
SQL Profiler	Generating alerts
System Monitor	Capturing and replaying trace activity
XEvents	Identifying cause of high page splits
	Troubleshooting cause of high page_io latch

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

Reference: <http://msdn.microsoft.com/en-us/library/bb630282.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms191246.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms181091.aspx>

Exam D

QUESTION 1

You administer a Windows Azure SQL Database database named Inventory that contains a stored procedure named p_AddInventory.

Users need to be able to SELECT from all tables in the database and execute the stored procedure.

You need to grant only the necessary permissions.

What should you do?

- A. Grant EXECUTE permission on p_AddInventory to all users. Grant VIEW DEFINITION to all users.
- B. Grant EXECUTE permission on p_AddInventory to all users. Add all users to the db_datawriter role.
- C. Add all users to the db_owner role.
- D. Grant EXECUTE permission on p_AddInventory to all users. Add all users to the db_datareader role.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 2

You administer a Windows Azure SQL Database database named Human_Resources. The database contains 2 tables named Employees and SalaryDetails.

You add two Windows groups as logins for the server:

- CORP\Employees - All company employees
- CORP\HRAdmins - HR administrators only
- HR Administrators are also company employees.

You need to grant users access according to the following requirements:

- CORP\Employees should have SELECT access to the Employees table.
- Only users in CORP\HRAdmins should have SELECT access to the SalaryDetails table.
- Logins are based only on Windows security groups.

What should you do?

- A. Create a database role called Employees.
Add CORP\Employees to the db_datareader role.
Add all company employees except HR administrators to the Employees role.
Deny SELECT access to the SalaryDetails table to the Employees role.
- B. Create a database role called HRAdmins.
Add all company employees except HR administrators to the db_datareader role, Add all HR administrators to the HRAdmins role.
Grant SELECT access to the SalaryDetails table to the HRAdmins role.
Deny SELECT access to the SalaryDetails table to the db_datareader role.
- C. Create two database roles: Employees and HRAdmins.
Add all company employees to the Employees role.
Add HR administrators to the HRAdmins role.
Grant SELECT access to all tables except SalaryDetails to the Employees role.
Grant SELECT access to the SalaryDetails table to the HRAdmins role.
Deny SELECT access to the SalaryDetails table to the Employees role.
- D. Create a database role called Employees.
Add all HR administrators to the db_datareader role.
Add all company employees to the Employees role.
Grant SELECT access to all tables except the SalaryDetails table to the Employees role.

Deny SELECT access to the SalaryDetails table to the Employees role.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 3

You administer a Microsoft SQL Server 2012 database. The database is currently configured to log ship to a secondary server.

You are preparing to cut over to the secondary server by stopping log-shipping and bringing the secondary database online. You want to perform a tail-log backup.

You need to leave the primary database in a restoring state.
Which option of the BACKUP LOG command should you use?

- A. NO_TRUNCATE
- B. NORECOVERY
- C. STANDBY
- D. FORMAT

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 4

You administer a Microsoft SQL Server 2012 database. You want to make a full backup of the database to a file on disk.

In doing so, you need to output the progress of the backup.

Which backup option should you use?

- A. STATS
- B. COMPRESSION
- C. CHECKSUM
- D. IN IT

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 5

You administer a Microsoft SQL Server 2012 database named Contoso on a server named Server01.

You need to track all SELECT statements issued in the Contoso database only by users in a role named Sales.

What should you create?

- A. An Alert
- B. A Resource Pool
- C. An Extended Event session
- D. A Server Audit Specification
- E. A SQL Profiler Trace
- F. A Database Audit Specification
- G. A Policy
- H. A Data Collector Set

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

QUESTION 6

You administer a Microsoft SQL Server 2012 database named Contoso that contains a single user-defined database role named BillingUsers.

All objects in Contoso are in the dbo schema.

You need to grant EXECUTE permissions for all stored procedures in Contoso to BillingUsers.

Which Transact-SQL statement should you use?

- A. EXEC sp_addrolemember 'db_procexecutor', 'BillingUsers'
- B. CREATE ROLE proc_caller GRANT EXECUTE ON ALL PROCEDURES TO proc_caller
ALTER MEMBER BillingUsers ADD TO ROLE proc_caller
- C. GRANT EXECUTE ON Schema::dbo TO BillingUsers
- D. GRANT EXECUTE ON Contoso::dbo TO BillingUsers

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 7

You administer a Microsoft SQL Server 2012 database called Human_Resources. The database contains a table named dbo.Salary.

You need to ensure that all read activity against dbo.Salary is audited and written to the Windows Security Log.

What should you do? (To answer, move the three appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Select and Place:

Enable the Audit and the Audit Specification.

Create a new Audit. For destination, choose **File**.

Create a new Audit. For destination, choose **Security Log**.

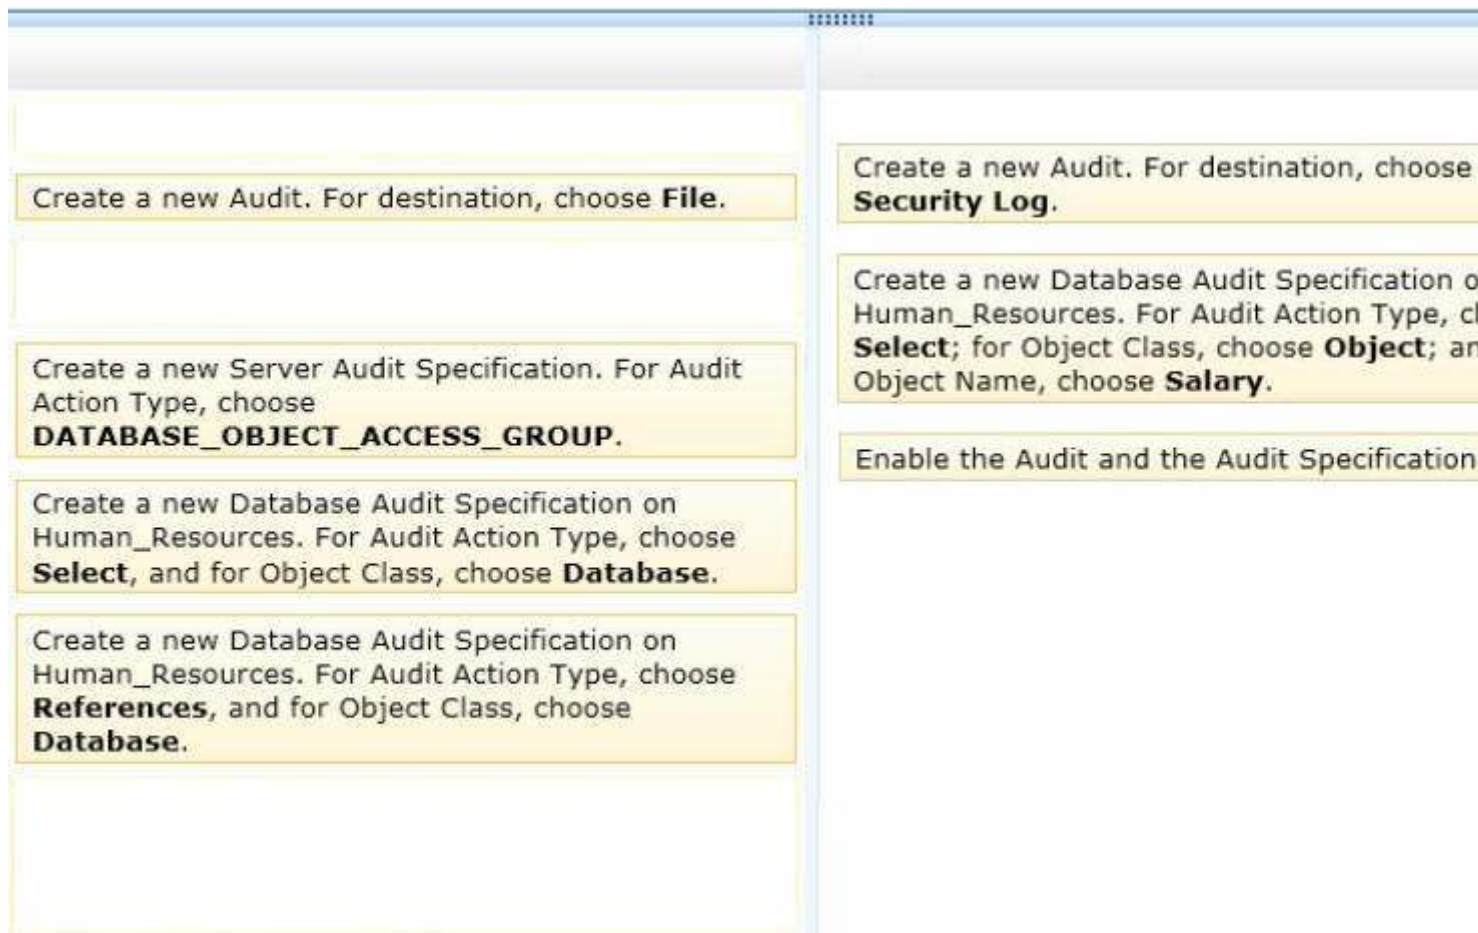
Create a new Server Audit Specification. For Audit Action Type, choose **DATABASE_OBJECT_ACCESS_GROUP**.

Create a new Database Audit Specification on Human_Resources. For Audit Action Type, choose **Select**, and for Object Class, choose **Database**.

Create a new Database Audit Specification on Human_Resources. For Audit Action Type, choose **References**, and for Object Class, choose **Database**.

Create a new Database Audit Specification on Human_Resources. For Audit Action Type, choose **Select**; for Object Class, choose **Object**; and for Object Name, choose **Salary**.

Correct Answer:



Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* An audit must exist before creating a server audit specification for it. When a server audit specification is created, it is in a disabled state.

* The general process for creating and using an audit is as follows.

1. Create an audit and define the target.
2. Create either a server audit specification or database audit specification that maps to the audit. Enable the audit specification.
3. Enable the audit.
4. Read the audit events by using the Windows Event Viewer, Log File Viewer, or the fn_get_audit_file function.

* Example (step 2) (we should audit SELECT on the Salary table):

```
USE AdventureWorks2012 ;
```

```
GO
```

```
-- Create the database audit specification.
```

```
CREATE DATABASE AUDIT SPECIFICATION Audit_Pay_Tables
```

```
FOR SERVER AUDIT Payrole_Security_Audit
ADD (SELECT , INSERT
ON HumanResources.EmployeePayHistory BY dbo )
WITH (STATE = ON) ;
GO
```

QUESTION 8

You administer a Microsoft SQL Server 2012 database.

The database contains a schema named CUSTOMER. CUSTOMER contains several tables and views with sensitive data, as well as various stored procedures and functions.

In order to configure security for CUSTOMER, you need to determine how to meet the security requirements listed in the answer area.

Which command or commands should you use? (To answer, drag the appropriate command or commands to their corresponding security requirement or requirements in the answer area. Answer choices may be used once, more than once, or not at all. Answer targets may be used once or not at all. Additionally, you may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Command	Security Requirement
DENY INSERT,UPDATE,DELETE ON customer.credit_card TO <user>	Allow users to run a stored procedure.
GRANT EXECUTE ON customer.getAddress TO <user>	Do not allow users the ability to change any data in a given table.
GRANT SELECT ON SCHEMA::CUSTOMER TO <user>	Grant users the ability to retrieve data from all tables and views in the schema.
GRANT SELECT ON customer.customer TO <user> WITH GRANT OPTION;	

Correct Answer:

Command	Security Requirement
	Allow users to run a stored procedure.
	Do not allow users the ability to change any data in a given table.
	Grant users the ability to retrieve data from all tables and views in the schema.
GRANT EXECUTE ON customer.getAdd TO <user> WITH GRANT OPTION;	GRANT EXECUTE ON customer.getAdd TO <user>
	DENY INSERT, UPDATE ON customer.credits TO <user>
	GRANT SELECT ON * TO <user>

Section: (none)

Explanation

Explanation/Reference:

QUESTION 9

You administer a Microsoft SQL Server database.

The database is in the Simple recovery mode.

You schedule the following backup plan:

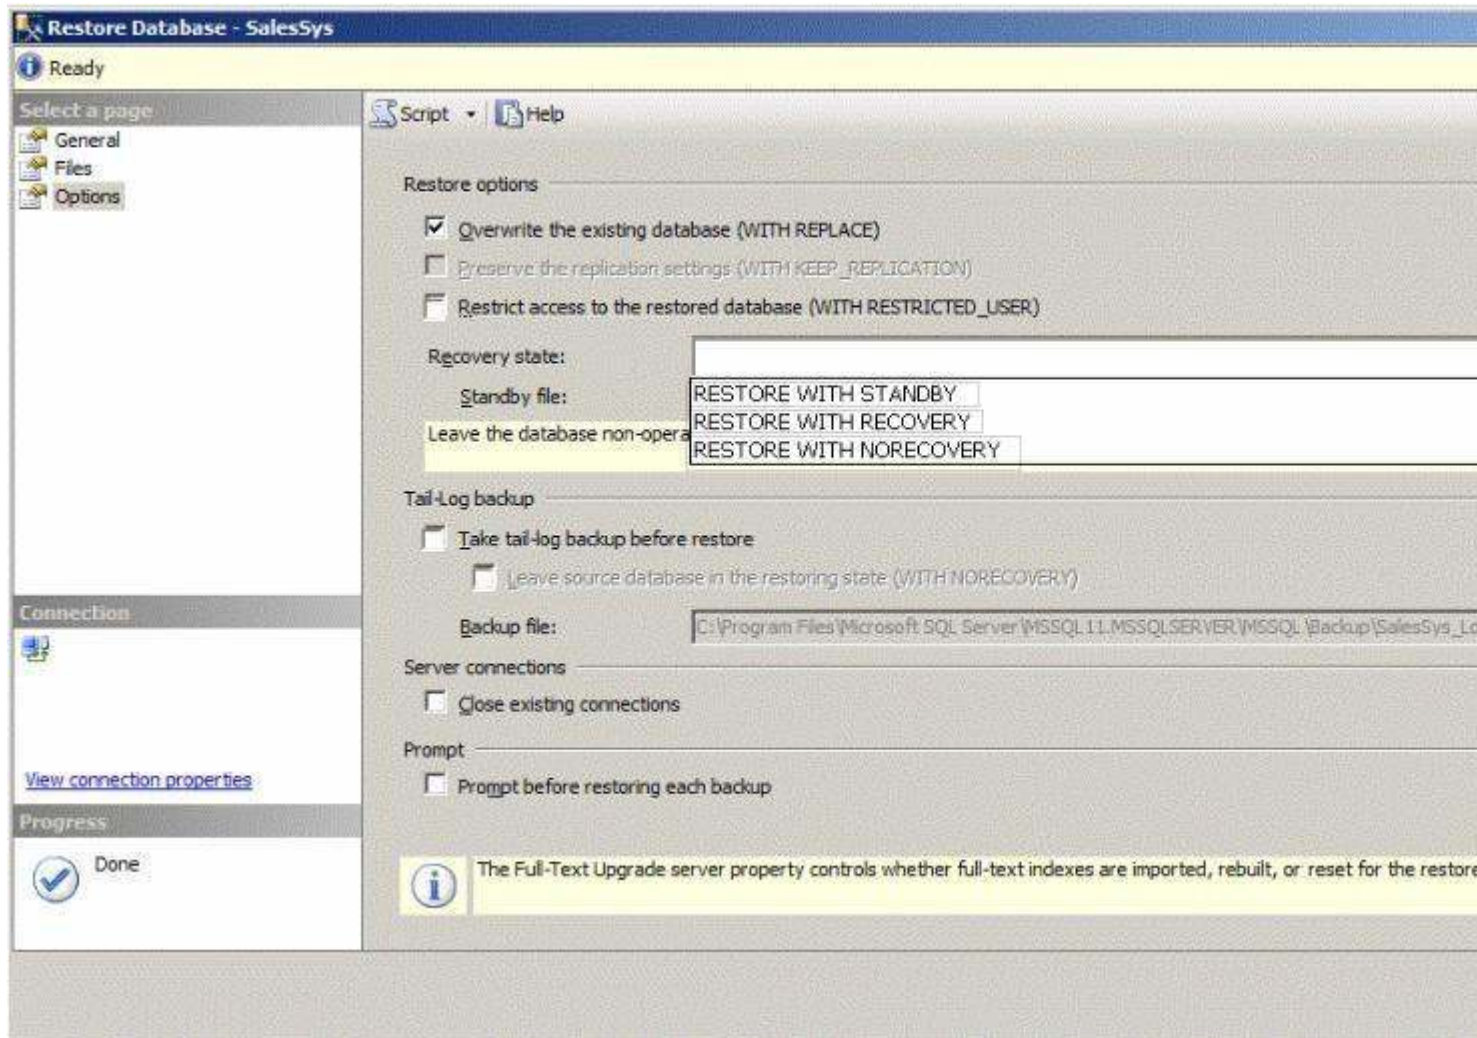
- Full backup every day at midnight
- Differential backups every hour on the hour, except at midnight

The database fails at 20:45 hours.

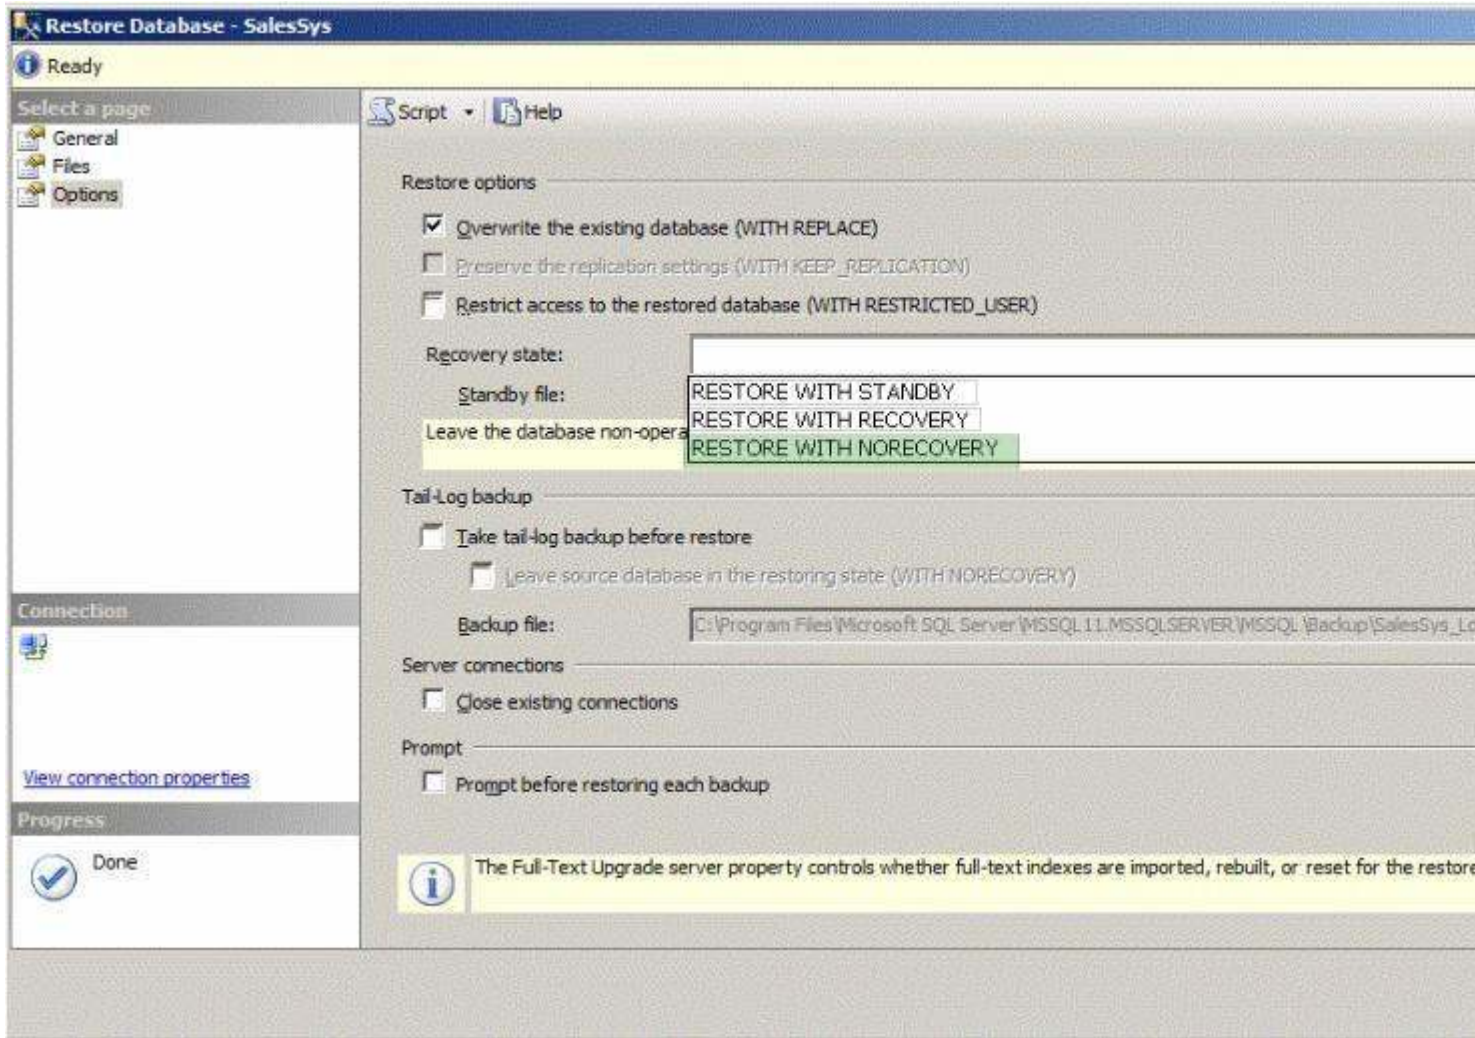
You need to use SQL Server Management Studio (SSMS) to begin restoring the database to a new server instance with a minimum amount of data loss.

Which options should you select on the Options page of the Restore Database window? (To answer, configure the appropriate option or options in the dialog box in the answer area.)

Hot Area:

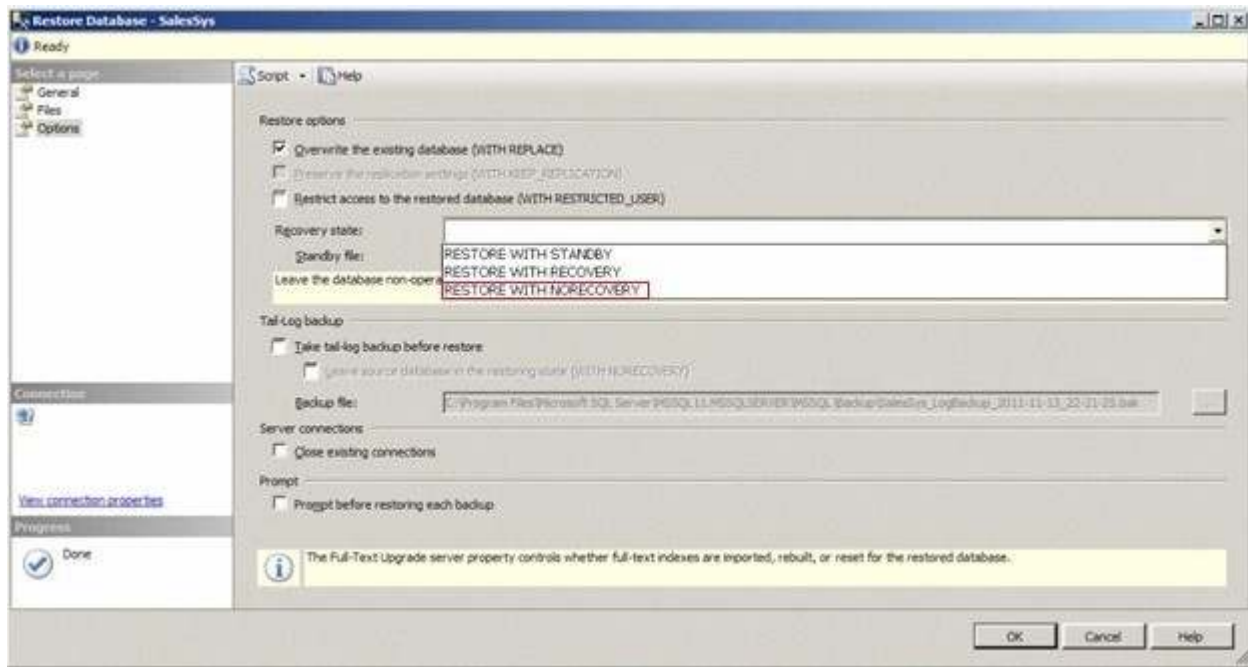


Correct Answer:



Section: (none)
Explanation

Explanation/Reference:
Explanation:



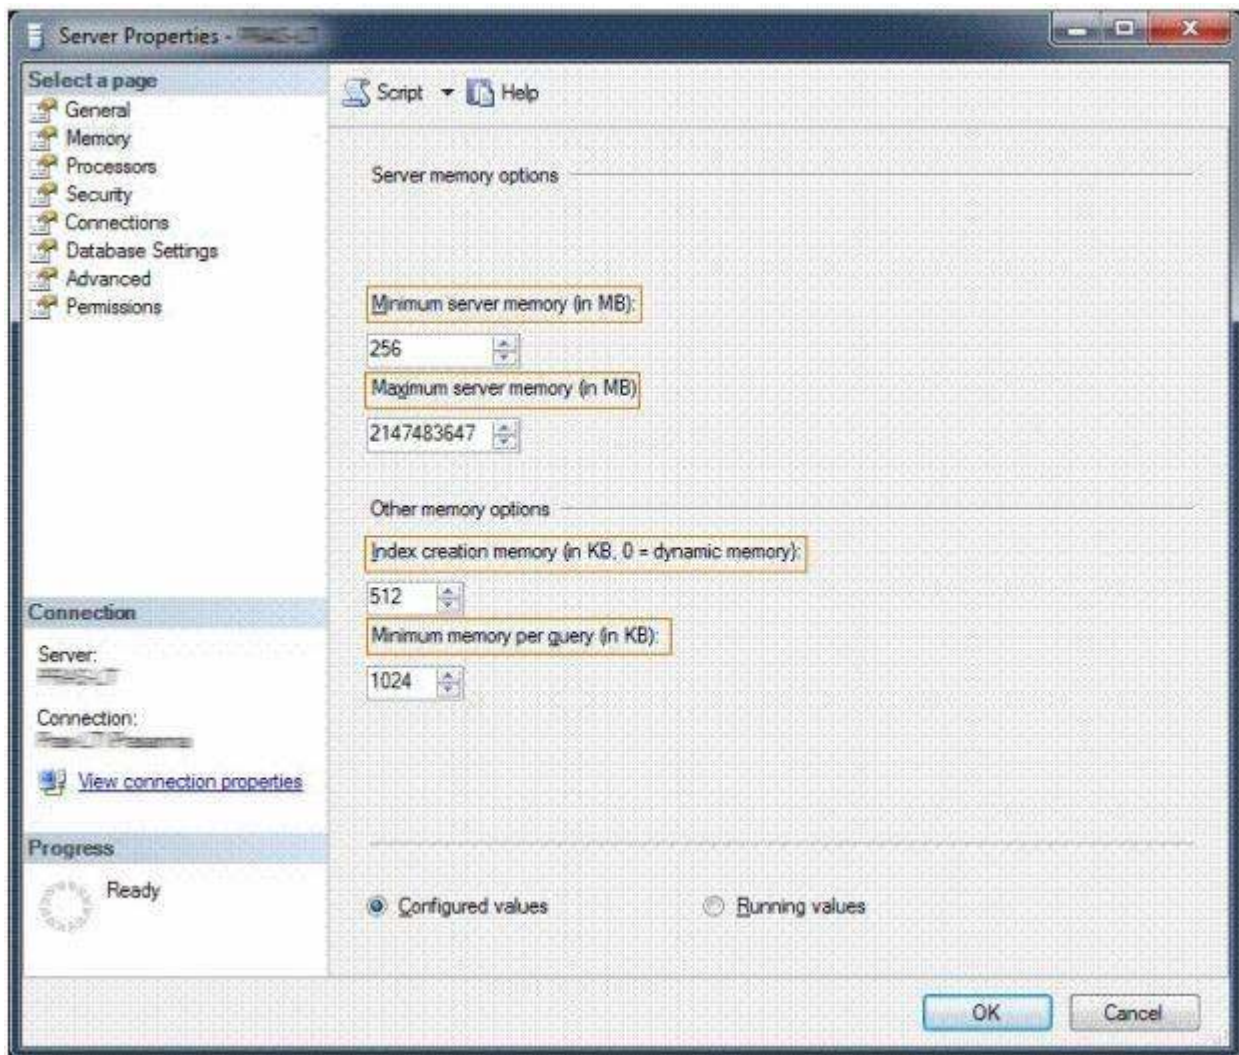
QUESTION 10

You administer a Microsoft SQL Server 2012 database instance. Other applications run on the server. Some of the applications are throwing errors because of insufficient memory.

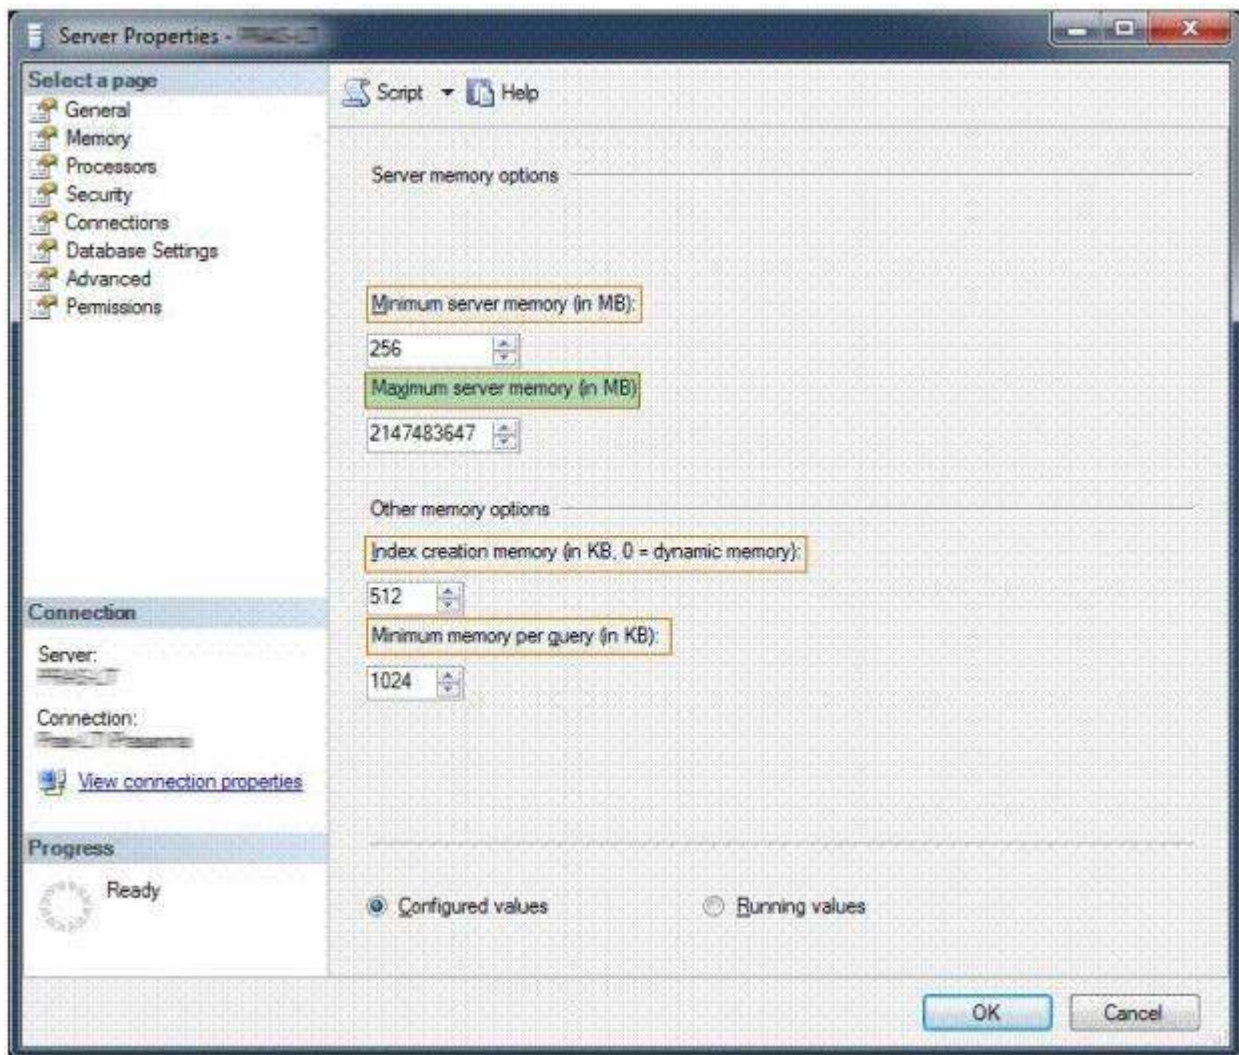
You need to ensure that the other applications have sufficient memory.

Which setting should you configure?

Hot Area:



Correct Answer:

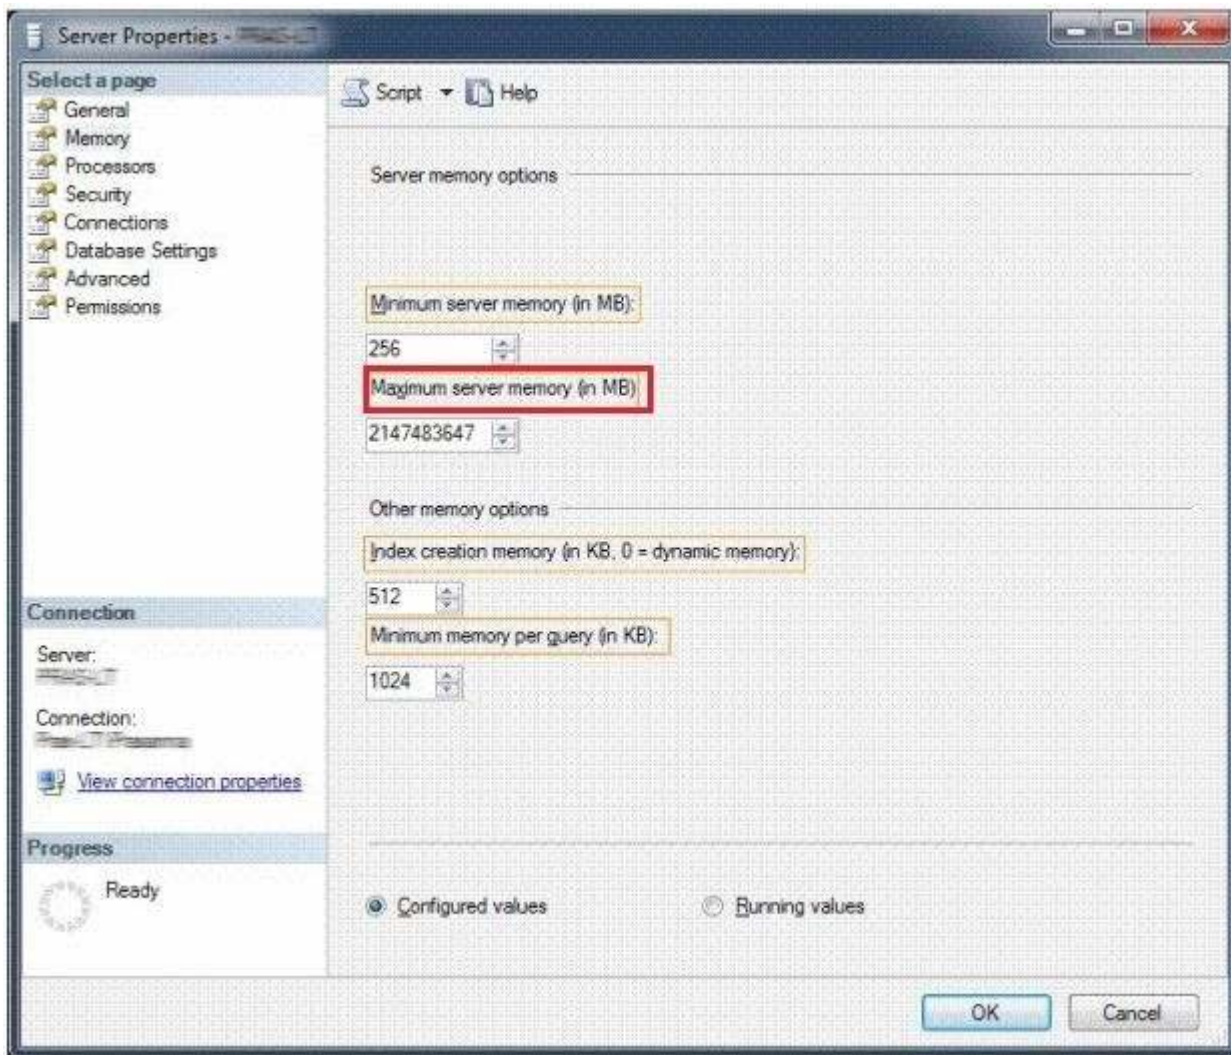


Section: (none)

Explanation

Explanation/Reference:

Explanation:



QUESTION 11

You are a database administrator for a Microsoft SQL Server 2012 database named AdventureWorks2012. You create an Availability Group defined by the following schema. (Line numbers are included for reference only.)

```

01 CREATE AVAILABILITY GROUP Group1
02 FOR DATABASE AdventureWorks2012
03 REPLICA ON 'SecondaryServer'
04 WITH (
05     ENDPOINT_URL = 'TCP://SecondaryServer:5022',
06     ...
07 );

```

You need to implement an AlwaysOnAvailability Group that will meet the following conditions:

- Production transactions should be minimally affected.
- The secondary server should allow reporting queries to be performed.

- If the primary server goes offline, the secondary server should not automatically take over.

Which Transact-SQL statement should you insert at line 06?

- A. `AVAILABILITY_MODE = SYNCHRONOUS_COMMIT,`
`FAILOVER_MODE = MANUAL`
`SECONDARY_ROLE (`
`ALLOW_CONNECTIONS = READ_ONLY,`
`READ_ONLY_ROUTING_URL = 'TCP://SecondaryServer:1433')`
`PRIMARY_ROLE (`
`ALLOW_CONNECTIONS = READ_WRITE,`
`READ_ONLY_ROUTING_LIST = NONE)`
- B. `AVAILABILITY_MODE = SYNCHRONOUS_COMMIT,`
`FAILOVER_MODE = MANUAL`
`SECONDARY_ROLE (`
`ALLOW_CONNECTIONS = READ_ONLY,`
`READ_ONLY_ROUTING_URL = 'TCP://SecondaryServer:1433')`
- C. `AVAILABILITY_MODE = ASYNCHRONOUS_COMMIT,`
`FAILOVER_MODE = MANUAL`
`SECONDARY_ROLE (`
`ALLOW_CONNECTIONS = READ_ONLY,`
`READ_ONLY_ROUTING_URL = 'TCP://SecondaryServer:1433')`
- D. `AVAILABILITY_MODE = ASYNCHRONOUS_COMMIT,`
`FAILOVER_MODE = MANUAL`
`SECONDARY_ROLE (`
`ALLOW_CONNECTIONS = YES,`
`READ_ONLY_ROUTING_URL = 'TCP://SecondaryServer:1433')`

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 12

You are migrating an OLTP database from Windows Azure SQL Database to on-premise. You are planning the installation of a Microsoft SQL Server 2012 server failover cluster. The server has the following partitions:

Disk Subsystem	Storage Location	RAID Level	Capacity
C:	Local	1	100GB
D:	Local SSD	1	128GB
Partition(1)	SAN	10	500GB
Partition(2)	SAN	5	500GB

You install the operating system on the C: drive.

The database solution will use row versioning, triggers, and cursors. The tempdb database cannot be placed on the same disk subsystem as the OLTP database.

The OLTP database needs to be on the fastest disk subsystem possible. The database is currently 200GB in size. It will not substantially increase in size in the next 3 years.

You need to plan the usage of the disk subsystems.

To which hard disk should each database belong? (To answer, drag the appropriate database or databases to

their corresponding hard disk or disks in the answer area. Answer choices may be used once, more than once, or not at all. Answer targets may be used once or not at all. Additionally, you may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Database	Hard Disk
C: drive	SQL Server binaries
D: drive	OLTP database
Partition(1)	tempdb database
Partition(2)	database backups

Correct Answer:

Database	Hard Disk
	SQL Server binaries
	OLTP database
	tempdb database
	database backups

Section: (none)

Explanation

Explanation/Reference:

QUESTION 13

You administer several Microsoft SQL Server 2012 servers.

You want to run scheduled checks to confirm that the databases on the servers are not configured to AutoShrink.

You need to create a policy to check the condition.

What should you do? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Select and Place:

Create a new condition to check that the Database facet SHRINKDB field is set to **TRUE**.

Create a new condition to check that the Database facet @Autoshrink field is set to **FALSE**.

Create a new policy to check the condition. Set the evaluation mode to **On Demand**.

Create a new policy to check the condition. Set the evaluation mode to **Schedule**.

Create a schedule and enable the policy.

Export the schedule, the policy, and the condition to SQL files using the management studio SCRIPT AS... CREATE TO option.

Right-click on the policy in Management Studio and select **Evaluate**.

Correct Answer:

Create a new condition to check that the Database facet SHRINKDB field is set to **TRUE**.

Create a new policy to check the condition. Set the evaluation mode to **Schedule**.

Export the schedule, the policy, and the condition to SQL files using the management studio SCRIPT AS... CREATE TO option.

Right-click on the policy in Management Studio and select **Evaluate**.

Create a new condition to check that the Database facet @Autoshrink field is set to **FALSE**.

Create a new policy to check the condition. Set the evaluation mode to **On Demand**.

Create a schedule and enable the policy.

Section: (none)**Explanation****Explanation/Reference:**

Explanation:

Note:

* @autoshrink is used to enable/disable the autoshrink feature.

* The Evaluation Mode defines the combination of 1) when the policy will be evaluated and 2) what the system will do in case of policy violation. The following evaluation modes are supported:

Evaluation Mode	When to Check	Action upon Violation
On Demand	On demand only	Users can choose to configure the system to comply to the policy
On Schedule	Periodically according to the specified schedule	log
On Change – Log only	When there is a relevant (DDL) change to the DBMS	log
On Change – Prevent	When there is a relevant (DDL) change to the DBMS	Roll back the transaction

QUESTION 14

You administer a Windows Azure SQL Database database used for data warehouse operations.

The database contains a table named OrdersHistory, defined as follows:

```
CREATE TABLE OrdersHistory
(
  OrderID bigint NOT NULL,
  CustomerID int NOT NULL,
  ShippingAddressID int NOT NULL,
  OrderDetails varchar(4000) NULL,
  ShipDate date NULL
)

CREATE NONCLUSTERED INDEX IX_OrdersHistory_OrderID
ON OrdersHistory (OrderID) INCLUDE (CustomerID, OrderDetails)

CREATE NONCLUSTERED INDEX IX_OrdersHistory_CustomerID_OrderID
ON OrdersHistory (CustomerID, OrderID) INCLUDE (OrderDetails)

CREATE NONCLUSTERED INDEX IX_OrdersHistory_Shipdate
ON OrdersHistory (ShipDate) INCLUDE (CustomerID, OrderID)
```

A weekly ETL (extract-transform-load) runs a large INSERT statement to add data into the OrdersHistory table. The process is taking a long time to complete. You discover that the bulk of the process is performing non-clustered index updates on the OrdersHistory table.

You need to improve the performance of the ETL to make the process faster. You need to meet the following requirements:

Avoid losing existing permissions on existing objects.

Use minimal administrative effort.

What should you do? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Select and Place:

	Answers
Change the INSERT statement to a SELECT INTO statement. Run the SELECT INTO statement.	
Run a CREATE INDEX statement for each index.	
Run a DROP INDEX statement for each index.	
Run the INSERT statement.	
Run the following statement: DROP TABLE OrdersHistory.	
Run the following statement: ALTER INDEX ALL on OrdersHistory DISABLE.	
Run the following statement: ALTER INDEX ALL on OrdersHistory REBUILD.	

Correct Answer:

<p>Change the INSERT statement to a SELECT INTO statement. Run the SELECT INTO statement.</p> <p>Run a CREATE INDEX statement for each index.</p>	<p>Run a DROP INDEX statement for each index.</p>
<p>Run the following statement: DROP TABLE OrdersHistory.</p>	<p>Run the following statement: ALTER INDEX OrdersHistory REBUILD.</p>
<p>Run the following statement: ALTER INDEX ALL on OrdersHistory DISABLE.</p>	<p>Run the INSERT statement.</p>

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* Drop and rebuild the index, then run the insert statement.

* Rebuilding an index can be executed online or offline. Reorganizing an index is always executed online. To achieve availability similar to the reorganize option, you should rebuild indexes online.

Incorrect:

* Dropping and rebuilding the OrdersHistory table would lose the permissions on this table.

* The SELECT INTO statement creates a new table and populates it with the result set of the SELECT statement.

QUESTION 15

You administer a Microsoft SQL Server 2012 database named Orders.

Orders contains a table named OrderShip that is defined as follows:

```
CREATE TABLE OrderShip
(OrderID bigint NOT NULL PRIMARY KEY,
 CustomerID int NOT NULL,
 ShipAddress nvarchar(500) NOT NULL,
 CountryCode tinyint NULL)
```

A NULL value represents a domestic order. Ninety percent of the values in CountryCode are NULL.

Customers require a procedure that will return orders for all customers from a specified country. You create a new procedure:

```
CREATE PROCEDURE p_GetIntlOrders
  (@countrycode tinyint)
AS
  SELECT DISTINCT CustomerID, ShipAddress
  FROM OrderShip
  WHERE CountryCode = @countrycode
GO
```

Performance on this procedure is slow.

You need to alter the schema to optimize this query. Objects created must use a minimum amount of resources.

Which Transact-SQL statement should you use?

- A. CREATE NONCLUSTERED INDEX IX_CountryCode ON OrderShip (CountryCode) WHERE CountryCode IS NOT NULL
- B. CREATE STATISTICS ST_CountryCode ON OrderShip (CountryCode) WHERE CountryCode IS NOT NULL
- C. CREATE CLUSTERED INDEX IX_CountryCode ON OrderShip (CountryCode)
- D. CREATE INDEX IX_CountryCode ON OrderShip (CustomerID) WHERE CountryCode IS NOT NULL

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

QUESTION 16

You administer a SQL Server 2012 server that contains a database named SalesDb. SalesDb contains a schema named Customers that has a table named Regions. A user named UserA is a member of a role named Sales. UserA is granted the Select permission on the Regions table. The Sales role is granted the Select permission on the Customers schema. You need to ensure that the following requirements are met:

- The Sales role does not have the Select permission on the Customers schema.
- UserA has the Select permission on the Regions table.

Which Transact-SQL statement should you use?

- A. DENY SELECT ON Object::Regions FROM Sales
- B. DENY SELECT ON Schema::Customers FROM Sales
- C. REVOKE SELECT ON Object::Regions FROM Sales
- D. REVOKE SELECT ON Schema::Customers FROM Sales
- E. DENY SELECT ON Object::Regions FROM UserA
- F. DENY SELECT ON Schema::Customers FROM UserA
- G. REVOKE SELECT ON Object::Regions FROM UserA
- H. REVOKE SELECT ON Schema::Customers FOR UserA
- I. EXEC sp_addrolemember 'Sales', 'UserA'

J. EXEC sp_droprolemember 'Sales', 'UserA'

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

According to these references, this answer looks correct.

References:

<http://msdn.microsoft.com/en-us/library/ms188369.aspx>

<http://msdn.microsoft.com/en-us/library/ms187750.aspx>

<http://msdn.microsoft.com/en-us/library/ff848791.aspx>

QUESTION 17

You develop a Microsoft SQL Server 2012 database that contains a heap named OrdersHistorical.

You write the following Transact-SQL query:

```
INSERT INTO OrdersHistorical  
SELECT * FROM CompletedOrders
```

You need to optimize transaction logging and locking for the statement. Which table hint should you use?

- A. HOLDLOCK
- B. ROWLOCK
- C. XLOCK
- D. UPDLOCK
- E. TABLOCK

Correct Answer: E

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://technet.microsoft.com/en-us/library/ms189857.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms187373.aspx>

QUESTION 18

Your database contains a table named Purchases. The table includes a DATETIME column named PurchaseTime that stores the date and time each purchase is made. There is a non-clustered index on the PurchaseTime column. The business team wants a report that displays the total number of purchases made on the current day. You need to write a query that will return the correct results in the most efficient manner. Which Transact-SQL query should you use?

- A.

```
SELECT COUNT(*)  
FROM Purchases  
WHERE PurchaseTime = CONVERT(DATE, GETDATE())
```
- B.

```
SELECT COUNT(*)  
FROM Purchases  
WHERE PurchaseTime = GETDATE()
```
- C.

```
SELECT COUNT(*)  
FROM Purchases  
WHERE CONVERT(VARCHAR, PurchaseTime, 112) = CONVERT(VARCHAR, GETDATE(), 112)
```
- D.

```
SELECT COUNT(*)  
FROM Purchases  
WHERE PurchaseTime >= CONVERT(DATE, GETDATE())  
AND PurchaseTime < DATEADD(DAY, 1, CONVERT(DATE, GETDATE()))
```

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

Two answers will return the correct results (the "WHERE CONVERT..." and "WHERE ... AND ..." answers). The correct answer for Microsoft would be the answer that is most "efficient". Anybody have a clue as to which is most efficient? In the execution plan, the one that I've selected as the correct answer is the query with the shortest duration. Also, the query answer with "WHERE CONVERT..." threw warnings in the execution plan...something about affecting CardinalityEstimate and SeekPlan.

I also found this article, which leads me to believe that I have the correct answer:
<http://technet.microsoft.com/en-us/library/ms181034.aspx>

QUESTION 19

You develop a database for a travel application. You need to design tables and other database objects. You need to store media files in several tables. Each media file is less than 1 MB in size. The media files will require fast access and will be retrieved frequently. What should you do?

- A. Use the CAST function.
- B. Use the DATE data type.
- C. Use the FORMAT function.
- D. Use an appropriate collation.
- E. Use a user-defined table type.
- F. Use the VARBINARY data type.
- G. Use the DATETIME data type.
- H. Use the DATETIME2 data type.
- I. Use the DATETIMEOFFSET data type.
- J. Use the TODATETIMEOFFSET function.

Correct Answer: F

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms188362.aspx>

QUESTION 20

You develop a database for a travel application. You need to design tables and other database objects. You create a view that displays the dates and times of the airline schedules on a report. You need to display dates and times in several international formats. What should you do?

- A. Use the CAST function.
- B. Use the DATE data type.
- C. Use the FORMAT function.
- D. Use an appropriate collation.
- E. Use a user-defined table type.
- F. Use the VARBINARY data type.
- G. Use the DATETIME data type.
- H. Use the DATETIME2 data type.
- I. Use the DATETIMEOFFSET data type.
- J. Use the TODATETIMEOFFSET function.

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/hh213505.aspx>

QUESTION 21

You have three tables that contain data for vendors, customers, and agents. You create a view that is used to look up telephone numbers for these companies. The view has the following definition:

```
Create view apt.vwCompanyPhoneList
(Source, CompanyID, CompanyNumber,
 LastName, FirstName, BusinessName, Phone)
as

SELECT 'Customer' as Source
    , CustomerID
    , CustomerNumber
    , CustomerLastName
    , CustomerFirstName
    , CustomerBusinessName
    , Phone
FROM apt.Customer
UNION ALL
SELECT 'Agent' as Source
    , AgentID
    , AgentNumber
    , AgentLastName
    , AgentFirstName
    , AgentBusinessName
    , Phone
FROM apt.Agent
UNION ALL
SELECT 'Vendor' as Source
    , VendorID
    , VendorNumber
    , VendorLastName
    , VendorFirstName
    , VendorBusinessName
    , Phone
FROM apt.Vendor
GO
```

You need to ensure that users can update only the phone numbers by using this view. What should you do?

- A. Alter the view. Use the EXPAND VIEWS query hint along with each SELECT statement.
- B. Drop the view. Re-create the view by using the SCHEMABINDING clause, and then create an index on the view.
- C. Create an AFTER UPDATE trigger on the view.

D. Create an INSTEAD OF UPDATE trigger on the view.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms187956.aspx>

QUESTION 22

You develop a Microsoft SQL Server 2012 database that contains tables named Employee and Person. The tables have the following definitions:

```
CREATE TABLE [dbo].[Employee] (
    [PersonId] [bigint] NOT NULL,
    [EmployeeNumber] [nvarchar](15) NOT NULL,
    CONSTRAINT [PK_Employee] PRIMARY KEY CLUSTERED
    (
        [PersonId] ASC
    ) ON [PRIMARY]
) ON [PRIMARY]
GO
```

```
CREATE TABLE [dbo].[Person] (
    [Id] [bigint] NOT NULL,
    [FirstName] [nvarchar](25) NOT NULL,
    [LastName] [nvarchar](25) NOT NULL,
    CONSTRAINT [PK_Person] PRIMARY KEY CLUSTERED
    (
        [Id] ASC
    ) ON [PRIMARY]
) ON [PRIMARY]
GO
```

You create a view named VwEmployee as shown in the following Transact-SQL statement.

```
CREATE VIEW [dbo].[VwEmployee]
AS
SELECT
Employee.EmployeeNumber,
    Person.FirstName,
    Person.LastName,
    Person.Id
FROM Employee
INNER JOIN Person
ON Employee.PersonId = Person.Id
GO
```

Users are able to use single INSERT statements or INSERT...SELECT statements into this view. You need to

ensure that users are able to use a single statement to insert records into both Employee and Person tables by using the VwEmployee view. Which Transact-SQL statement should you use?

- A. CREATE TRIGGER TrgVwEmployee
ON VwEmployee
FOR INSERT
AS
BEGIN
INSERT INTO Person(Id, FirstName, LastName)
SELECT Id, FirstName, LastName, FROM inserted

INSERT INTO Employee(PersonId, EmployeeNumber)
SELECT Id, EmployeeNumber FROM inserted

END
- B. CREATE TRIGGER TrgVwEmployee
ON VwEmployee
INSTEAD OF INSERT
AS
BEGIN
INSERT INTO Person(Id, FirstName, LastName)
SELECT Id, FirstName, LastName, FROM inserted

INSERT INTO Employee(PersonId, EmployeeNumber)
SELECT Id, EmployeeNumber FROM inserted

END
- C. CREATE TRIGGER TrgVwEmployee
ON VwEmployee
INSTEAD OF INSERT
AS
BEGIN

DECLARE @ID INT, @FirstName NVARCHAR(25), @LastName NVARCHAR(25), @PersonID
INT, @EmployeeNumber NVARCHAR(15)

SELECT @ID = ID, @FirstName = FirstName, @LastName = LastName, @EmployeeNumber
= EmployeeNumber
FROM inserted

INSERT INTO Person(Id, FirstName, LastName)
VALUES(@ID, @FirstName, @LastName)

INSERT INTO Employee(PersonID, EmployeeNumber)
VALUES(@PersonID, @EmployeeNumber

End
- D. CREATE TRIGGER TrgVwEmployee
ON VwEmployee
INSTEAD OF INSERT
AS
BEGIN
INSERT INTO Person(Id, FirstName, LastName)
SELECT Id, FirstName, LastName FROM VwEmployee

INSERT INTO Employee(PersonID, EmployeeNumber)
SELECT Id, EmployeeNumber FROM VwEmployee

End

Correct Answer: B

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

QUESTION 23

You develop a Microsoft SQL Server 2012 database. You create a view from the Orders and OrderDetails tables by using the following definition.

```
CREATE VIEW vOrders
WITH SCHEMABINDING
AS
SELECT o.ProductID,
       o.OrderDate,
       SUM(od.UnitPrice * od.OrderQty) AS Amount
FROM OrderDetails AS od INNER JOIN
     Orders AS o ON od.OrderID = o.OrderID
WHERE od.SalesOrderID = o.SalesOrderID
GROUP BY o.OrderDate, o.ProductID
GO
```

You need to improve the performance of the view by persisting data to disk. What should you do?

- A. Create an INSTEAD OF trigger on the view.
- B. Create an AFTER trigger on the view.
- C. Modify the view to use the WITH VIEW_METADATA clause.
- D. Create a clustered index on the view.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms188783.aspx>

QUESTION 24

Your database contains tables named Products and ProductsPriceLog. The Products table contains columns named ProductCode and Price. The ProductsPriceLog table contains columns named ProductCode, OldPrice, and NewPrice. The ProductsPriceLog table stores the previous price in the OldPrice column and the new price in the NewPrice column. You need to increase the values in the Price column of all products in the Products table by 5 percent. You also need to log the changes to the ProductsPriceLog table. Which Transact-SQL query should you use?

- A. UPDATE Products SET Price = Price * 1.05
OUTPUT inserted.ProductCode, deleted.Price, inserted.Price
INTO ProductsPriceLog(ProductCode, OldPrice, NewPrice)
- B. UPDATE Products SET Price = Price * 1.05
OUTPUT inserted.ProductCode, inserted.Price, deleted.Price
INTO ProductsPriceLog(ProductCode, OldPrice, NewPrice)
- C. UPDATE Products SET Price = Price * 1.05
OUTPUT inserted.ProductCode, deleted.Price, inserted.Price *
INTO ProductsPriceLog(ProductCode, OldPrice, NewPrice)

D. UPDATE Products SET Price = Price * 1.05
INSERT INTO ProductsPriceLog (ProductCode, CldPnce, NewPrice;
SELECT ProductCode, Price, Price * 1.05 FROM Products

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

Reference: <http://msdn.microsoft.com/en-us/library/ms177564.aspx>

QUESTION 25

A table named Profits stores the total profit made each year within a territory. The Profits table has columns named Territory, Year, and Profit. You need to create a report that displays the profits made by each territory for each year and its previous year. Which Transact-SQL query should you use?

- A. SELECT Territory, Year, Profit,
LEAD(Profit, 1, 0) OVER (PARTITION BY Territory ORDER BY Year) AS PrevProfit
FROM Profits
- B. SELECT Territory, Year, Profit,
LAG(Profit, 1, 0) OVER (PARTITION BY Year ORDER BY Territory) AS PrevProfit
FROM Profits
- C. SELECT Territory, Year, Profit,
LAG(Profit, 1, 0) OVER (PARTITION BY Territory ORDER BY Year) AS PrevProfit
FROM Profits
- D. SELECT Territory, Year, Profit,
LEAD(Profit, 1, 0) OVER (PARTITION BY Year ORDER BY Territory) AS PrevProfit
FROM Profits

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

Reference: <http://msdn.microsoft.com/en-us/library/hh231256.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/hh213125.aspx>

QUESTION 26

Your database contains a table named SalesOrders. The table includes a DATETIME column named OrderTime that stores the date and time each order is placed. There is a non-clustered index on the OrderTime column. The business team wants a report that displays the total number of orders placed on the current day. You need to write a query that will return the correct results in the most efficient manner. Which Transact-SQL query should you use?

- A. SELECT COUNT(*) FROM SalesOrders
WHERE OrderTime = CONVERT(DATE, GETDATE())
- B. SELECT COUNT(*) FROM SalesOrders
WHERE OrderTime = GETDATE()
- C. SELECT COUNT(*) FROM SalesOrders
WHERE CONVERT(VARCHAR, OrderTime, 112) = CONVERT(VARCHAR, GETDATE(I, 112))
- D. SELECT COUNT(*) FROM SalesOrders
WHERE OrderTime >= CONVERT(DATE, GETDATE())
AND OrderTime < DATEADD(DAY, CONVERT(DATE, GETDATE()))

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

QUESTION 27

You use Microsoft SQL Server 2012 to develop a database application. You create a stored procedure named `dbo.ModifyData` that can modify rows. You need to ensure that when the transaction fails, `dbo.ModifyData` meets the following requirements:

- Does not return an error
- Closes all opened transactions

Which Transact-SQL statement should you use?

- A. `BEGIN TRANSACTION`
`BEGIN TRY`
`EXEC dbo.ModifyData`
`COMMIT TRANSACTION`
`END TRY`
`BEGIN CATCH`
`IF @@ TRANCOUNT = 0`
`ROLLBACK TRANSACTION;`
`END CATCH`
- B. `BEGIN TRANSACTION`
`BEGIN TRY`
`EXEC dbo.ModifyData`
`COMMIT TRANSACTION`
`END TRY`
`BEGIN CATCH`
`IF @@ERROR != 0`
`ROLLBACK TRANSACTION;`
`THROW;`
`END CATCH`
- C. `BEGIN TRANSACTION`
`BEGIN TRY`
`EXEC dbo.ModifyData`
`COMMIT TRANSACTION`
`END TRY`
`BEGIN CATCH`
`IF @@TRANCOUNT = 0`
`ROLLBACK TRANSACTION;`
`THROW;`
`END CATCH`
- D. `BEGIN TRANSACTION`
`BEGIN TRY`
`EXEC dbo.ModifyData`
`COMMIT TRANSACTION`
`END TRY`
`BEGIN CATCH`
`IF @@ERROR != 0`
`ROLLBACK TRANSACTION;`
`END CATCH`

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

QUESTION 28

You are developing a database application by using Microsoft SQL Server 2012. You have a query that runs slower than expected. You need to capture execution plans that will include detailed information on missing indexes recommended by the query optimizer. What should you do?

- A. Add a HASH hint to the query.
- B. Add a LOOP hint to the query.
- C. Add a FORCESEEK hint to the query.
- D. Add an INCLUDE clause to the index.
- E. Add a FORCESCAN hint to the Attach query.
- F. Add a columnstore index to cover the query.
- G. Enable the optimize for ad hoc workloads option.
- H. Cover the unique clustered index with a columnstore index.
- I. Include a SET FORCEPLAN ON statement before you run the query.
- J. Include a SET STATISTICS PROFILE ON statement before you run the query.
- K. Include a SET STATISTICS SHOWPLAN_XML ON statement before you run the query.
- L. Include a SET TRANSACTION ISOLATION LEVEL REPEATABLE READ statement before you run the query.
- M. Include a SET TRANSACTION ISOLATION LEVEL SNAPSHOT statement before you run the query.
- N. Include a SET TRANSACTION ISOLATION LEVEL SERIALIZABLE statement before you run the query.

Correct Answer: K

Section: (none)

Explanation

Explanation/Reference:

QUESTION 29

You are developing a database application by using Microsoft SQL Server 2012. An application that uses a database begins to run slowly. You discover that a large amount of memory is consumed by single-use dynamic queries. You need to reduce procedure cache usage from these statements without creating any additional indexes. What should you do?

- A. Add a HASH hint to the query.
- B. Add a LOOP hint to the query.
- C. Add a FORCESEEK hint to the query.
- D. Add an INCLUDE clause to the index.
- E. Add a FORCESCAN hint to the Attach query.
- F. Add a columnstore index to cover the query.
- G. Enable the optimize for ad hoc workloads option.
- H. Cover the unique clustered index with a columnstore index.
- I. Include a SET FORCEPLAN ON statement before you run the query.
- J. Include a SET STATISTICS PROFILE ON statement before you run the query.
- K. Include a SET STATISTICS SHOWPLAN_XML ON statement before you run the query.
- L. Include a SET TRANSACTION ISOLATION LEVEL REPEATABLE READ statement before you run the query.
- M. Include a SET TRANSACTION ISOLATION LEVEL SNAPSHOT statement before you run the query.
- N. Include a SET TRANSACTION ISOLATION LEVEL SERIALIZABLE statement before you run the query.

Correct Answer: G

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/cc645587.aspx>

QUESTION 30

You create a table that has the StudentCode, SubjectCode, and Marks columns to record mid-year marks for students. The table has marks obtained by 50 students for various subjects. You need to ensure that the top half of the students arranged by their average marks must be given a rank of 1 and the remaining students must be given a rank of 2. Which Transact-SQL query should you use?

- A.

```
SELECT StudentCode as Code,
       RANK() OVER (ORDER BY AVG (Marks) DESC) AS Value
FROM StudentMarks
GROUP BY StudentCode
```
- B.

```
SELECT Id, Name, Marks,
       DENSE_RANK() OVER (ORDER BY Marks DESC) AS Rank
FROM StudentMarks
```
- C.

```
SELECT StudentCode as Code,
       DENSE_RANK() OVER (ORDER BY AVG (Marks) DESC) AS Value
FROM StudentMarks
GROUP BY StudentCode
```
- D.

```
SELECT StudentCode as Code,
       NTILE (2) OVER (ORDER BY AVG (Marks) DESC) AS Value
FROM StudentMarks
GROUP BY StudentCode
```
- E.

```
SELECT StudentCode AS Code,Marks AS Value FROM (
    SELECT StudentCode, Marks AS Marks,
           RANK() OVER (PARTITION BY SubjectCode ORDER BY Marks ASC) AS Rank
    FROM StudentMarks) tmp
WHERE Rank = 1
```
- F.

```
SELECT StudentCode AS Code,Marks AS Value FROM (
    SELECT StudentCode, Marks AS Marks,
           RANK() OVER (PARTITION BY SubjectCode ORDER BY Marks DESC) AS Rank
    FROM StudentMarks) tmp
WHERE Rank = 1
```
- G.

```
SELECT StudentCode AS Code,Marks AS Value FROM (
    SELECT StudentCode, Marks AS Marks,
           RANK () OVER (PARTITION BY StudentCode ORDER BY Marks ASC) AS Rank
    FROM StudentMarks) tmp
WHERE Rank = 1
```
- H.

```
SELECT StudentCode AS Code,Marks AS Value FROM (
    SELECT StudentCode, Marks AS Marks,
           RANXO OVER (PARTITION BY StudentCode ORDER BY Marks DESC) AS Rank
    FROM StudentMarks) tmp
WHERE Rank = 1
```

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 31

You develop a database for a travel application. You need to design tables and other database objects. You create the Airline_Schedules table. You need to store the departure and arrival dates and times of flights along

with time zone information. What should you do?

- A. Use the CAST function.
- B. Use the DATE data type.
- C. Use the FORMAT function.
- D. Use an appropriate collation.
- E. Use a user-defined table type.
- F. Use the VARBINARY data type.
- G. Use the DATETIME data type.
- H. Use the DATETIME2 data type.
- I. Use the DATETIMEOFFSET data type.
- J. Use the TODATETIMEOFFSET function.

Correct Answer: I

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ff848733.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/bb630289.aspx>

QUESTION 32

You develop a database for a travel application. You need to design tables and other database objects. You create a stored procedure. You need to supply the stored procedure with multiple event names and their dates as parameters. What should you do?

- A. Use the CAST function.
- B. Use the DATE data type.
- C. Use the FORMAT function.
- D. Use an appropriate collation.
- E. Use a user-defined table type.
- F. Use the VARBINARY data type.
- G. Use the DATETIME data type.
- H. Use the DATETIME2 data type.
- I. Use the DATETIMEOFFSET data type.
- J. Use the TODATETIMEOFFSET function.

Correct Answer: E

Section: (none)

Explanation

Explanation/Reference:

QUESTION 33

You develop a Microsoft SQL Server 2012 database. The database is used by two web applications that access a table named Products. You want to create an object that will prevent the applications from accessing the table directly while still providing access to the required data. You need to ensure that the following requirements are met:

- Future modifications to the table definition will not affect the applications' ability to access data.
- The new object can accommodate data retrieval and data modification.
- You need to achieve this goal by using the minimum amount of changes to the applications.

What should you create for each application?

- A. Synonyms
- B. Common table expressions
- C. Views
- D. Temporary tables

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms190174.aspx>

QUESTION 34

You want to add a new GUID column named BookGUID to a table named dbo.Book that already contains data. BookGUID will have a constraint to ensure that it always has a value when new rows are inserted into dbo.Book. You need to ensure that the new column is assigned a GUID for existing rows. Which four Transact-SQL statements should you use? (To answer, move the appropriate SQL statements from the list of statements to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div style="border: 1px solid #ccc; height: 200px; margin-top: 10px;"></div>	<div style="border: 1px solid #ccc; padding: 5px; margin-top: 10px;"> newid() newguid() WITH VALUES WITH EXISTING CONSTRAINT CK_BookGuid CHECK CONSTRAINT DF_BookGuid DEFAULT ALTER TABLE dbo.Book ADD BookGuid VARCHAR(10) NOT NULL ALTER TABLE dbo.Book ADD BookGuid Uniqueidentifier NOT NULL </div>
<div style="display: flex; justify-content: center; gap: 20px;"> <div style="border: 1px solid #ccc; padding: 2px 10px; background-color: #f0f0f0;"><< Move</div> <div style="border: 1px solid #ccc; padding: 2px 10px; background-color: #f0f0f0;">Remove >></div> </div>	

Correct Answer:

```

ALTER TABLE dbo.Book
ADD BookGuid Uniqueidentifier NOT NULL
CONSTRAINT DF_BookGuid DEFAULT
newid()
WITH VALUES

```

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct. Actually, in the real world, you don't have to use WITH VALUES at the end of the statement and it works just as well. But because the question specifically states which FOUR TSQL statements

to use, we have to include it.

Reference: <http://msdn.microsoft.com/en-us/library/ms190273.aspx>

QUESTION 35

You create a view based on the following statement:

```
CREATE VIEW dbo.vwItemList
AS
SELECT
    b.BatchID
    , b.MailItemID
    , c.ContractNum
    , c.FirstName + ' ' + c.LastName as ContractName
    , a.Address1
    , a.City + ', ' + a.State + ' ' + a.Zip
FROM BatchLog b
join Contract c on b.MailItemID = c.ContractID
join Address a on a.ContractID = c.ContractID
WHERE
    b.ProcessDate >= dateadd(d, 1, EOMONTH(GETDATE(), -2));
```

You grant the Select permission to User1 for this view. You need to change the view so that it displays only the records that were processed in the month prior to the current month. You need to ensure that after the changes, the view functions correctly for User1. Which four Transact-SQL statements should you use? (To answer, move the appropriate SQL statements from the list of statements to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲</div><div>▼</div></div> <div></div>	<pre>DROP VIEW dbo.wvltemList; GO CREATE VIEW dbo.wvltemList AS ALTER VIEW dbo.wvltemList AS WHERE b.ProcessDate >= dateadd(d, 1, EOMONTH (GETDATE(), -2)) AND b.ProcessDate <= EOMONTH(GETDATE(), -1); WHERE b.ProcessDate >= dateadd(d, 1, EOMONTH (GETDATE(), -2)) AND b.ProcessDate < dateadd(d, 1, EOMONTH (GETDATE(), -1)) SELECT b.BatchID , b.MailltemID , c.ContractNum , c.FirstName + ' ' + c.LastName as ContractName , a.Address1 , a.City + ', ' + a.State + ' ' + a.Zip FROM BacthLog b JOIN Contract c ON b.MailltemID = c.ContractID JOIN Address a ON a.ContractID = c.ContractID GO GRANT SELECT ON SCHEMA::wvltemList TO User1;</pre>
	<div><< Move</div> <div>Remove >></div>

Correct Answer:

```

ALTER VIEW dbo.wvltemList
AS
SELECT
    b.BatchID
    , b.MailltemID
    , c.ContractNum
    , c.FirstName + ' ' + c.LastName as
ContractName
    , a.Address1
    , a.City + ' ' + a.State + ' ' + a.Zip
FROM BacthLog b
JOIN Contract c ON b.MailltemID =
c.ContractID
JOIN Address a ON a.ContractID =
c.ContractID
WHERE b.ProcessDate >= dateadd(d, 1,
EOMONTH (GETDATE(), -2))
AND b.ProcessDate < dateadd(d, 1,
EOMONTH (GETDATE(), -1))

```

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/hh213020.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms186819.aspx>

Reference: <http://msdn.microsoft.com/en-us/library/ms173846.aspx>

QUESTION 36

You use a Microsoft SQL Server 2012 database. You need to create an indexed view within the database for a report that displays Customer Name and the total revenue for that customer. Which four T-SQL statements should you use? (To answer, move the appropriate SQL statements from the list of statements to the answer area and arrange them in the correct order.)

Build List and Reorder:

Ordered List Title	Answer Choices Title
<div><div>▲</div><div>▼</div></div> <div></div>	<div>CREATE VIEW Sales.vwCustomerRevenue AS WITH SCHEMABINDING CREATE VIEW Sales.vwCustomerRevenue WITH SCHEMABINDING AS SELECT O.CustomerID , C.CustomerName , SUM(O.SubTotal) AS CustomerTotal , COUNT_BIG(*) AS RecCount FROM Sales.SalesOrderHeader AS O JOIN Sales.Customer AS C ON C.CustomerID = O.CustomerID GROUP BY O.CustomerID , C.CustomerName GO CREATE UNIQUE CLUSTERED INDEX idx_vwCustomerRevenue ON Sales.vwCustomerRevenue (CustomerID); GO CREATE UNIQUE INDEX idx_vwCustomerRevenue ON Sales.vwCustomerRevenue (CustomerID);</div>

<< Move

Remove >>

Correct Answer:

```
CREATE VIEW Sales.wwCustomerRevenue
WITH SCHEMABINDING
AS
SELECT
O.CustomerID
, C.CustomerName
, SUM(O.SubTotal) AS CustomerTotal
, COUNT_BIG(*) AS RecCount
FROM Sales.SalesOrderHeader AS O
JOIN Sales.Customer AS C ON C.CustomerID
= O.CustomerID
GROUP BY
O.CustomerID
, C.CustomerName
GO
CREATE UNIQUE CLUSTERED INDEX
idx_wwCustomerRevenue
ON Sales.wwCustomerRevenue (CustomerID);
```

Section: (none)

Explanation

Explanation/Reference:

Reference: <http://msdn.microsoft.com/en-us/library/ms191432.aspx>

QUESTION 37

You use Microsoft SQL Server 2012 to develop a database application. You need to implement a computed column that references a lookup table by using an INNER JOIN against another table. What should you do?

- A. Reference a user-defined function within the computed column.
- B. Create a BEFORE trigger that maintains the state of the computed column.
- C. Add a default constraint to the computed column that implements hard-coded values.
- D. Add a default constraint to the computed column that implements hard-coded CASE statements.

Correct Answer: A

Section: (none)


Explanation

Explanation/Reference:

Verified answer as correct.

QUESTION 38

You use a Microsoft SQL Server 2012 database that contains two tables named SalesOrderHeader and SalesOrderDetail. The indexes on the tables are as shown in the exhibit. (Click the Exhibit button.)



You write the following Transact-SQL query:

```
SELECT h.SalesOrderID, h.TotalDue, d.OrderQty
FROM Sales.SalesOrderHeader AS h
     INNER JOIN Sales.SalesOrderDetail AS d
     ON h.SalesOrderID = d.SalesOrderID
WHERE h.TotalDue > 100
AND (d.OrderQty > 5 OR d.LineTotal < 1000.00);
```

You discover that the performance of the query is slow. Analysis of the query plan shows table scans where the estimated rows do not match the actual rows for SalesOrderHeader by using an unexpected index on SalesOrderDetail. You need to improve the performance of the query. What should you do?

- A. Use a FORCESCAN hint in the query.
- B. Add a clustered index on SalesOrderID in SalesOrderHeader.
- C. Use a FORCESEEK hint in the query.
- D. Update statistics on SalesOrderID on both tables.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

References: <http://msdn.microsoft.com/en-us/library/ms187348.aspx>

QUESTION 39

You are a database developer for an application hosted on a Microsoft SQL Server 2012 server. The database contains two tables that have the following definitions:

```
CREATE TABLE Customer
(CustomerID int NOT NULL PRIMARY KEY,
 CustomerName varchar(50) NOT NULL)

CREATE TABLE Orders
(OrderID int NOT NULL PRIMARY KEY,
 CustomerID int NOT NULL FOREIGN KEY REFERENCES Customer (CustomerID),
 OrderAmount money NOT NULL,
 ShippingCountry varchar(50) NOT NULL)
```

Global customers place orders from several countries. You need to view the country from which each customer has placed the most orders. Which Transact-SQL query do you use?

- A.

```
SELECT c.CustomerID, c.CustomerName, o.ShippingCountry
FROM Customer c
INNER JOIN
  (SELECT CustomerID, ShippingCountry,
   RANK() OVER (PARTITION BY CustomerID
    ORDER BY COUNT(OrderAmount) DESC) AS Rnk
  FROM Orders
  GROUP BY CustomerID, ShippingCountry) AS o
ON c.CustomerID = o.CustomerID
WHERE o.Rnk = 1
```
- B.

```
SELECT c.CustomerID, c.CustomerName, o.ShippingCountry
FROM
  (SELECT c.CustomerID, c.CustomerName, o.ShippingCountry,
   RANK() OVER (PARTITION BY CustomerID
    ORDER BY COUNT(o.OrderAmount) ASC) AS Rnk
  FROM Customer c
  INNER JOIN Orders o
  ON c.CustomerID = o.CustomerID
  GROUP BY c.CustomerID, c.CustomerName, o.ShippingCountry) cs
WHERE Rnk = 1
```
- C.

```
SELECT c.CustomerID, c.CustomerName, o.ShippingCountry
FROM Customer c
INNER JOIN
  (SELECT CustomerID, ShippingCountry,
   RANK() OVER (PARTITION BY CustomerID
    ORDER BY OrderAmount DESC) AS Rnk
  FROM Orders
  GROUP BY CustomerID, ShippingCountry) AS o
ON c.CustomerID = o.CustomerID
WHERE o.Rnk = 1
```
- D.

```
SELECT c.CustomerID, c.CustomerName, o.ShippingCountry
FROM Customer c
INNER JOIN
```

```
(SELECT CustomerID, ShippingCountry,  
COUNT(OrderAmount) DESC) AS OrderAmount  
FROM Orders  
GROUP BY CustomerID, ShippingCountry) AS o  
ON c.CustomerID = o.CustomerID  
ORDER BY OrderAmount DESC
```

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

Verified answer as correct.

QUESTION 40

You administer a Microsoft SQL Server 2012 database that includes a table named dbo.Log. This table contains millions of records about user activity in an application.

Records in dbo.Log that are more than 90 days old are purged nightly. When records are purged, table locks are causing contention with inserts.

You need to be able to modify dbo.Log without requiring any changes to the applications that utilize dbo.Log. Which type of solution should you use?

- A. Extended events
- B. Columnstore index
- C. Partitioned tables
- D. Read committed snapshot

Correct Answer: C

Section: (none)

Explanation

Explanation/Reference:

QUESTION 41

You administer a Microsoft SQL Server 2012 server along with a Windows Azure SQL Database database.

For both servers, you need to grant users the ability to create logins and create databases. You need to determine which permission to grant users in each instance.

Which permission or permissions should you choose? (To answer, drag the appropriate permission or permissions to their corresponding instance type or types in the answer area. Answer choices may be used once, more than once, or not at all. Answer targets may be used once or not at all. Additionally, you may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Instance Type	Permissions
local instance	dbmanager
Azure instance	dbcreator
	db_accessadmin
	loginmanager
	securityadmin

Correct Answer:

Instance Type	Permissions
local instance	Azure instance
Azure instance	local instance
	local instance
	Azure instance
	local instance

Section: (none)

Explanation

Explanation/Reference:

QUESTION 42

You administer a Microsoft SQL Server 2012 instance.

You need to configure an existing SQL Authenticated Login that meets the following requirements:

- Enforces password policy
- No password expiration enforcement
- Resets the default database to AdventureWorks database

Which option or options should you choose in the Login Properties of the login for User1? (To answer,

configure the appropriate option or options in the dialog box in the answer areaA.)

Hot Area:

Login Properties - User1

Select a page: General, Server Roles, User Mapping, Securables, Status

Script Help

Login name: User1 Search...

☐ Windows authentication

☒ SQL Server authentication

Password: Confirm password:

☐ Specify old password

Old password:

☐ Enforce password policy

☐ Enforce password expiration

☐ User must change password at next login

☐ Mapped to certificate

☐ Mapped to asymmetric key

☐ Map to Credential

Mapped Credentials

Credential	Provider
Cred	

Add Remove

Cancel

Connection: Server: Connection: View connection properties

Progress: Ready

Default database: Default language:

ReportServer\$SQL2012TempDB
AdventureWorks2008R2
master
model
msdb
ReportServer\$SQL2012
AdventureWorks
tempdb

Correct Answer:

Login Properties - User1

Select a page: General, Server Roles, User Mapping, Securables, Status

Script Help

Login name: User1 Search...

☐ Windows authentication
☒ SQL Server authentication

Password:
 Confirm password:
☐ Specify old password
 Old password:

☒ Enforce password policy
☐ Enforce password expiration
☐ User must change password at next login

☐ Mapped to certificate
☐ Mapped to asymmetric key
☒ Map to Credential

Mapped Credentials

Credential	Provider
Cred	

Add

Remove

Cancel

Default database: ReportServer\$SQL2012TempDB

Default language: AdventureWorks

Connection

Server:

Connection:

[View connection properties](#)

Progress

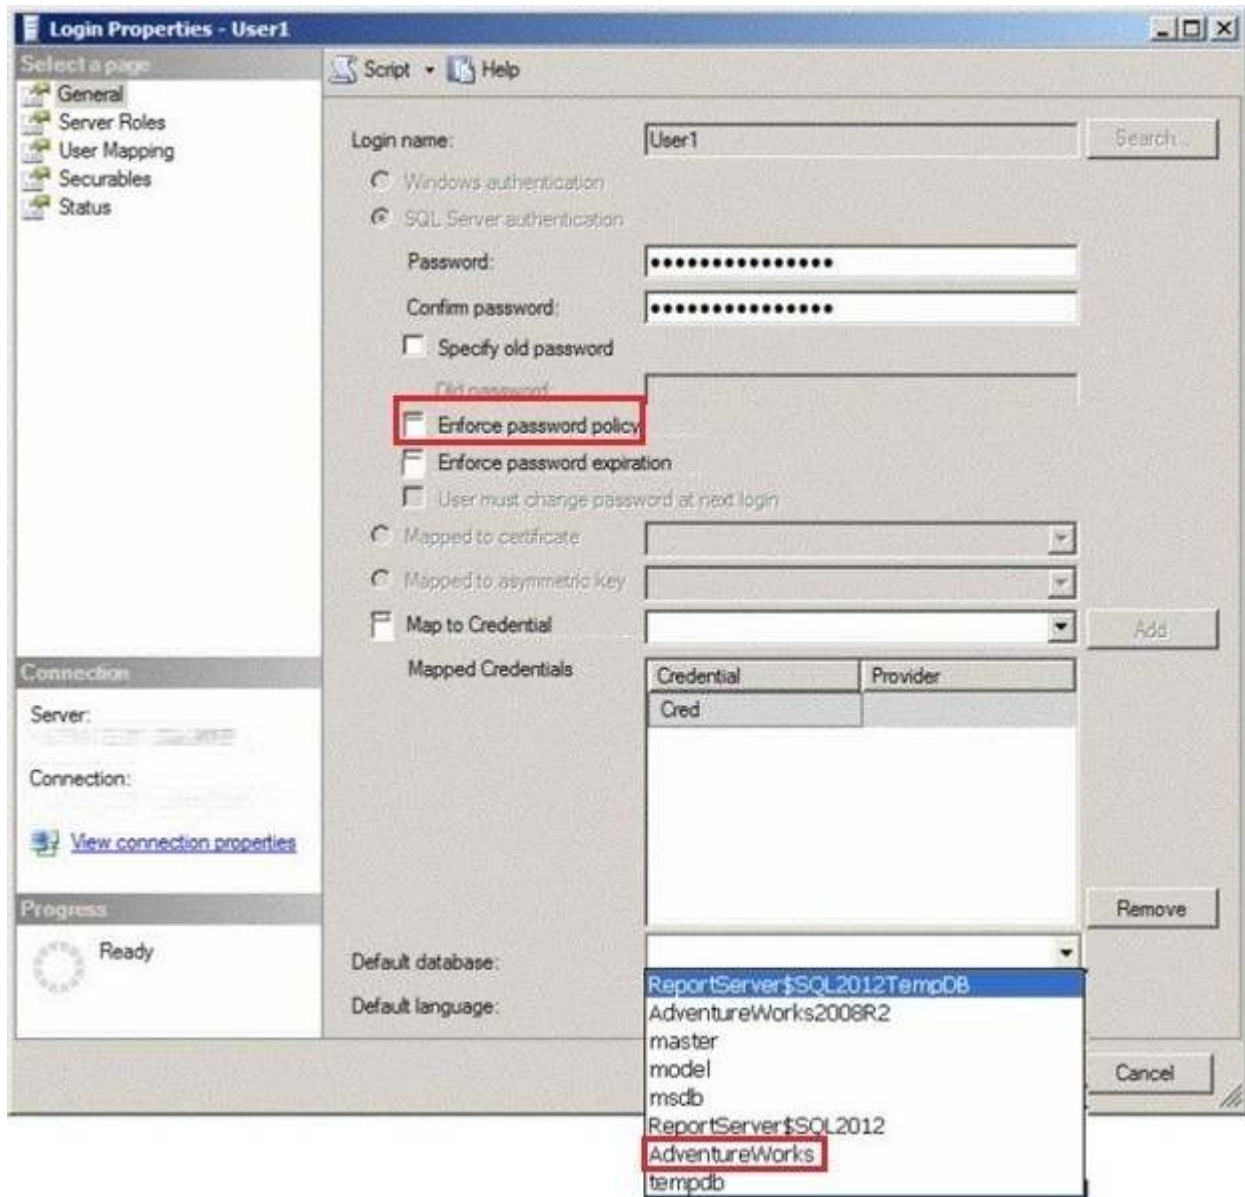
Ready

Section: (none)

Explanation

Explanation/Reference:

Explanation:



QUESTION 43

You administer a database that has User A, B and C.

User A must be able to create new tables and stored procedures, User B must be able to select, update, delete and insert data, User C must be able to provide windows logins to the database

db_accessadmin

db_ddladmin

db_datawriter

db_datareader

db_owner

To which role or roles should you add the users?

- A. ddladmin
- B. datawriter, datareader
- C. accessadmin
- D.

Correct Answer: A

Section: (none)

Explanation

Explanation/Reference:

QUESTION 44

You administer two Microsoft SQL Server 2012 servers named ProdSrv1 and ProdSrv2. ProdSrv1 is configured as a Distributor. Both servers are configured to use the Windows NT Service virtual accounts for all SQL Services.

You are configuring snapshot replication from ProdSrv1 to ProdSrv2 by using ProdSrv2 as a pull subscriber.

The distribution agent on ProdSrv2 regularly fails, displaying the following error message:

"Cannot access the file. Operating system error code 5 (Access is denied.)."

You need to configure the distribution agent by granting only the minimum required access to all accounts.

What should you do?

- A. Configure the Subscriber to use the Local System account.
- B. Configure the SQL Server Agent service to run under the Local System account. Configure the Subscriber to use the SQL Server Agent service account.
- C. Configure the SQL Server Agent service to run under a Windows domain account. Configure the Subscriber to use the SQL Server Agent service account. Grant FULL CONTROL access for the domain account to the ReplData share on ProdSrv1.
- D. Configure the Subscriber to use a Windows domain account. Grant READ access for the domain account to the ReplData share on ProdSrv1.

Correct Answer: D

Section: (none)

Explanation

Explanation/Reference:

QUESTION 45

You administer a Microsoft SQL Server 2012 database named Contoso on a server named Server01.

You need to collect data for a long period of time to troubleshoot wait statistics when querying Contoso. You also need to ensure minimum impact to the server.

What should you create?

- A. An Alert
- B. A Resource Pool
- C. An Extended Event session
- D. A Server Audit Specification
- E. A SQL Profiler Trace
- F. A Database Audit Specification
- G. A Policy

Correct Answer: C

Section: (none)

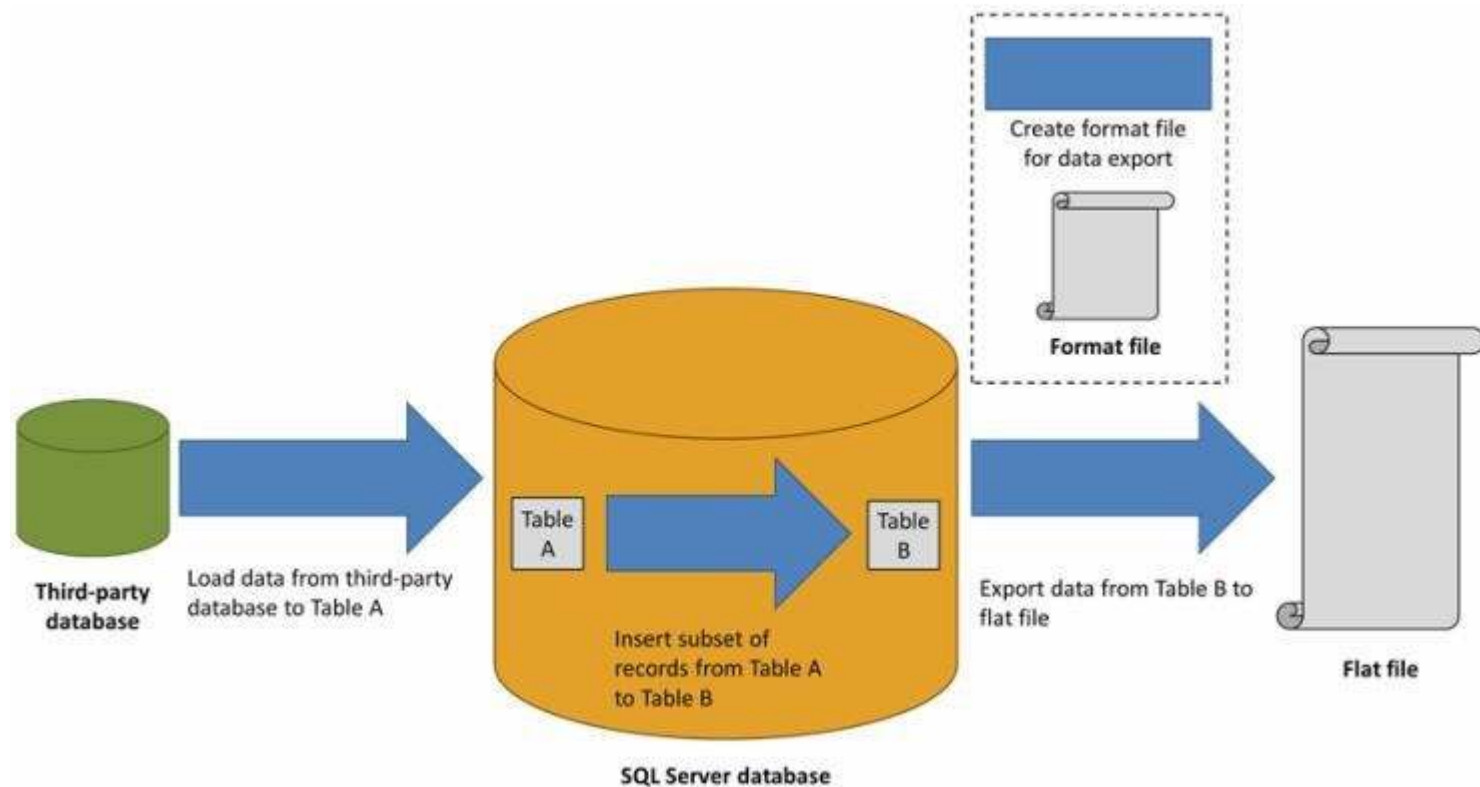
Explanation

Explanation/Reference:

QUESTION 46

You administer a Microsoft SQL Server 2012 database. The database contains two tables named Table A and Table B.

You perform an import and export operation by using the following steps in sequence, as shown in the exhibit. (Click the Exhibit button.)



Load data from a third-party database to Table A. Insert a subset of records from Table A to Table B. Export the data from Table B to a flat file.

You also create a format file for the data export.

You need to choose the correct command or tool for each task. You need to ensure that each task is performed efficiently.

Which command(s) or tool(s) should you use? (To answer, drag the appropriate command(s) or tool(s) to their corresponding task or tasks in the answer area. Answer choices may be used once, more than once, or not at all. Answer targets may be used once or not at all. Additionally, you may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Command/Tool	Task	
INSERT...SELECT	Load data from third-party database to Table A	
INSERT...SELECT... OPENROWSET	Insert subset of records from Table A to Table B	
BCP	Export data from Table B to flat file	
Extended Events	Create format file for data export	
SELECT...INTO		

Correct Answer:

Command/Tool	Task	
INSERT...SELECT	Load data from third-party database to Table A	INSERT...SELECT... OPENROWSET
INSERT...SELECT... OPENROWSET	Insert subset of records from Table A to Table B	INSERT...SELECT
BCP	Export data from Table B to flat file	BCP
Extended Events	Create format file for data export	BCP
SELECT...INTO		

Section: (none)

Explanation

Explanation/Reference:

QUESTION 47

You are a database administrator for a Microsoft SQL Server 2012 environment that contains two SQL Server instances named Server01 and Server02. The Contoso database is located on Server01. Through database mirroring, a copy of the Contoso database is on Server02.

During the last manual failover of database mirroring, the recovery took longer than usual.

You need to determine where the synchronization bottleneck is for the database mirroring topology. Which performance counter or counters should you use for each server? (To answer, drag the appropriate performance counter or counters to their corresponding server or servers in the answer area. Answer choices

may be used once, more than once, or not at all. Answer targets may be used once or not at all. Additionally, you may need to drag the split bar between panes or scroll to view content.)

Select and Place:

Performance Counter	Server
Log Send Queue KB	Server01
Redo Queue KB	Server02
Output Queue Length	
Processor Queue Length	

Correct Answer:

Performance Counter	Server
	Server01
	Server02
Output Queue Length	Log Send Queue KB
Processor Queue Length	Redo Queue KB

Section: (none)

Explanation

Explanation/Reference:

QUESTION 48

You administer a Microsoft SQL Server 2012 instance.

An application executes a large volume of dynamic queries.

You need to reduce the amount of memory used for cached query plans.

What should you do? (To answer, move the appropriate statements from the list of statements to the answer area and arrange them in the correct order.)

Select and Place:


```
EXEC sp_configure 'ad hoc distributed
queries', 1

EXEC sp_configure 'optimize for ad hoc
workloads', 1

EXEC sp_configure 'show advanced options', 1;
RECONFIGURE

EXEC sp_updatestats

EXEC sp_configure 'recovery interval', 75

RECONFIGURE

DBCC DROPCLEANBUFFERS
```

Correct Answer:

```
EXEC sp_configure 'ad hoc distributed
queries', 1

EXEC sp_updatestats

EXEC sp_configure 'recovery interval', 75

DBCC DROPCLEANBUFFERS
```

```
EXEC sp_configure 'show advanced op
RECONFIGURE

EXEC sp_configure 'optimize for ad
workloads', 1

RECONFIGURE
```

Section: (none)

Explanation

Explanation/Reference:

Explanation:

Note:

* SQL SERVER 2008 Optimize for Ad hoc Workloads Advance Performance Optimization

Every batch (T-SQL, SP etc) when ran creates execution plan which is stored in system for re-use. Due to this reason large number of query plans are stored in system. However, there are plenty of plans which are only used once and have never re-used again. One time ran batch plans wastes memory and resources.

* Let us now enable the option of optimizing ad hoc workload. This feature is available in all the versions of SQL Server 2008.

```
sp_CONFIGURE 'show advanced options',1  
RECONFIGURE
```

GO

```
sp_CONFIGURE `optimize for ad hoc workloads`,1  
RECONFIGURE
```

GO

SQL Server 2008 has feature of optimizing ad hoc workloads.

QUESTION 49

You administer two Microsoft SQL Server 2012 databases named Contoso and ContosoWarehouse.

You plan to replicate tables from the Contoso database to the ContosoWarehouse database.

Indexes will be added to the tables in ContosoWarehouse to improve the performance of reports.

You need to ensure that the following requirements are met:

- Indexes are not modified when the subscriptions are reinitialized.
- Only the clustered index from Contoso will be replicated.
- Minimal transaction log activity during the reinitialization.

What should you do? Choose the correct option(s).

Hot Area:

Properties for All Table Articles

Article name: <Default is based on table name>

Properties:

☒ **Copy Objects and Settings to Subscriber**

Copy foreign key constraints	False
Copy check constraints	False
Copy clustered index	True
Copy nonclustered indexes	
Copy default value specifications	True
Copy user triggers	False
Copy extended properties	False
Copy collation	True
Copy INSERT, UPDATE and DELETE	True
Copy unique key constraints	True
Copy file group associations	False
Copy table partitioning schemes	False
Copy index partitioning schemes	False
Copy user-defined statistics	False
Copy default bindings	False
Copy rule bindings	False
Copy full text indexes	False
Copy XML XSD	True
Copy XML indexes	False
Copy XML indexes	True
Copy permissions	False
Copy spatial indexes	
Copy filtered indexes	
Copy data compression attribute	True
Copy sparse column attribute	False

☒ **Destination Object**

Destination object name	<Default is based on table name>
Destination object owner	<source Table owner>
Action if name is in use	
Convert data types	Drop existing object and create a new one
Convert TIMESTAMP to BINARY	Delete data. If article has a row filter, delete only data that matches the filter
Create schemas at Subscriber	Truncate all data in the existing object
Convert XML to NTEXT	False
Convert MAX data types to NTEXT	False
Convert new datetime to NVARCHAR	False

Description
The description of the article.

OK Cancel Help

Correct Answer:

Properties for All Table Articles

Article name:

Properties:

☒ **Copy Objects and Settings to Subscriber**

Copy foreign key constraints	False
Copy check constraints	False
Copy clustered index	True
Copy nonclustered indexes	<input type="text" value=""/>
Copy default value specifications	True
Copy user triggers	False
Copy extended properties	False
Copy collation	True
Copy INSERT, UPDATE and DELETE	True
Copy unique key constraints	True
Copy file group associations	False
Copy table partitioning schemes	False
Copy index partitioning schemes	False
Copy user-defined statistics	False
Copy default bindings	False
Copy rule bindings	False
Copy full text indexes	False
Copy XML XSD	True
Copy XML indexes	False
Copy XML indexes	True
Copy permissions	False
Copy spatial indexes	<input type="text" value=""/>
Copy filtered indexes	<input type="text" value=""/>
Copy data compression attribute	True
Copy sparse column attribute	False

☒ **Destination Object**

Destination object name	<Default is based on table name>
Destination object owner	<source Table owner>
Action if name is in use	<input type="text" value=""/>
Convert data types	Drop existing object and create a new one
Convert TIMESTAMP to BINARY	Delete data. If article has a row filter, delete only data that matches the filter
Create schemas at Subscriber	Truncate all data in the existing object
Convert XML to NTEXT	False
Convert MAX data types to NTEXT	False
Convert new datetime to NVARCHAR	False

Description
The description of the article.

OK Cancel Help

Section: (none)
Explanation

Explanation/Reference:

Explanation:

Properties for All Table Articles

Article name: <Default is based on table name>

Properties:

Copy Objects and Settings to Subscriber

Copy foreign key constraints	False
Copy check constraints	False
Copy clustered index	True
Copy nonclustered indexes	
Copy default value specifications	True
Copy user triggers	False
Copy extended properties	False
Copy collation	True
Copy INSERT, UPDATE and DELETE triggers	True
Copy unique key constraints	True
Copy file group associations	False
Copy table partitioning schemes	False
Copy index partitioning schemes	False
Copy user-defined statistics	False
Copy default bindings	False
Copy rule bindings	False
Copy full text indexes	False
Copy XML XSD	True
Copy XML indexes	False
Copy XML schemas	True
Copy permissions	False
Copy spatial indexes	
Copy filtered indexes	
Copy data compression attribute	True
Copy sparse column attribute	False

Destination Object

Destination object name	<Default is based on table name>
Destination object owner	<source Table owner>
Action if name is in use	
Convert data types	Drop existing object and create a new one
Convert TIMESTAMP to BINARY	Delete data. If article has a row filter, delete only data that matches the filter
Create schemas at Subscriber	Truncate all data in the existing object
Convert XML to NTEXT	False
Convert MAX data types to NTEXT	False
Convert new datetime to NVARCHAR	False

Description

The description of the article.

OK Cancel Help

QUESTION 50

You administer a Microsoft SQL Server 2012 database instance.

The instance is running on a server with the following configuration:

- 1TB RAM
- SAN storage for database and log files
- 4 quad-core processors
- 64-Bit Windows 2008R2 operating system

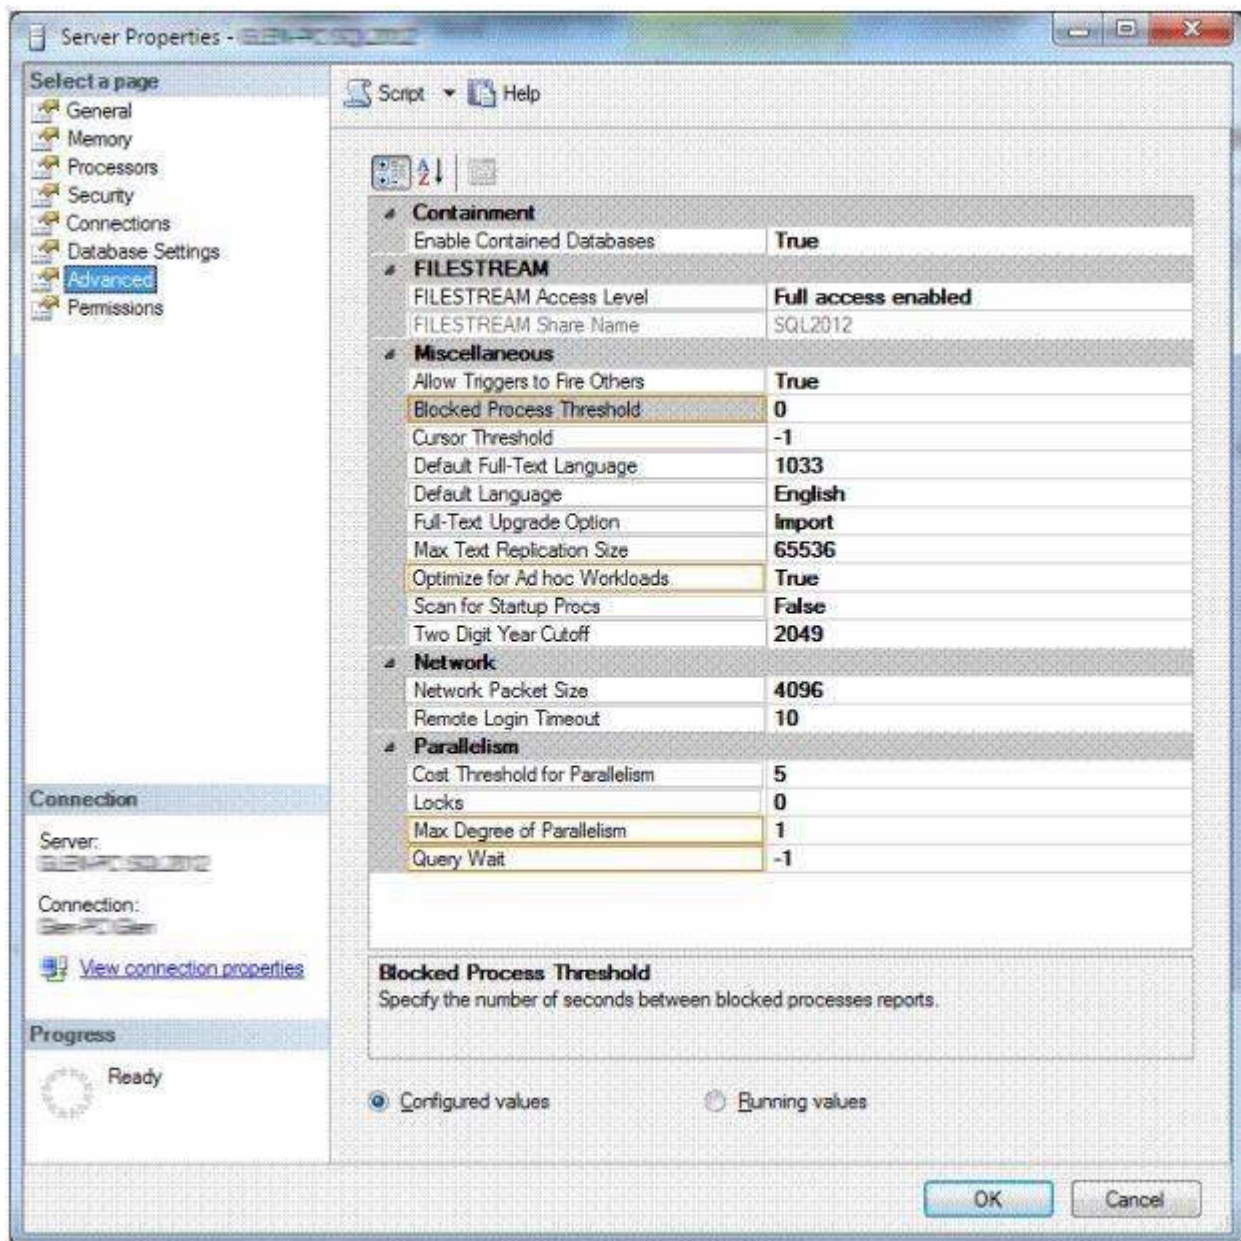
This instance hosts a database with large partitioned tables. Users report that complex queries are taking a long time to complete.

While troubleshooting, you discover that CPU utilization is low (less than 20 percent), disk activity is low (little or no waiting processes), and no significant blocking is occurring.

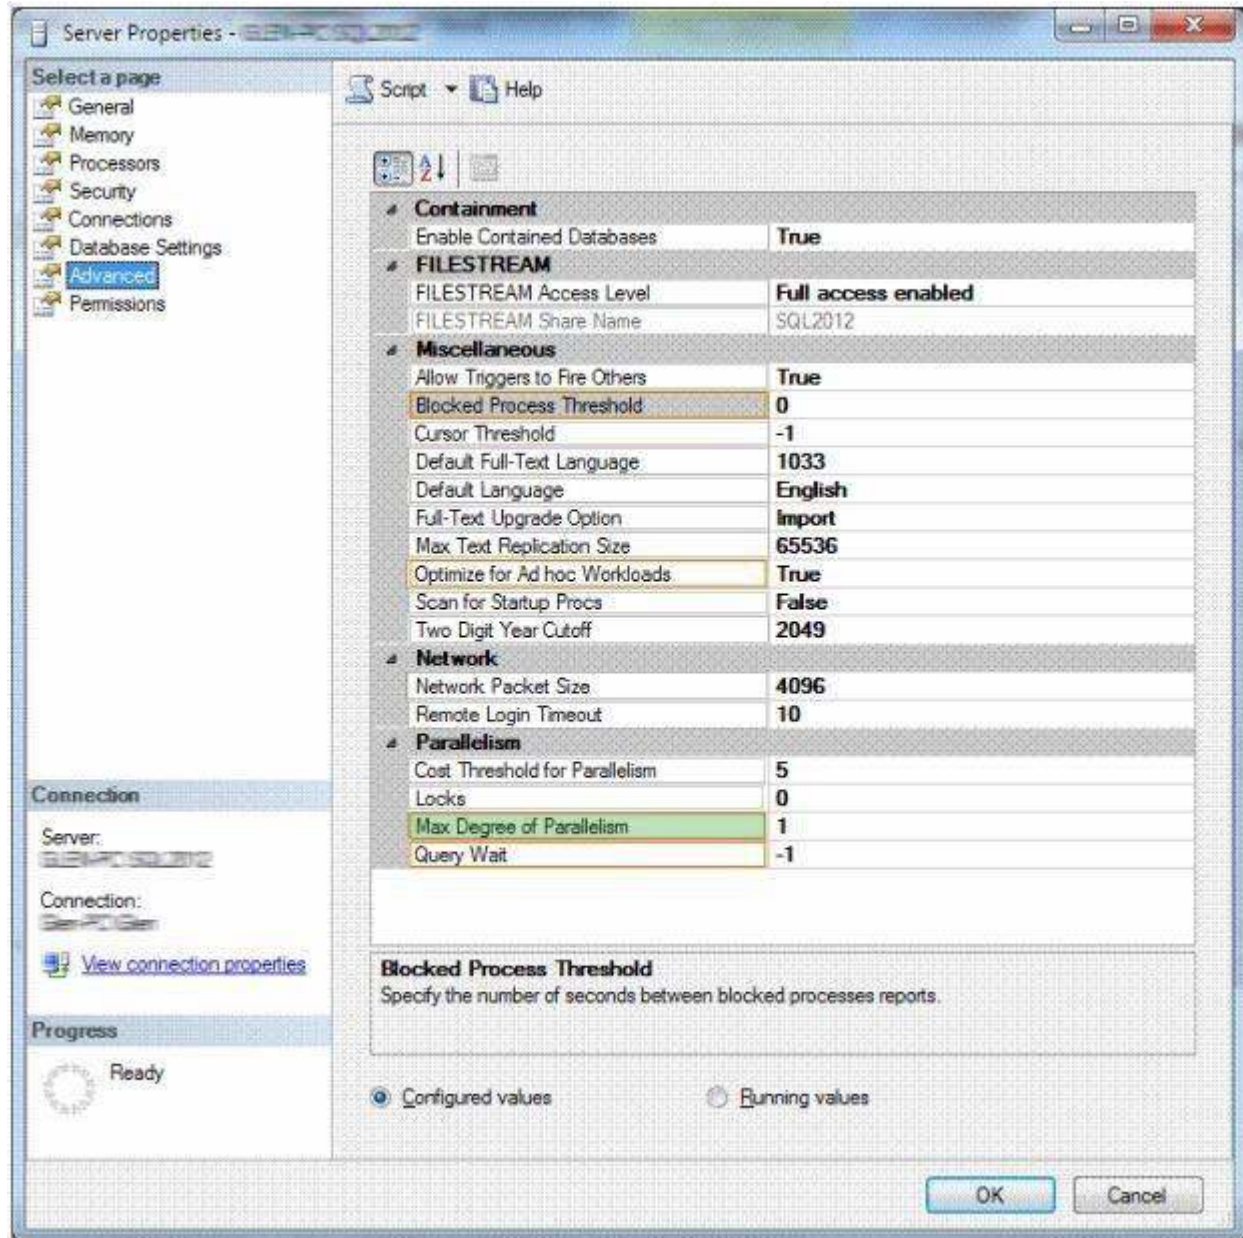
You need to ensure that the instance can process queries as quickly and efficiently as possible.

Which setting should you configure?

Hot Area:



Correct Answer:



Section: (none)

Explanation

Explanation/Reference:

Explanation:

